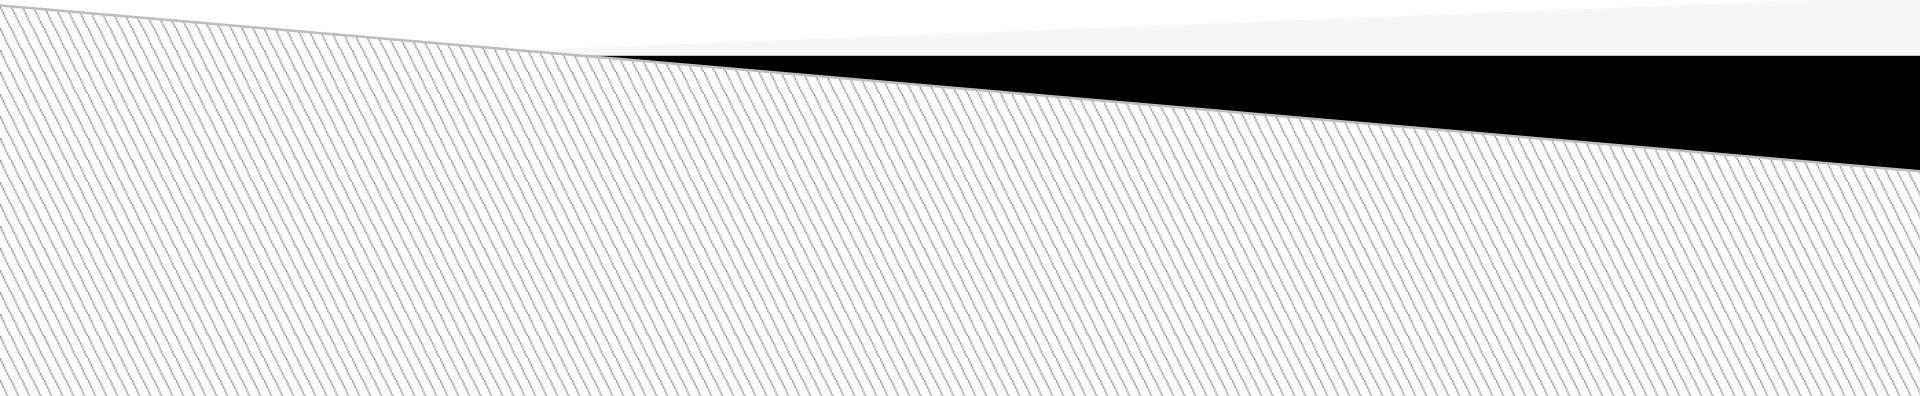


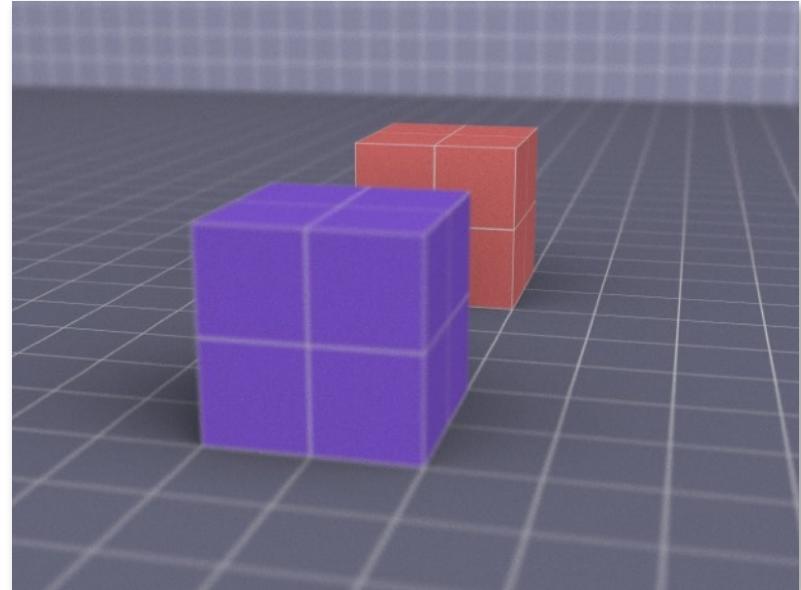
# Global Illumination II



# Our story so far

« Distributed Ray Tracing » Cook et al. (1984)

- ▶ Soft shadows
  - Several rays for each extended light source
- ▶ Anti-aliasing
  - Several rays for each pixel
- ▶ Glossy Reflection
  - Several rays are reflected
- ▶ Motion blur
  - Several rays during time
- ▶ Depth of field
  - Several rays per pixel,  
focusing with a lens



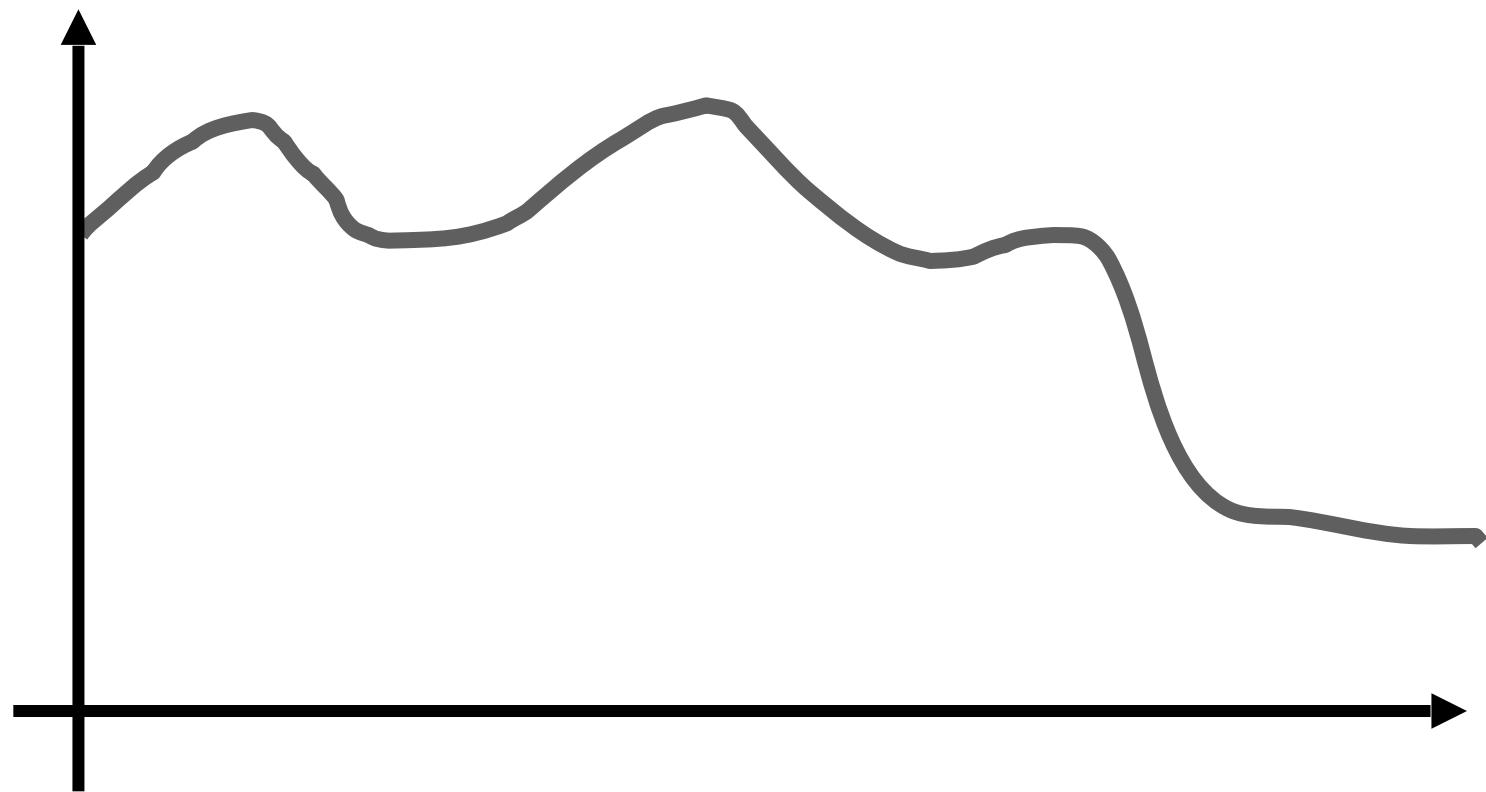
# Ray tracing = integrating!

- ▶ Integrating what?
  - light sources: soft shadows
  - pixels: anti-aliasing
  - BRDF: glossy reflections
  - over Time: motion blur
  - over the lens: depth of field
  - over the **hemisphere**: indirect lighting
  - over **light paths**: global illumination
- ▶ Generic method for computing multi-dimensional integrals:

Monte Carlo Integration

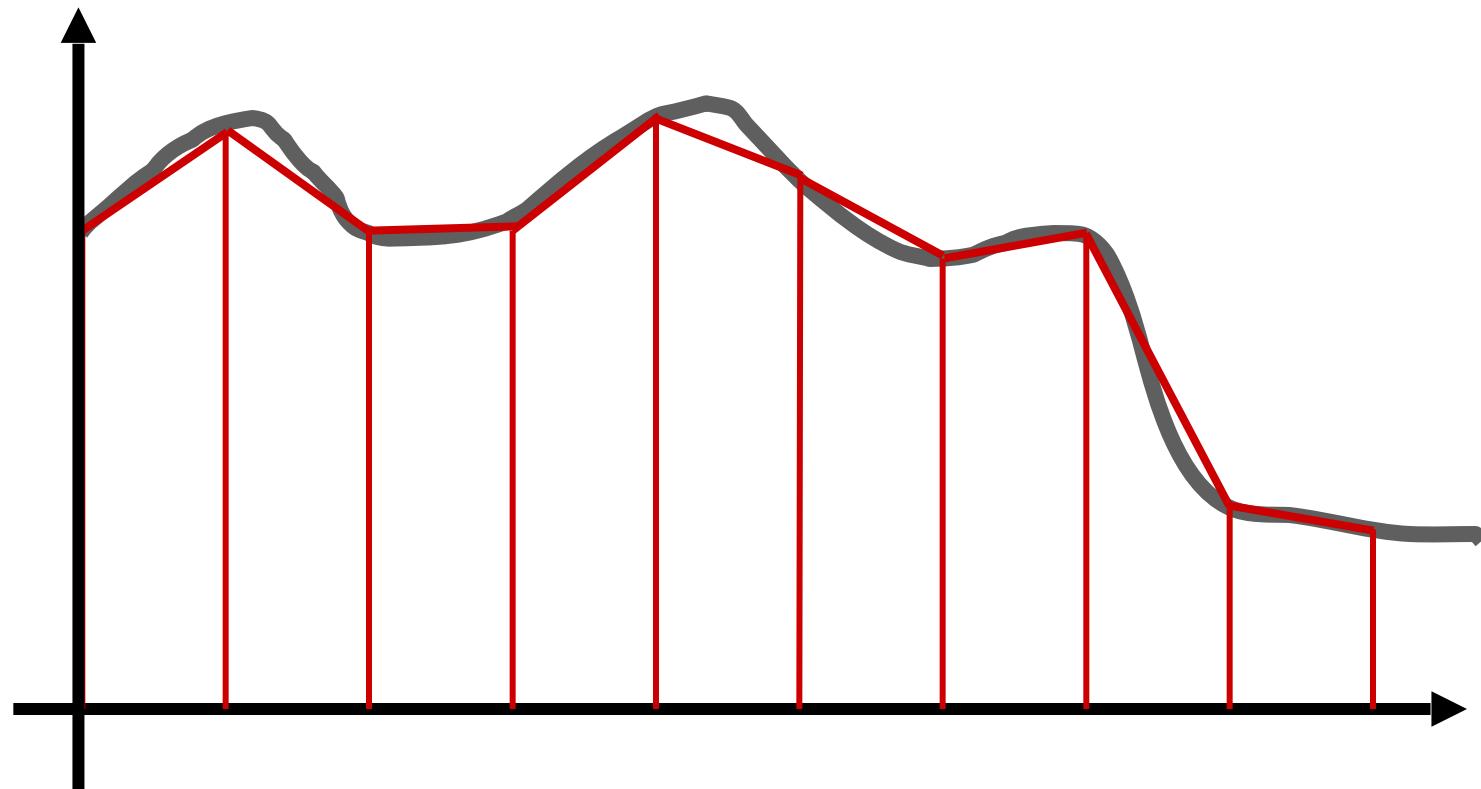
# 1D Integrals

- ▶ Integral of an arbitrary function
- ▶ Continuous problem  $\Rightarrow$  discretization



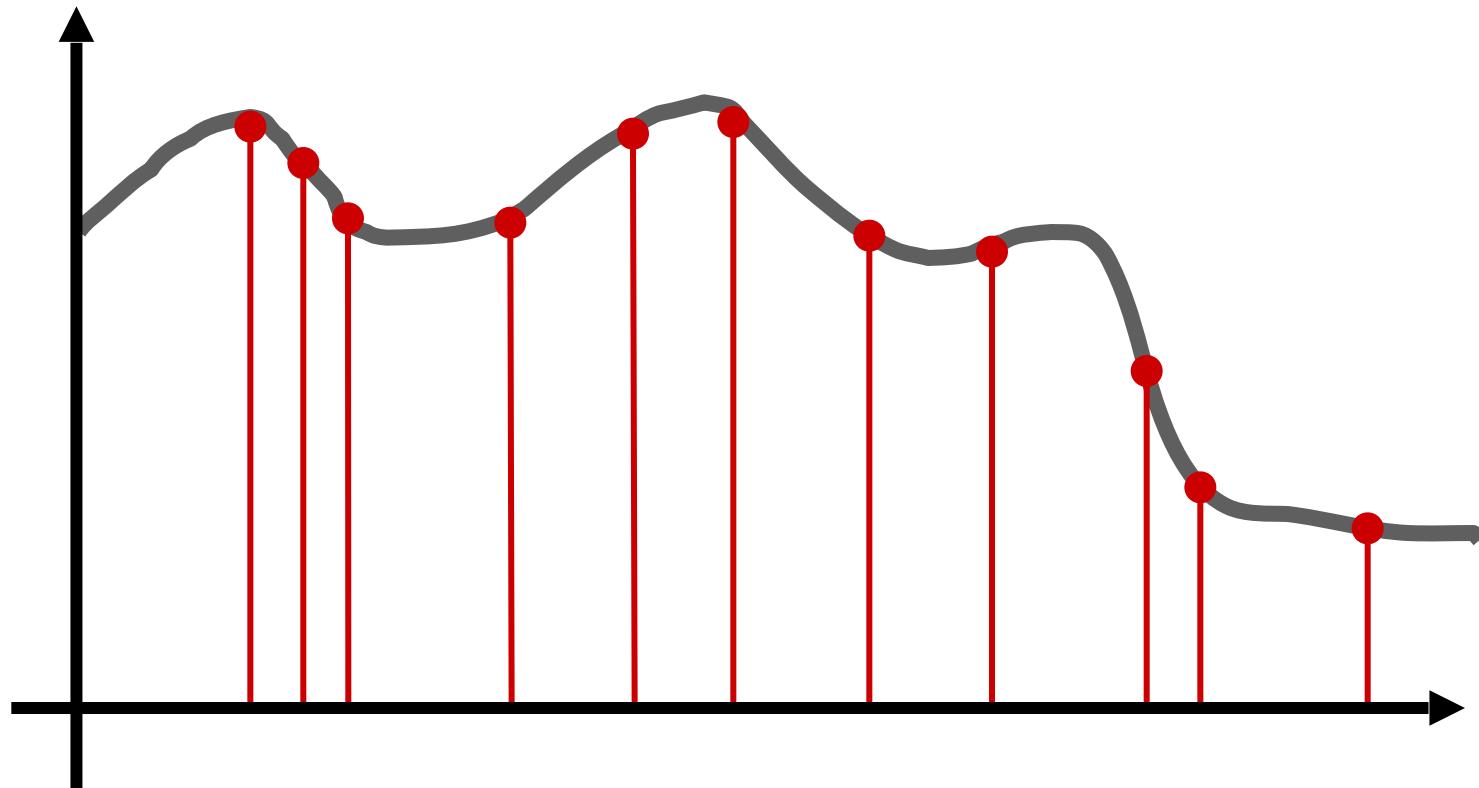
# 1D Integrals

- ▶ Trapezoidal approximation:
  - Also Simpson's rule, midpoint rule...



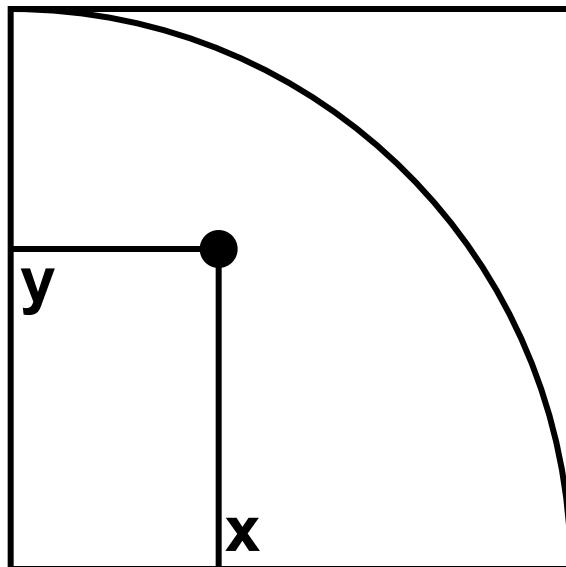
# 1D Integrals

- ▶ Monte Carlo: random sampling
  - Don't keep the distance between the  $n$  samples
  - But on average, expect it to be  $1/n$



# Monte Carlo: computing $\pi$

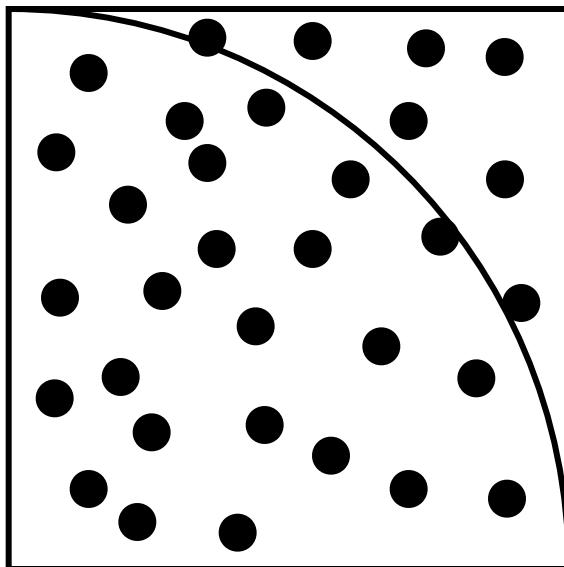
- ▶ Take a square
- ▶ Take a random point  $(x,y)$  in the square
- ▶ Test whether it is inside the  $\frac{1}{4}$  disc  $(x^2+y^2 < 1)$
- ▶ Probability is  $\pi / 4$



Integral of the function  
equal to 1 on the disc, 0  
outside

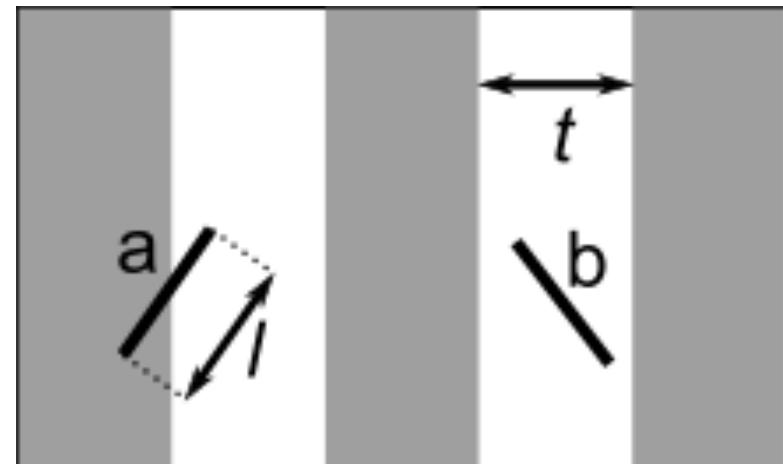
# Monte Carlo: computing $\pi$

- ▶ Probability is  $\pi / 4$
- ▶  $n = \# \text{ points inside} / \# \text{ total points}$
- ▶  $\pi \approx n * 4$
- ▶ Error depends on the number of samples



# See also: Buffon's needle

- ▶ Floor made of parallel strips of wood
- ▶ Throw needles on the floor
  - The needle either crosses the line or it doesn't
  - Count number of time it crosses the line
  - Divide by total number of throws
- ▶ Result is connected to  $\pi$
- ▶  $P = 2l/t\pi$



# Why not use Trapezoidal rule?

- ▶ To compute  $\pi$ , Monte Carlo is not highly efficient
- ▶ But convergence rate independent from dimension
- ⇒ Monte Carlo integration very efficient for higher dimensions

# Continuous random variables

- ▶ Random variable  $x$
- ▶ Probability distribution:  $p(x)$ 
  - Probability that this variable is between  $x$  and  $x+dx$  is  $p(x) dx$

# Expected value (mean)

$$E[x] = \int_{-\infty}^{\infty} xp(x)dx$$

$$E[f(x)] = \int_{-\infty}^{\infty} f(x)p(x)dx$$

- ▶ Expected value is linear:

$$E[f_1(x) + a f_2(x)] = E[f_1(x)] + a E[f_2(x)]$$

# Monte Carlo integration

- ▶ Take the function  $f(x)$  with  $x$  in  $[a b]$
- ▶ We want to compute: 
$$I = \int_a^b f(x)dx$$
- ▶ Take a random variable  $x$
- ▶ If  $x$  has a uniform distribution,  $I=E[f(x)]$ 
  - By definition of expected value

# Sum of random variables

- ▶ Take N random variables, independent, identical distribution (IID)  $x_i$  (N échantillons)
  - Same probability (here uniform)
- ▶ Define:

$$F_N = \frac{1}{N} \sum_{j=1}^n f(x_i) \quad \text{Monte Carlo estimator}$$

- ▶ By linearity of expected value:  
 $E[F_N] = E[f(x)]$

# Variance

$$\sigma^2 = E[(x - E[x])^2] = \int_{-\infty}^{\infty} (x - E[x])^2 p(x) dx$$

- ▶ Measures the distance to expected value
- ▶ Standard deviation  $\sigma$ : square root of variance
- ▶ Properties:
  - $\sigma^2[x+y] = \sigma^2[x] + \sigma^2[y] + 2 \text{Cov}[x,y]$
  - $\sigma^2[ax] = a^2 \sigma^2[x]$

# About the variance

$$\sigma^2[F_N] = \sigma^2 \left[ \sum_{j=1}^n \frac{f(x_i)}{N} \right]$$

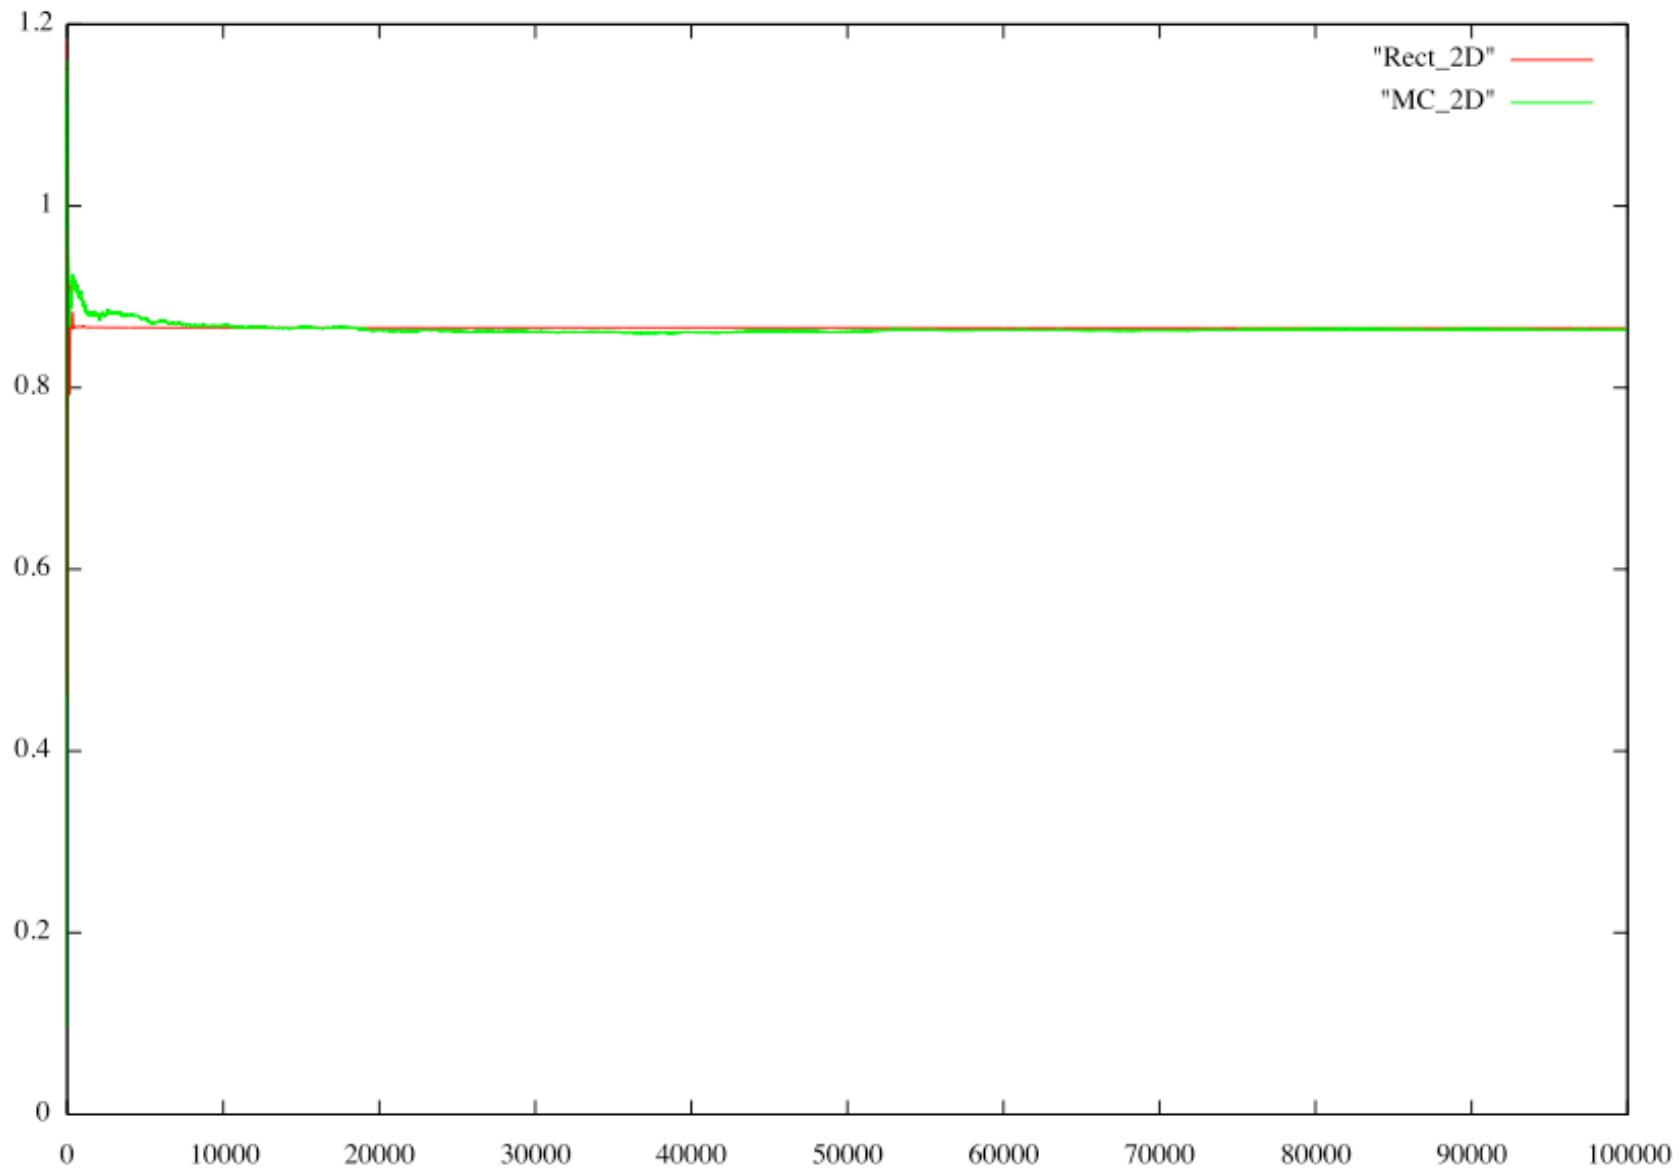
- ▶ Independent variables  $\Rightarrow \text{Cov}[x_i, x_j] = 0$  si  $i \neq j$

$$\sigma^2[F_N] = \frac{\sigma^2[f(x)]}{N}$$

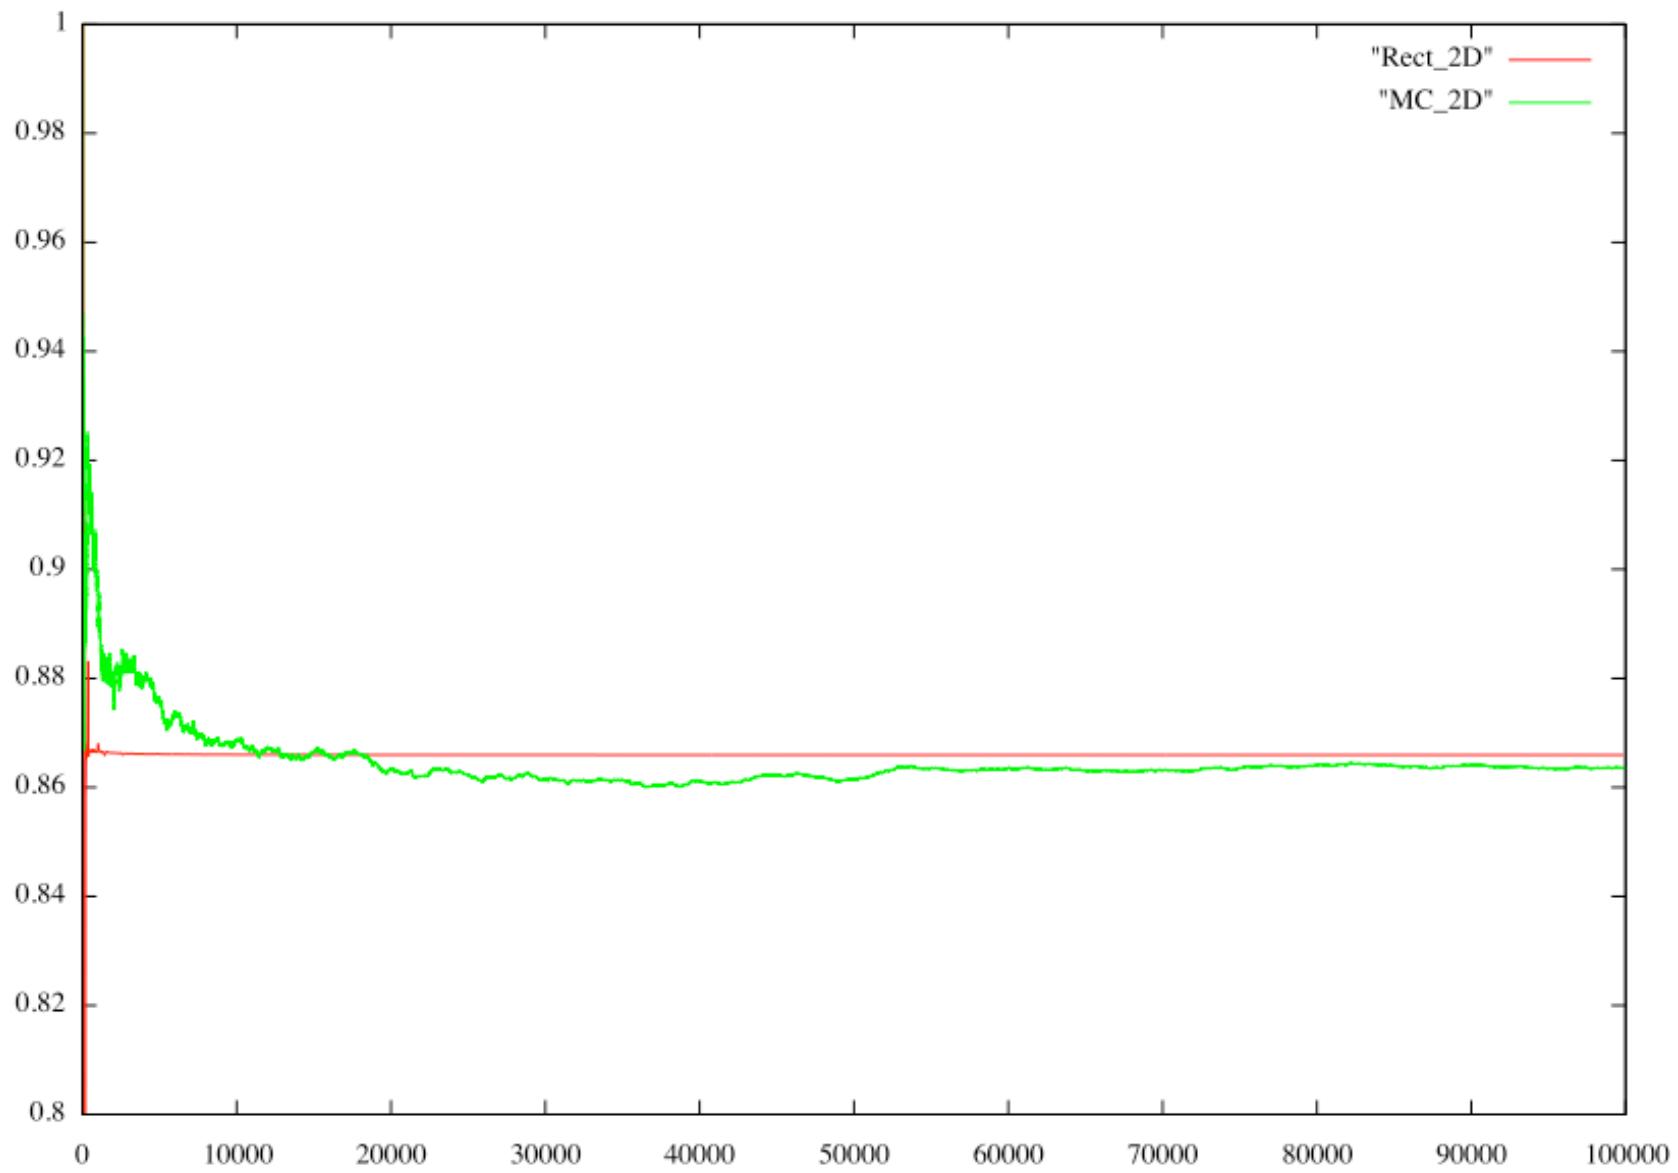
- ▶ thus  $\sigma$  (error) decreases with  
 $\Rightarrow$  slow convergence

$$\sqrt{N}$$

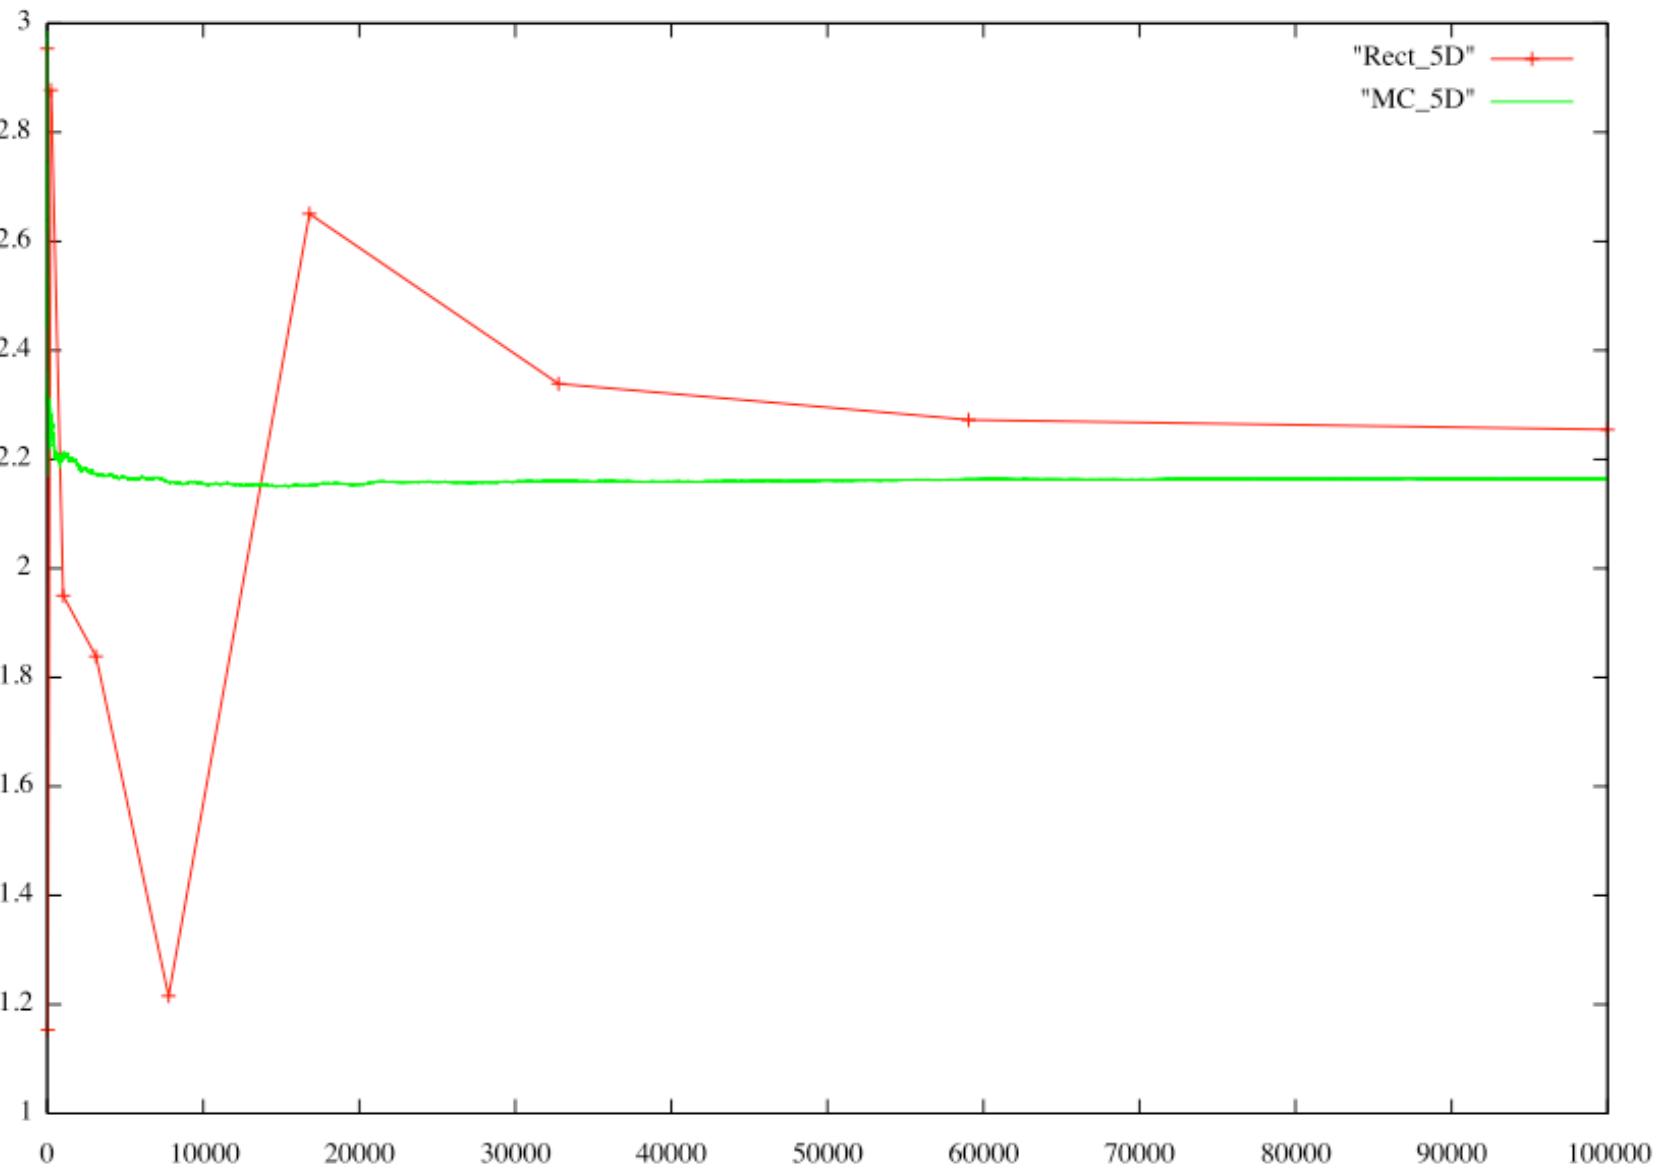
# Ex.: Integral of a 2D function



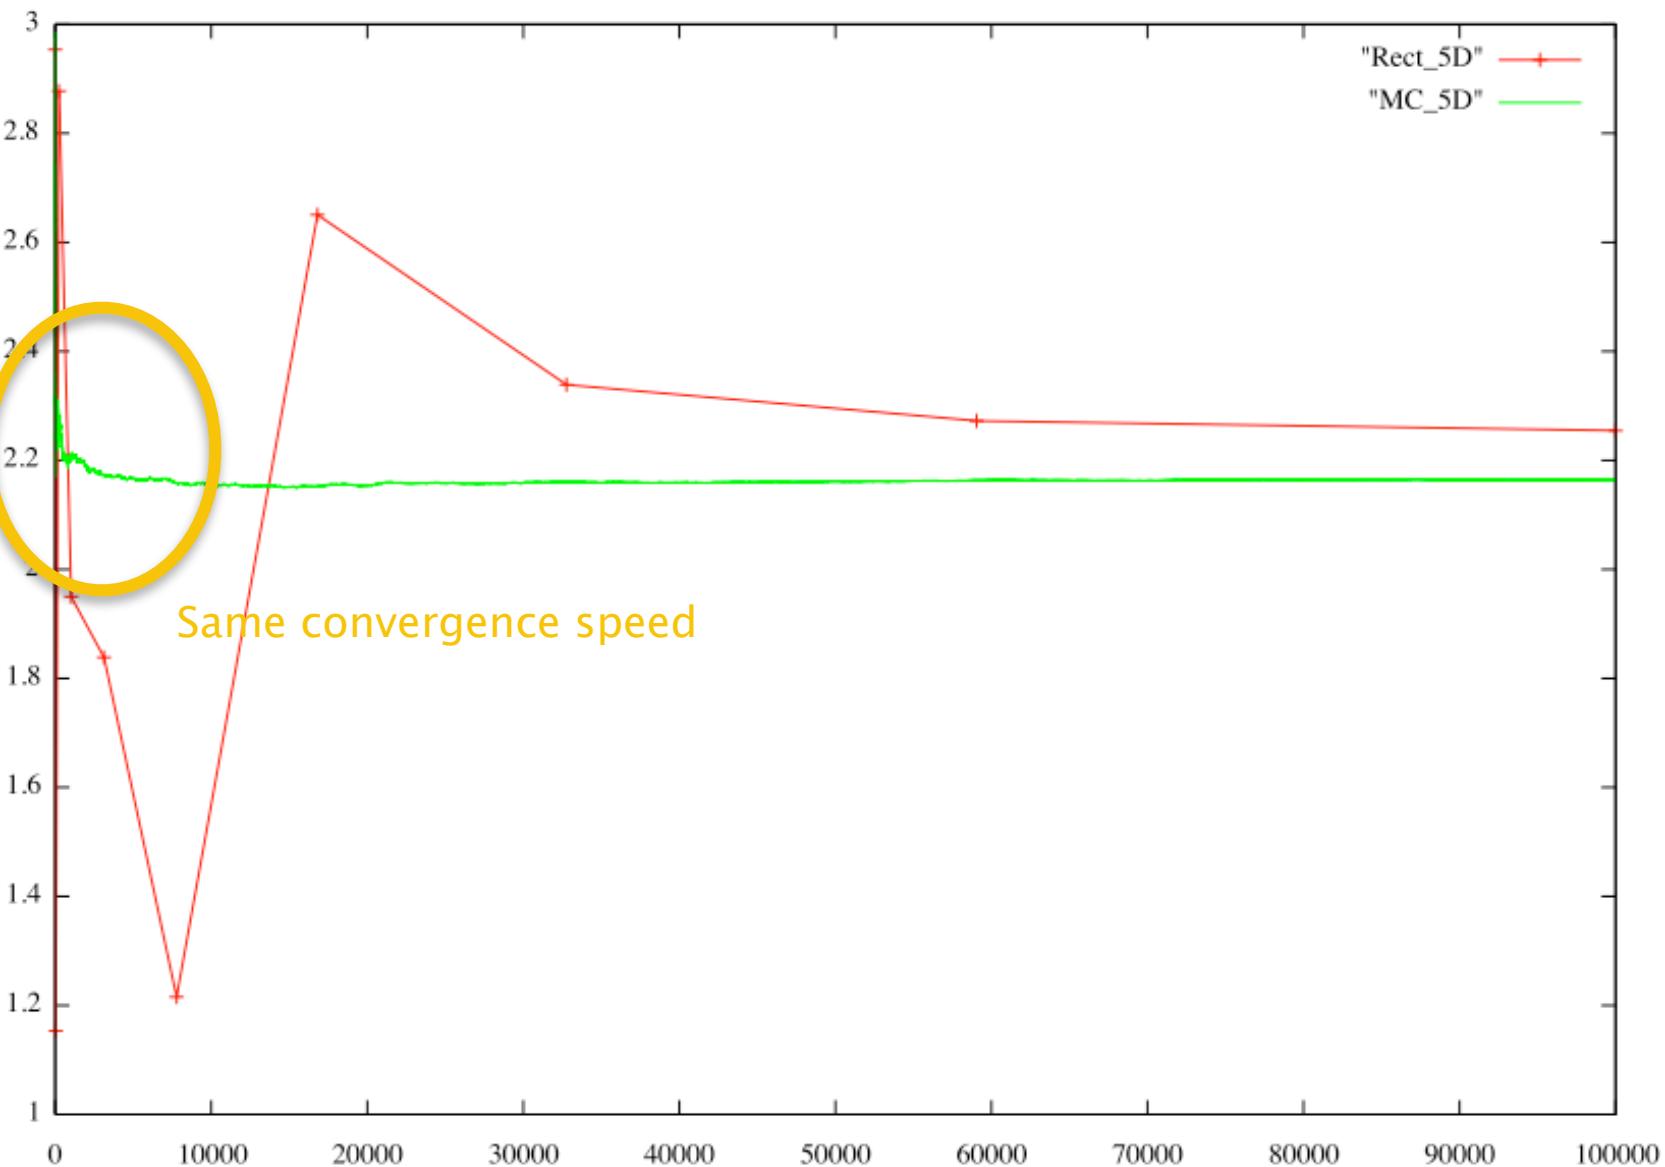
# 2D function: zoom-in



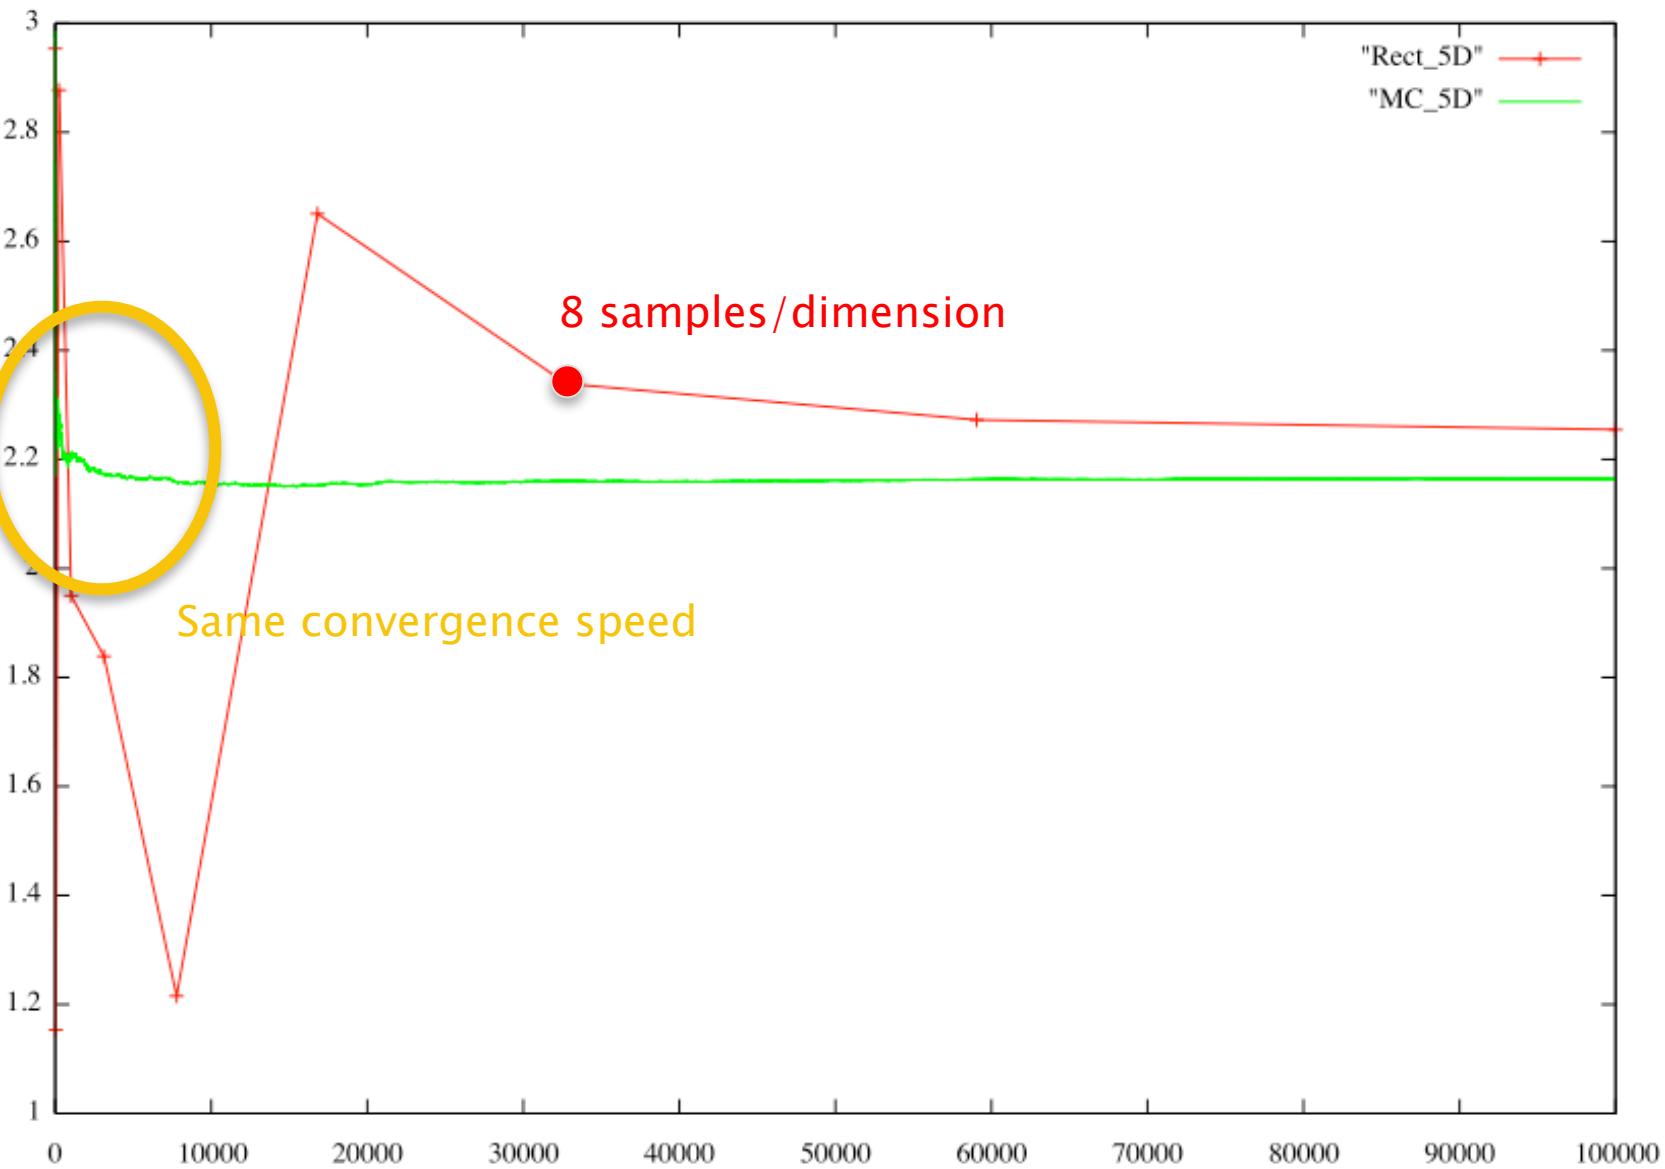
# Ex.: Integral of a 5D fonction



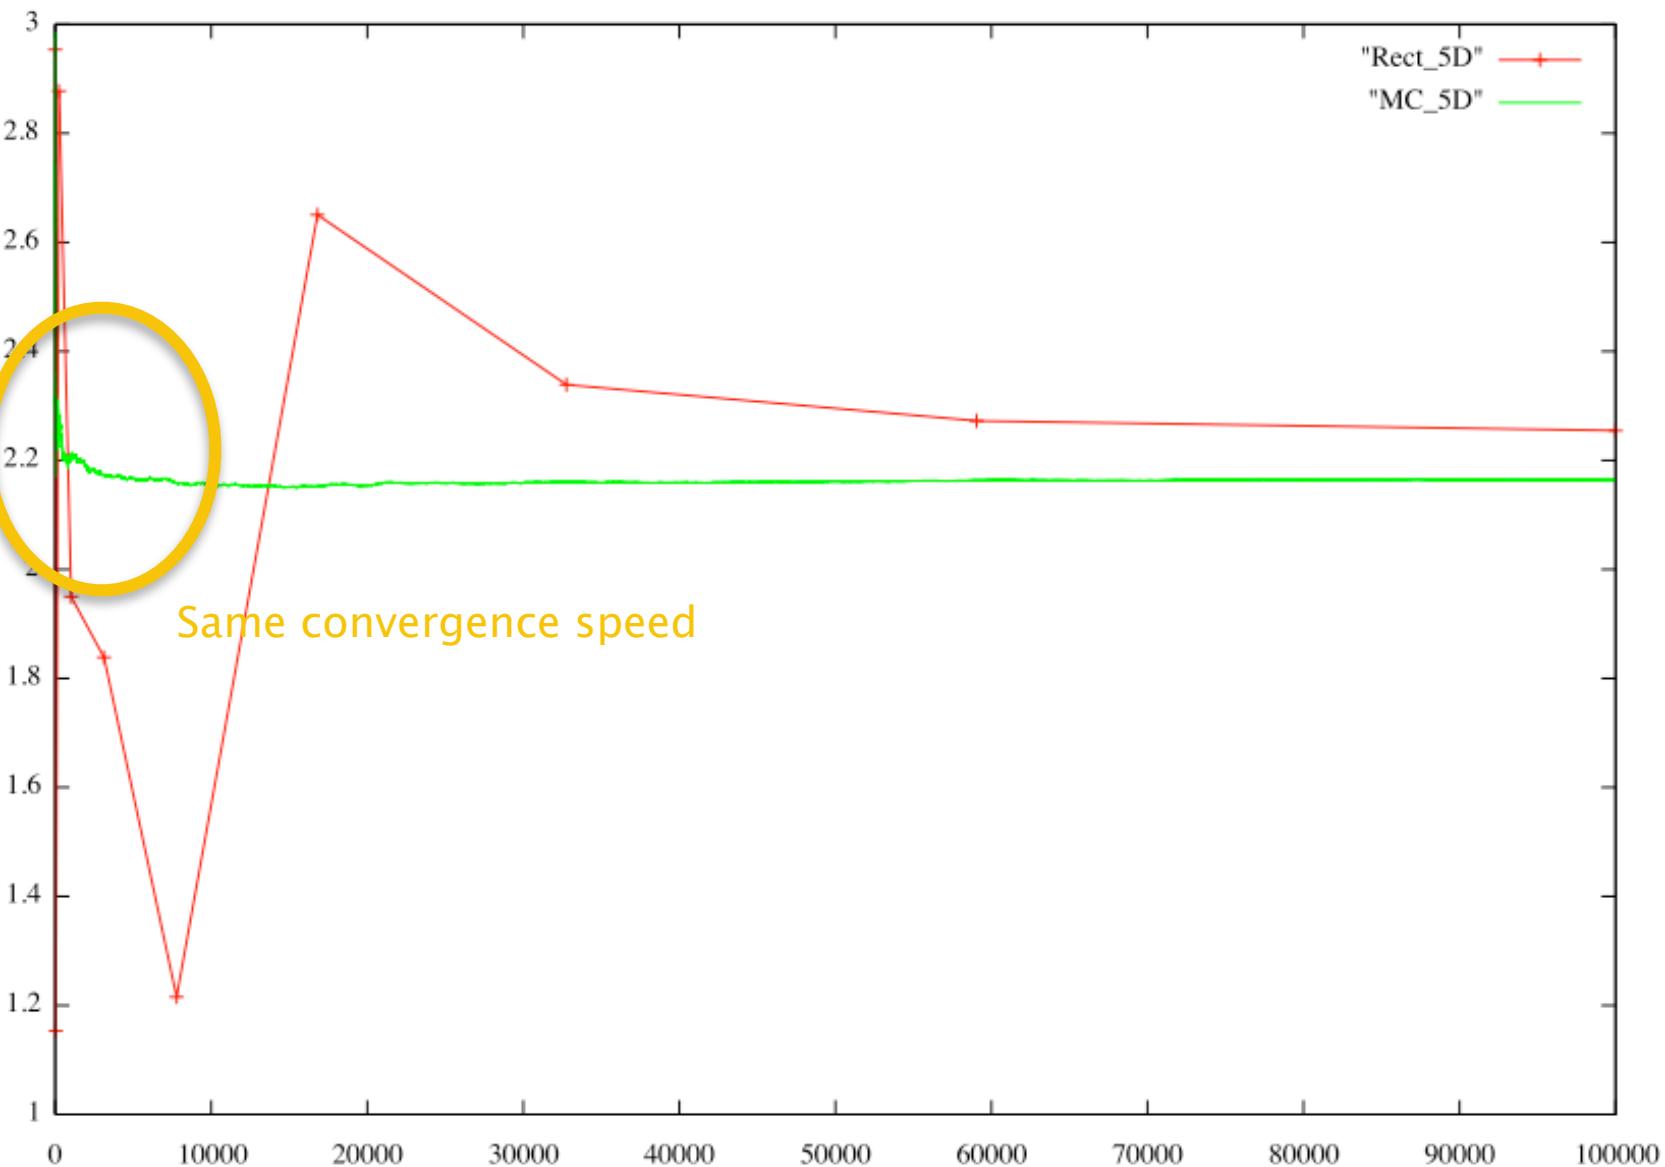
# Ex.: Integral of a 5D fonction



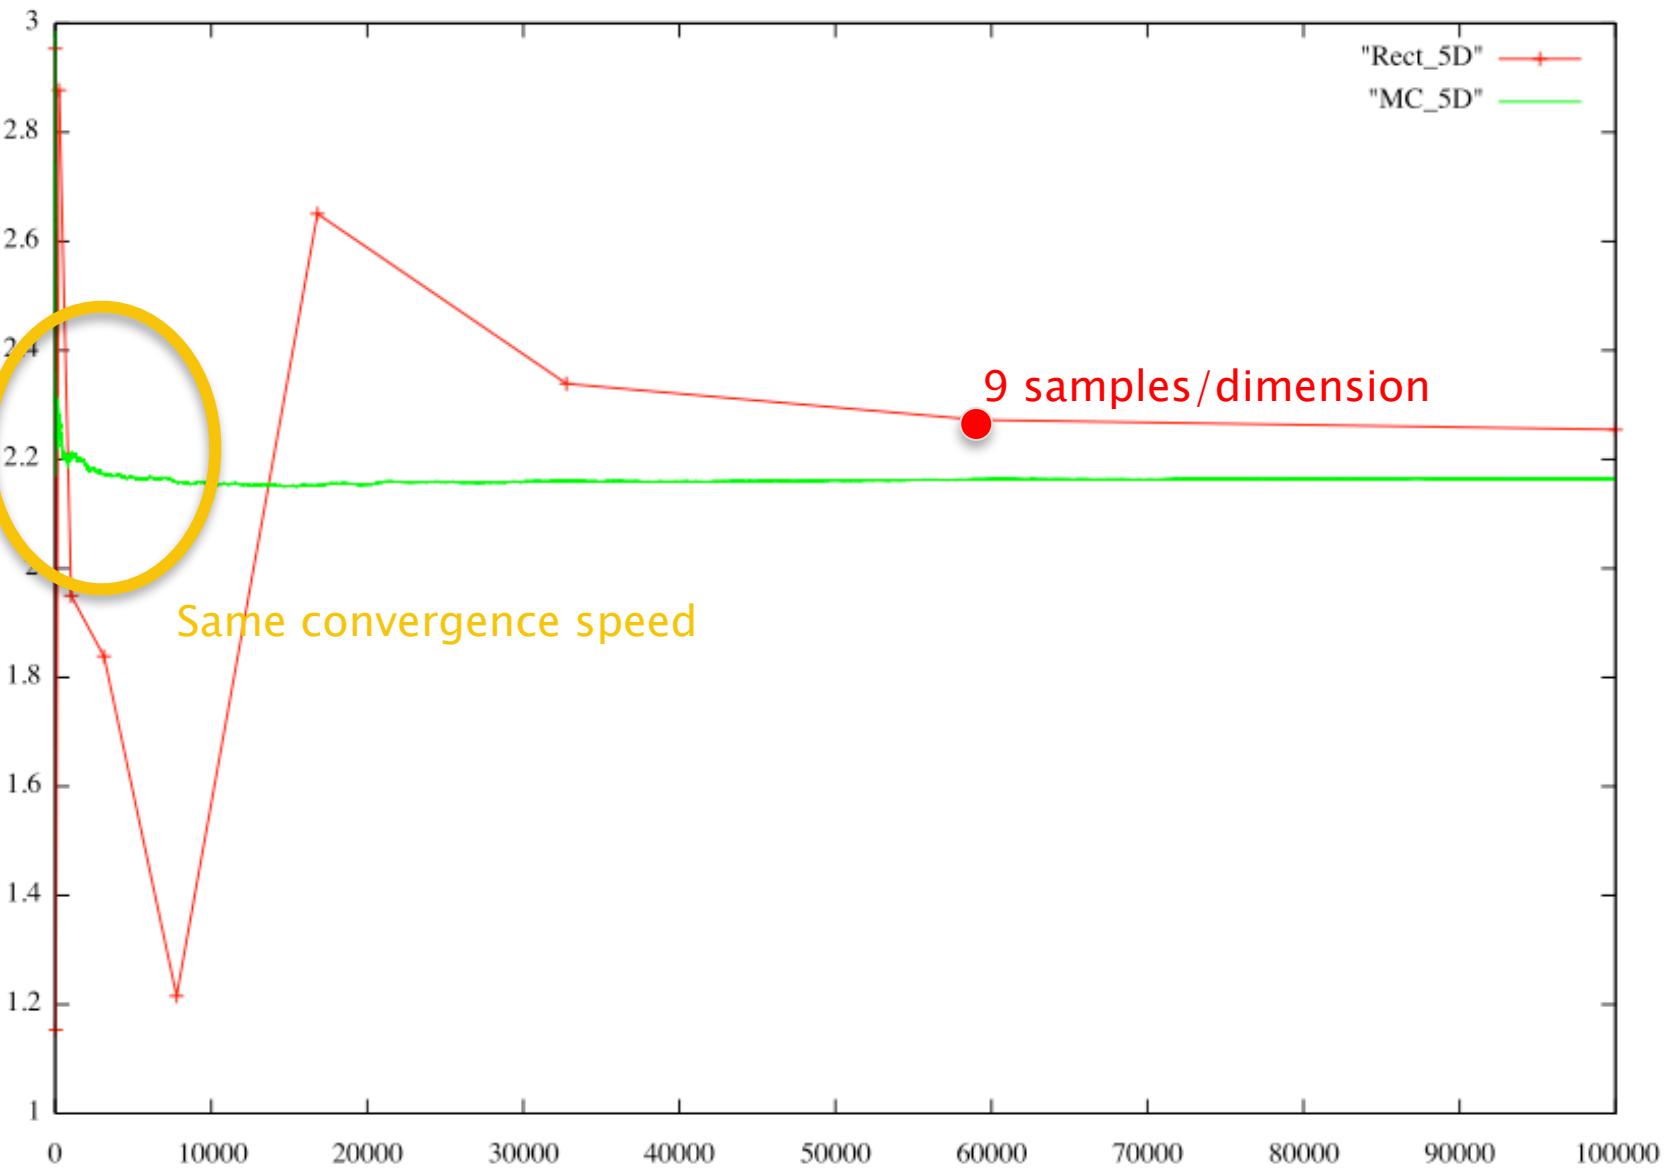
# Ex.: Integral of a 5D fonction



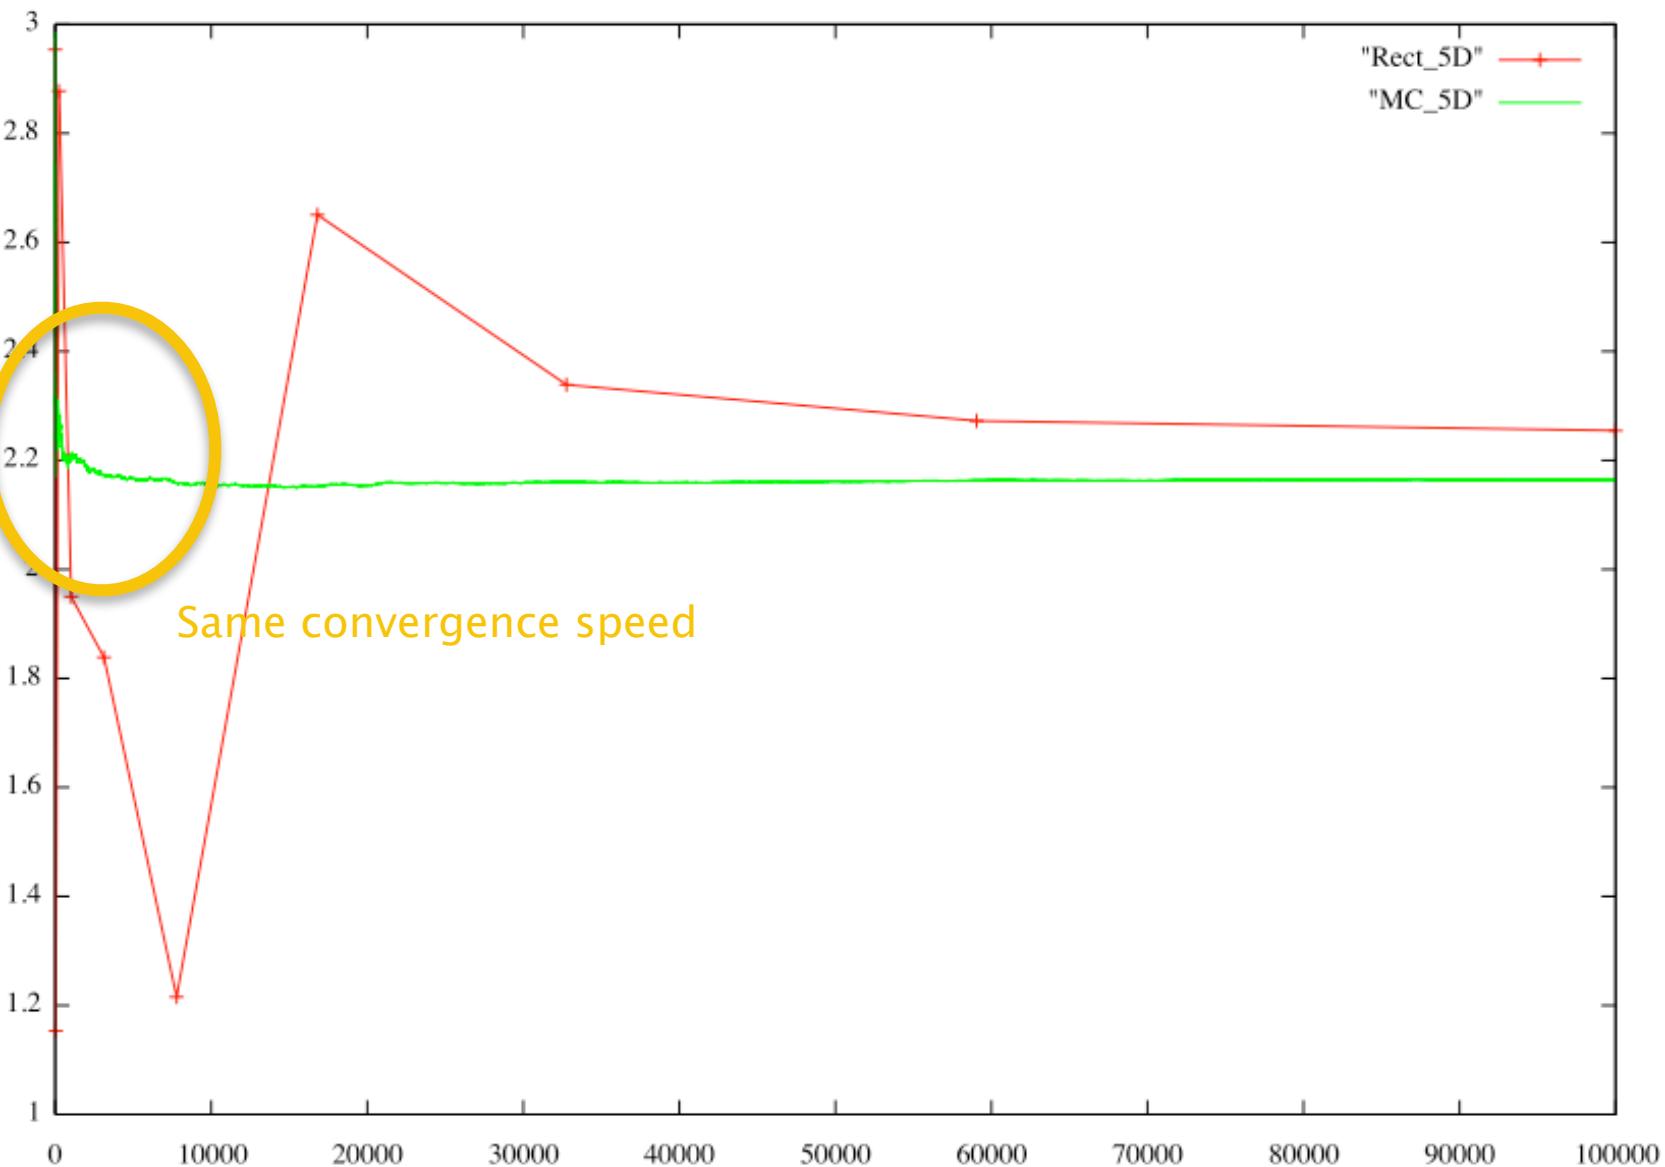
# Ex.: Integral of a 5D fonction



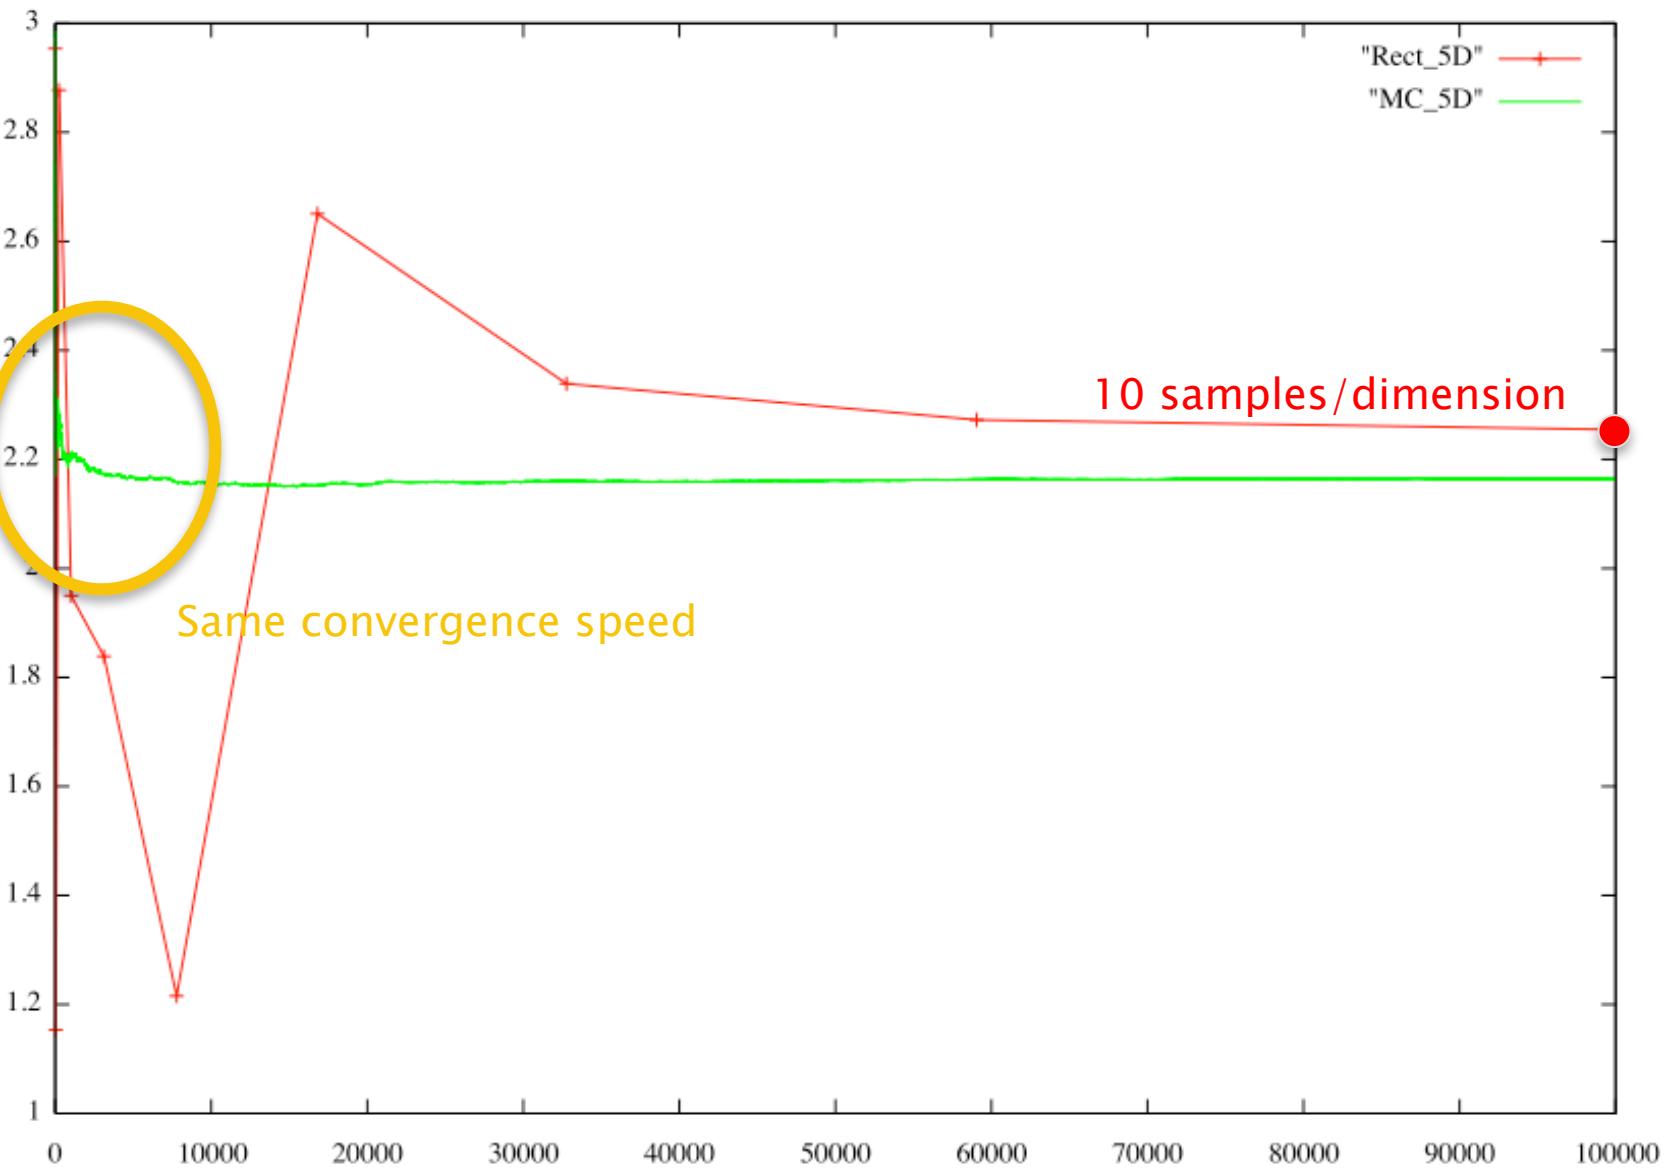
# Ex.: Integral of a 5D fonction



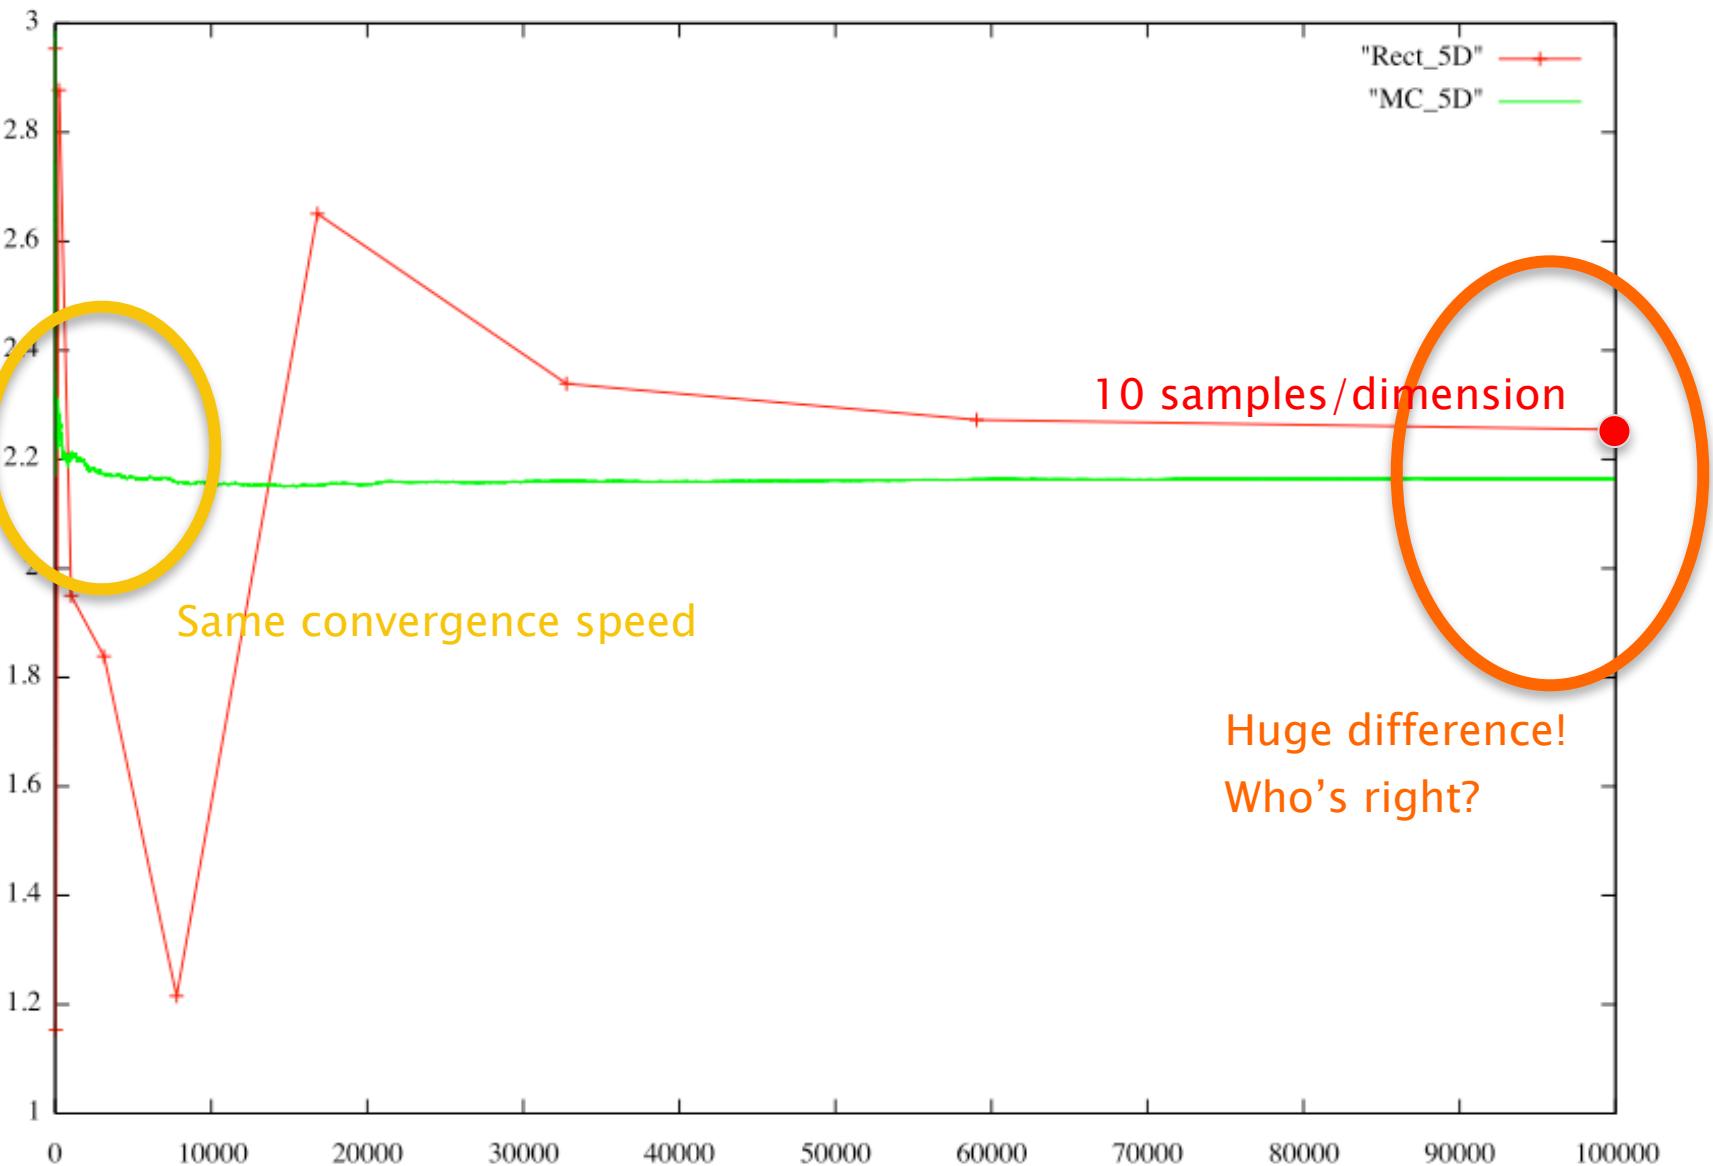
# Ex.: Integral of a 5D fonction

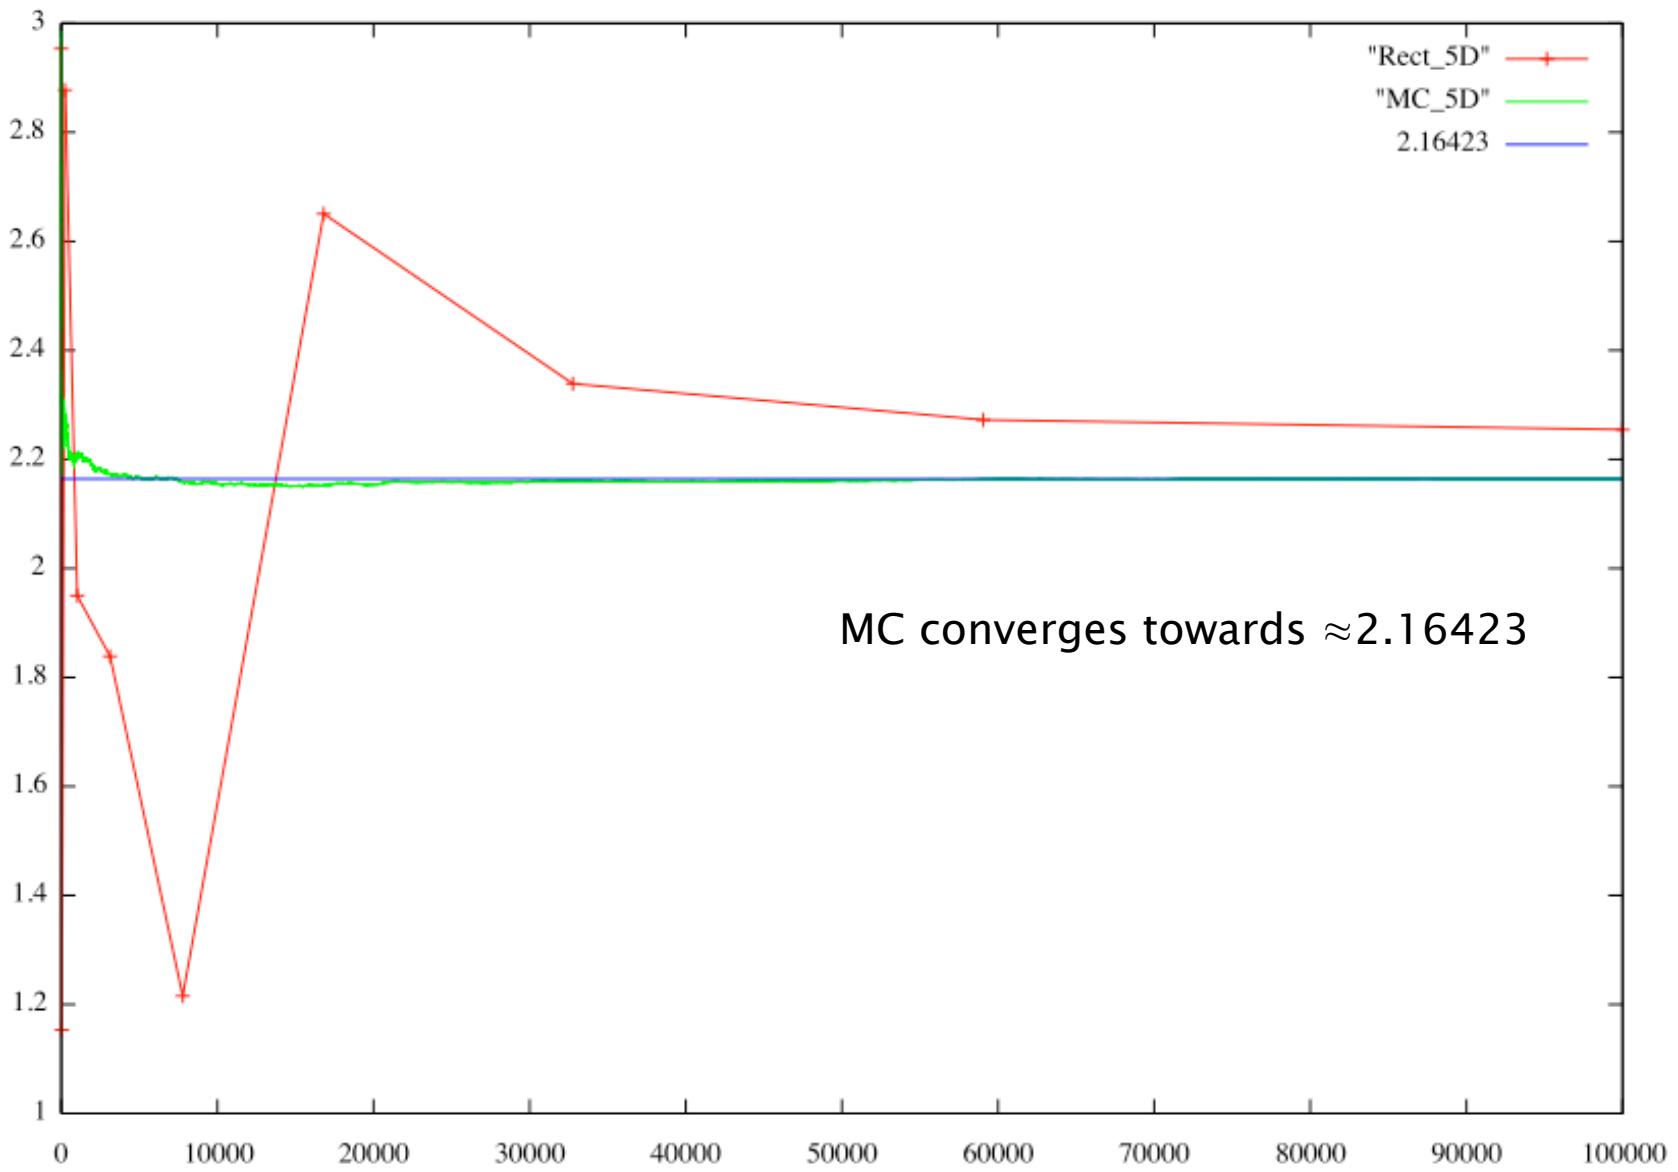


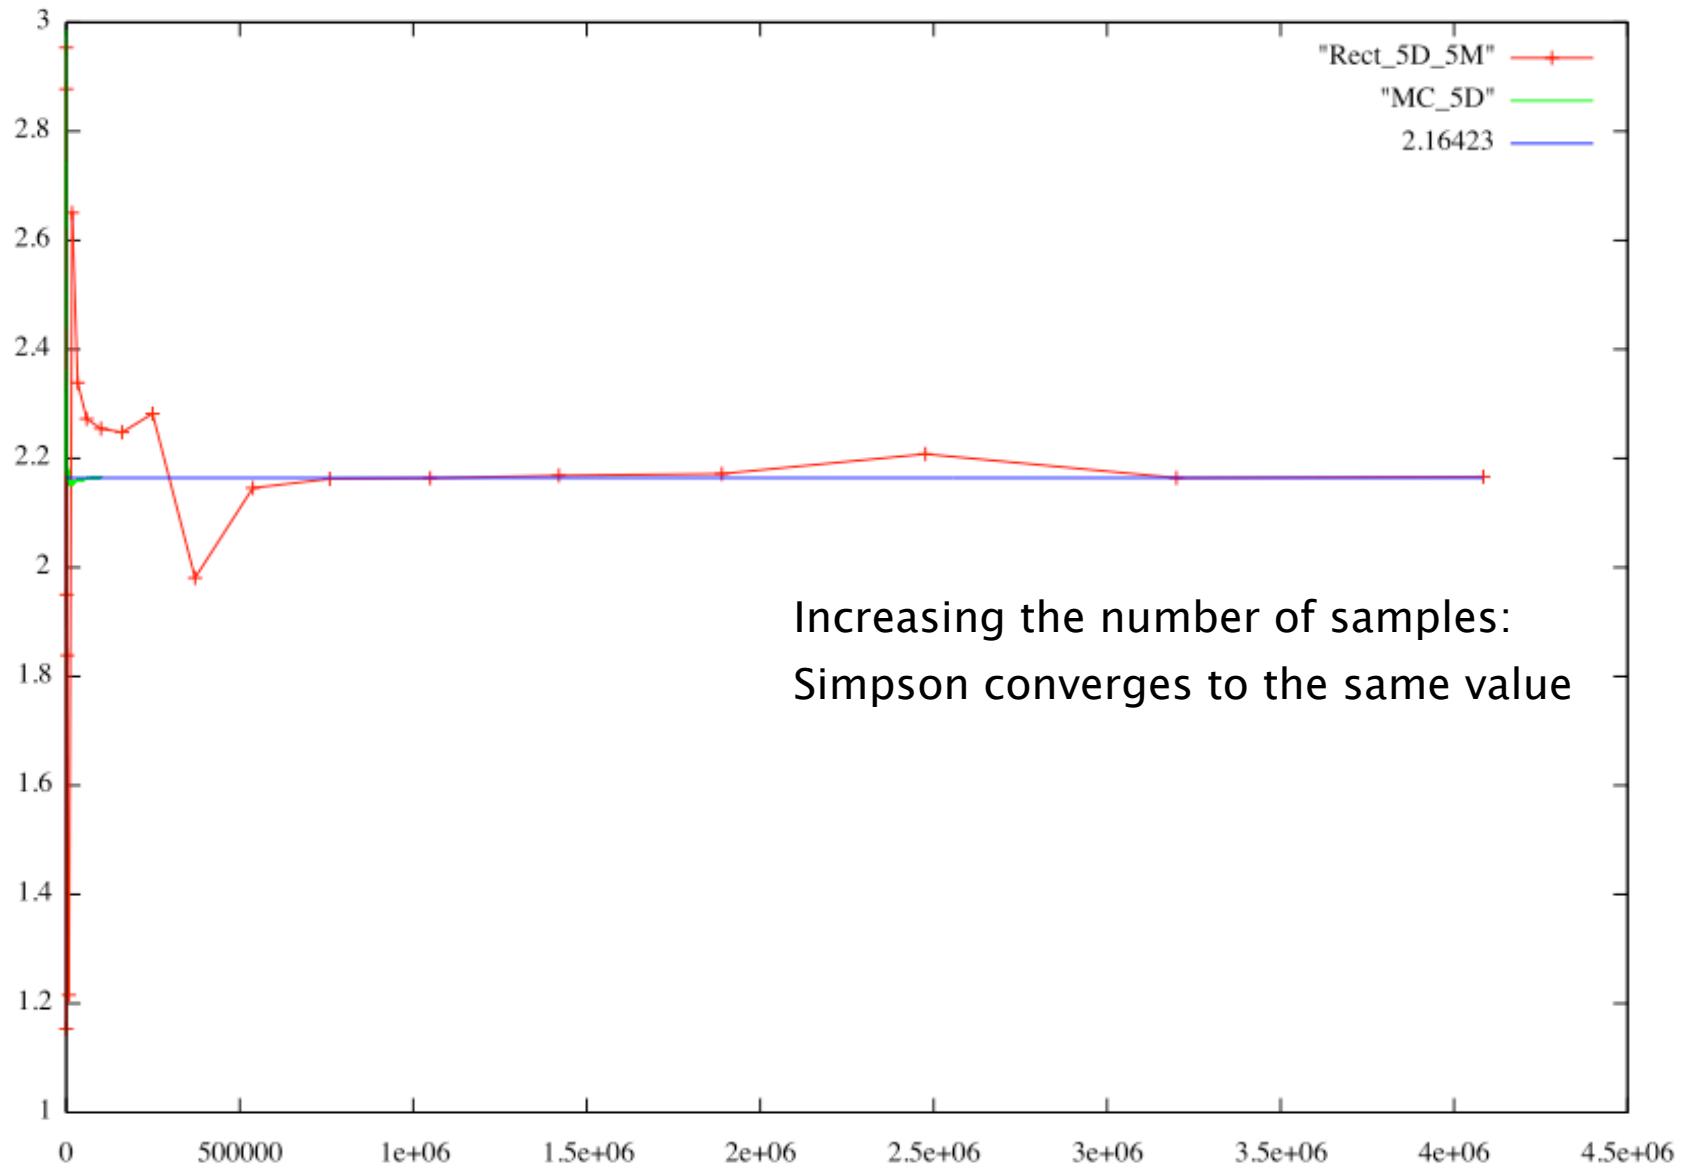
# Ex.: Integral of a 5D fonction

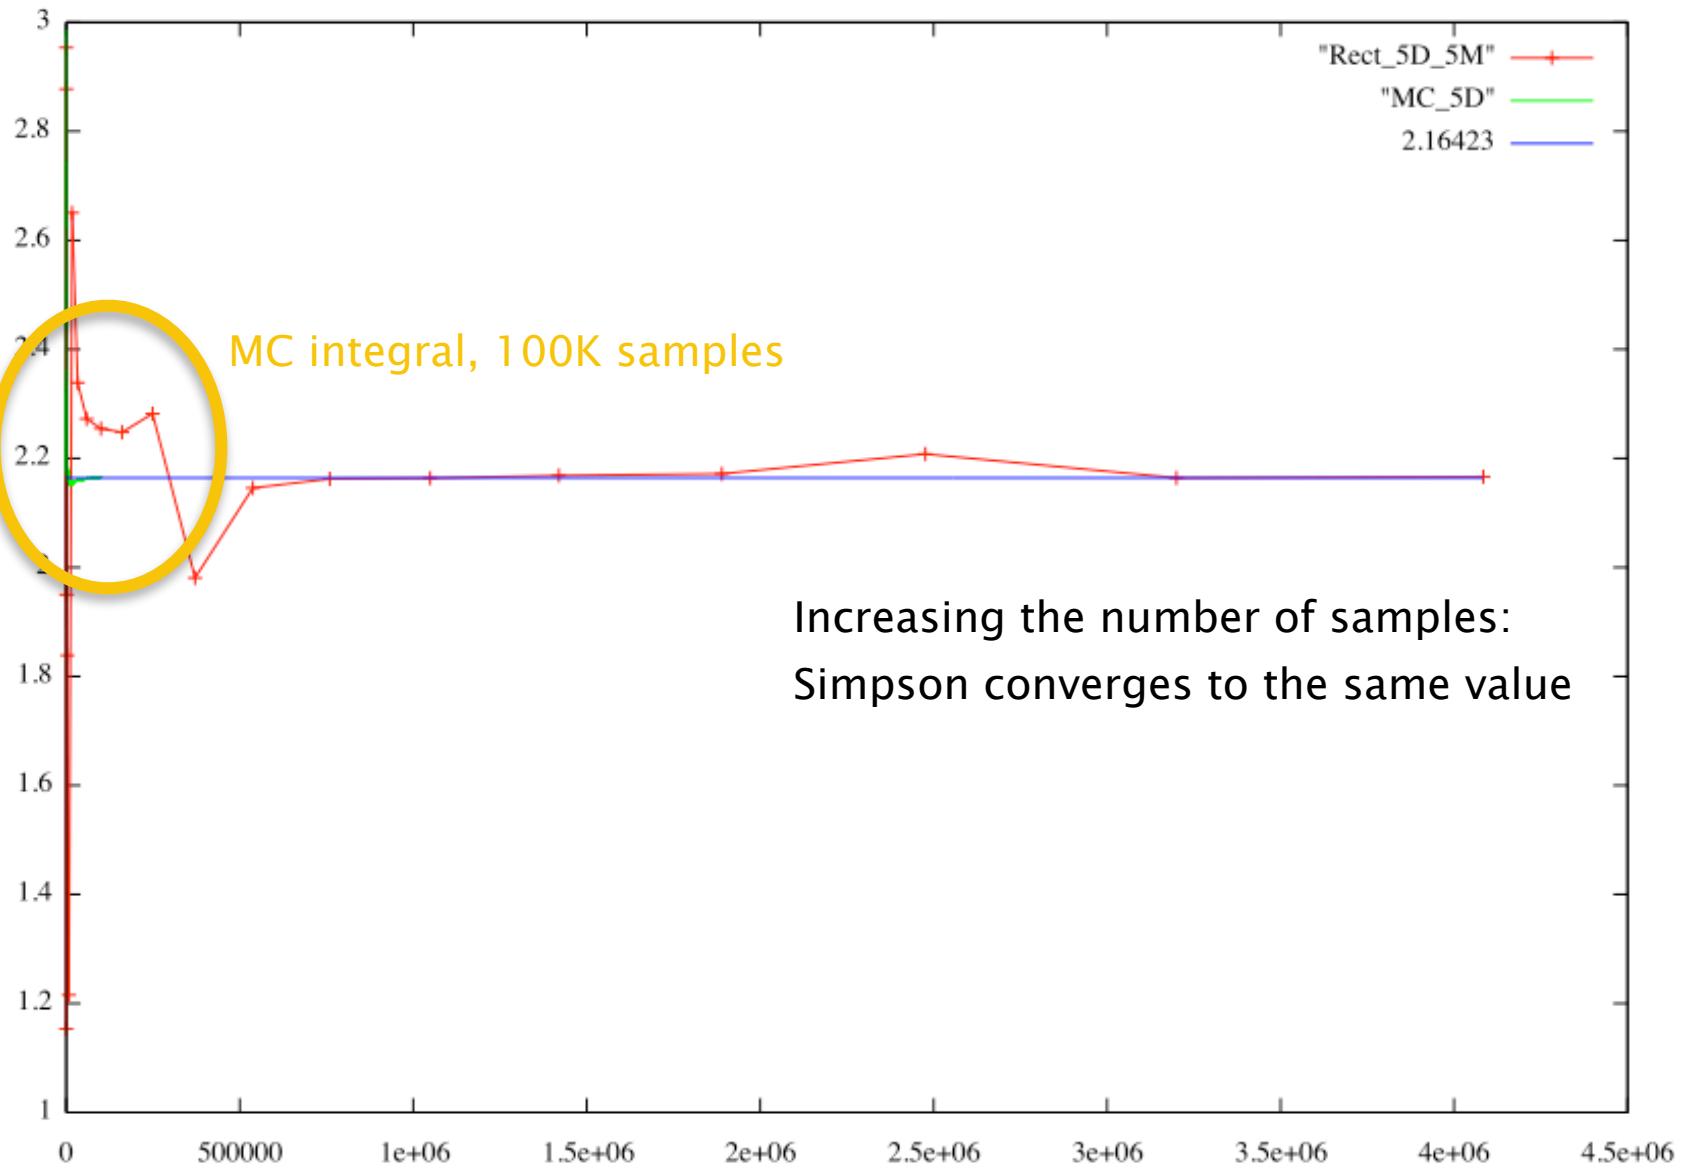


# Ex.: Integral of a 5D fonction









# Monte Carlo integration: pros

- ▶ Few restrictions on the function to integrate
  - No requirements on continuity, regularity ...
  - Only needs sampling on a single point
- ▶ Same convergence rate on higher dimensions
- ▶ Very simple

# Monte Carlo integration: cons

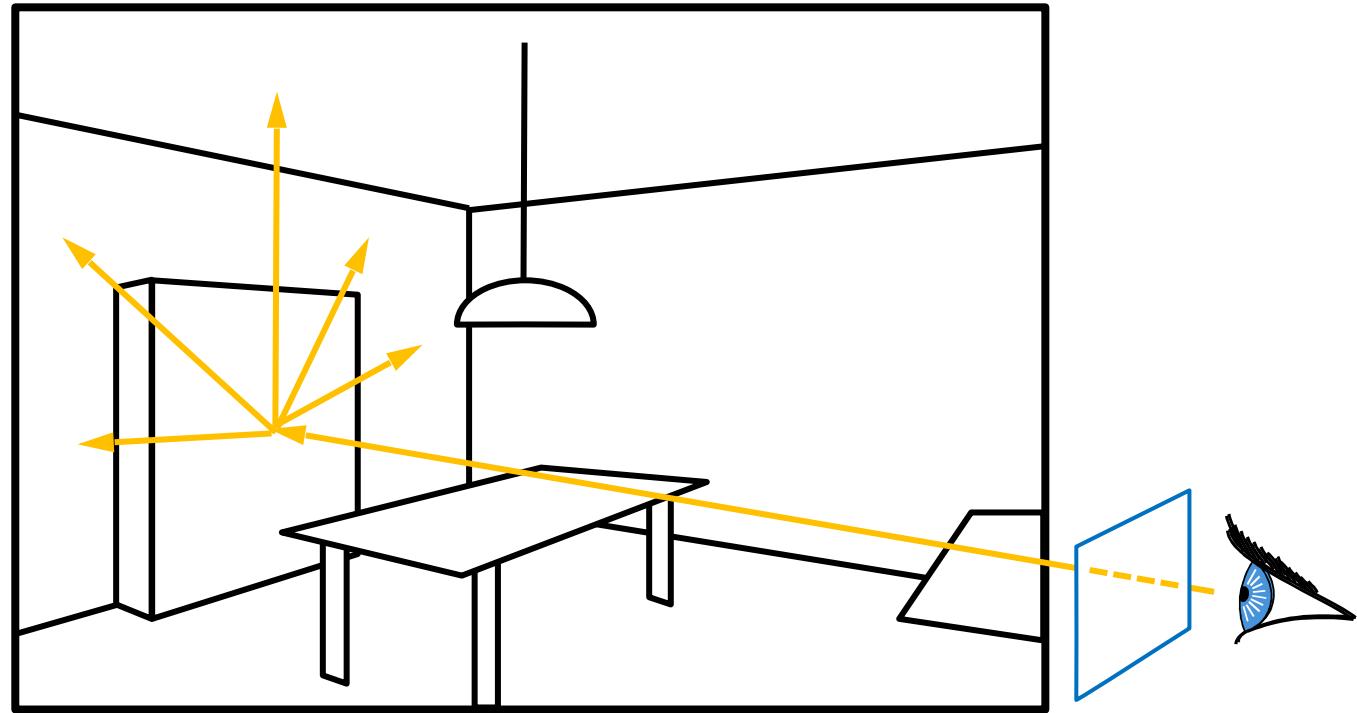
- ▶ Noise
- ▶ Slow convergence
- ▶ Efficient implementation harder

# Global Illumination Simulation

- ▶ Monte-Carlo methods for illumination
- ▶ How?
- ▶ Each pixel is an integral

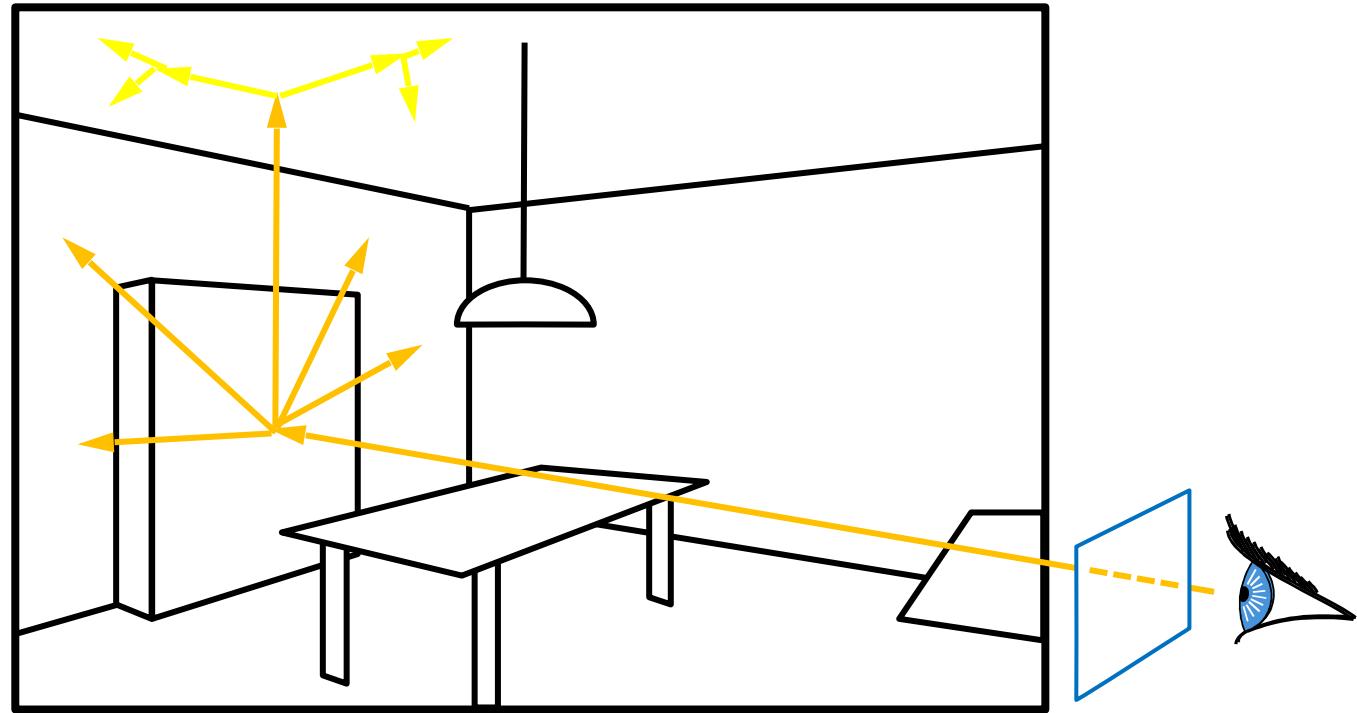
# Monte Carlo methods

- ▶ One ray for every pixel
- ▶ For each visible point: random sampling of rays, accumulate radiance



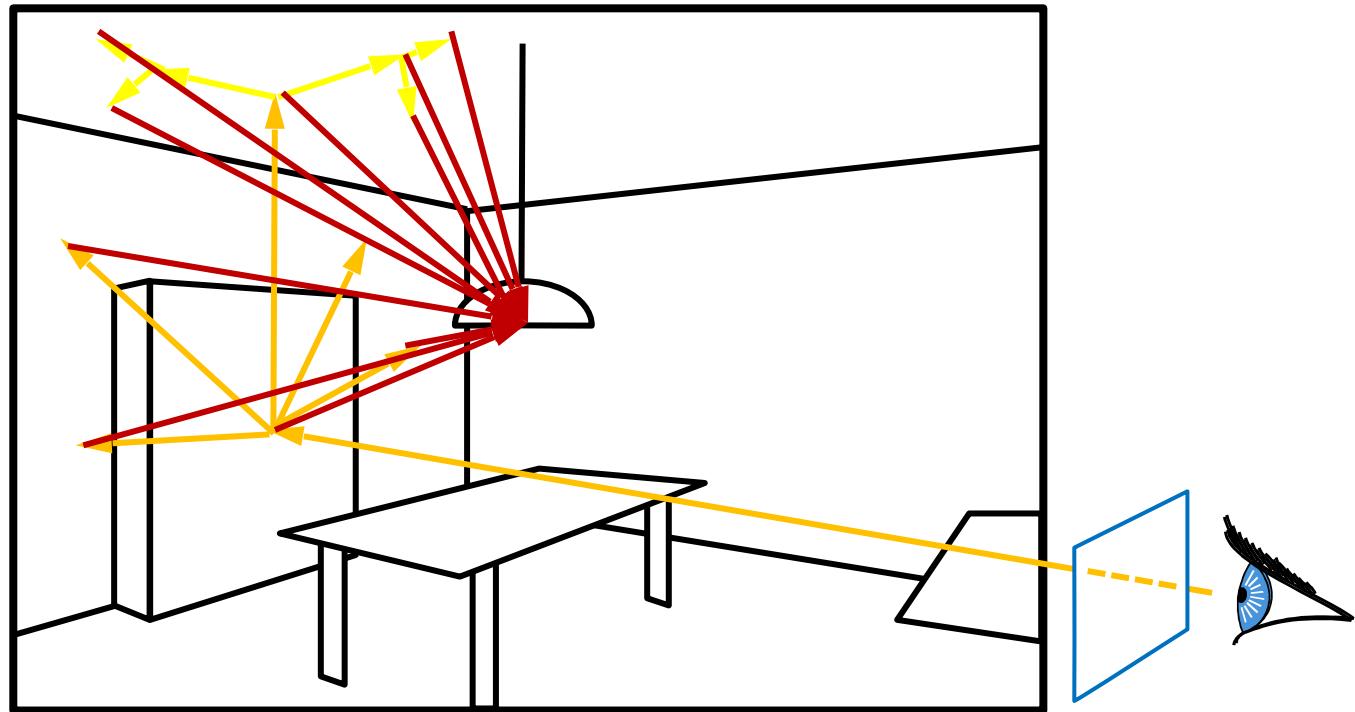
# Monte Carlo methods

- ▶ One ray for every pixel
- ▶ For each visible point : random sampling of rays, accumulate radiance
- ▶ Keep going, recursively



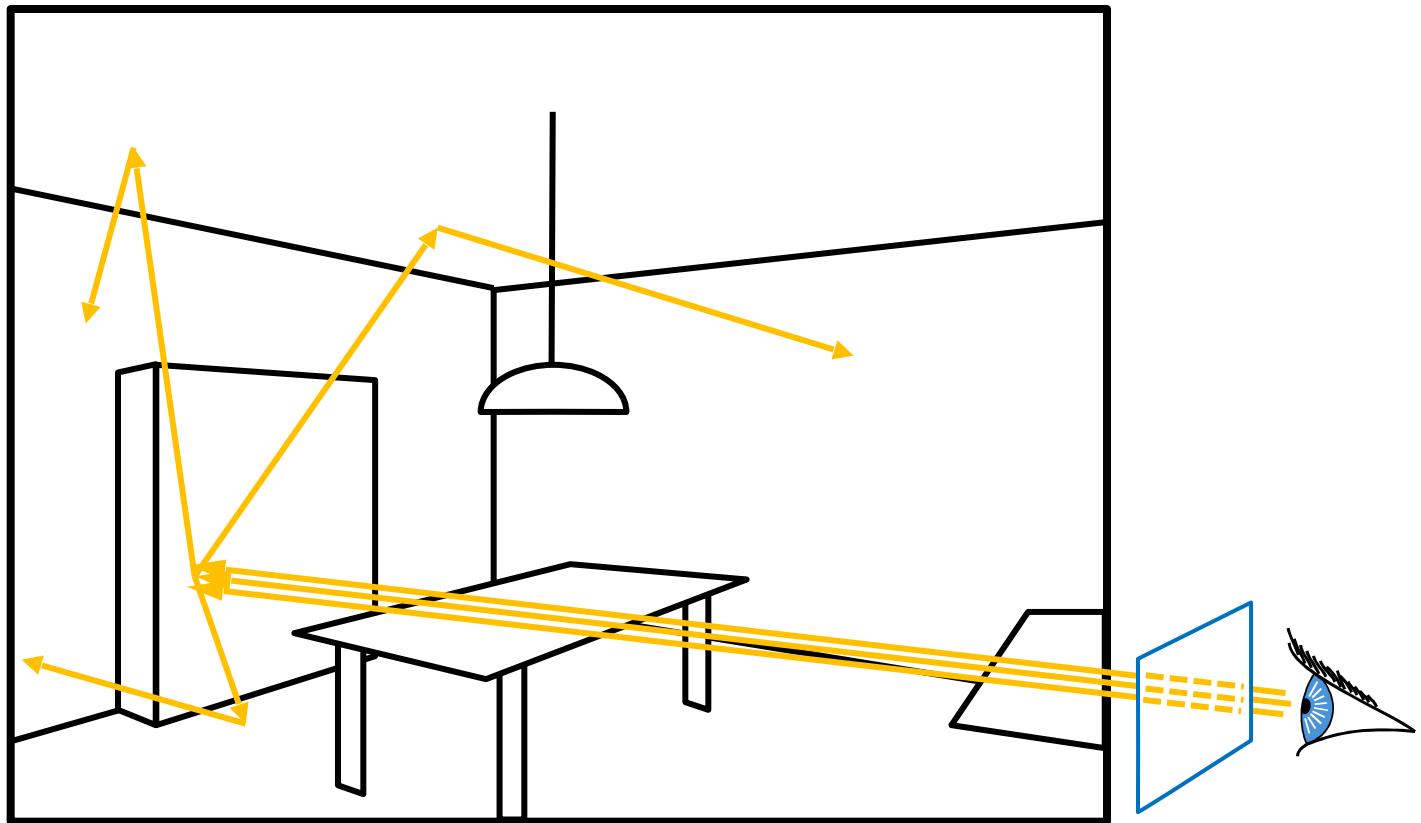
# Monte Carlo methods

- ▶ One ray for every pixel
- ▶ For each visible point : random sampling of rays, accumulate radiance
- ▶ Keep going, recursively
- ▶ Sample the light source each time



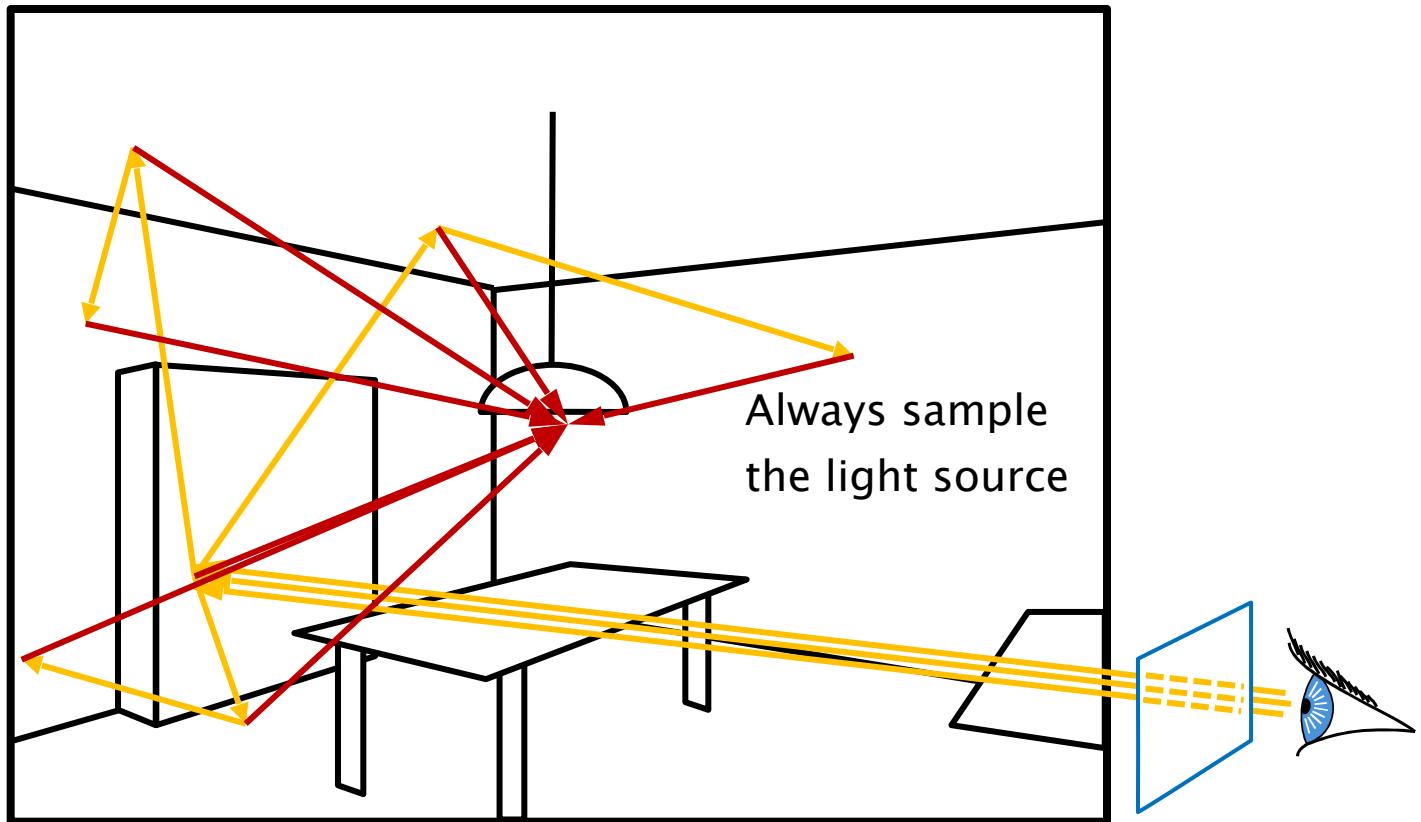
# Monte Carlo Path Tracing

- ▶ Trace only one ray at each recursion
- ▶ But trace several (hundreds of) primary rays for each pixel

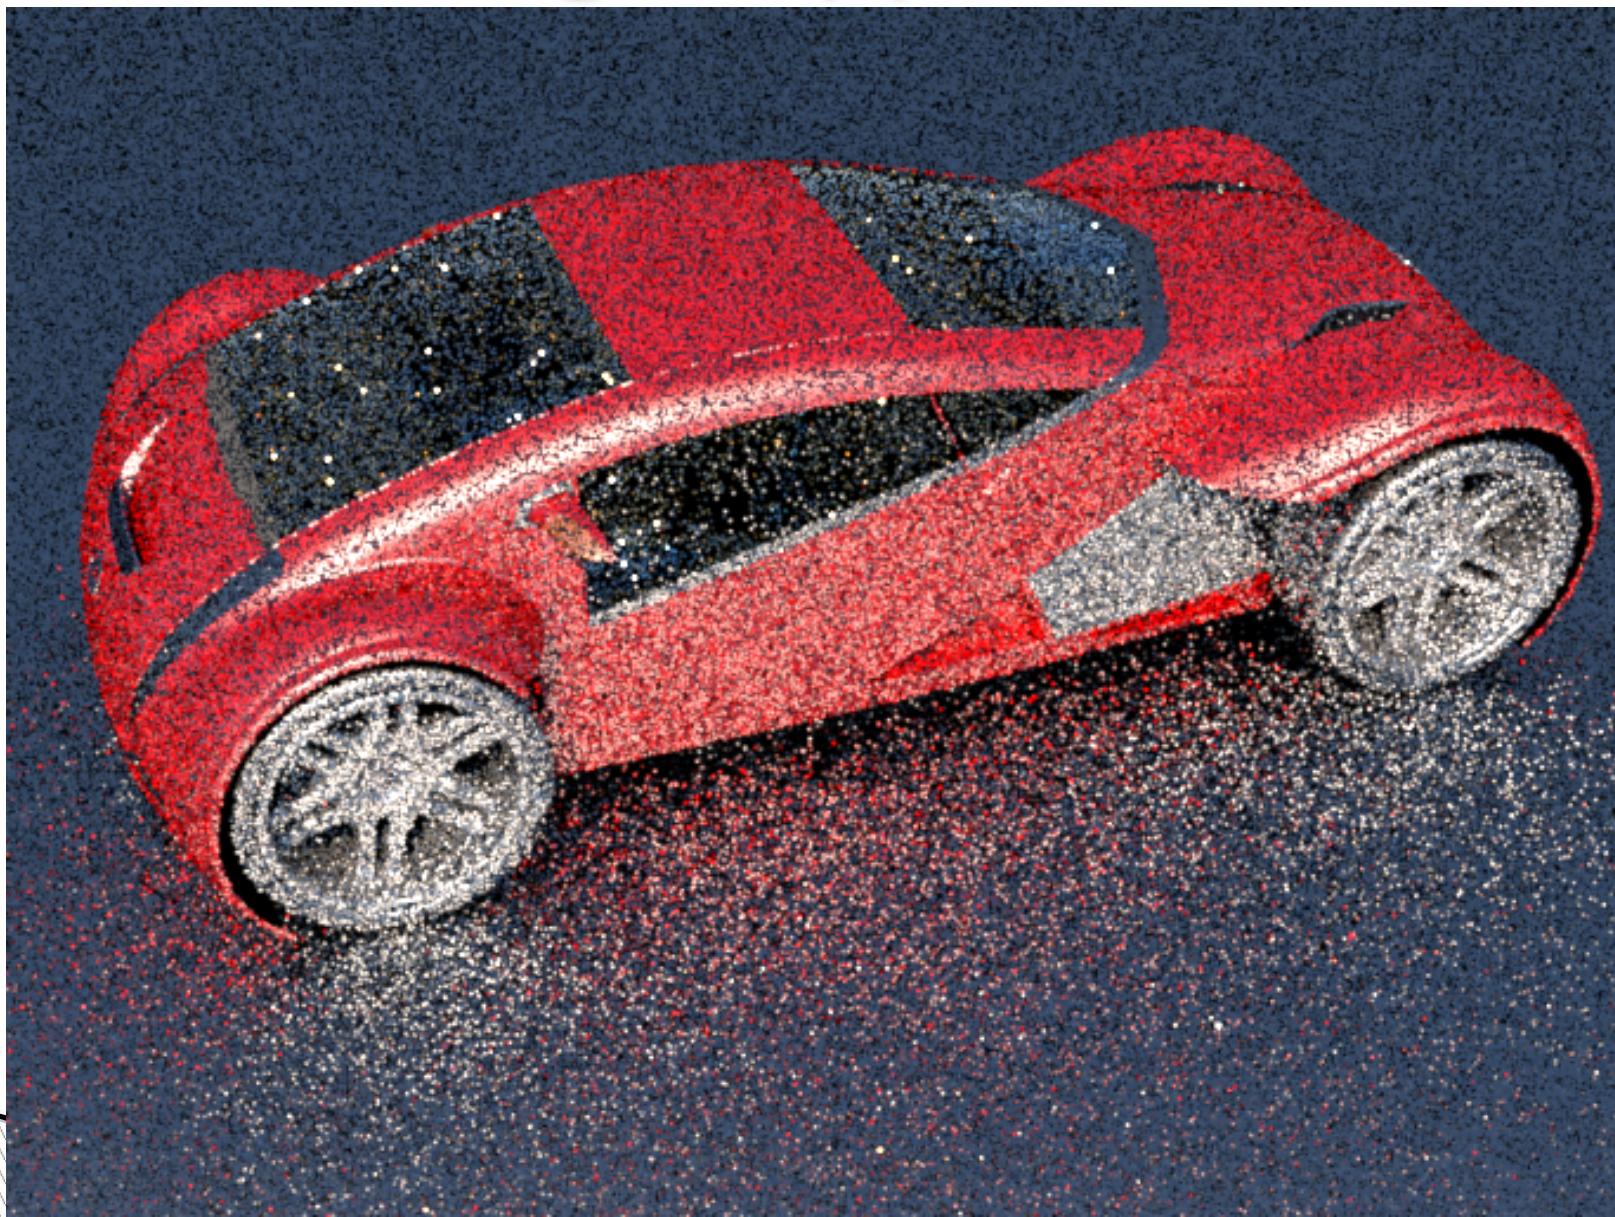


# Monte Carlo Path Tracing

- ▶ Trace only one ray at each recursion
- ▶ But trace several (hundreds of) primary rays for each pixel



# Path Tracing, 1 spp (0.7 s)



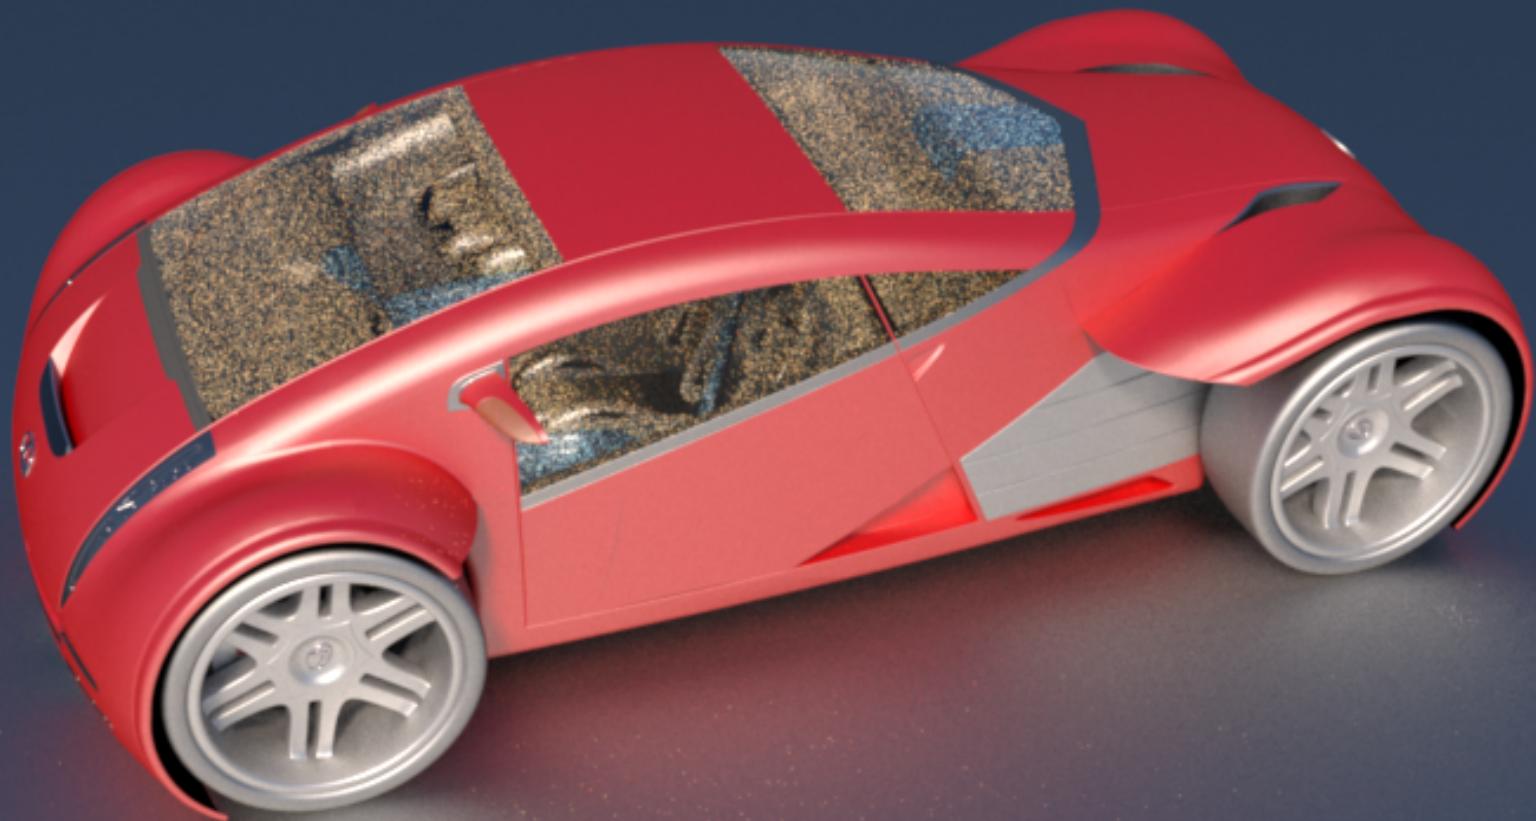
# Path Tracing, 32 spp (21 s)



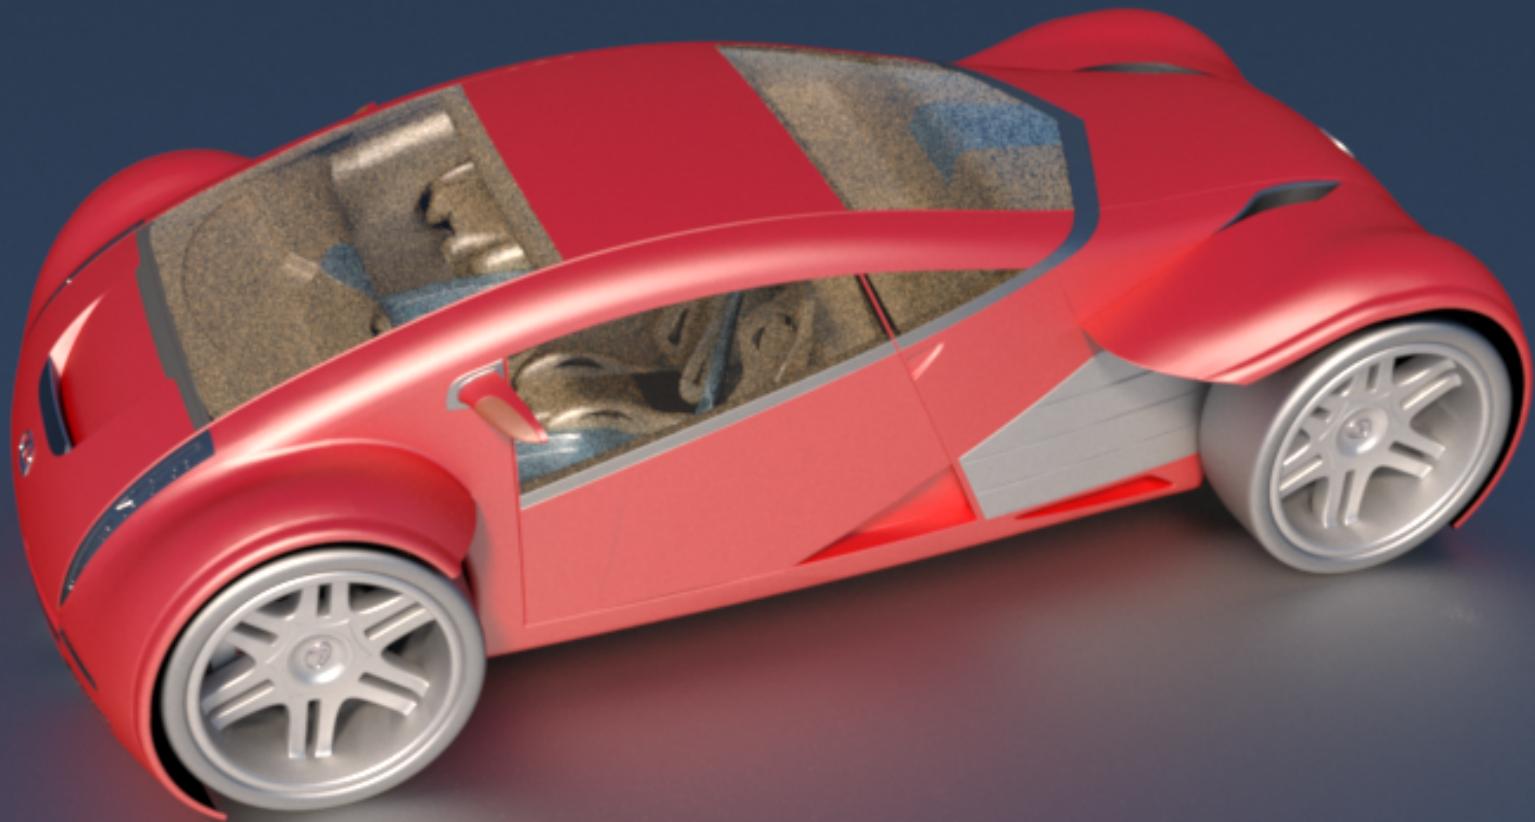
# Path Tracing, 256 spp (3 mn)



# Path Tracing, 1024 spp (12.4 mn)



PT, adaptive, 256-8192 spp (53 mn)

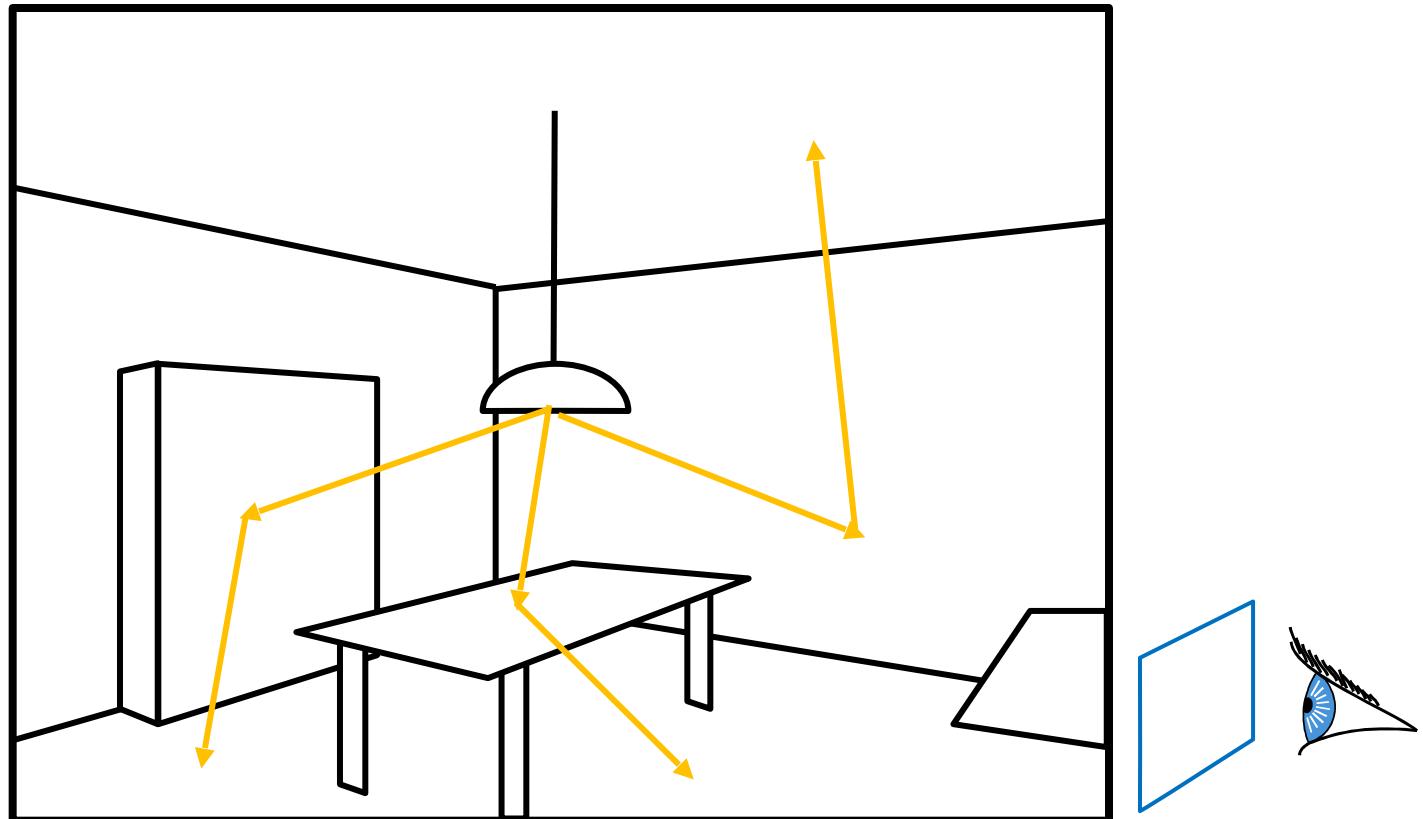


# What's happening?

- ▶ Light source: sun + sky
- ▶ Fast convergence outside
- ▶ Much slower inside
  - Light is refracted by windows
  - No direct path towards the light source

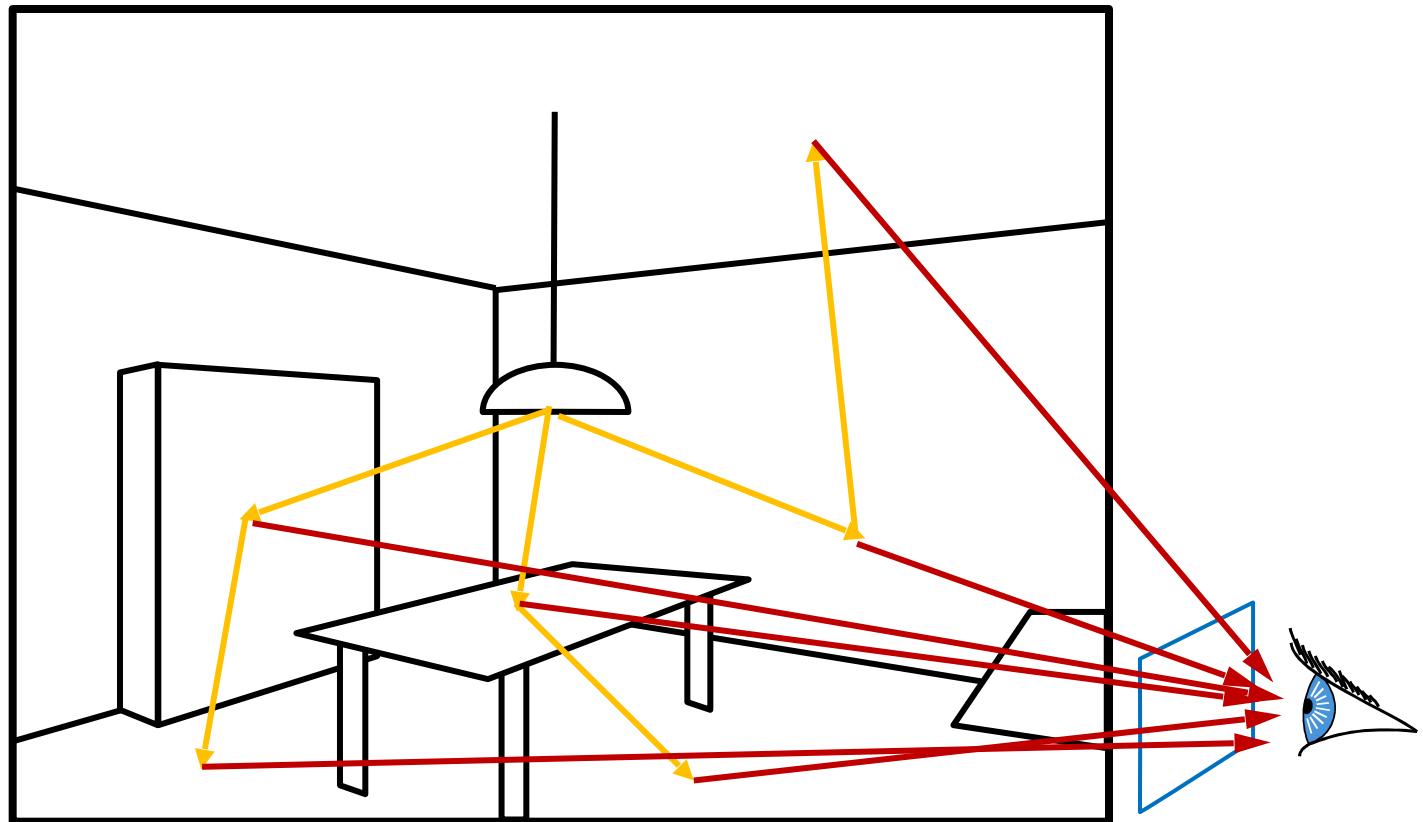
# Monte Carlo Light Path Tracing

- ▶ Start from the light source
- ▶ Send particles, let them bounce around

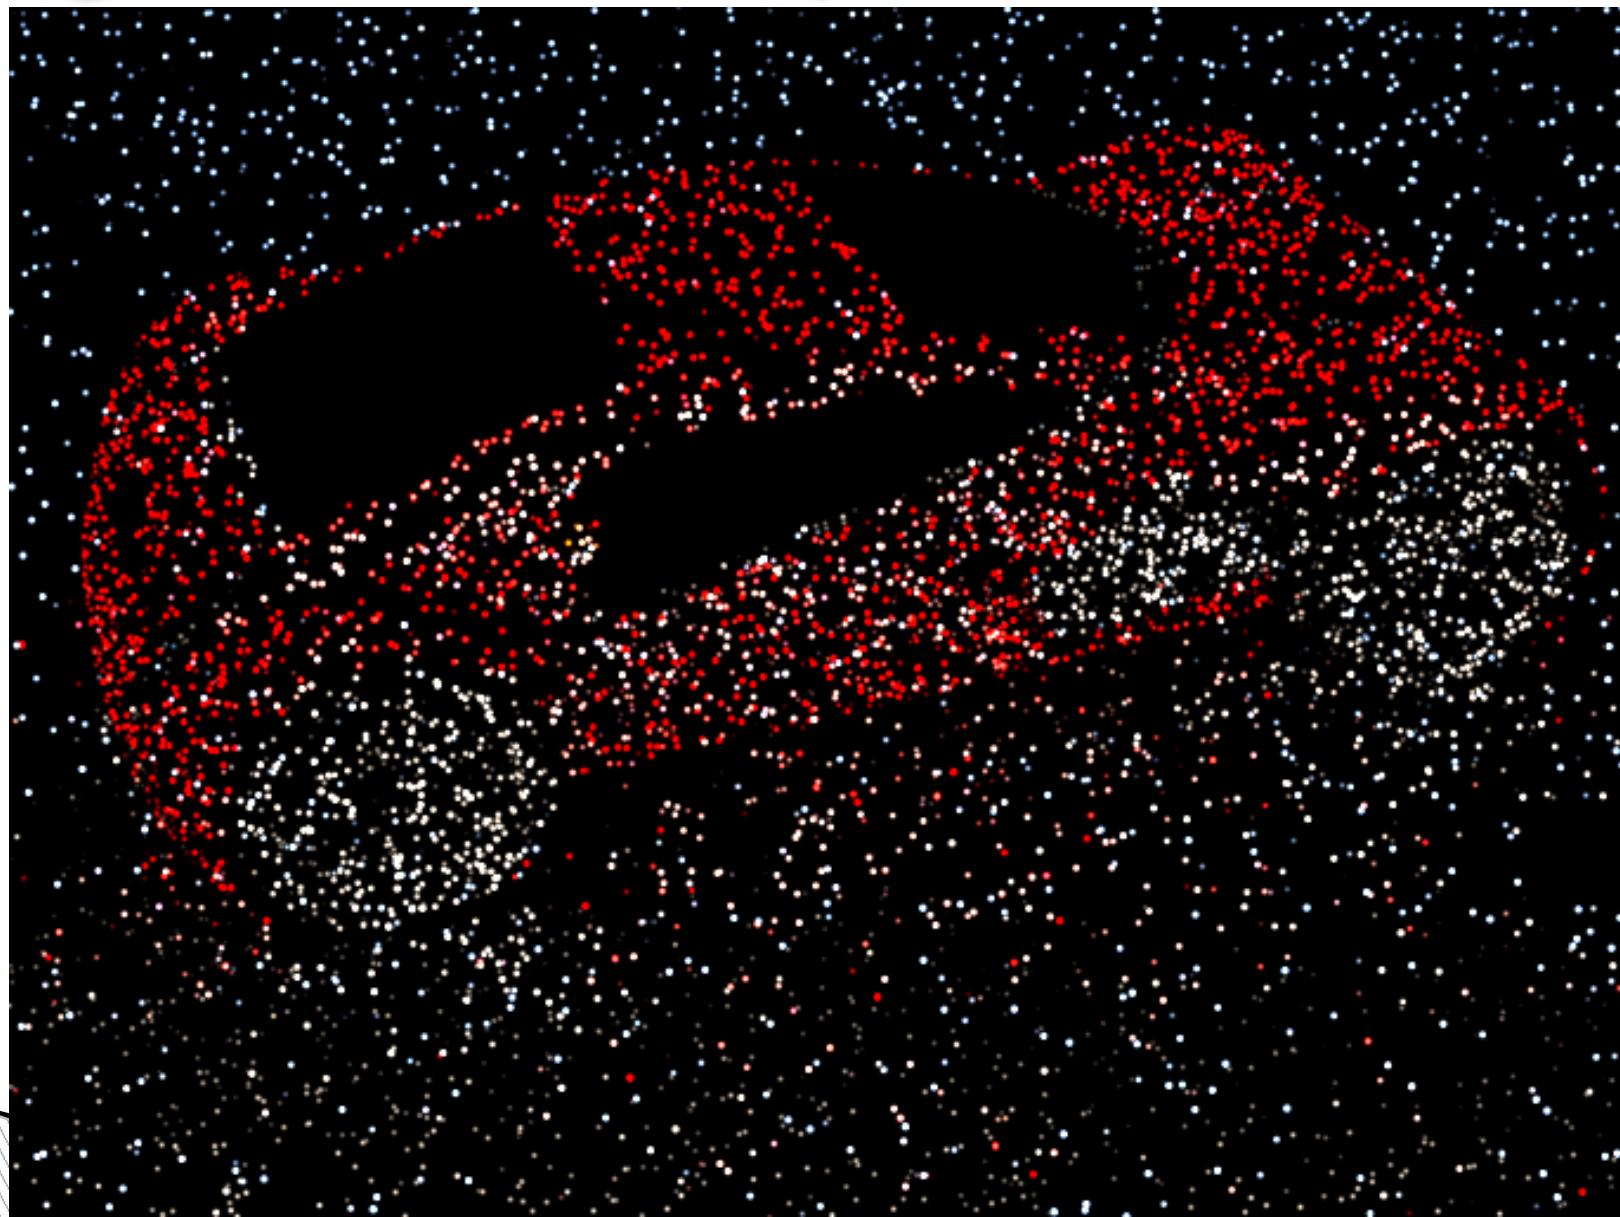


# Monte Carlo Light Path Tracing

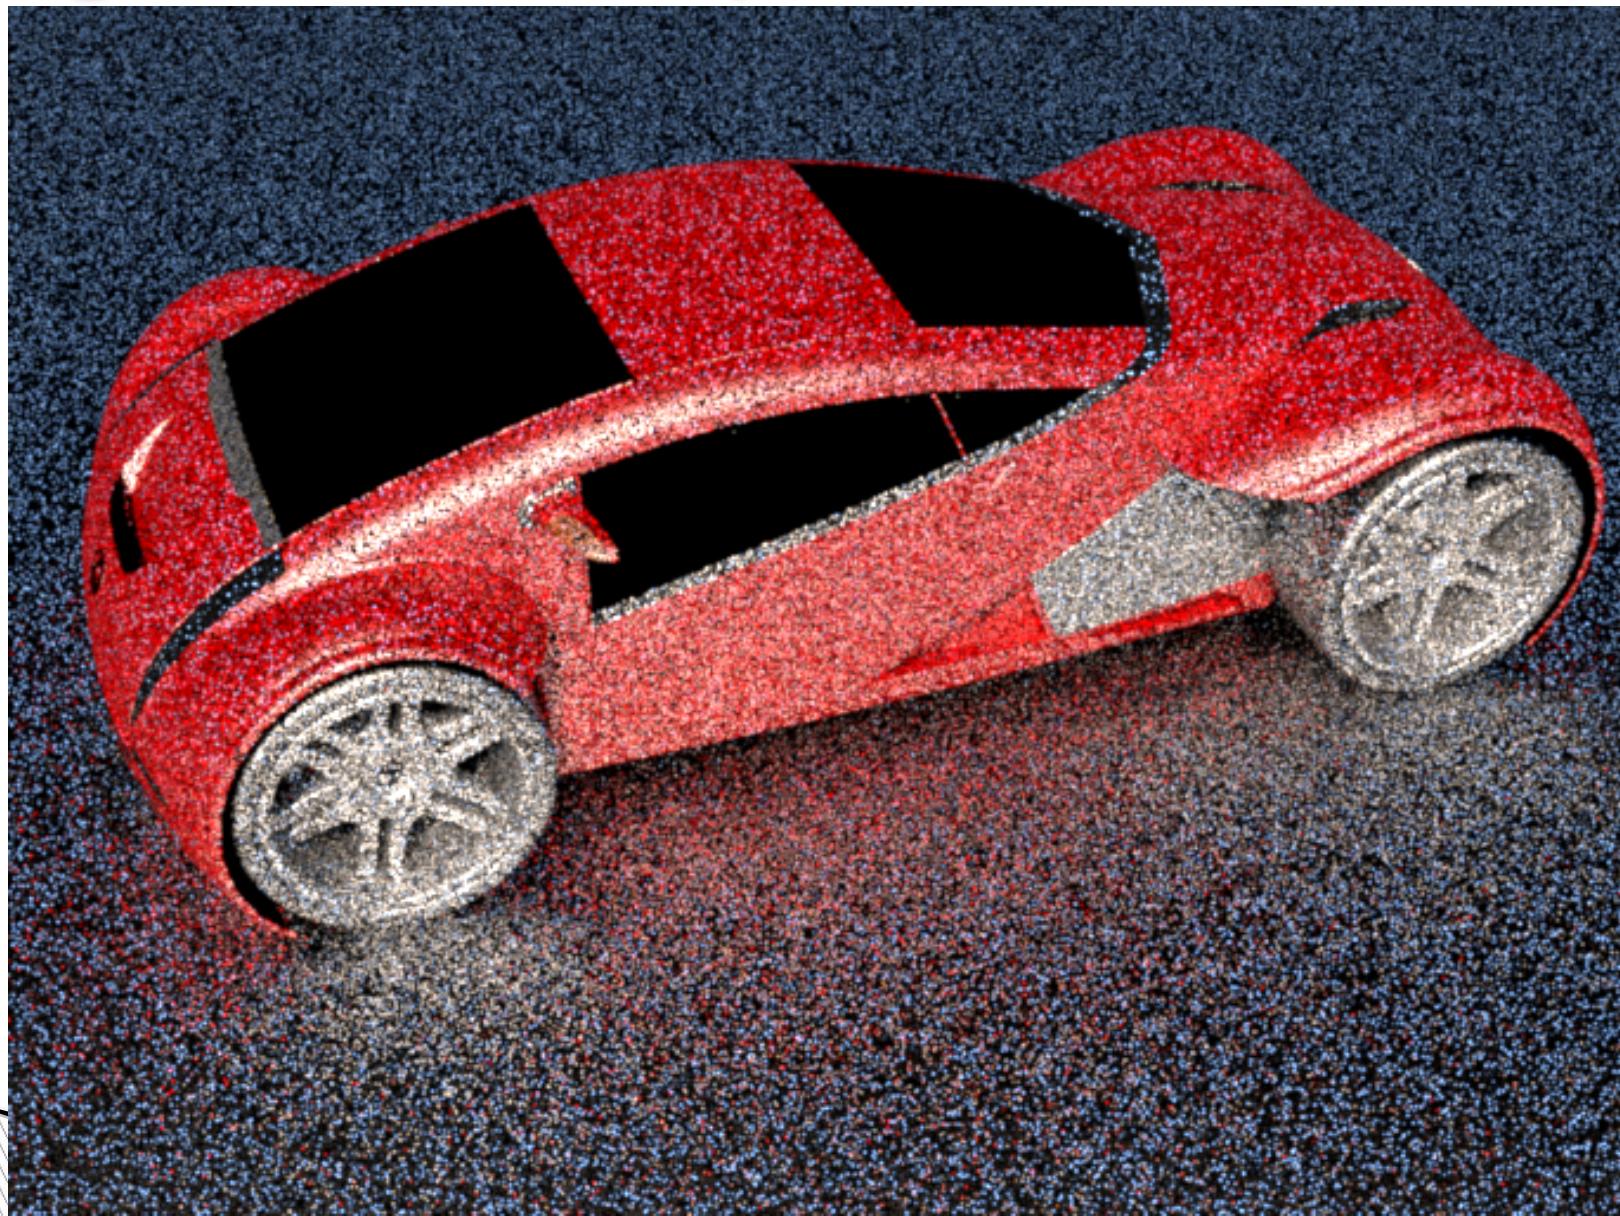
- ▶ Start from the light source
- ▶ Send particles, let them bounce around
- ▶ Sample the camera directly



# Light PT: 700M particles, 4 mn

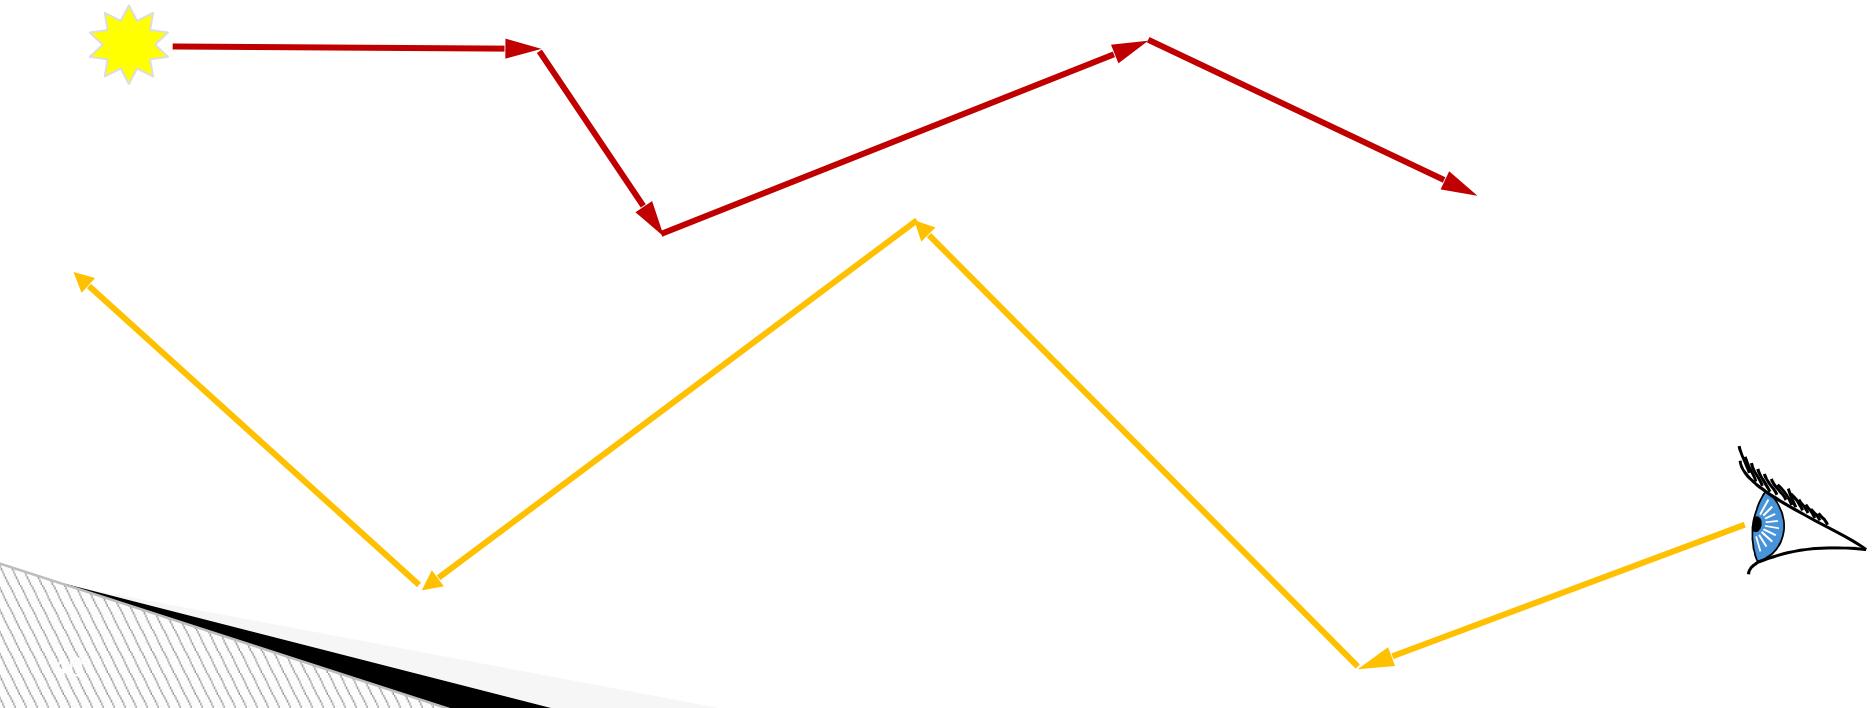


# Light PT: 70G particles, 6.7 h



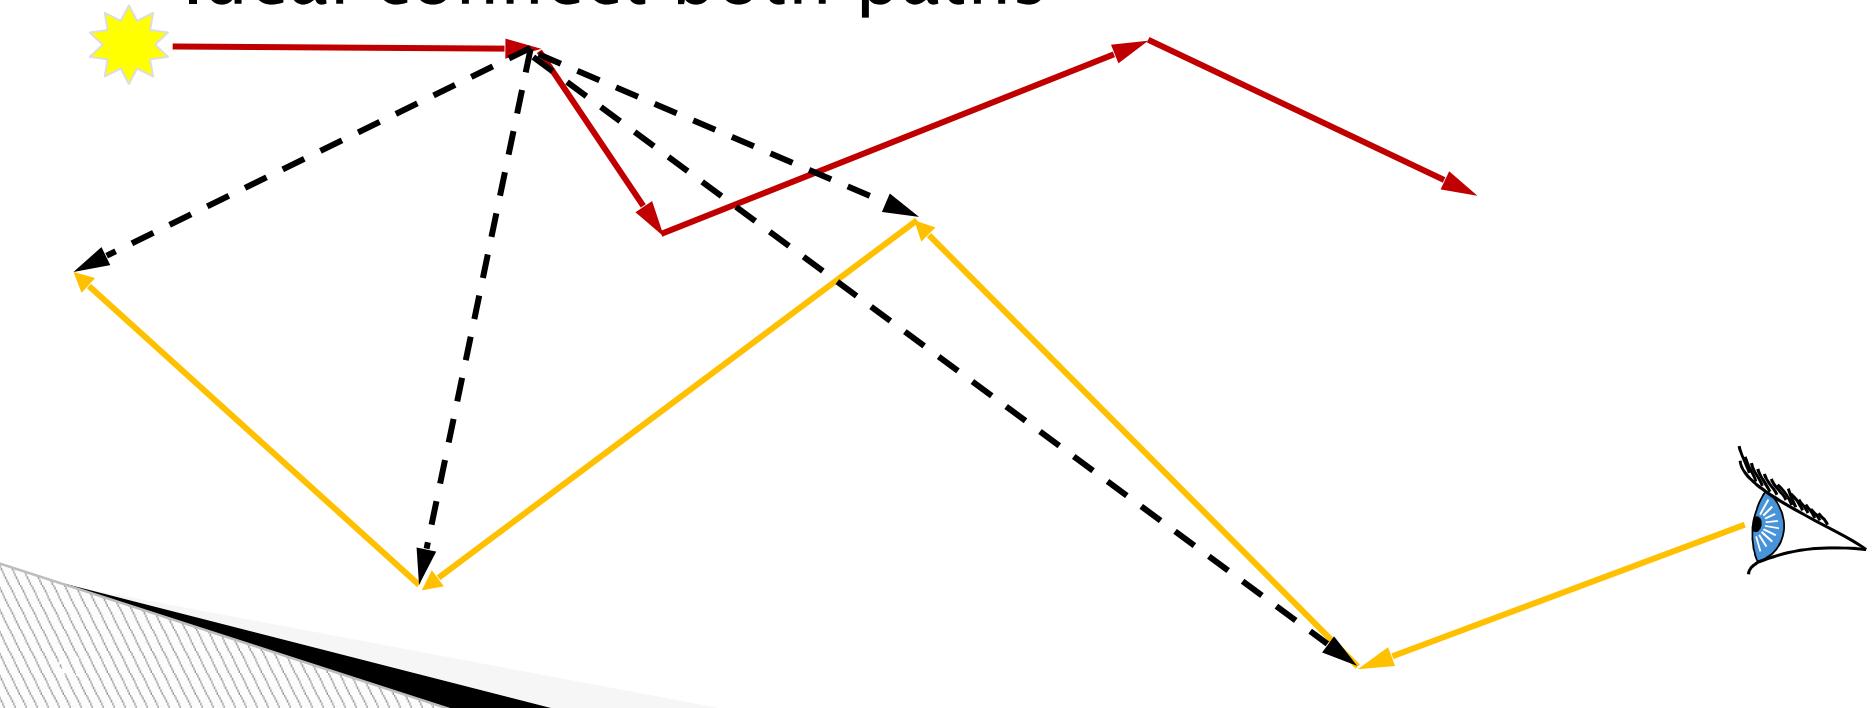
# Bi-directional path tracing

- ▶ Path-tracing is not enough
  - Need a connection towards light source, camera
- ▶ Connection between eye-paths and light-paths?
- ▶ Idea: connect both paths



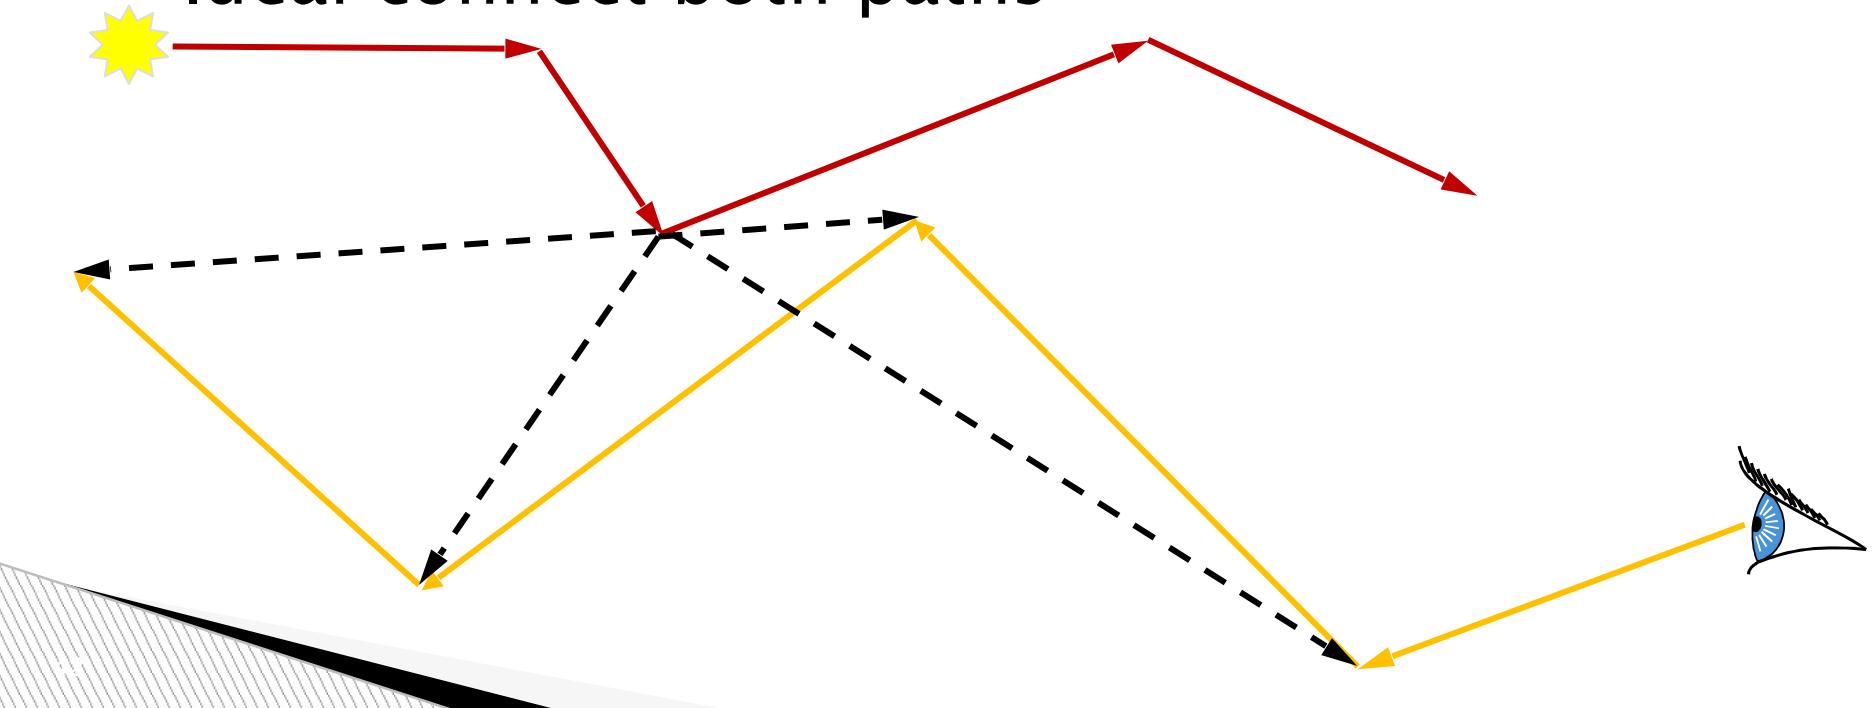
# Bi-directional path tracing

- ▶ Path-tracing is not enough
  - Need a connection towards light source, camera
- ▶ Connection between eye-paths and light-paths?
- ▶ Idea: connect both paths



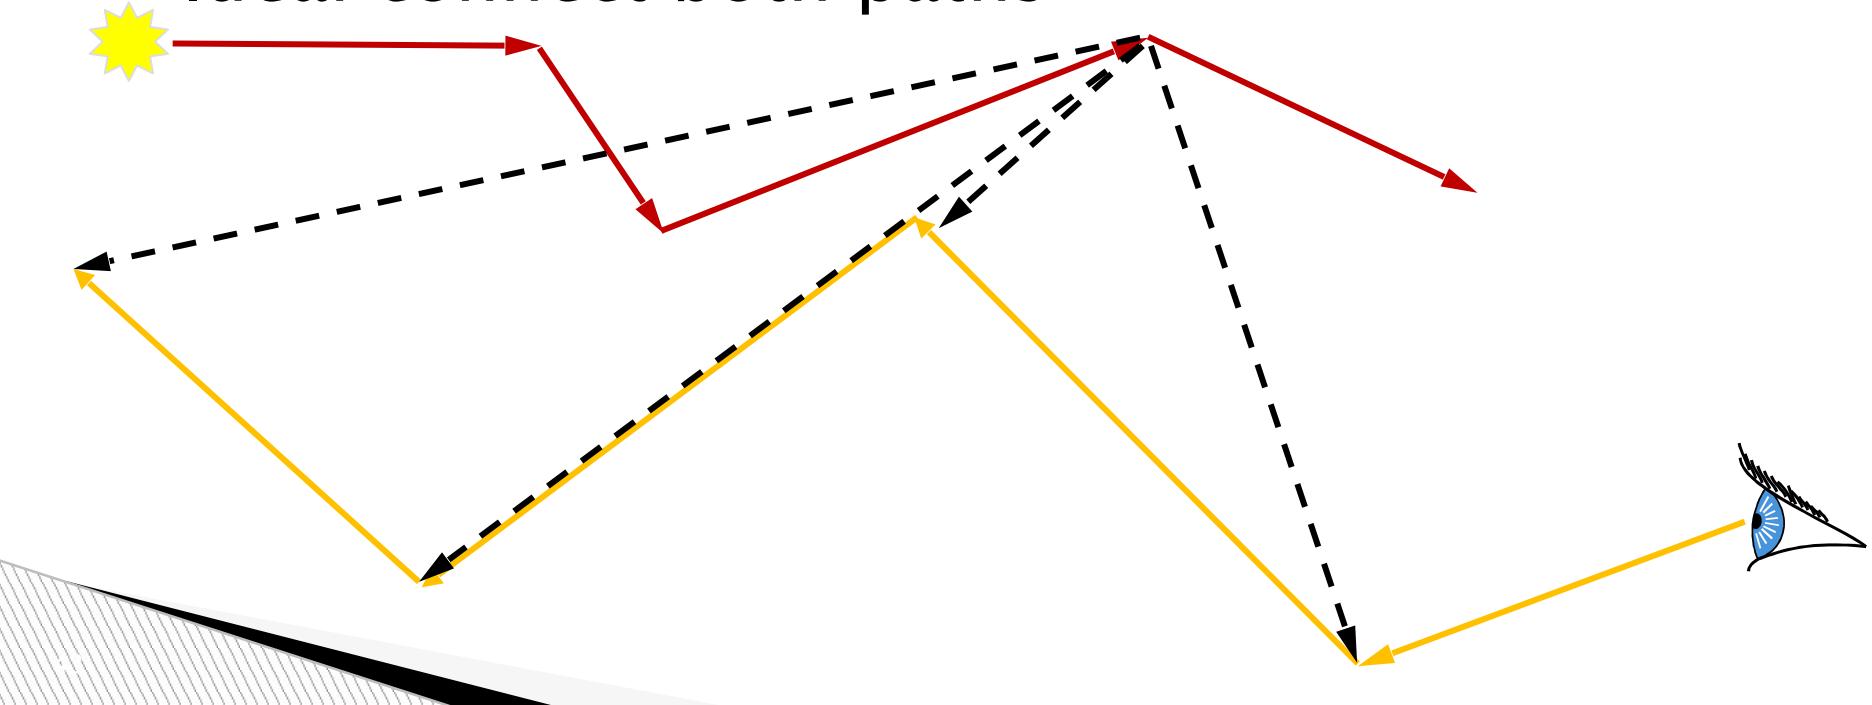
# Bi-directional path tracing

- ▶ Path-tracing is not enough
  - Need a connection towards light source, camera
- ▶ Connection between eye-paths and light-paths?
- ▶ Idea: connect both paths



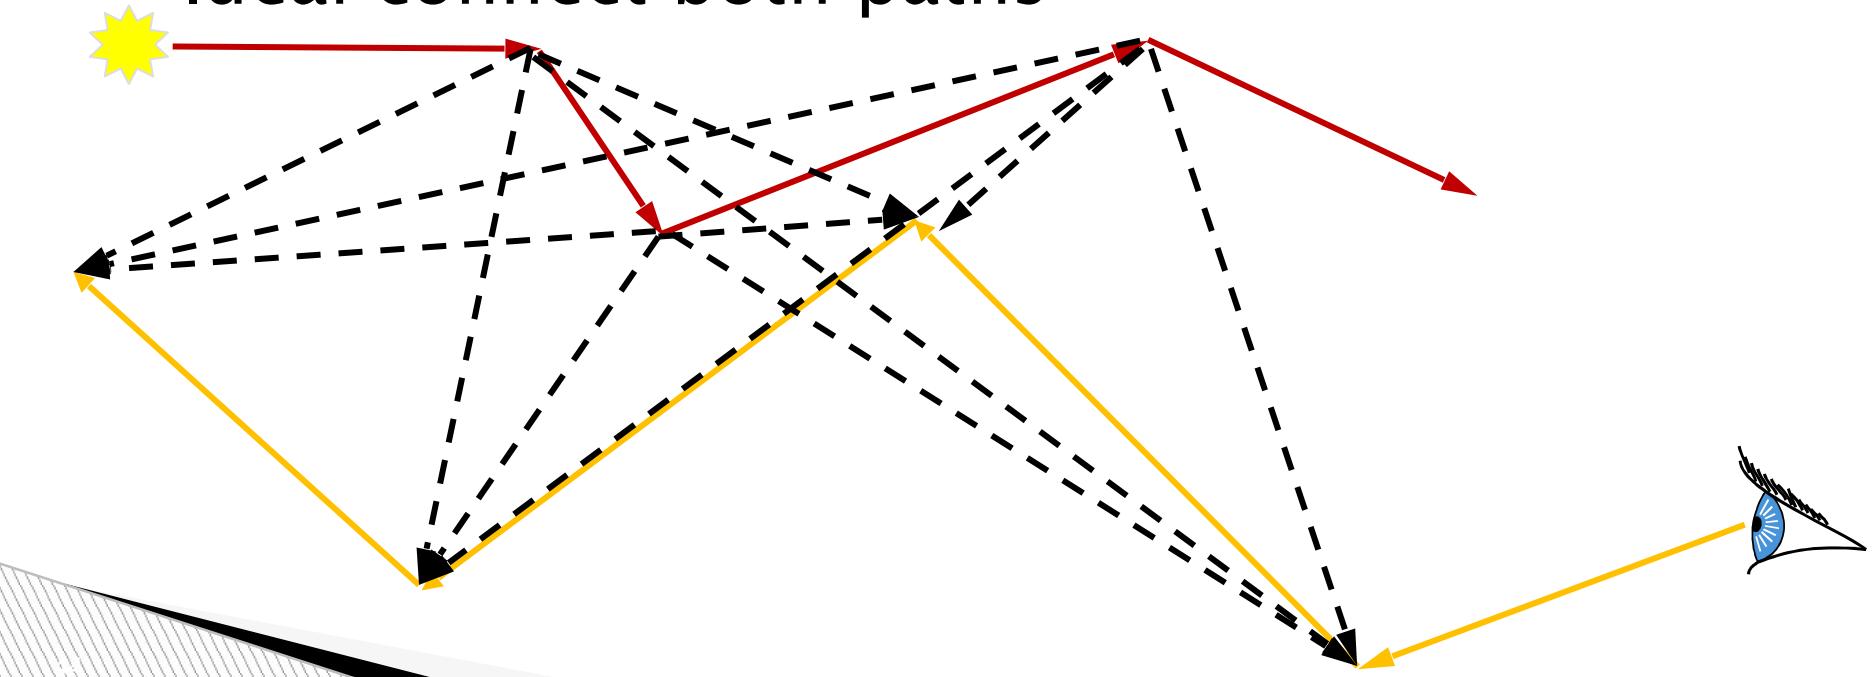
# Bi-directional path tracing

- ▶ Path-tracing is not enough
  - Need a connection towards light source, camera
- ▶ Connection between eye-paths and light-paths?
- ▶ Idea: connect both paths



# Bi-directional path tracing

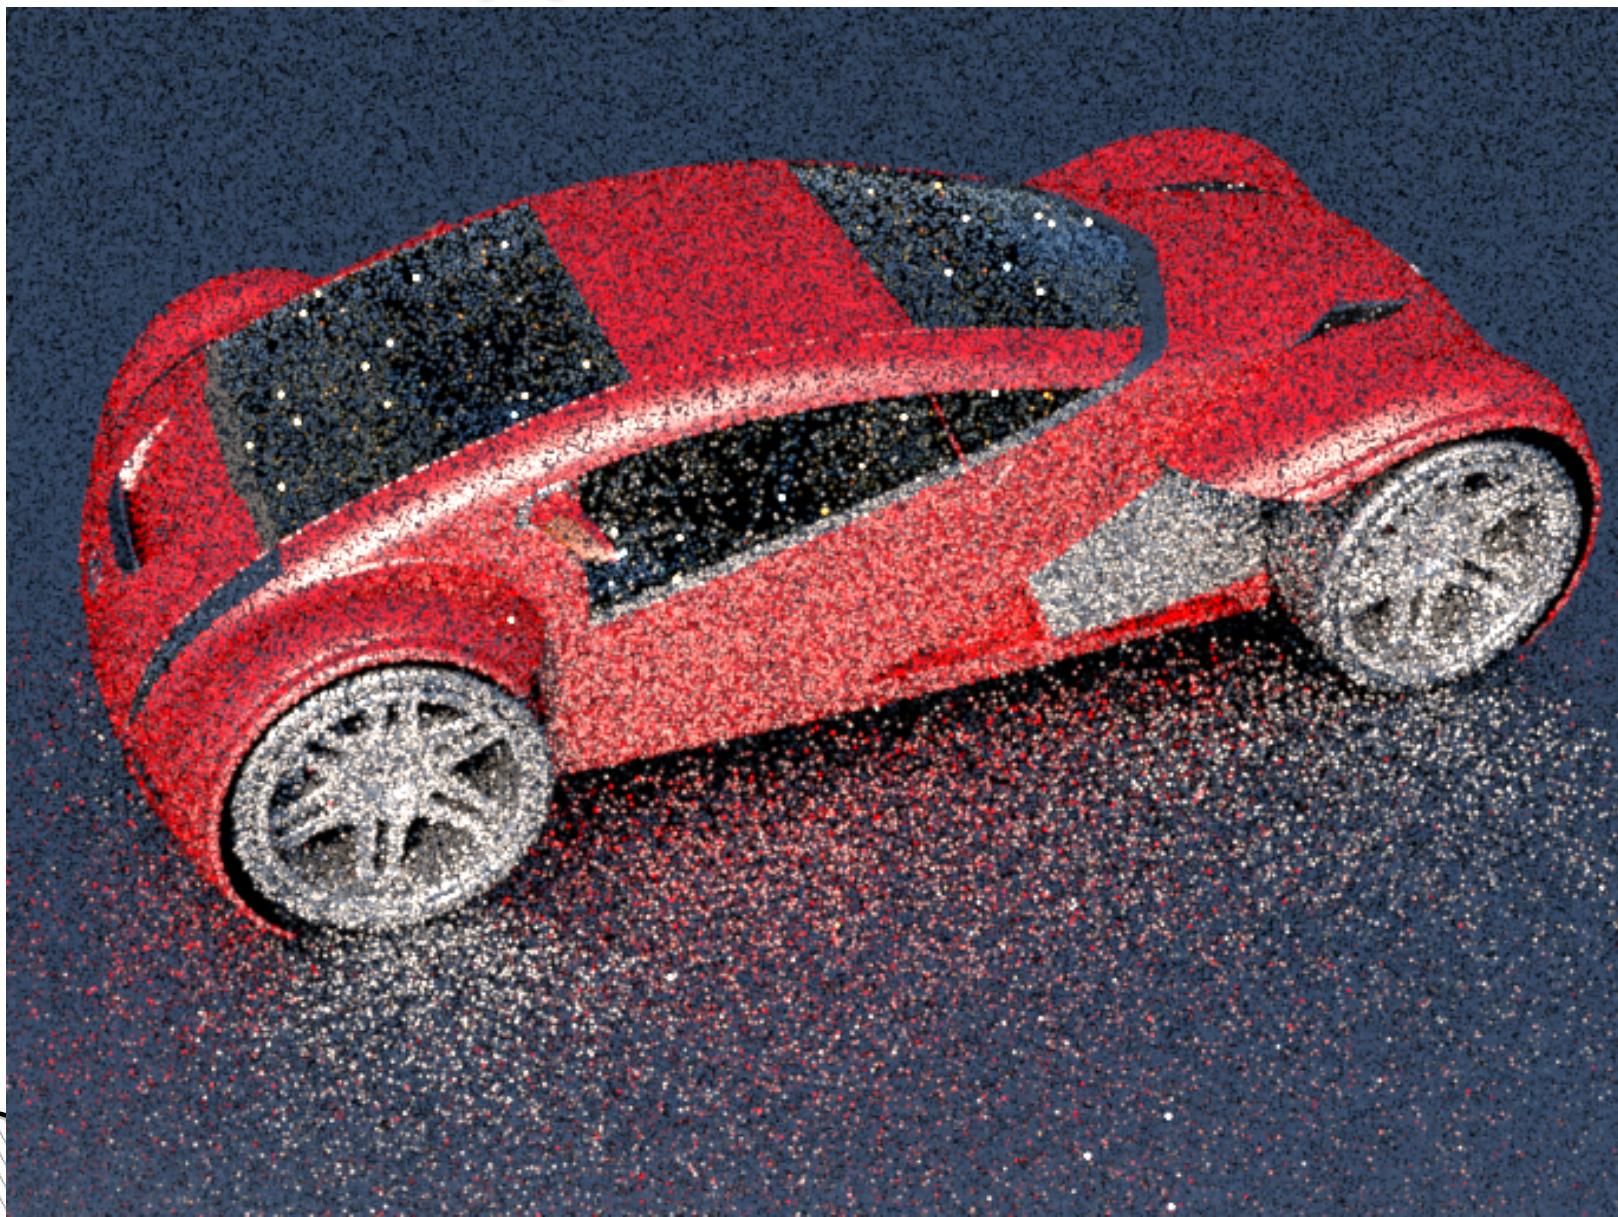
- ▶ Path-tracing is not enough
  - Need a connection towards light source, camera
- ▶ Connection between eye-paths and light-paths?
- ▶ Idea: connect both paths



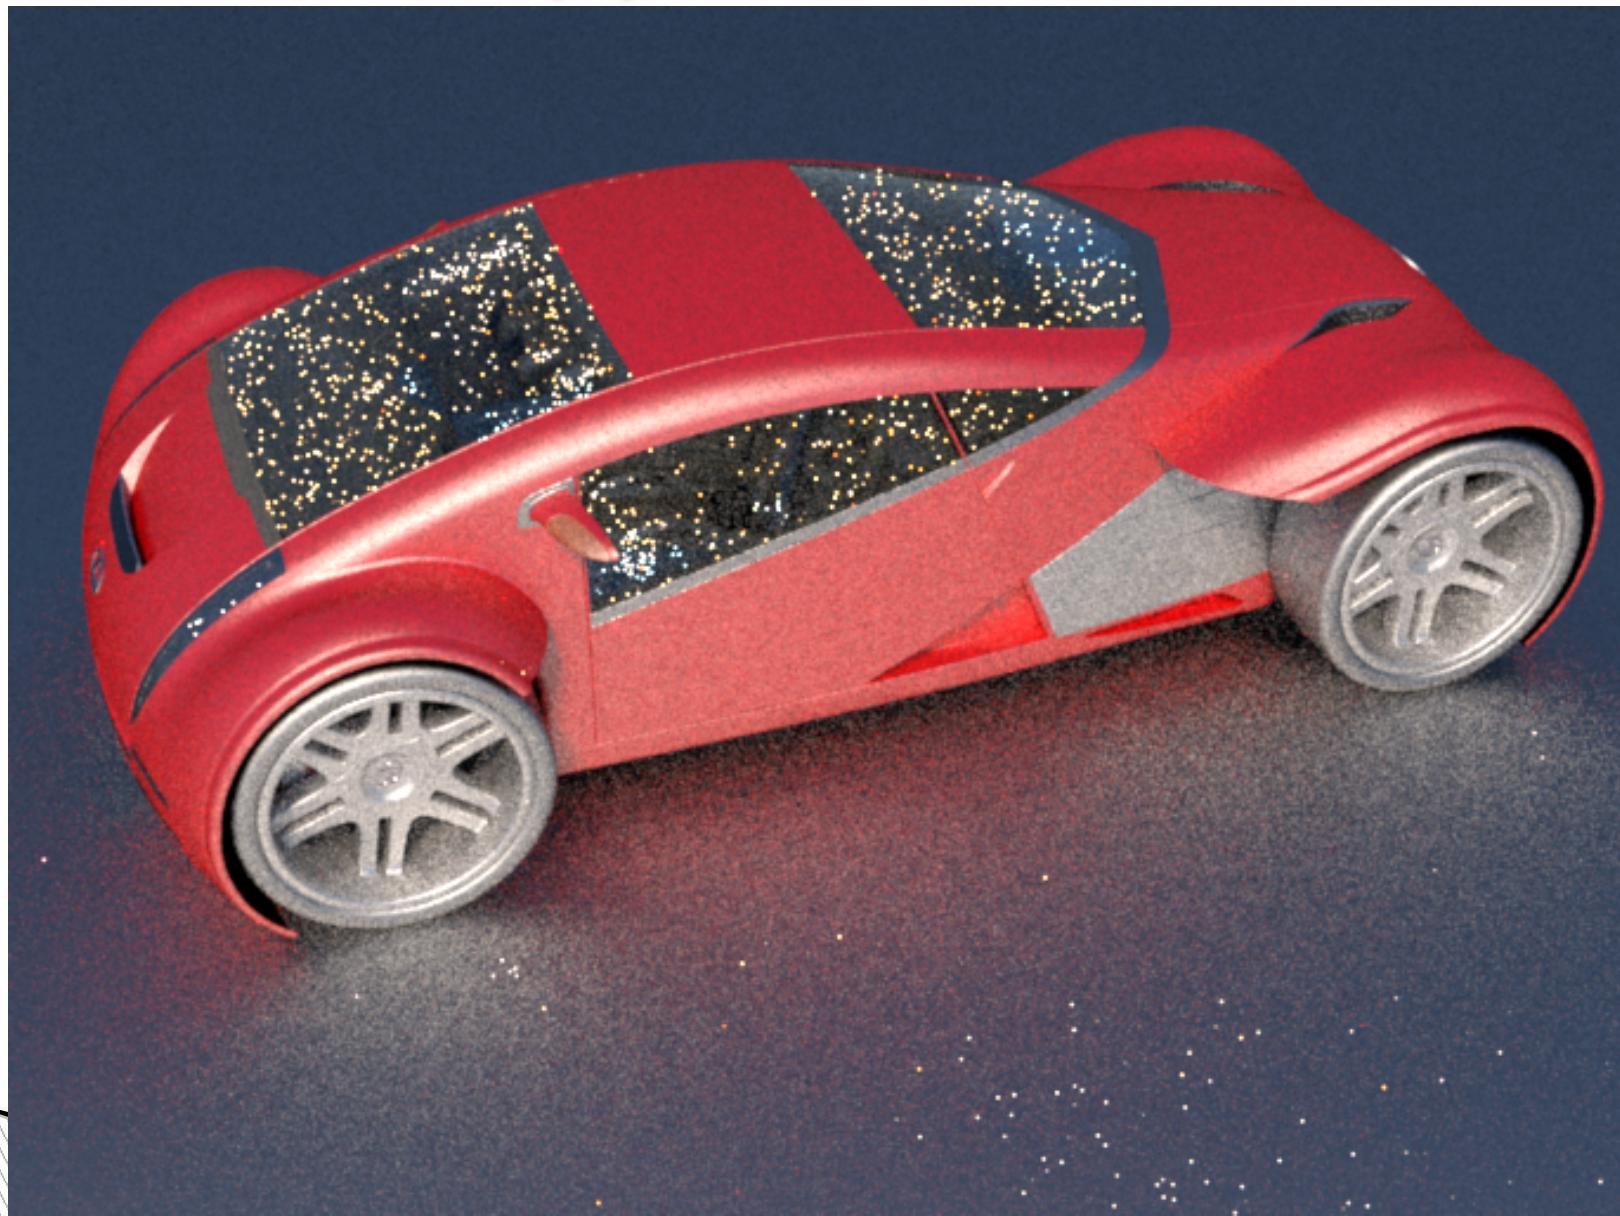
# Bi-directional path-tracing

- ▶ Tracer light-paths and eye-paths
- ▶ Connect these paths together
  - Only for non-specular path vertices
- ▶ Probability? Weight?
- ▶ Express things in path-space
- ▶ Multiple Importance-sampling

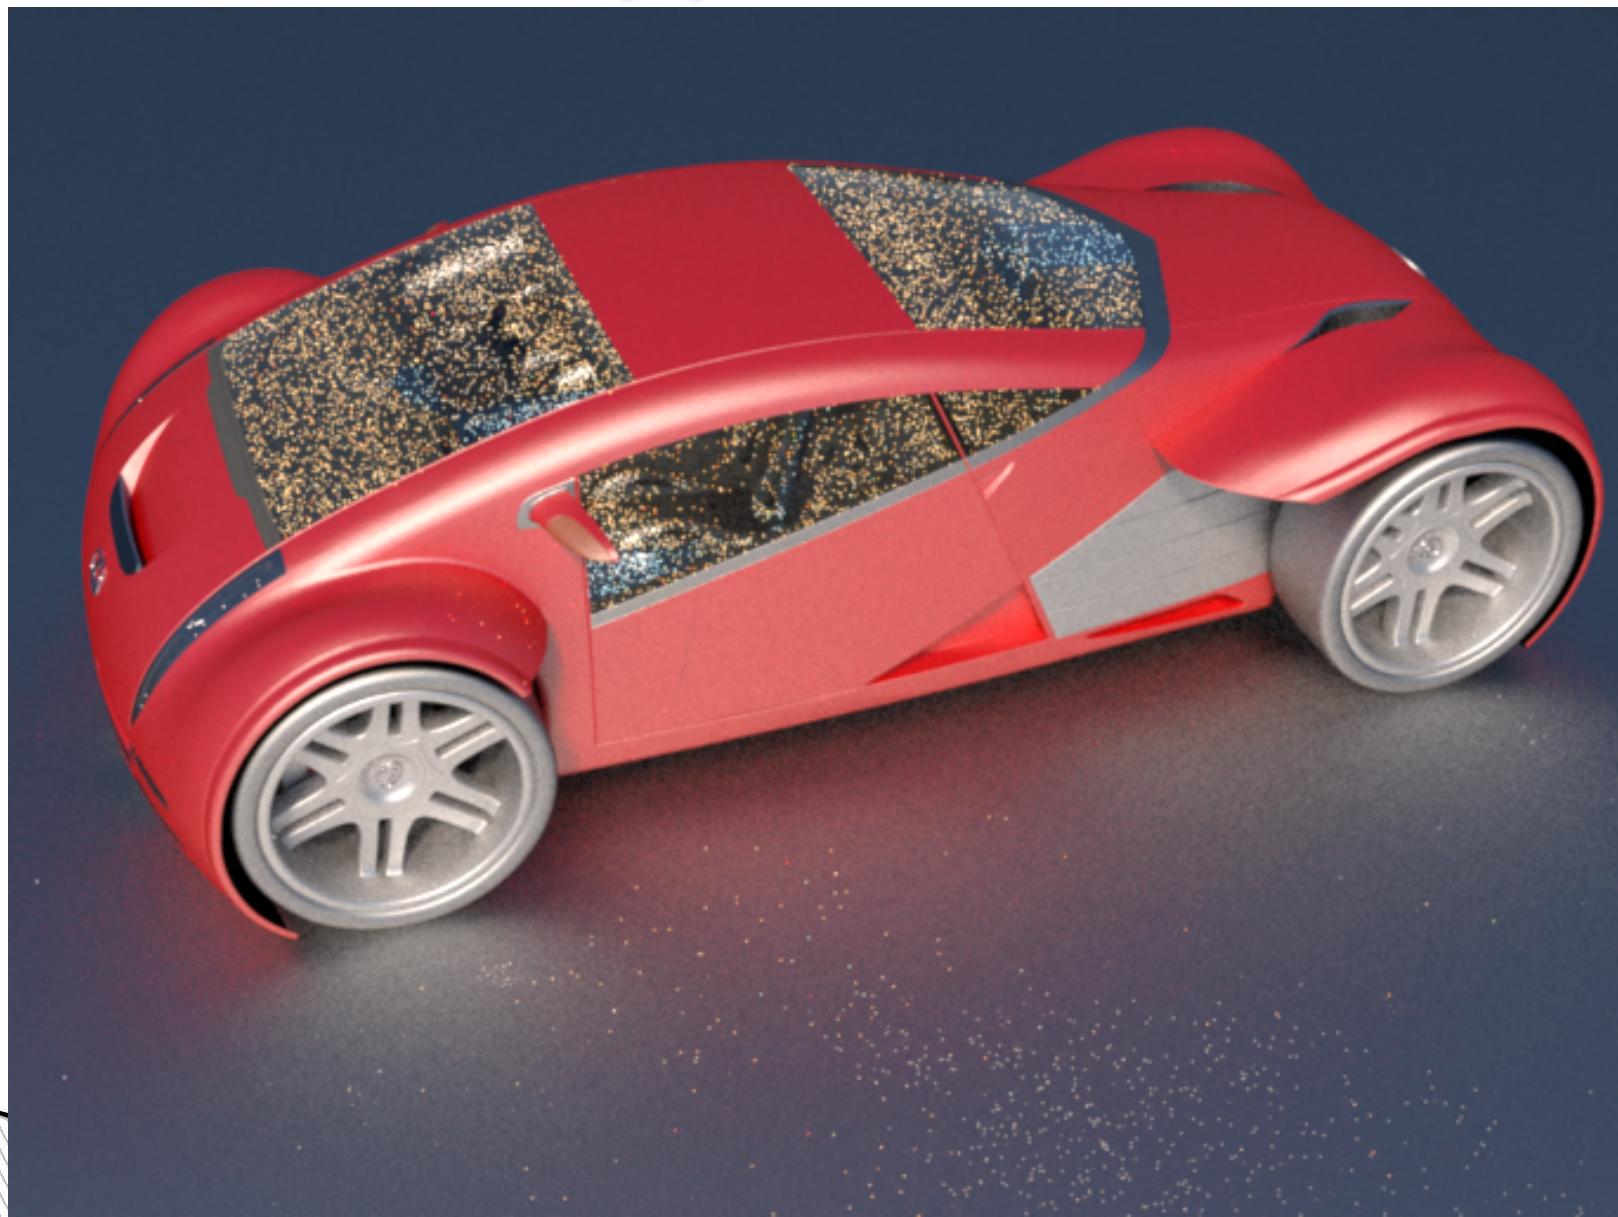
BDPT, 1 spp, 2s



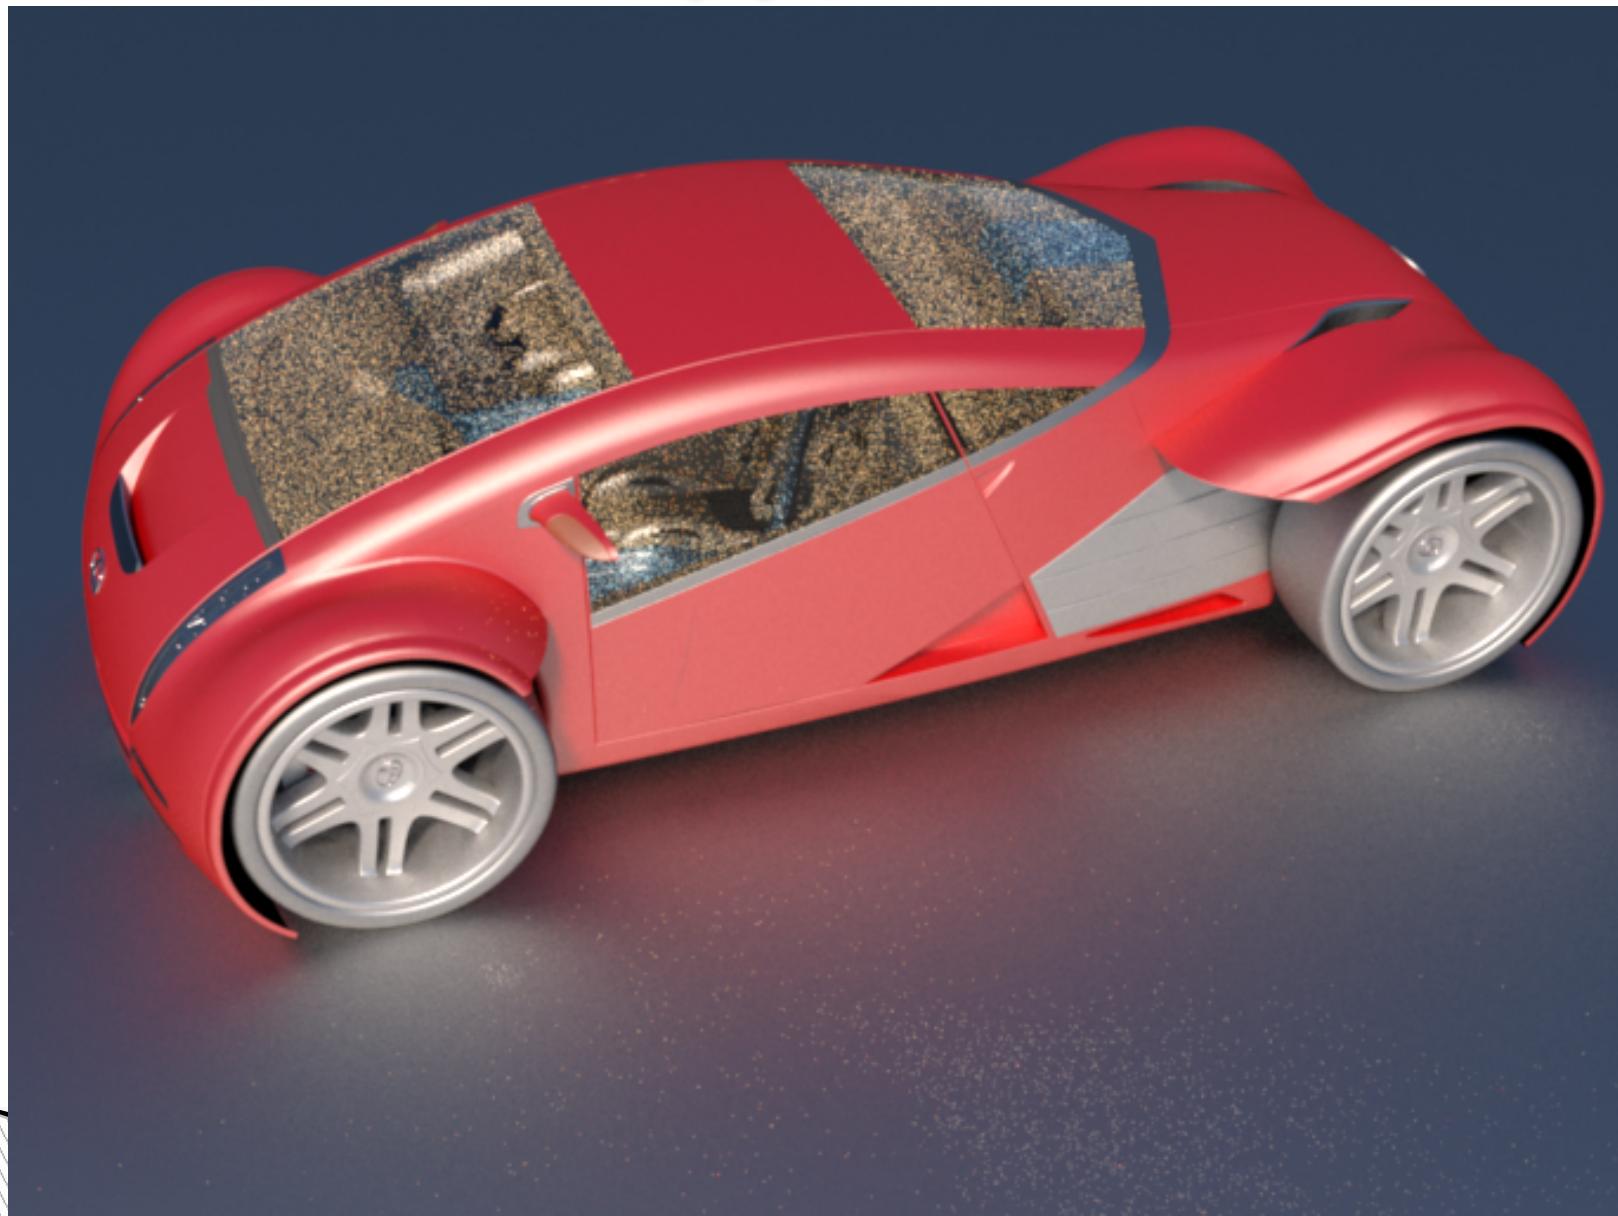
BDPT, 32 spp, 42 s



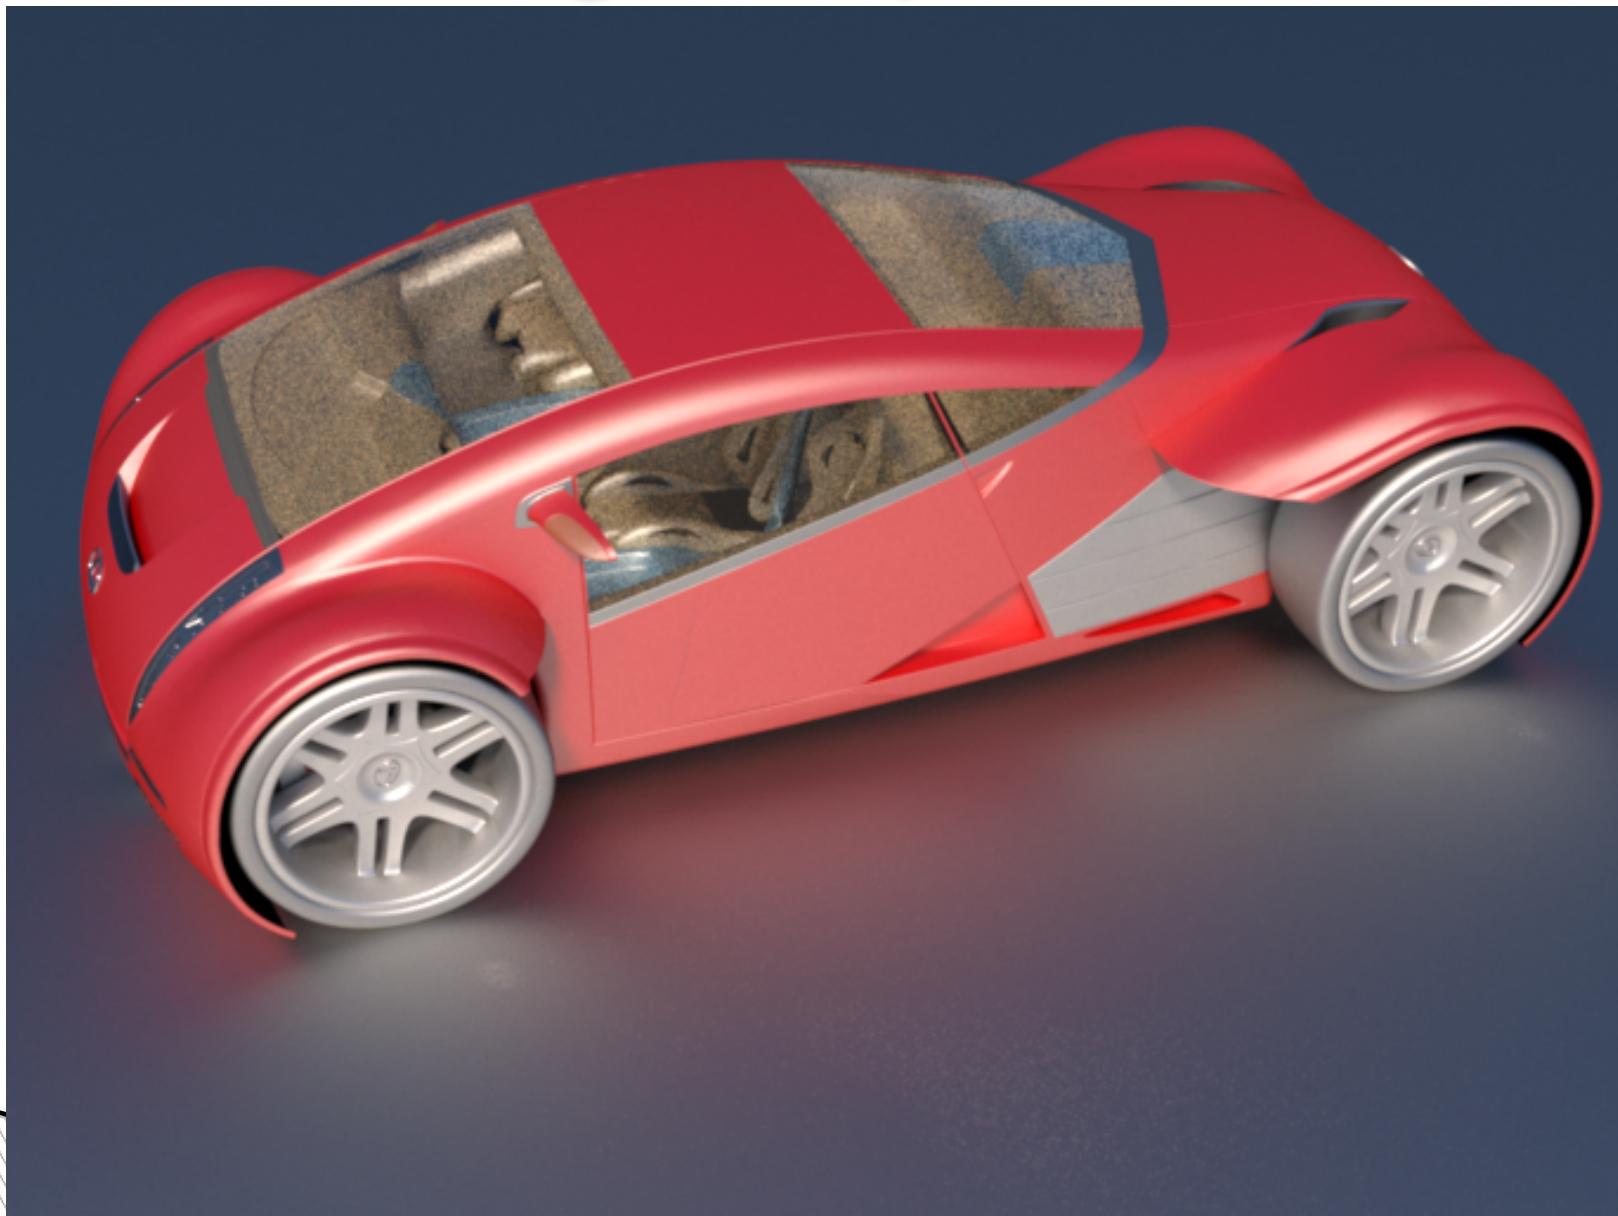
BDPT, 256 spp, 5.4 mn



BDPT, 1024 spp, 23 mn



# Path-Tracing, adaptive, 53 mn



# Bi-Directional Path Tracing

- ▶ Faster convergence than Path-Tracing
- ▶ Less noise
- ▶ Inside the car:
  - Light being refracted by the windows
  - Hard to find connections
  - Reverts to path-tracing
  - BDPT better than PT for small light sources

# Bi-Directional Path Tracing



Path Tracing



Bi-Directional Path Tracing  
(Veach and Guibas 1995)

# Advanced Monte-Carlo techniques

- ▶ Random samples generation
- ▶ Importance Sampling
- ▶ Irradiance Caching
- ▶ Photon Mapping
- ▶ Metropolis Light Transport

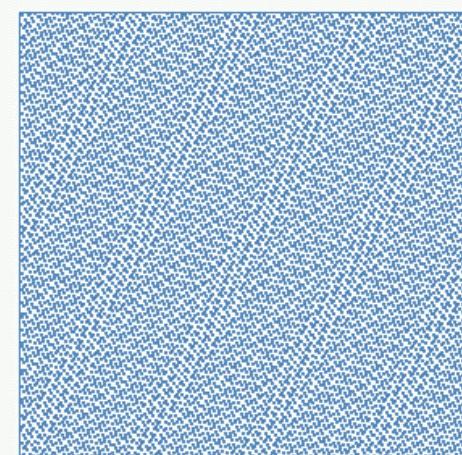
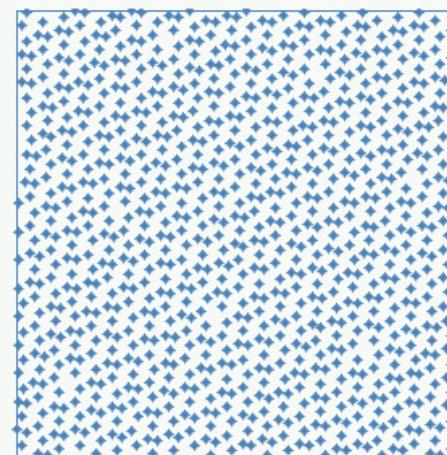
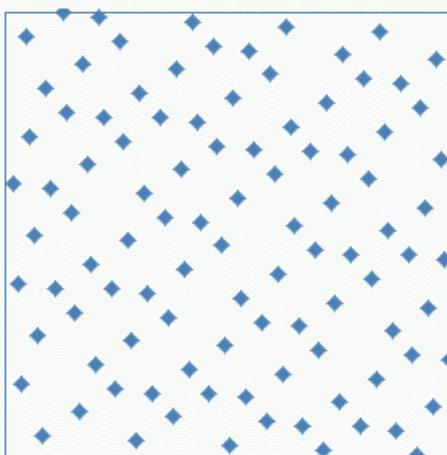
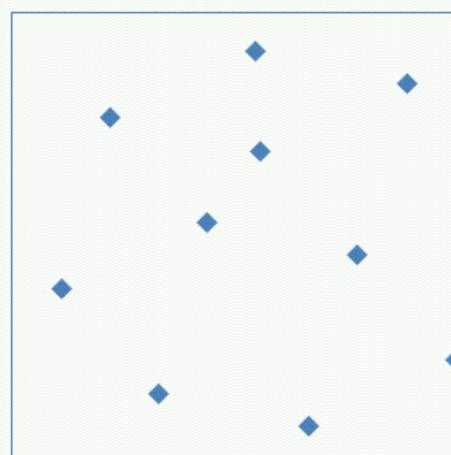
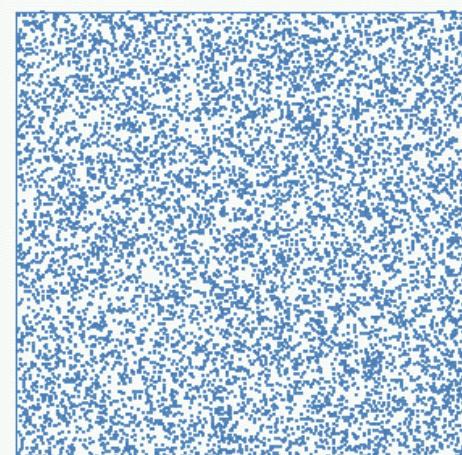
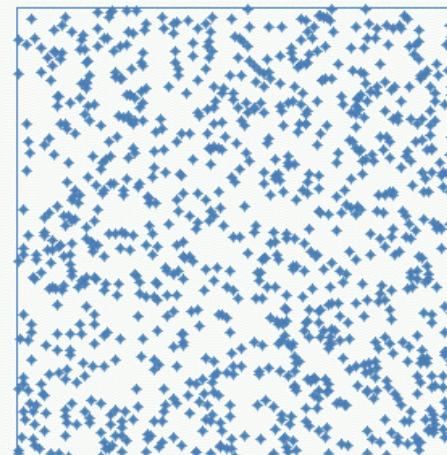
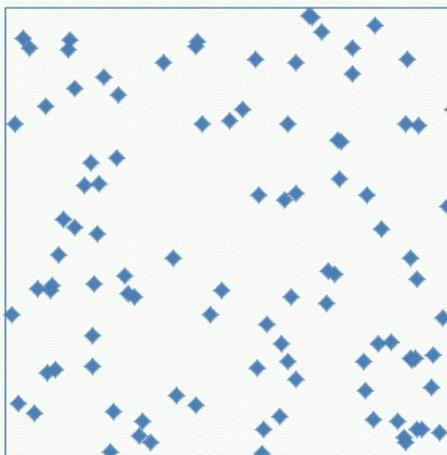
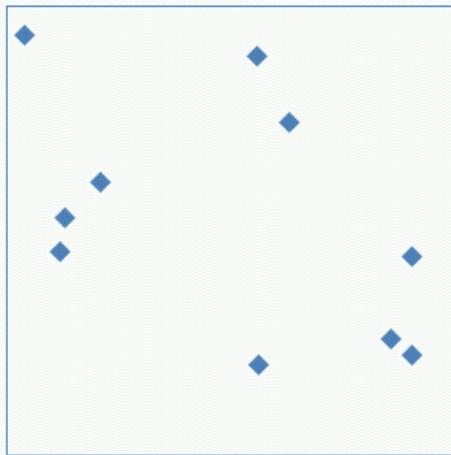
# Random samples generation

- ▶ Pseudo-random sampling sequence
- ▶ The structure appears:



# Quasi-random samples

Fully random sequences

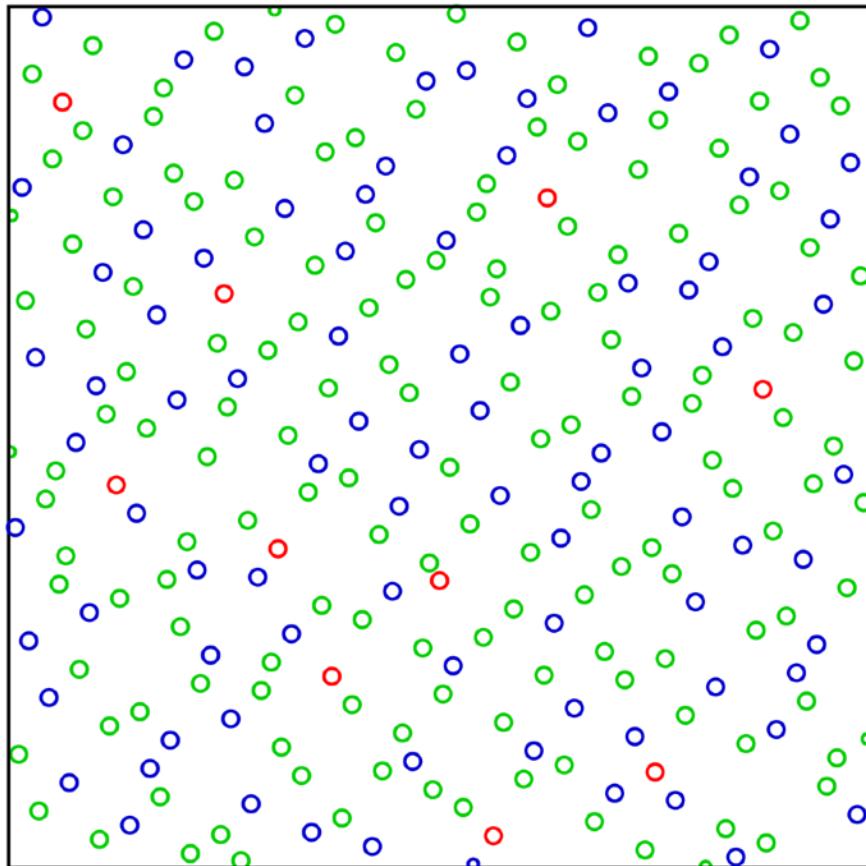


Quasi random sequences

# Quasi-random sequences

- ▶ Fully random sample points:
  - ▶ Easy to generate continuously
  - ▶ No uniform coverage of sample space
  - ▶ Discrepancy: slow convergence
- ▶ Quasi-random sequences:
  - ▶ Uniform coverage
  - ▶ Low discrepancy: converges faster
  - ▶ Many versions: Halton, Sobol, Hammersley
  - ▶ Usually built so you can use a subset

# Example: Halton sequence



Red= 1..10

Blue= 11..100

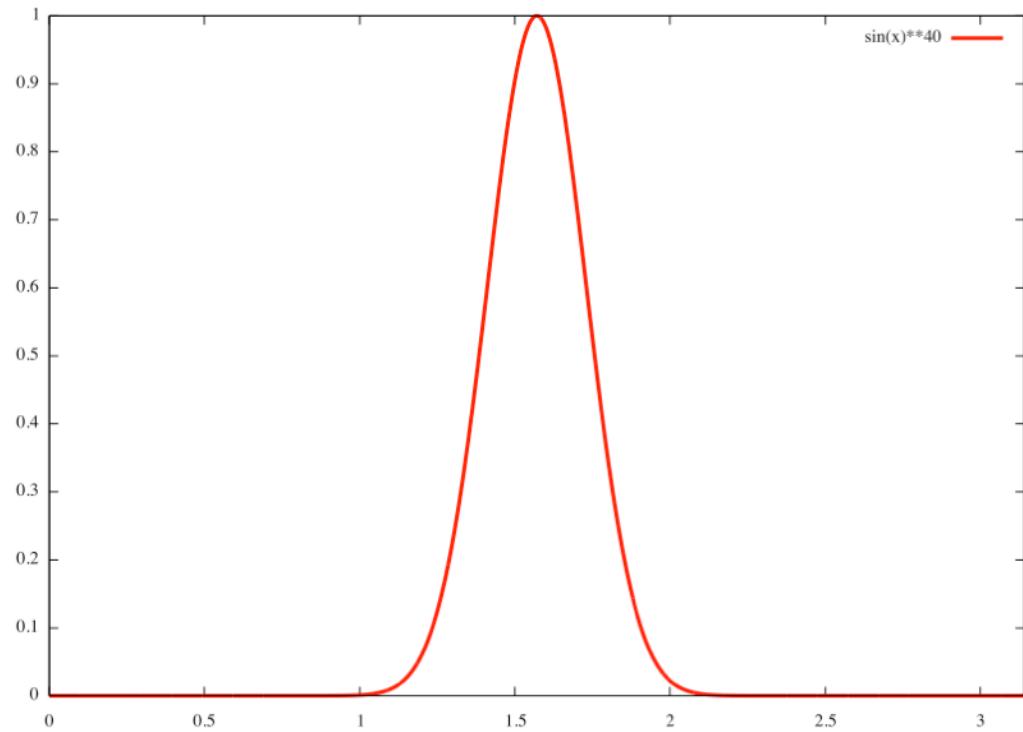
Green= 101..256

# Quasi-random sequences

- ▶ There is a structure!
- ▶ It can become visible
- ▶ Don't use always the same sequence:
  - ▶ Different starting point
  - ▶ Rotated 90°
  - ▶ Mirrored
  - ▶ ...

# Importance Sampling

- ▶ Functions with tight lobes/small support
- ▶ Random points outside the lobe: useless
- ▶ Can we guide randomness?



# Non-uniform distribution

- ▶ N samples with probability  $p(x)$
- ▶ Monte Carlo estimator becomes:

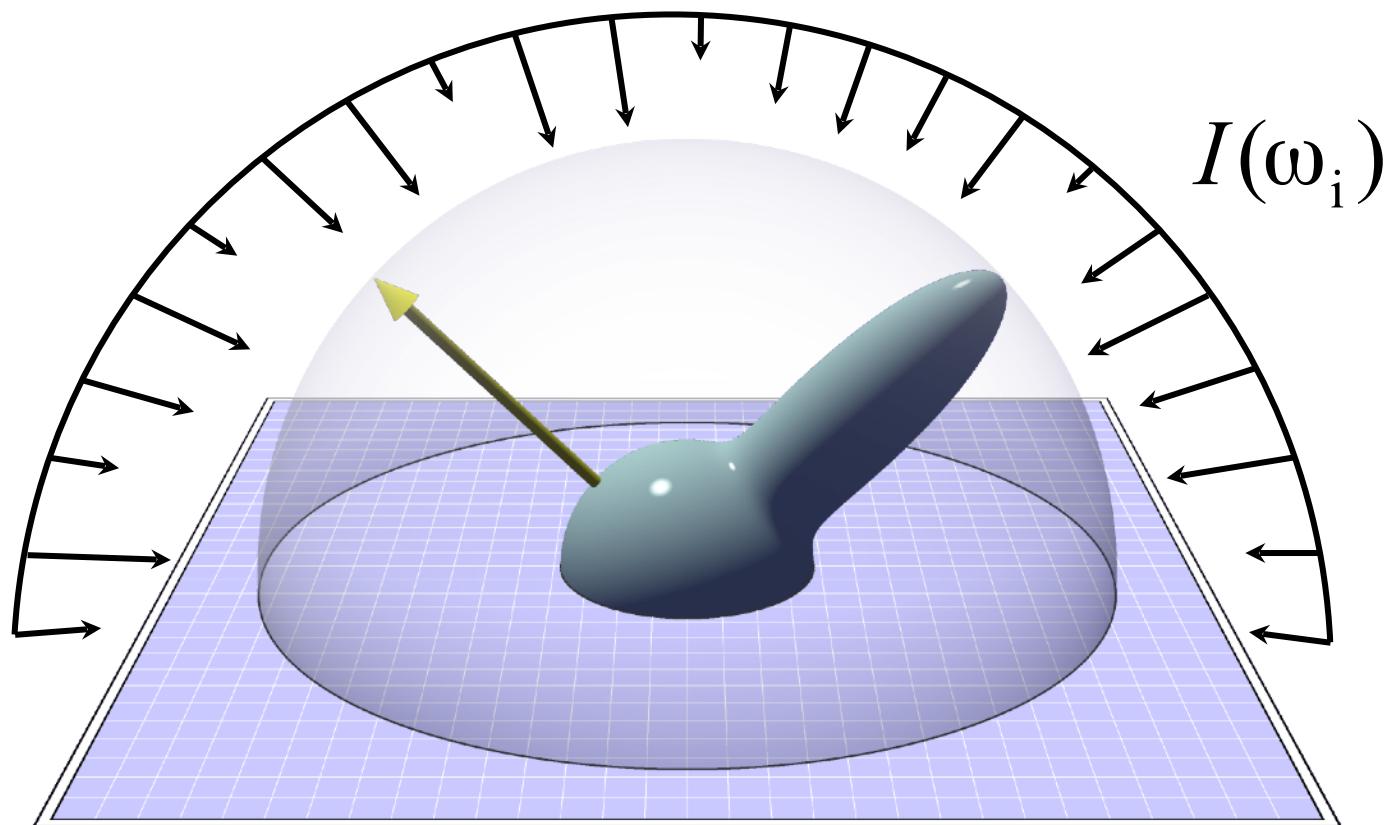
$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)}$$

- ▶ Probability  $p$  allows a better sampling of the domain

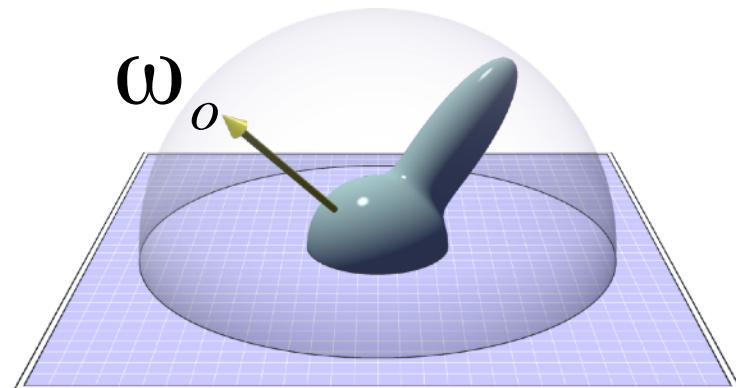
How can we choose  $p$ ?

# Example: glossy reflections

- ▶ Integral over hemisphere of directions
- ▶ BRDF x cosine x incoming light

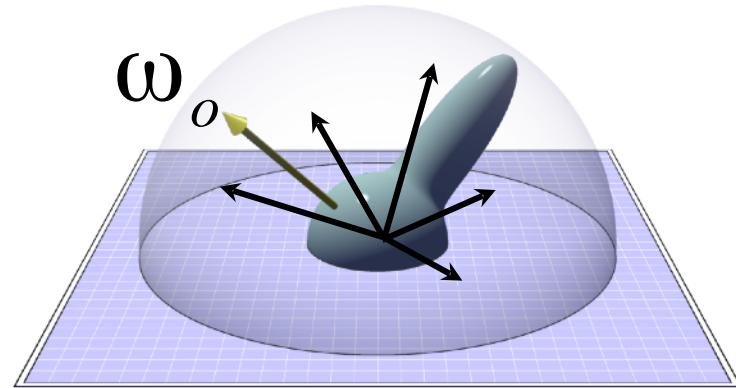


# Sampling a BRDF



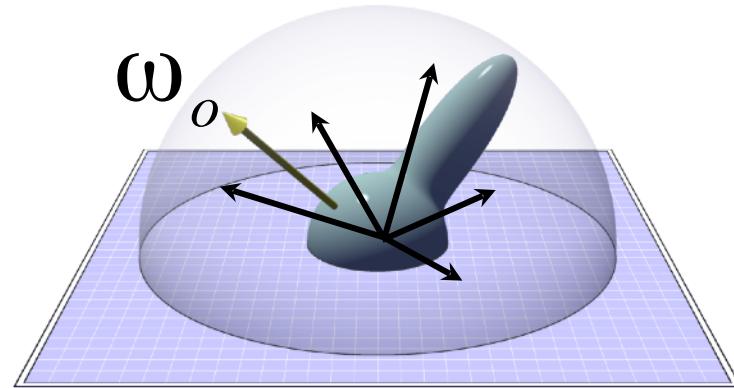
# Sampling a BRDF

$$U(\omega_i)$$

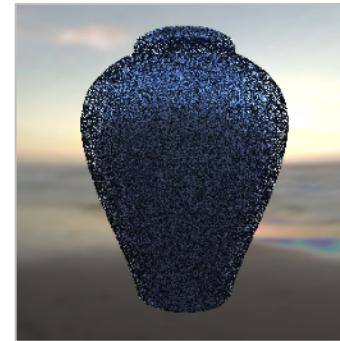


# Sampling a BRDF

$U(\omega_i)$

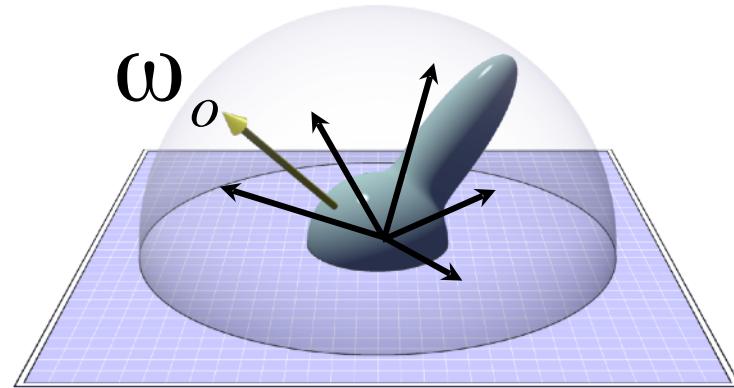


5 Samples/Pixel

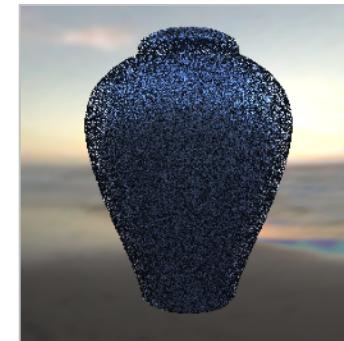


# Sampling a BRDF

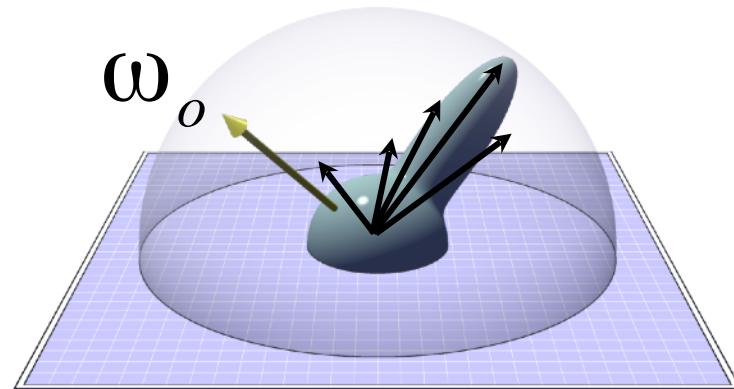
$U(\omega_i)$



5 Samples/Pixel

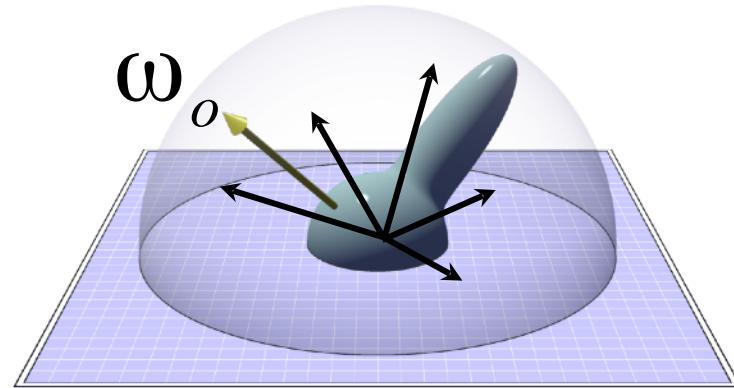


$P(\omega_i)$

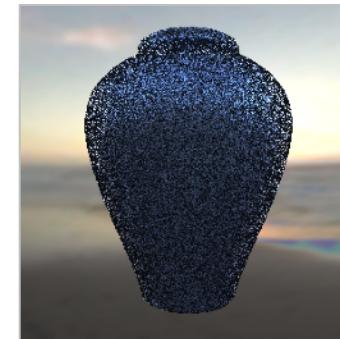


# Sampling a BRDF

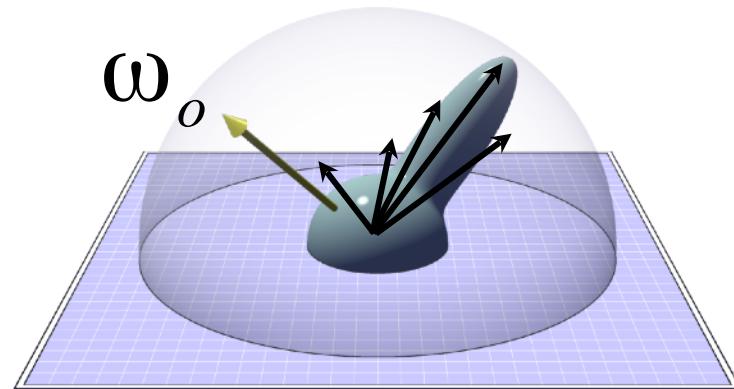
$U(\omega_i)$



5 Samples/Pixel

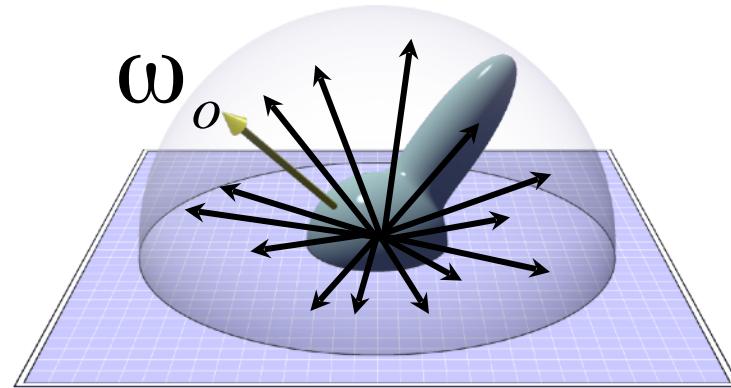


$P(\omega_i)$

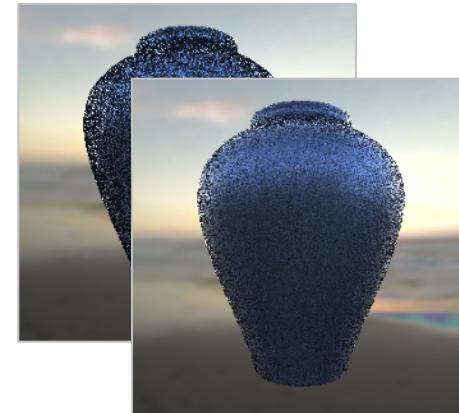


# Sampling a BRDF

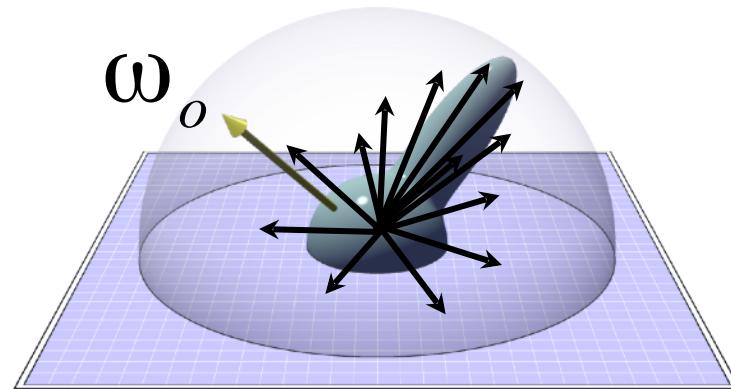
$U(\omega_i)$



25 Samples/Pixel

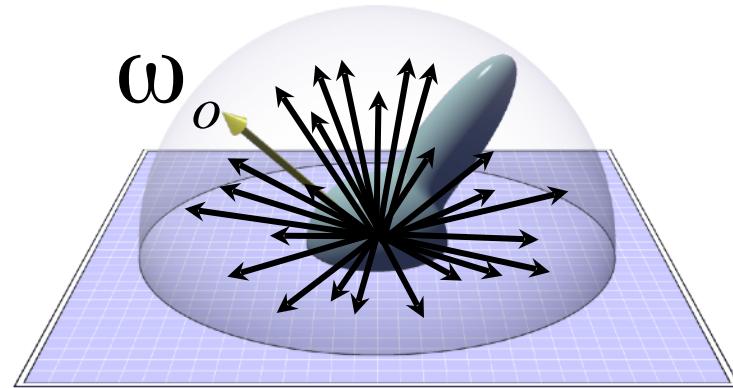


$P(\omega_i)$

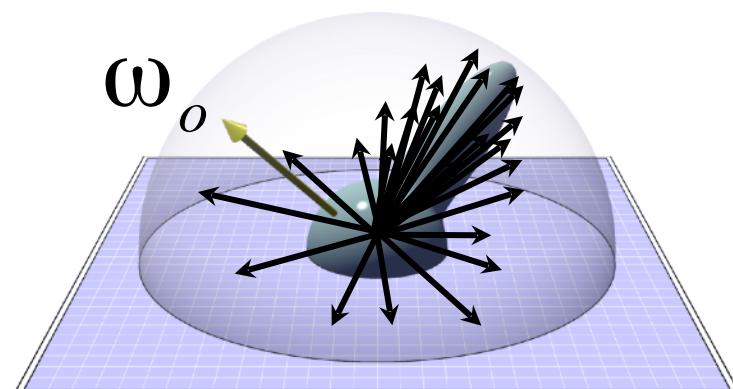


# Sampling a BRDF

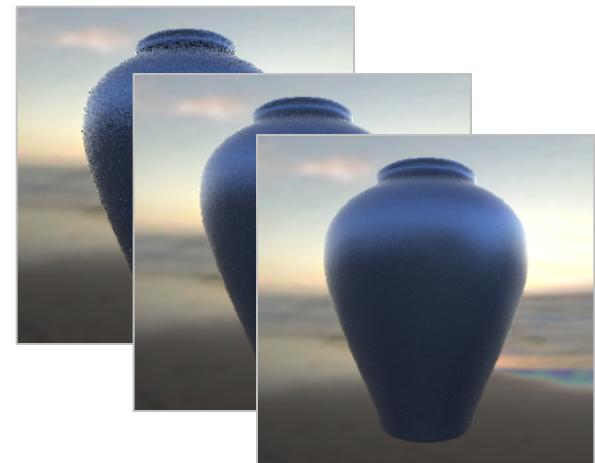
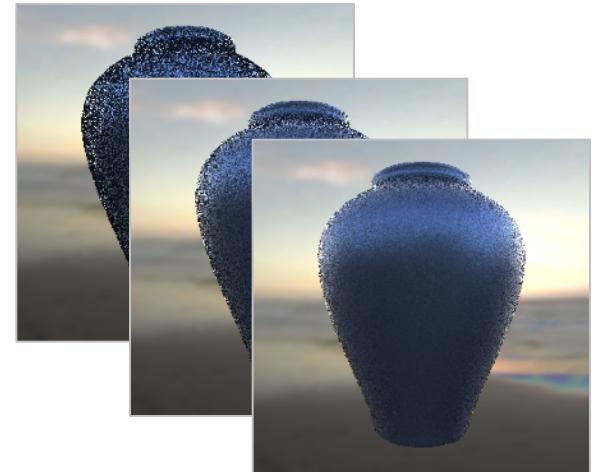
$U(\omega_i)$



$P(\omega_i)$



75 Samples/Pixel



# Importance sampling

$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)}$$

- ▶ Choose  $p$  wisely to reduce variance:
  - $p$  must look like  $f$
  - Doesn't change convergence with  $\sqrt{N}$   
(reduces the constant)
- ▶ Is present in all path-tracers

# Importance sampling in practice



Phong + environment map, on the GPU

100 spp

# Importance sampling in practice

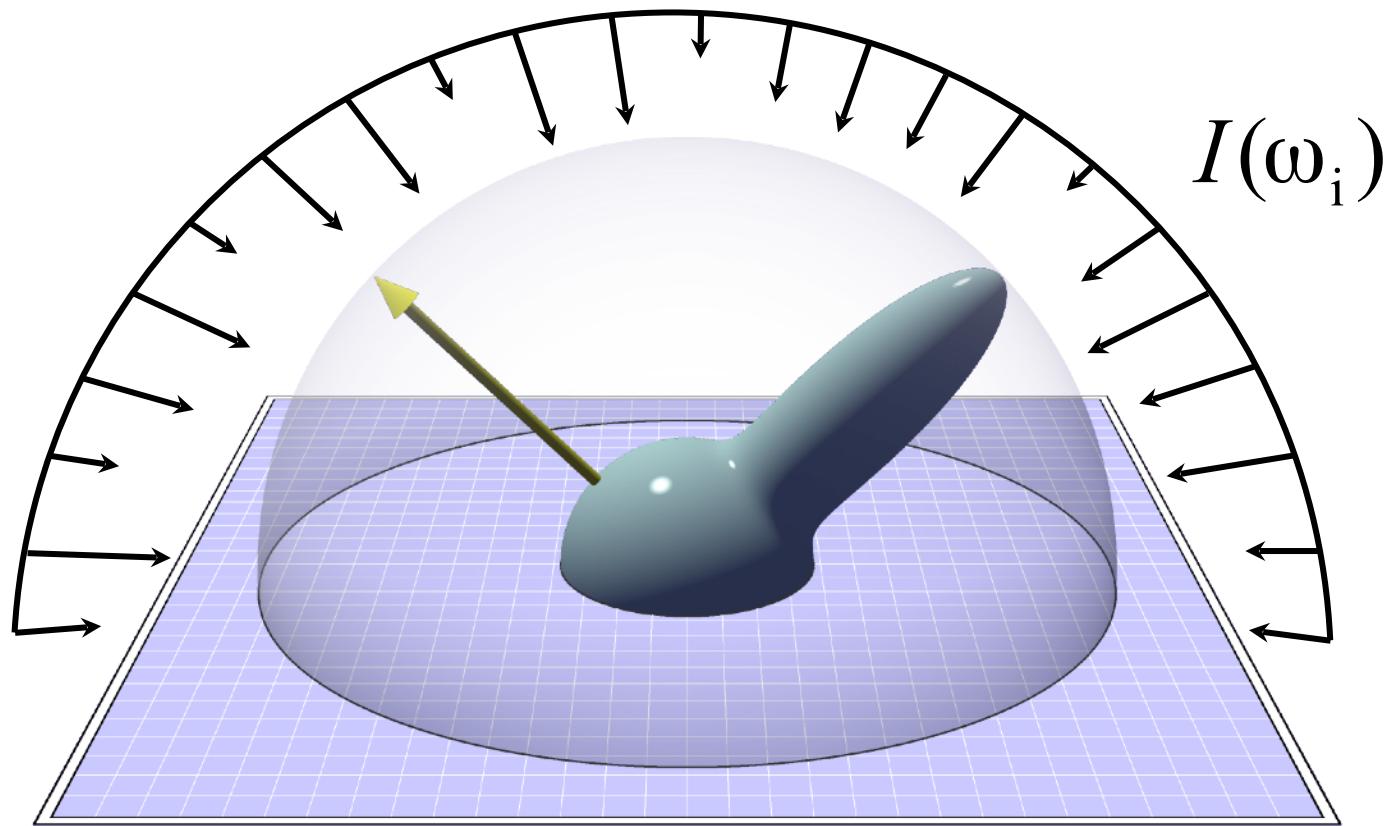


Phong + environment map, on the GPU

1000 spp

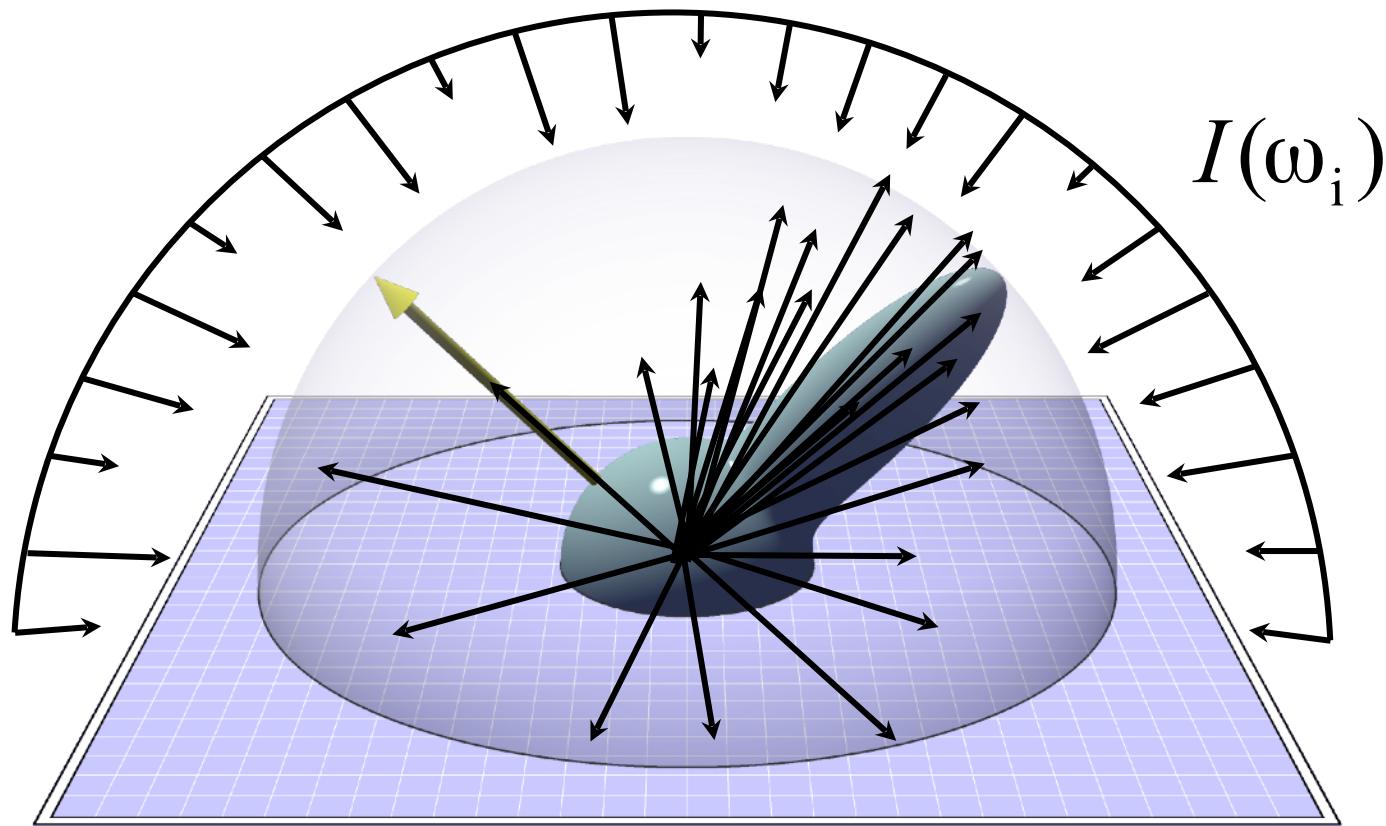
# Importance sampling Phong

- ▶ Input: ray from the eye/camera
- ▶ Goal: outgoing rays, generated with IS



# Importance sampling Phong

- ▶ Input: ray from the eye/camera
- ▶ Goal: outgoing rays, generated with IS



# Importance sampling Phong

- ▶ Actually: generate new normal  $m$  randomly
- ▶ Incoming ray + reflection by  $m$ 
  - Outgoing direction
- ▶ New normal  $m$  :
  - 2 random numbers ( $u_1, u_2$ )
  - $\phi = 2 \pi u_1$
  - $\theta = \arccos(u_2^{(s+1)/2})$
  - $m = \text{vector } (\theta, \phi)$
  - Probability density function for  $m = \cos(\theta)^s$
  - $\theta = \text{angle } (n, m) = \text{angle a ray/specular direction}$

# Importance sampling Phong

- ▶ Contribution from one sample:

$$P += F(x_i)/p(x_i)$$

$$P += \text{envmap(reflected)} * \text{Phong(reflected, n)} / \\ \text{probabilité(reflected)}$$

$$P += \text{envmap(reflected)} * k_s (\cos \alpha)^s / (\cos \alpha)^s$$

$$P += \text{envmap(reflected)} * k_s$$

# Importance sampling Phong

- ▶ Contribution from one sample:

$$P += F(x_i)/p(x_i)$$

$$P += \text{envmap(reflected)} * \text{Phong(reflected, n)} / \\ \text{probabilité(reflected)}$$

$$P += \text{envmap(reflected)} * k_s (\cos \alpha)^s / (\cos \alpha)^s$$

$$P += \text{envmap(reflected)} * k_s$$

Wait, what?

# Importance sampling Phong

- ▶ What if there is a diffuse component?
  - Or several specular components...
- ▶ Select the lobe using Monte-Carlo :
  - Energy for each lobe:  $k_d, k_s, \dots$
  - Random sample  $u_3$  in  $[0,1]$
  - Select sampled lobe:
    - $u_3 < k_d/(k_s + k_d)$  : diffuse, otherwise specular
  - Find sampling direction using this lobe
  - Compute full BRDF (diffuse + specular)
  - Divide by full PDF (diffuse + spéculaire)
  - No simplification anymore

# Importance sampling Phong

- ▶ What if there is an ambient component?
  - Doesn't depend from incoming light...
  - So no need for Monte-Carlo sampling

# Monte-Carlo + Phong on the GPU

- ▶ Function random() doesn't exist in GLSL
- ▶ Pseudo-random function using sin:

```
highp float rand(vec2 co) {  
    highp float a = 12.9898;  
    highp float b = 78.233;  
    highp float c = 43758.5453;  
    highp float dt = dot(co.xy, vec2(a,b));  
    highp float sn = mod(dt, 3.14);  
    return fract(sin(sn) * c);  
}
```

Source: [byteblacksmith.com](http://byteblacksmith.com)

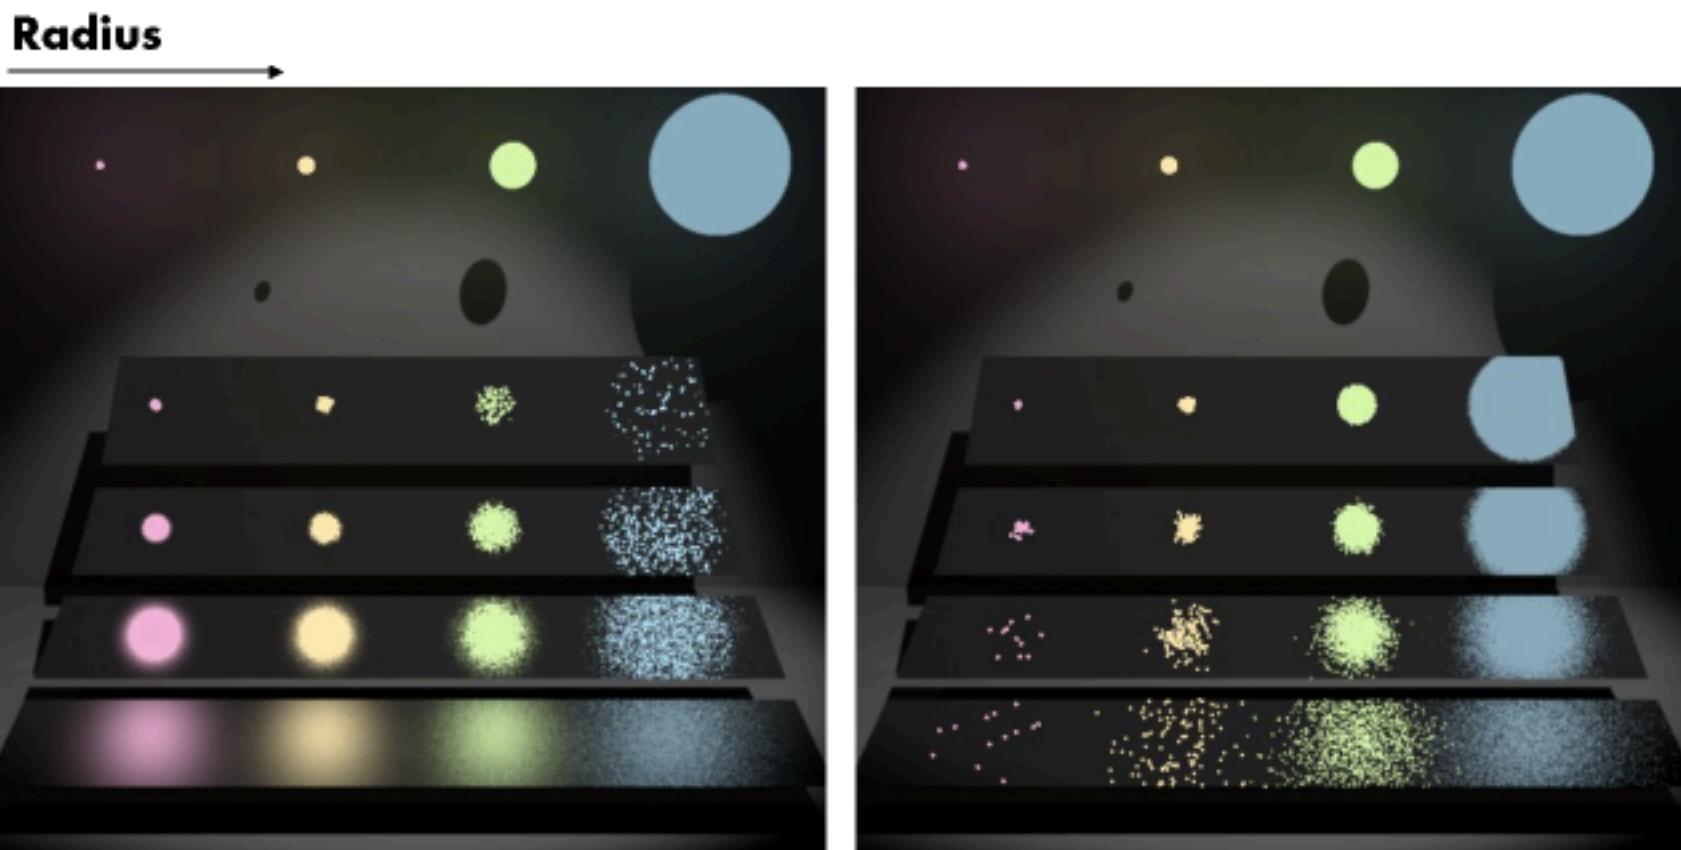
- ▶ Seed (co) :
  - Essential
  - Must be different for  $u_1$  and  $u_2$
  - Use your imagination

# Monte-Carlo + Phong on the GPU

- ▶ New normal  $m$  : coordinates relative to  $n$ 
  - Local coordinate system
  - Must move to global coordinates
  - Need a coordinate transform matrix
  - Vertical axis ( $z$ ) =  $n$  (point normal)
  - Other axes ( $x,y$ ) = anything
  - $M = (t,b,n)$

# Multiple Importance Sampling

Reflection of a circular light source by a rough surface



Sampling the light source

$$\int f(x)g(x)dx$$

Sampling the BRDF

# Multiple Importance Sampling

---

## Two sampling techniques

$$\begin{array}{ll} X_{1,i} \sim p_1(x) & X_{2,i} \sim p_2(x) \\ Y_{1,i} = \frac{f(X_{1,i})}{p_1(X_{1,i})} & Y_{2,i} = \frac{f(X_{2,i})}{p_2(X_{2,i})} \end{array}$$

## Form weighted combination of samples

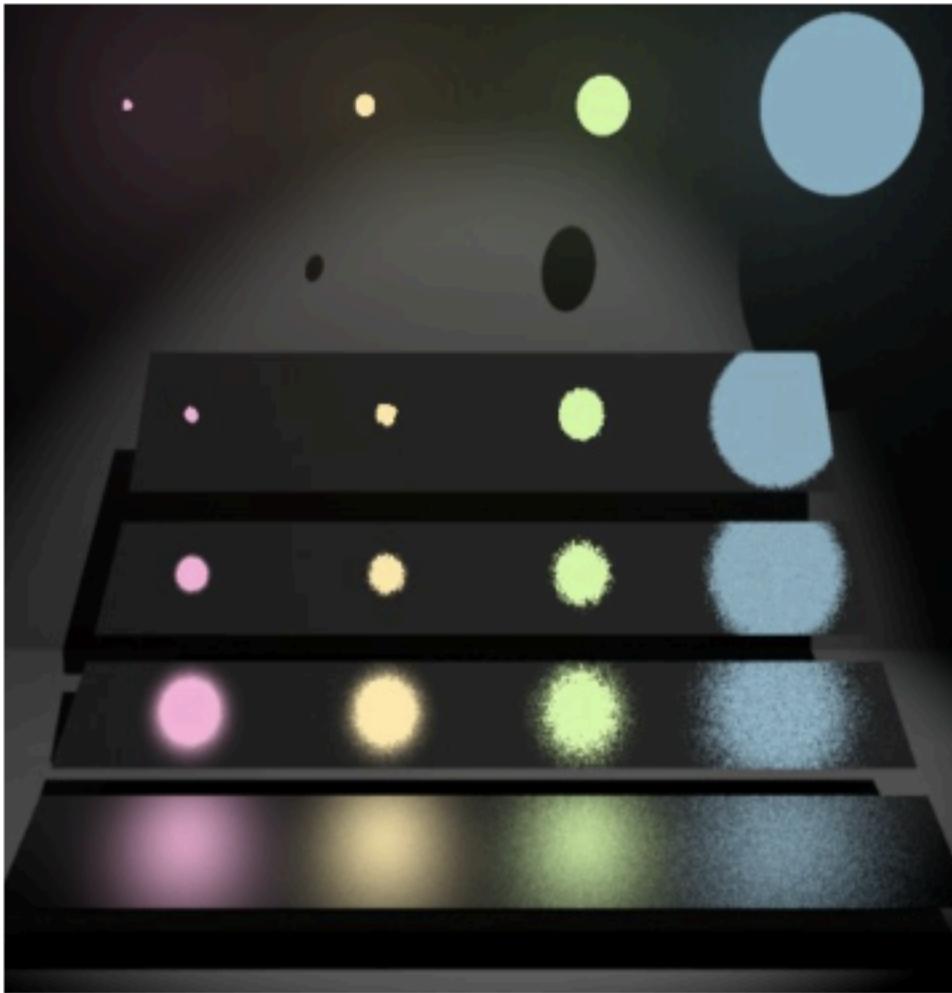
$$Y_i = w_1 Y_{1,i} + w_2 Y_{2,i}$$

## The balance heuristic

$$w_i(x) = \frac{p_i(x)}{p_1(x) + p_2(x)} \Rightarrow p(x) = w_1(x)p_1(x) + w_2(x)p_2(x)$$

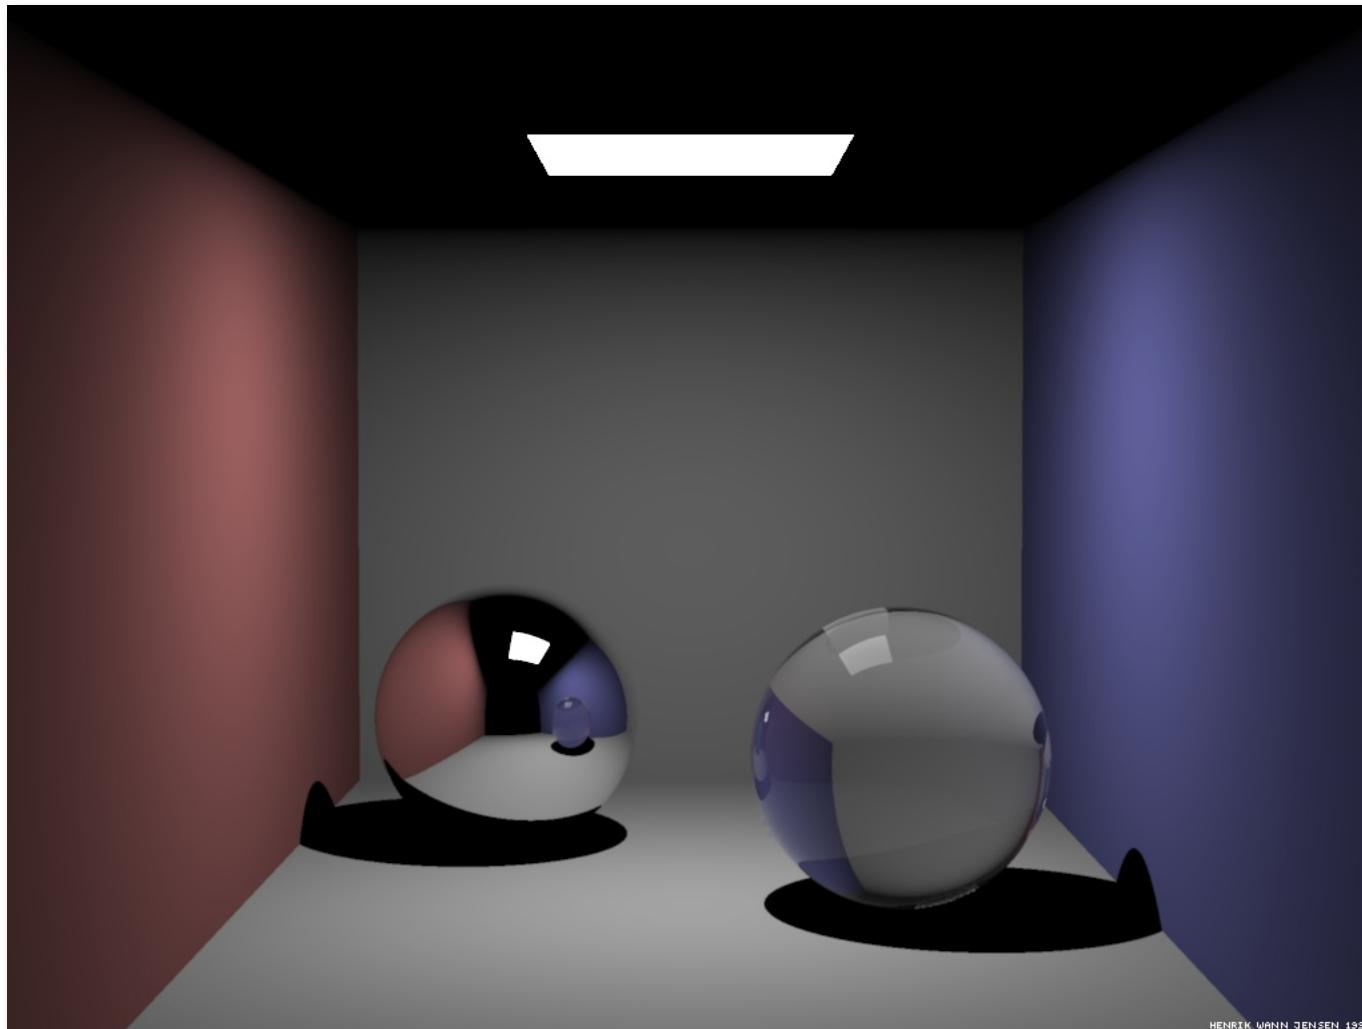
# Multiple Importance Sampling

---



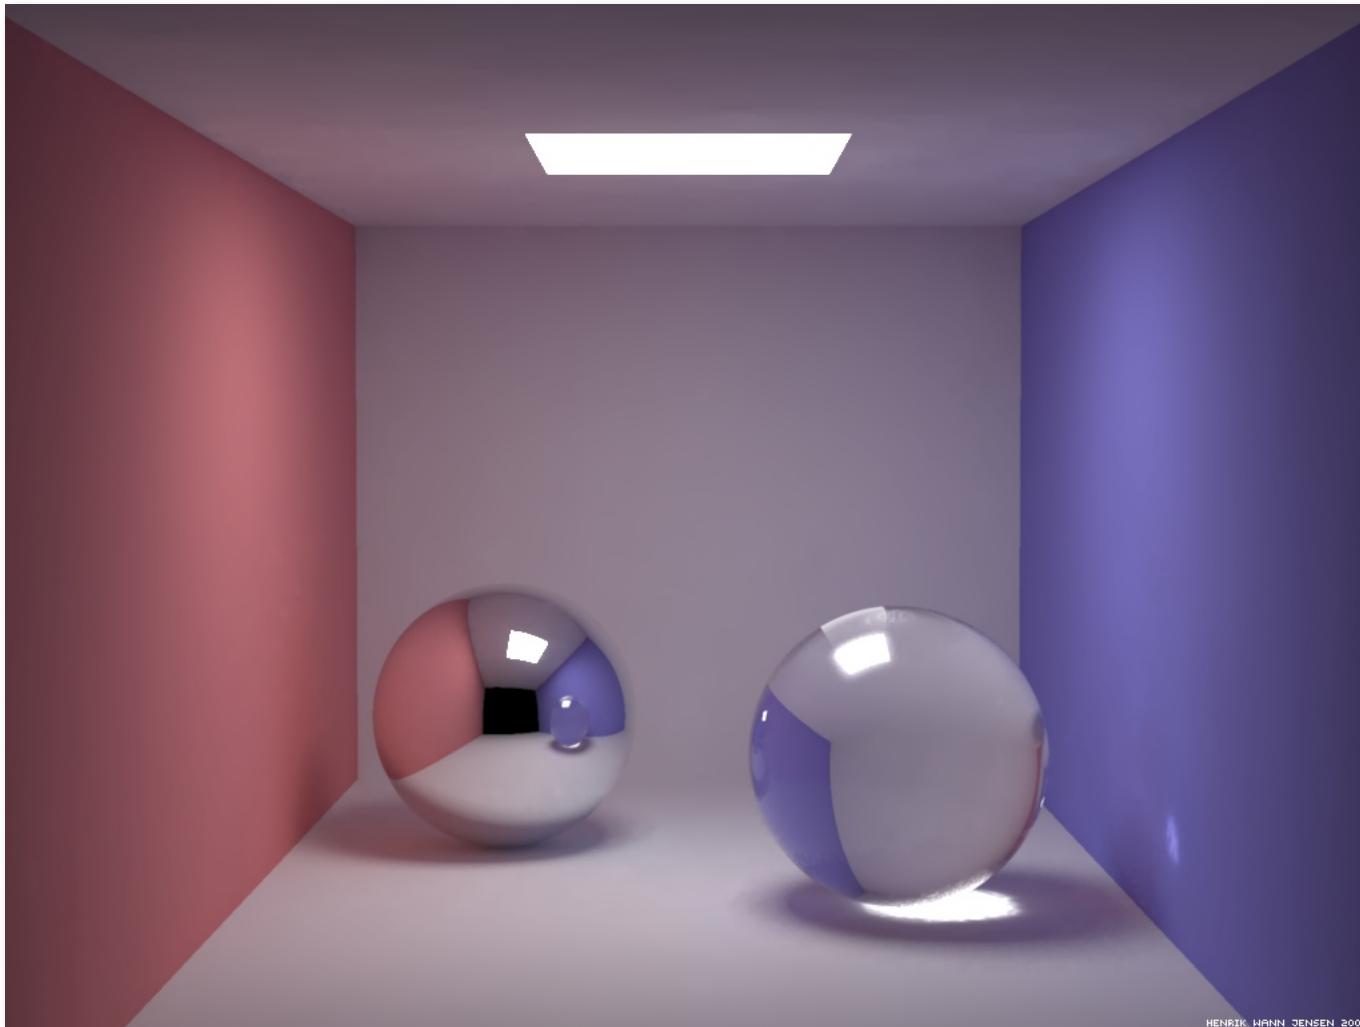
**Source: Veach and Guibas**

# Direct lighting



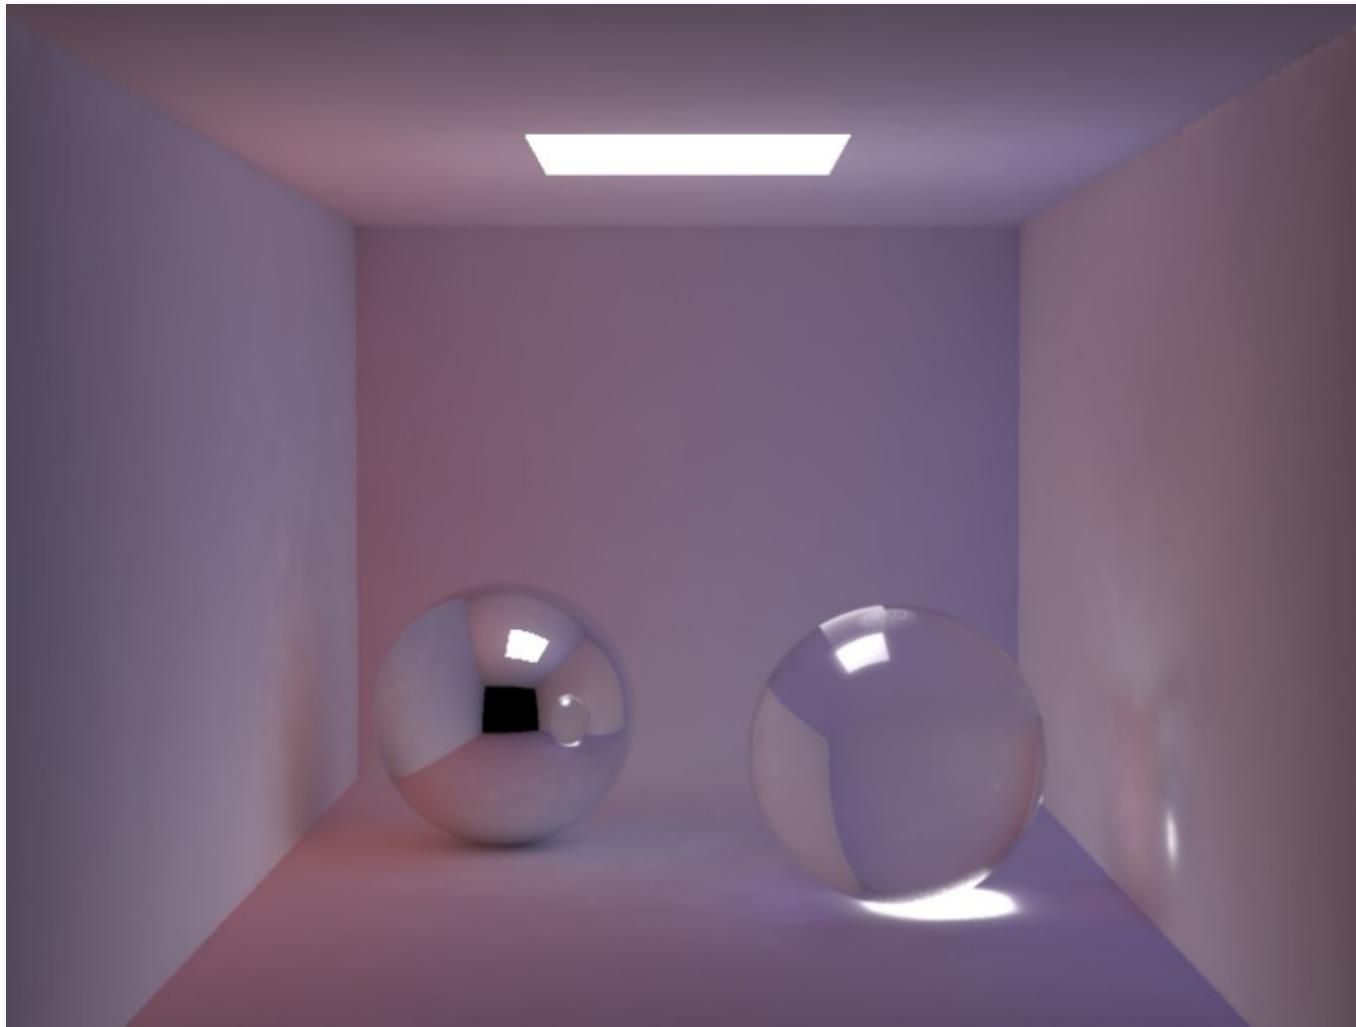
HENRIK WANN JENSEN 1999

# Global illumination



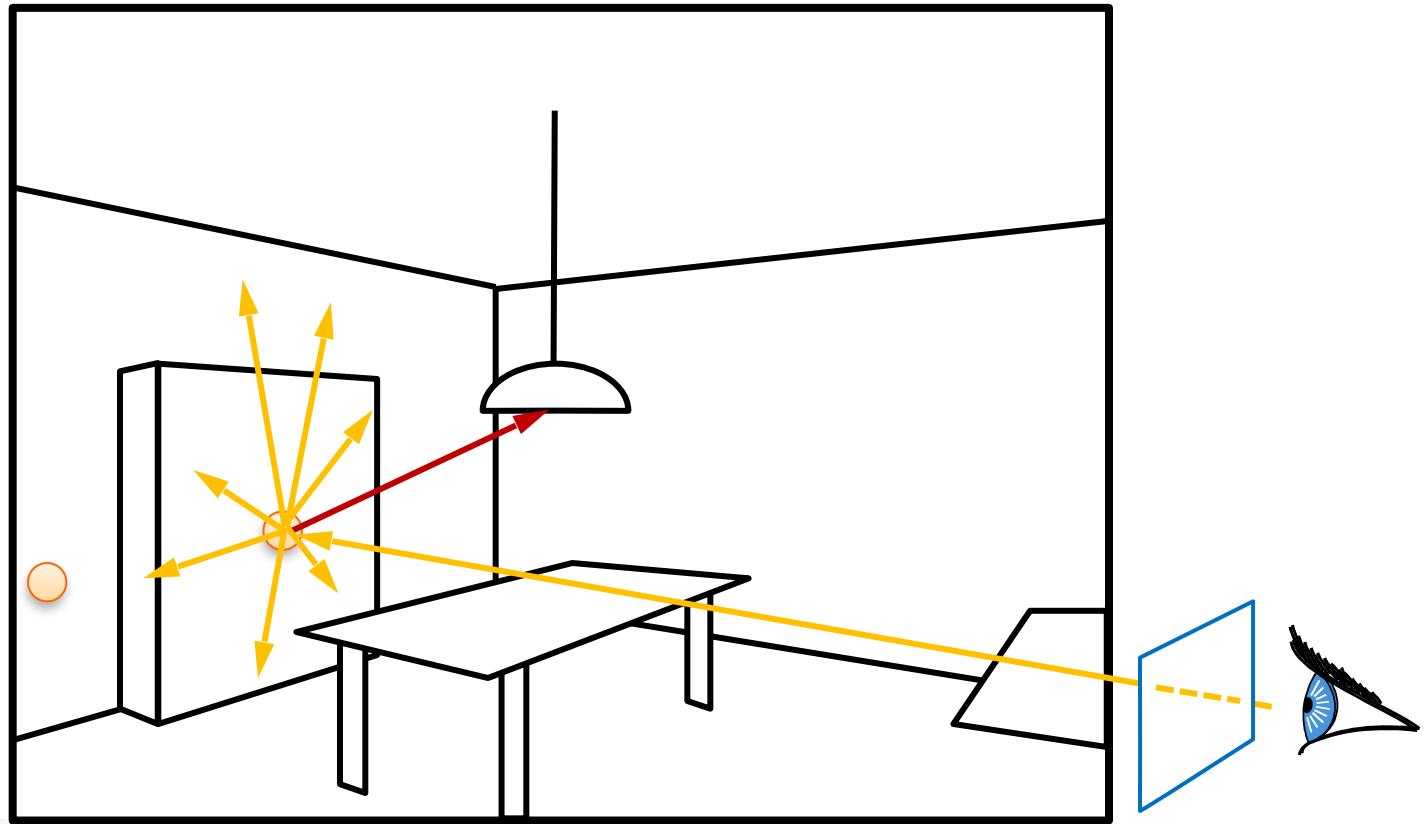
HENRIK HANN JENSEN 2000

# Indirect lighting only



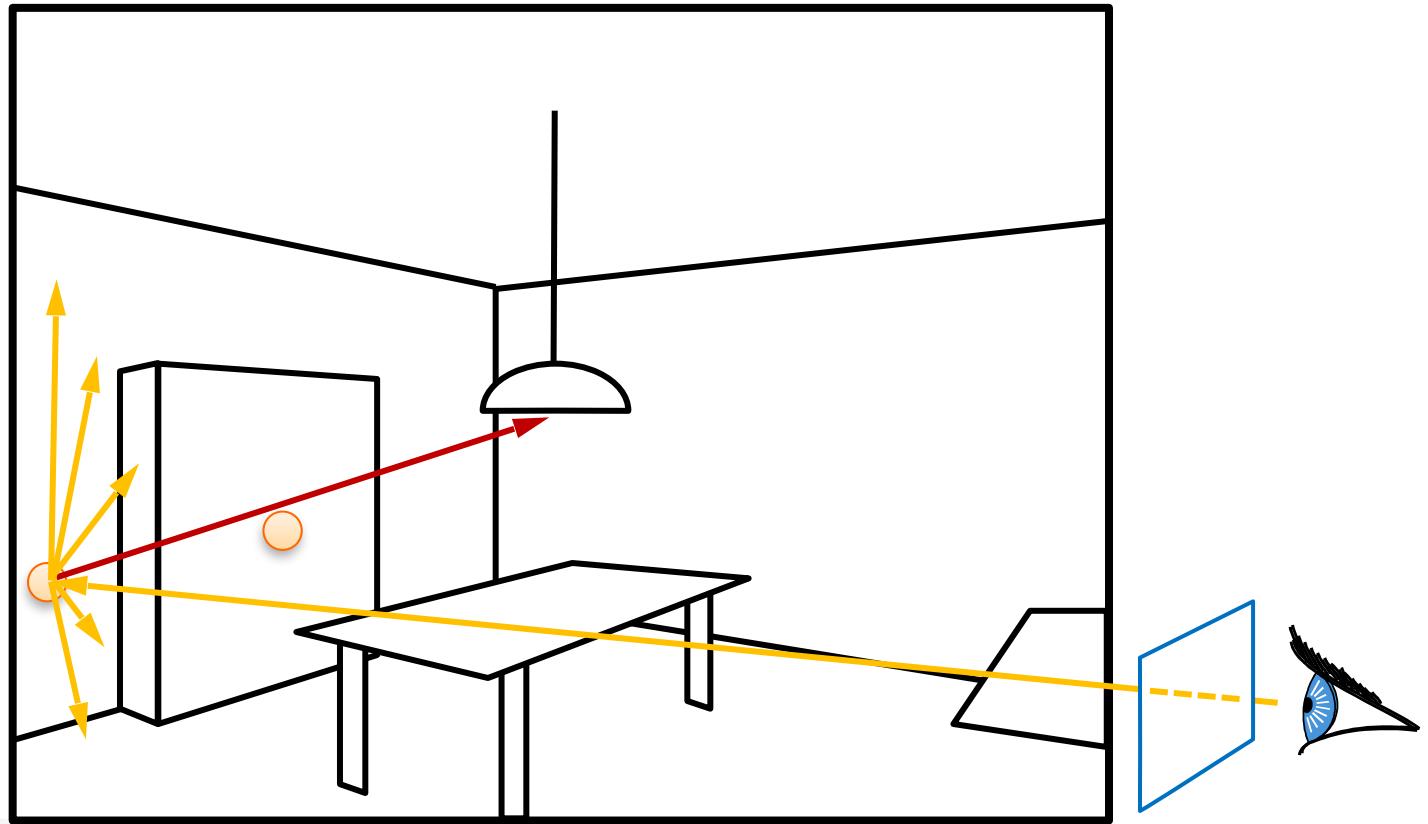
# Irradiance cache

- Indirect lighting changes slowly in space



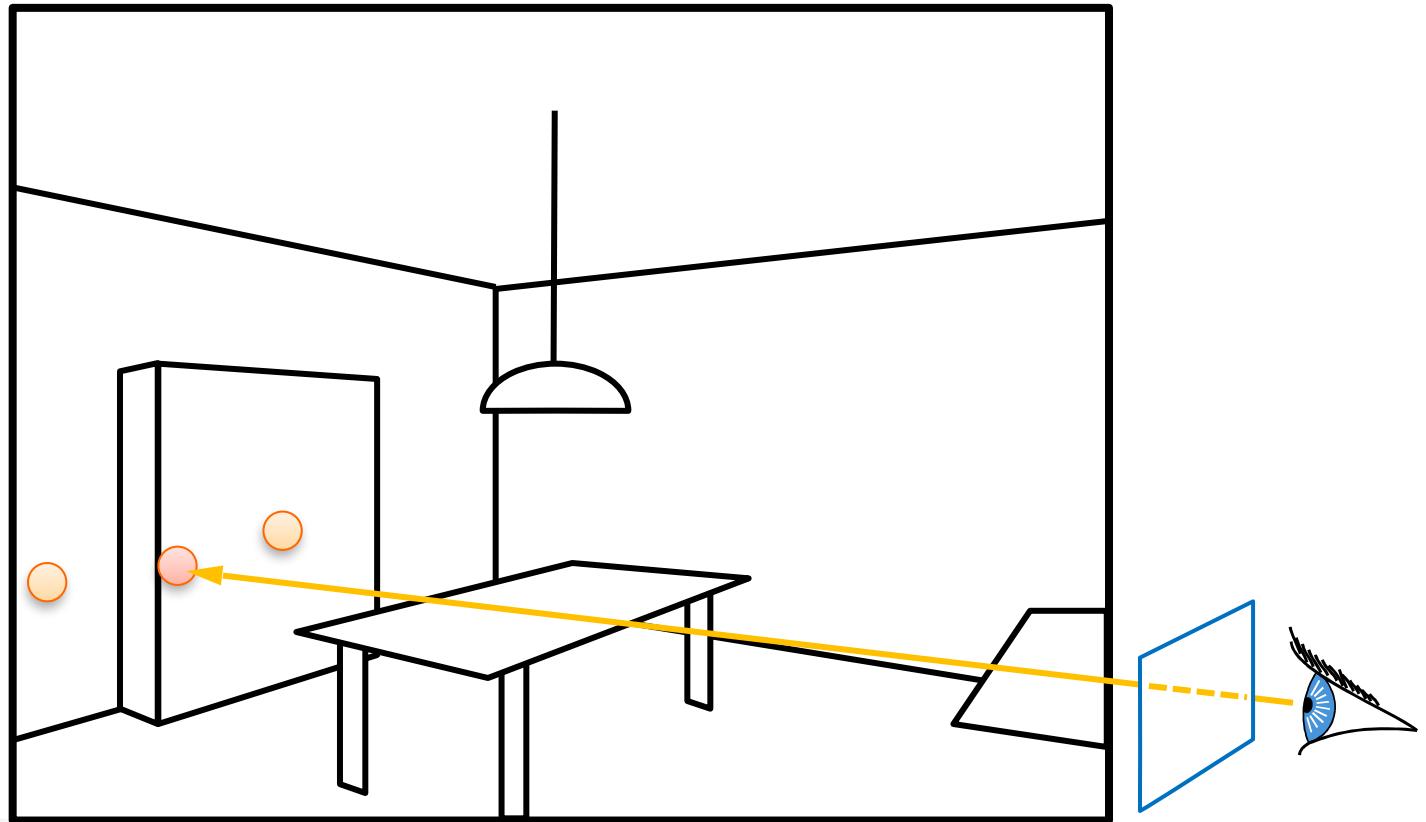
# Irradiance cache

- ▶ Indirect lighting changes slowly in space
- ▶ Particularly with diffuse surfaces



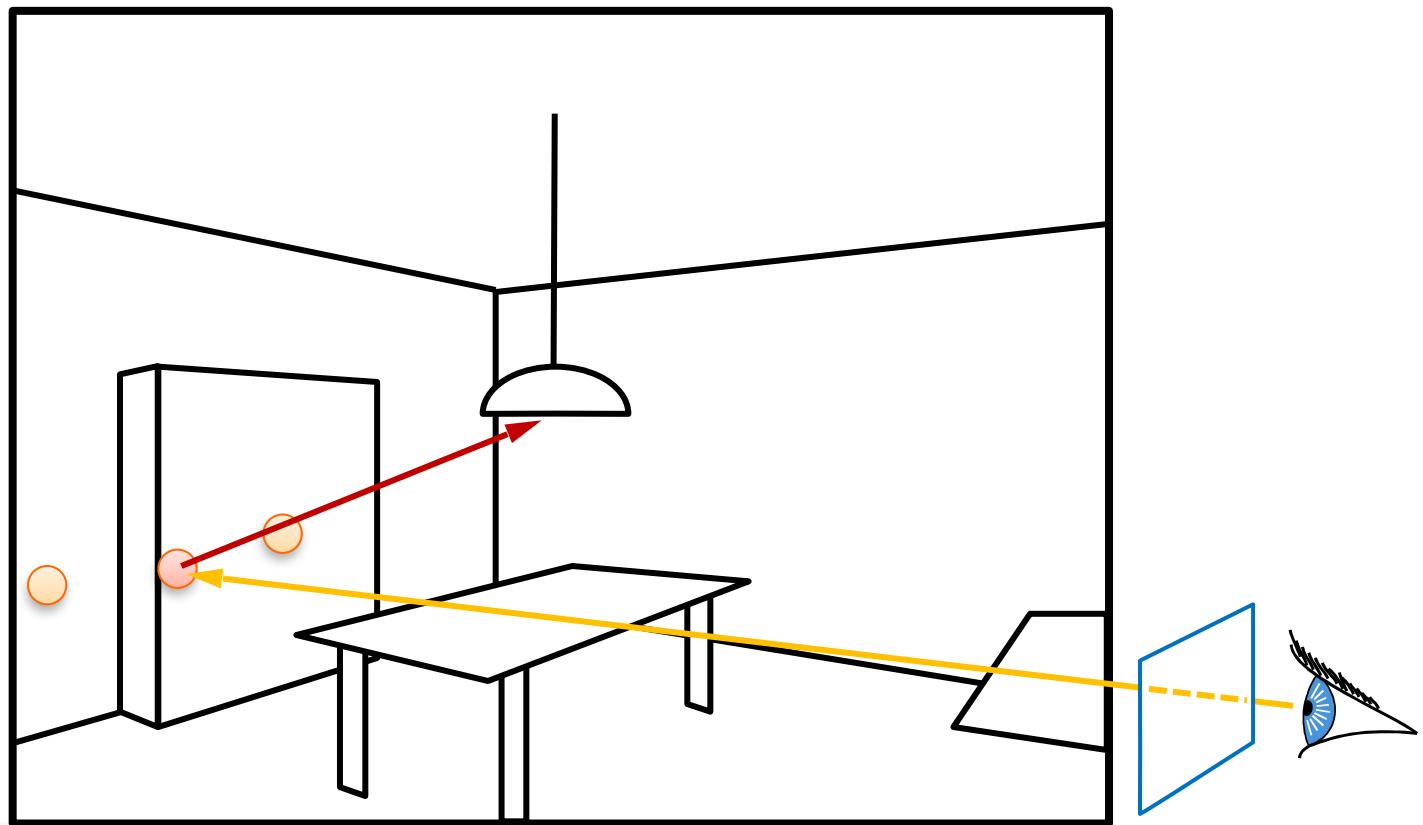
# Irradiance cache

- ▶ Indirect lighting changes slowly in space
- ▶ Interpolate between neighboring values



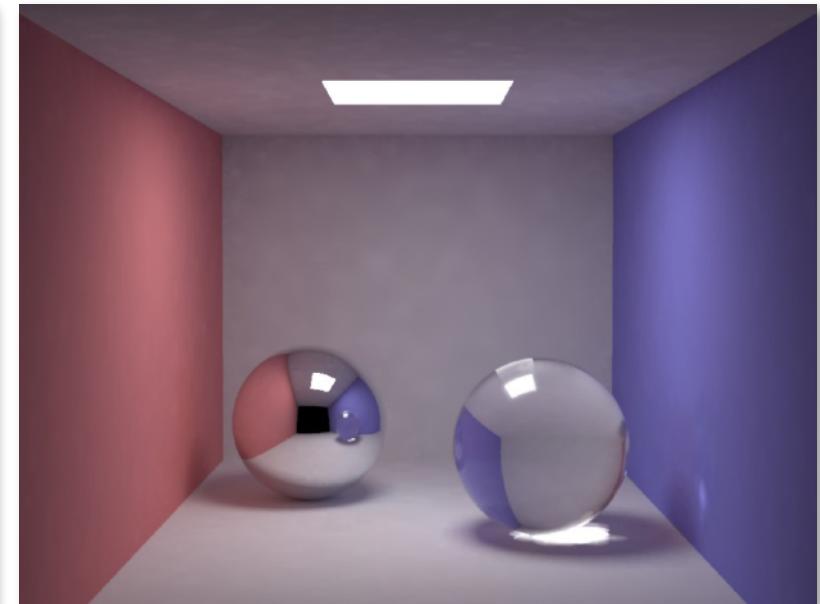
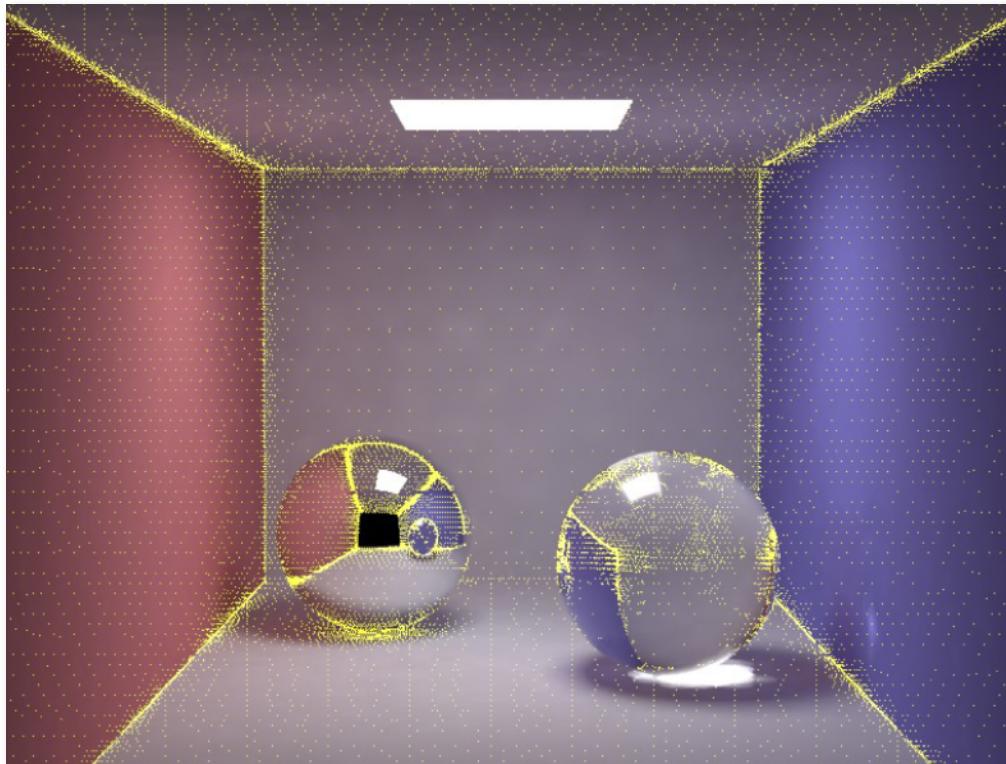
# Irradiance cache

- ▶ Indirect lighting changes slowly in space
- ▶ Interpolate between cached values
- ▶ But full computation for direct lighting

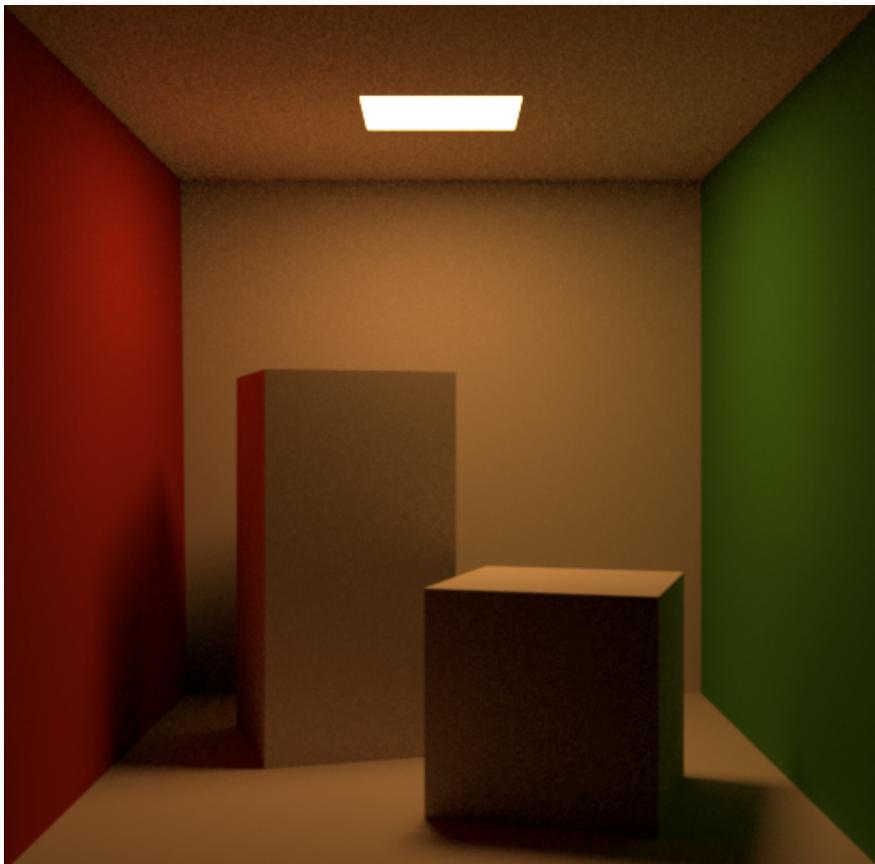


# Irradiance cache

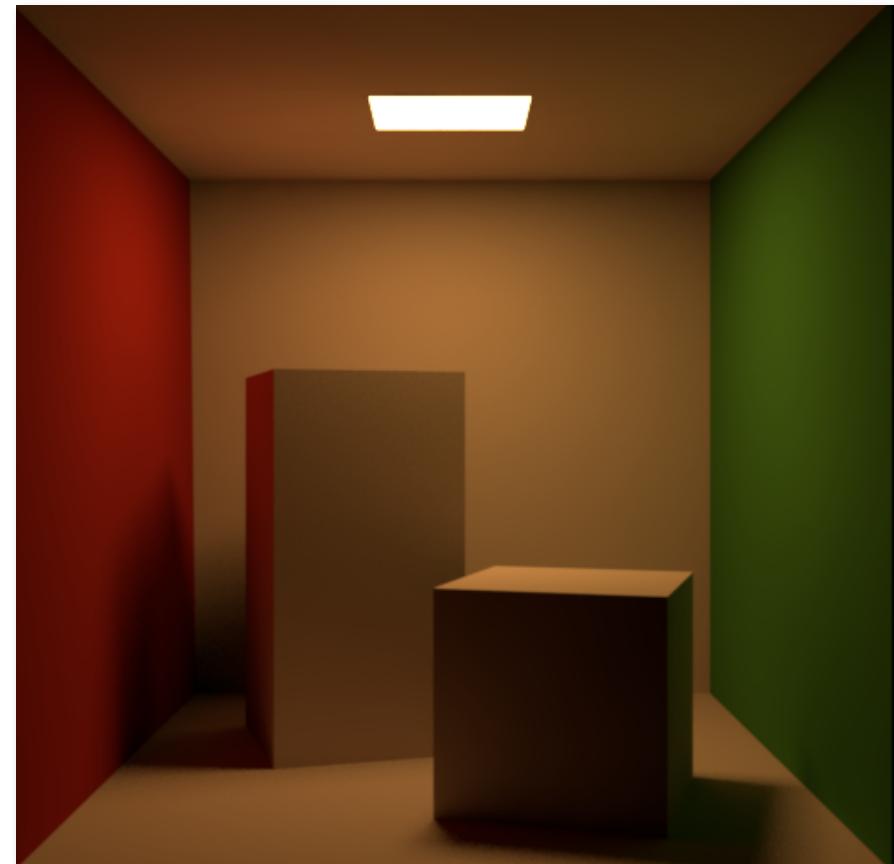
- ▶ Yellow dots: computation of indirect lighting



# Irradiance cache



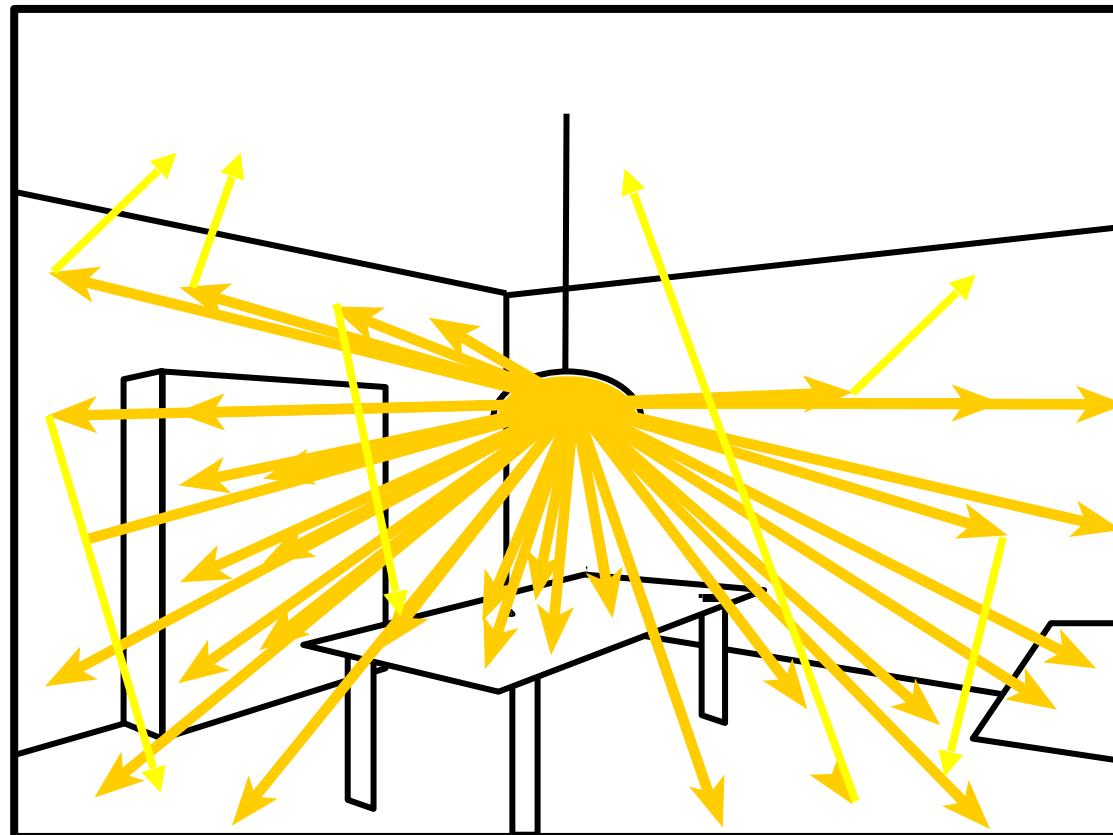
Path Tracing



Path tracing + Irradiance cache

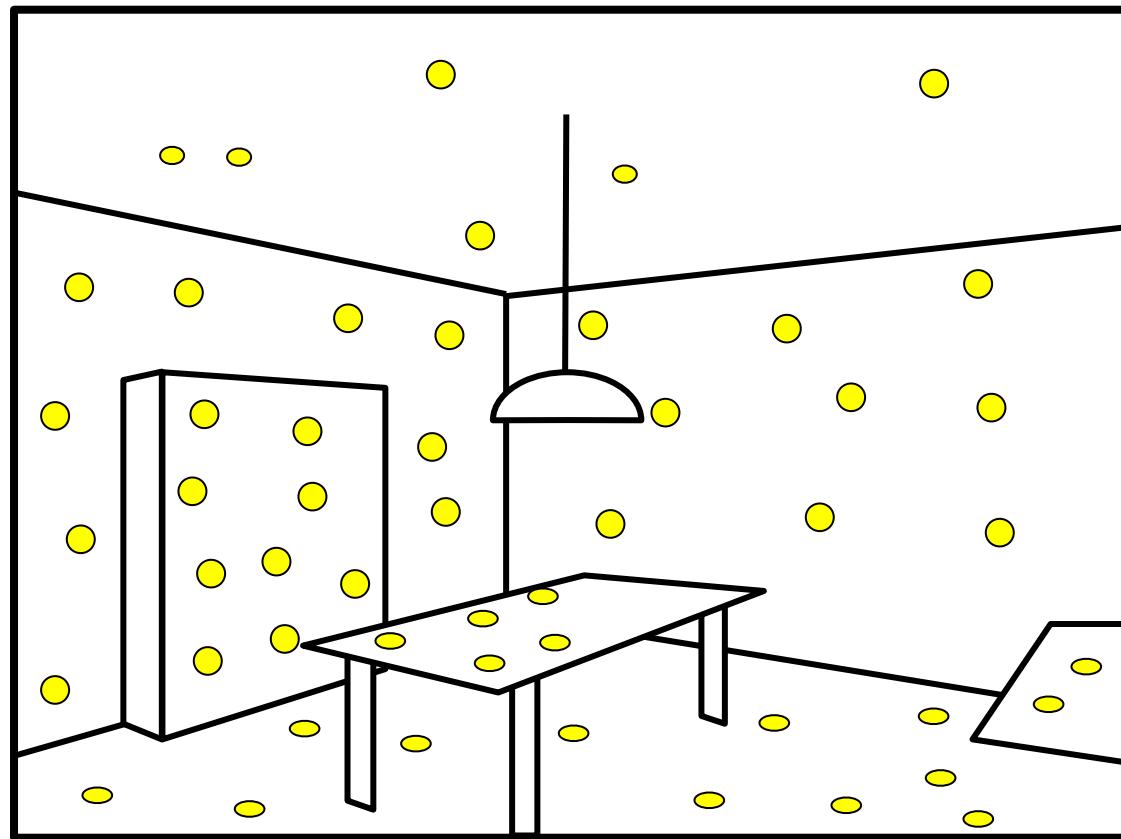
# Photon mapping

- ▶ Pre-computation: throw rays from the light sources



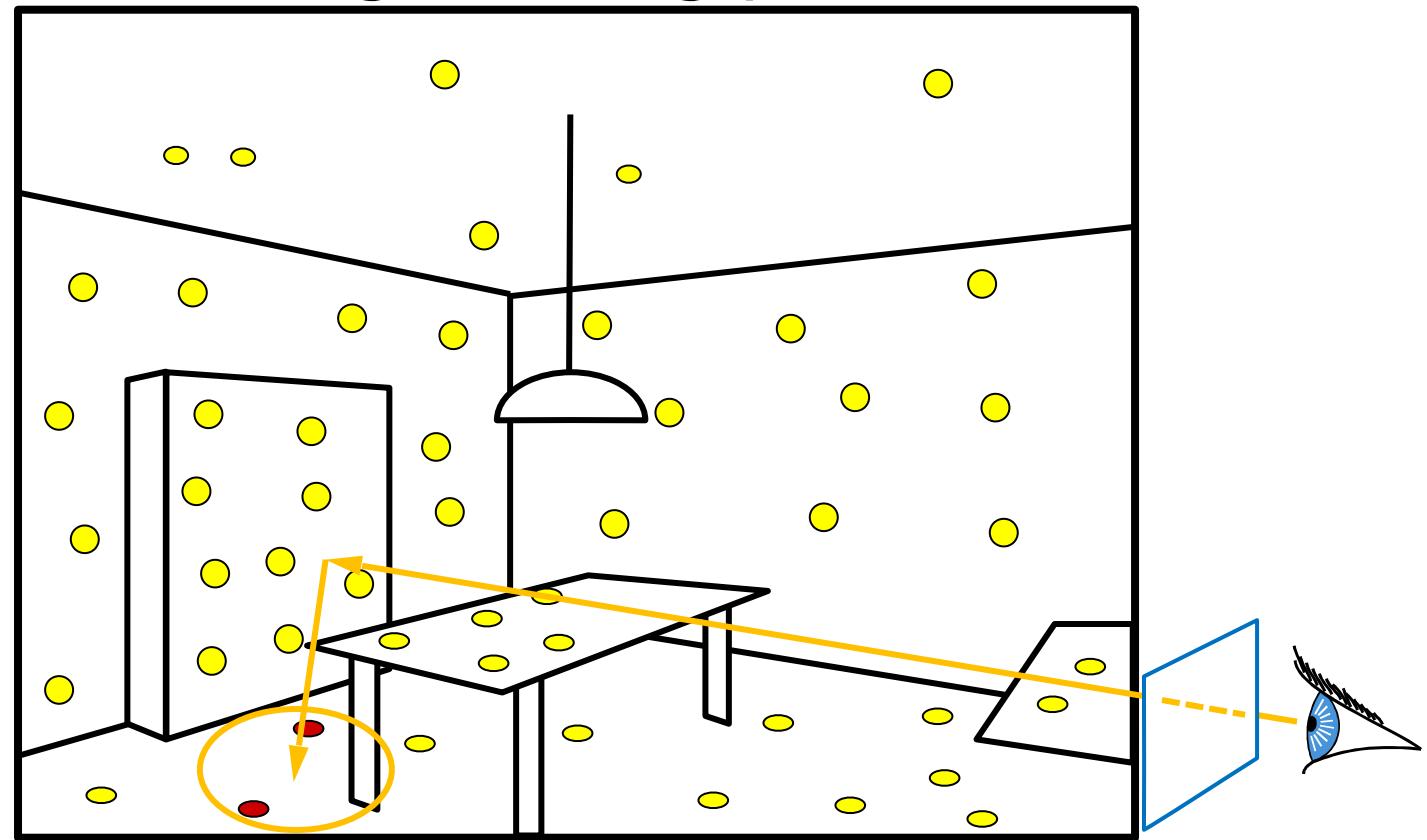
# Photon mapping

- ▶ Store photons (position + intensity + direction) on the geometry or inside a data structure

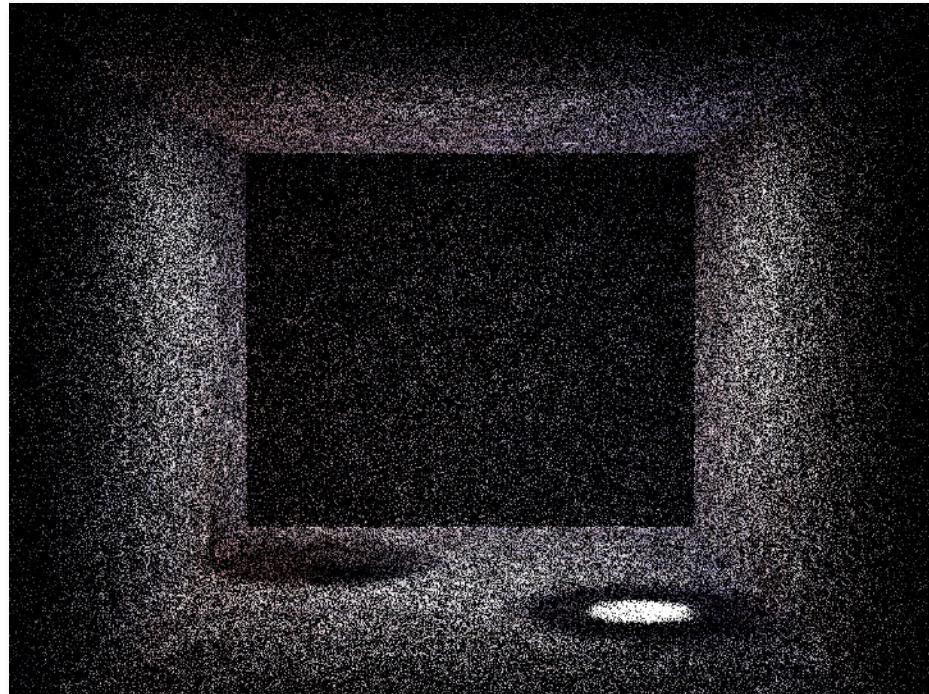


# Photon mapping - rendering

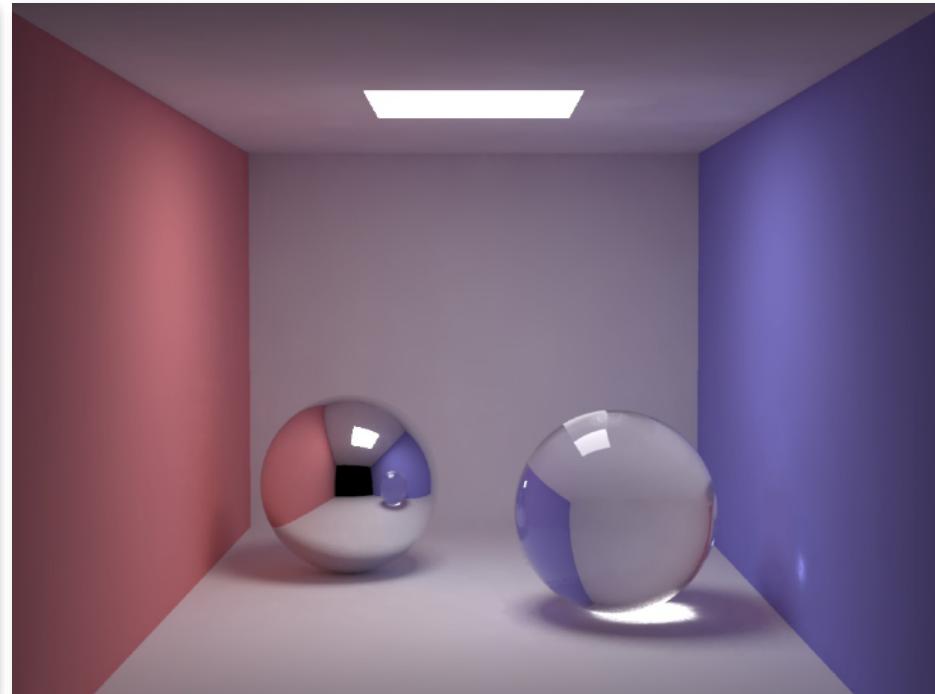
- ▶ Throw primary rays
- ▶ Radiance for secondary rays by gathering radiance from neighbouring photons



# Results



Photon map



Final rendering

# Results

- ▶ Jensen (1996)
  - Direct visualization of the photon map: 6min



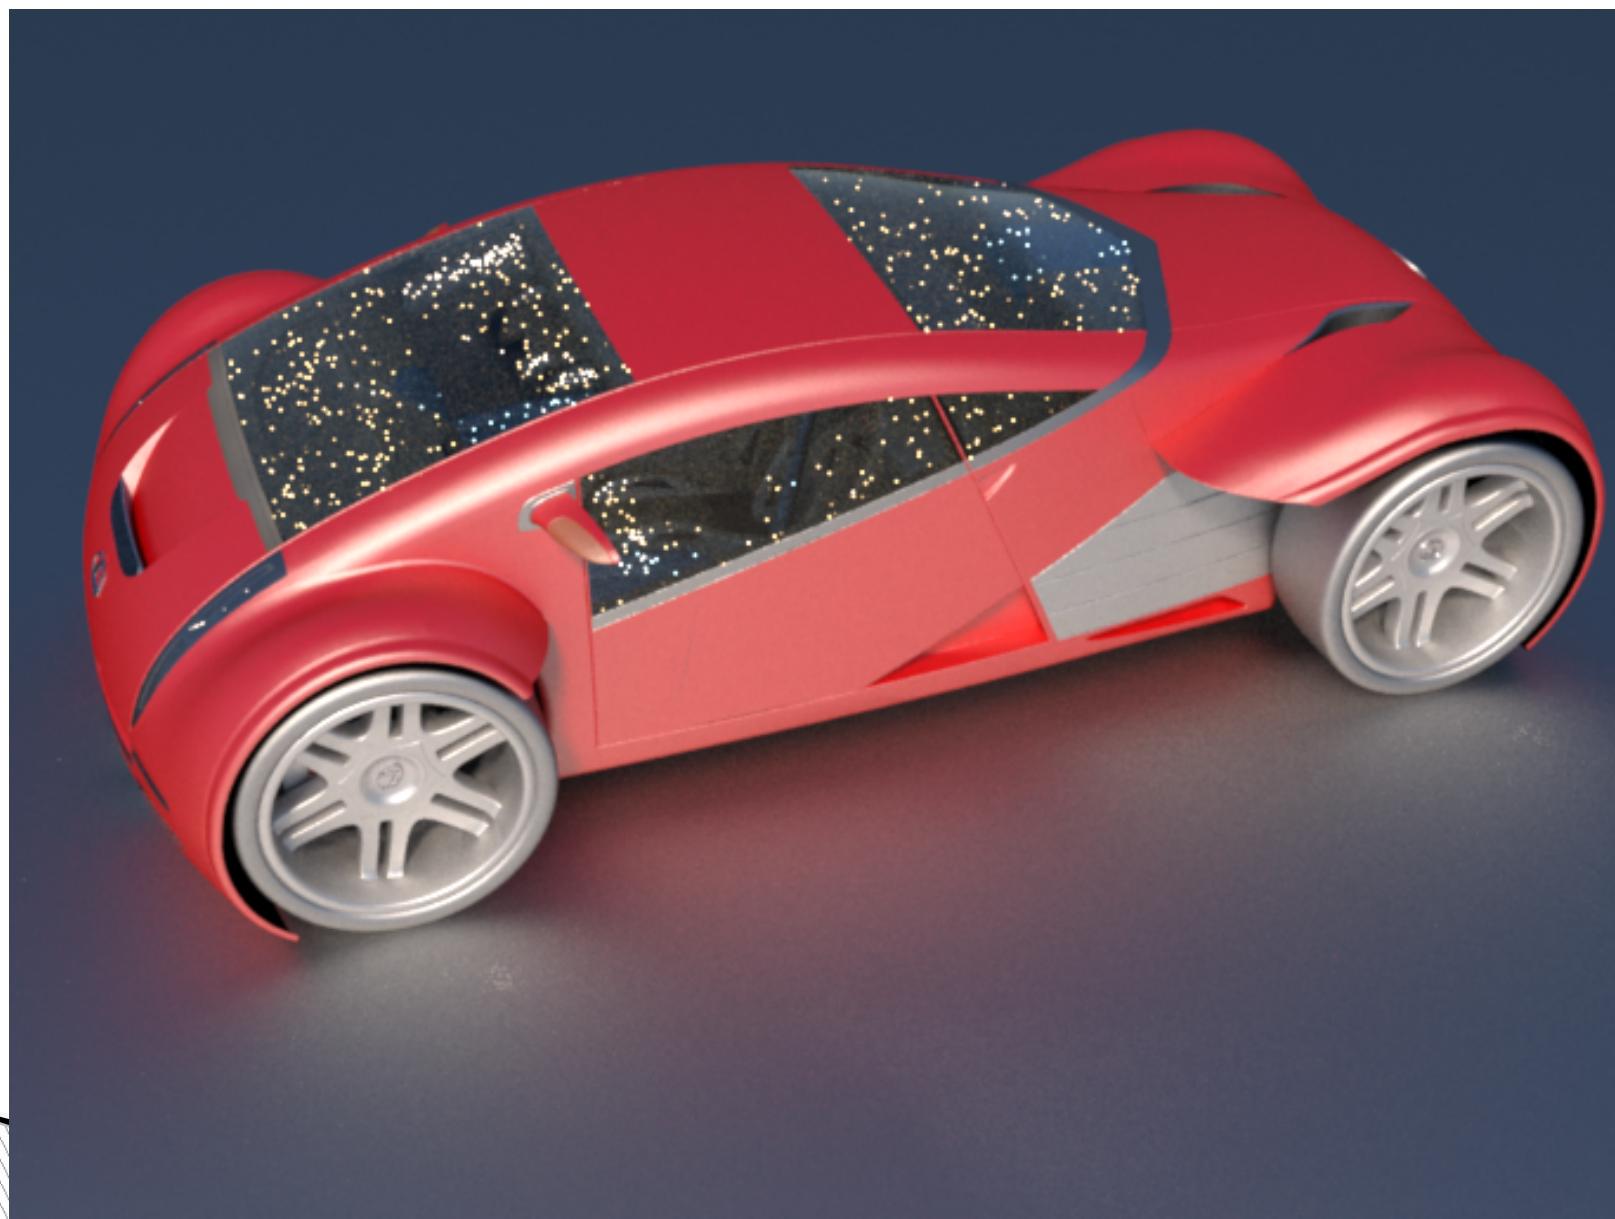
HENRIK WANN JENSEN 1996

# Results

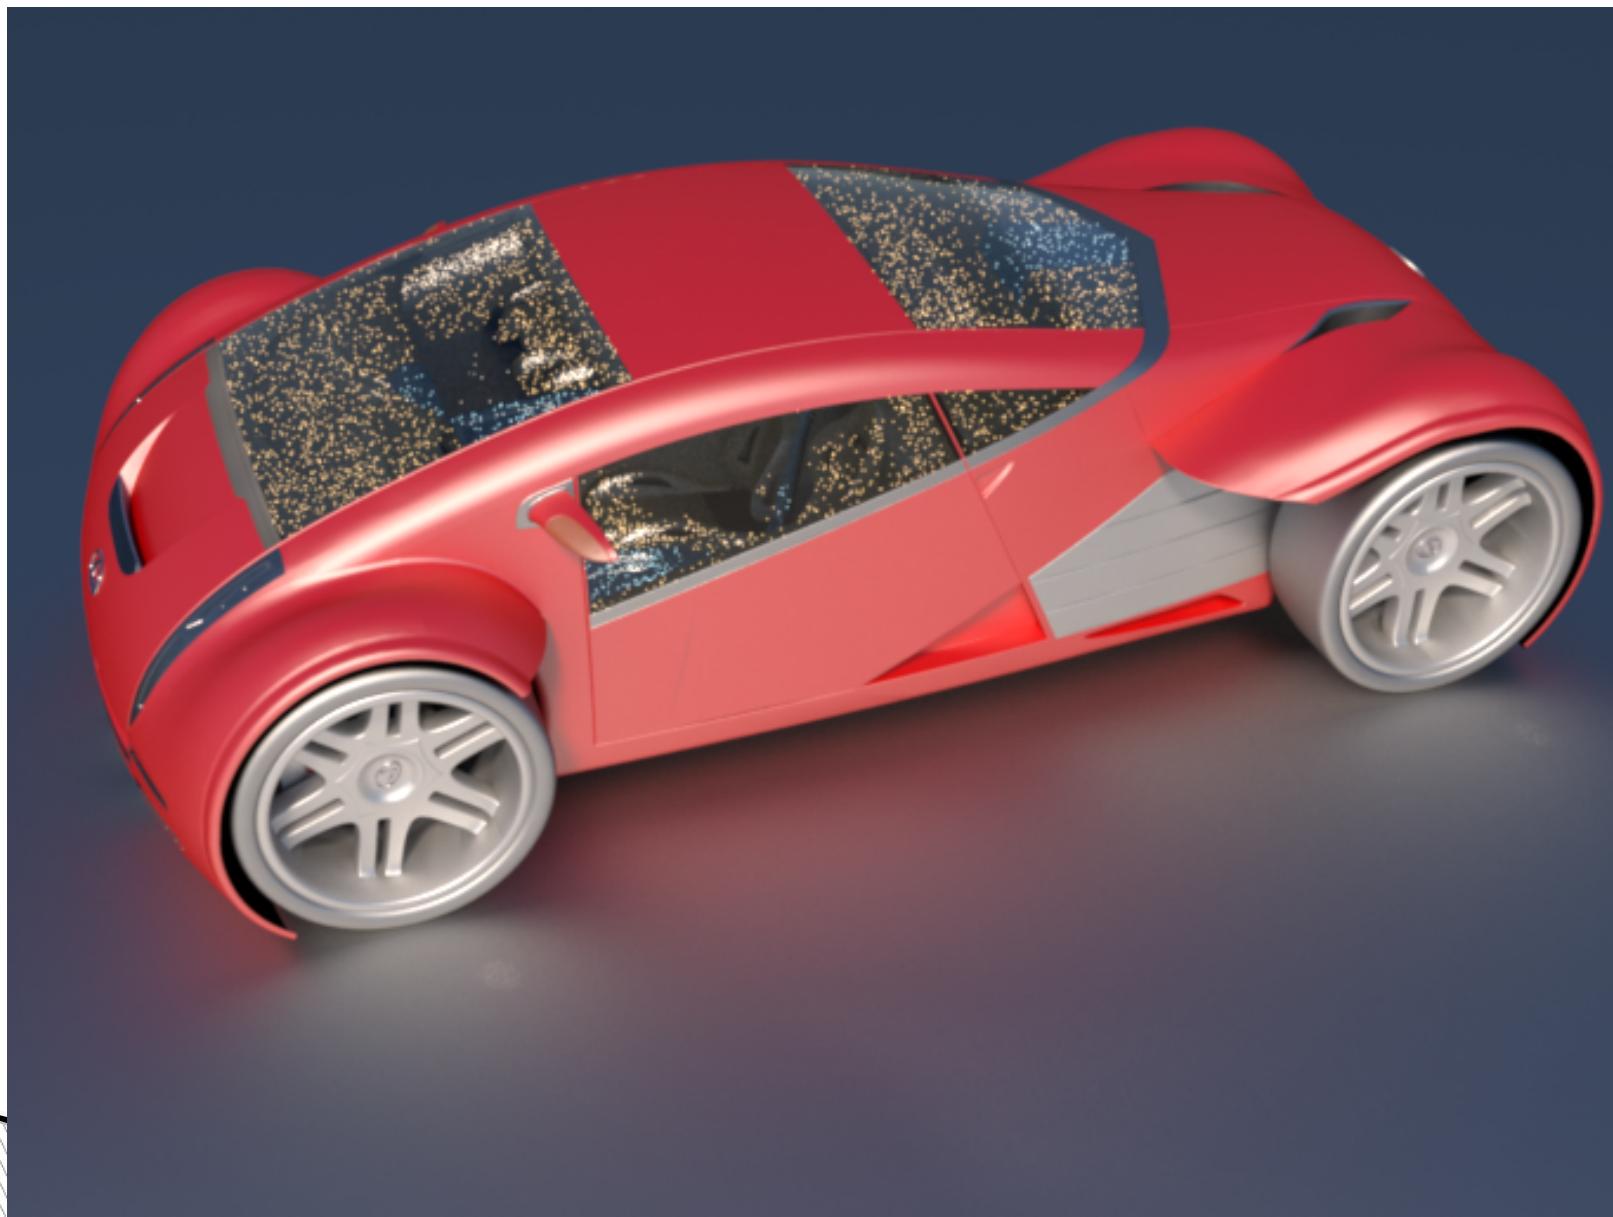
- ▶ Jensen (1996)
  - Final gather pass: + 51 mn



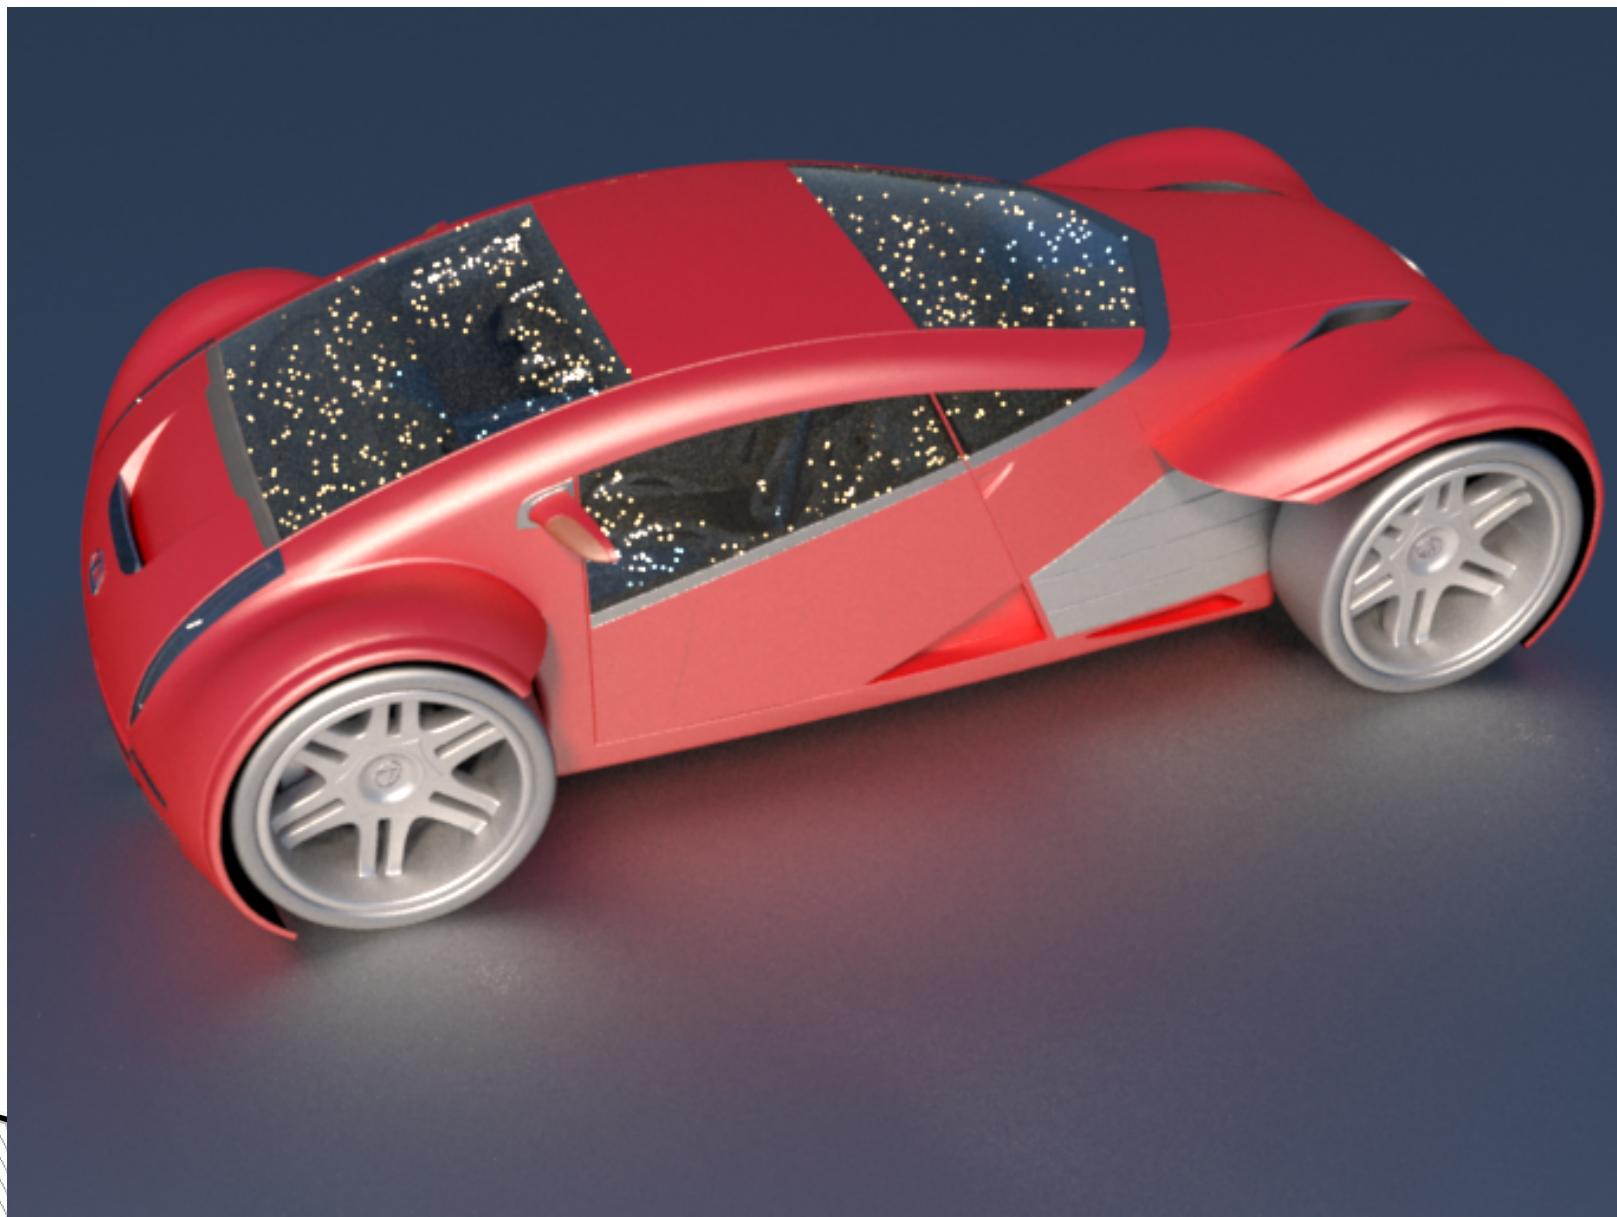
PM: 10k p, 10k cp, 32 spp, 20 mn



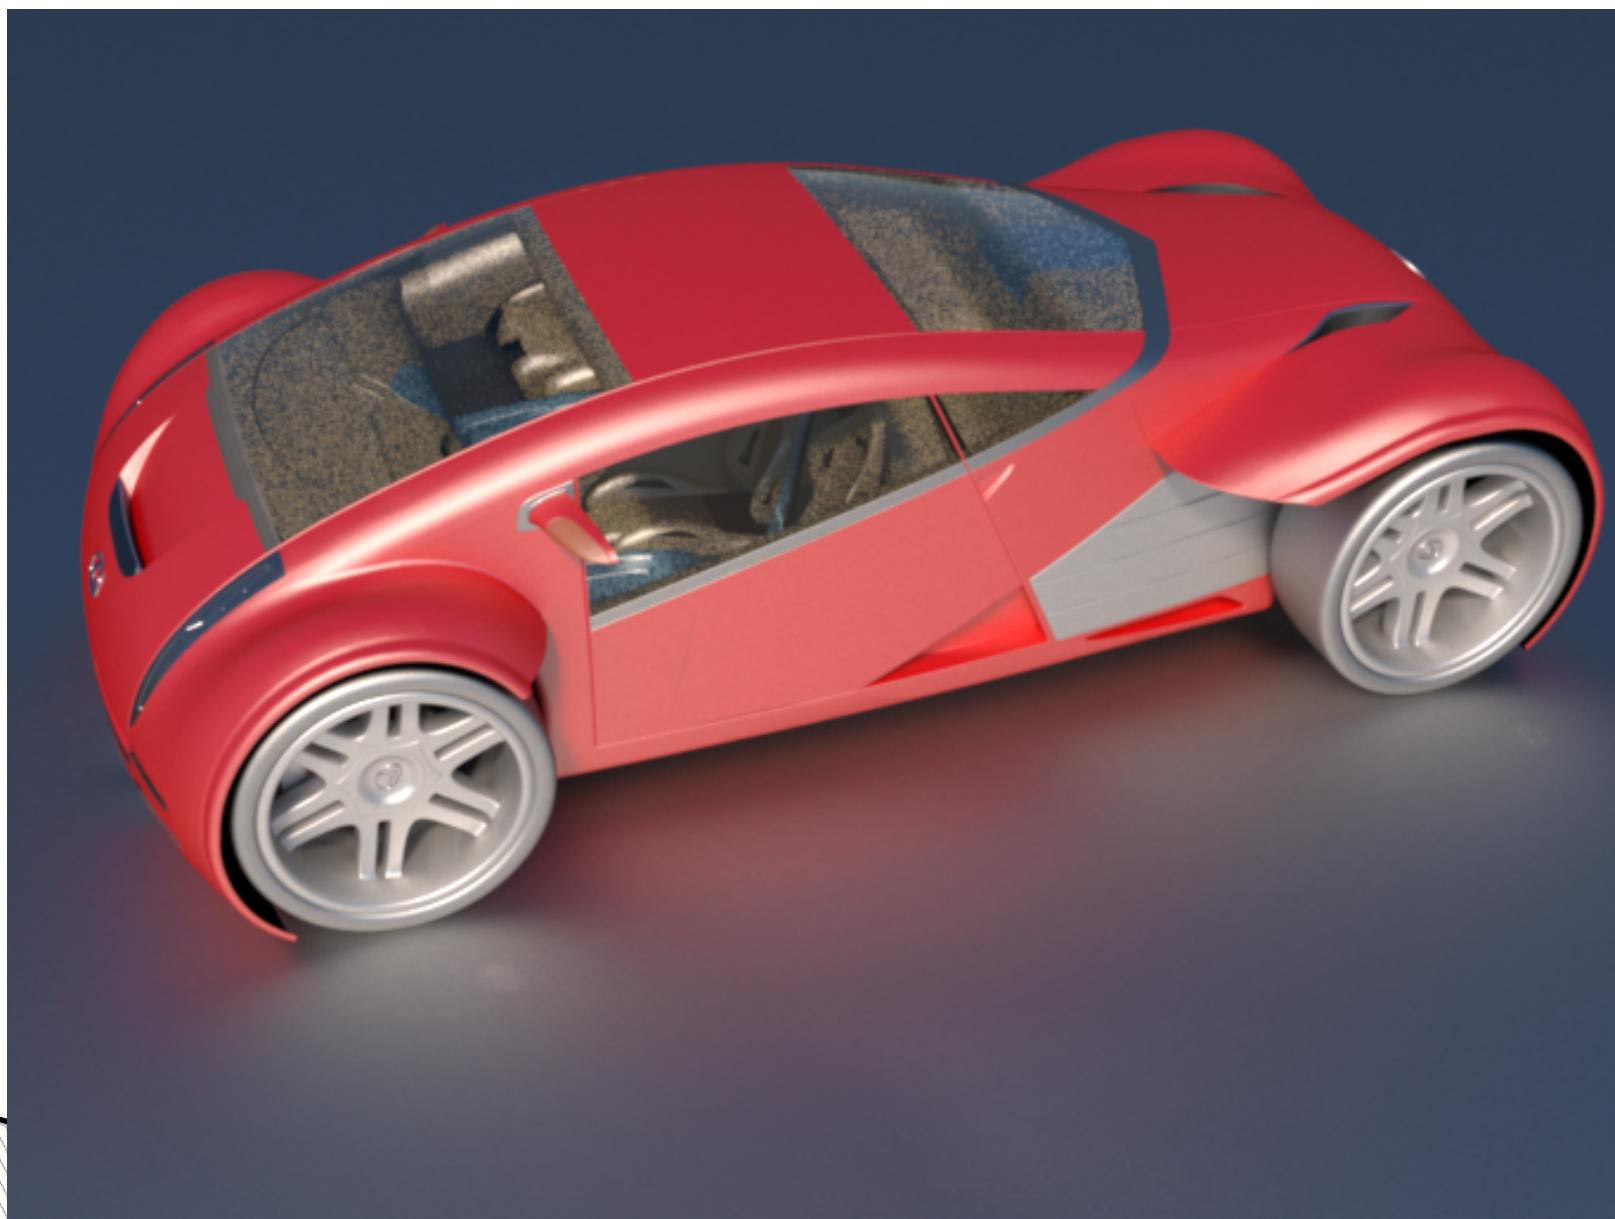
PM: 10k p, 10k cp, 256 spp, 1.3h



PM: 100k p, 100k cp, 32 spp, 1.8h

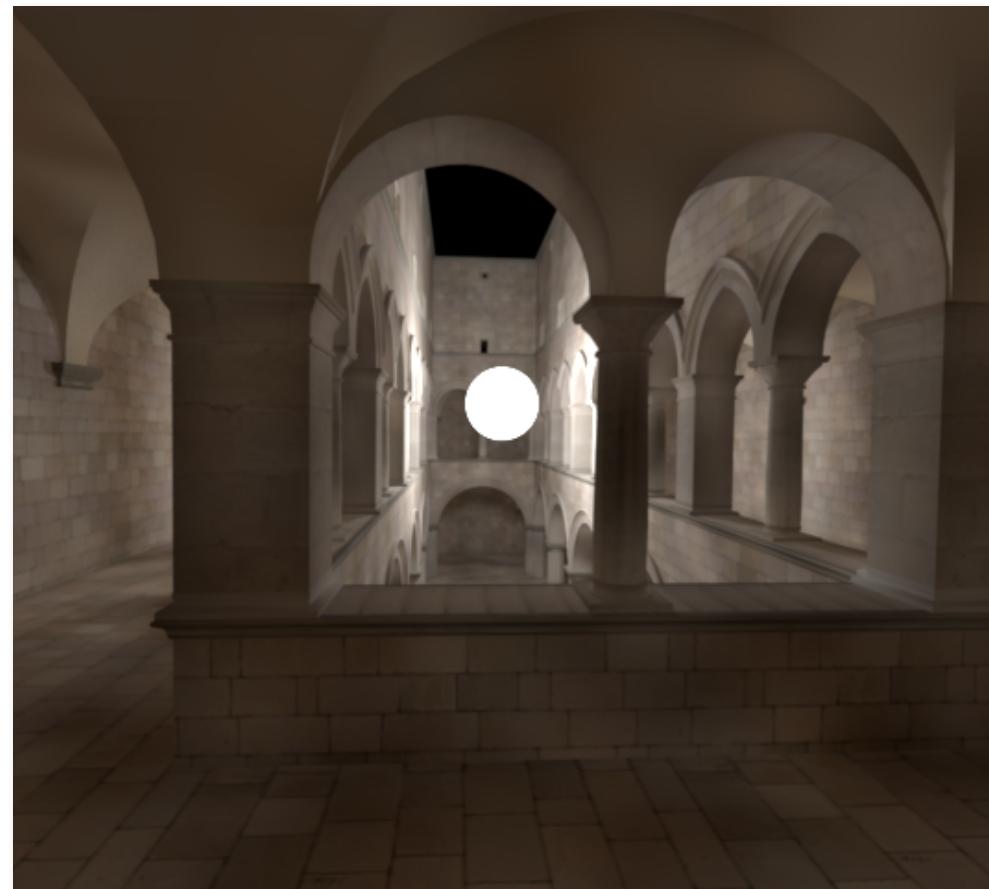


PM: 100k p, 100k cp, 256-8192 spp, 2.3 h



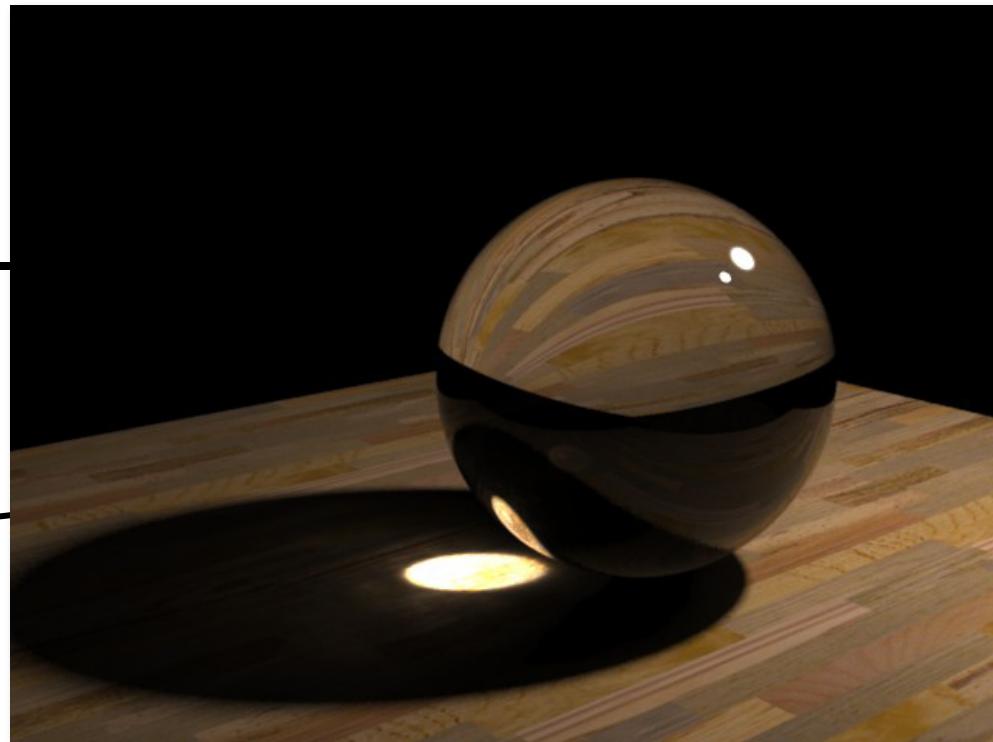
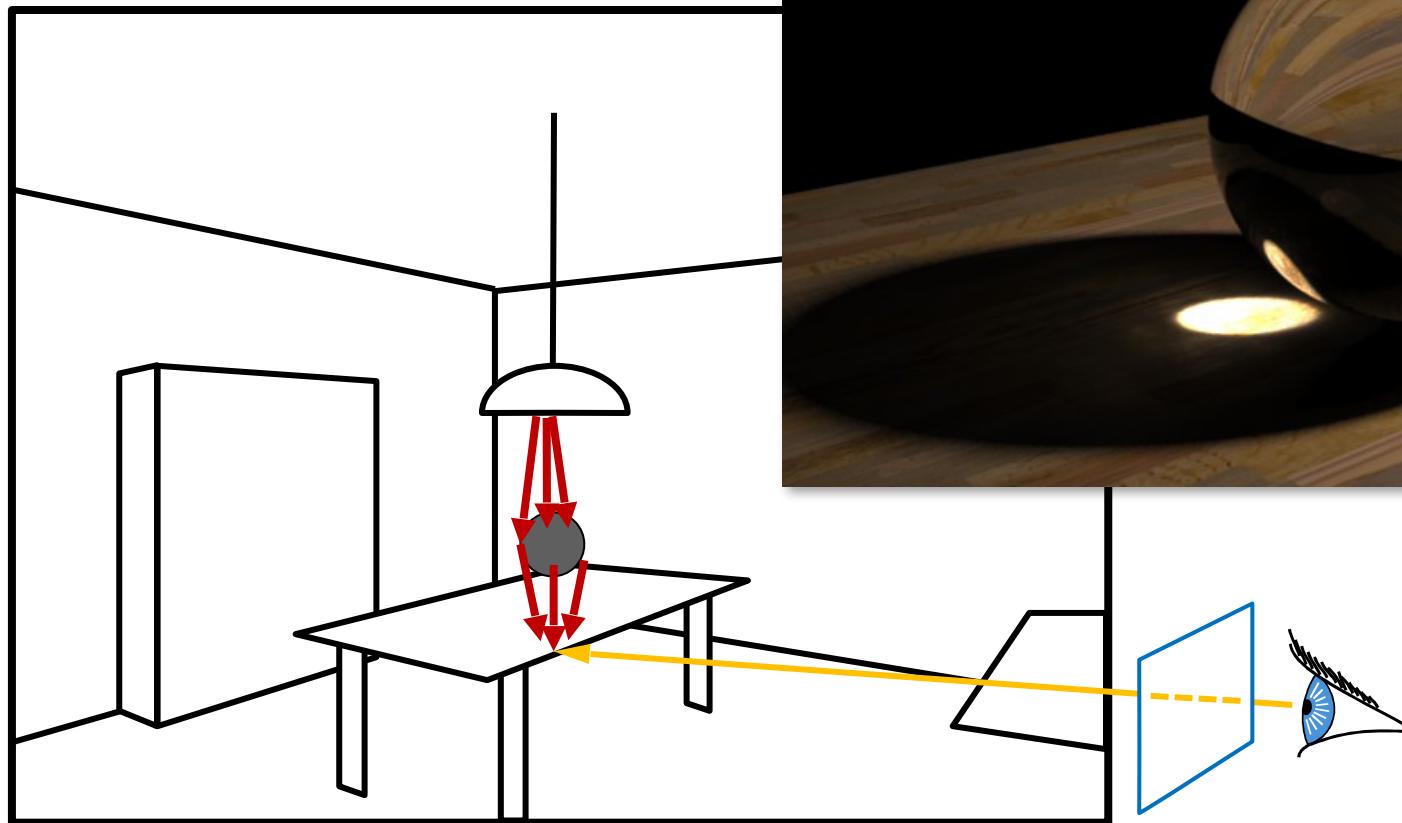
# Results

Mitsuba (<http://www.mitsuba-renderer.org/>)



# Caustics

- ▶ Separate Photon map for refraction



# Photon Mapping: summary

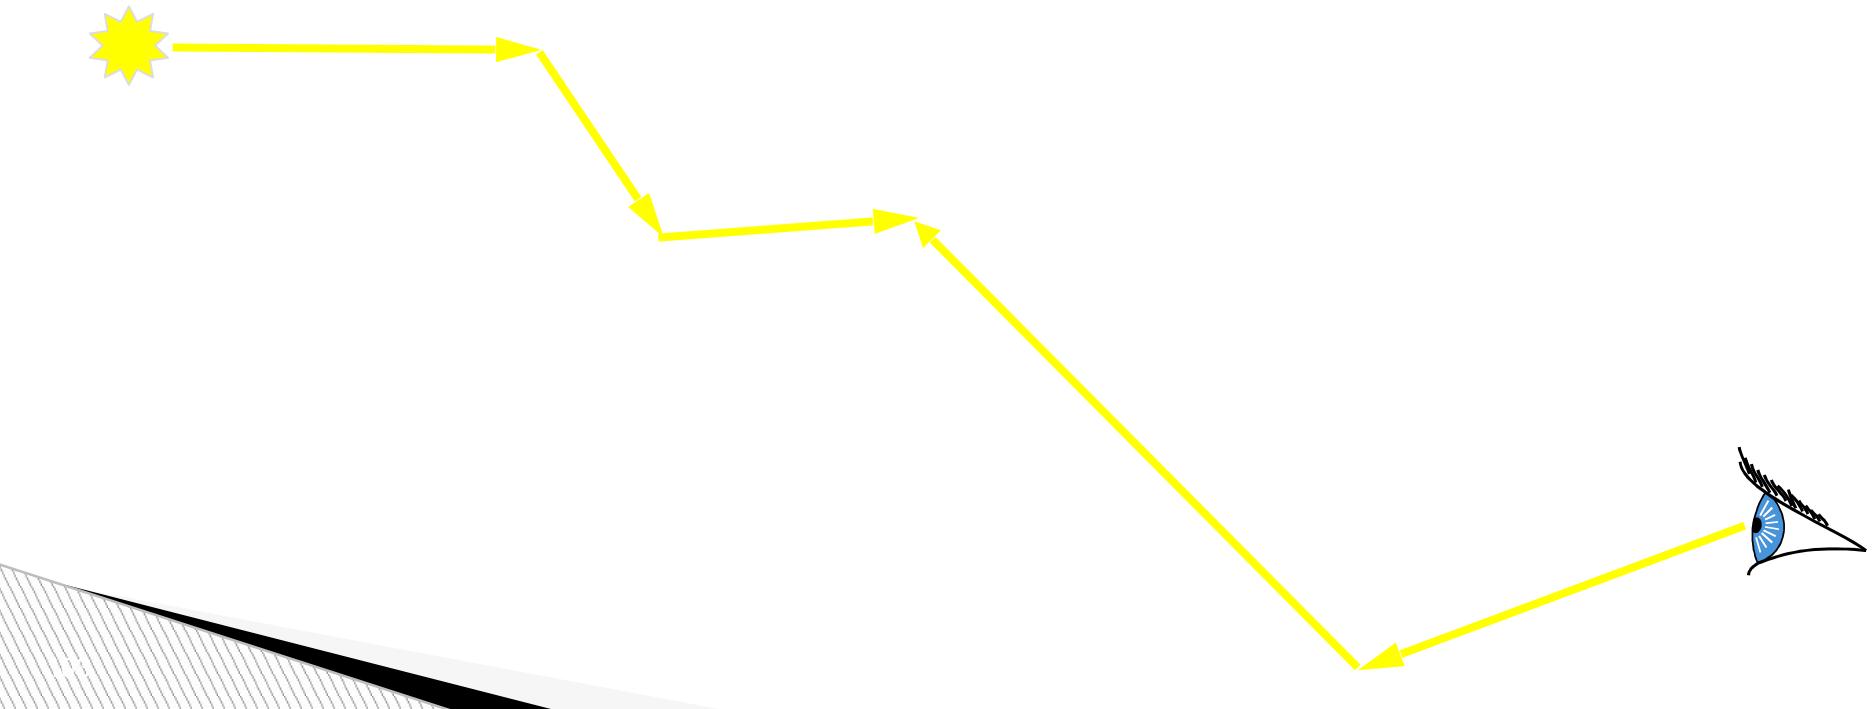
- ▶ Pre-computation view-point independent
  - Storage on the surfaces
- ▶ Good representation for caustics
- ▶ Noisy: smoothing based on the samples
  - Reconstruction of the radiance function
- ▶ Can be coded in two passes with a ray-tracer
  - One pass for each direction

# Metropolis Light Transport

- ▶ Some paths are hard to find
- ▶ Explorer path space
  - BDPT: treat each path independently
  - Metropolis: use existing paths (if they're good)
- ▶ Start with a given path from eye to light
- ▶ Mutations + perturbations
  - Accept mutations: with probability
  - Importance sampling for energy carried

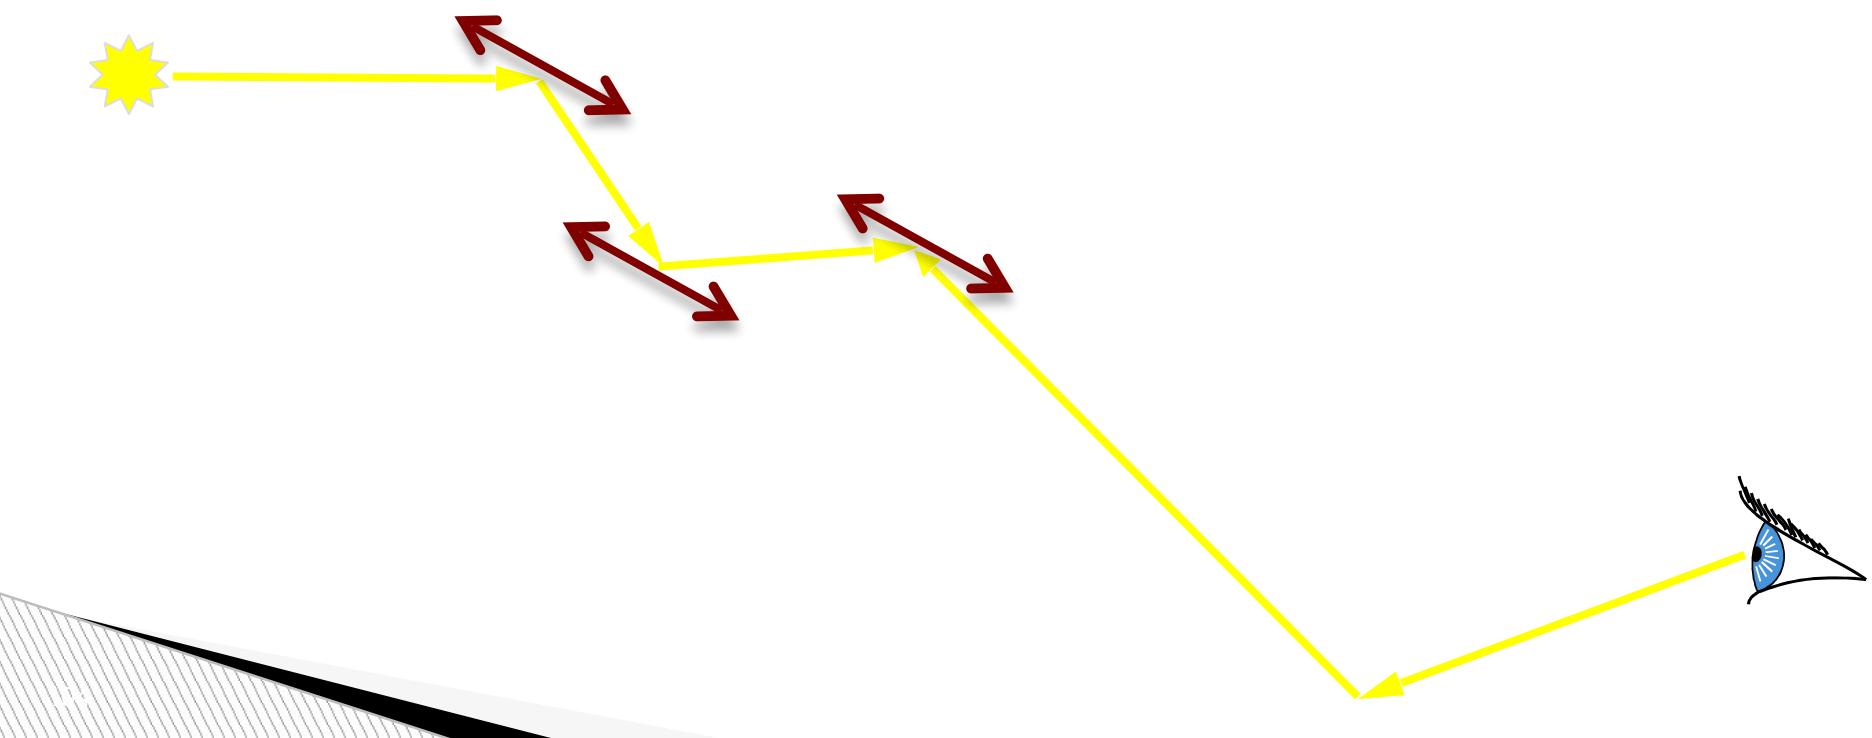
# Metropolis Light Transport

- ▶ Start with a light path from eye to source:



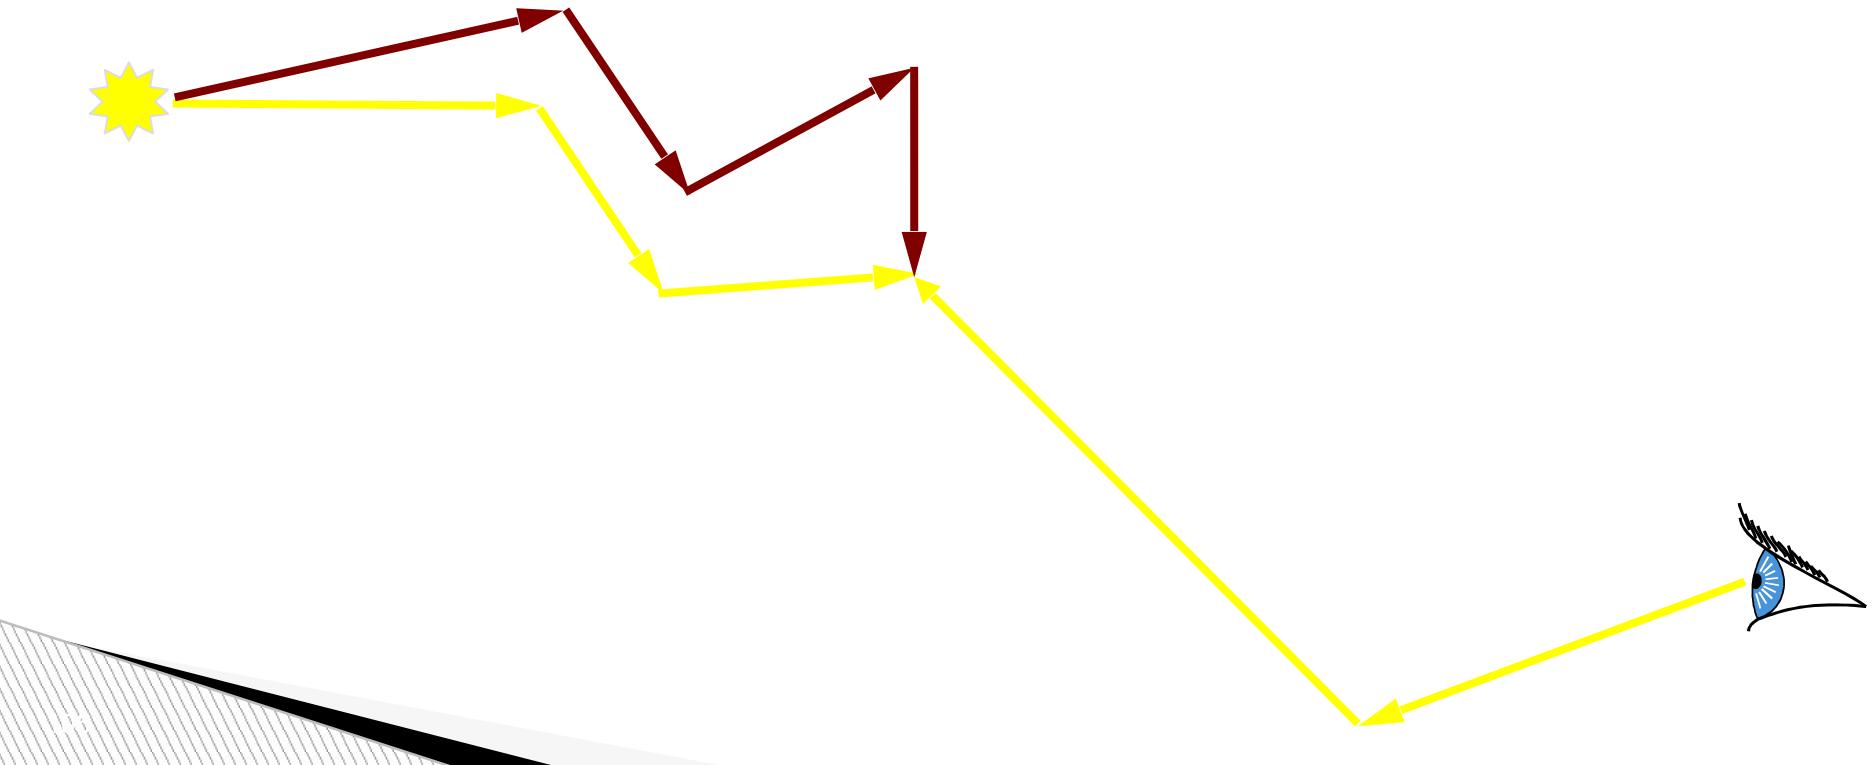
# Metropolis Light Transport

- ▶ Start with a light path from eye to source
- ▶ Random perturbations
  - Keep / don't keep



# Metropolis Light Transport

- ▶ Start with a light path from eye to source
- ▶ Random mutations (adding vertices)
  - Keep / don't keep



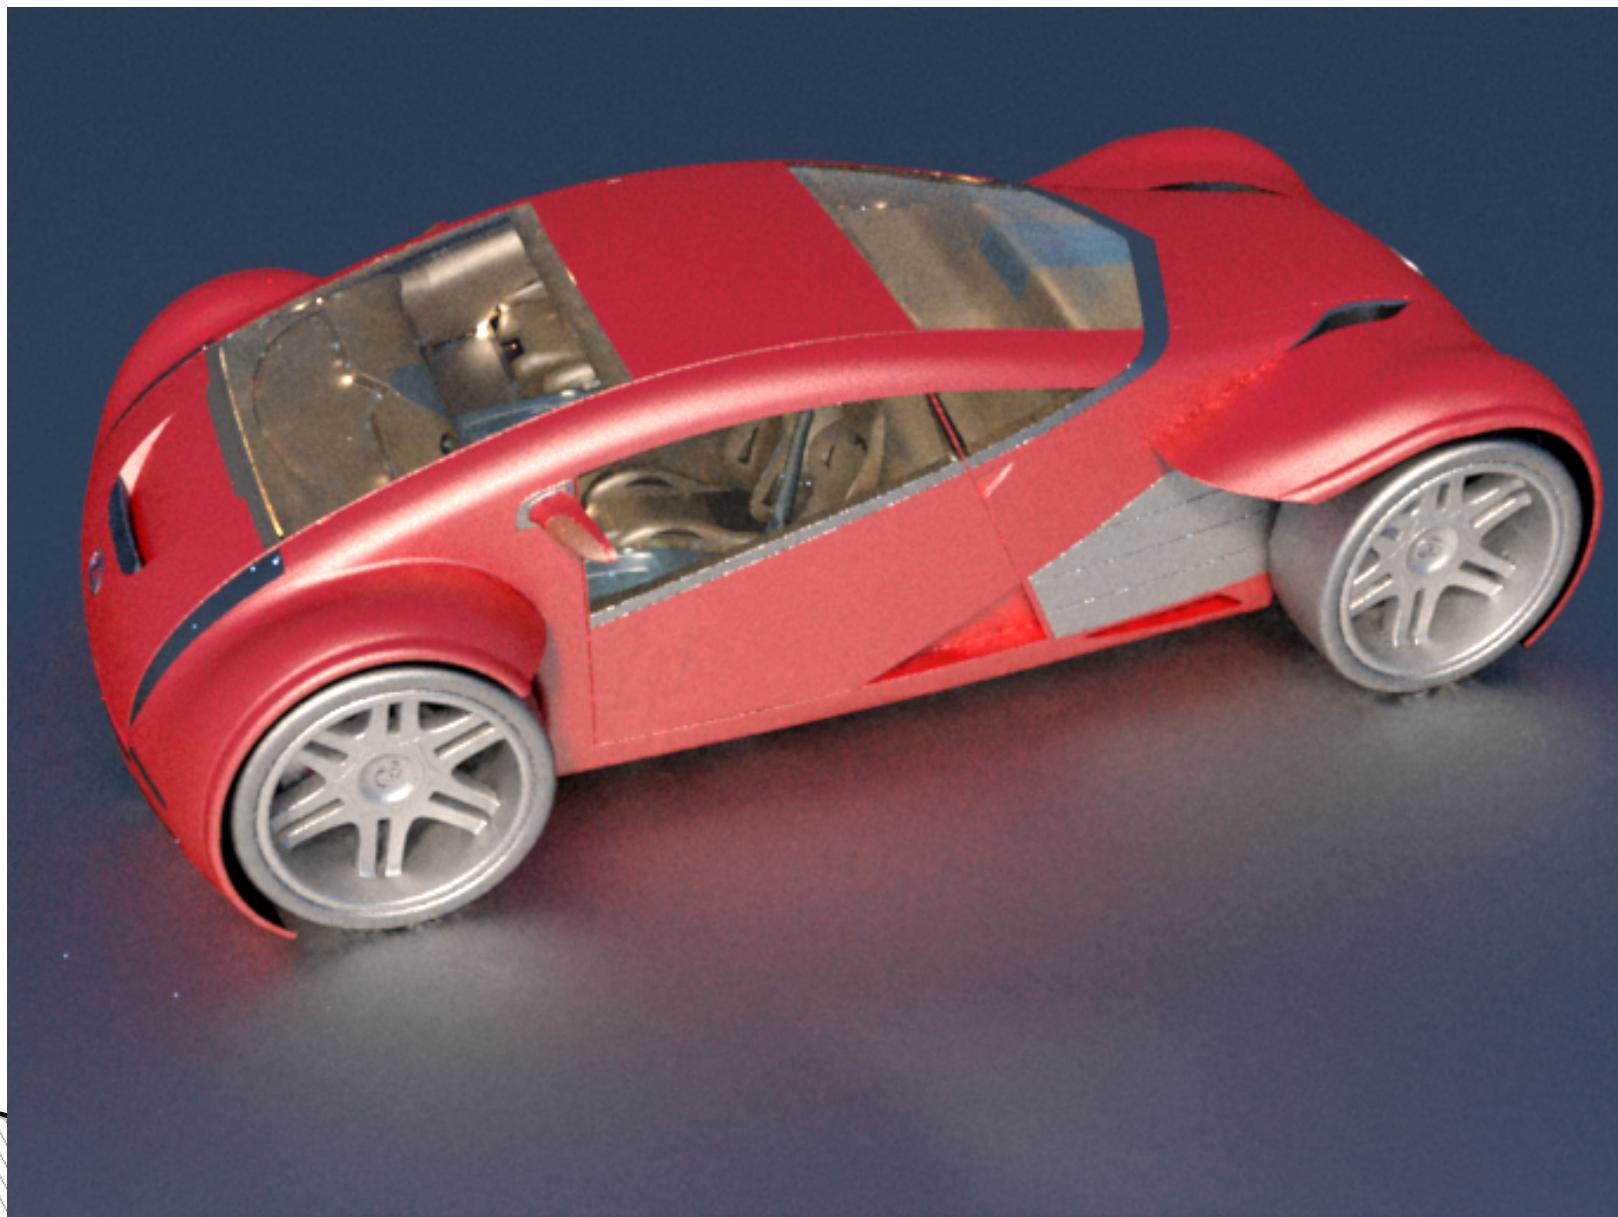
# Bi-Directional Path Transport



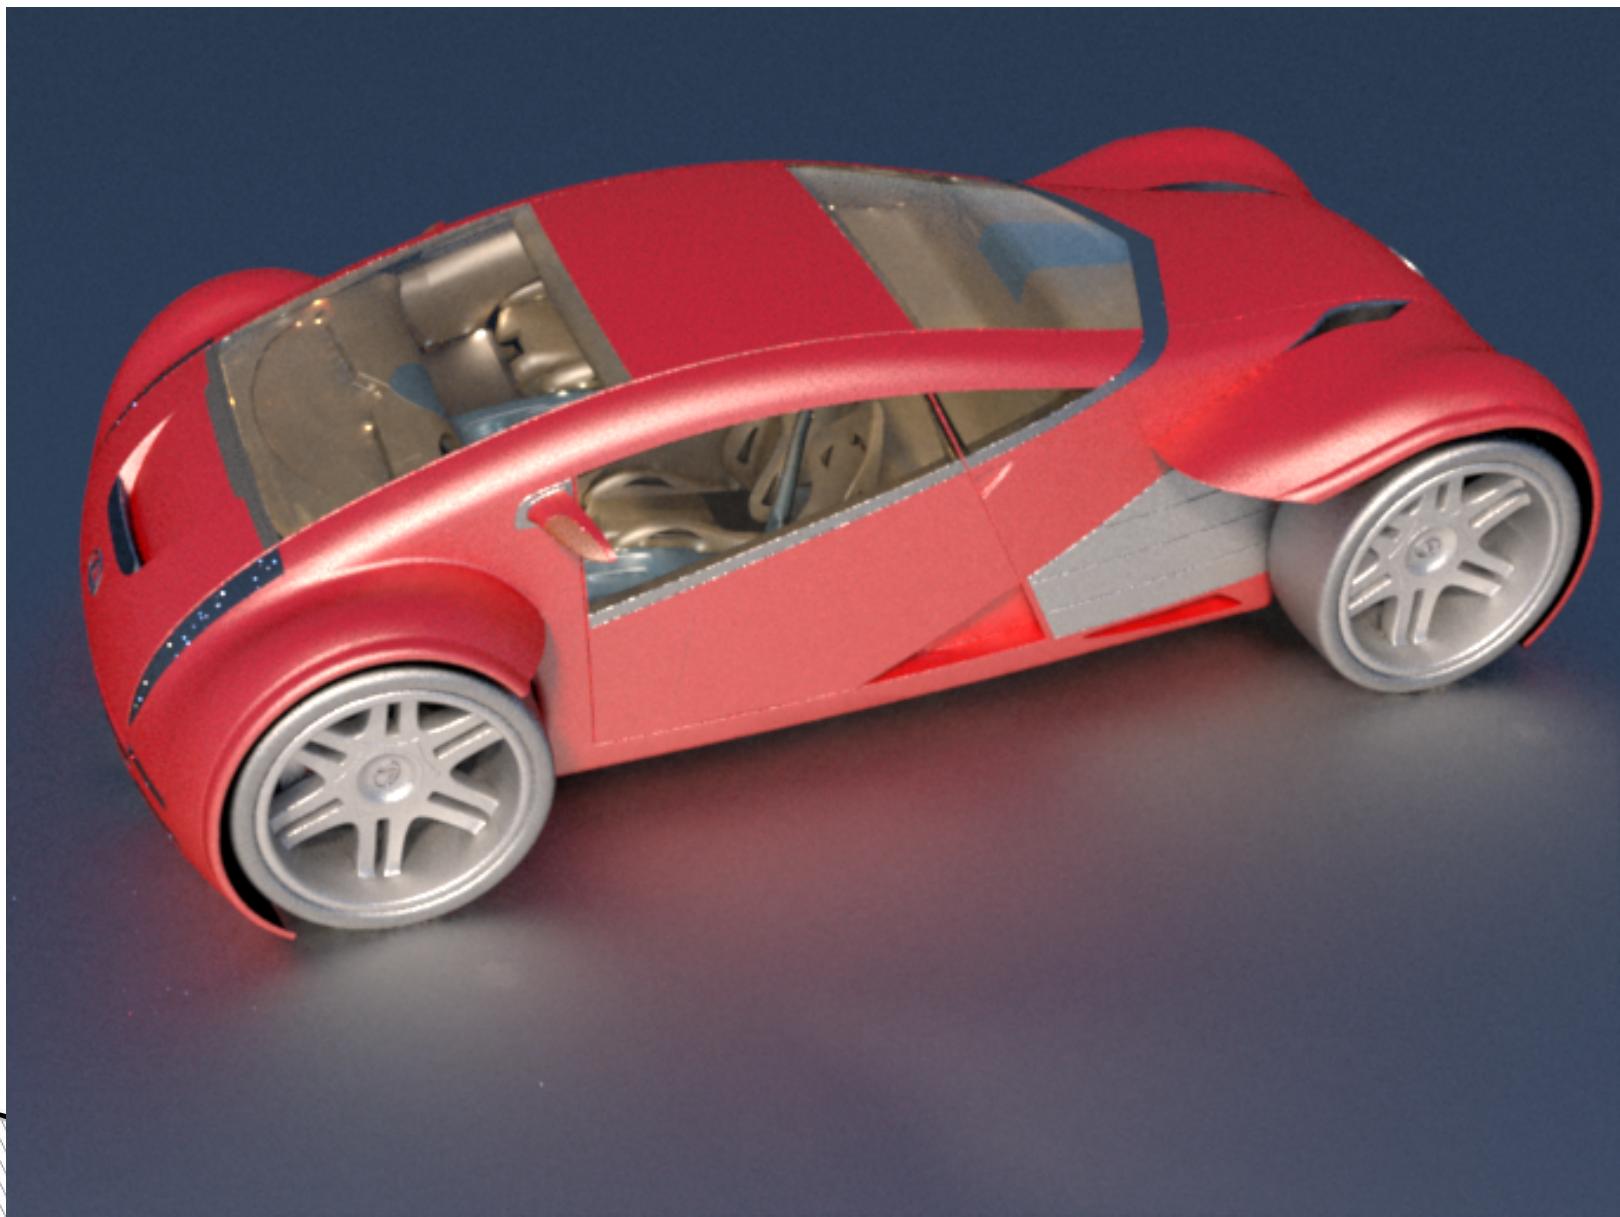
# Metropolis light transport



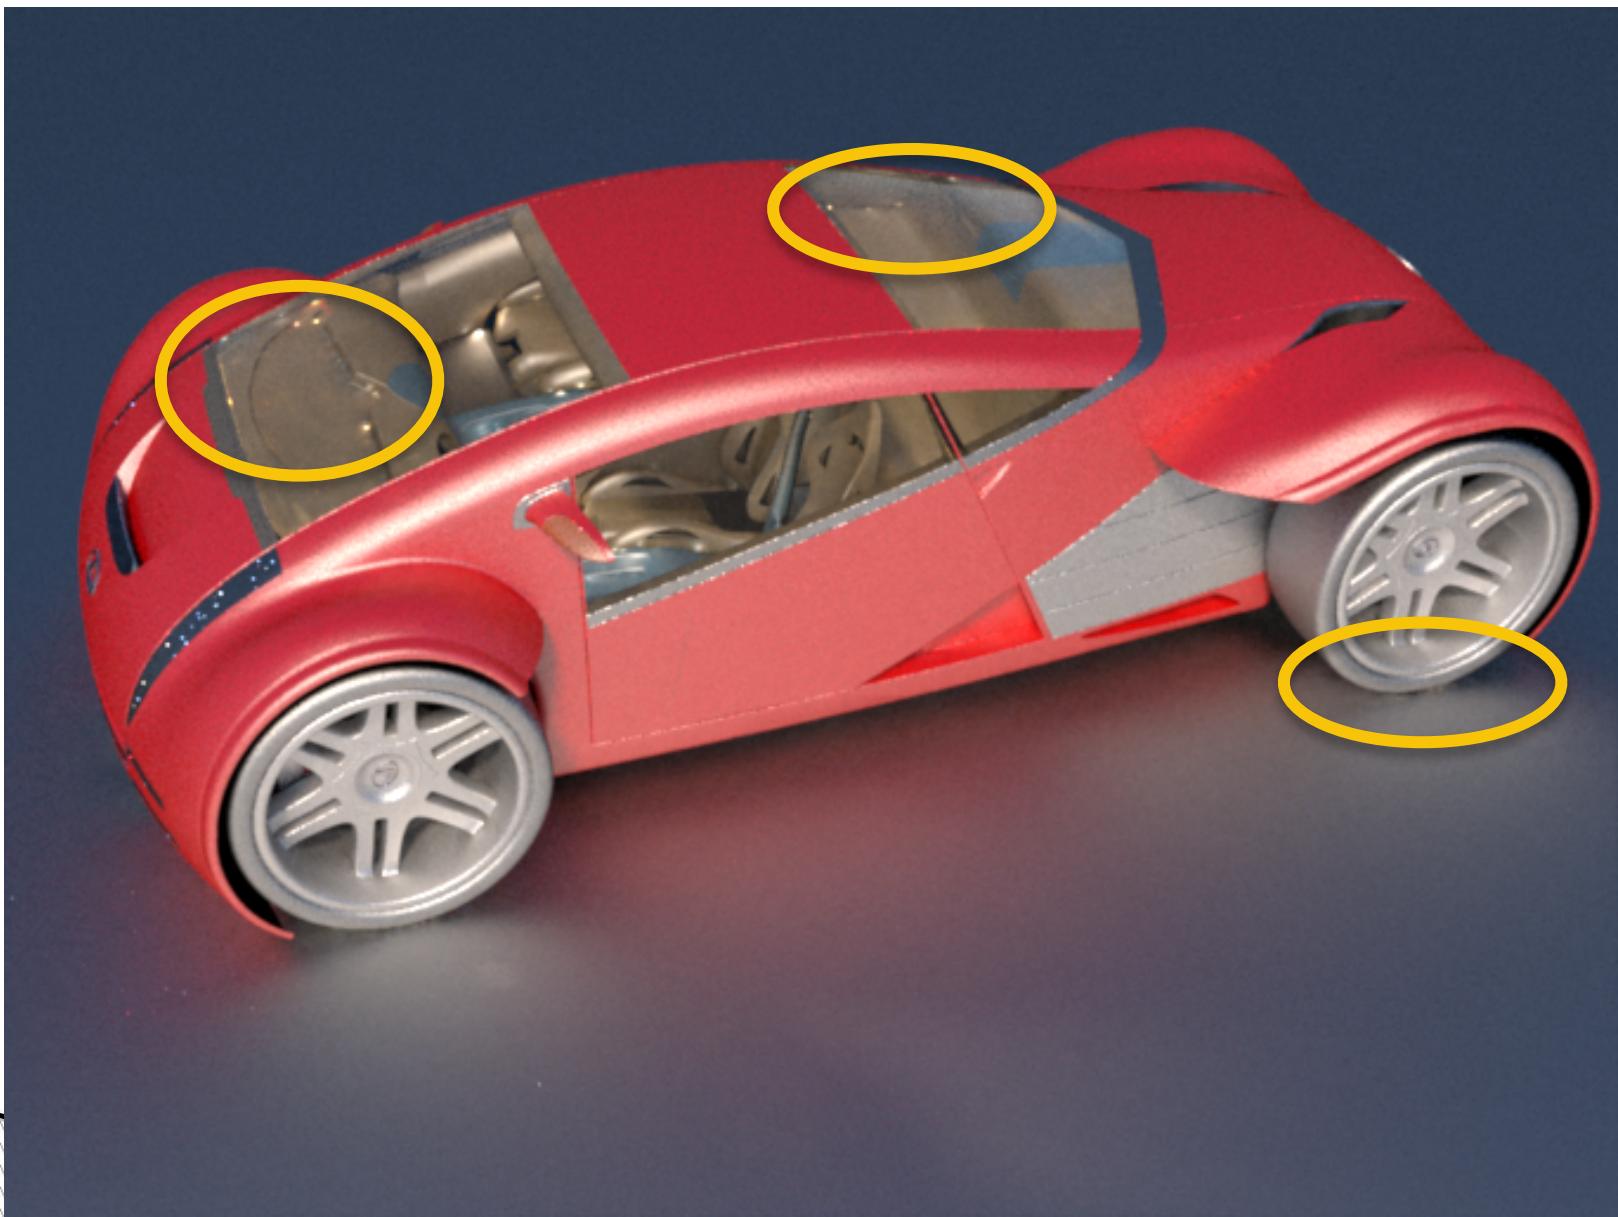
# MLT, 256 mutations/pixel (2.9 mn)



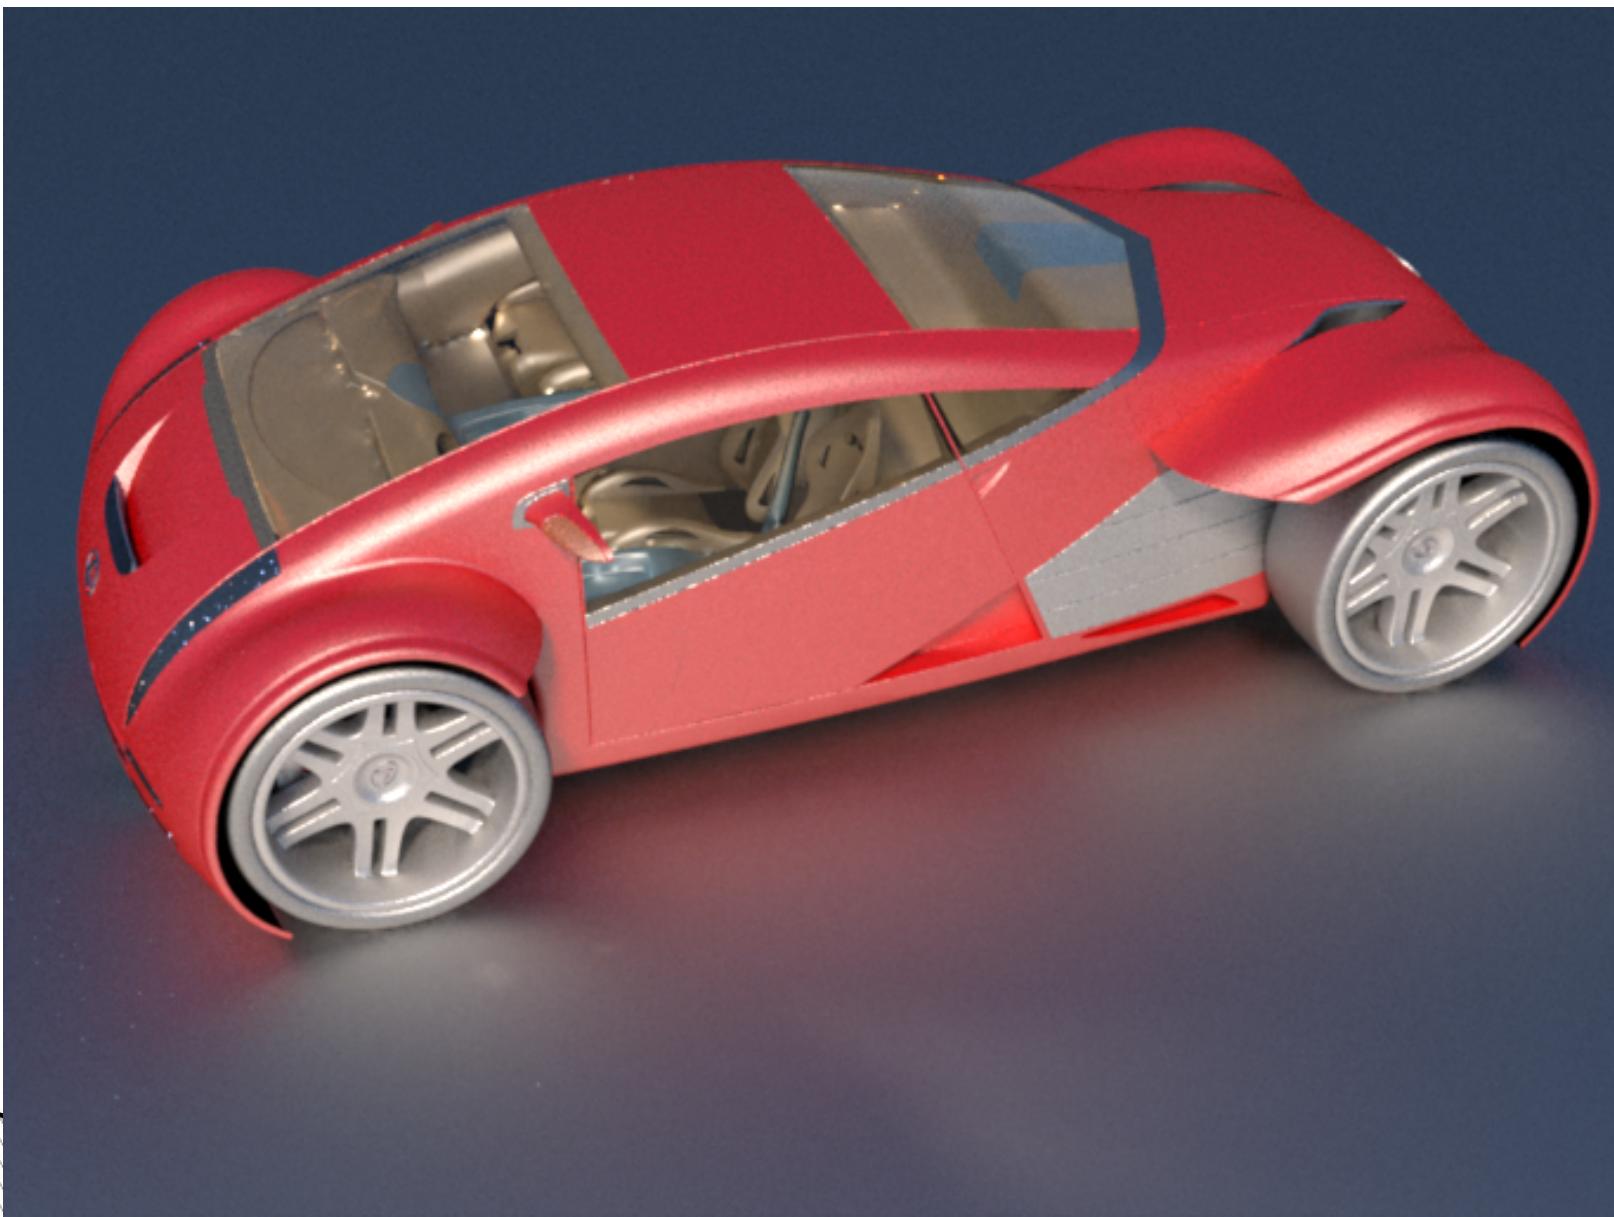
MLT, 1024 mutations/pixel (12 mn)



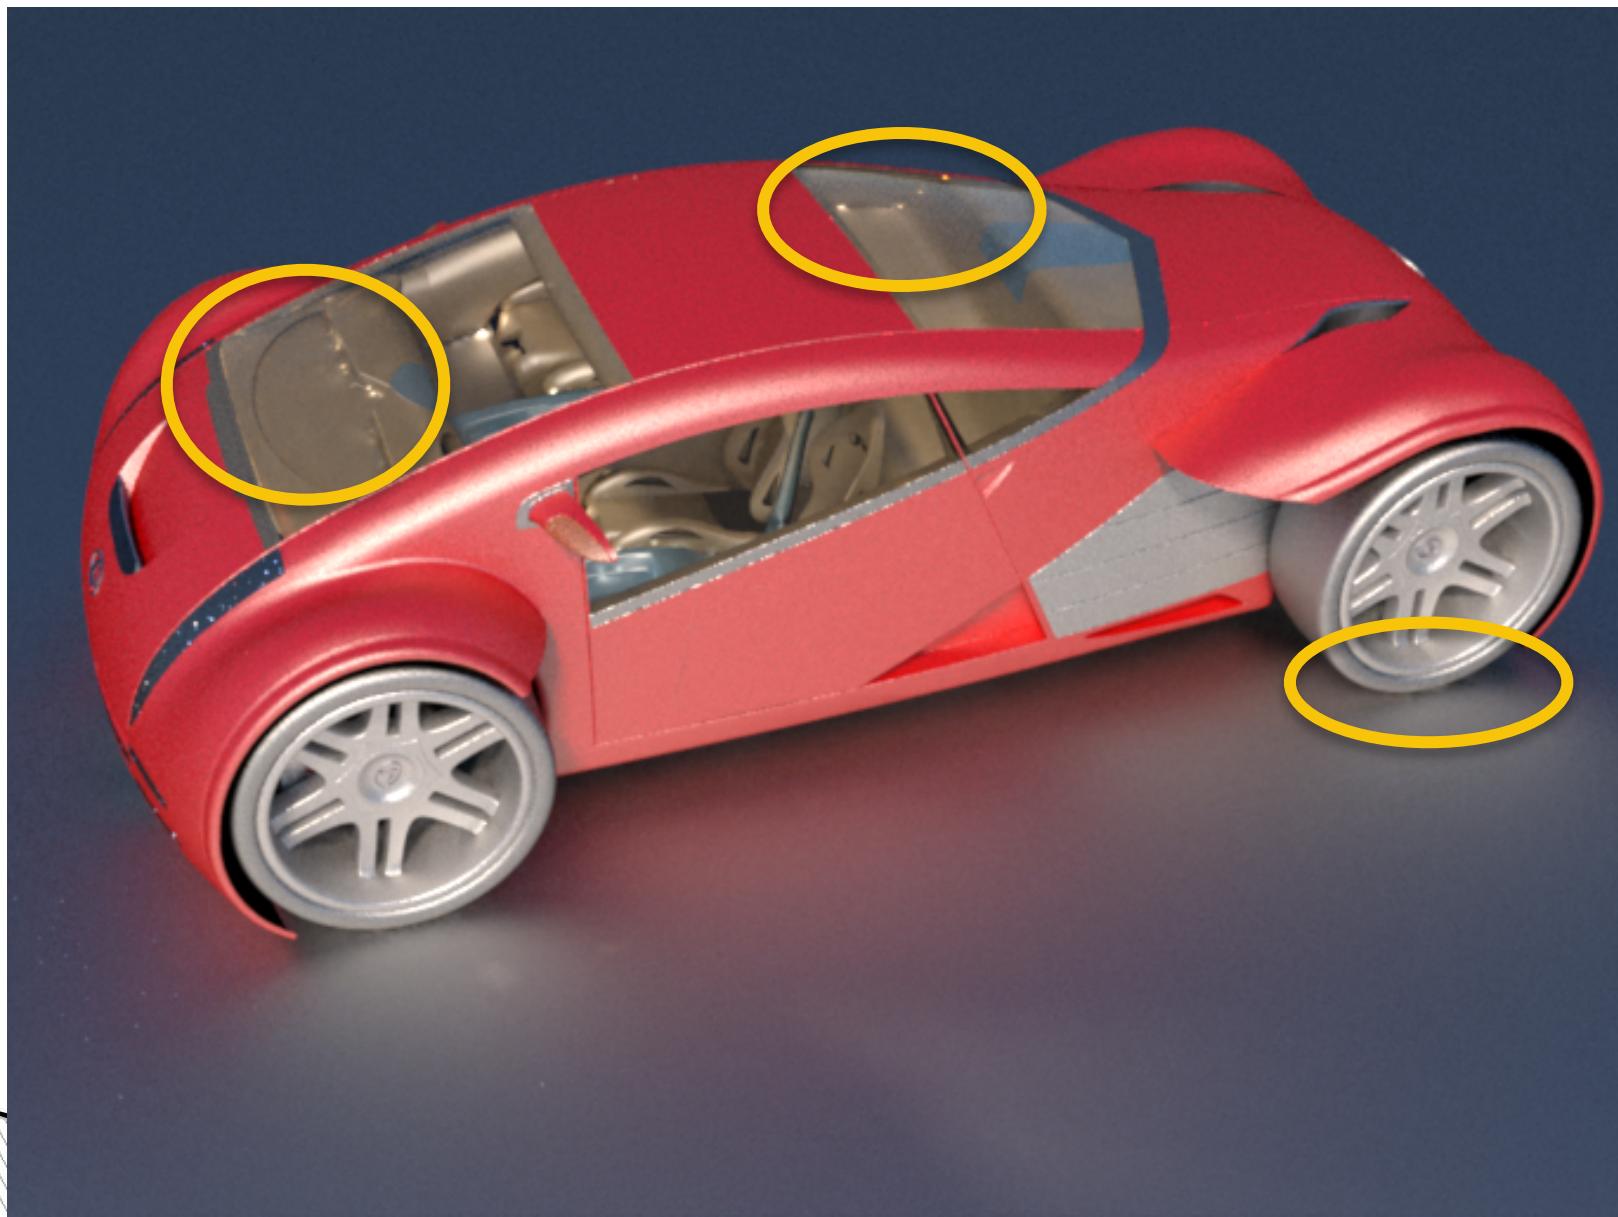
MLT, 1024 mutations/pixel (12 mn)



# MLT, 2048 mutations/pixel (23 mn)



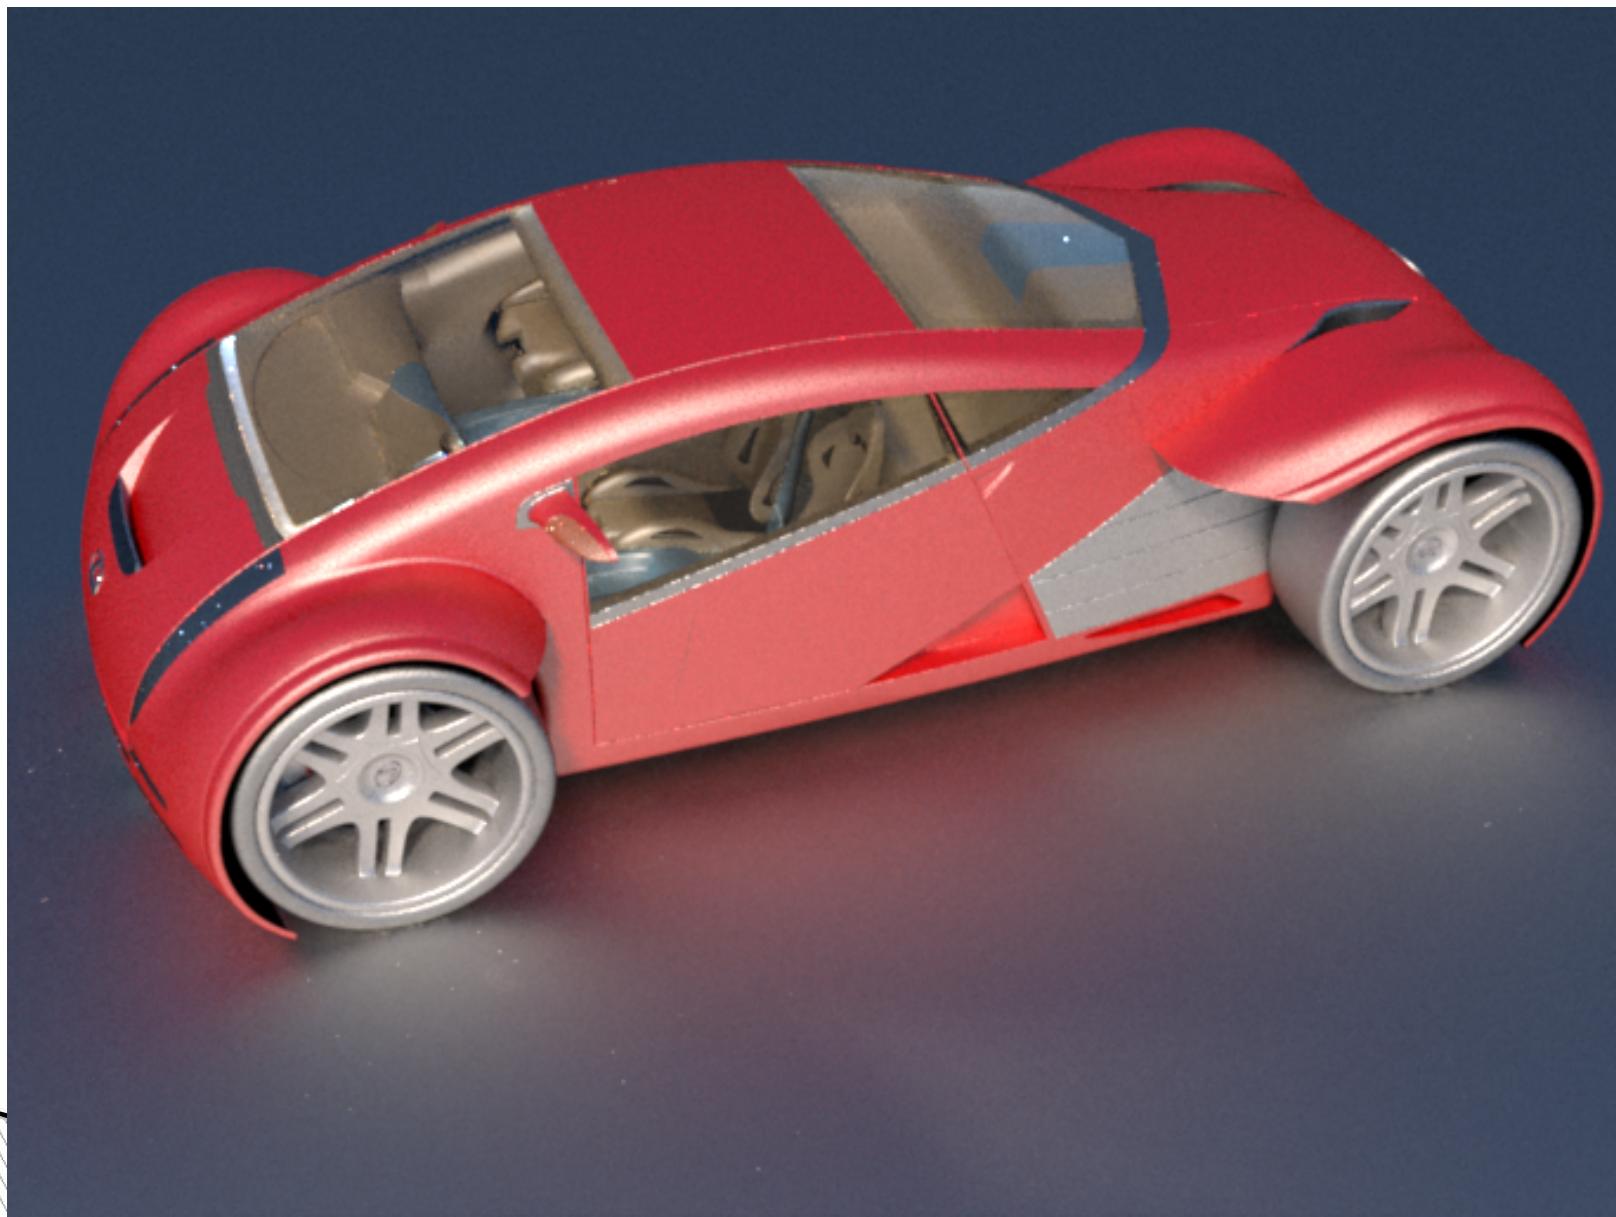
# MLT, 2048 mutations/pixel (23 mn)



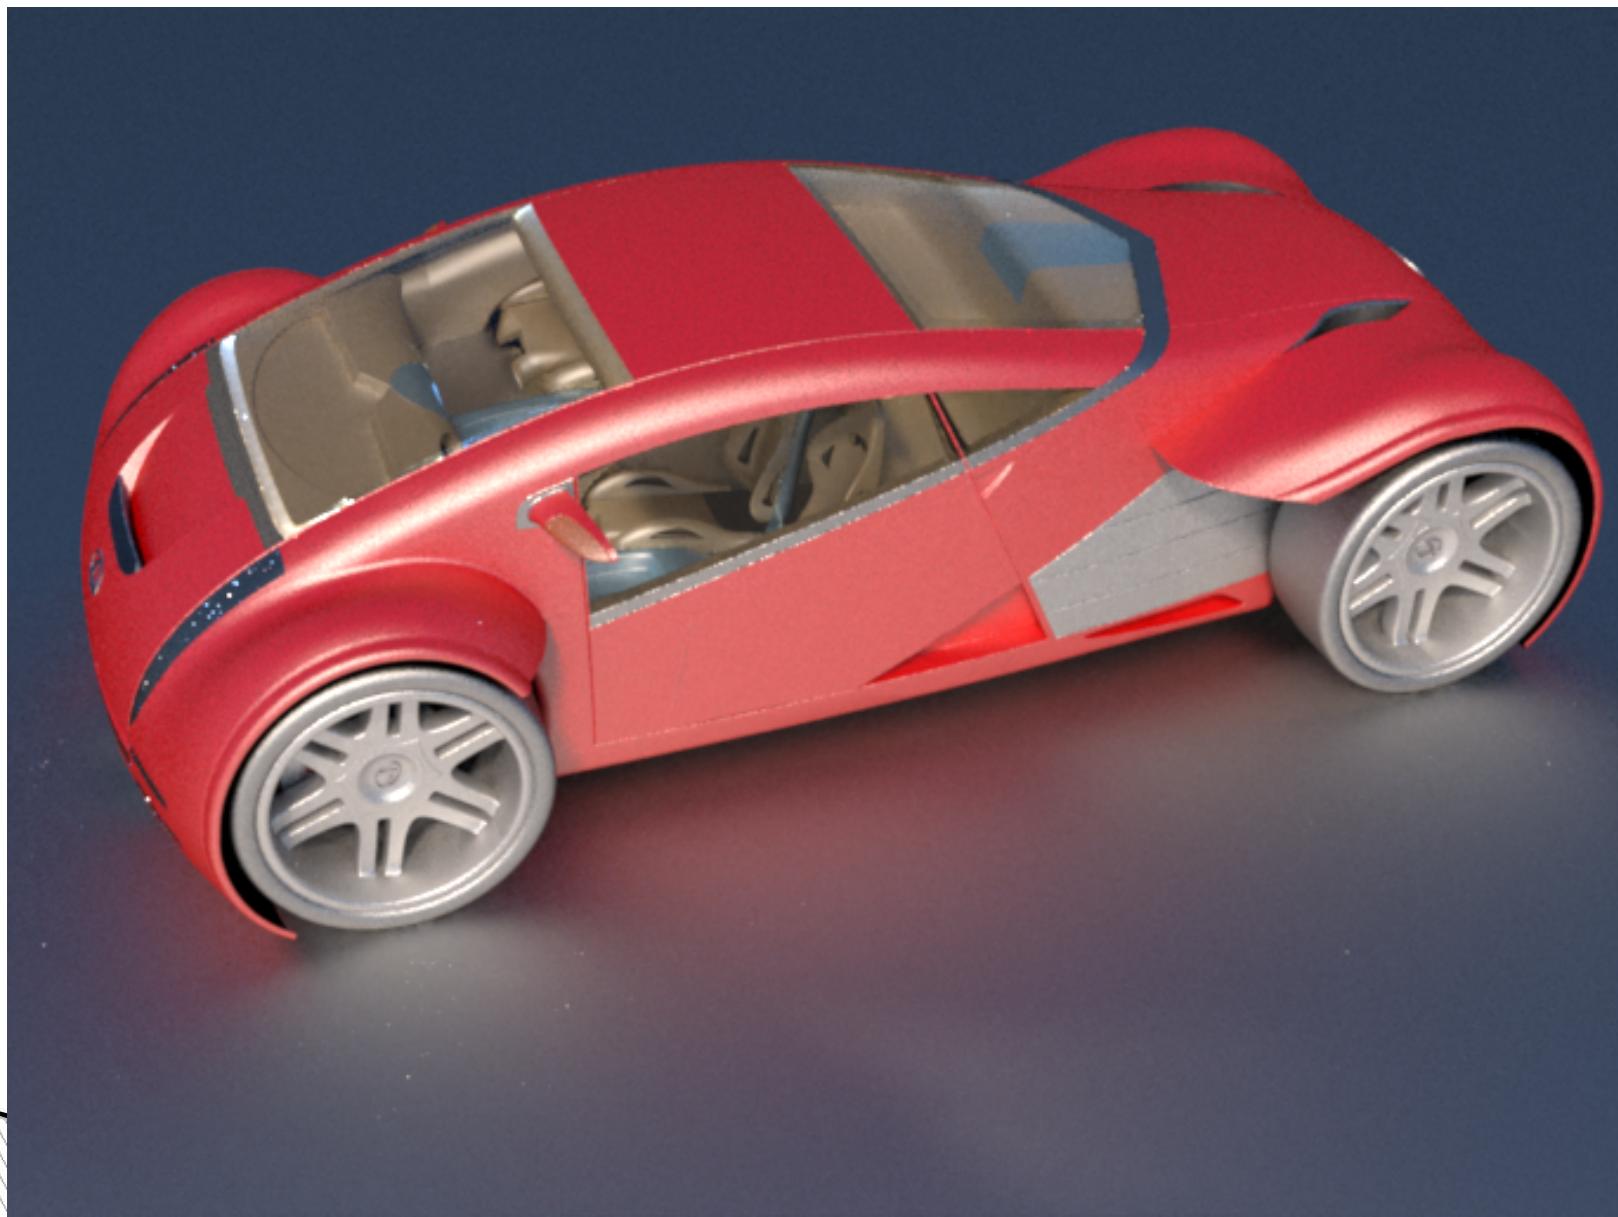
# MLT: what's happening?

- ▶ Test scene n° 1: tailored for MLT/BDPT
- ▶ Scène n° 2:
  - Some paths are hard to find
  - When they're found, higher importance
- ▶ Solution: manifold exploration
  - Interesting paths = manifold (in path space)
  - Stay on this manifold
  - Constrained exploration of path space
  - Extension: Half-Vector Space exploration
- ▶ Aliasing

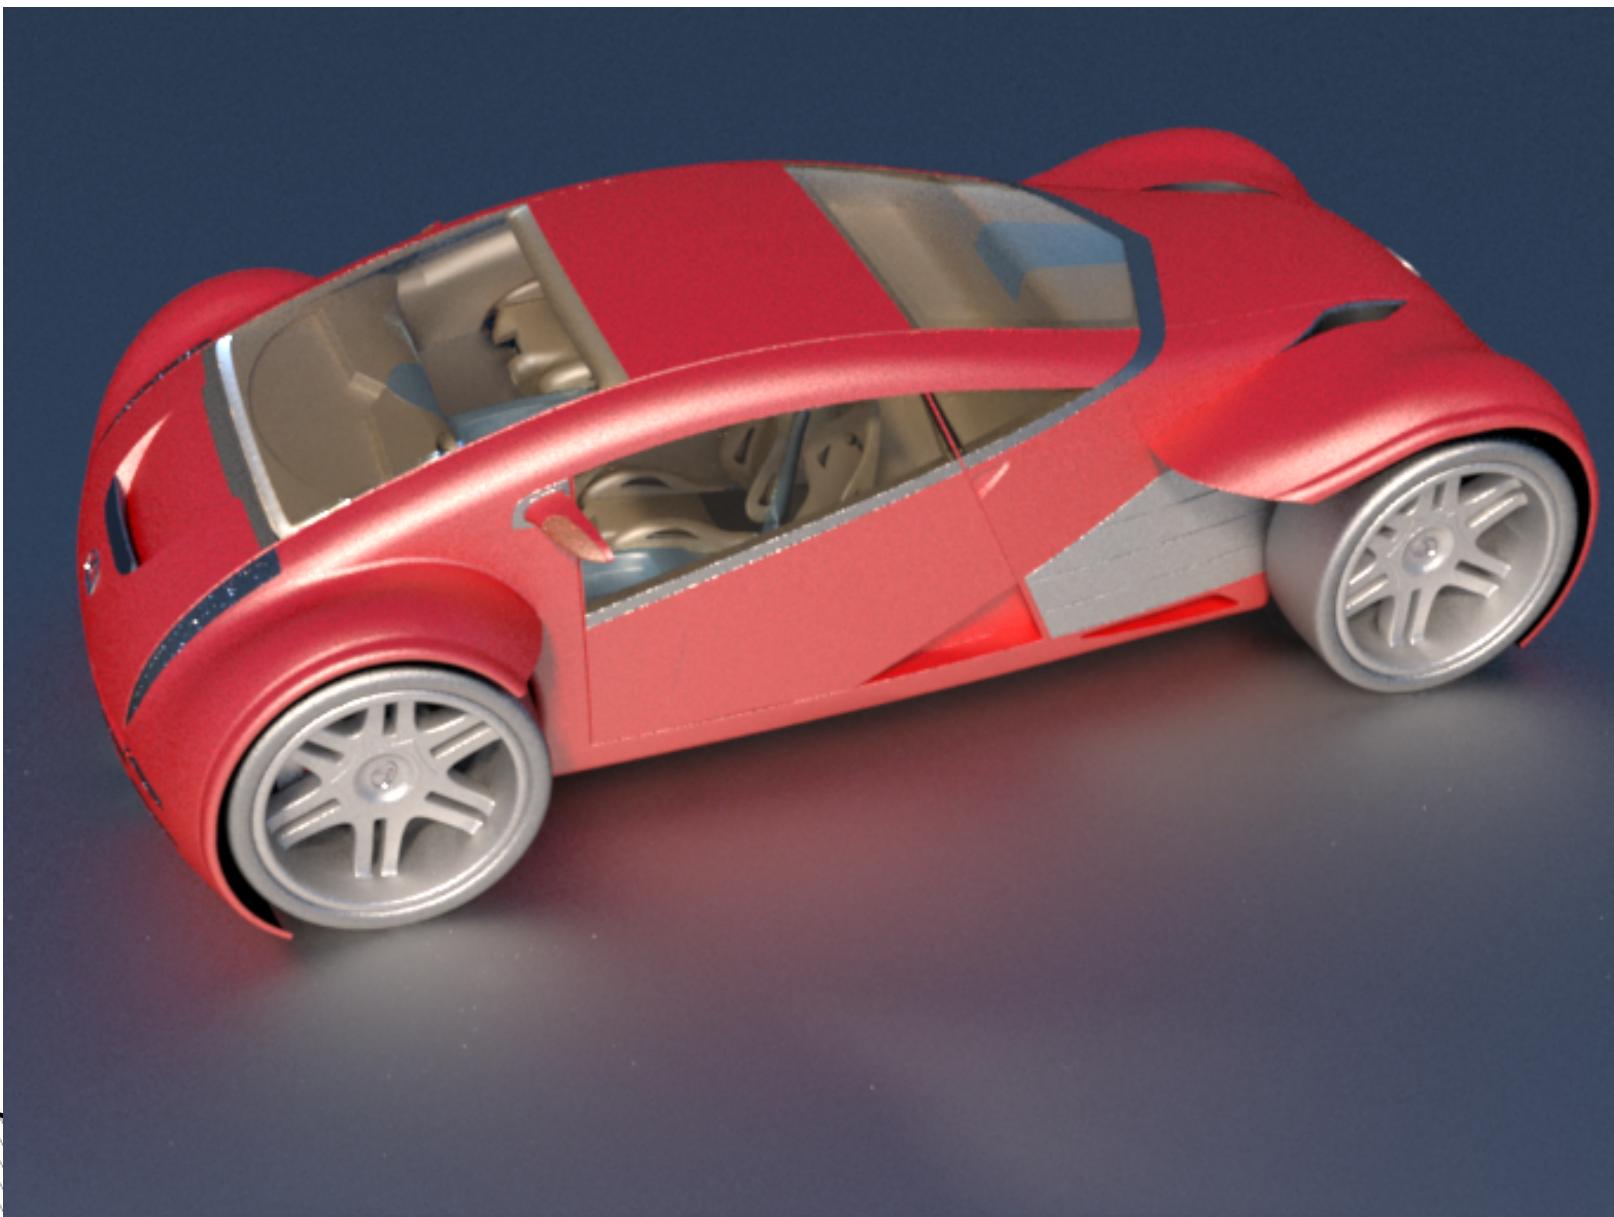
MEMLT, 1024 mut./pixel, 17 mn



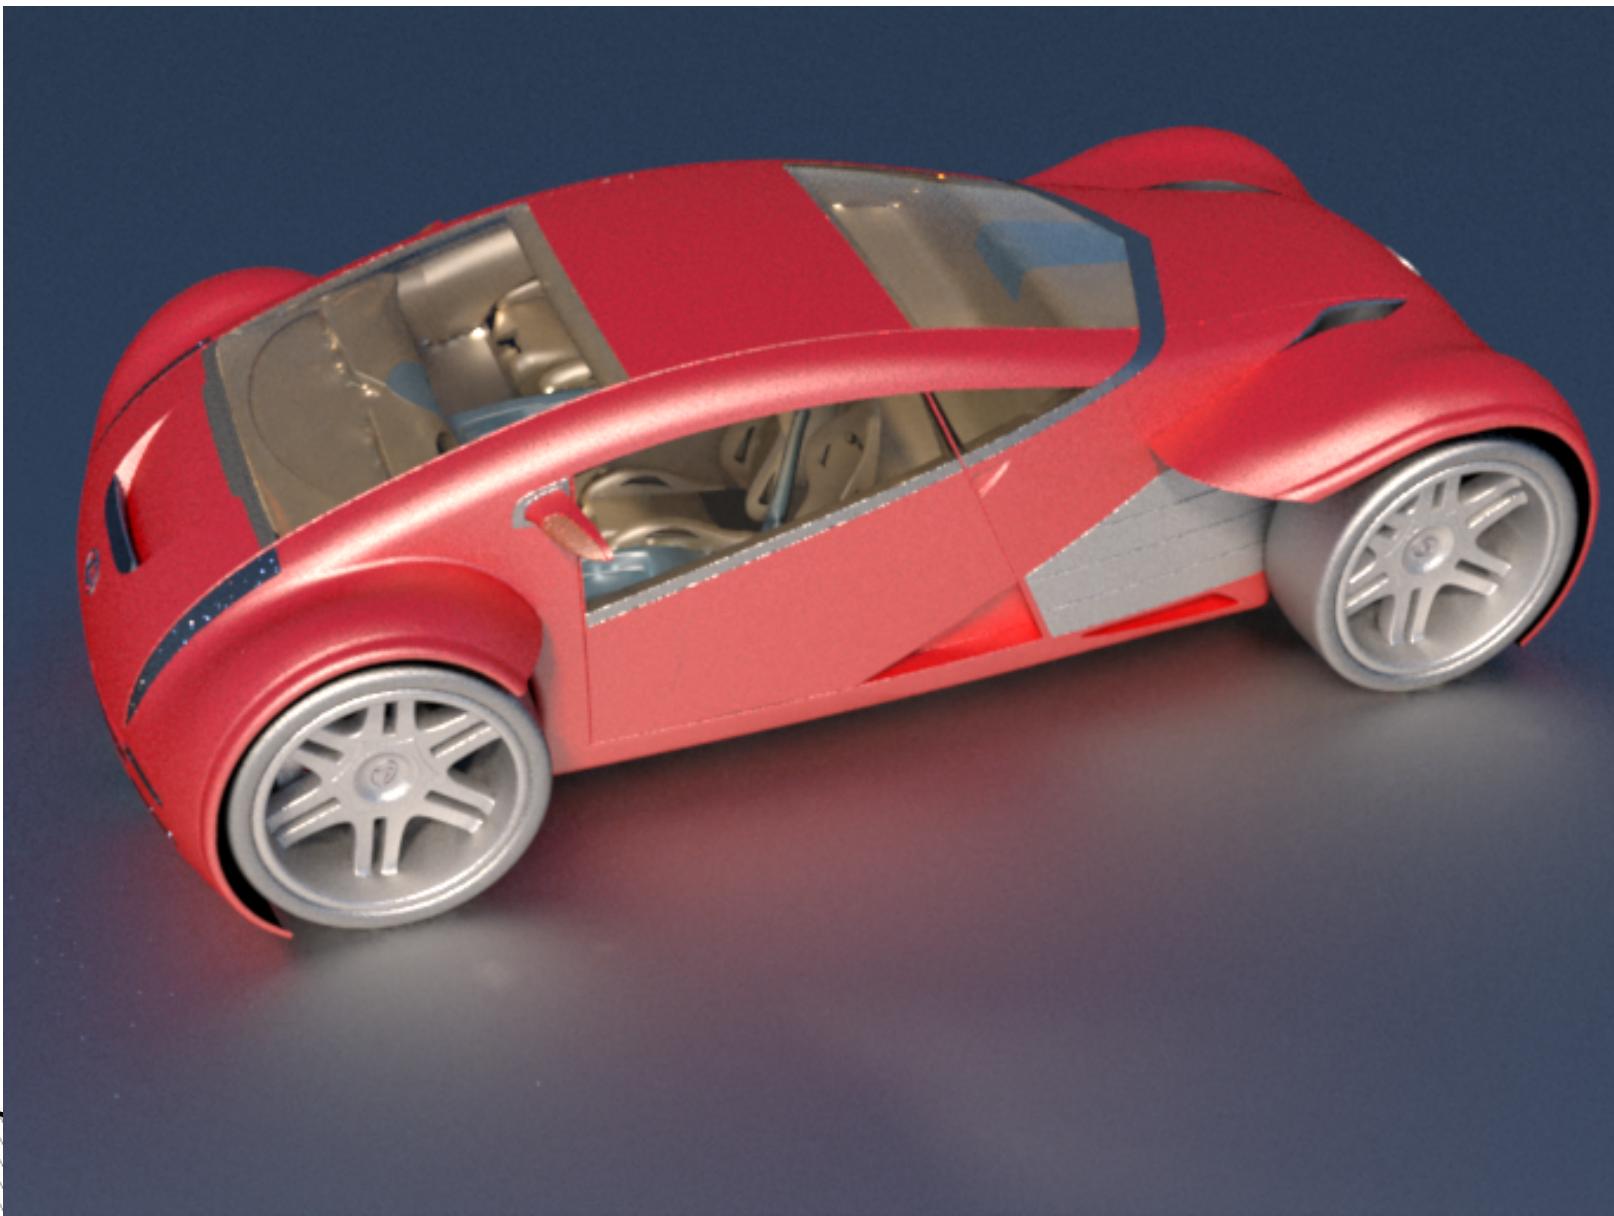
# MEMLT, 2048 mut./pixel (27 mn)



# MEMLT, 4096 mut./pixel (54 mn)



# MLT, 2048 mutations/pixel (23 mn)



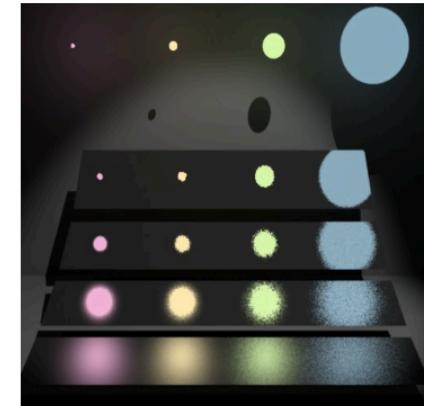
# Rendering in the movie industry

(As of August 2019; sources: Weta/Disney/Sony)

- ▶ Temporal coherence is essential:
  - ▶ Similar image for following frames
  - ▶ Not necessarily *exact*, but same bugs
- ▶ Frame *consistency* also important:
  - ▶ Same time for each frame
- ▶ So... back to path tracing:
  - ▶ Path guiding / importance resampling
  - ▶ Denoising

# Resampling Importance Sampling

- ▶ Importance sampling: guide rays based on BRDF
- ▶ Multiple Importance sampling:
  - ▶ Combine BRDF-sampled rays and light-sampled rays
  - ▶ Can be inefficient
- ▶ Importance resampling:
  - ▶ Trace a few rays
  - ▶ Get an approximate version of outgoing illumination
  - ▶ Importance sample based on that



# Resampling Importance Sampling

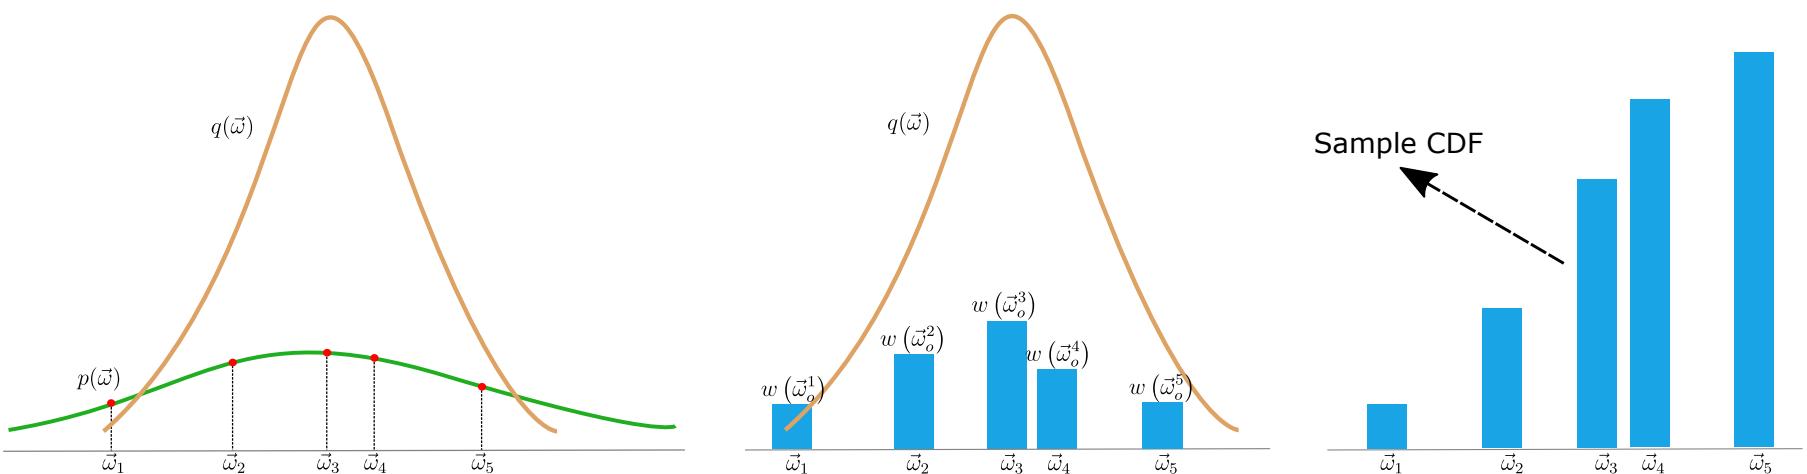


Figure 4. Example of resampling in flatland curves. To sample target function (orange curve), we first choose a function as the source function (green), and then sample the source function and get 5 outgoing direction candidates. In the second step, we evaluate the weight of each candidate which is target function divided by source function. Finally, we compute the Cumulative Distribution Function (CDF) of the weights of the candidates and use it to get the final sample.

# Resampling Importance Sampling

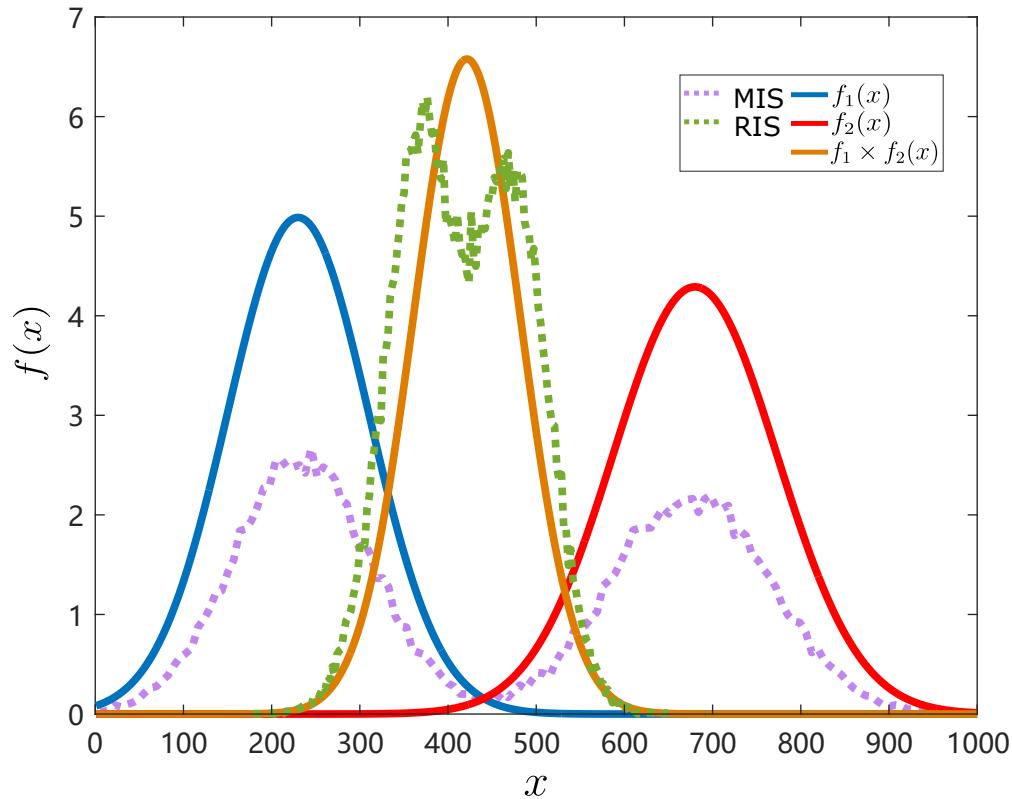


Figure 5. Comparison between MIS and RIS of two function  $f_1$  and  $f_2$  product. The RIS produces result much closer to the target function (product of two functions) than MIS.

# Path guiding

- ▶ Slides from Thomas Müller, “Practical path guiding in production”, 2019.

# 'Practical Path Guiding' in Production

Thomas Müller

Affiliation:  NVIDIA.

Work done while  ETH Zürich





1024 spp

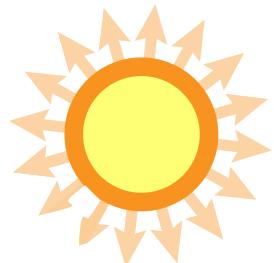
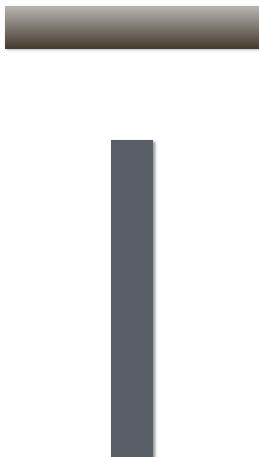
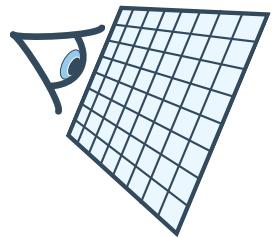
# Path tracing

122

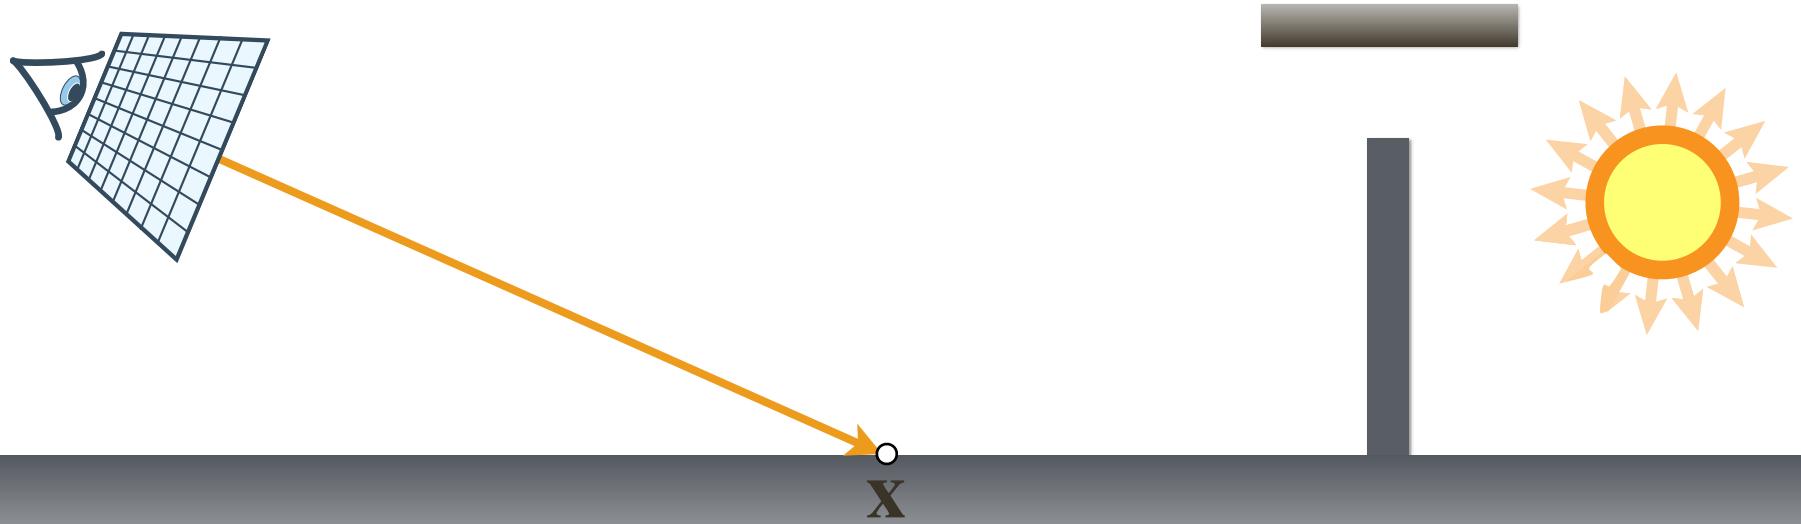


1024 spp "Practical path guiding" with improvements 123

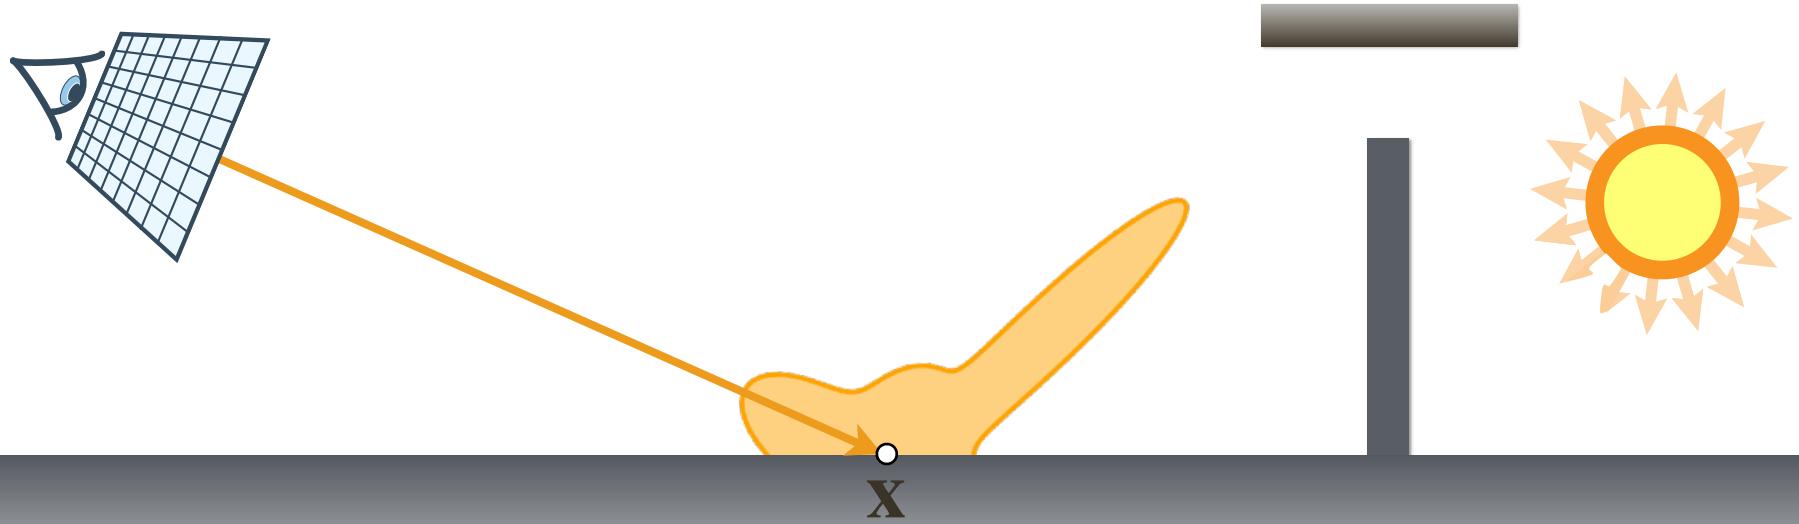
# Path guiding



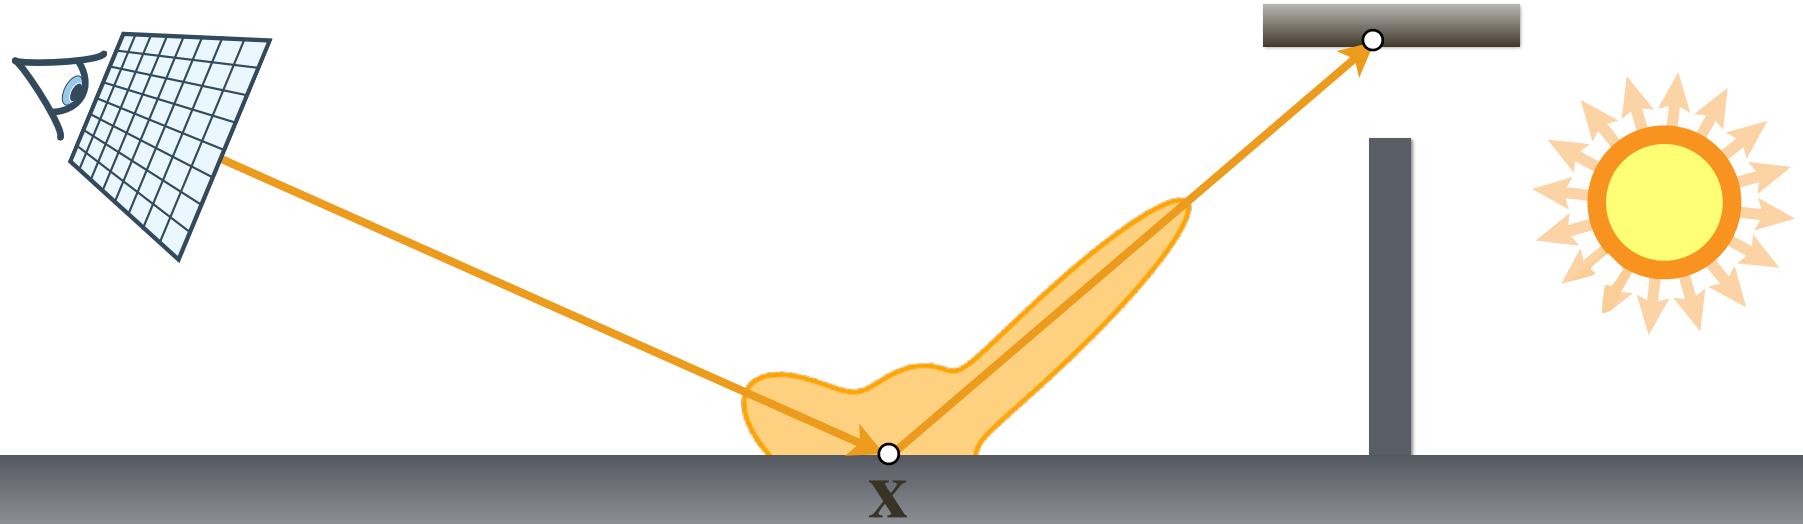
# Path guiding



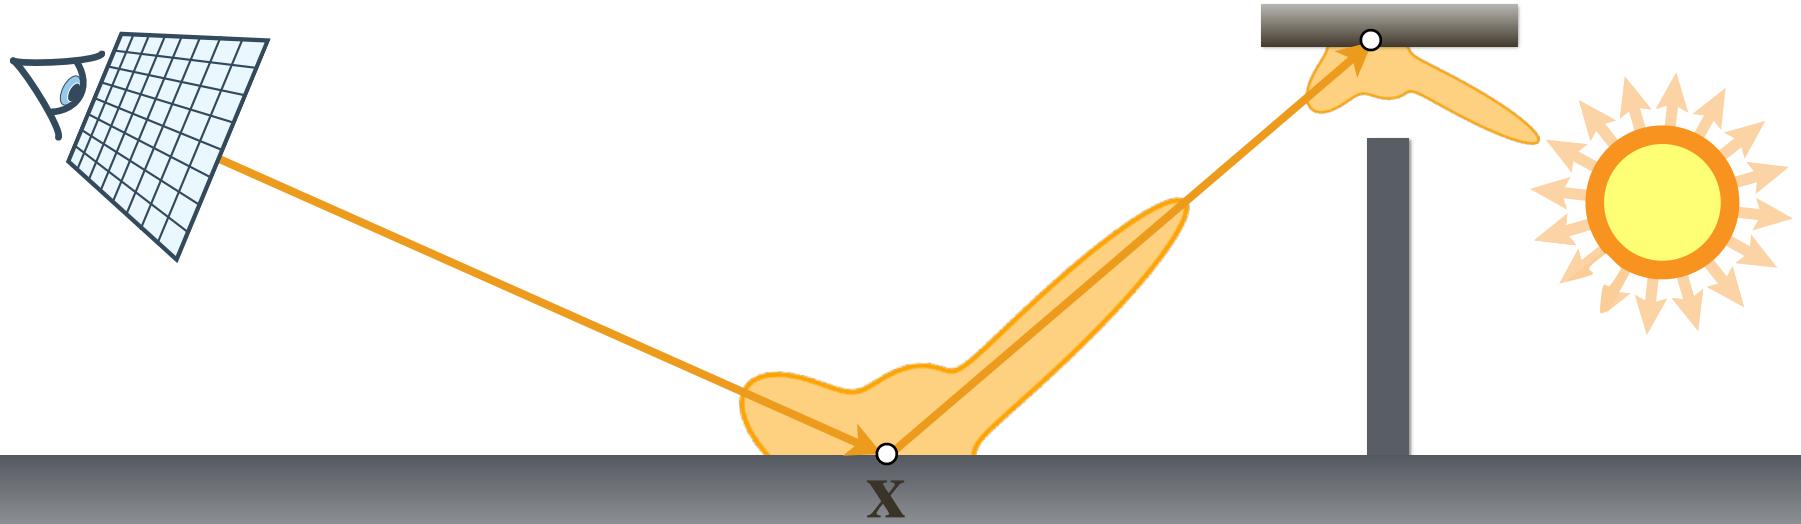
# Path guiding



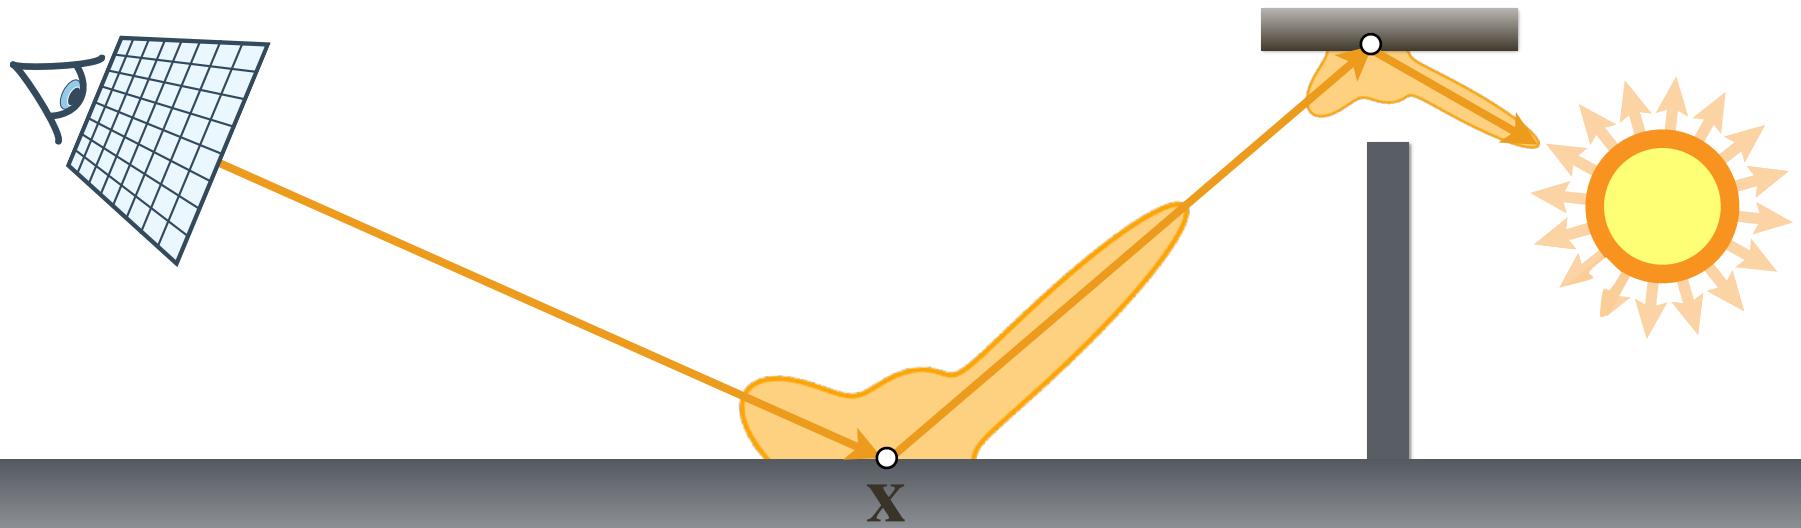
# Path guiding



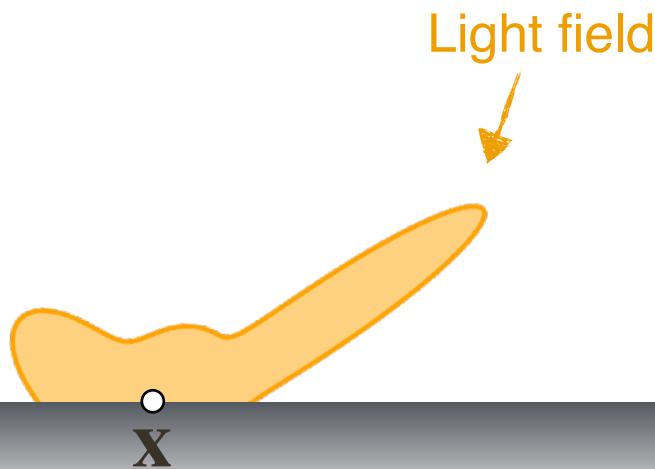
# Path guiding



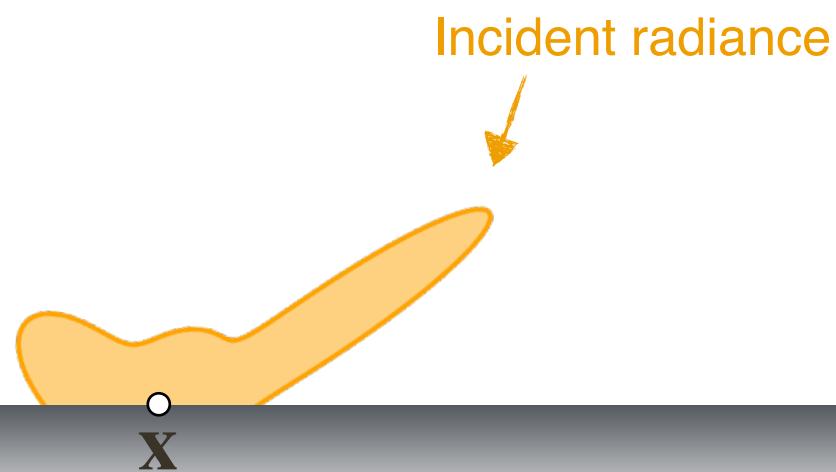
# Path guiding



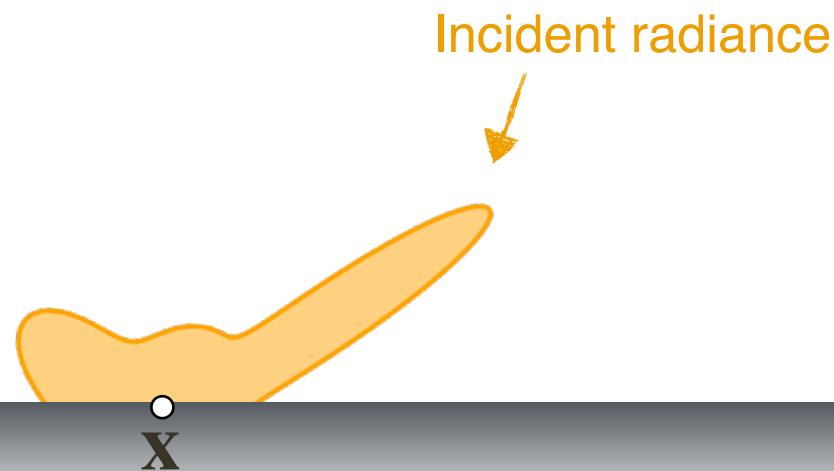
The goal is to learn the  $5D = 3D + 2D$  light field



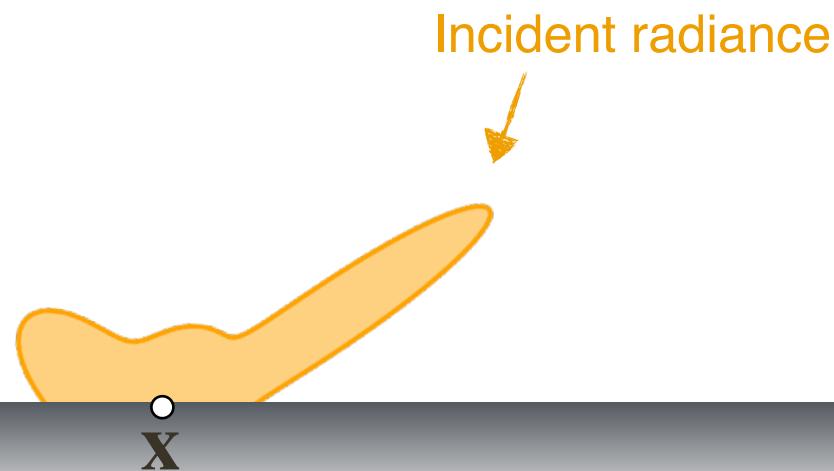
The goal is to learn the  $5D = 3D + 2D$  light field



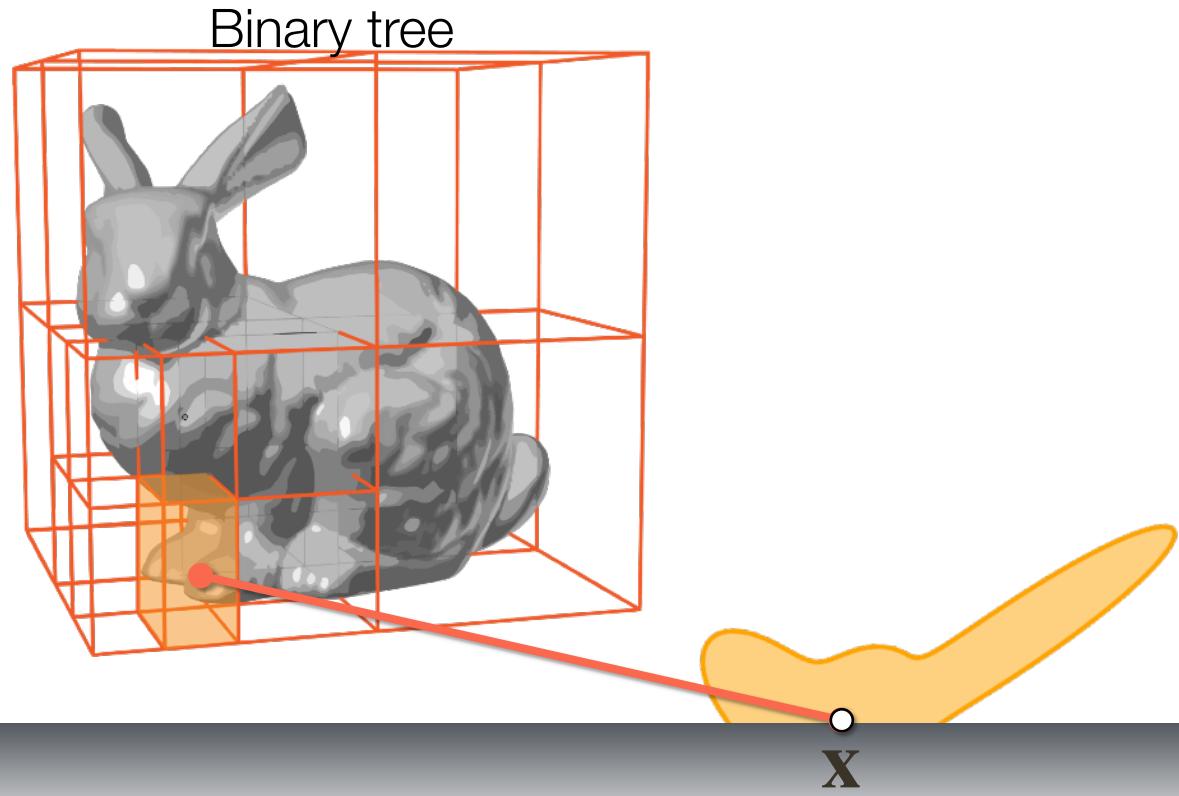
The goal is to learn the  $5D = 3D + 2D$  light field



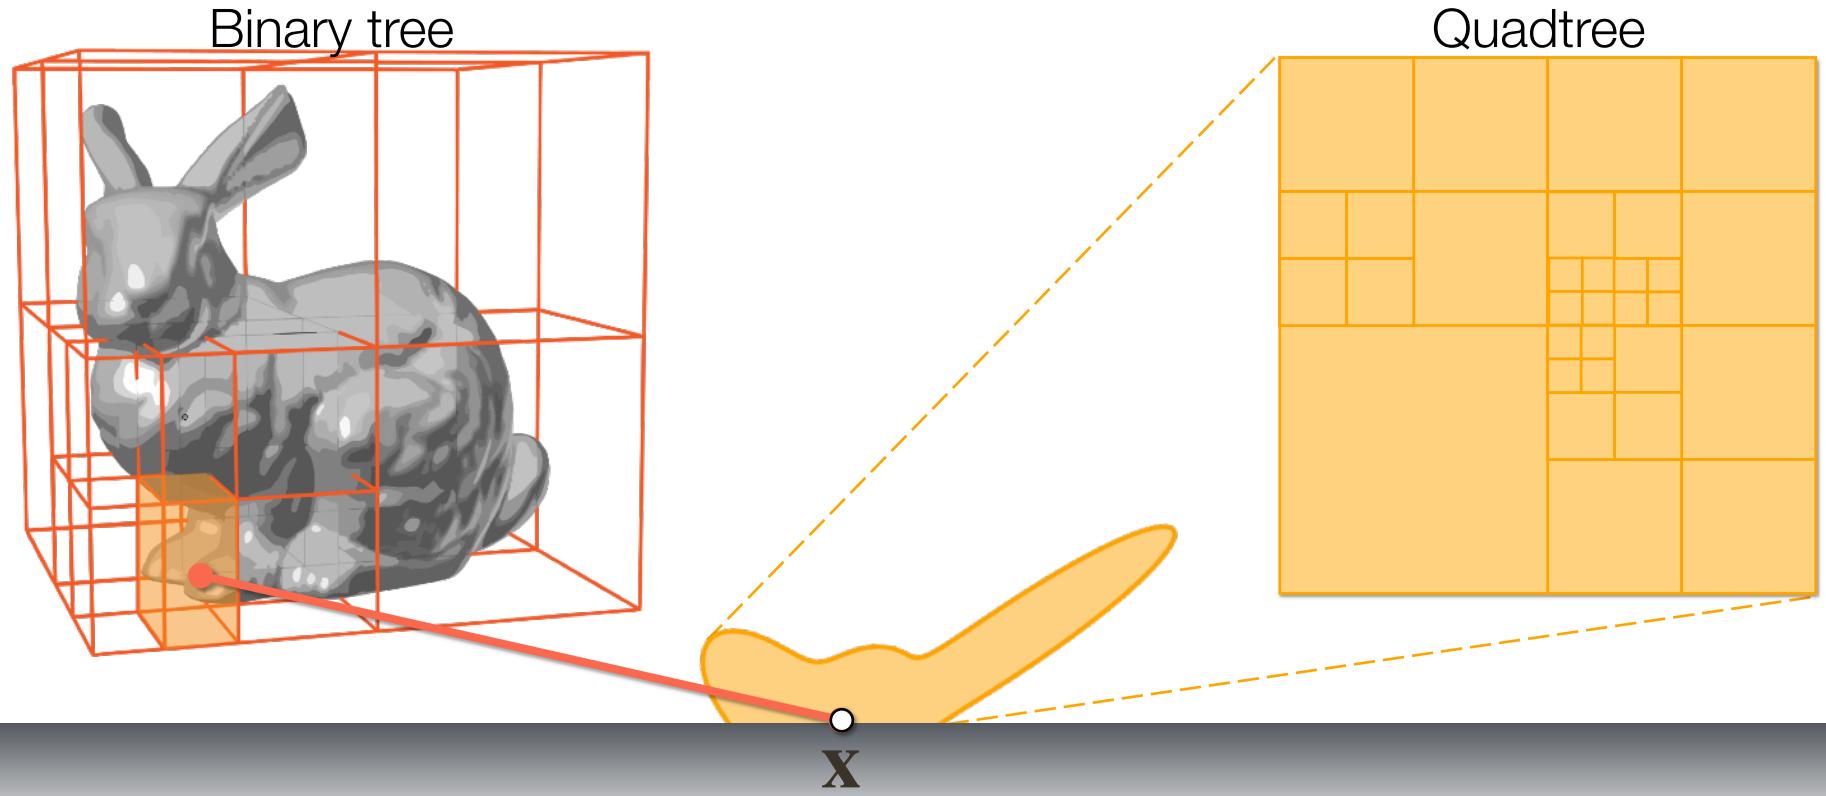
The goal is to learn the  $5D = 3D + 2D$  light field



# Representing the light field in a 3D + 2D tree



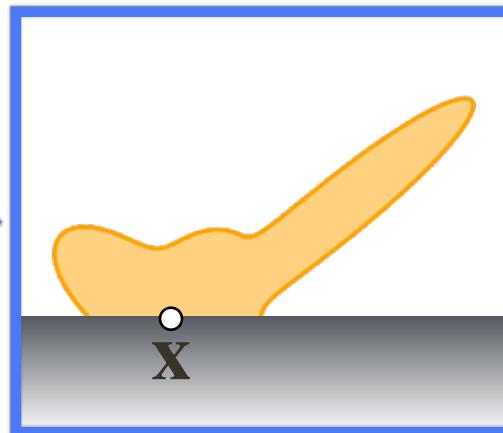
# Representing the light field in a 3D + 2D tree



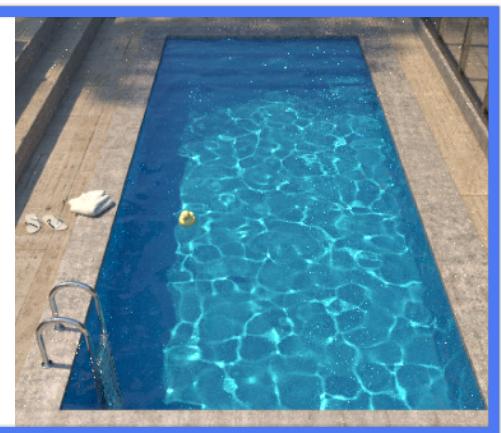
# Iterative learning during rendering



Load scene



Learn incident radiance



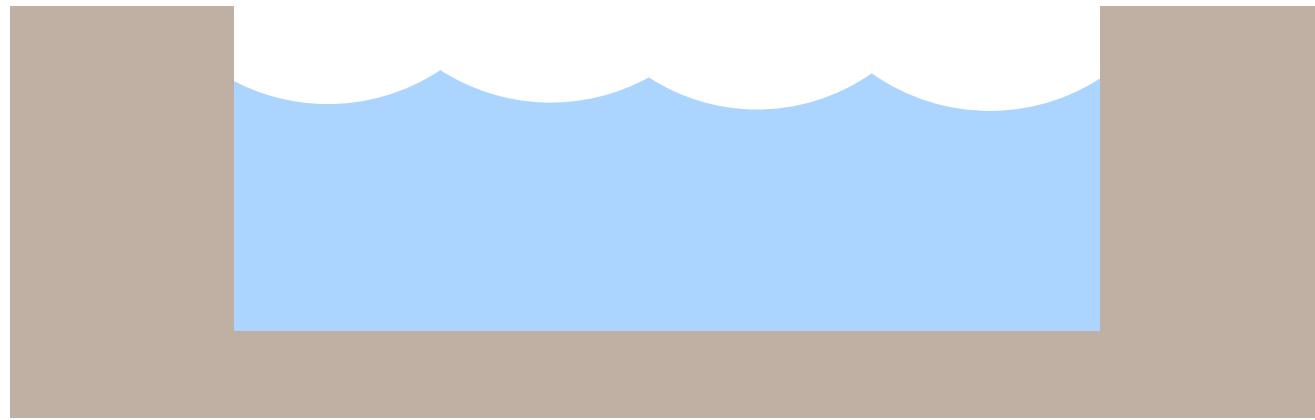
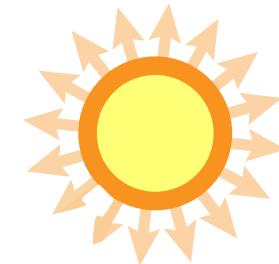
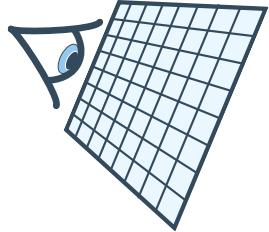
Render image



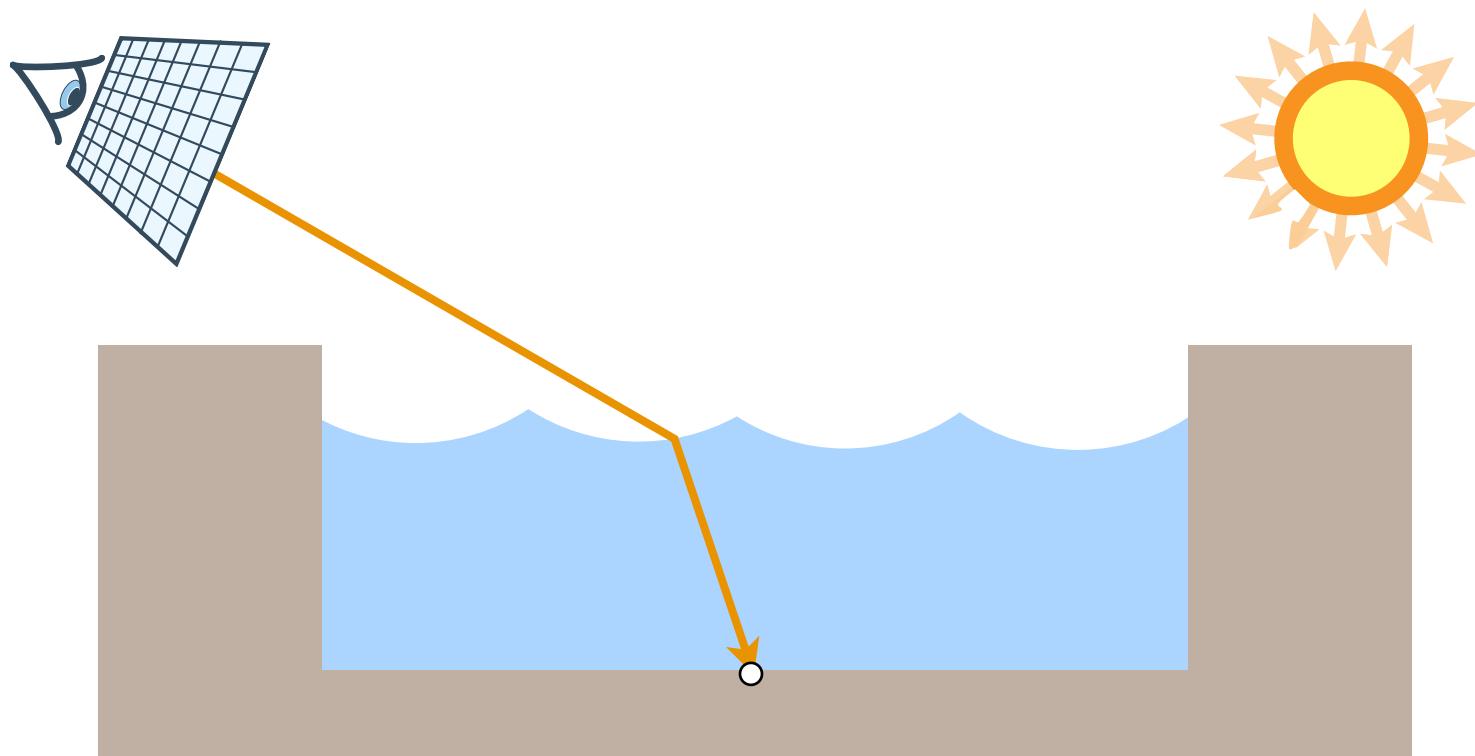
# Iterative learning during rendering



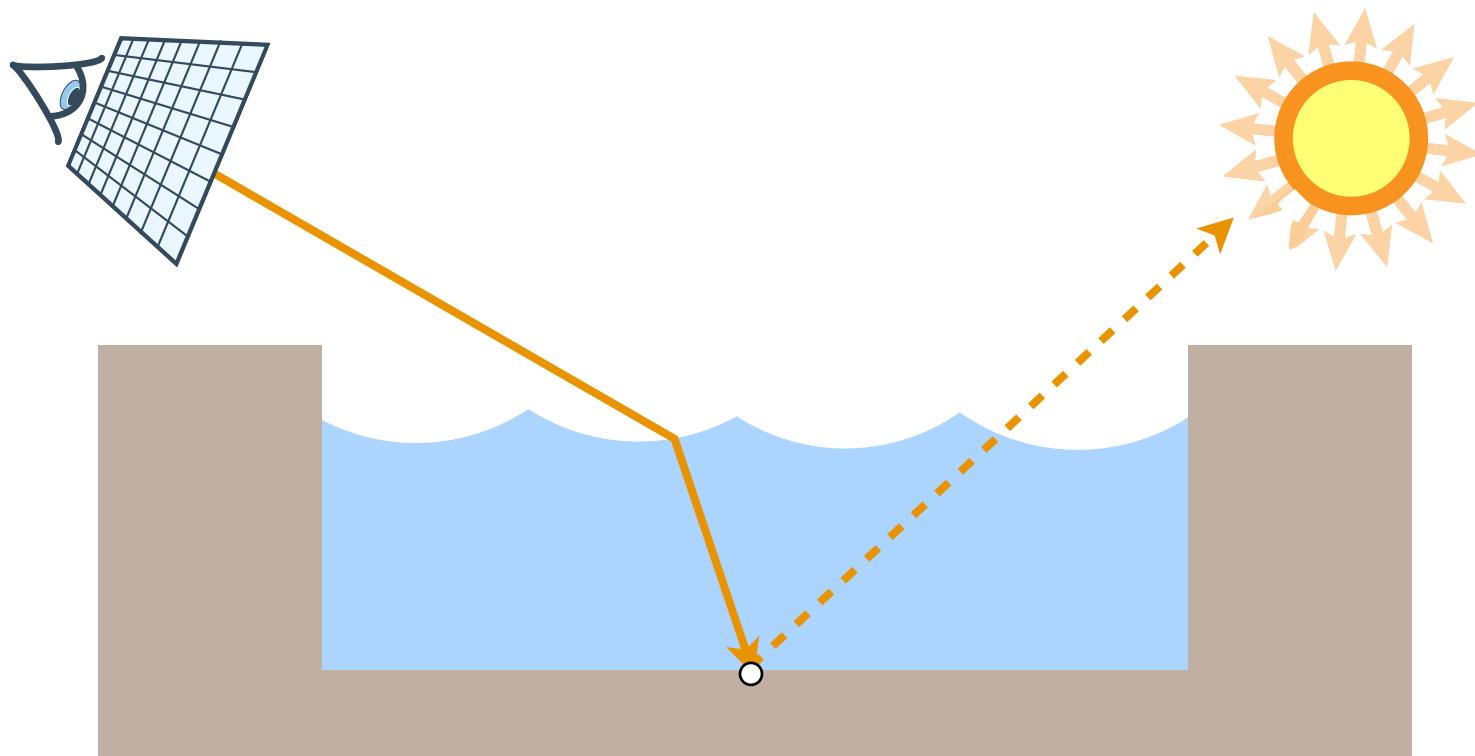
# Iterative learning during rendering



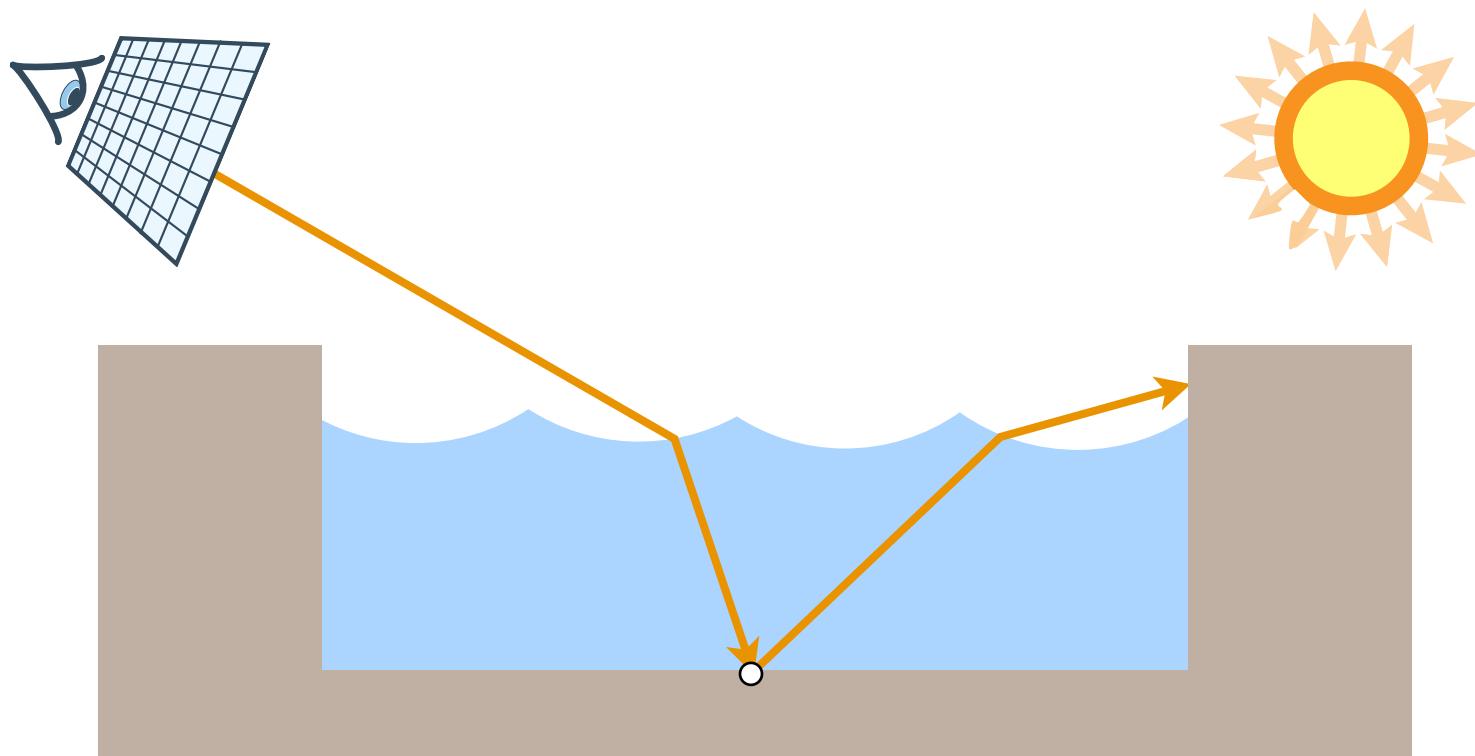
# Iterative learning during rendering



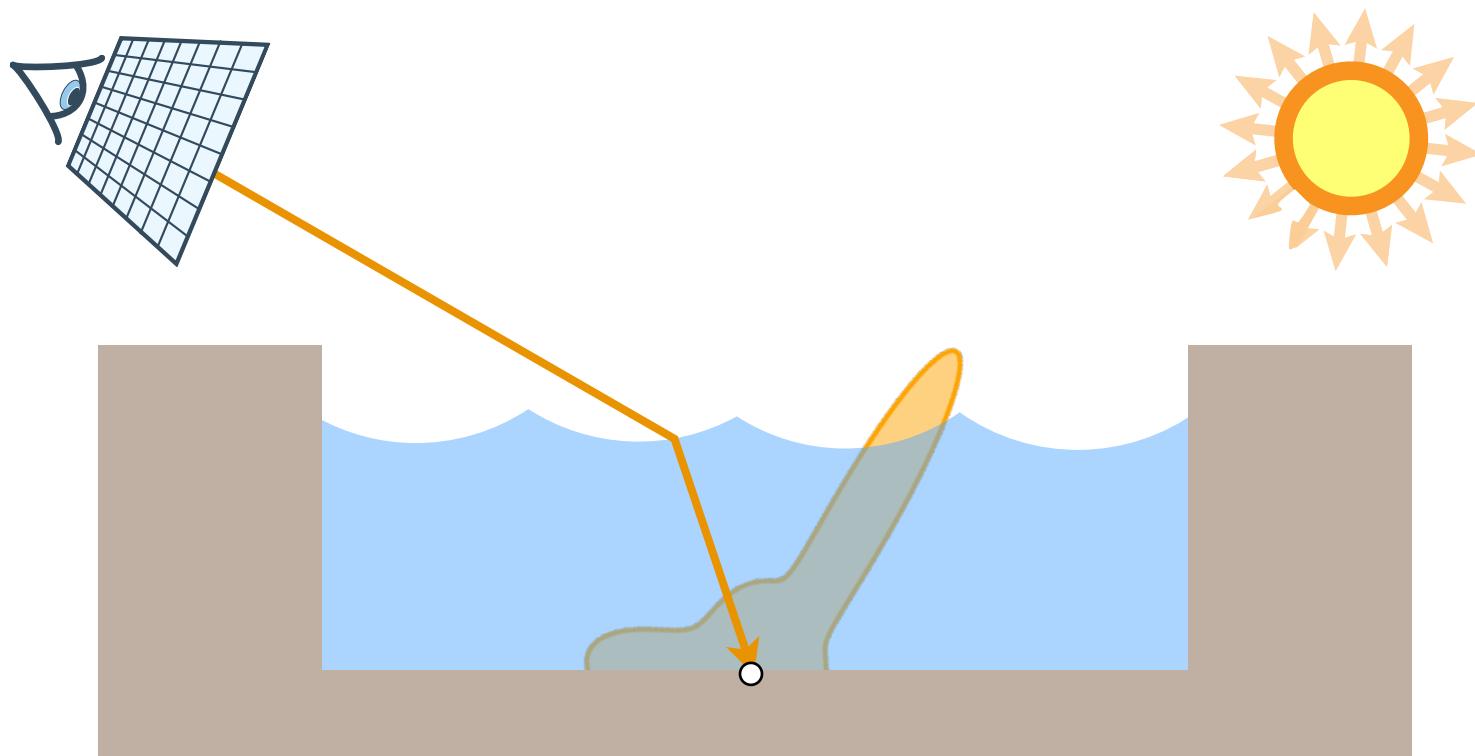
# Iterative learning during rendering



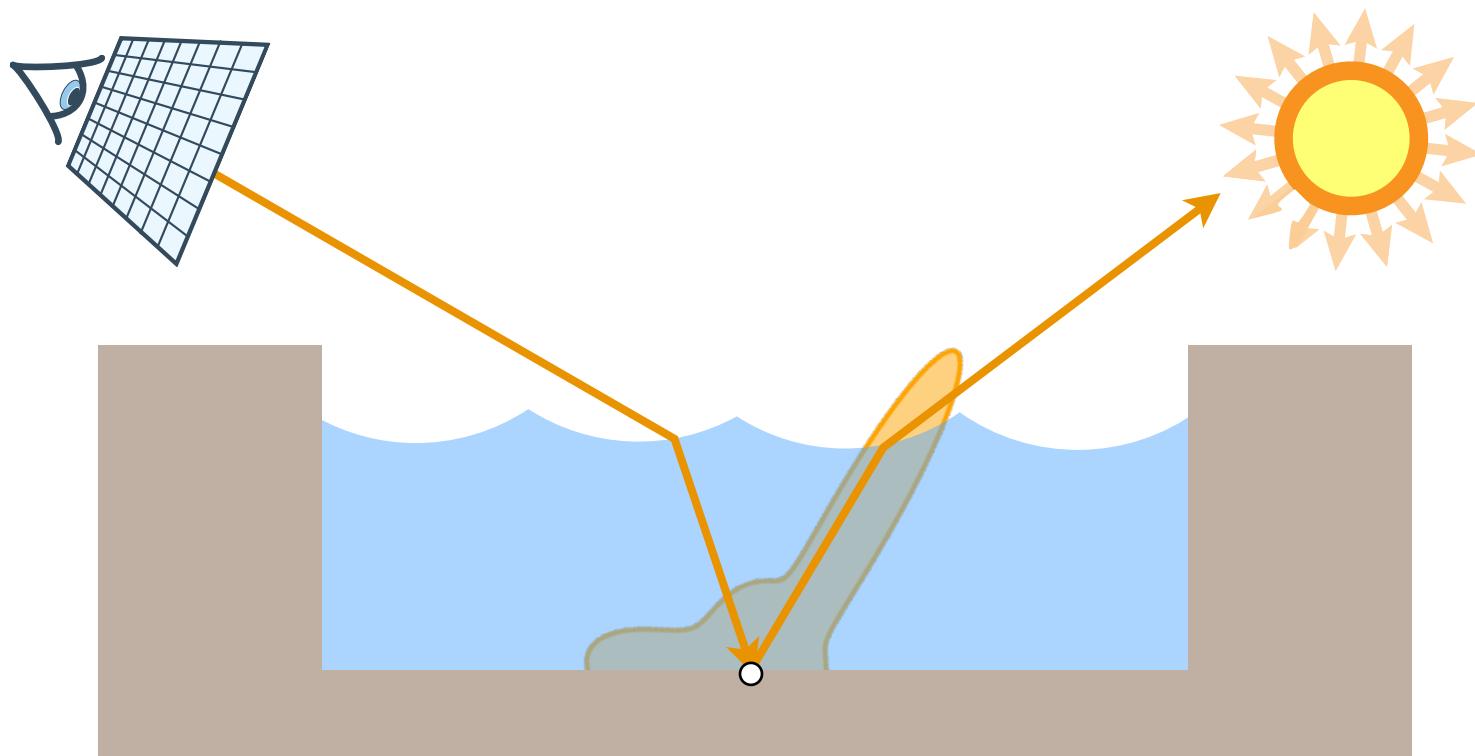
# Iterative learning during rendering



# Iterative learning during rendering

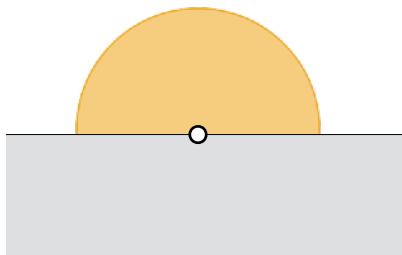


# Iterative learning during rendering

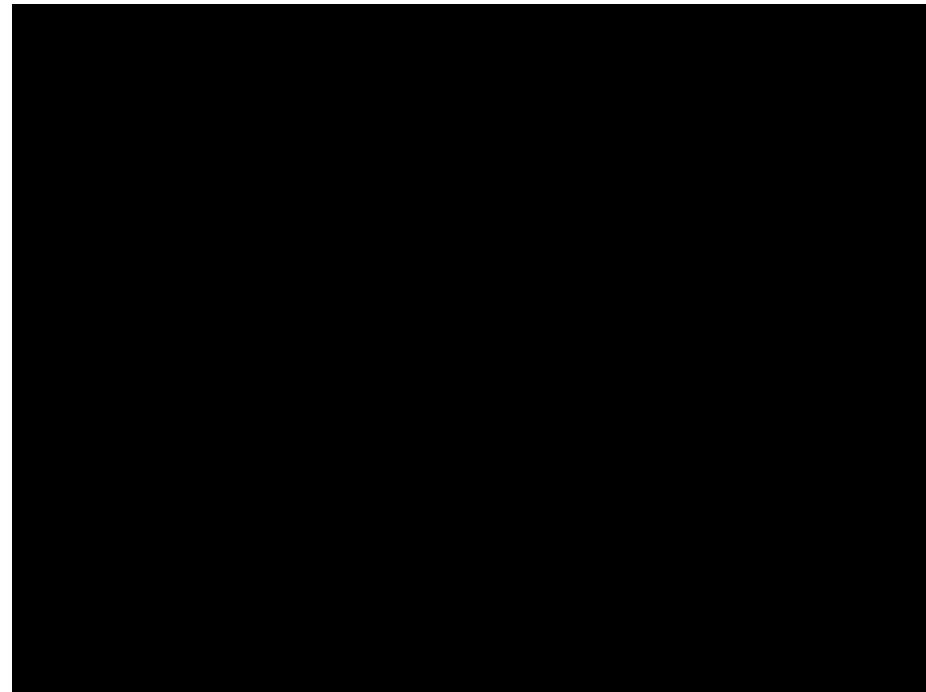


# Iterative learning during rendering

Illustration  
of data  
structure

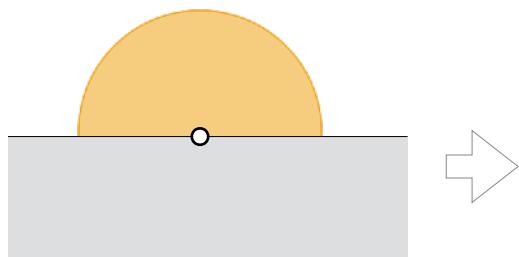


Iteration 1

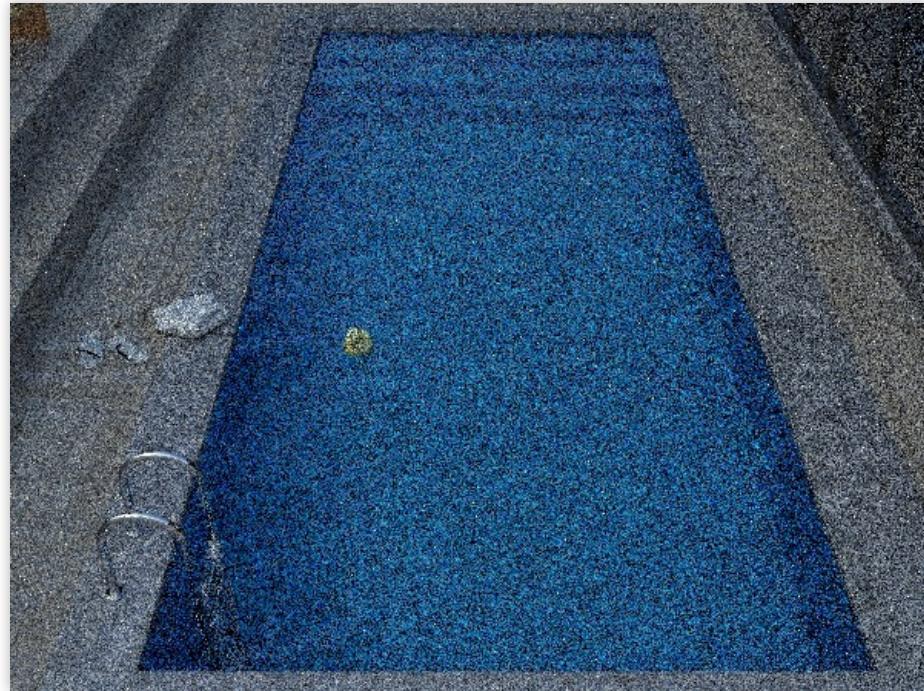


# Iterative learning during rendering

Illustration  
of data  
structure

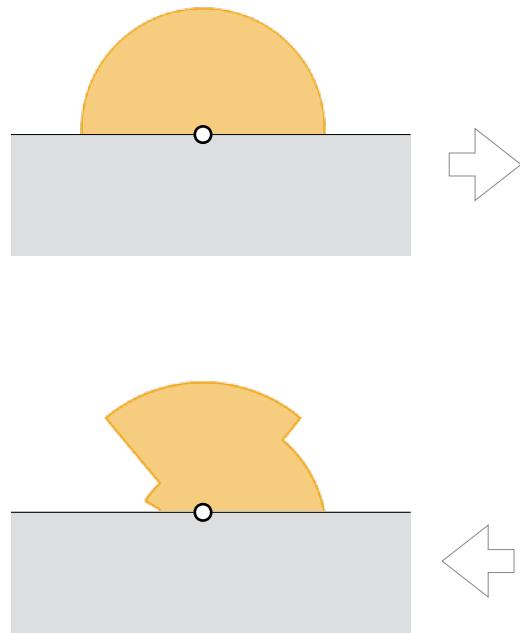


Iteration 1



# Iterative learning during rendering

Illustration  
of data  
structure

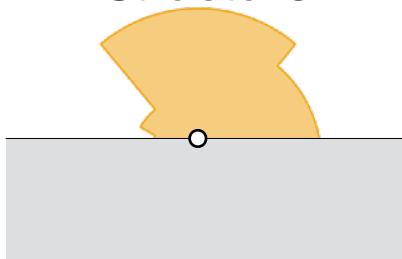


Iteration 1

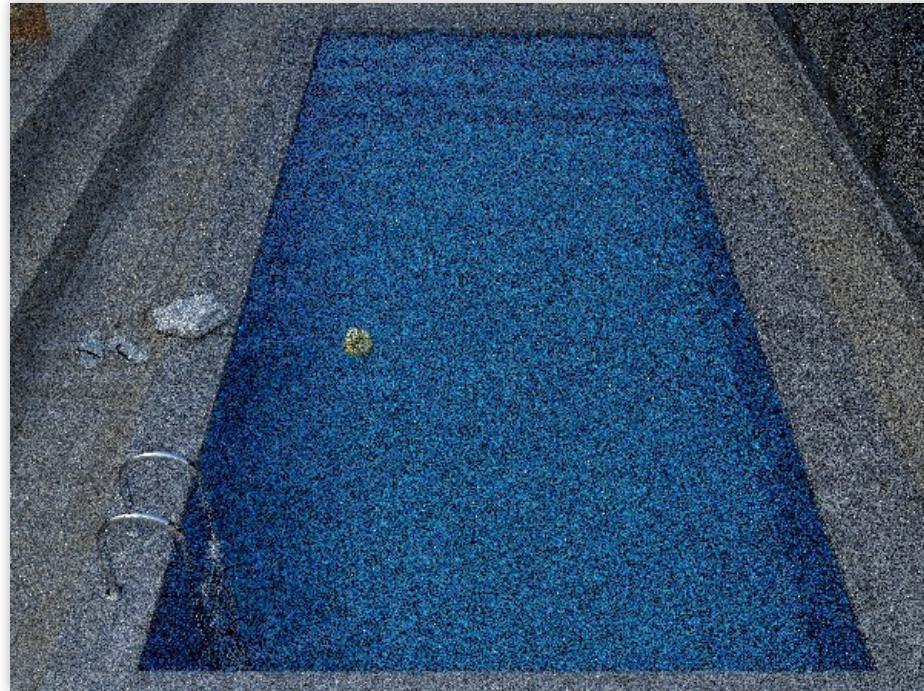


# Iterative learning during rendering

Illustration  
of data  
structure

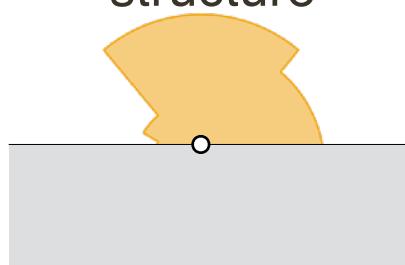


Iteration 2

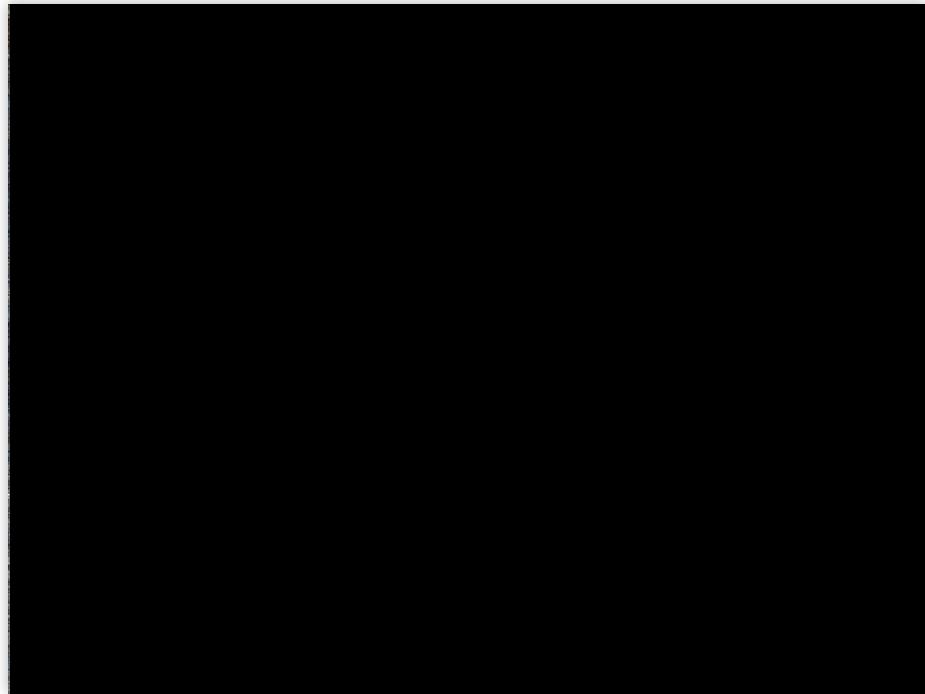


# Iterative learning during rendering

Illustration  
of data  
structure

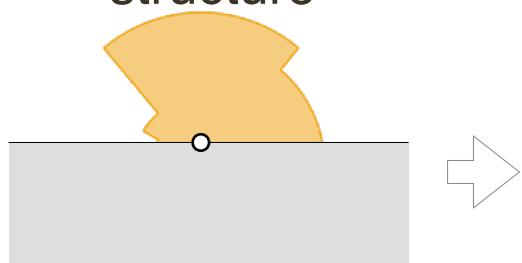


Iteration 2



# Iterative learning during rendering

Illustration  
of data  
structure

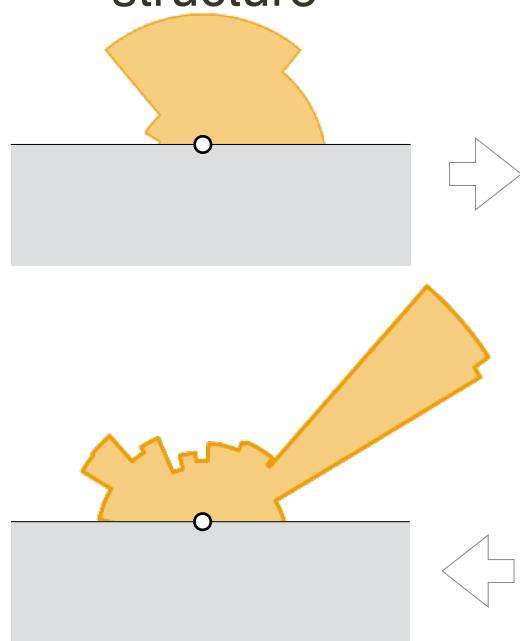


Iteration 2



# Iterative learning during rendering

Illustration  
of data  
structure

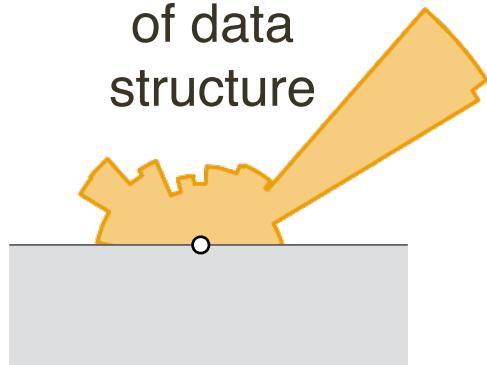


Iteration 2



# Iterative learning during rendering

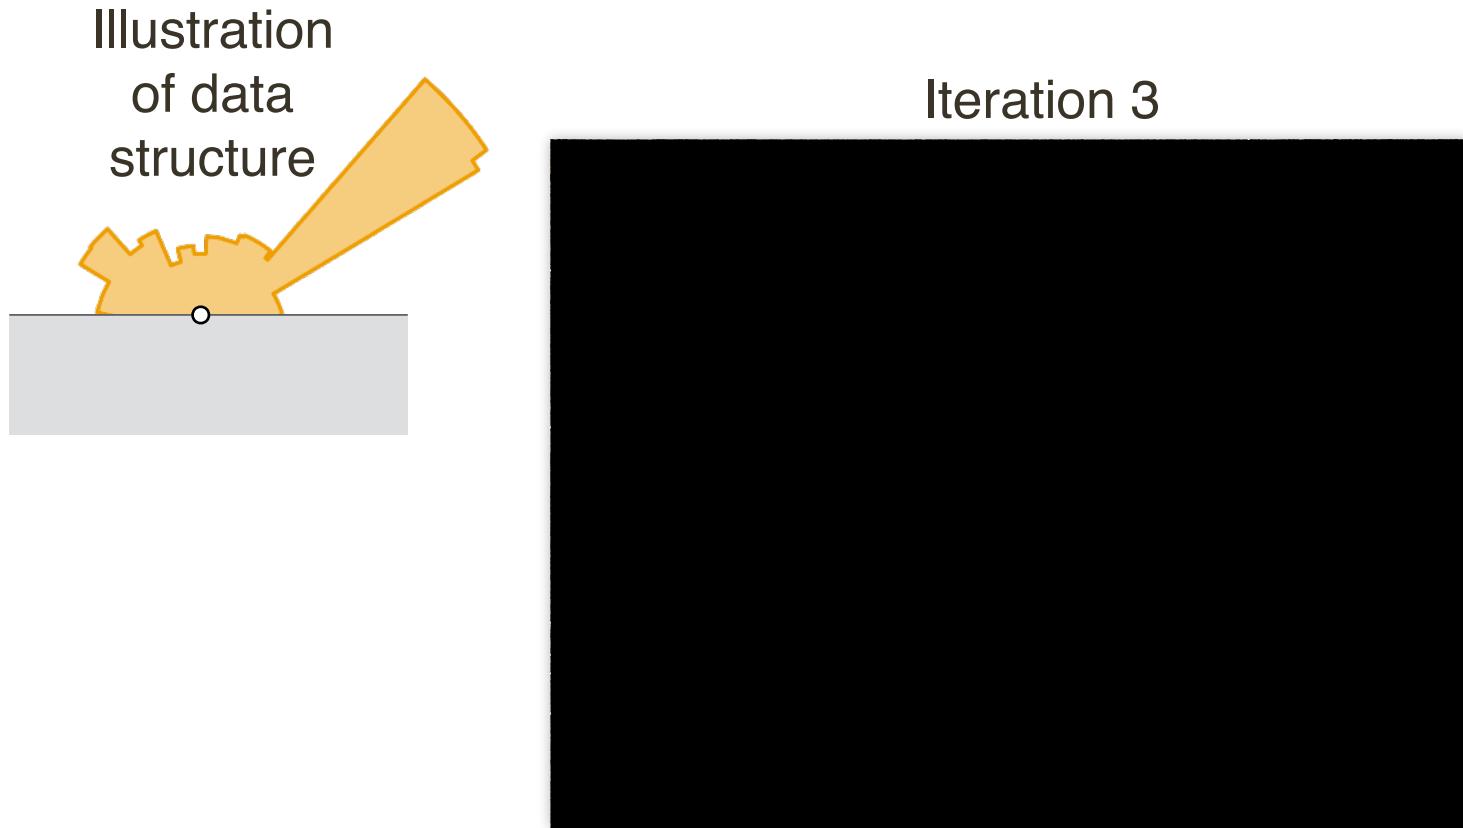
Illustration  
of data  
structure



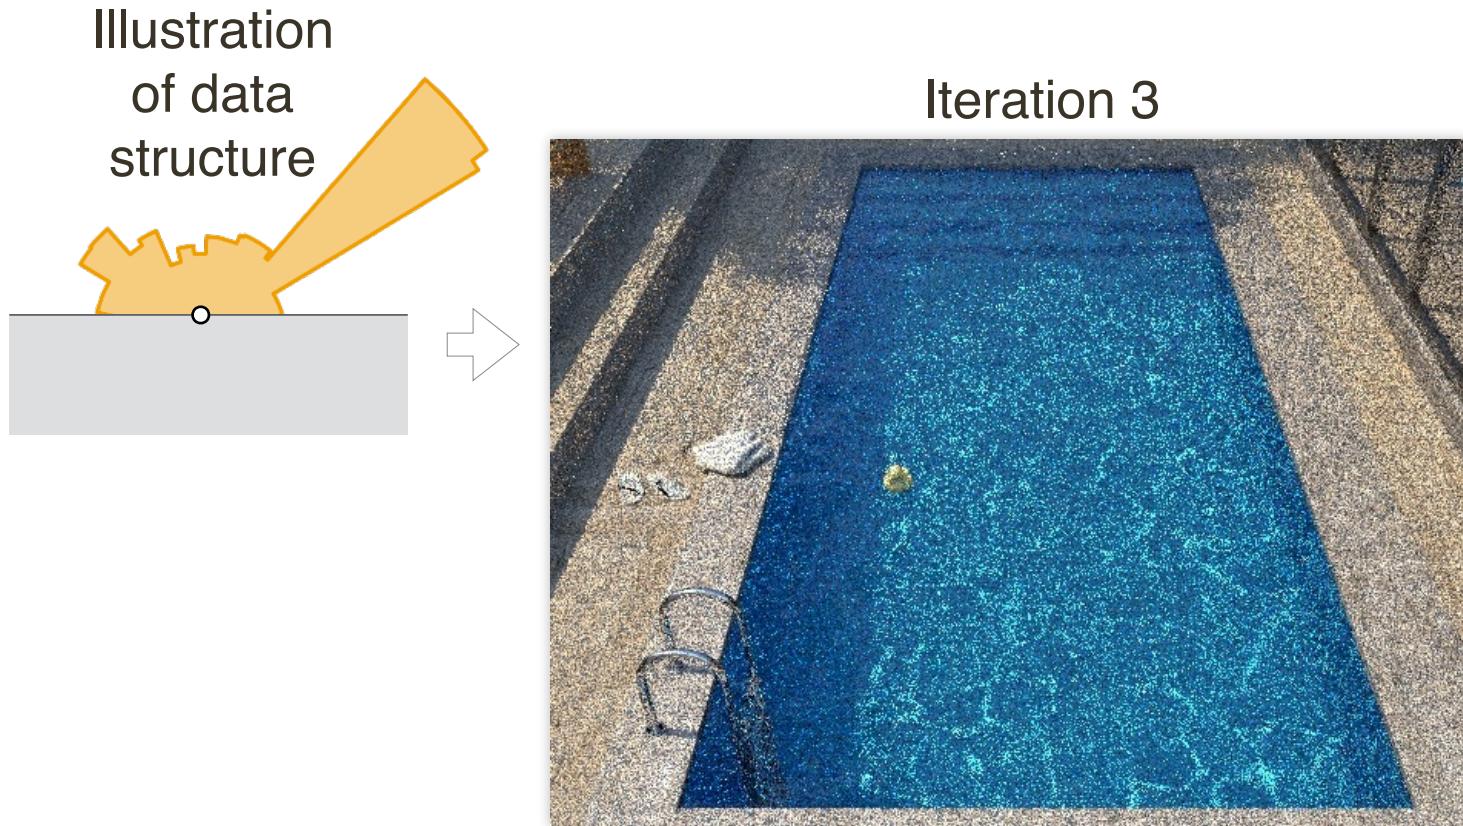
Iteration 3



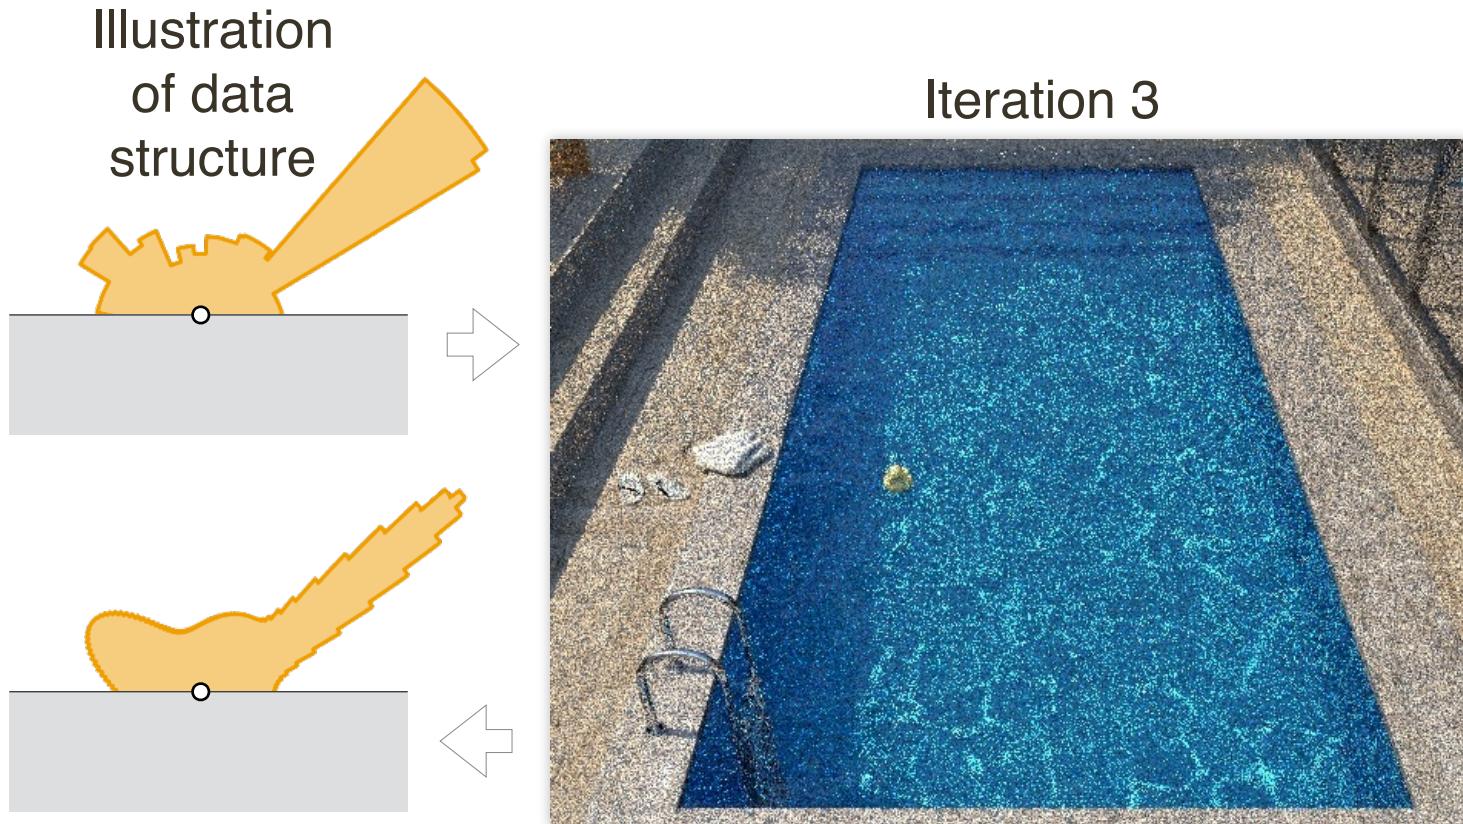
# Iterative learning during rendering



# Iterative learning during rendering

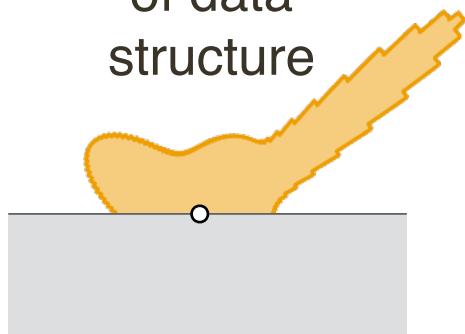


# Iterative learning during rendering

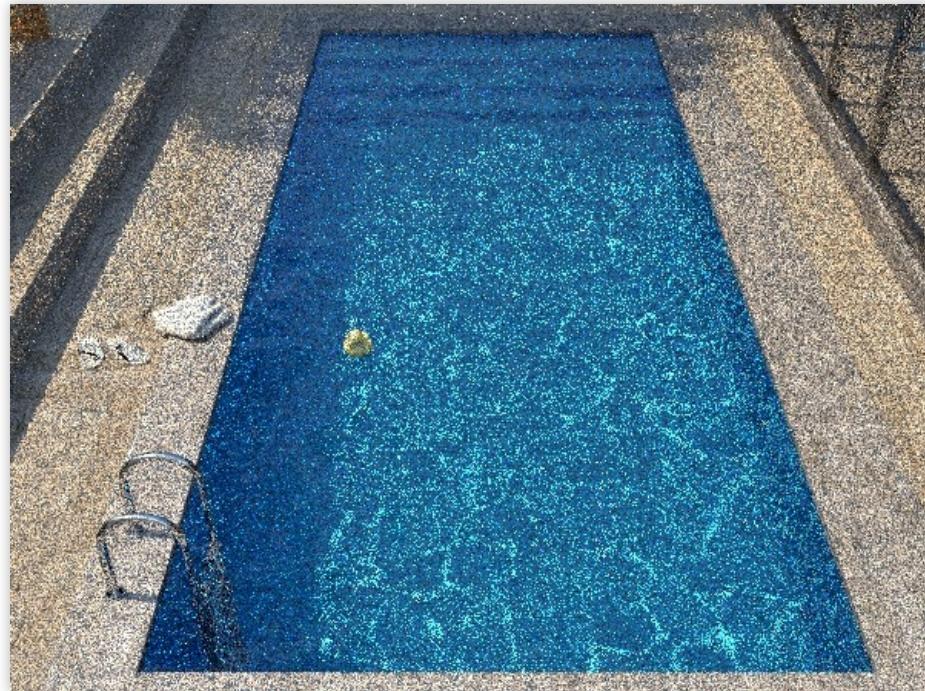


# Iterative learning during rendering

Illustration  
of data  
structure

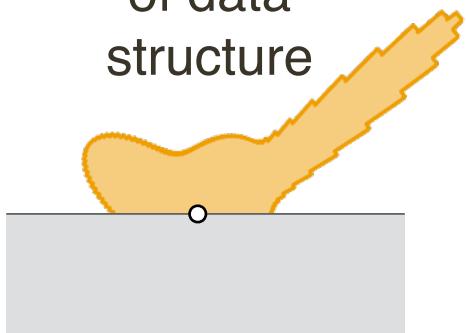


Iteration 4

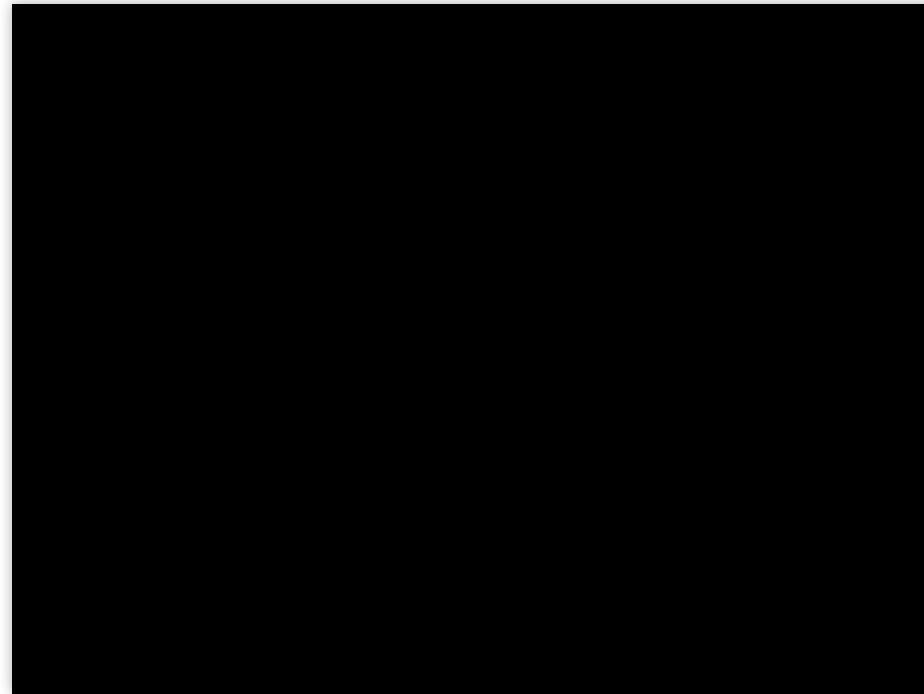


# Iterative learning during rendering

Illustration  
of data  
structure

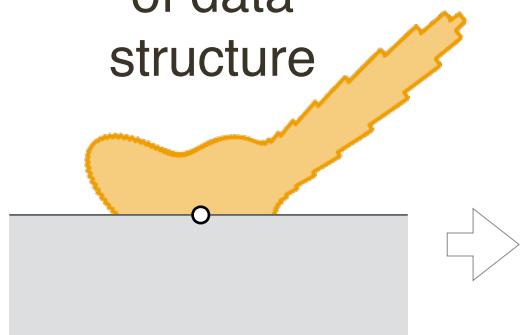


Iteration 4

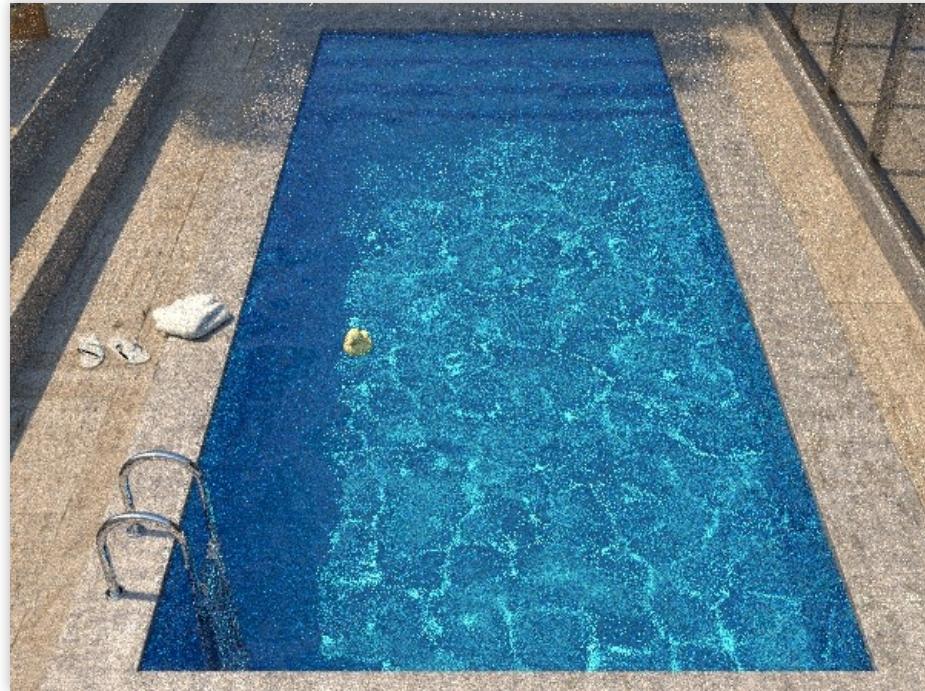


# Iterative learning during rendering

Illustration  
of data  
structure

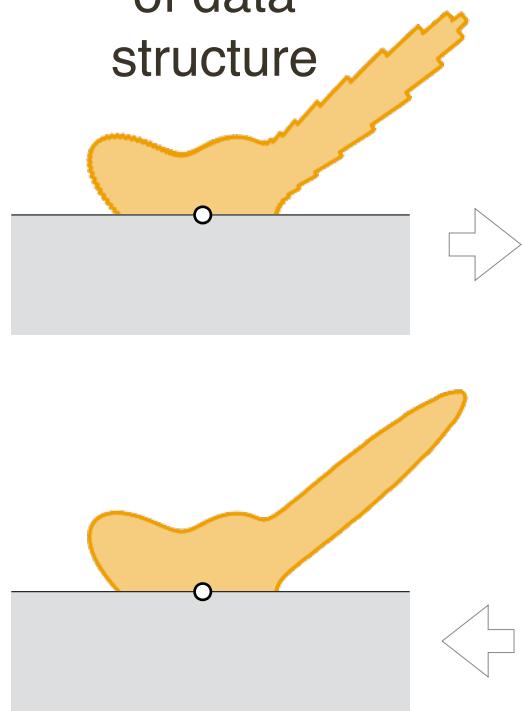


Iteration 4

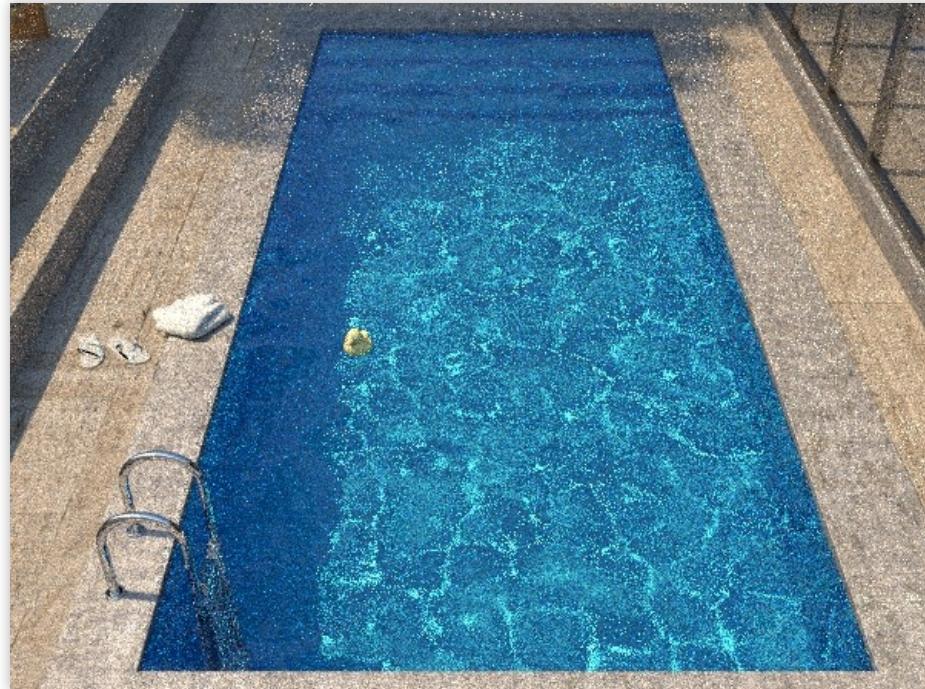


# Iterative learning during rendering

Illustration  
of data  
structure



Iteration 4





1024 spp

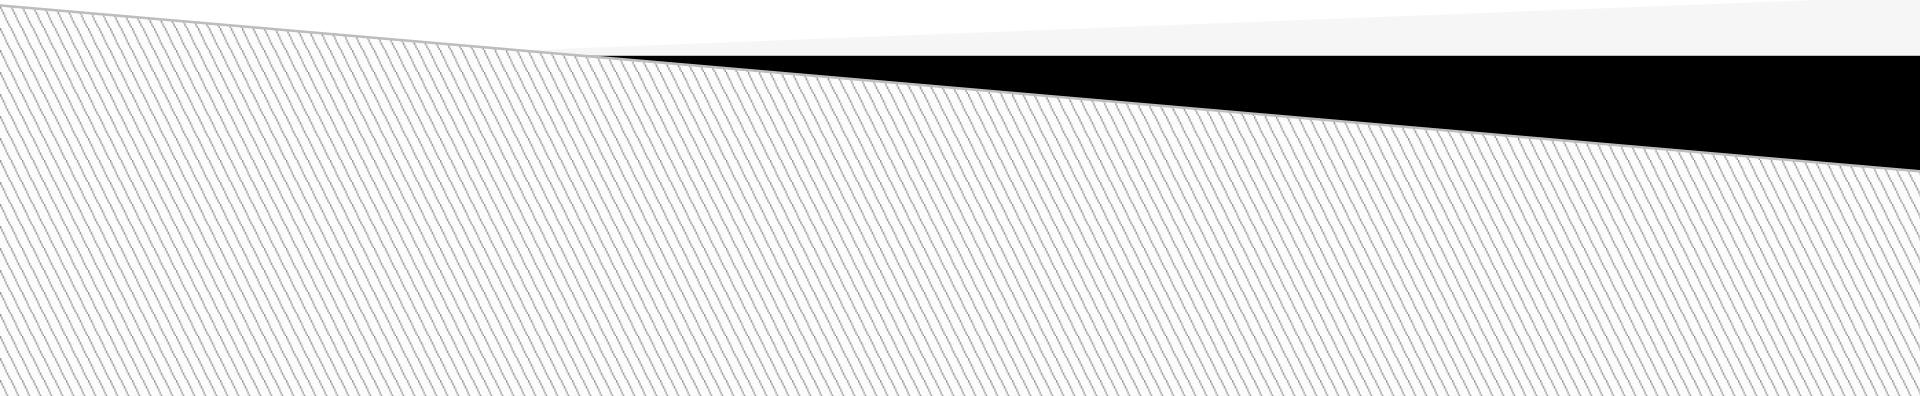
# Path tracing

135



1024 spp "Practical path guiding" with improvements 136

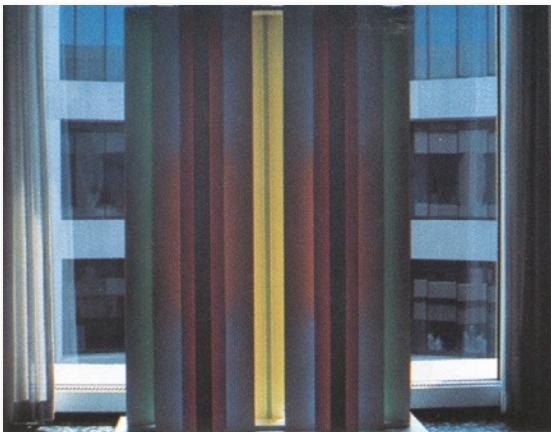
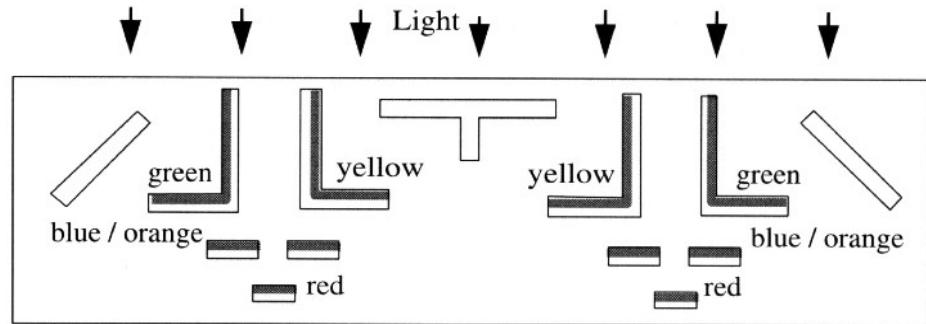
# Finite-Element Methods



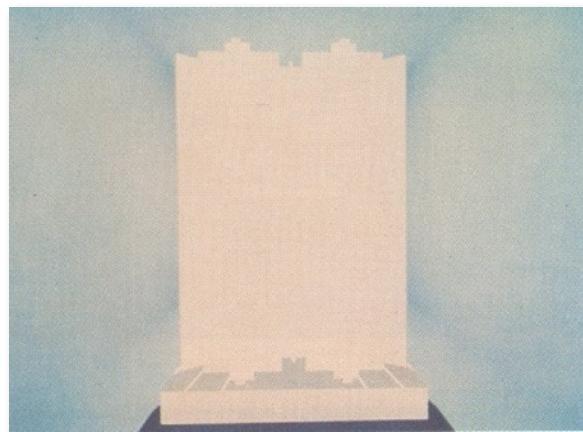
# Radiosity

- ▶ Taking into account all inter-reflections

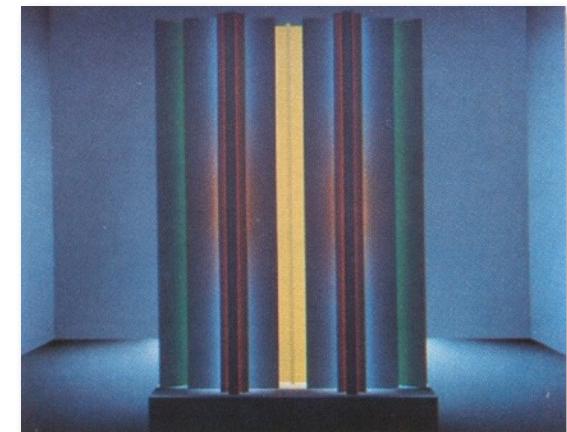
Sculpture by John Ferren



Photo



Ray-tracing



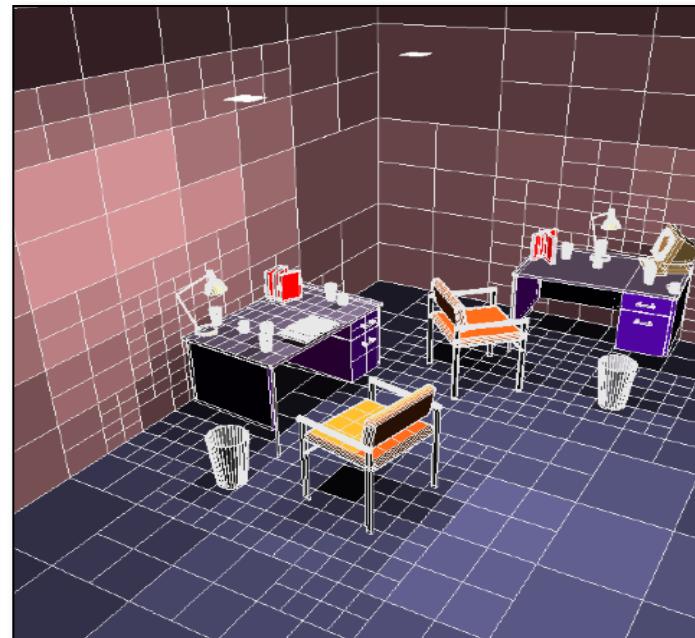
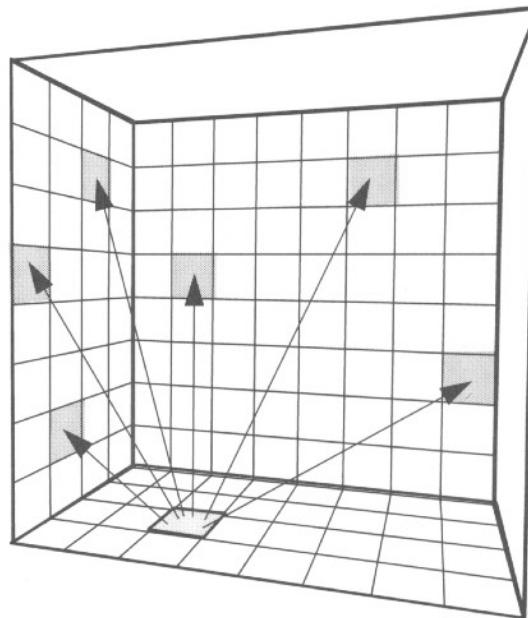
Radiosité

# Radiosity methods [1984]

- ▶ Hypothesis: diffuse materials
- ▶ Radiance, BRDF... are independent from the direction
  - ⇒ Simplifies the rendering equation
- ▶ Radiosity method:
  - Discretize this equation in object space
    - (viewpoint independent)
  - Solve the discretized equation
  - Render the scene using the illumination

# Radiosity Equation

- ▶ Discretize scene into patches



# Simplification and discretization

$$L(x,d) = E(x,d) + \int \rho(x,d,d') v(x,x') G(x,x') dA$$

► Simplified:

$$B(x) = E(x) + \rho_x \int B(x') v(x,x') G(x,x') dA$$

► Discrete version: Form factor

$$B_i = E_i + \rho_i \sum_j F_{ji} B_j A_j / A_i$$

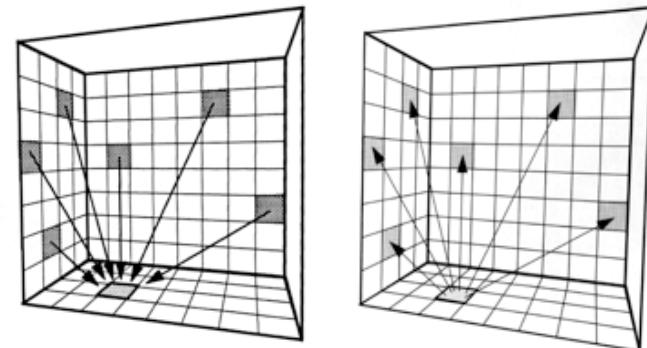
- $B_i$ ,  $B_j$ : radiosity for patches i et j (in W/m<sup>2</sup>)
- $E_i$ : emission for patch i
- $F_{ji}$  form factor, characterizes proportion of energy leaving patch j arriving on i
- $A_i$  et  $A_j$  : areas for patches i and j

# Matrix representation

- ▶ Grouping together all elements:

$$\begin{pmatrix} B_0 \\ B_n \end{pmatrix} = \begin{pmatrix} E_0 \\ E_n \end{pmatrix} + \left( \rho_i F_{ji} \right) \begin{pmatrix} B_0 \\ B_n \end{pmatrix} \Leftrightarrow B = E + MB$$

- ▶ Matrix equation, to solve iteratively
  - Relaxation methods ( gathering / shooting )



# Form factor

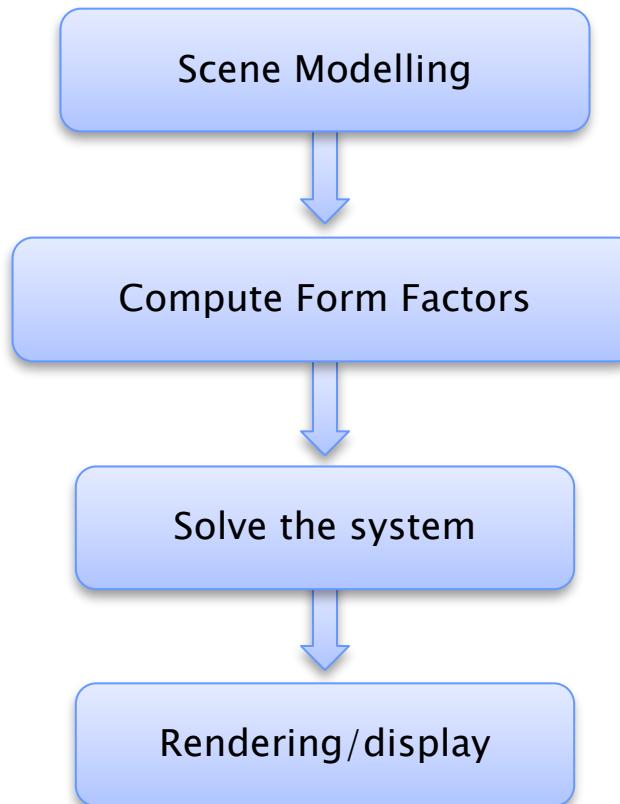
- ▶ Form factor  $F_{ij}$  from a patch  $A_i$  towards a patch  $A_j$  :

$$F_{ij} = \int_{A_i} \int_{A_j} v(x, x') \frac{\cos(\theta) \cos(\theta')}{\pi r^2} dx dx'$$

- ▶ Problem: computing this integral (4D). No analytical solution
  - ⇒ Approximated solutions: projection on a hemisphere or a hemicube.

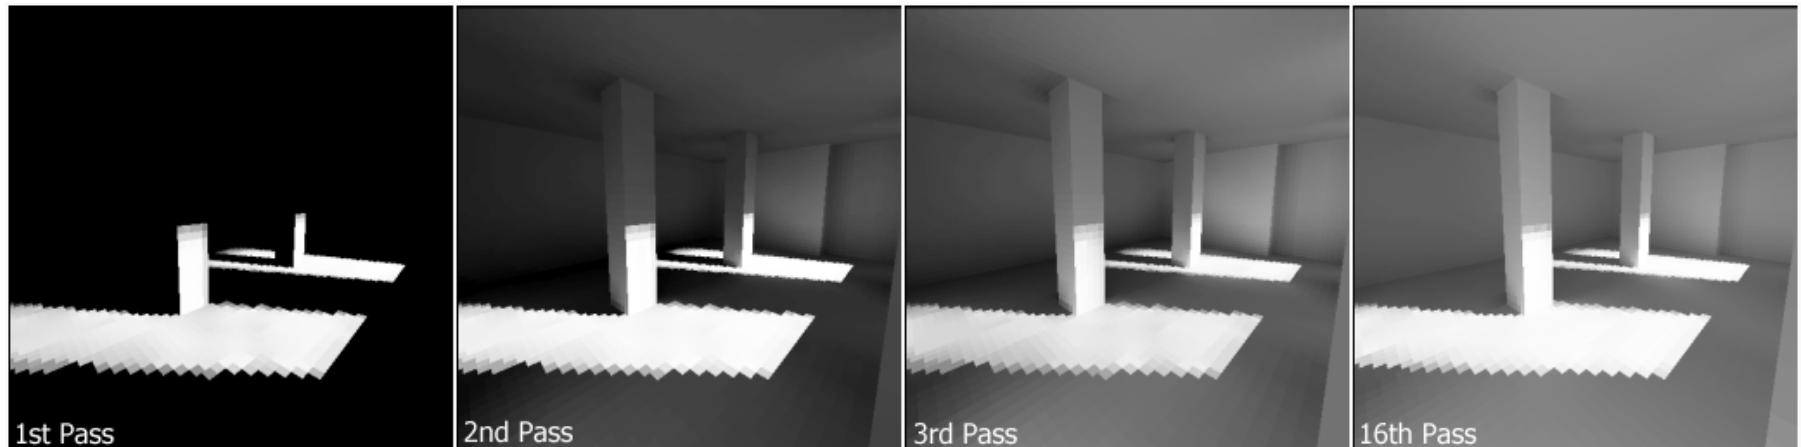
# Solving radiosity

- ▶ Pipeline for computing global illumination :



# Solving radiosity

- ▶ Iterative solving:

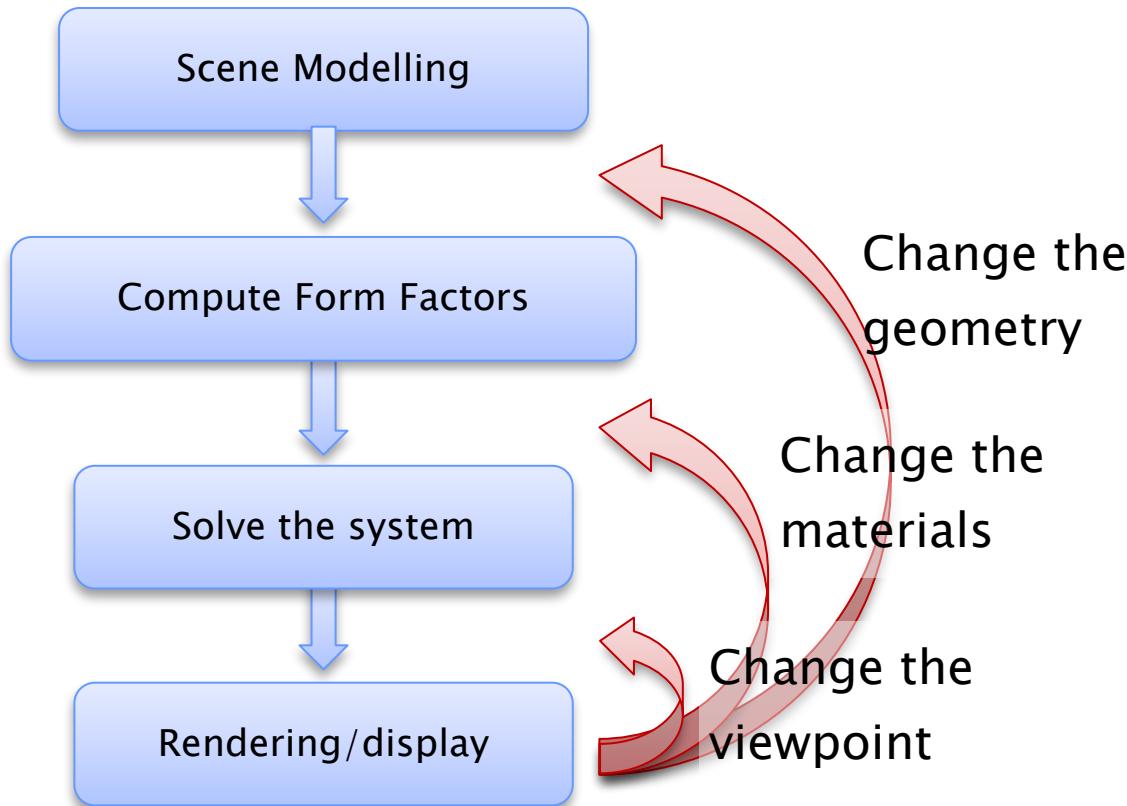


# Question – 1 mn

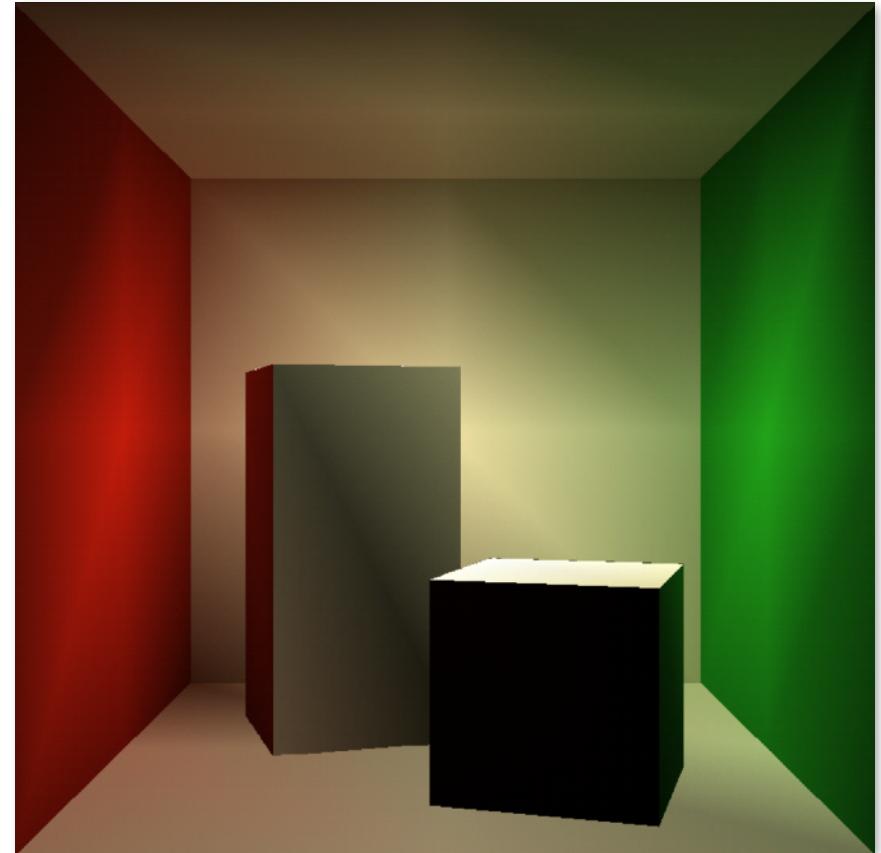
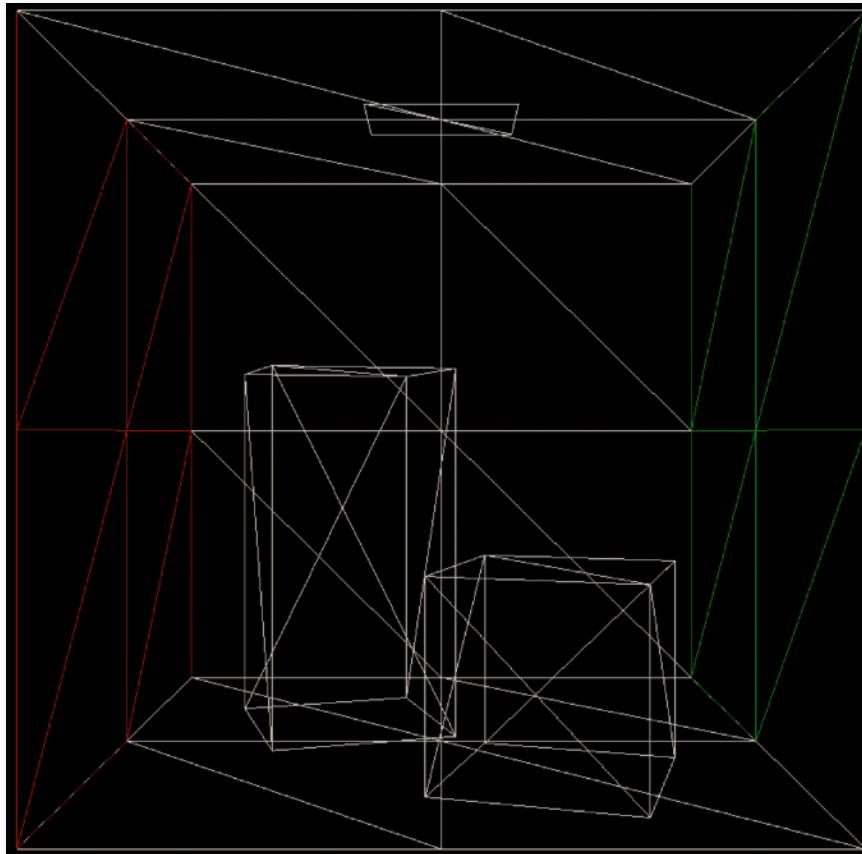


- ▶ What must be recomputed if something changes in the scene?
  - Geometry
  - Reflectance properties
  - Viewpoint

# Solving radiosity



# Solving radiosity



# Solving radiosity



Museum simulation. Cornell University. 50,000 patches.

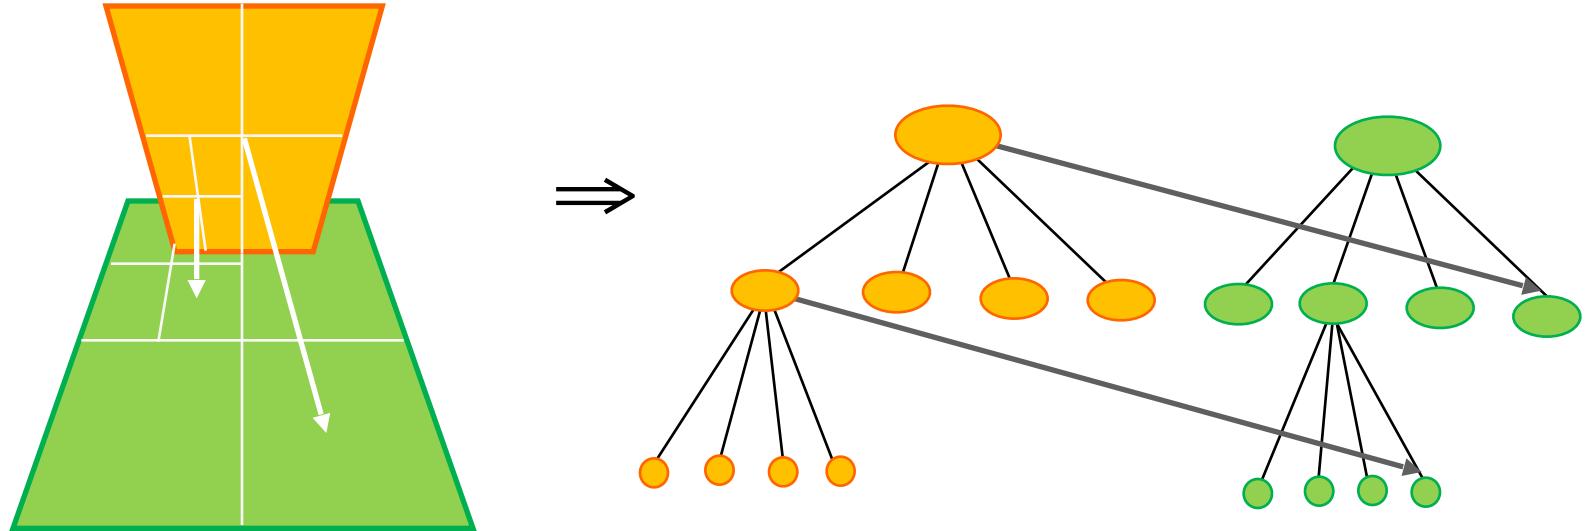
# Radiosity: pros

- ▶ Computations are viewpoint independent
- ▶ Okay for complex scenes
- ▶ Solving light exchanges
  - Interactive rendering

# Radiosity: cons

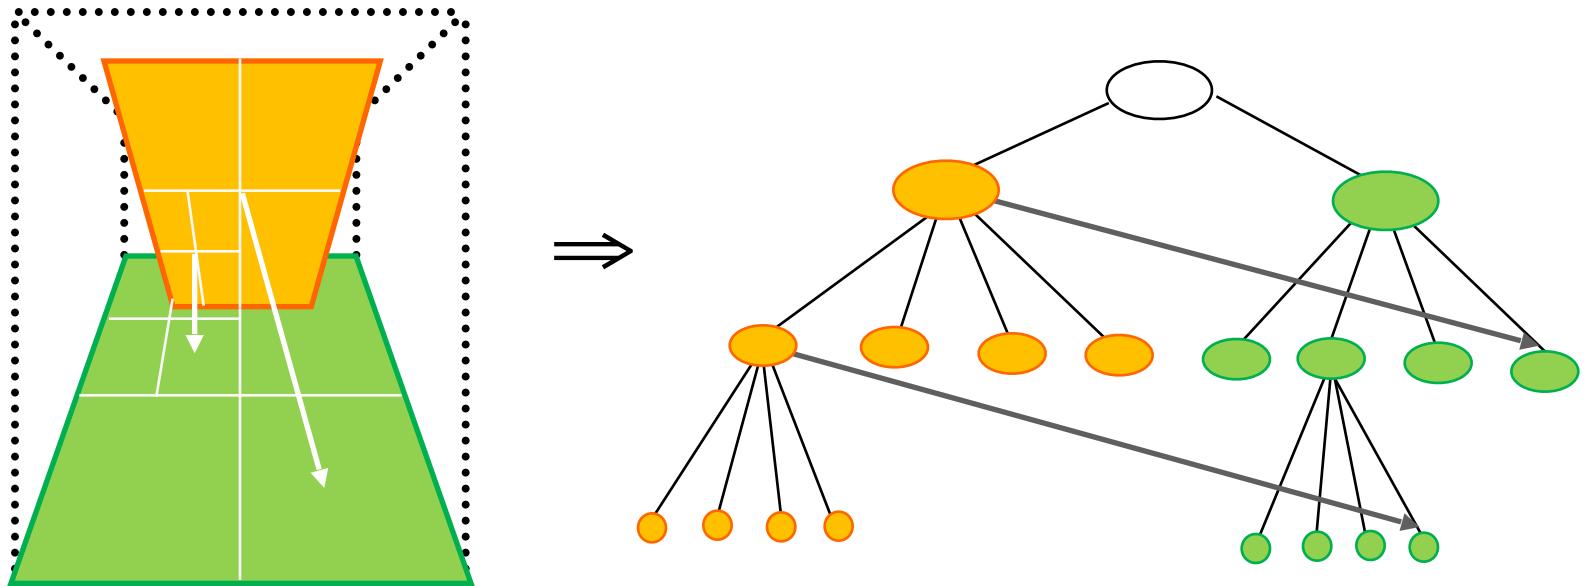
- ▶ Memory cost
- ▶ Only diffuse materials
  - « final gather » using Ray-Tracing
- ▶ Meshing
  - Discontinuity mesh
- ▶ Long pre-computations
  - Possible speed-up: hierarchical radiosity

# Hierarchical radiosity [Hanrahan91]



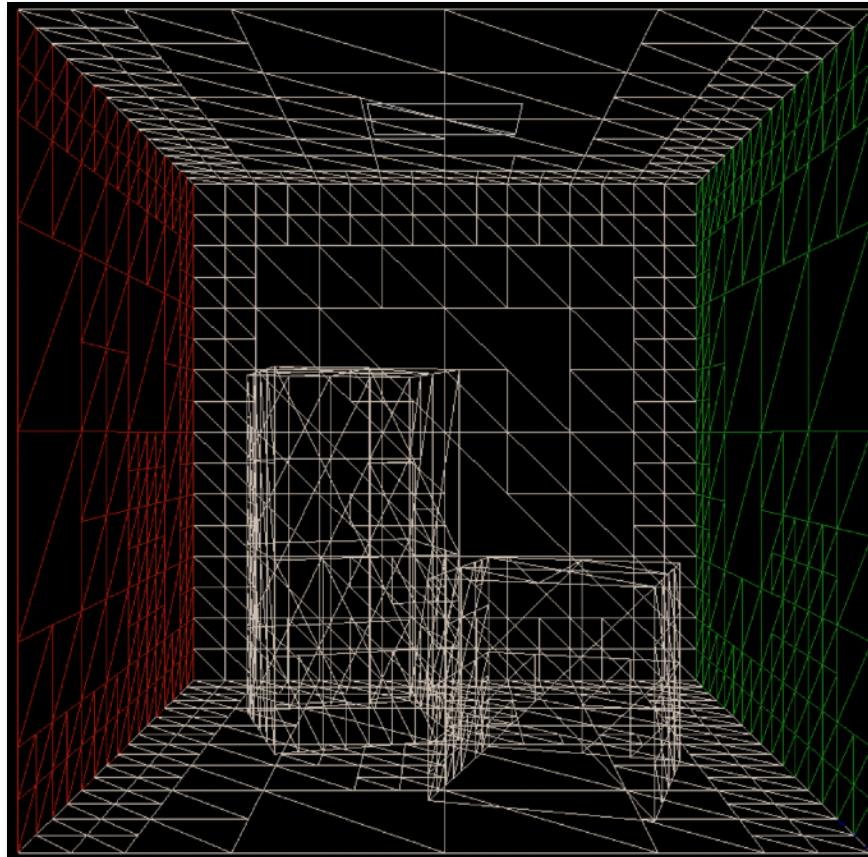
- Computation at different hierarchical levels
- Push-pull

# Hierarchical radiosity [Hanrahan91]

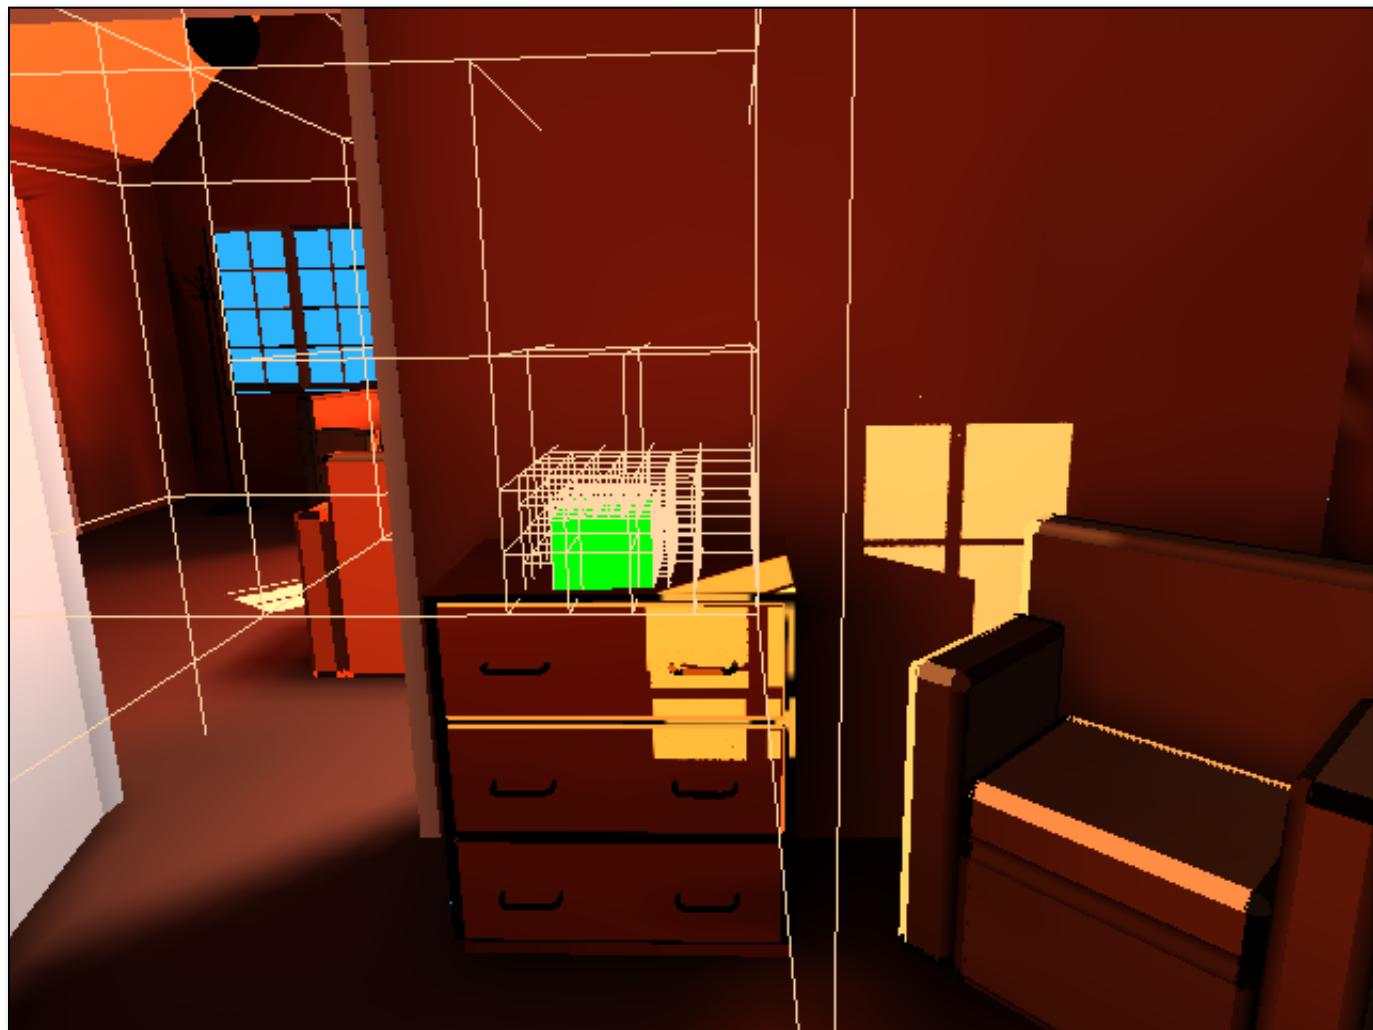


- ▶ Computation at different hierarchical levels
- ▶ Push-pull

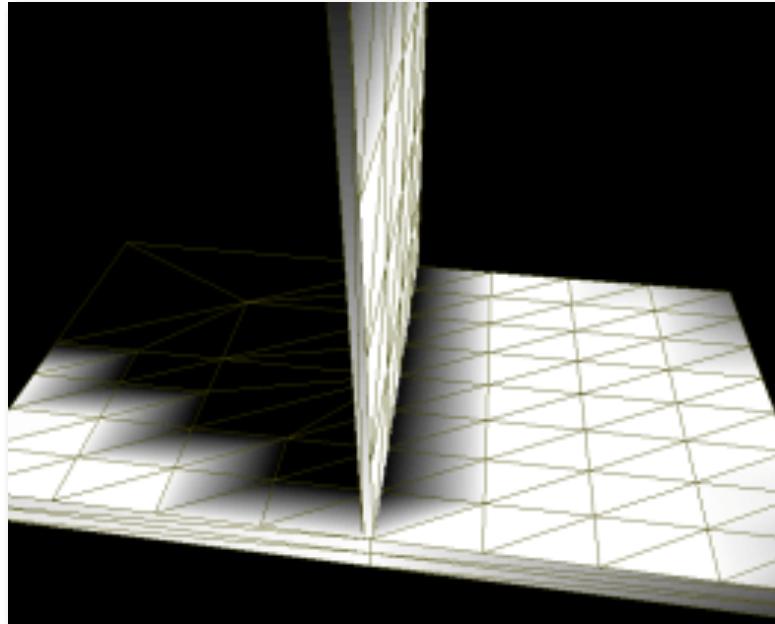
# Hierarchical radiosity



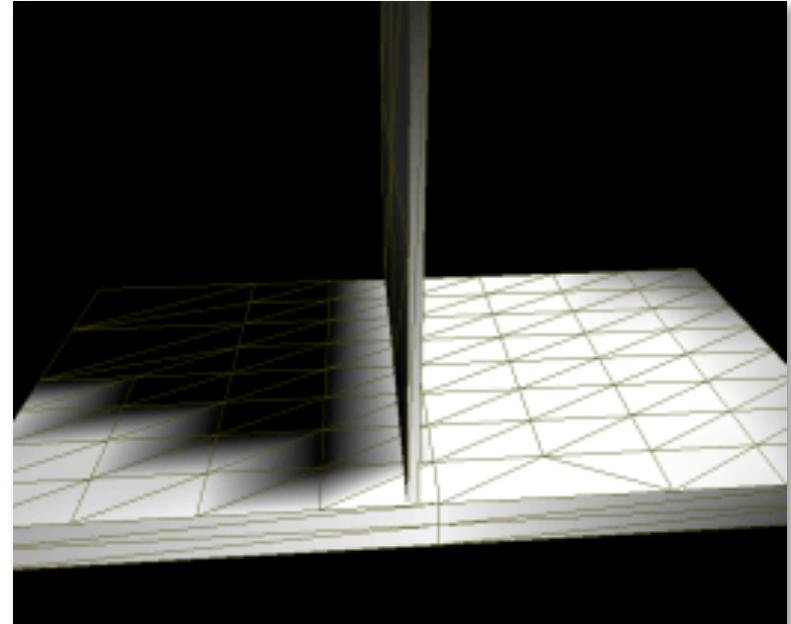
# Hierarchical radiosity



# Mesh quality

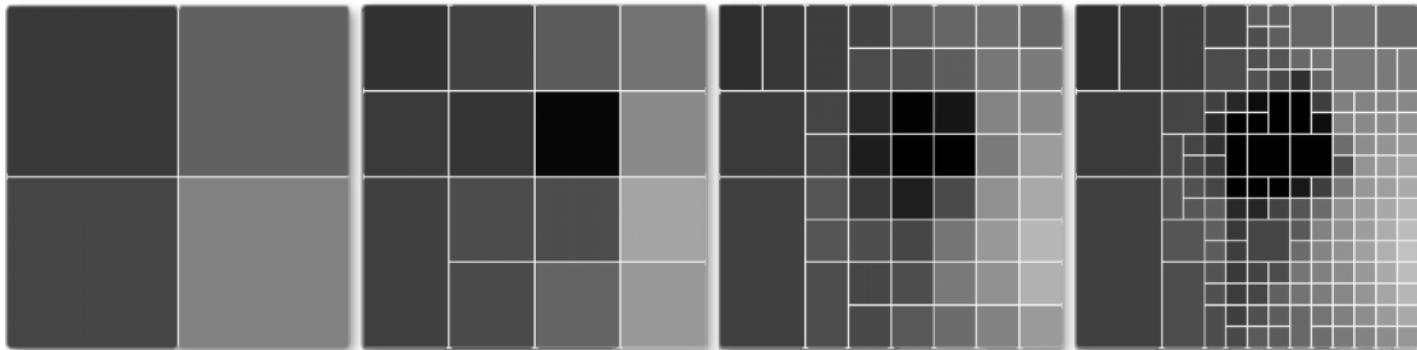


**Shadow leak**

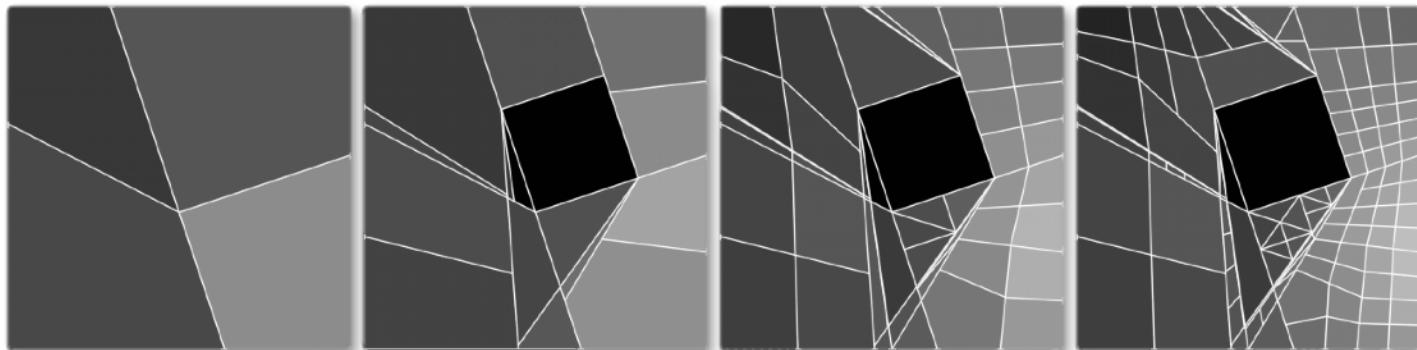


**Light leak**

# Discontinuity meshing



Regular subdivision



Discontinuity meshing

# Discontinuity meshing



# Discontinuity meshing



# Discontinuity meshing



# Résultats et comparaison



Rendu Scan-line.  
(3DS MAX)

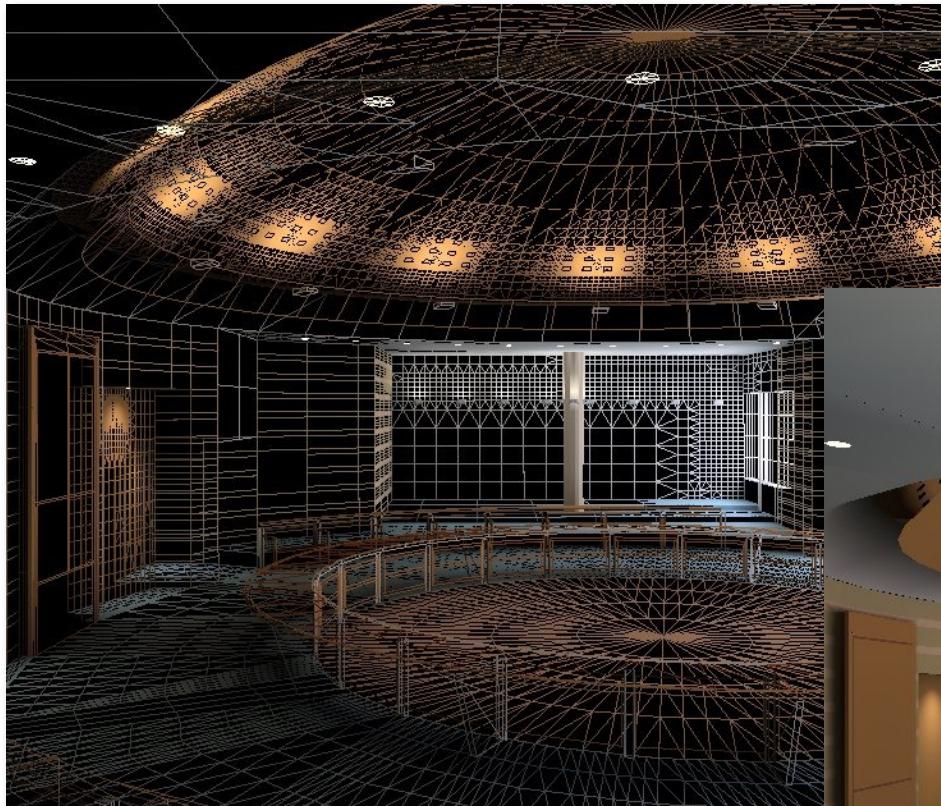
Rendu en radiosité et  
lancer de rayons.



# Results and comparison



# Results and comparison



# Results and comparison



# Results and comparison



# Radiosity today

- ▶ Used by architects (Lightscape)
- ▶ Used to precompute diffuse lighting for some video games (light maps)
- ▶ Not an active research topic anymore
  - Monte-Carlo is more generic
  - But pre-computed radiance transfer is similar  
(used e.g. in Max Payne 2)
- Complexity =  $O(n^2)$  where  $n$  = geometry

# Virtual Point Lights

- ▶ Initially, a technique for interactive rendering
  - “instant radiosity”
- ▶ First step:
  - Propagate light in the scene
  - Store as virtual point lights (VPL)
- ▶ Second step:
  - Direct illumination from the VPL
  - Can be done on the GPU (100–1000 VPLs/s)
  - Needs shadows

# Virtual Point Lights

- ▶ Also Point-Based Global Illumination
- ▶ Separate illumination from scene geometry
- ▶ Group point lights for more efficiency:
  - ▶ Light cuts, row-column matrix sampling



# Row–Column Matrix Sampling

- ▶ VPL as a method for global illumination
- ▶ Many light sources (100k)
- ▶ Matrix interpretation:
  - Columns= contribution from one light source
  - Rows= one pixel
  - Huge matrix (2m x 100k)
- Don't compute the whole matrix!
- ▶ Following slides: presentation of RCMS, Siggraph 2007,  
Miloš Hašan.

# Indirect Illumination → Many Lights



=



$\Sigma$  (

)

100,000 point lights

# Environment Map → Many Lights

---



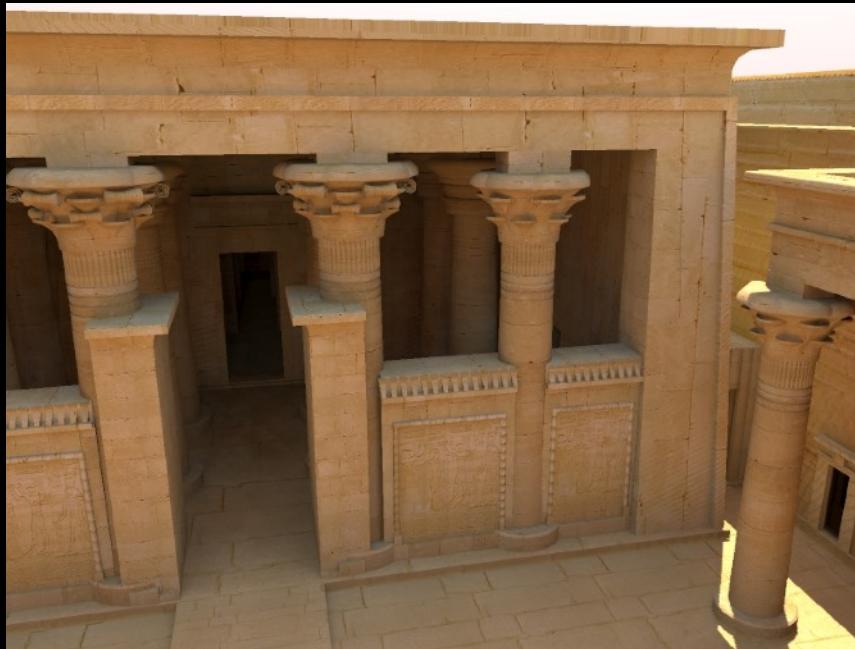
=

$$\sum \left( \begin{array}{c} \text{[Image of a bunny with point lights]} \\ \text{[Image of a bunny with point lights]} \\ \text{[Image of a bunny with point lights]} \\ \text{[Image of a bunny with point lights]} \end{array} \right )$$

100,000 point lights

# Sun, Sky, Indirect → Many Lights

---



=

$$\Sigma \left( \begin{array}{c} \text{[Image of a dark scene with a few point lights]} \\ \text{[Image of a scene with many small point lights]} \\ \text{[Image of a scene with many small point lights]} \\ \text{[Image of a scene with many small point lights]} \end{array} \right )$$

100,000 point lights

# Brute Force Takes Minutes

---

- Why not sum all columns?
  - With 100,000 lights, still several minutes



10 min



13 min



20 min

# Our Contribution

- Fast, accurate, GPU-based approximation

**Brute  
force:**



**10 min**



**13 min**



**20 min**

**Our  
result:**



**3.8 sec**



**13.5 sec**



**16.9 sec**

- Application: Preview for lighting design

# Conversion to Many Lights

---

- Area, indirect, sun/sky



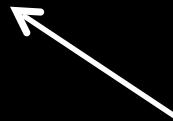
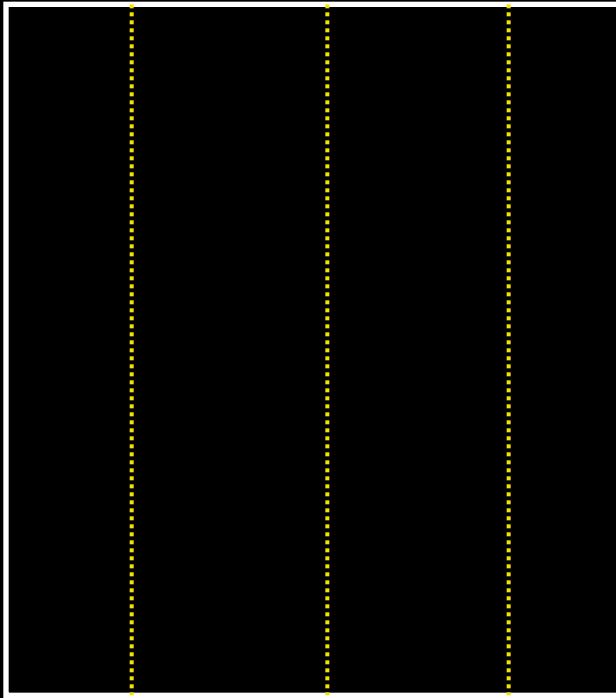
Courtesy Walter et al., Lightcuts, SIGGRAPH 05/06

# A Matrix Interpretation

---

**Lights (100,000)**

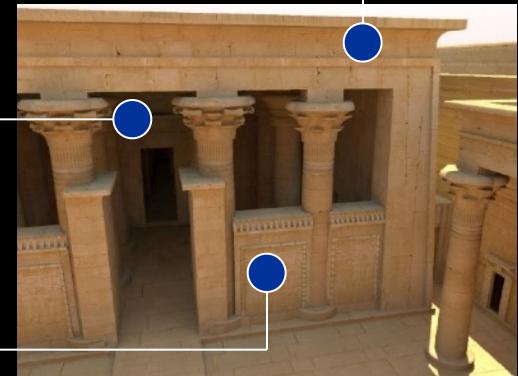
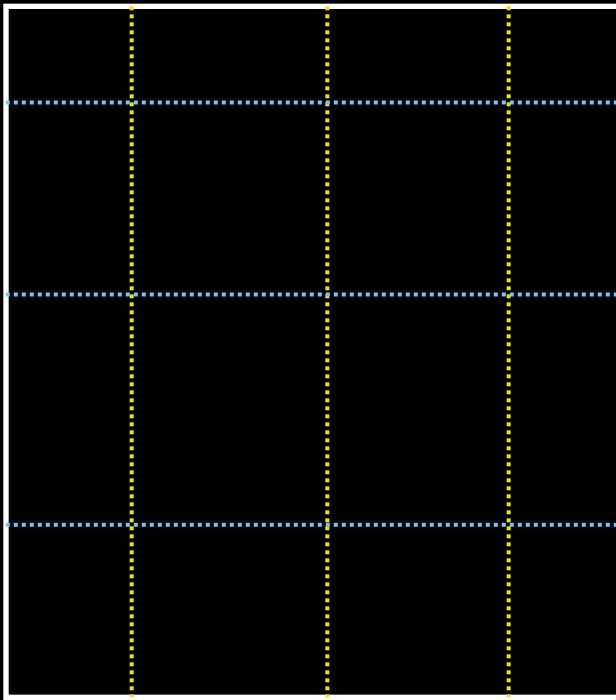
**Pixels  
(2,000,000)**



# A Matrix Interpretation

Lights (100,000)

Pixels  
(2,000,000)



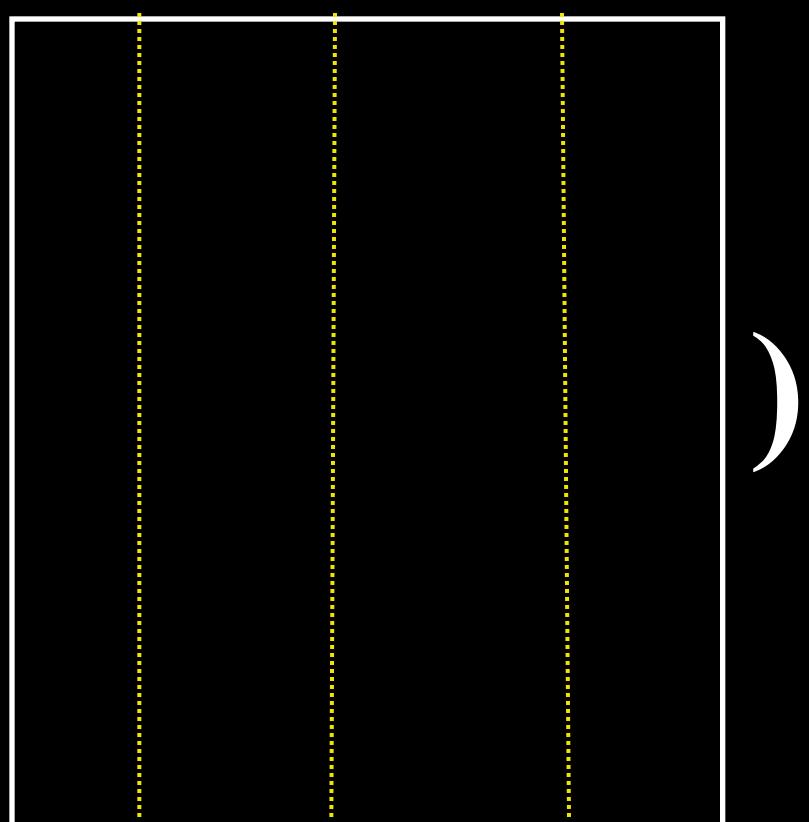
# Problem Statement

---



=  $\Sigma$  (

Pixels



Lights

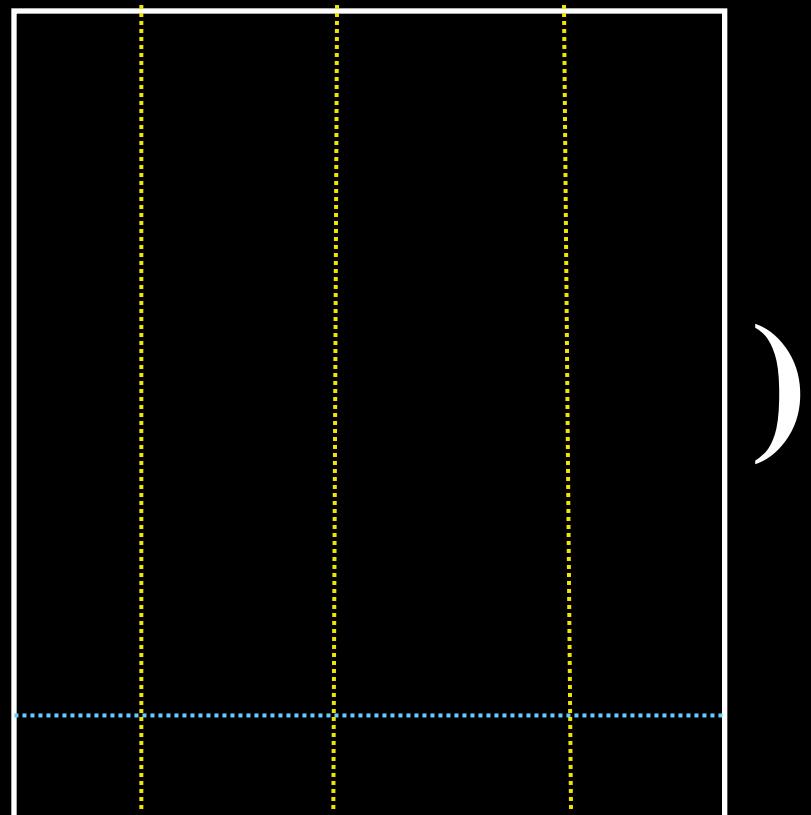
# Problem Statement

---



=  $\sum$  (Pixels

Lights



# Problem Statement

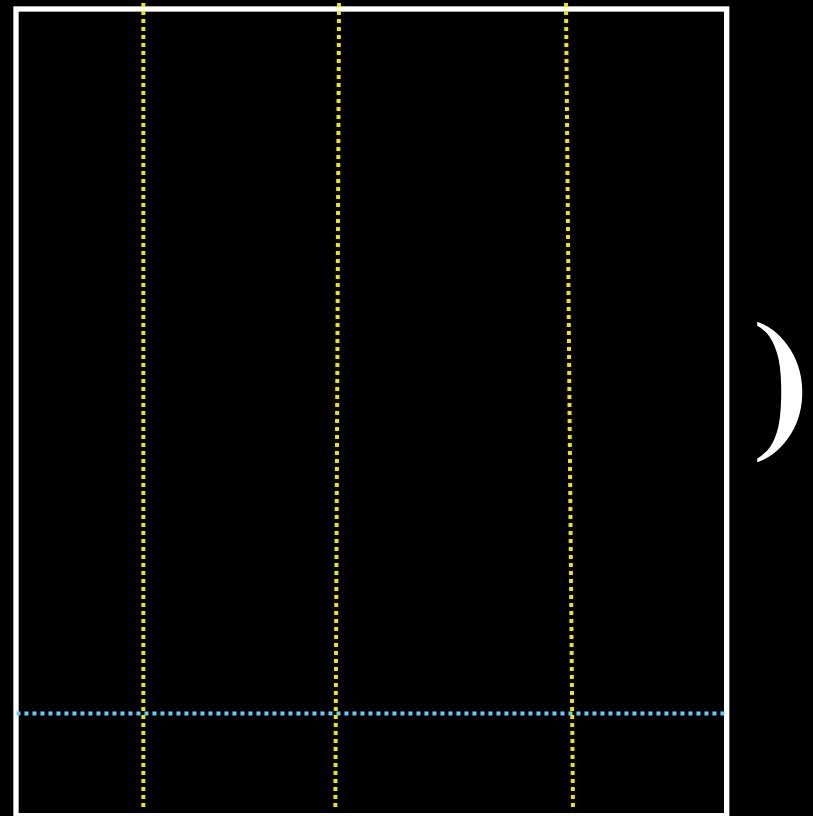
---

- Compute sum of columns

Lights



$$= \sum (\text{Pixels})$$



- Note: We don't have the matrix data

# Insight #1: Matrix has structure

---

- Compute small subset of elements
- Reconstruct

**643 lights**



**A simple scene**

**30 x 30 image**

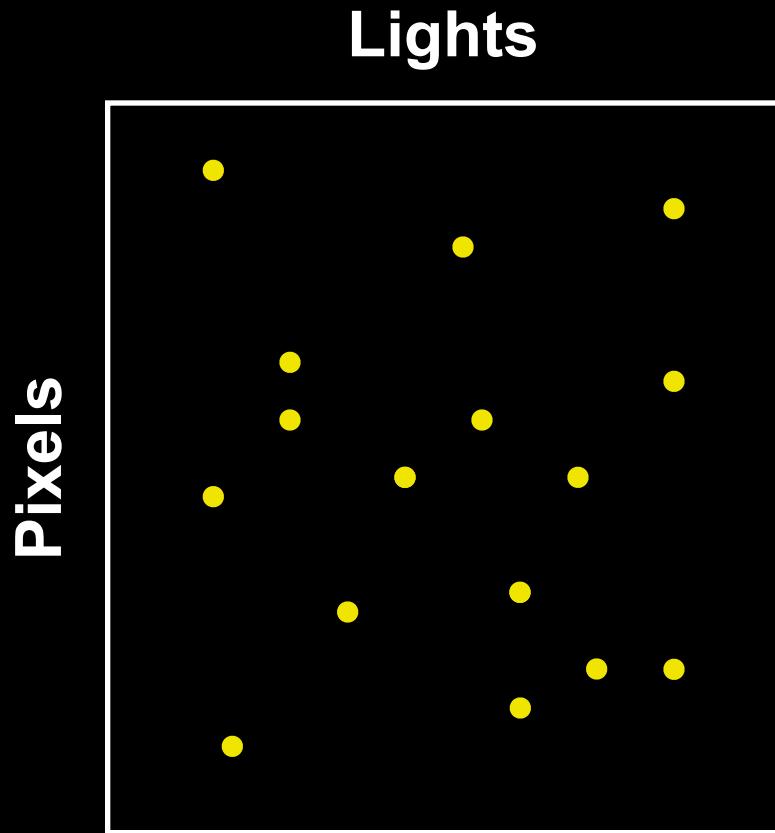
**900 pixels**



**The matrix**

# Insight #2: Sampling Pattern Matters

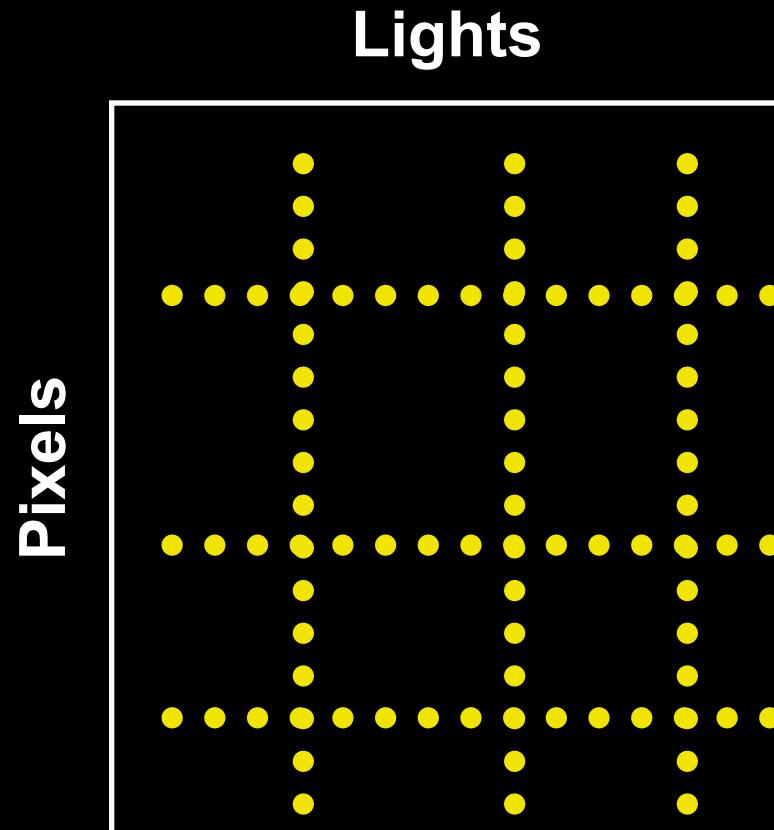
---



**Point-to-point visibility: Ray-tracing**

# Insight #2: Sampling Pattern Matters

---

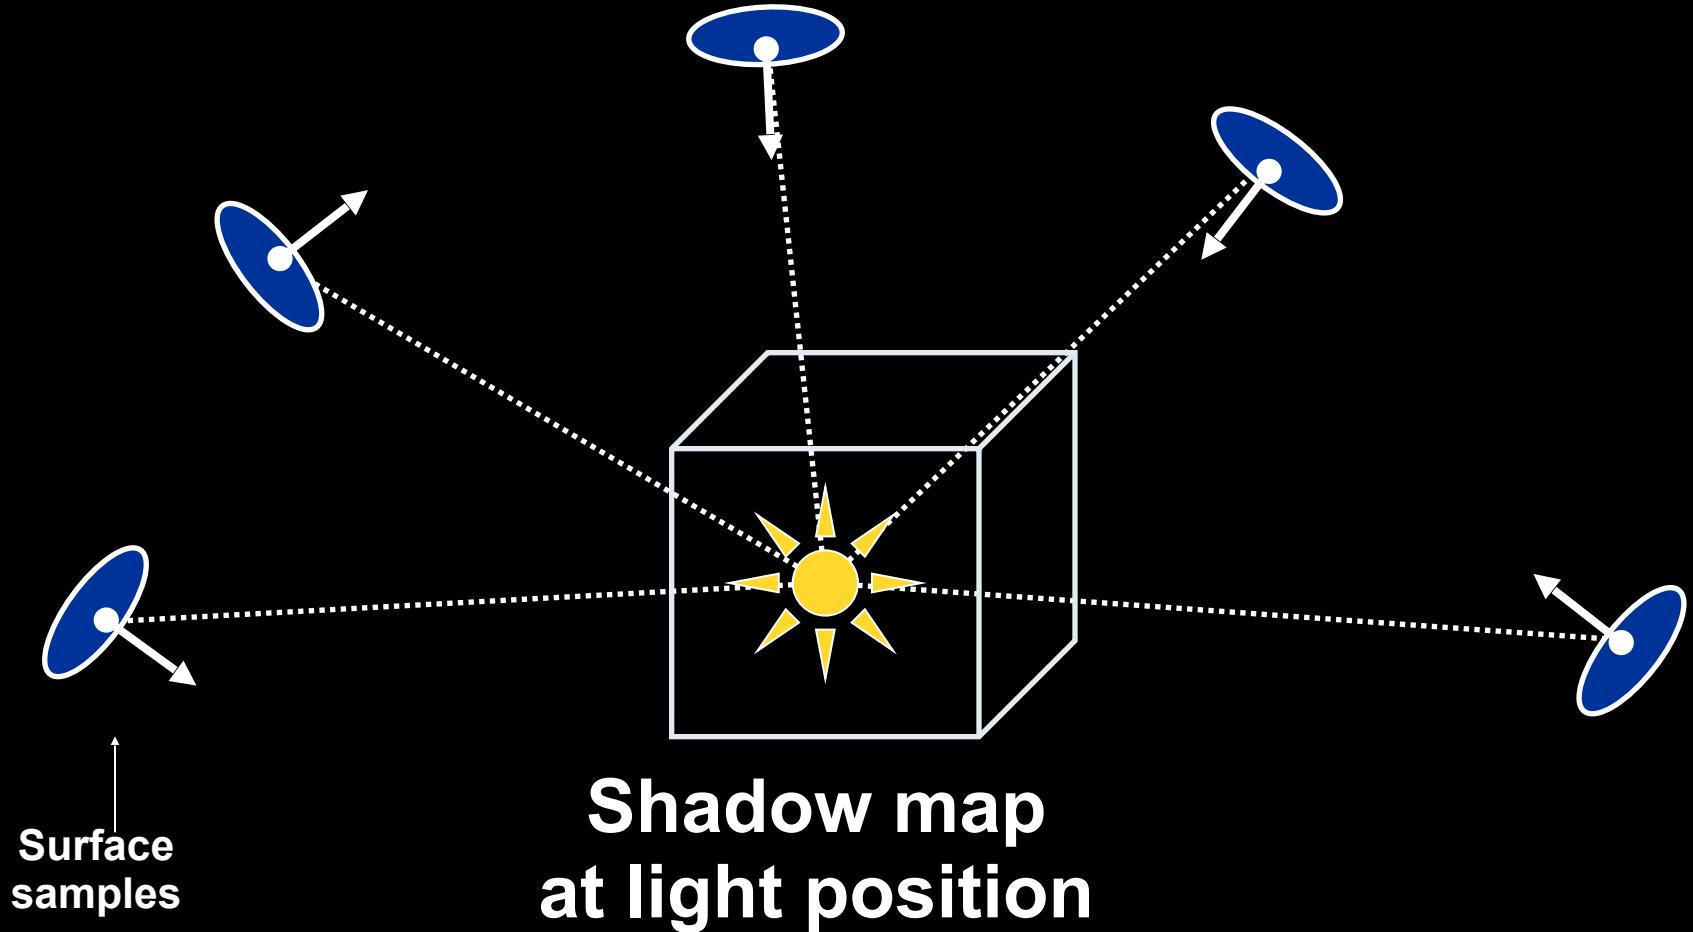


**Point-to-many-points visibility: Shadow-mapping**

# Row-Column Duality

---

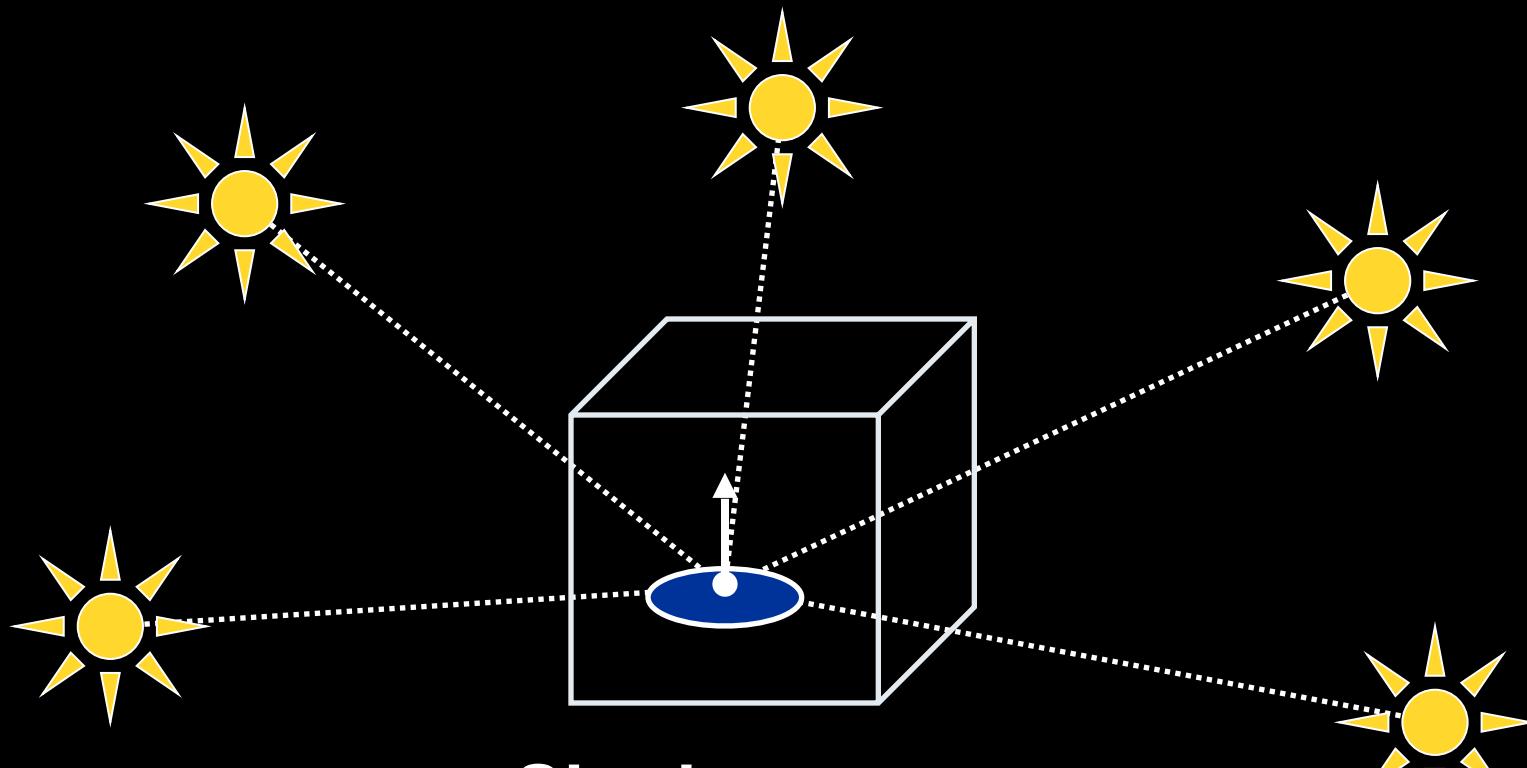
- Columns: Regular Shadow Mapping



# Row-Column Duality

---

- Rows: Also Shadow Mapping!

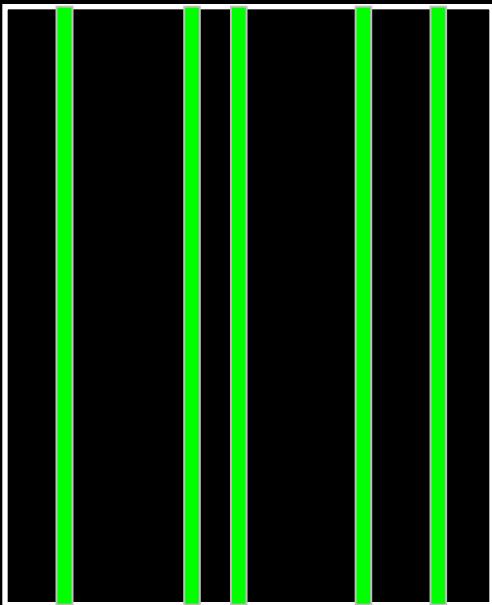


**Shadow map  
at sample position**

# Image as a Weighted Column Sum

---

- The following is possible:



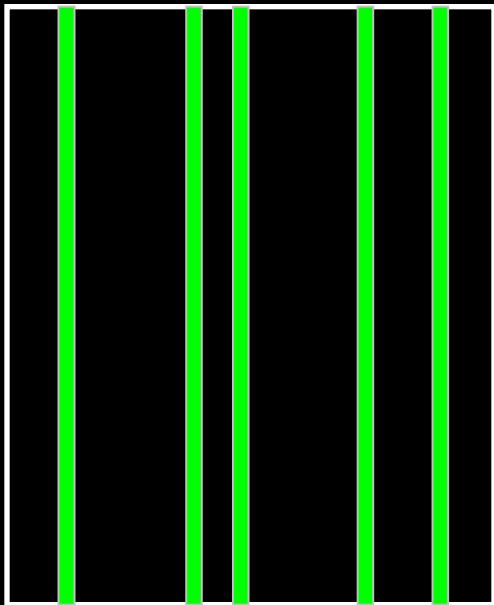
**compute very small  
subset of columns**

**compute  
weighted sum**

# Image as a Weighted Column Sum

---

- The following is possible:



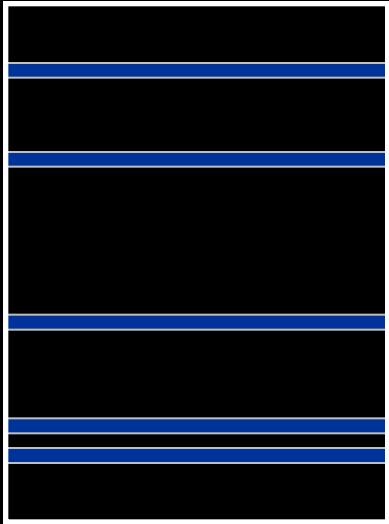
**compute very small  
subset of columns**

**compute  
weighted sum**

- Use rows to choose a good set of columns!

# Exploration and Exploitation

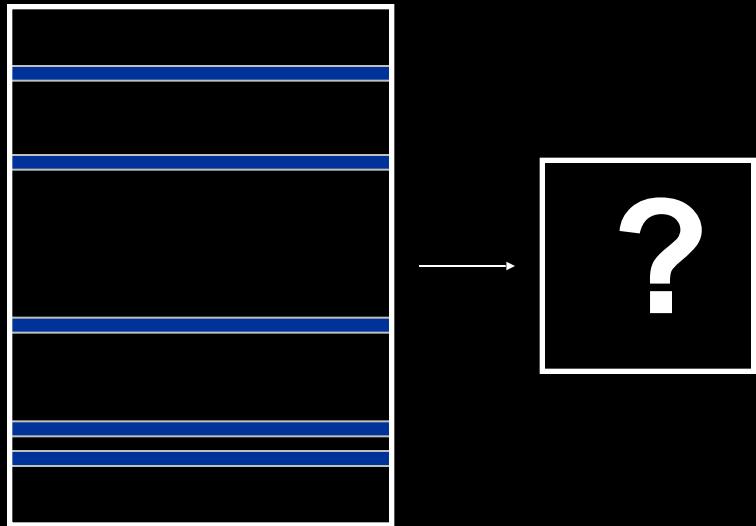
---



**compute rows  
(explore)**

# Exploration and Exploitation

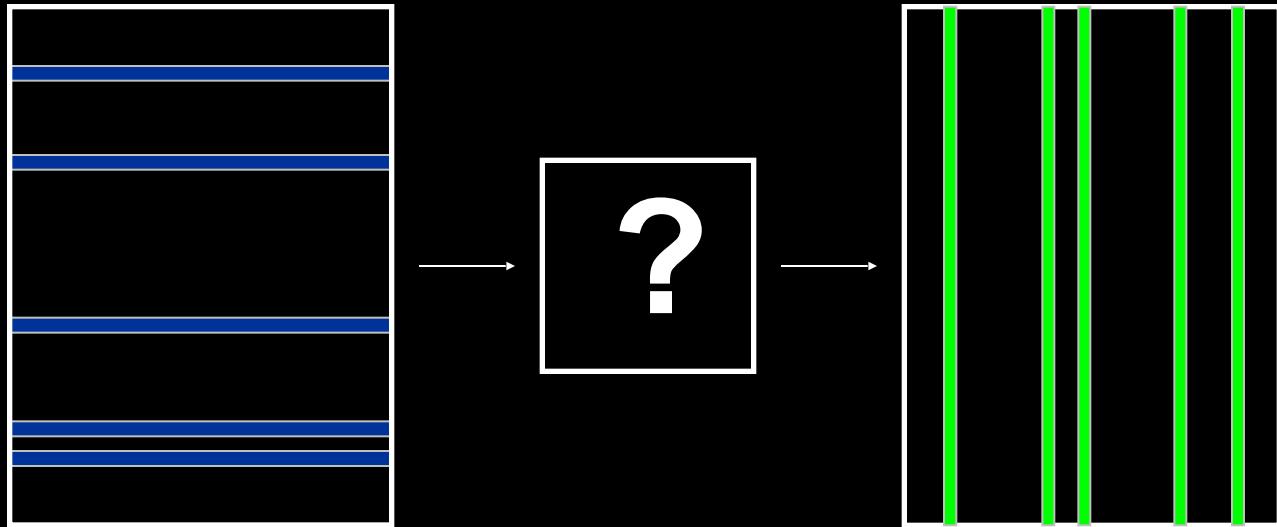
---



**compute rows choose columns  
(explore) and weights**

# Exploration and Exploitation

---



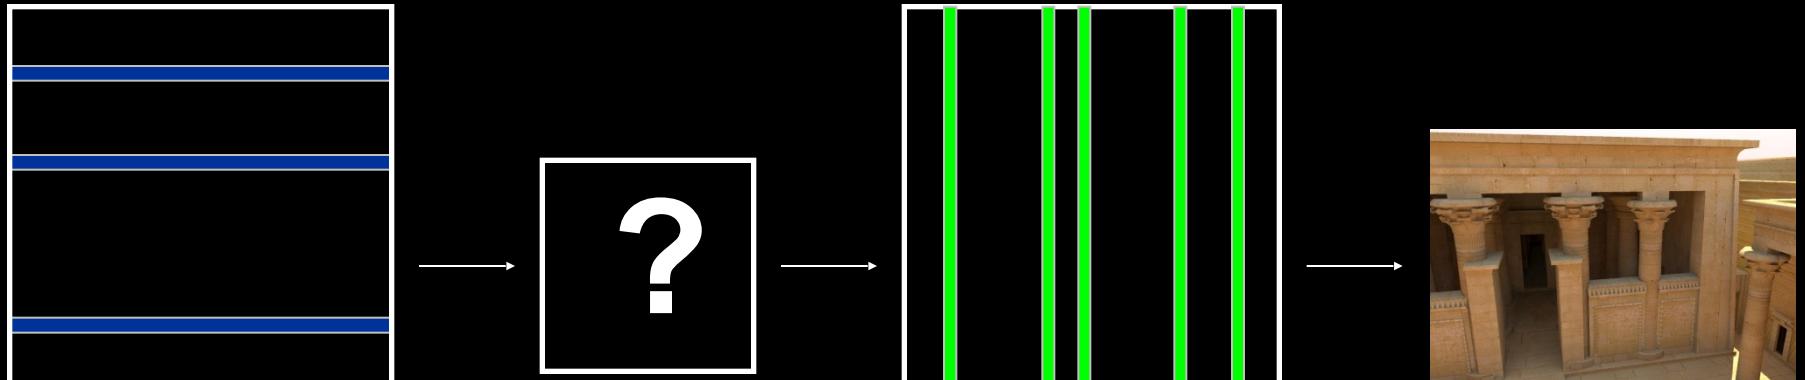
**compute rows  
(explore)**

**choose columns  
and weights**

**compute  
columns (exploit)**

# Exploration and Exploitation

---



**compute rows  
(explore)**

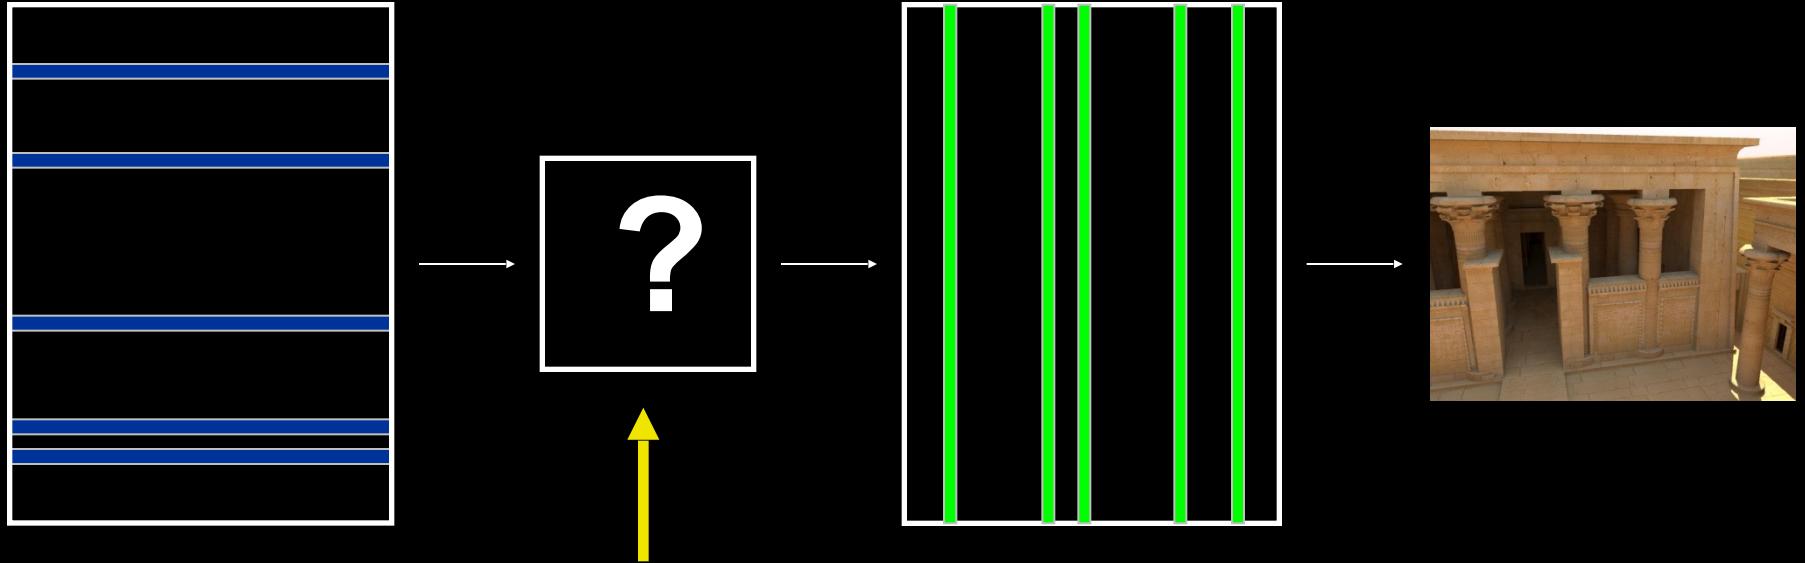
**choose columns  
and weights**

**compute  
columns (exploit)**

**weighted  
sum**

# Exploration and Exploitation

---



compute rows  
(explore)

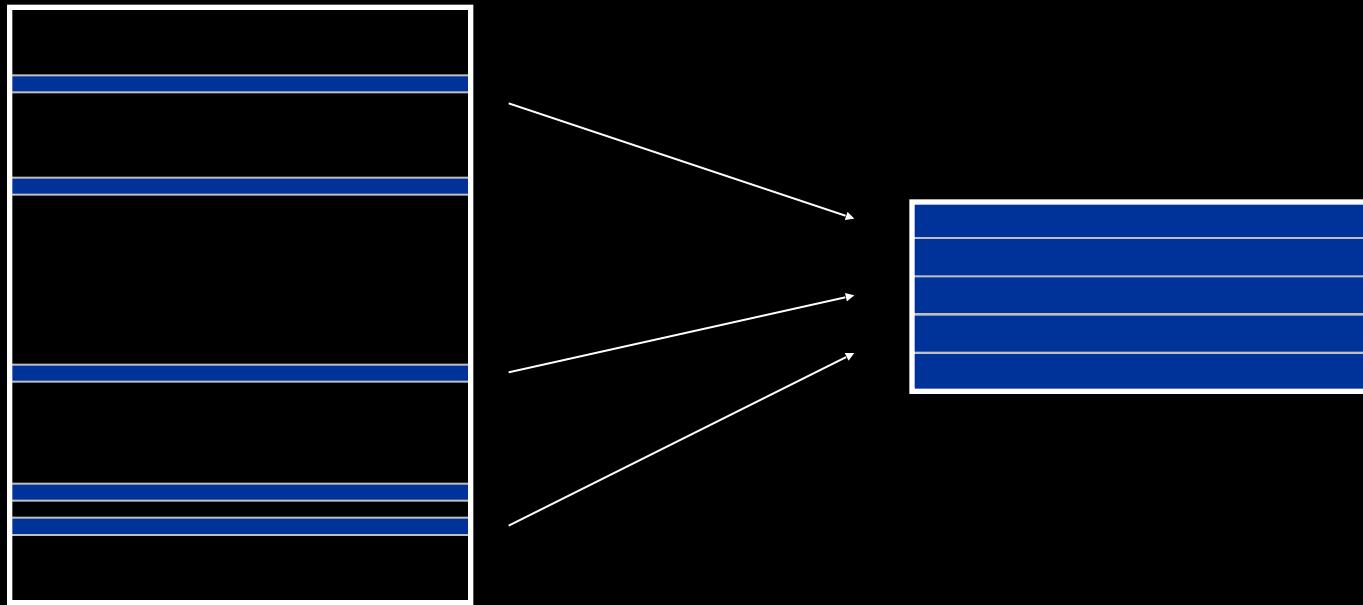
how to choose  
columns and  
weights?

compute  
columns (exploit)

weighted  
sum

# Reduced Matrix

---

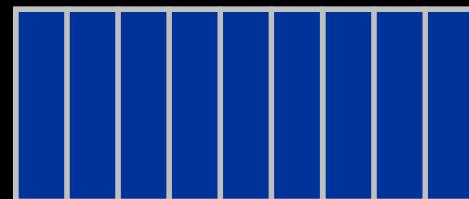


# Reduced Matrix

---

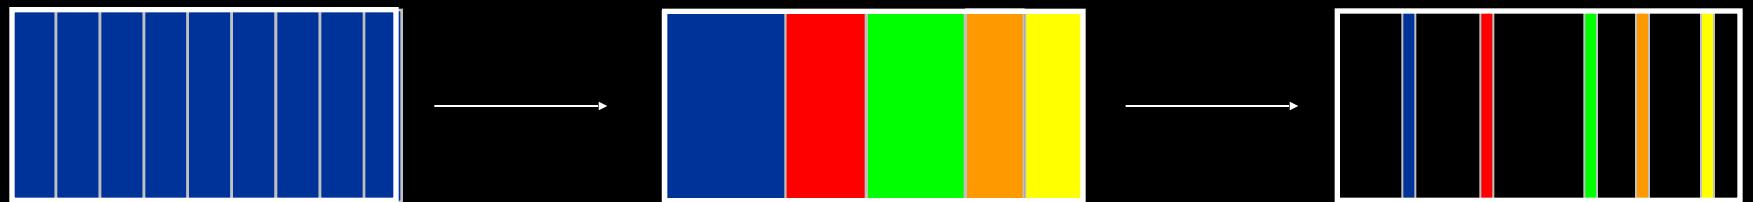


Reduced  
columns



# Clustering Approach

---

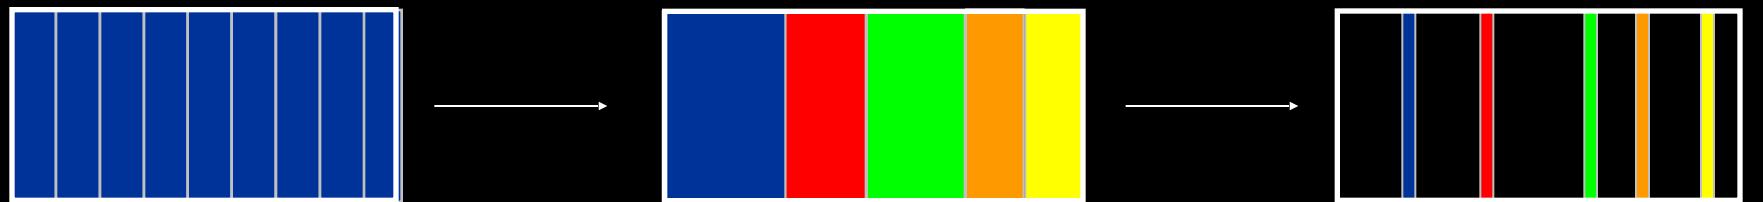


**Choose k clusters**

**Choose  
representative  
columns**

# Clustering Approach

---



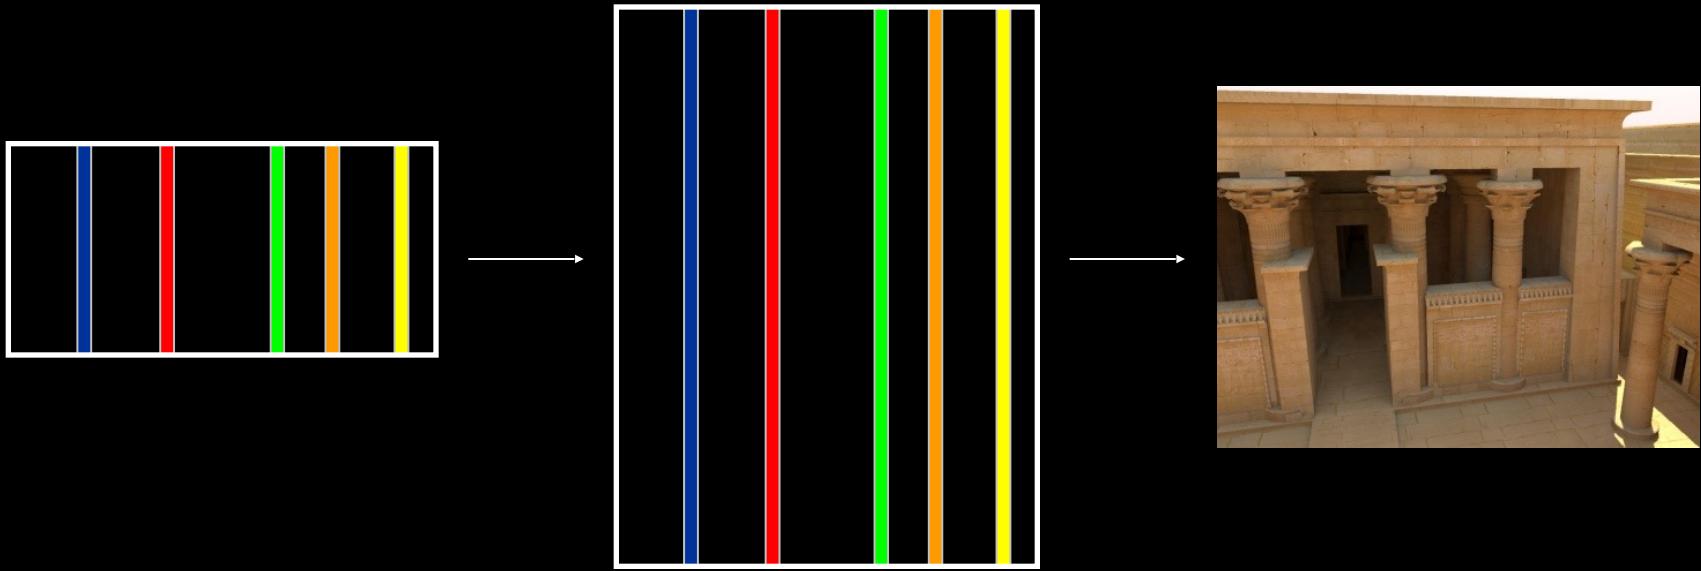
Reduced  
columns

Choose k clusters

Choose  
representative  
columns

# Reduced → Full

---



**Representative  
columns**

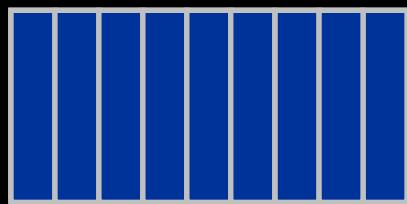
**Use the same  
representatives  
for the full matrix**

**Weighted  
sum**

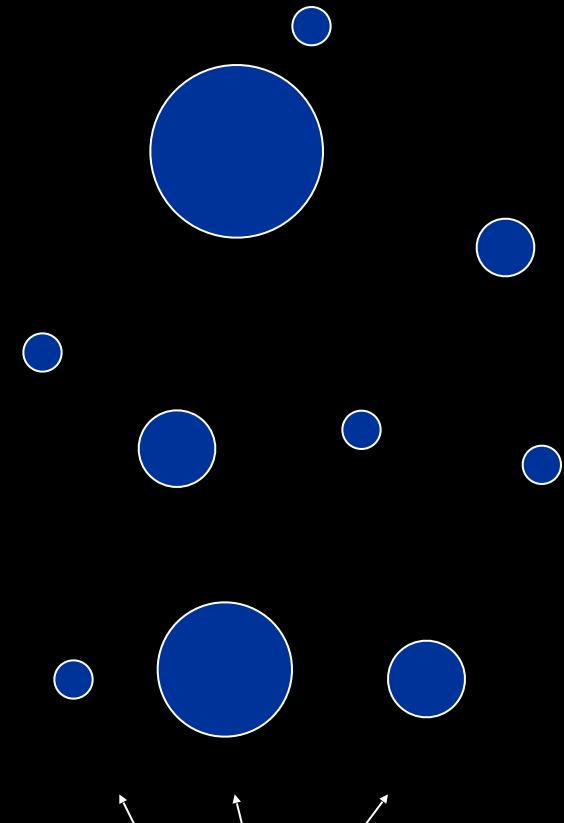
# Visualizing the Reduced Columns

---

**Reduced columns:**  
vectors in high-dimensional space



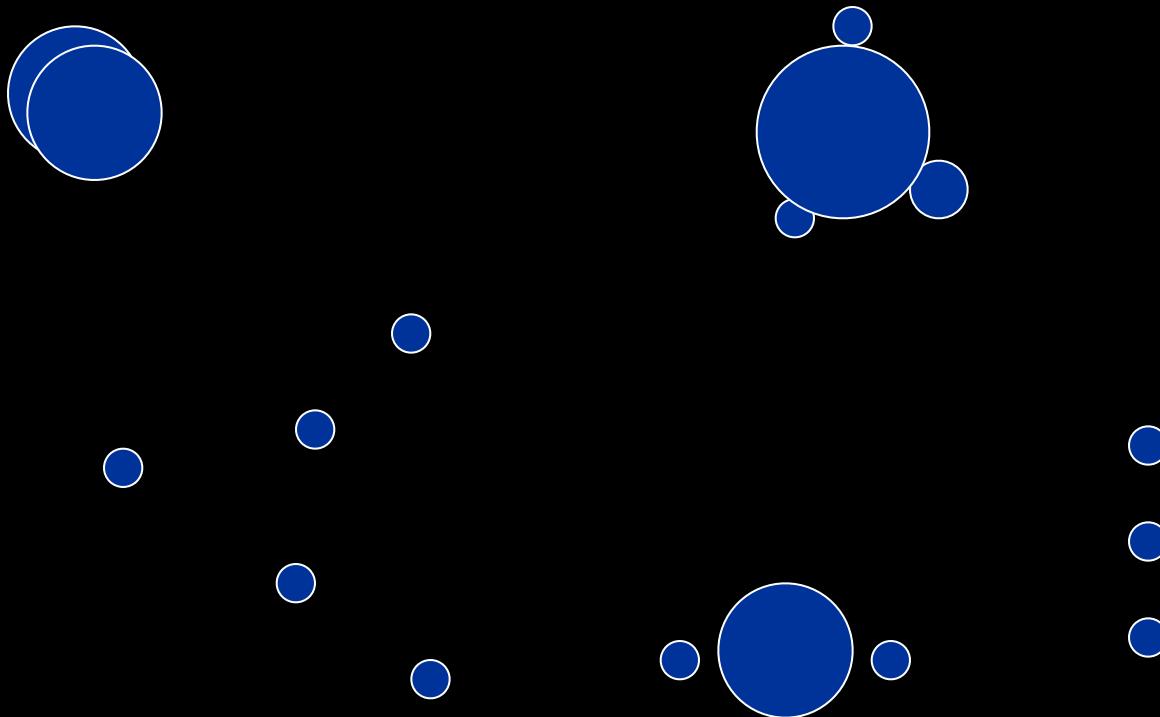
visualize as ...



radius = norm

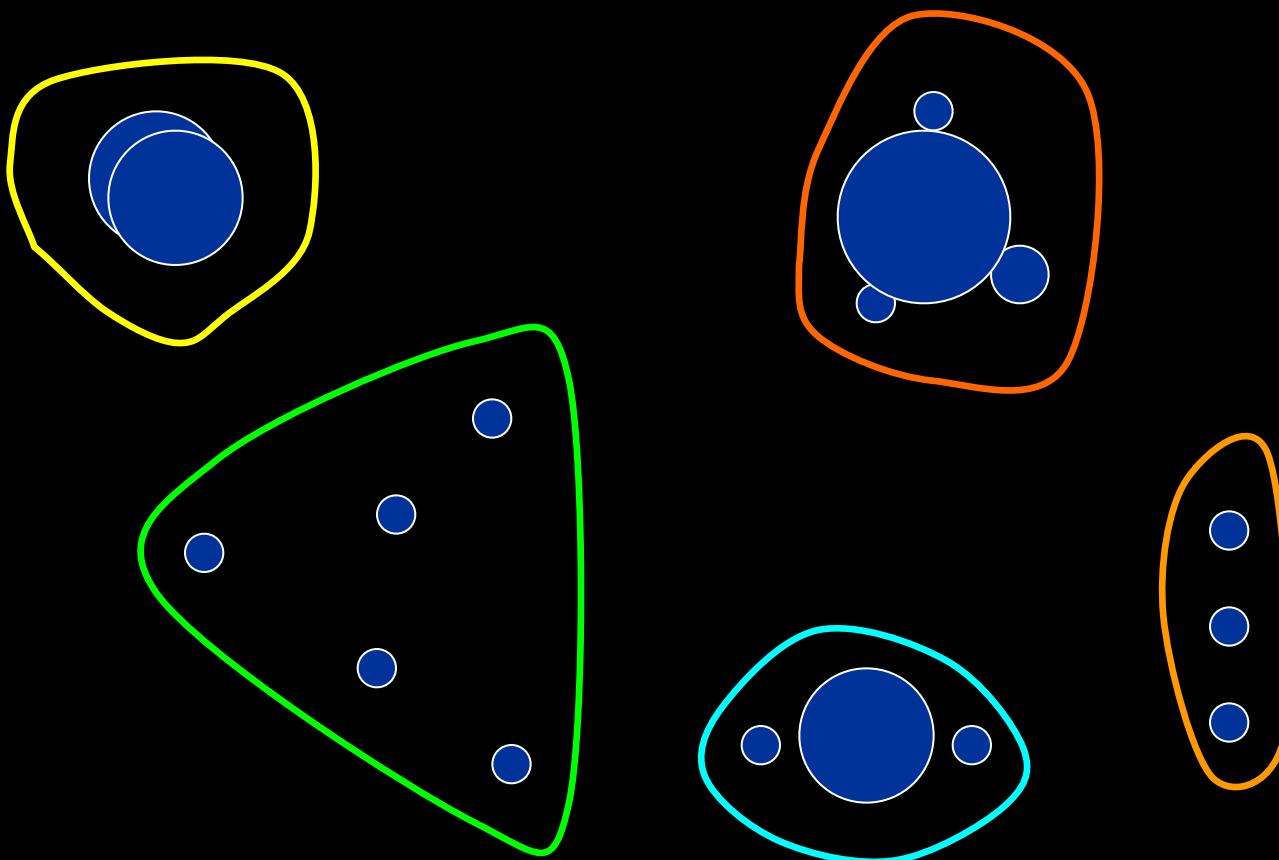
# Clustering Illustration

---



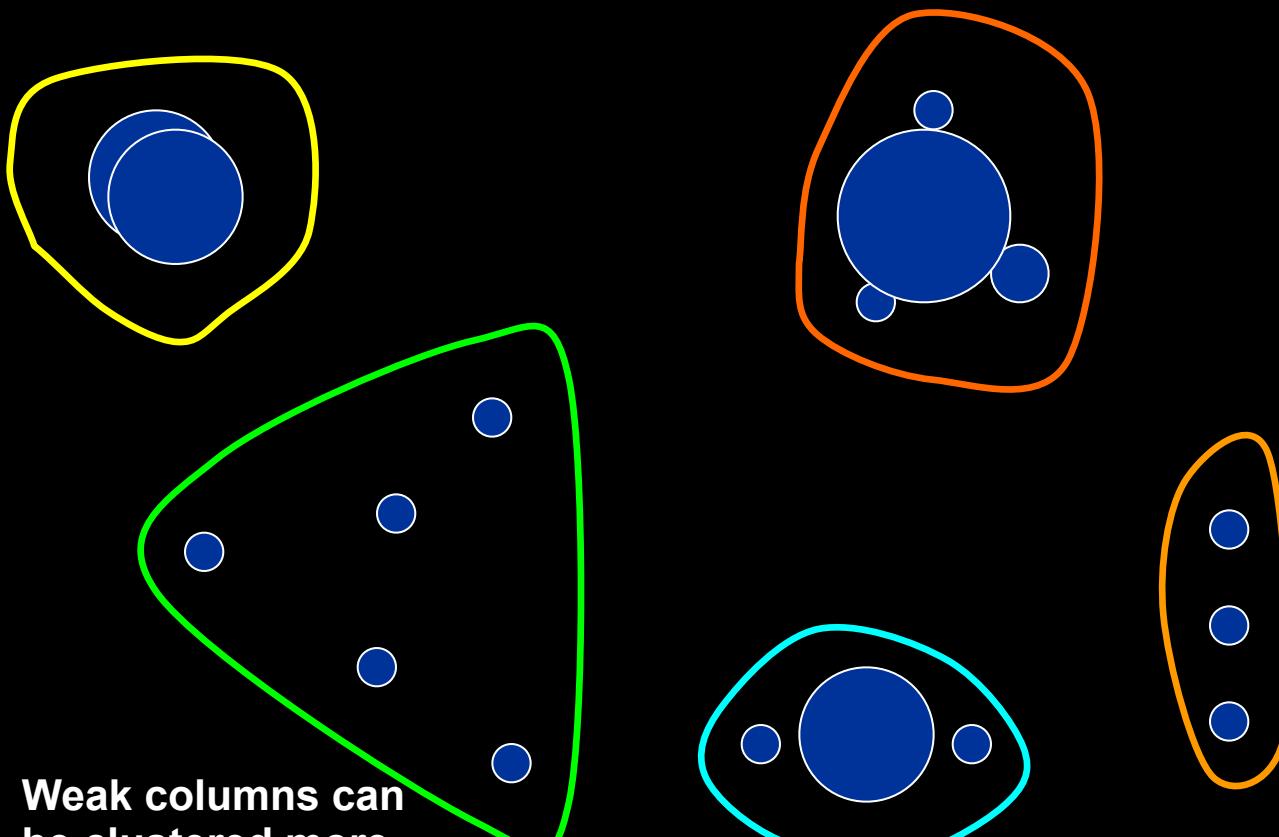
# Clustering Illustration

---



# Clustering Illustration

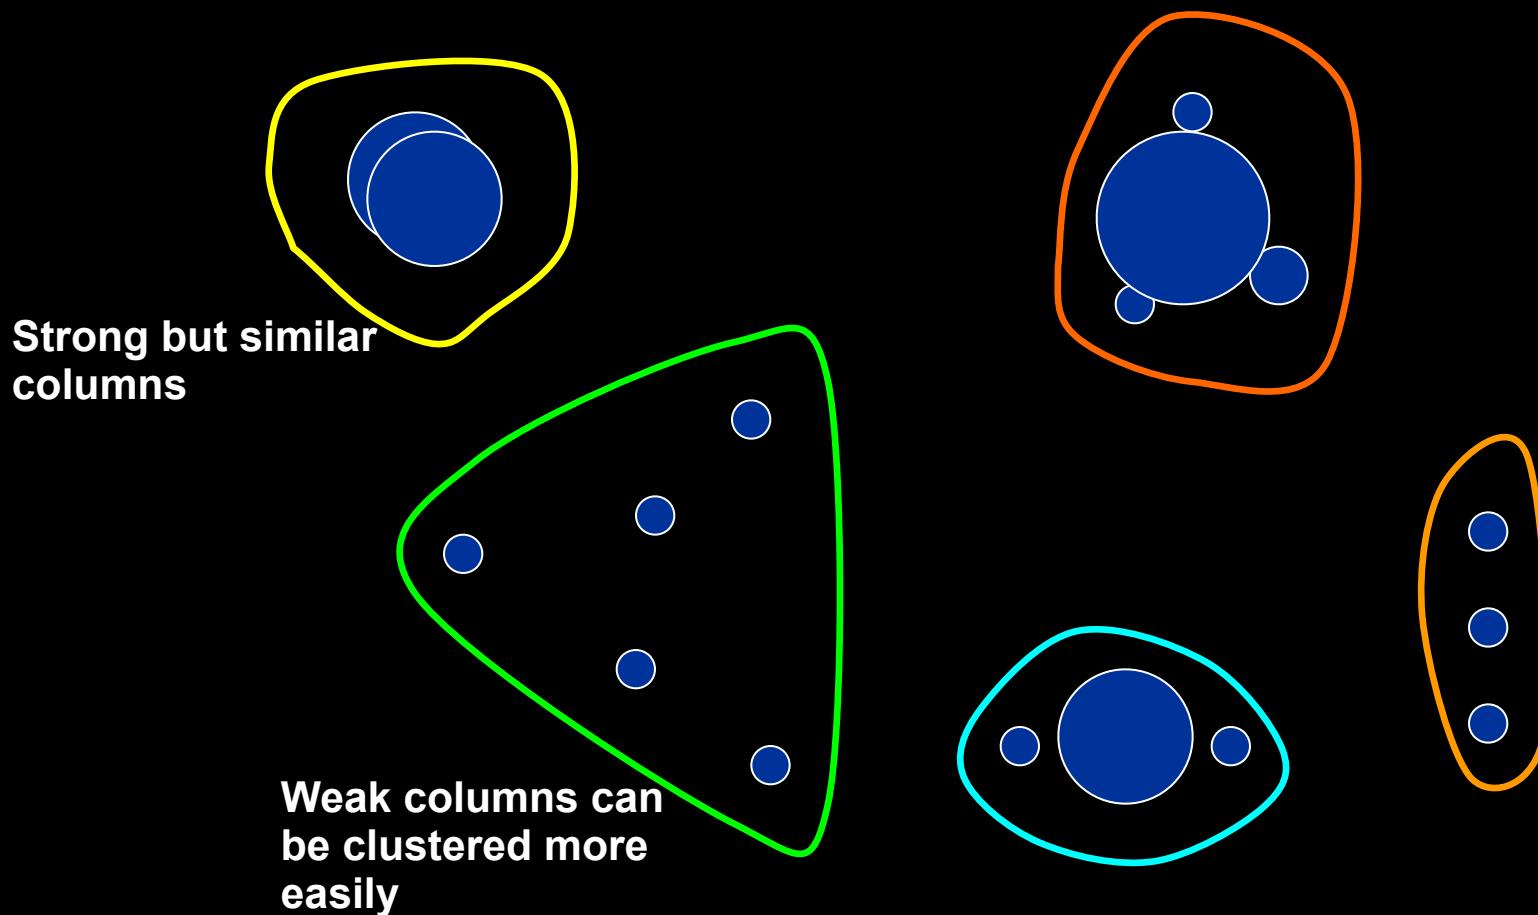
---



**Weak columns can  
be clustered more  
easily**

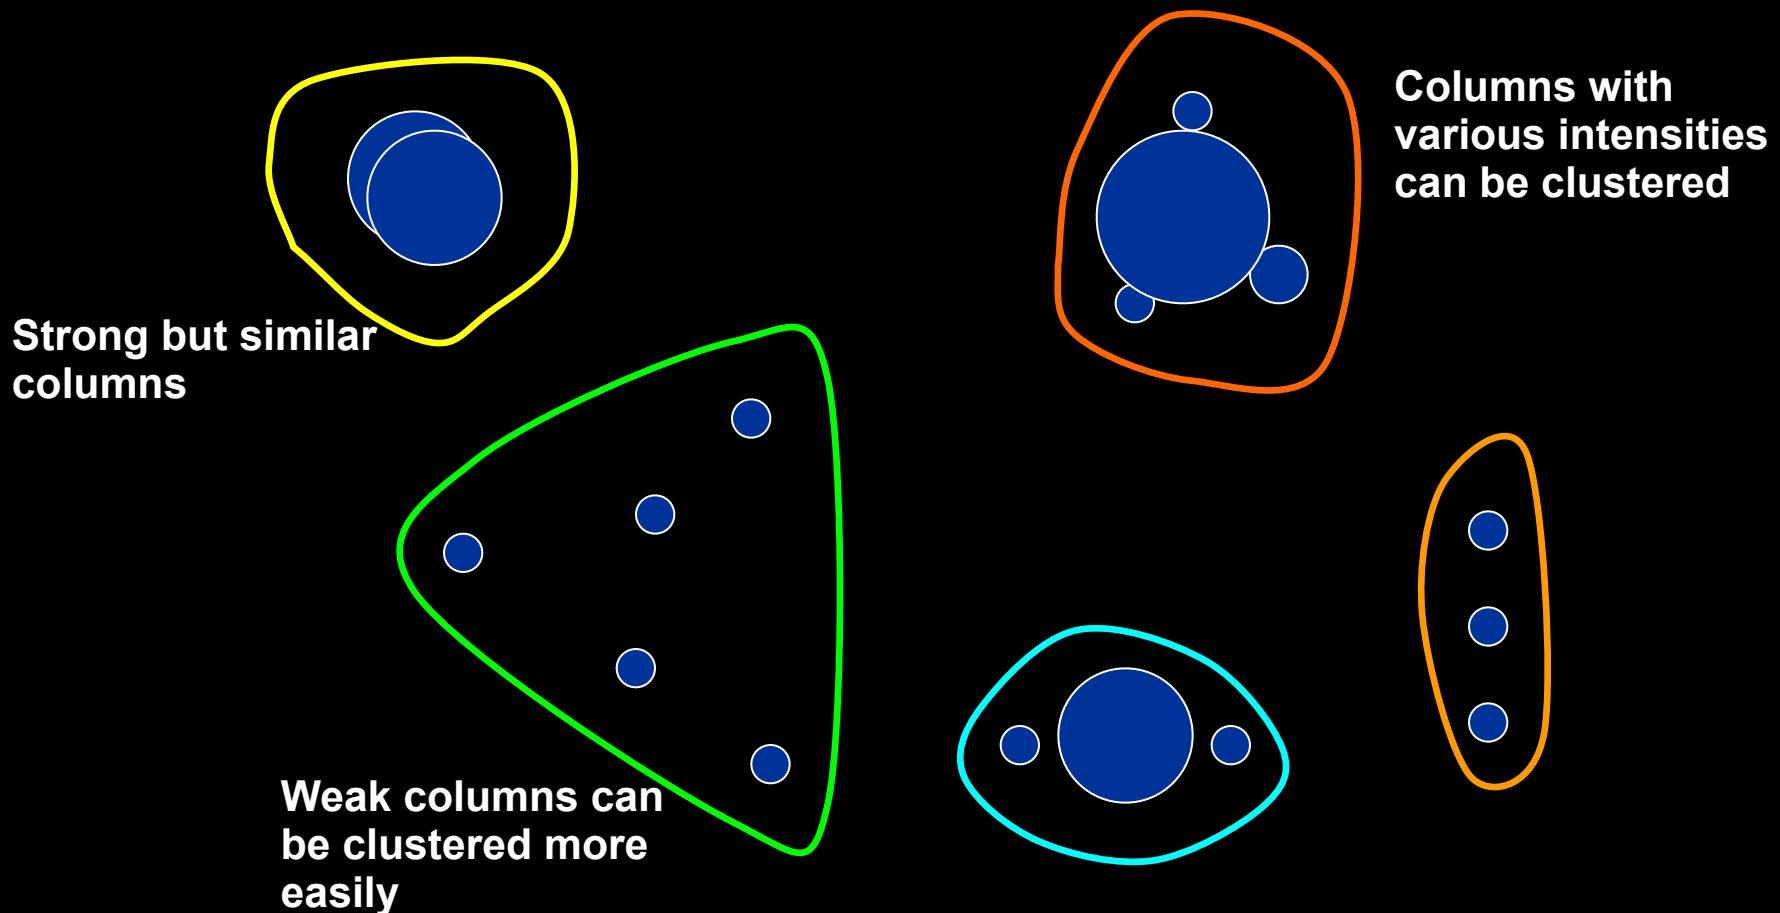
# Clustering Illustration

---



# Clustering Illustration

---



# The Clustering Metric

---

- Minimize:  $\sum_{p=1,\dots,k} cost(C_p)$
- where:  $cost(C) = \sum_{i,j \in C} w_i \ w_j \ \|x_i - x_j\|^2$

# The Clustering Metric

---

- Minimize:  $\sum_{p=1,\dots,k} cost(C_p)$   
  
**total cost of all clusters**
- where:  $cost(C) = \sum_{i,j \in C} w_i \ w_j \ \|x_i - x_j\|^2$

# The Clustering Metric

---

- Minimize:

$$\sum_{p=1,\dots,k} cost(C_p)$$

total cost of all clusters

- where:

$$cost(C) = \sum_{i,j \in C} w_i w_j \|x_i - x_j\|^2$$

cost of a cluster

sum over all pairs in it

norms of the squared distance between reduced columns

between normalized reduced columns

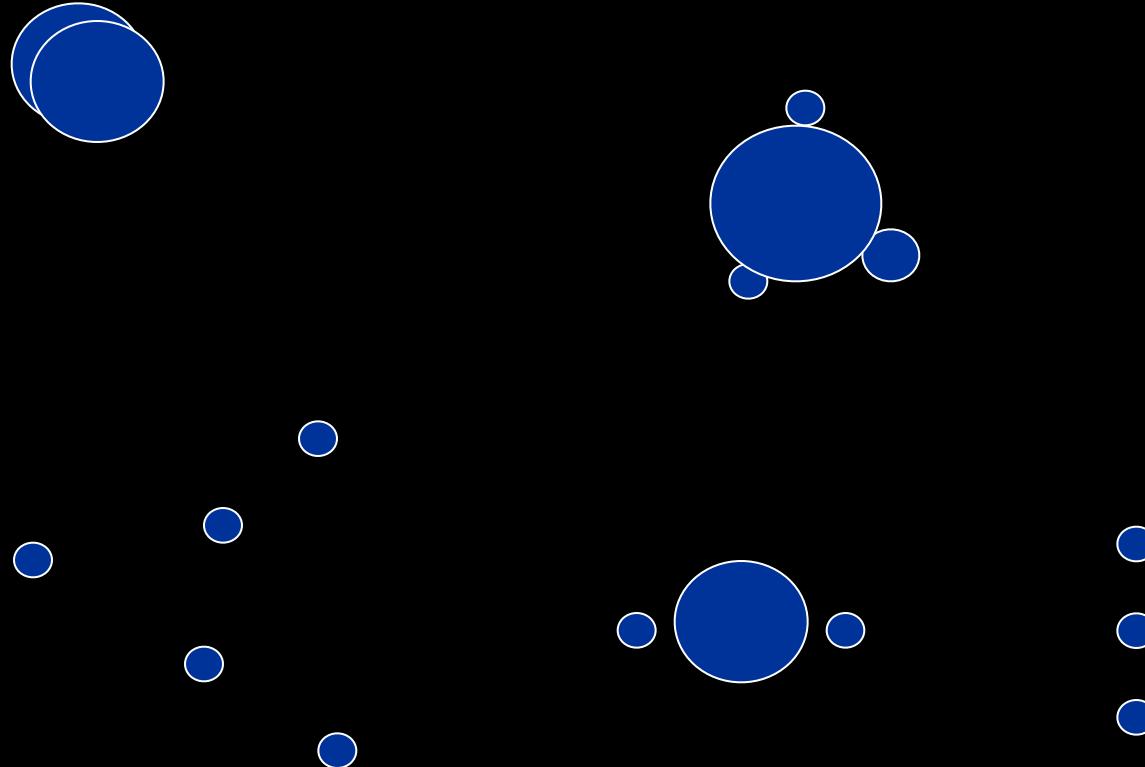
# How to minimize?

---

- Problem is NP-hard
- Not much previous research
- Should handle large input:
  - 100,000 points
  - 1000 clusters
- We introduce 2 heuristics:
  - Random sampling
  - Divide & conquer

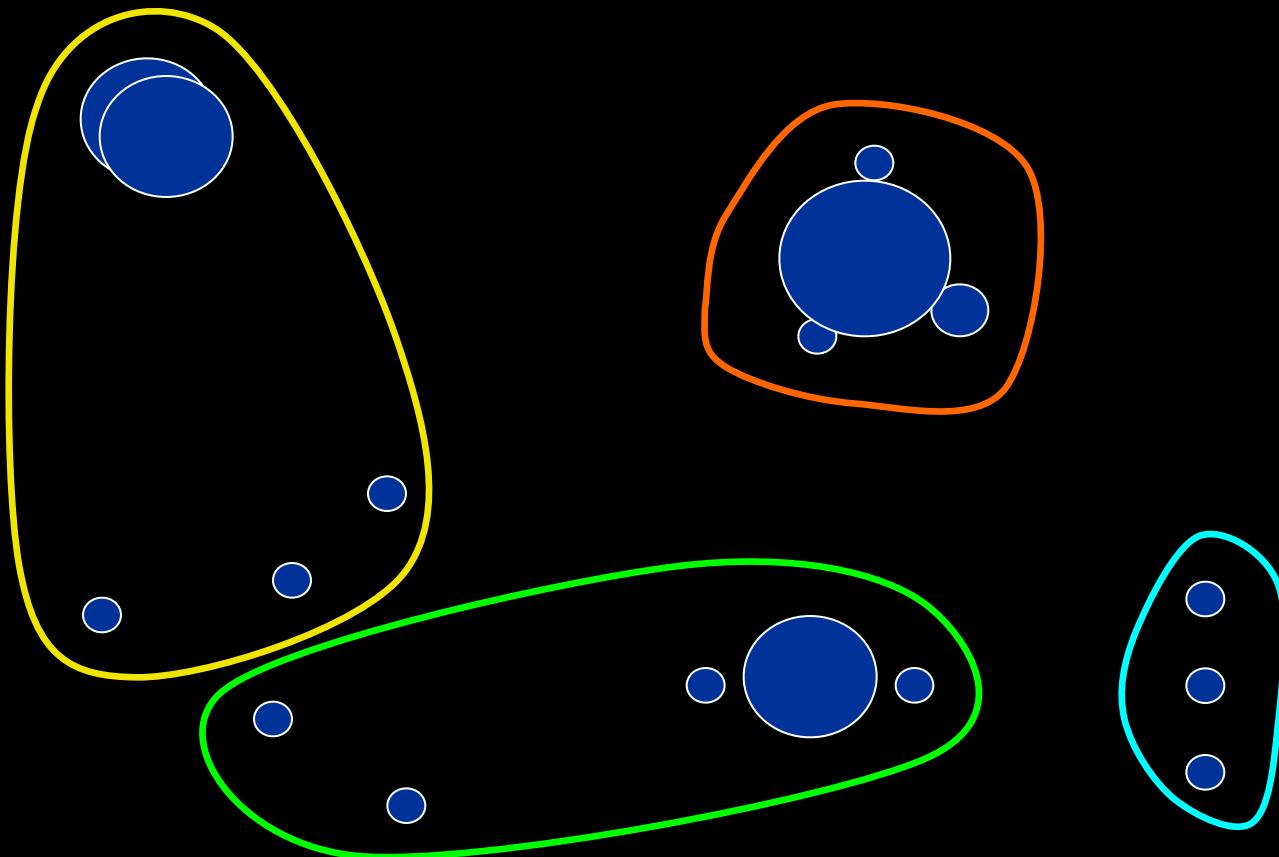
# Clustering by Random Sampling

---



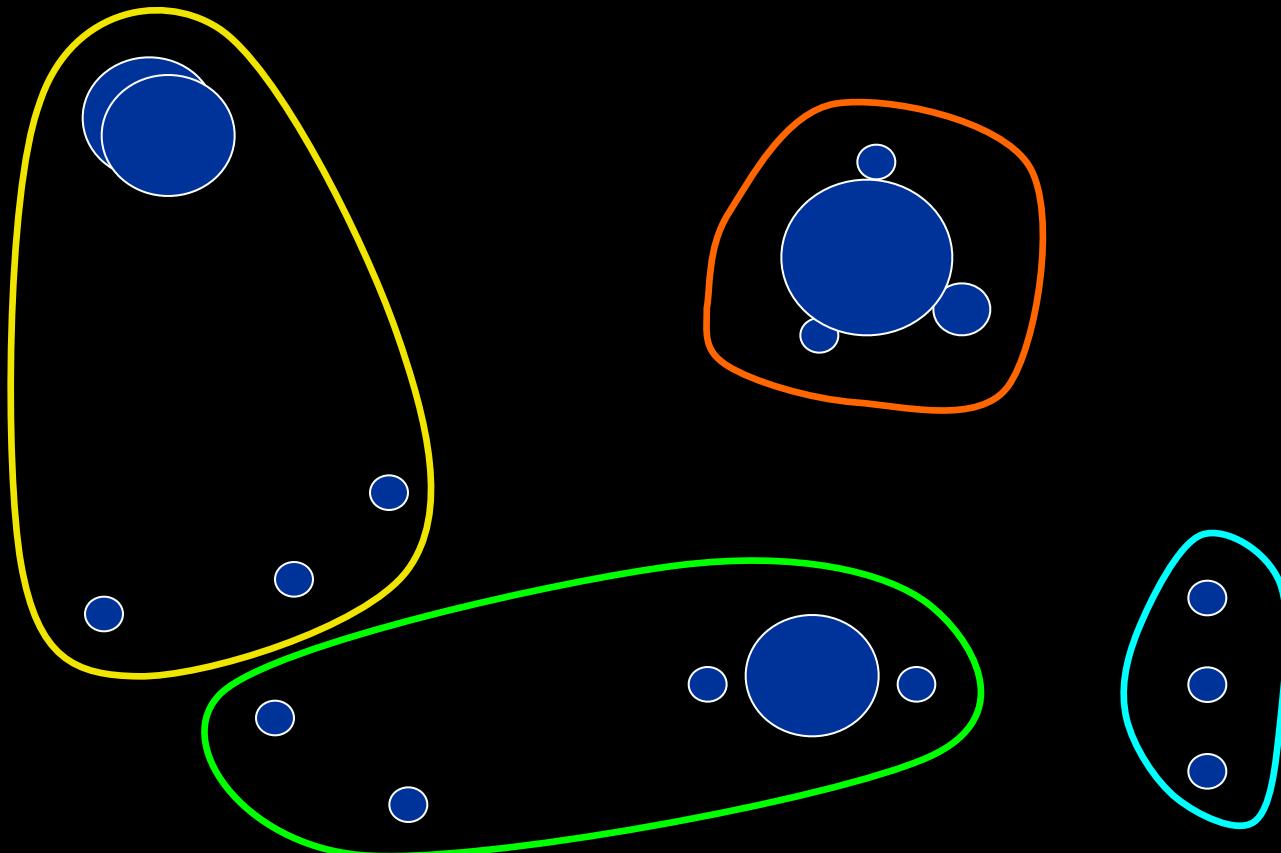
# Clustering by Random Sampling

---



# Clustering by Random Sampling

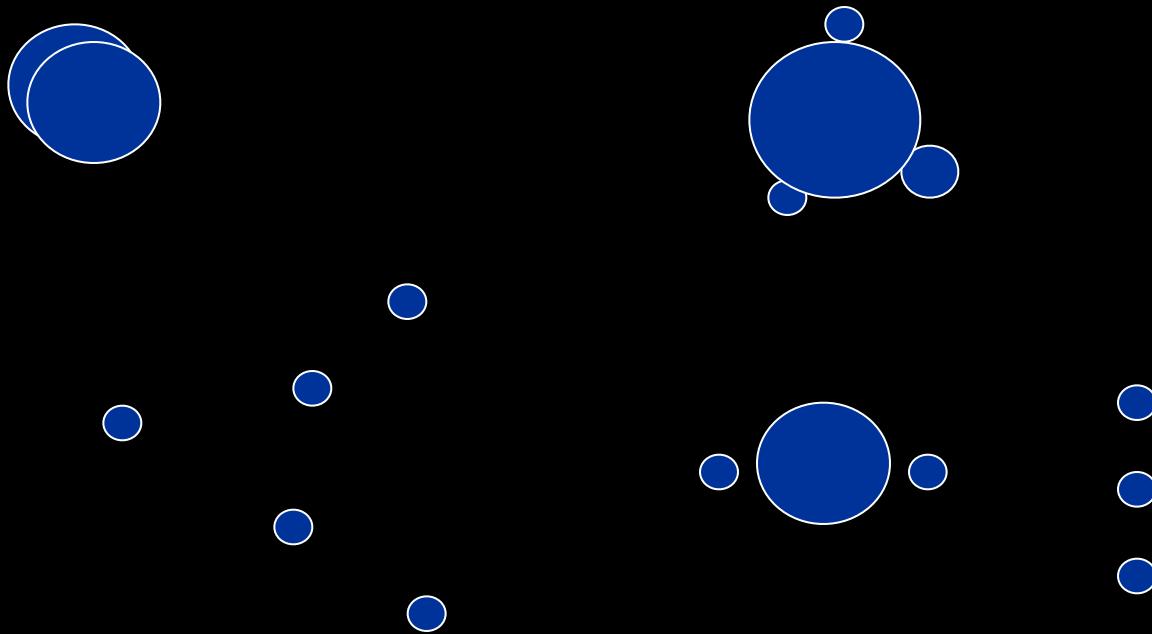
---



- + Very fast (use optimized BLAS)
- Some clusters might be too small / large

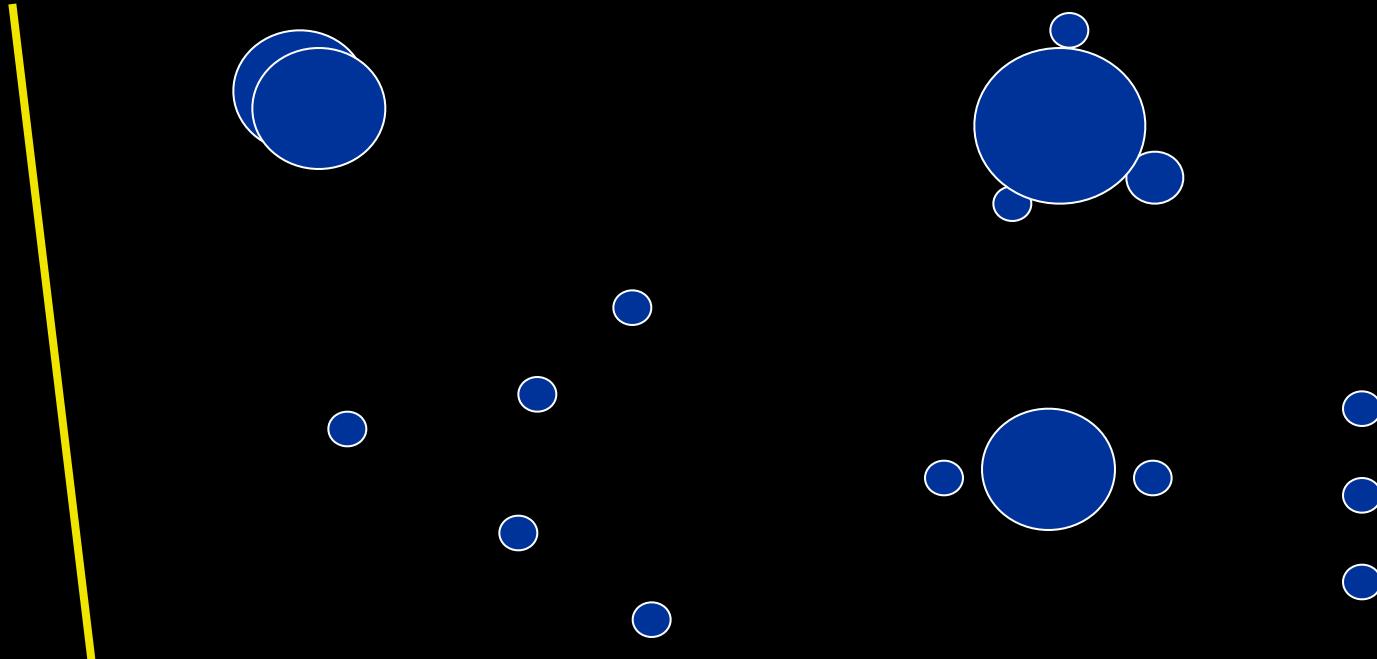
# Clustering by Divide & Conquer

---



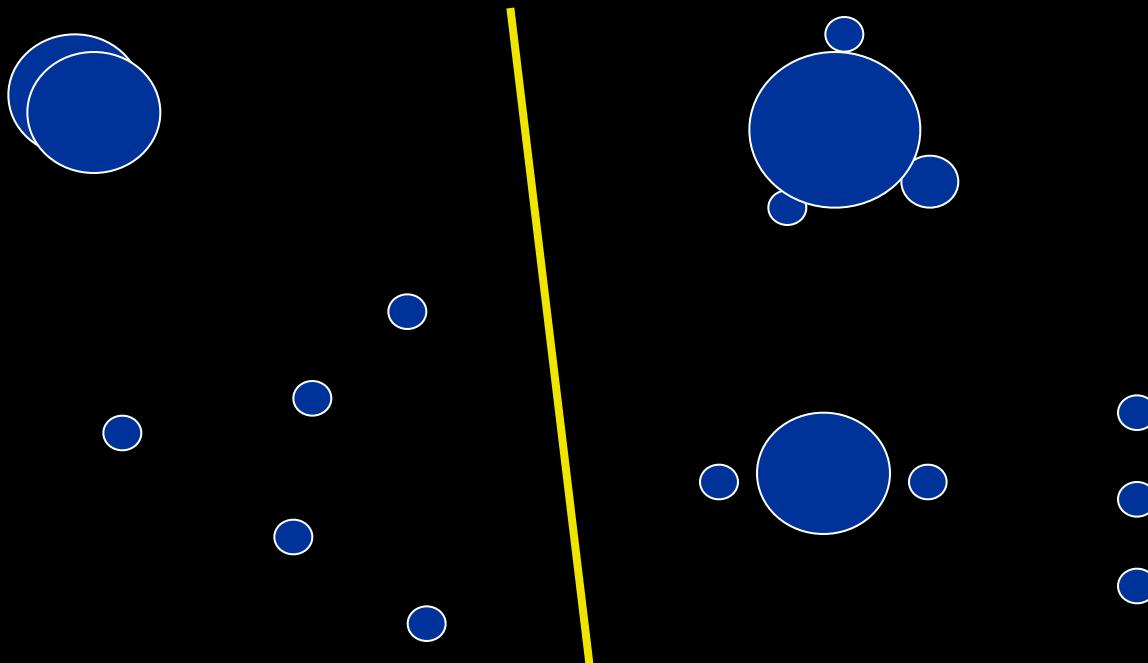
# Clustering by Divide & Conquer

---



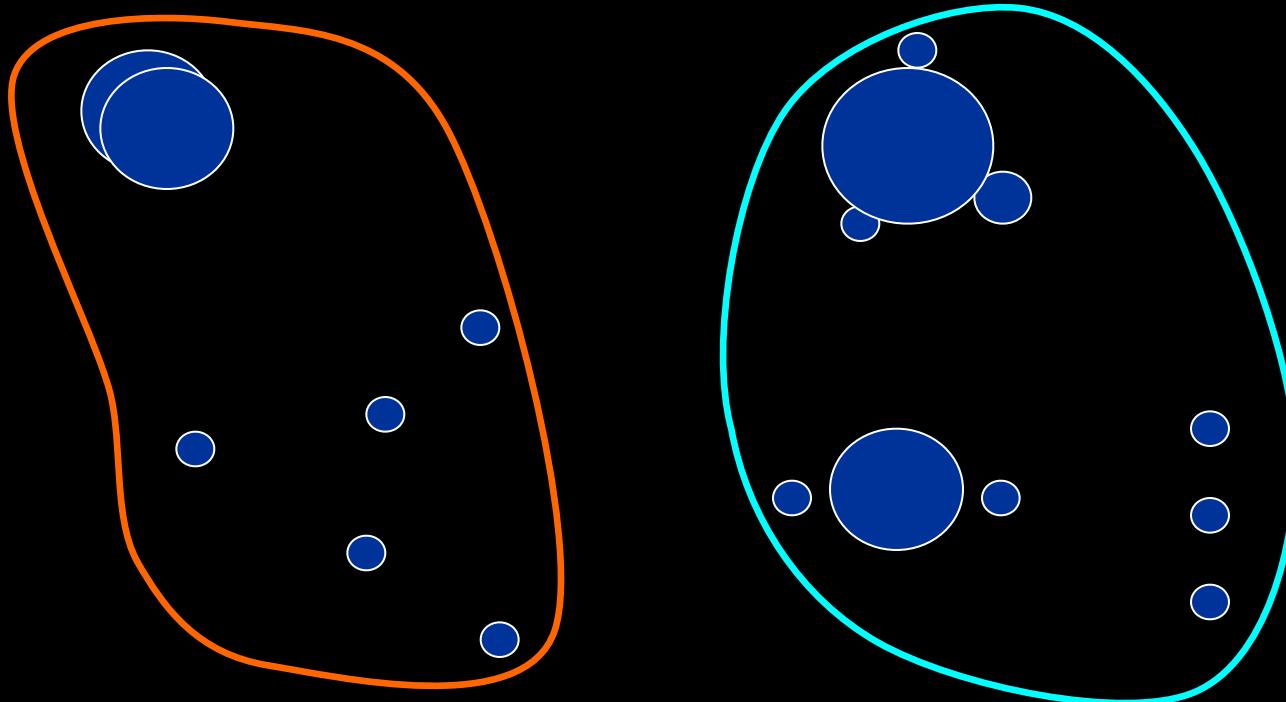
# Clustering by Divide & Conquer

---



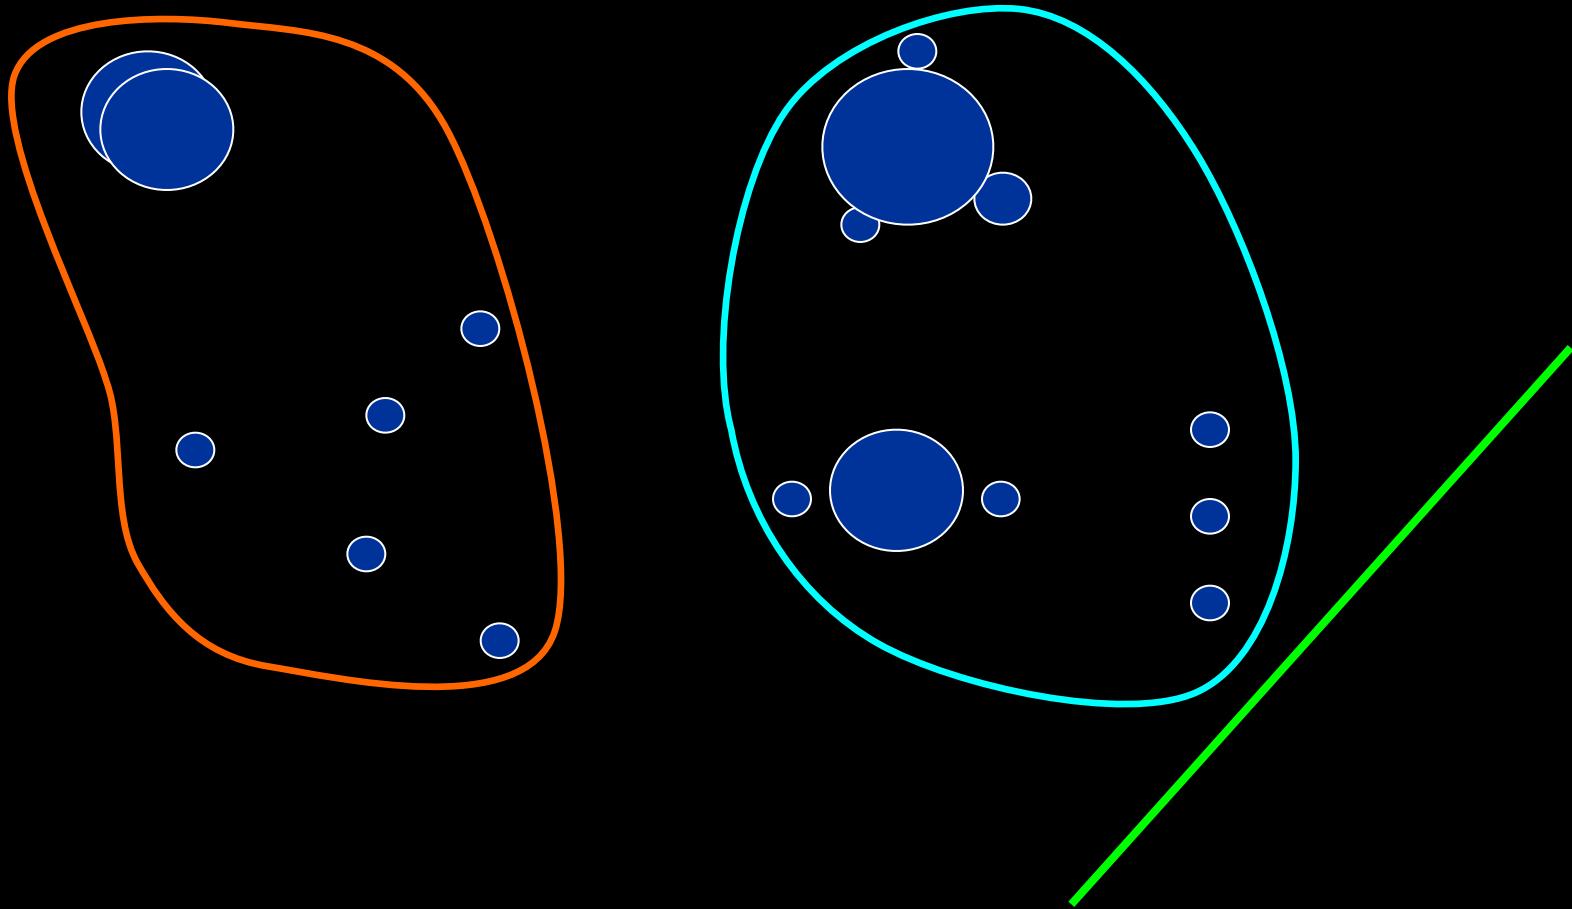
# Clustering by Divide & Conquer

---



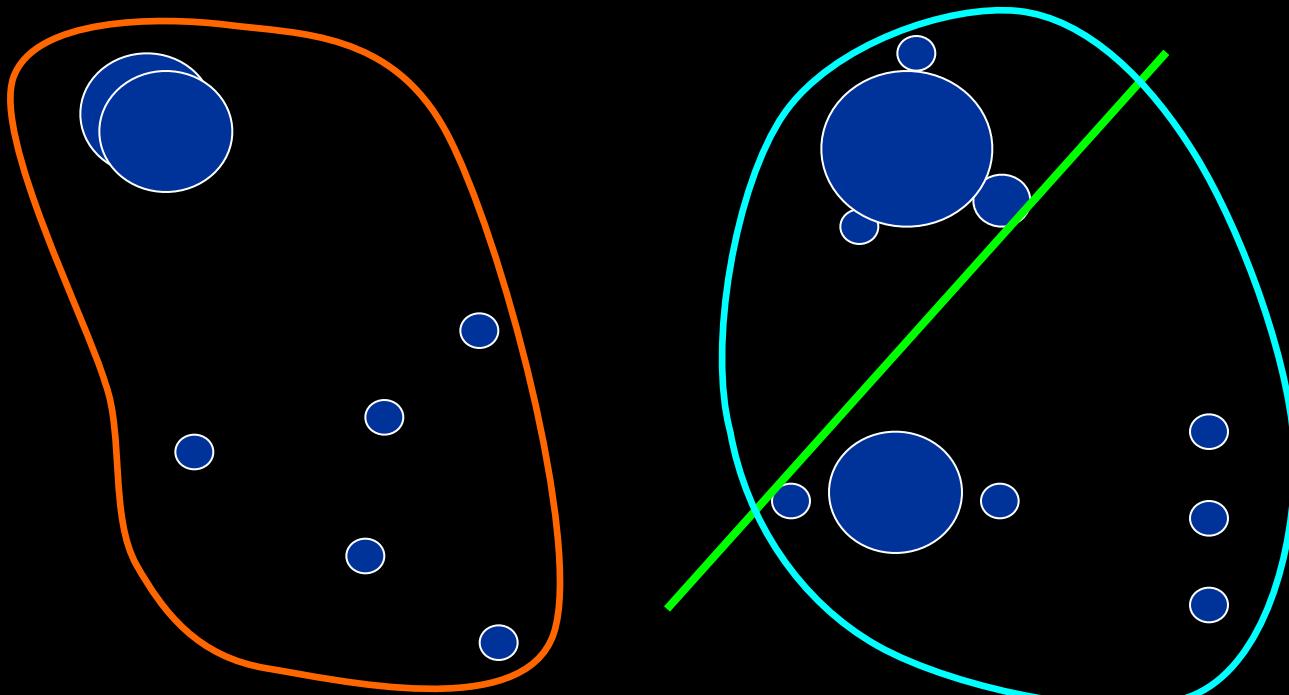
# Clustering by Divide & Conquer

---



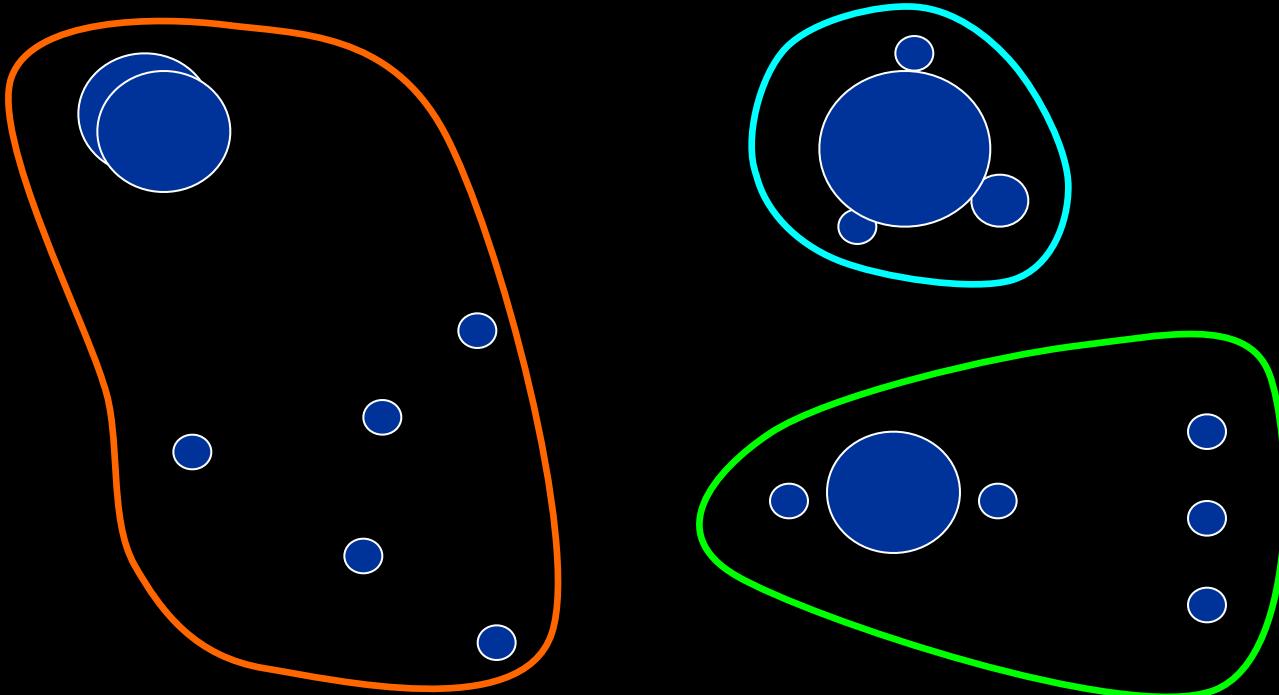
# Clustering by Divide & Conquer

---



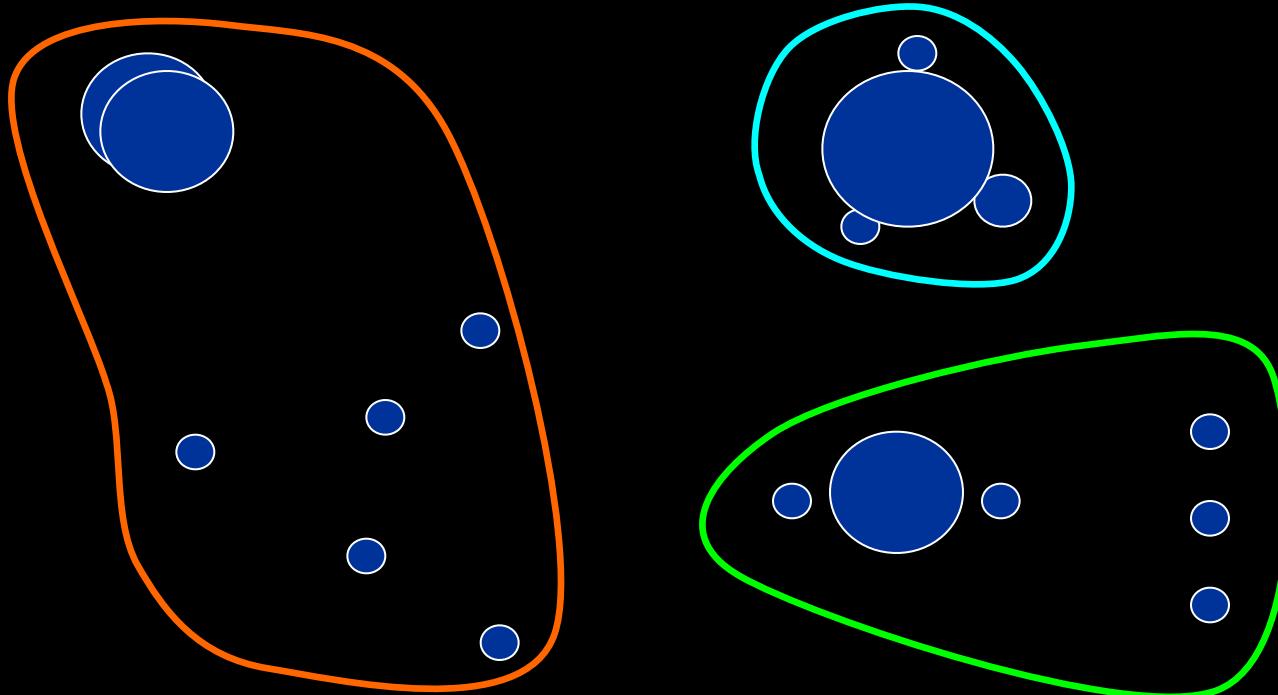
# Clustering by Divide & Conquer

---



# Clustering by Divide & Conquer

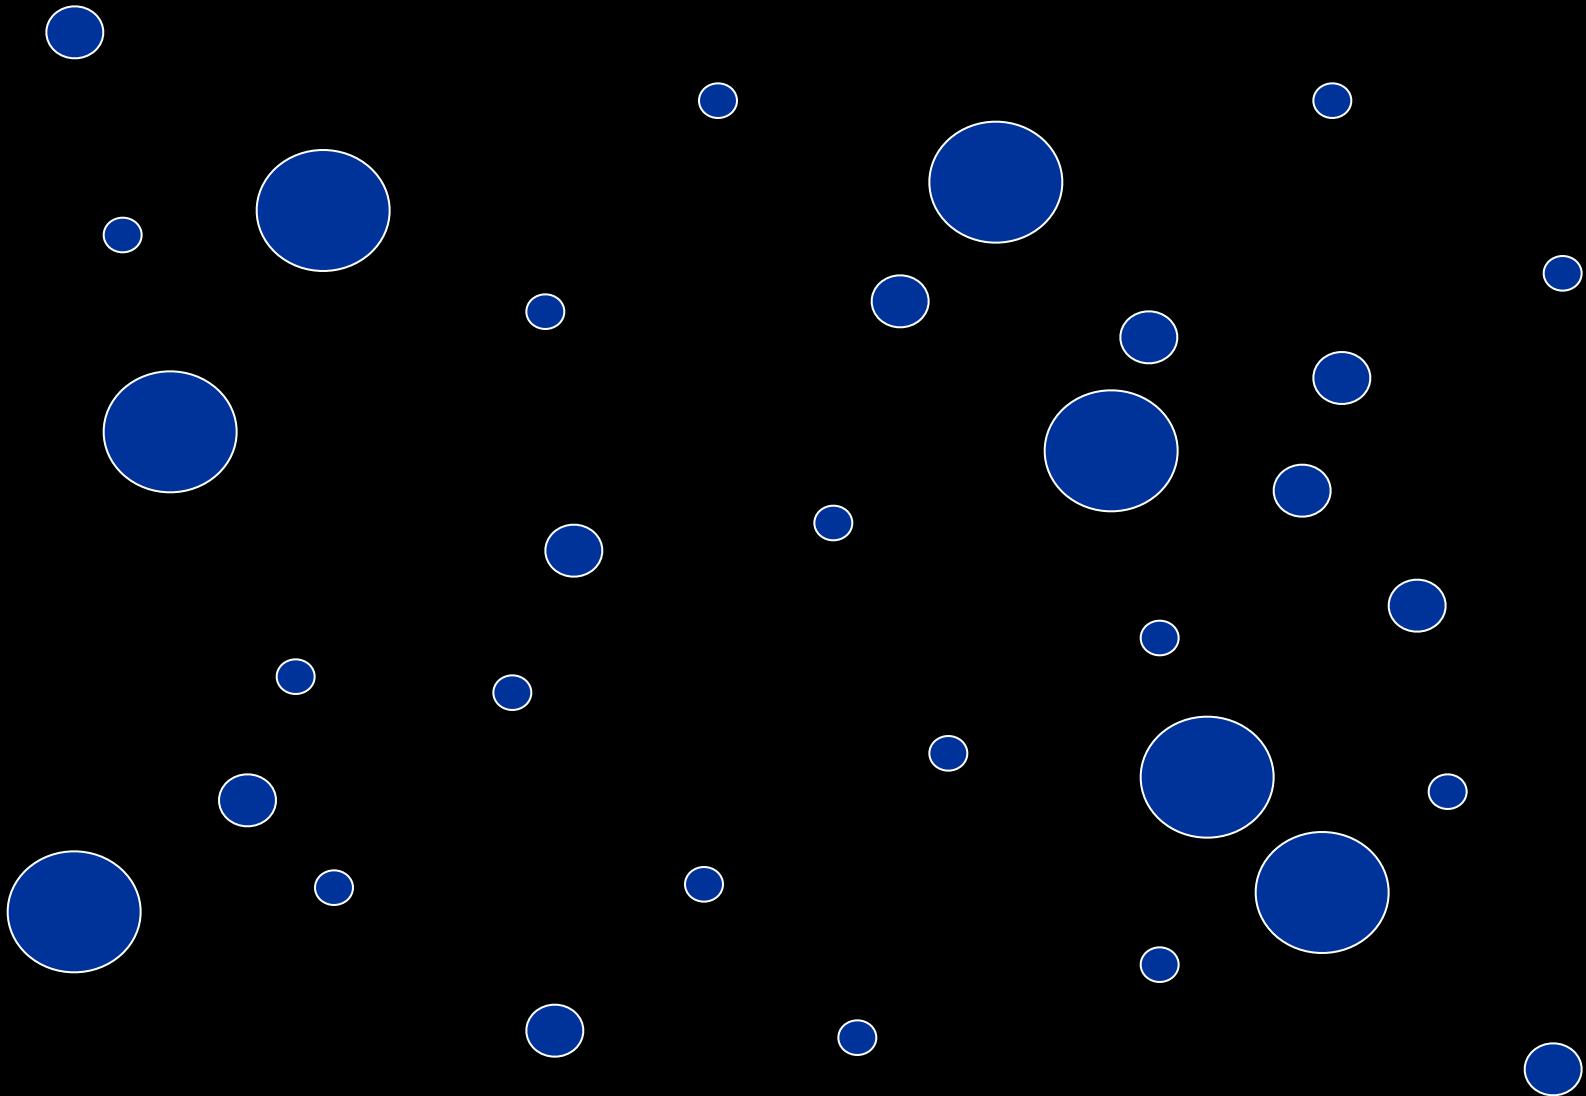
---



- Splitting small clusters is fast**
- Splitting large clusters is slow**

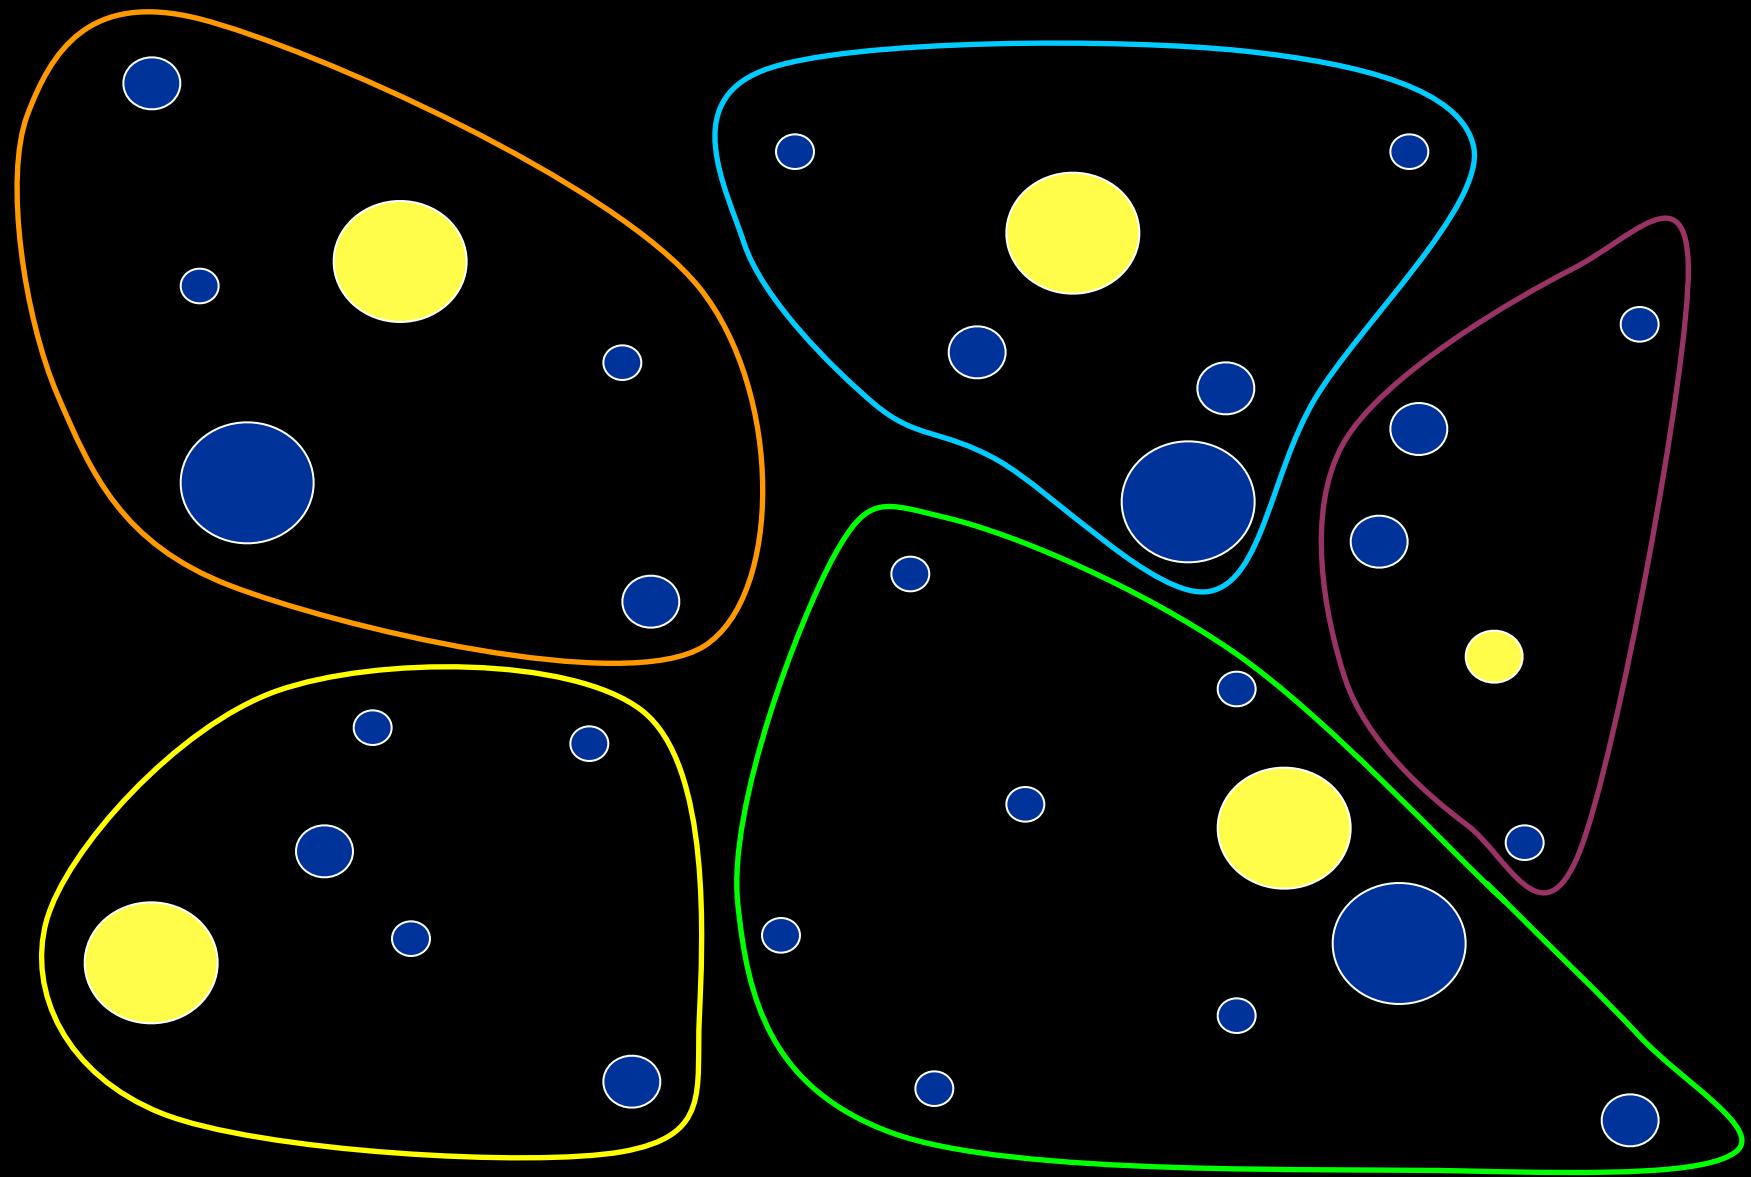
# Combined Clustering Algorithm

---



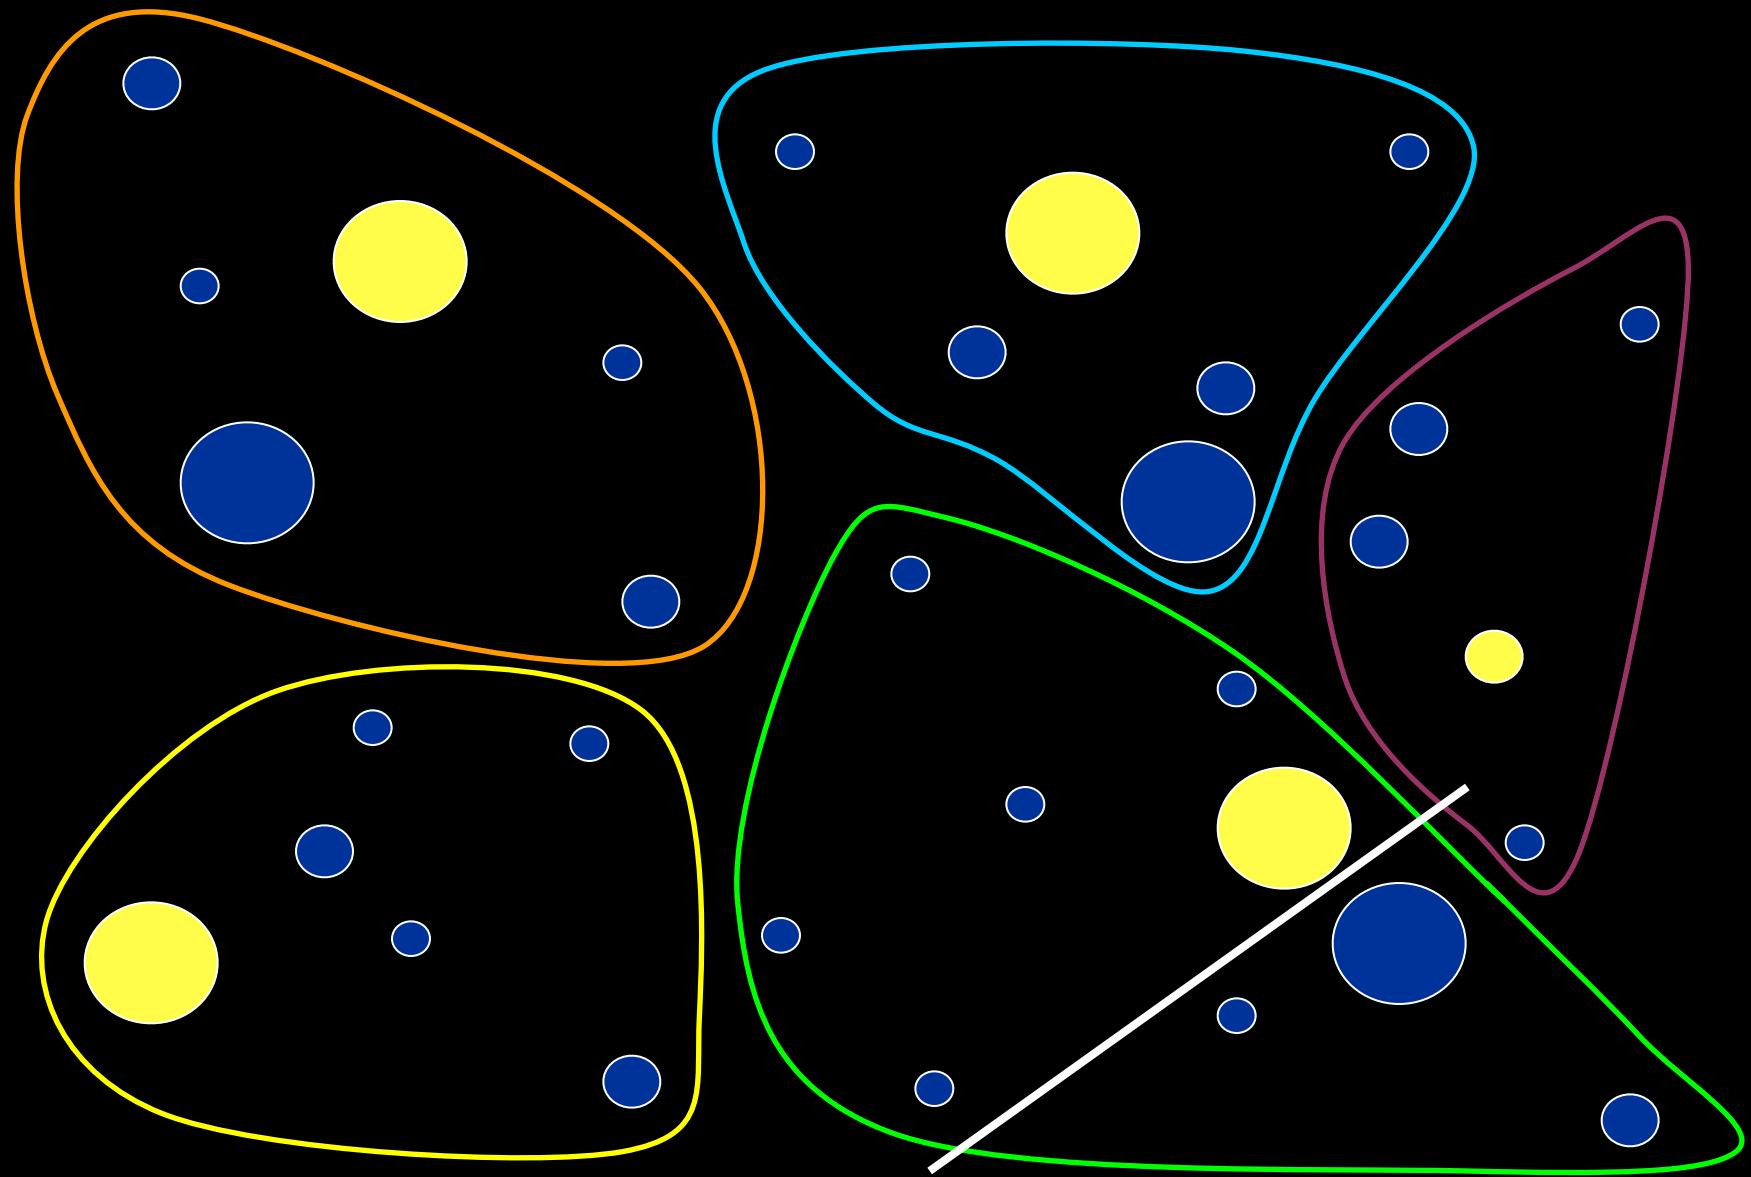
# Combined Clustering Algorithm

---



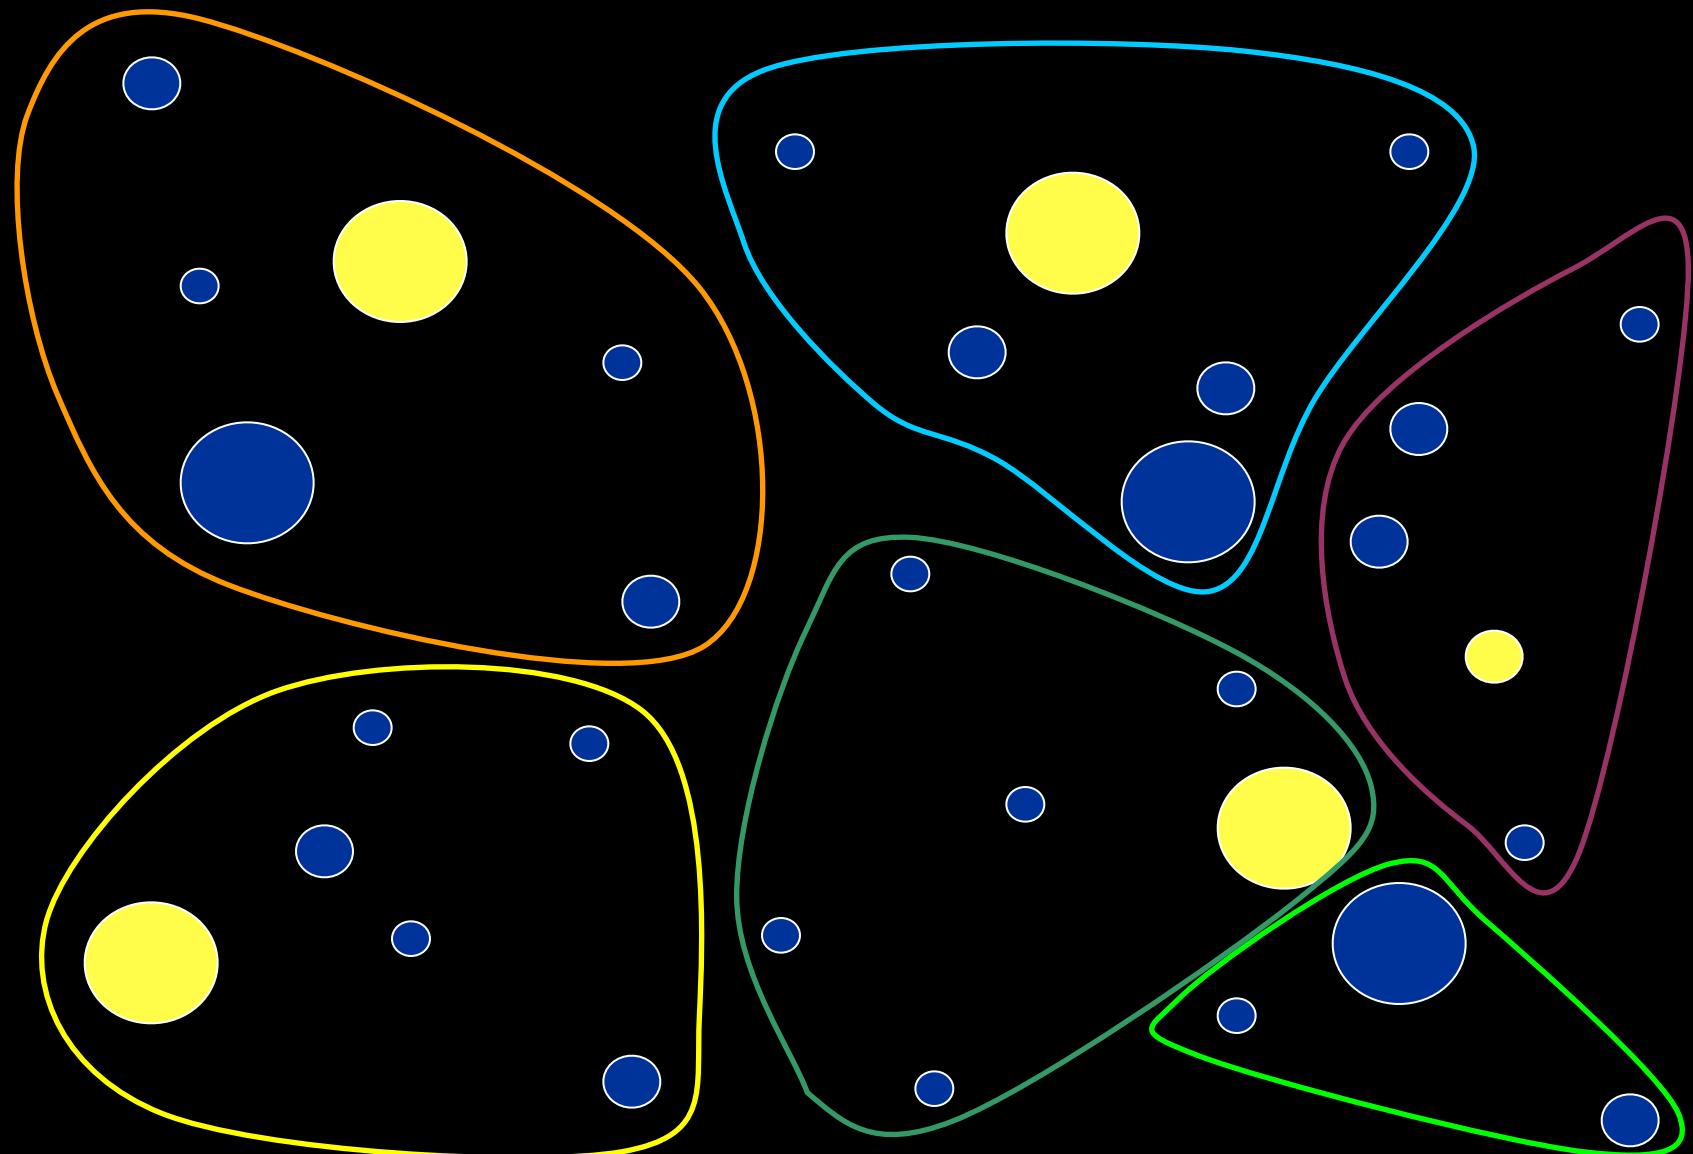
# Combined Clustering Algorithm

---



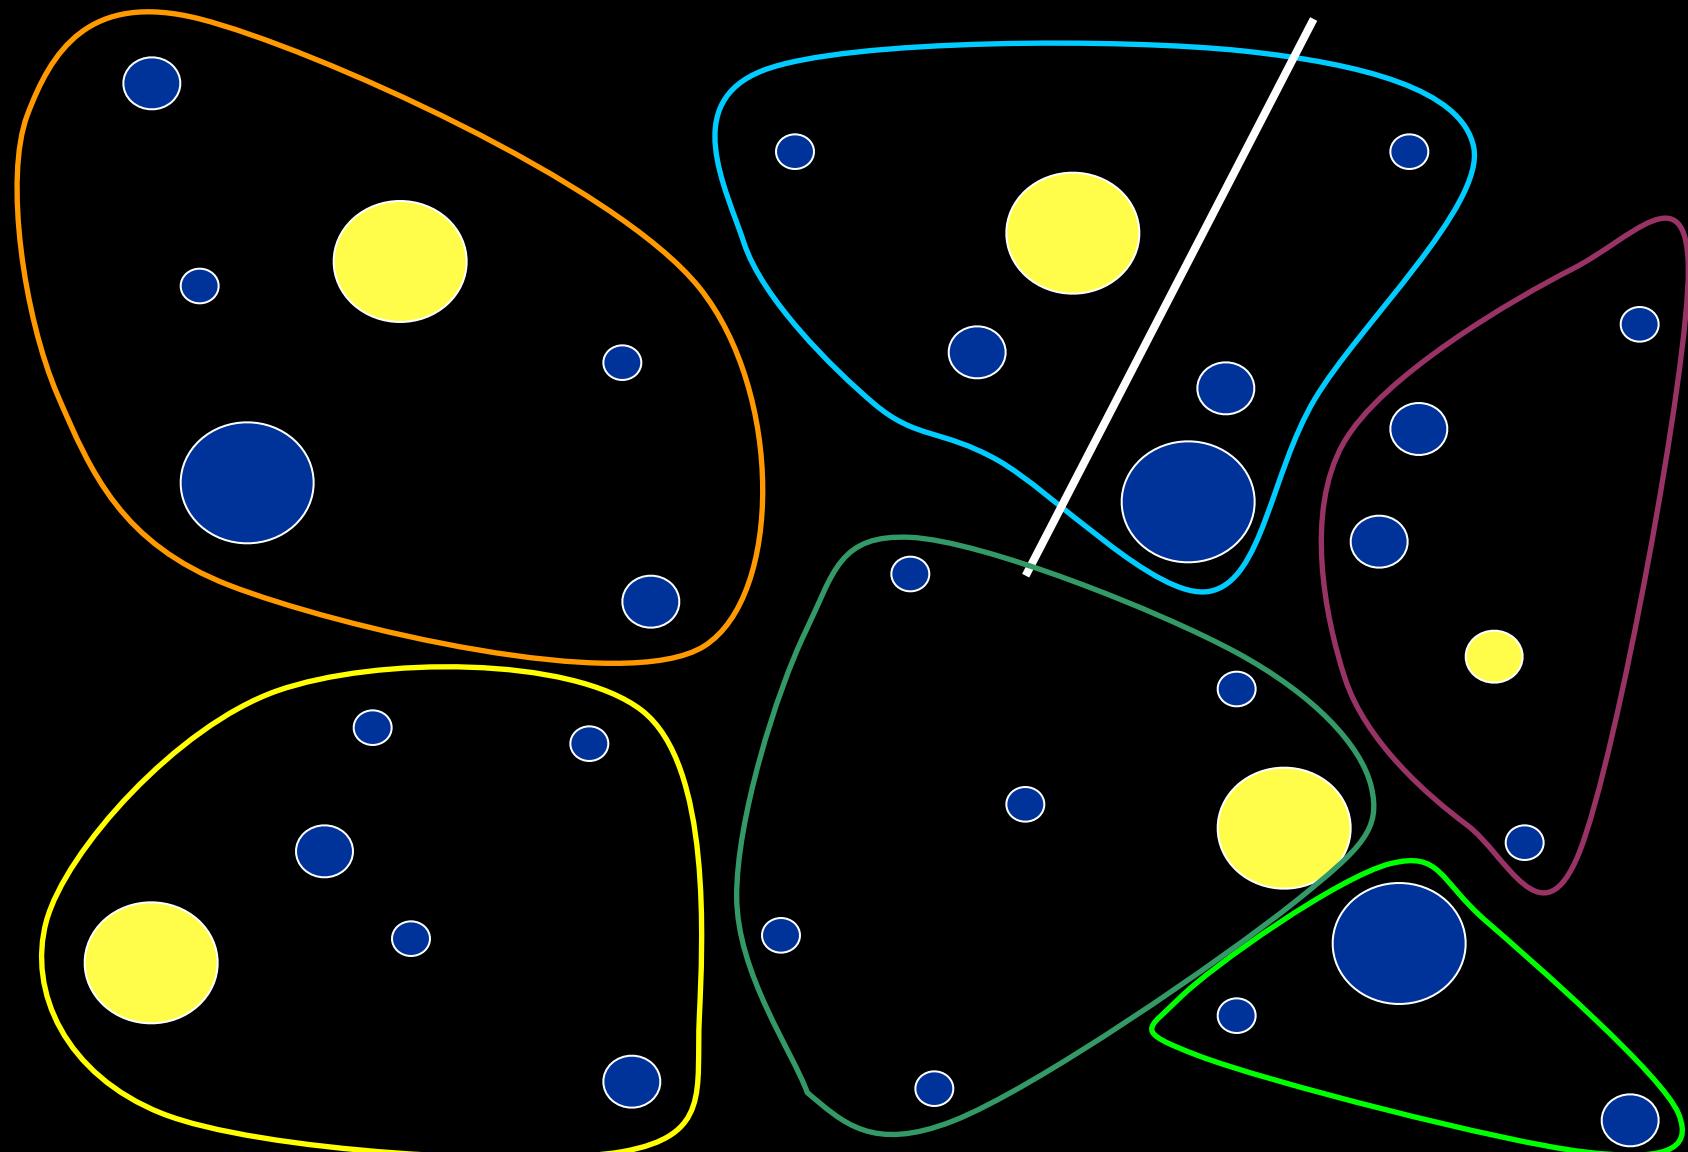
# Combined Clustering Algorithm

---



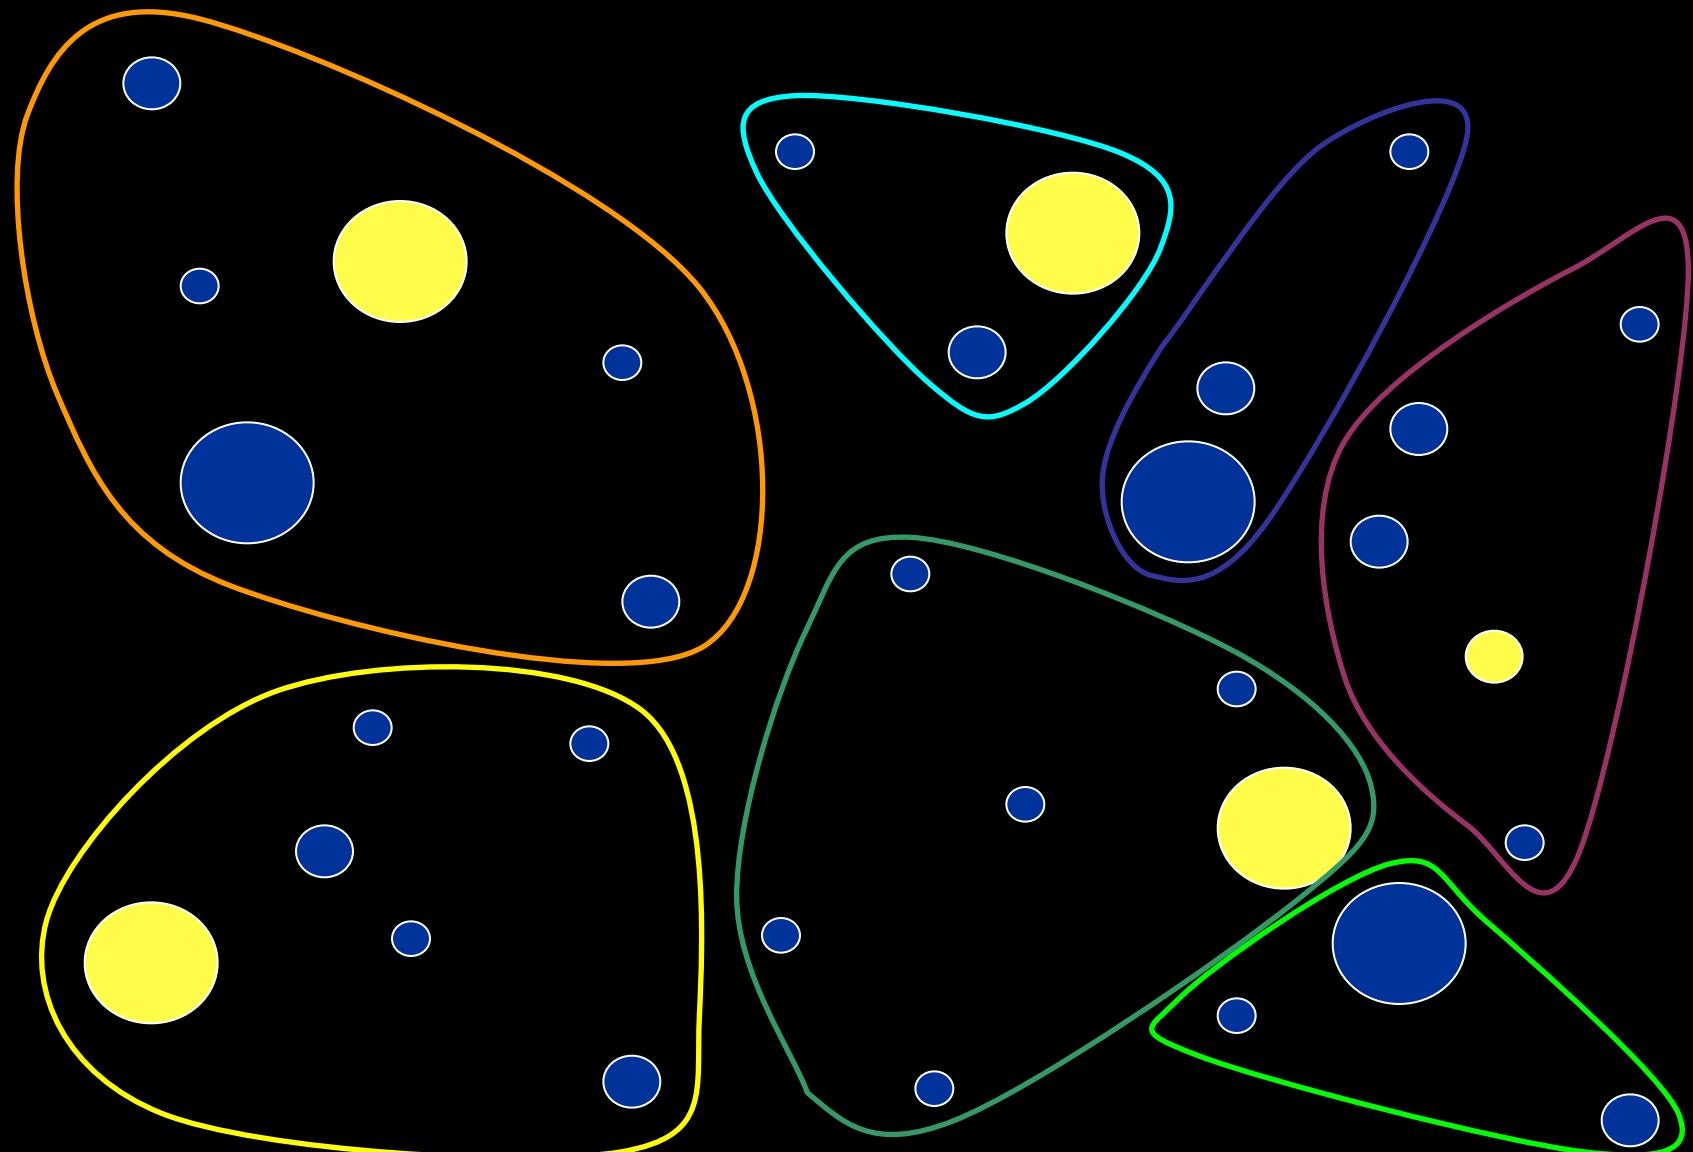
# Combined Clustering Algorithm

---



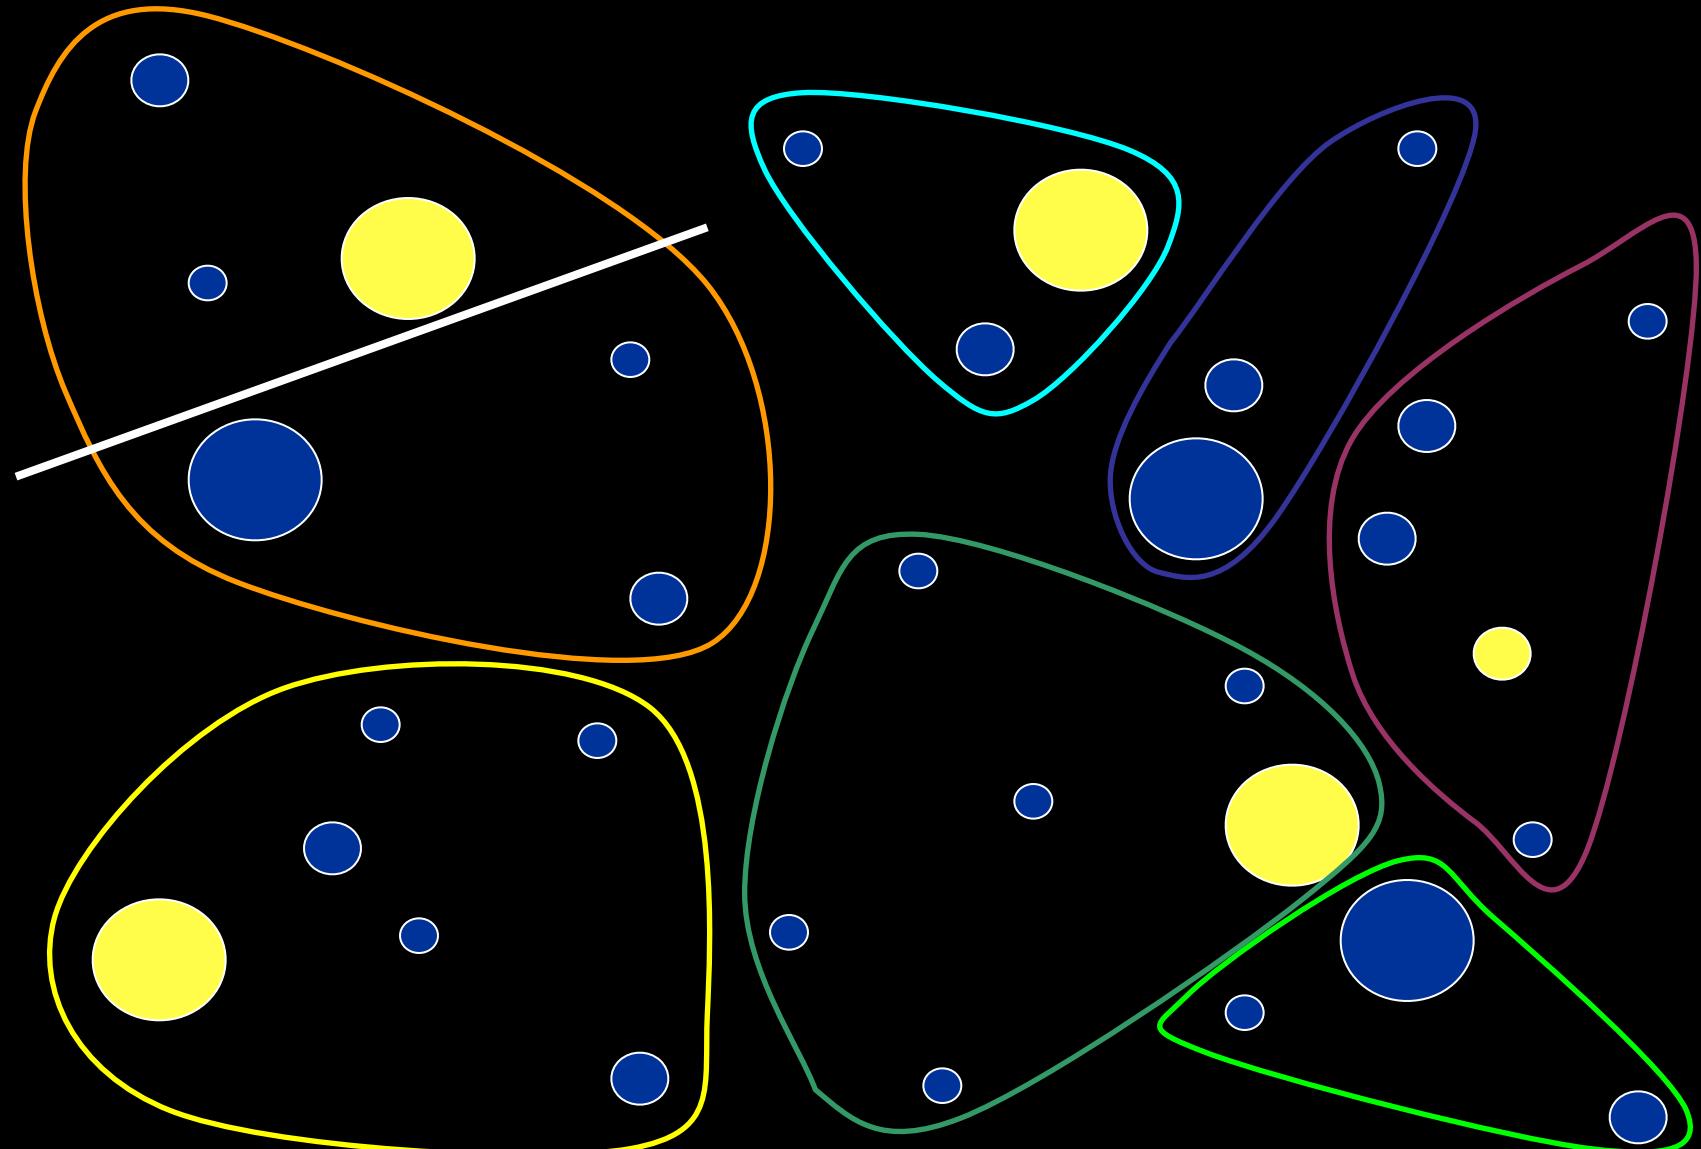
# Combined Clustering Algorithm

---



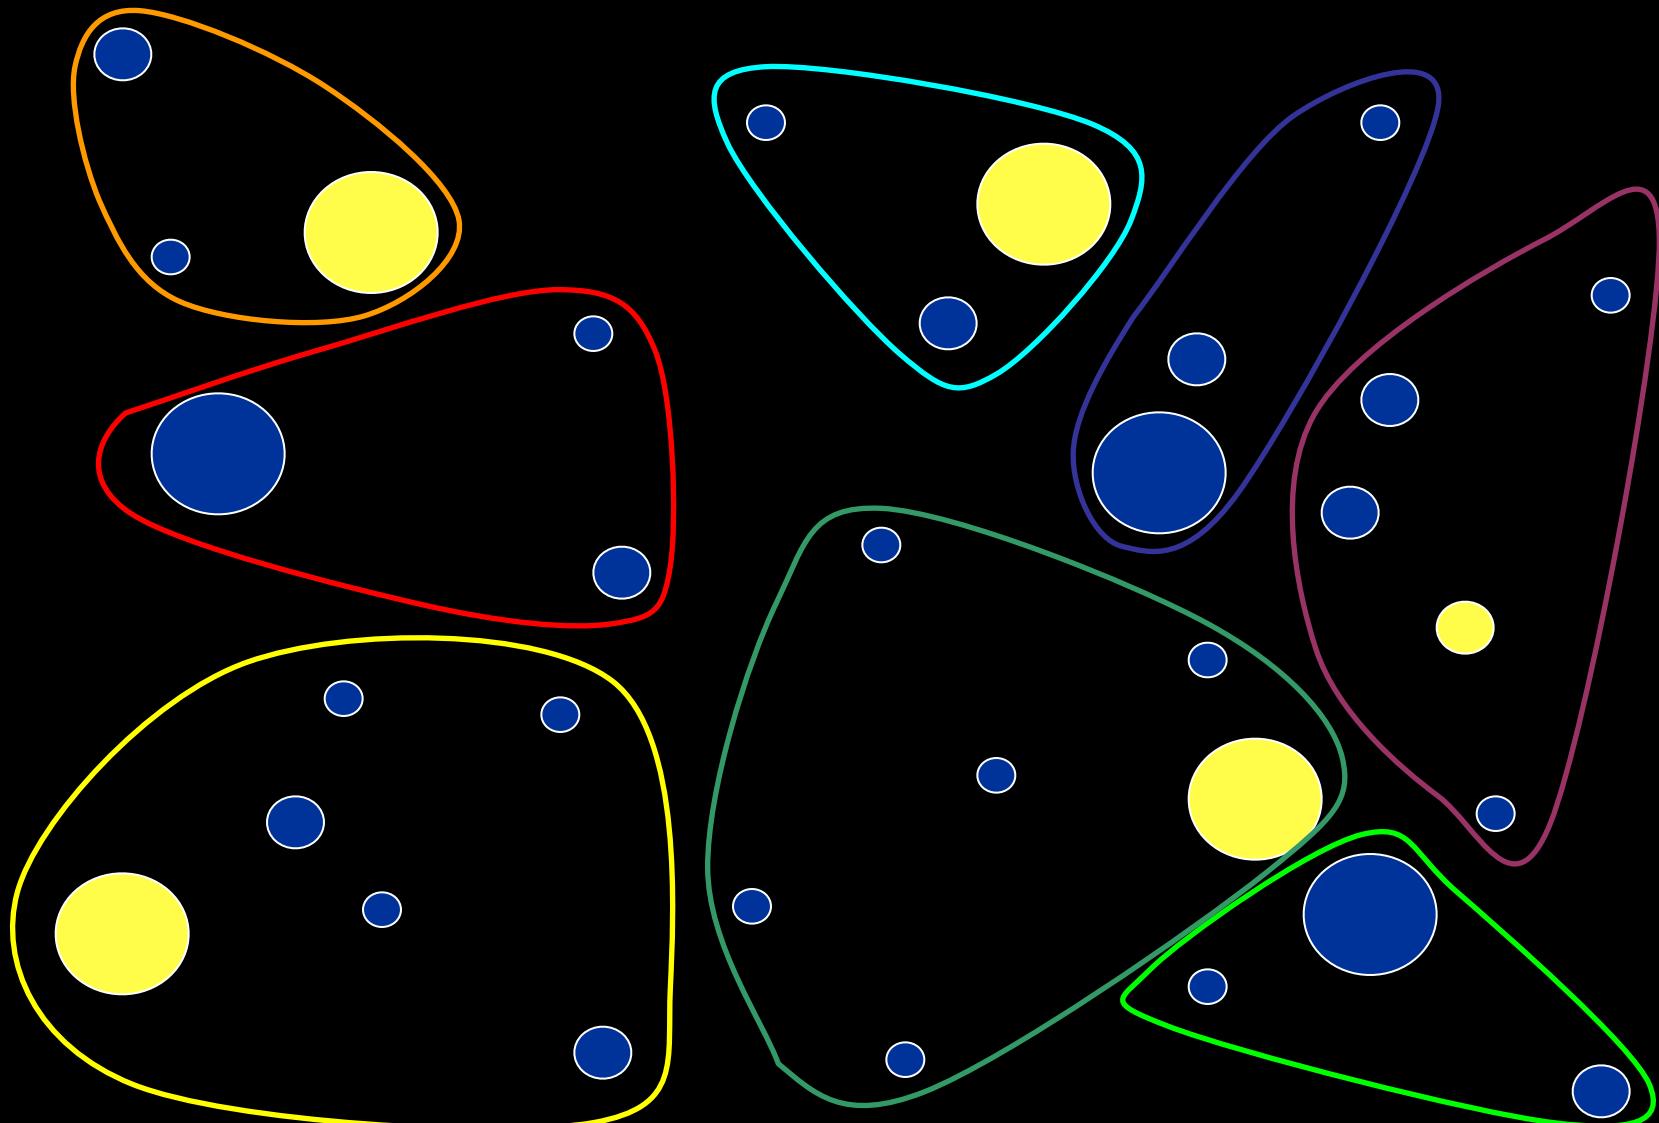
# Combined Clustering Algorithm

---

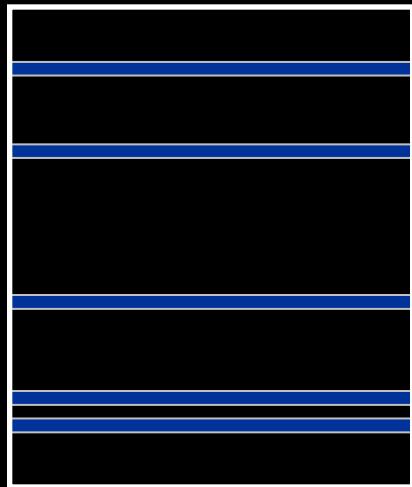


# Combined Clustering Algorithm

---



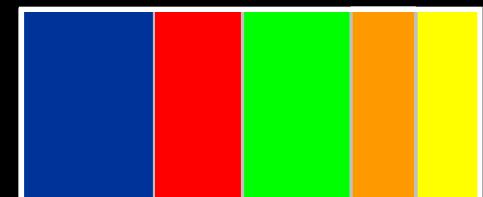
# Full Algorithm



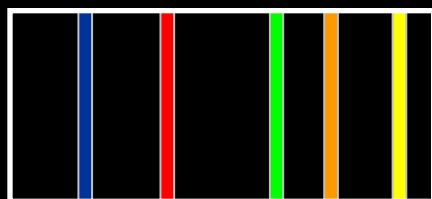
Compute rows  
(GPU)



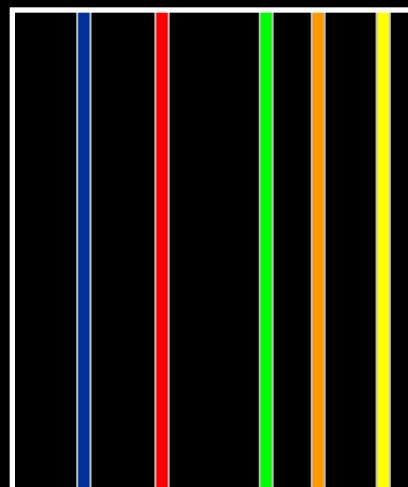
Assemble rows into  
reduced matrix



Cluster reduced  
columns



Choose  
representatives



Compute columns  
(GPU)



Weighted sum

# Results

---

- We show 5 scenes:



Kitchen



Temple



Trees



Bunny

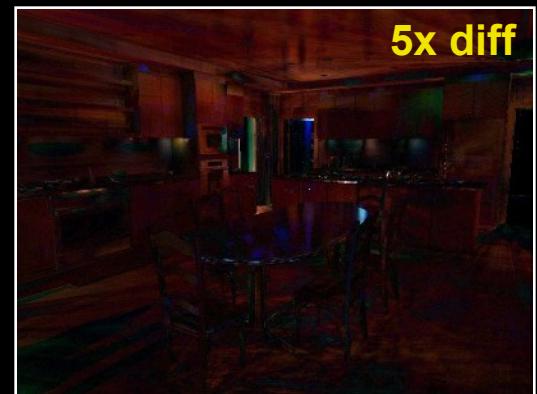


Grand Central

- Show reference and 5x difference image
- All scenes have 100,000+ lights
- Timings
  - NVidia GeForce 8800 GTX
  - Light / surface sample creation not included

# Results: Kitchen

- 388k polygons
- Mostly indirect illumination
- Glossy surfaces
- Indirect shadows



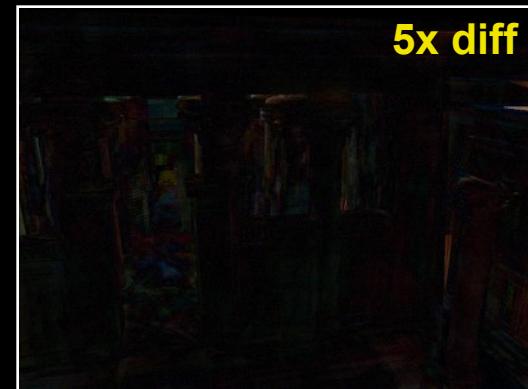
**Our result: 13.5 sec**  
(432 rows + 864 columns)



**Reference: 13 min**  
(using all 100k lights)

# Results: Temple

- 2.1m polygons
- Mostly indirect & sky illumination
- Indirect shadows

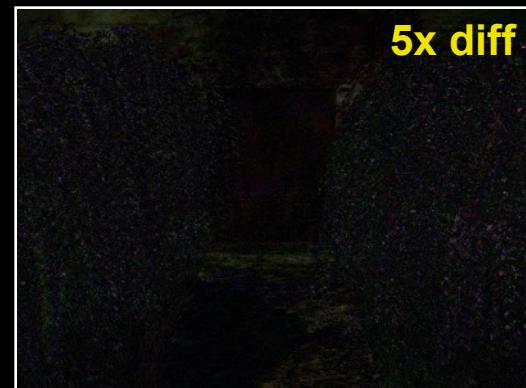


**Our result: 16.9 sec**  
(300 rows + 900 columns)

**Reference: 20 min**  
(using all 100k lights)

# Results: Trees

- 328k polygons
- Complex incoherent geometry



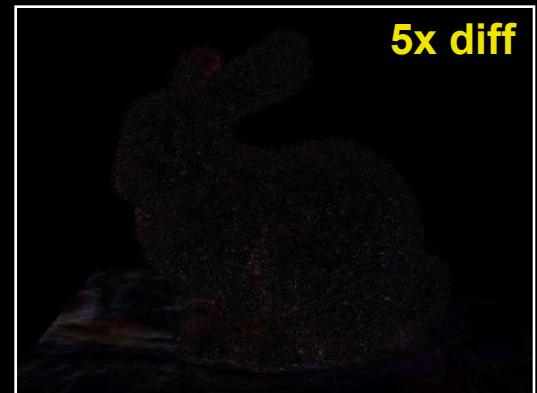
**Our result: 2.9 sec**  
(100 rows + 200 columns)



**Reference: 14 min**  
(using all 100k lights)

# Results: Bunny

- 869k polygons
- Incoherent geometry
- High-frequency lighting
- Kajiya-Kay hair shader



**Our result: 3.8 sec**  
(100 rows + 200 columns)



**Reference: 10 min**  
(using all 100k lights)

# Results: Grand Central

- 1.5m polygons
- Point lights between stone blocks



**Our result: 24.2 sec**  
(588 rows + 1176 columns)



**Reference: 44 min**  
(using all 100k lights)

# The Value of Exploration

---



**Our result**

(432 rows + 864 columns)



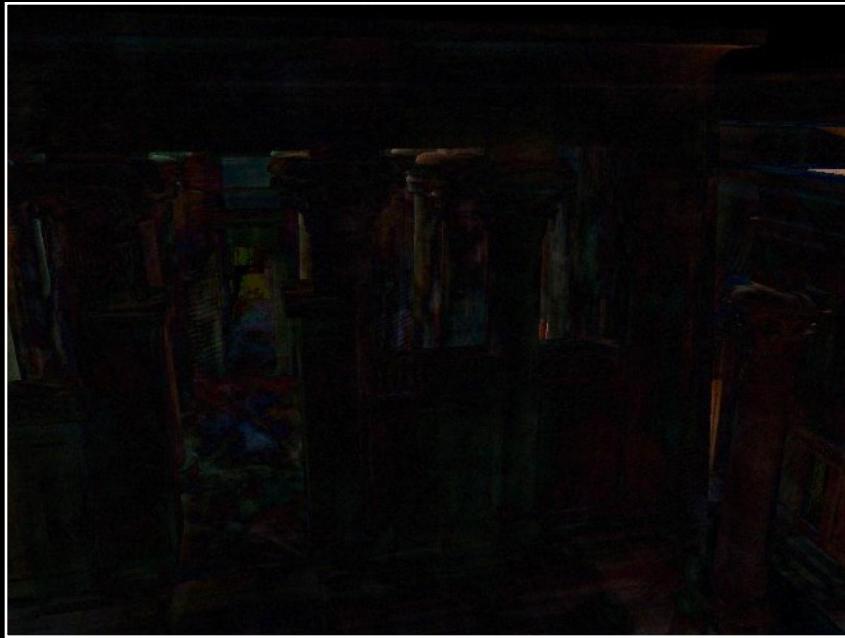
**No exploration**

(Using 1455 lights)

**Equal time comparison**

# The Value of Exploration

---



**Our result**



**No exploration**

**Equal time comparison:  
5x difference from  
reference**

# Conclusion

---

- Fast, high quality approximation for many lights
  - GPU-oriented
  - Sample rows to explore low-rank structure
  - Sample well-chosen columns
- Application: Preview for lighting design
  - Indirect illumination
  - Environment maps
  - Arbitrary lights and shaders

# Future Work

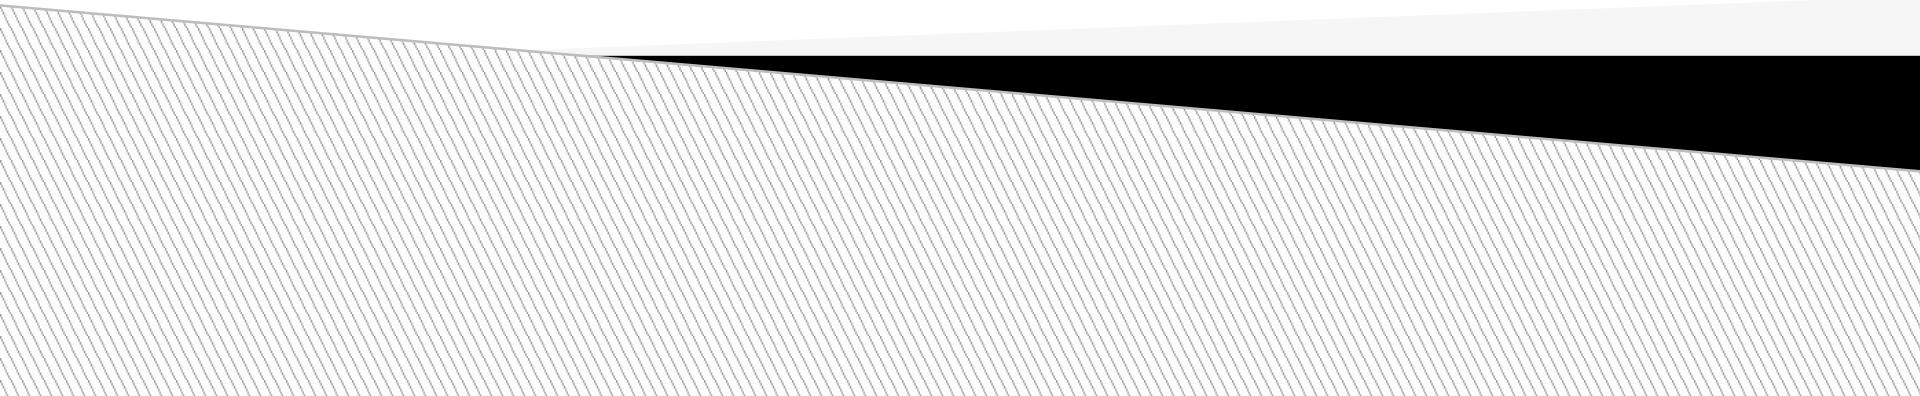
---

- How many rows + columns?
  - Pick automatically
- Row / column alternation
- Progressive algorithm:
  - stop when user likes the image
- Render multiple frames at once?

# Row-Column Matrix Sampling

- ▶ Visible surfaces are diffuse / low specular
  - Problems with glossy/specular surfaces
  - Extensions to spherical light sources
- ▶ Pre computations (for the light sources)
- ▶ Light transport operator
  - Part of it: last transport step
  - Operator compression
  - Highly efficient. Why?

# Global Illumination II



Global Illumination II

The End