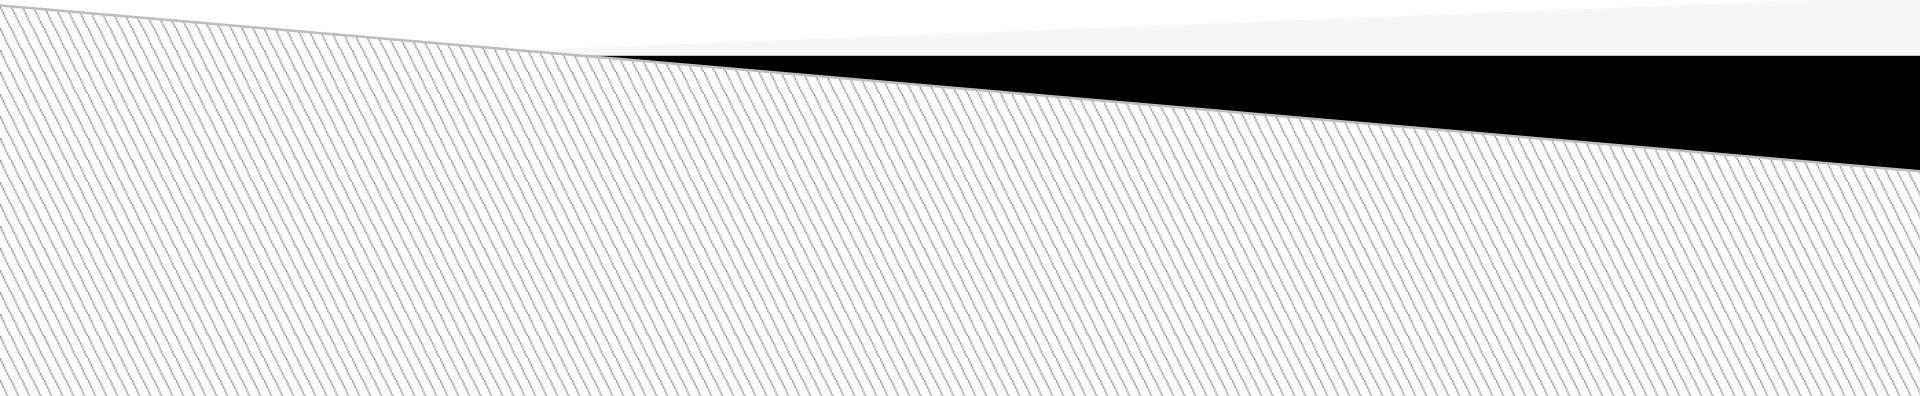


Shadows



Light and shadows

- ▶ Shadows increase realism:



Zaxxon (1982)

Cry Engine



Light and shadows

- ▶ Shadows increase realism
- ▶ Shadows help you perceive:
 - hidden objects



© 2003 - Artis



© 2003 - Artis

Light and shadows

- ▶ Shadows increase realism
- ▶ Shadows help you perceive:
 - hidden objects
 - the relative position of objects



Light and shadows

- ▶ Shadows increase realism
- ▶ Shadows help you perceive:
 - hidden objects
 - the relative position of objects
 - the object shape



Light and shadows

► Constraints for real-time shadows

- Light sources Dynamic
 - Shadow Casters Dynamic
 - Shadow Receivers Dynamic



Doom 3 (2004)

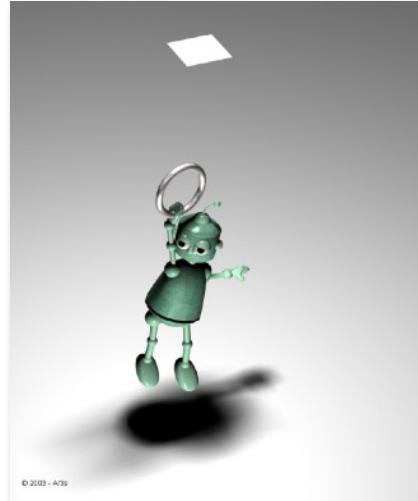
Light and shadows

- ▶ 2 kind of shadows:

- Hard shadows
 - Point light source

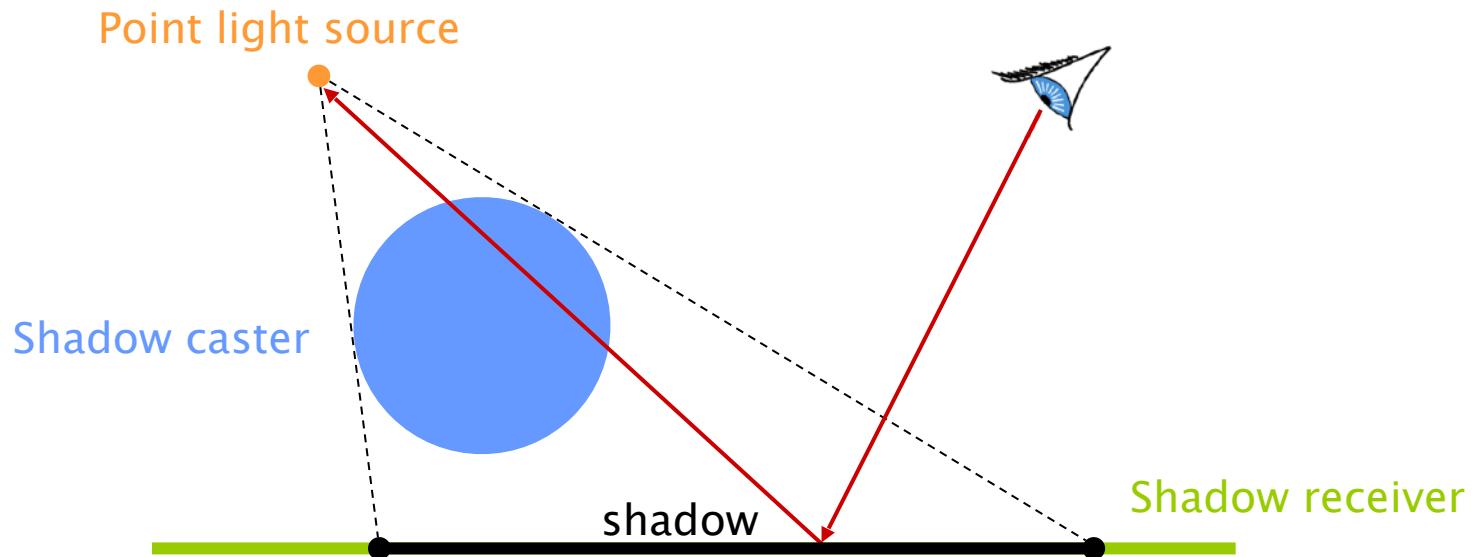


- Soft shadows
 - Extended light source



Hard shadow

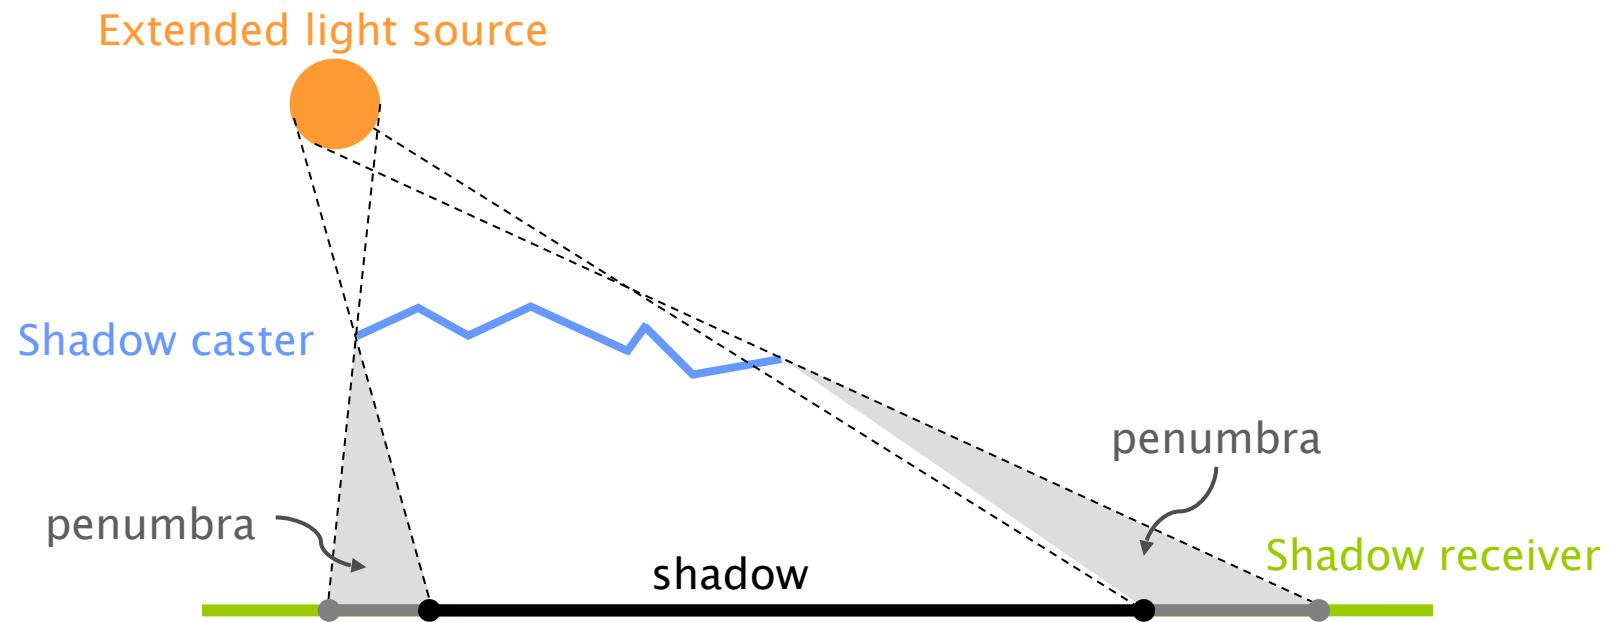
- ▶ Point light source
- ▶ A point is in shadow if it is not visible from the light source

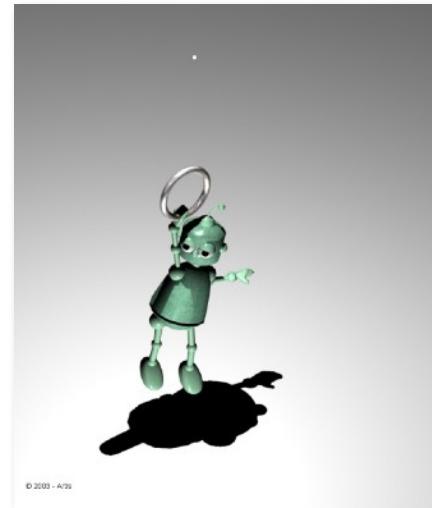


Soft shadow

- ▶ 3 areas:

- Shadow: light source completely hidden
- Penumbra: light source partially hidden
- Lit: light source completely visible

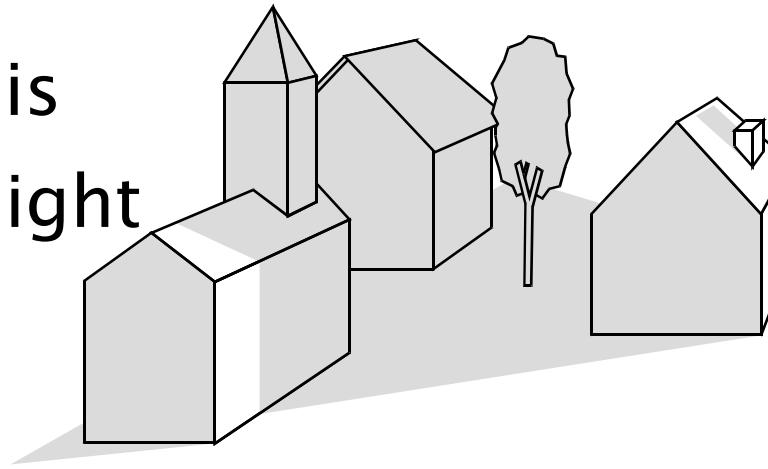




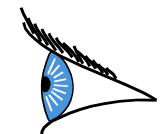
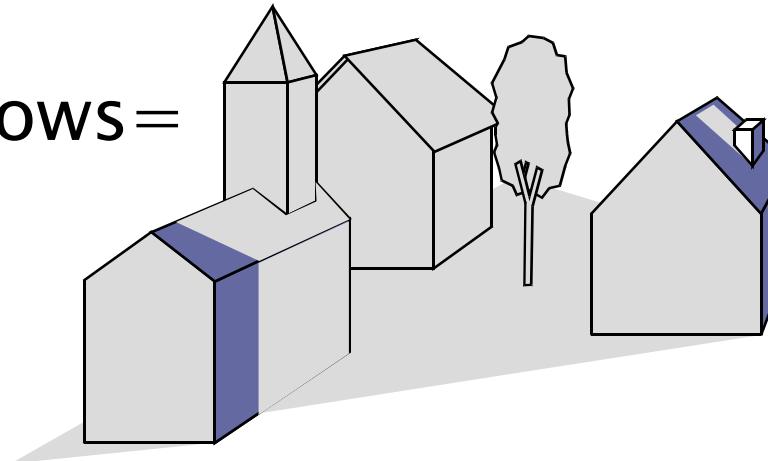
Computing hard shadows

Shadows / visibility

- ▶ A point is lit if it is visible from the light source

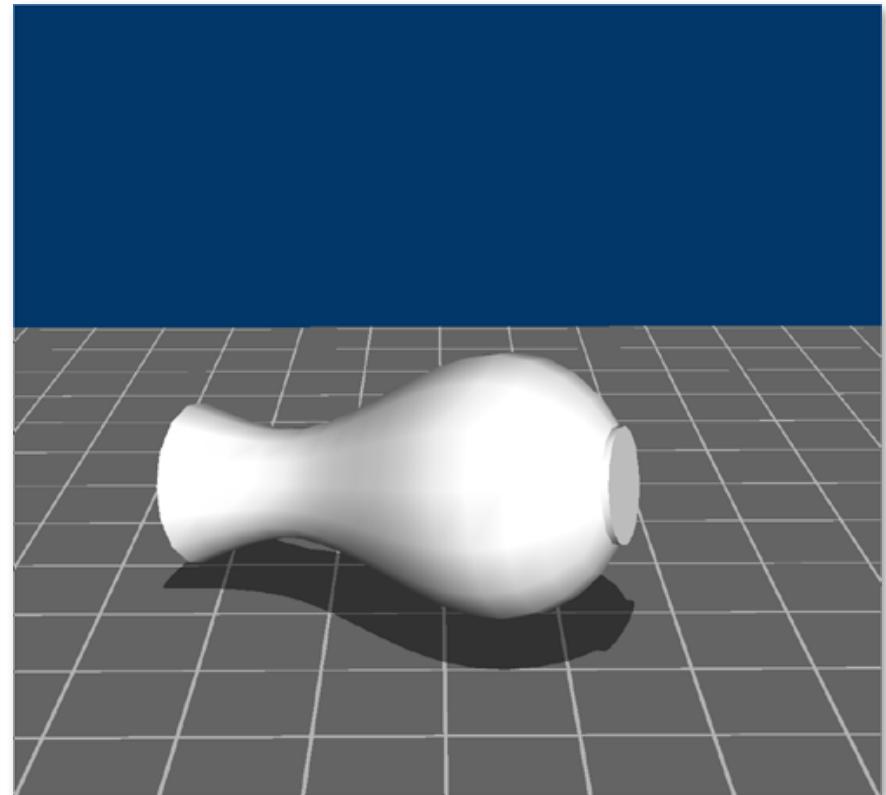
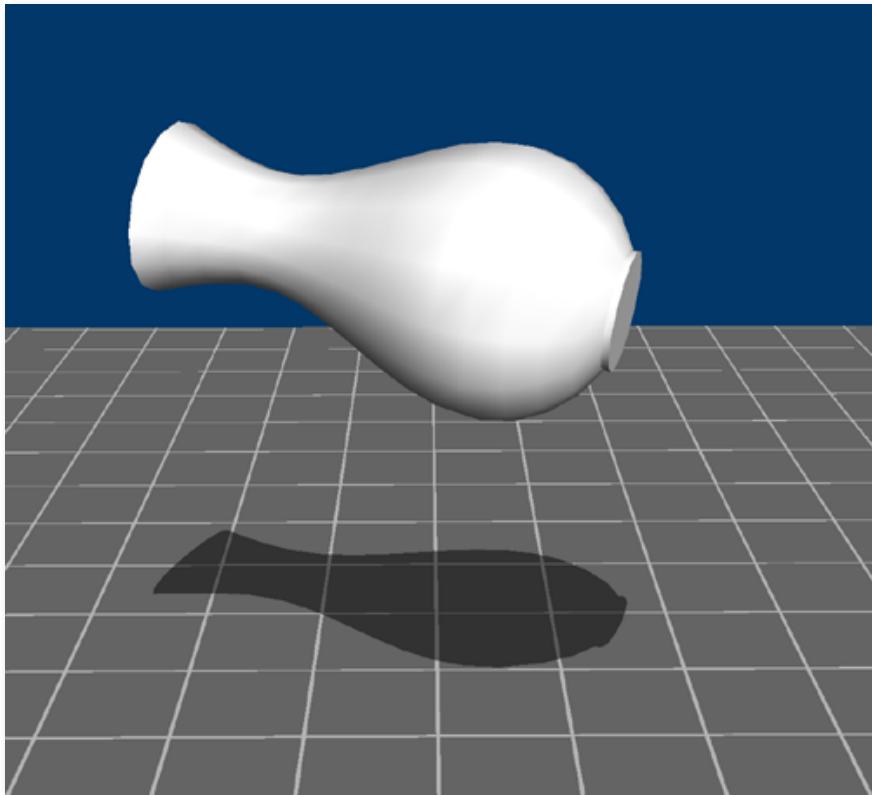


- ▶ Computing shadows = visible surface determination



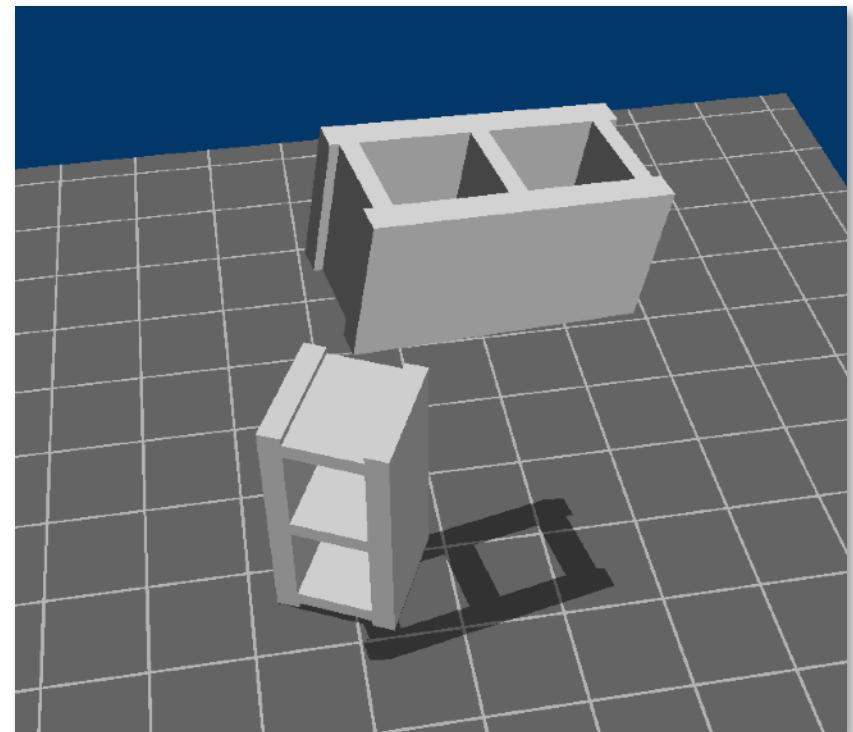
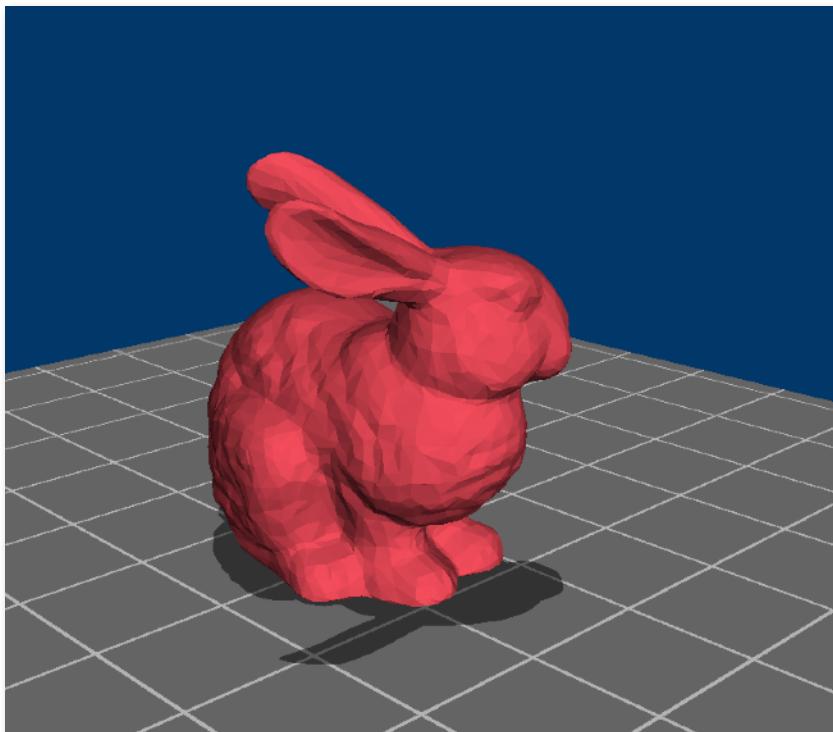
Flat shadows

- ▶ Draw the graphics primitives again, projected on the ground



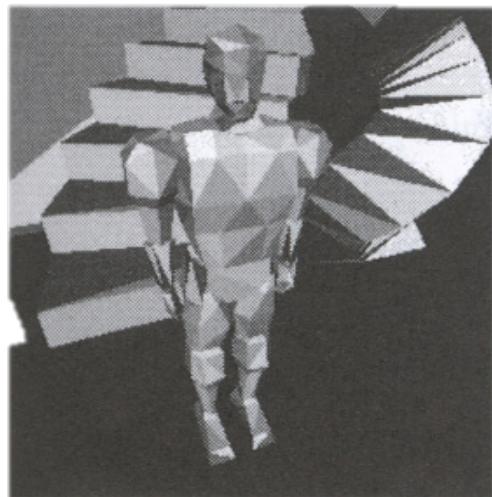
Flat shadows + / -

- + Fast, easy to code
- No self shadows, no shadows on curved surfaces, no shadows on other objects

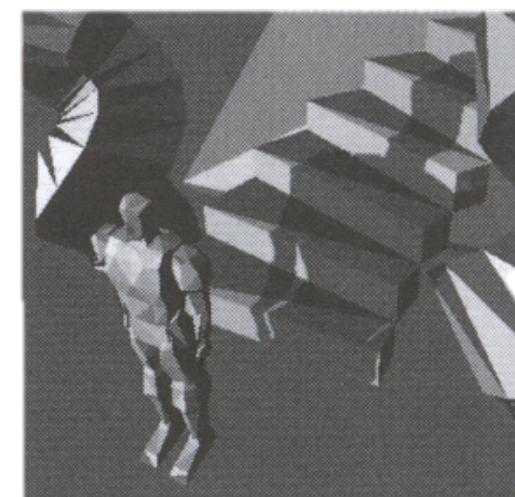


Using textures

- ▶ Separate between occluder and receiver
- ▶ Draw a picture of the occluder, seen from the light source
- ▶ Use it as a texture on the receiver



From the light source



From the viewpoint

Modern shadow algorithms

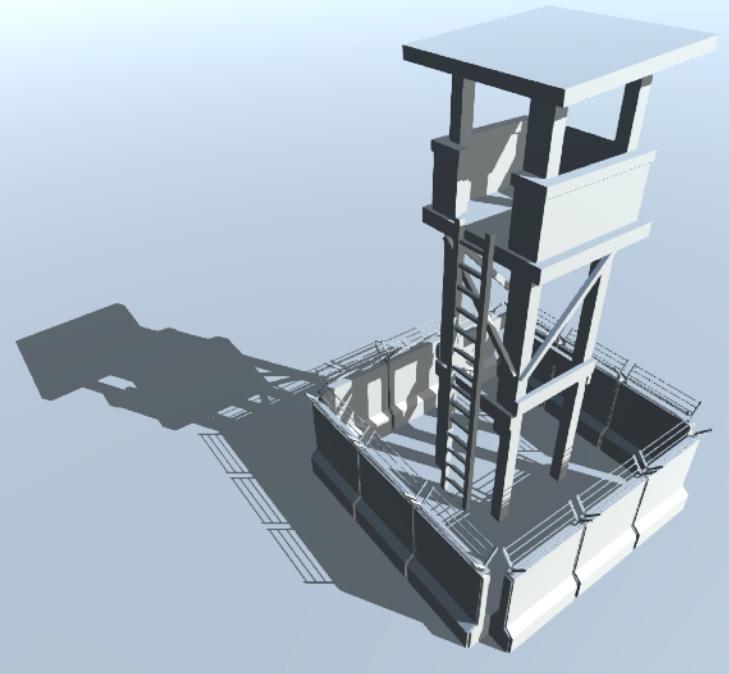


▶ Shadow Maps

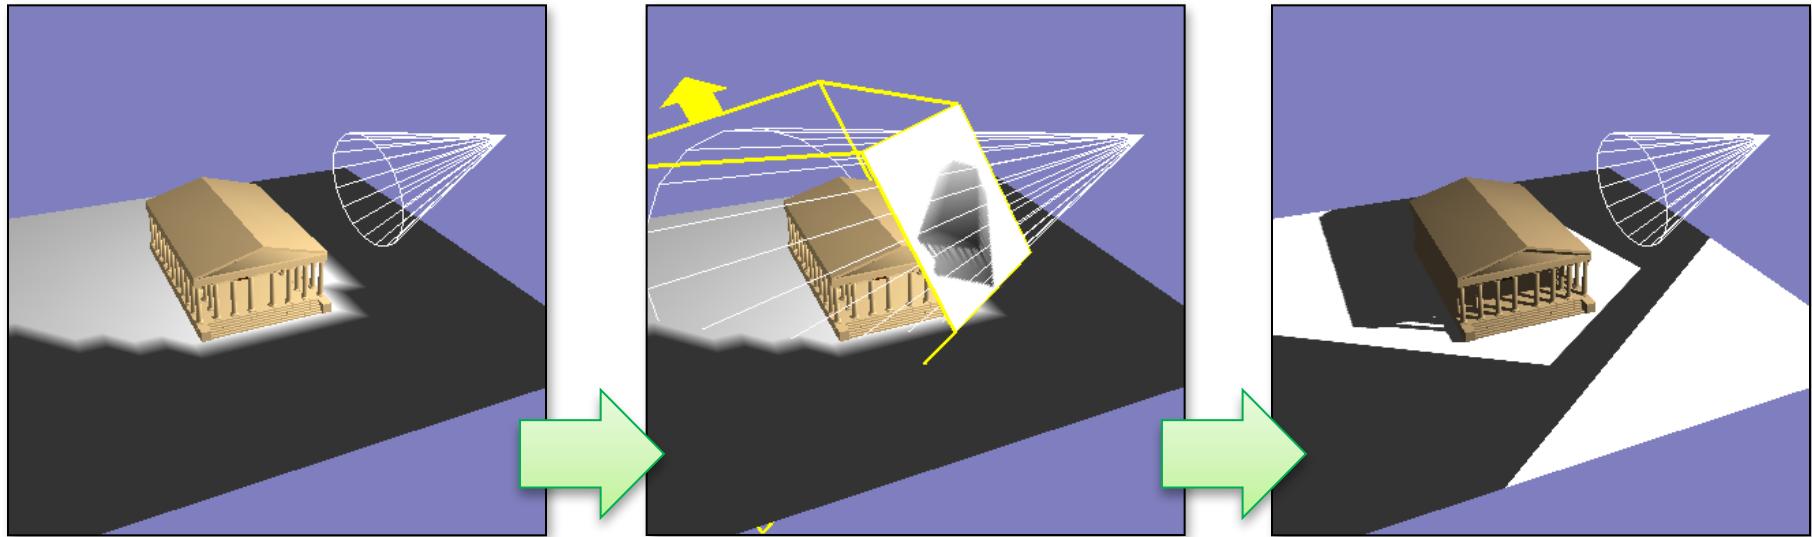
- Image space approach

▶ Shadow Volumes

- Object space approach



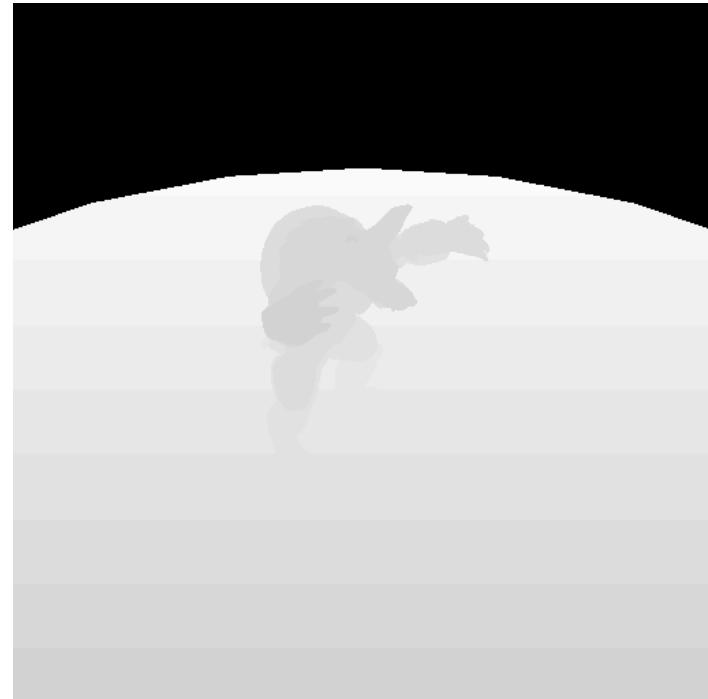
Shadow maps



1. Offscreen rendering **from the light source**
 - Keep z-buffer in a texture
2. Rendering **from the view point**
 - Transform current pixel into light space coord.
 - Compare current depth with depth in texture
 - Change lighting depending on visibility test

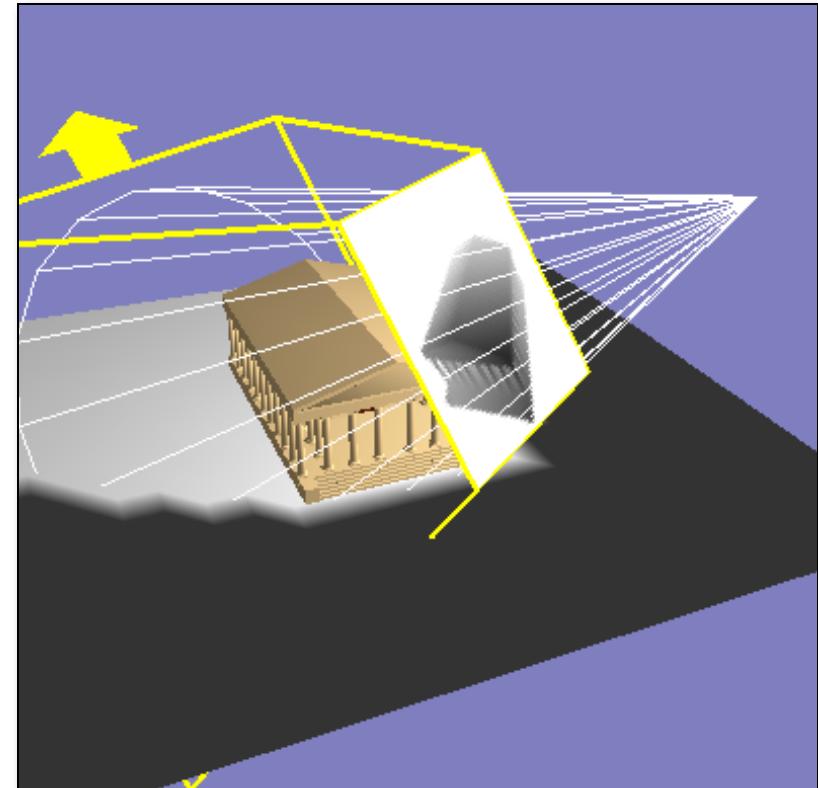
Shadow maps : step 1

- ▶ offscreen rendering from the light source:
 - Transformation + projection matrix
 - Light space coordinates
 - Store depth into an FBO
- ▶ FBO -> texture



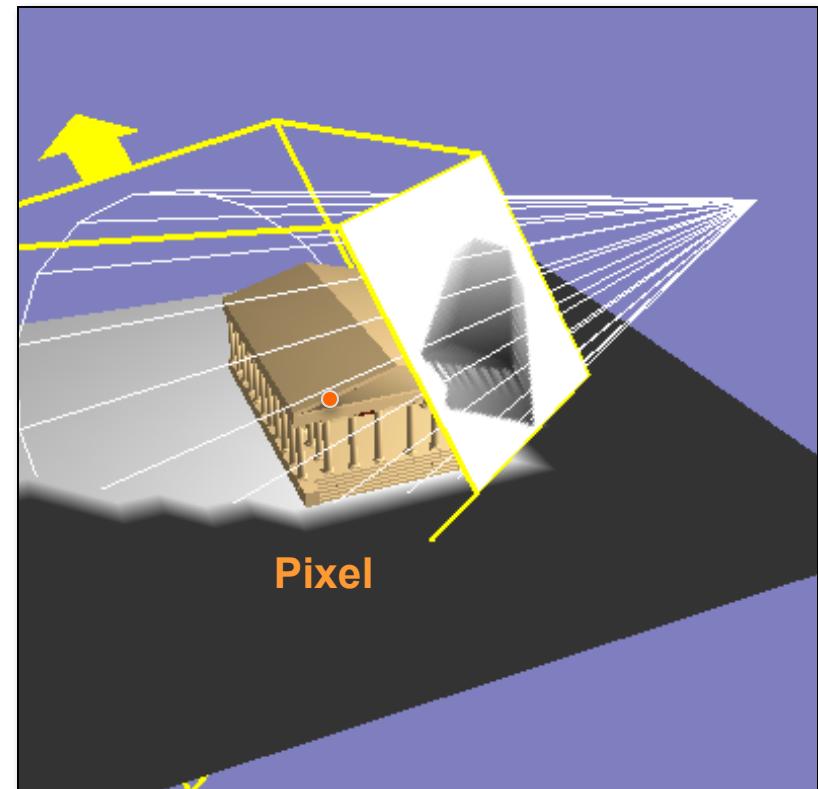
Shadow maps : step 2

- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $=[0,1]^2$



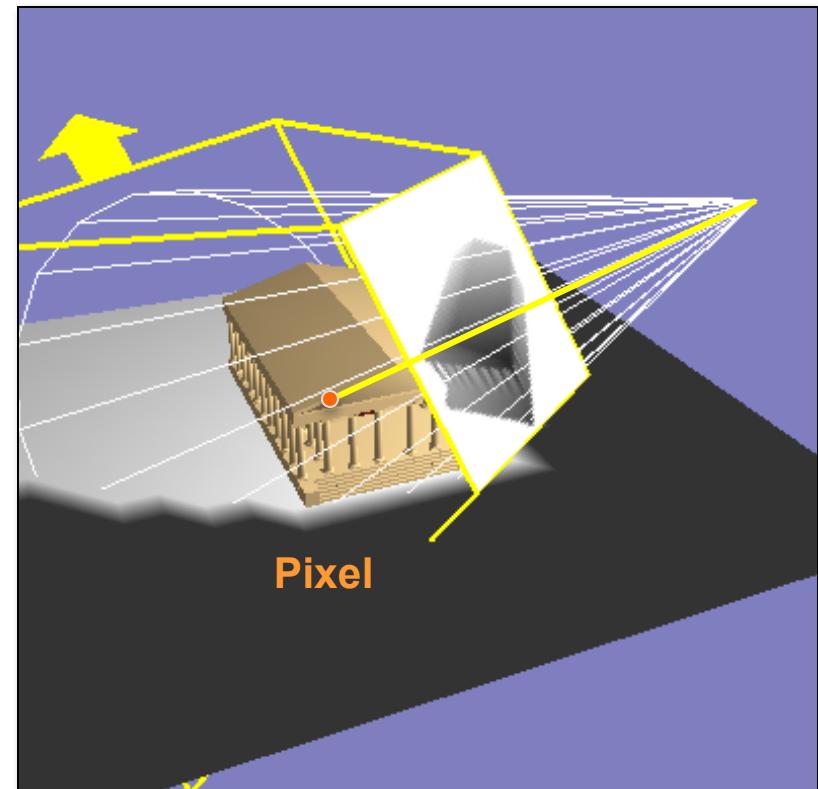
Shadow maps : step 2

- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $= [0, 1]^2$



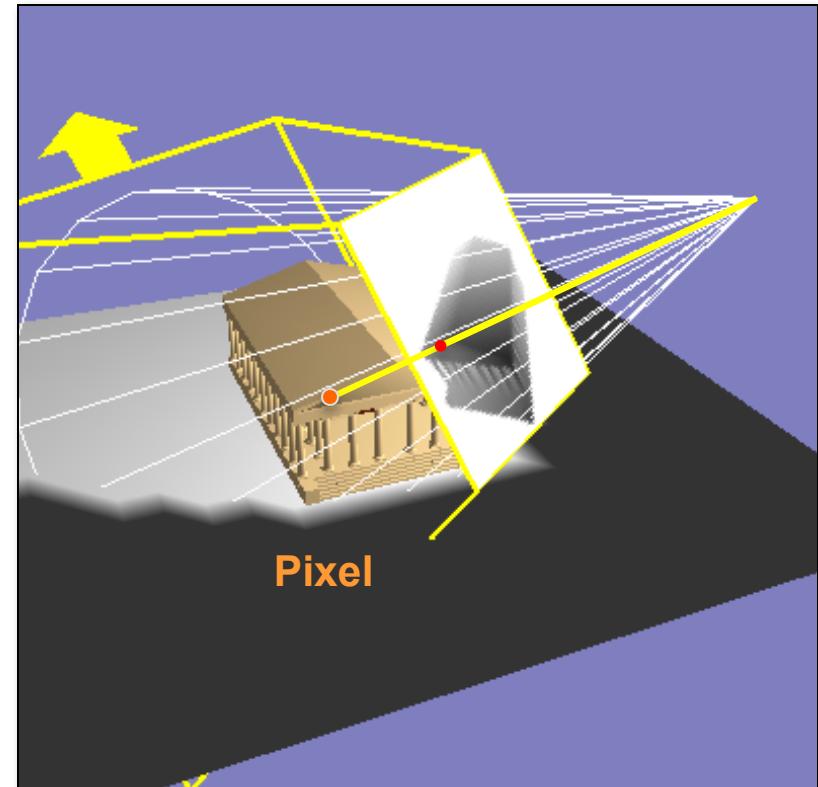
Shadow maps : step 2

- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $= [0, 1]^2$



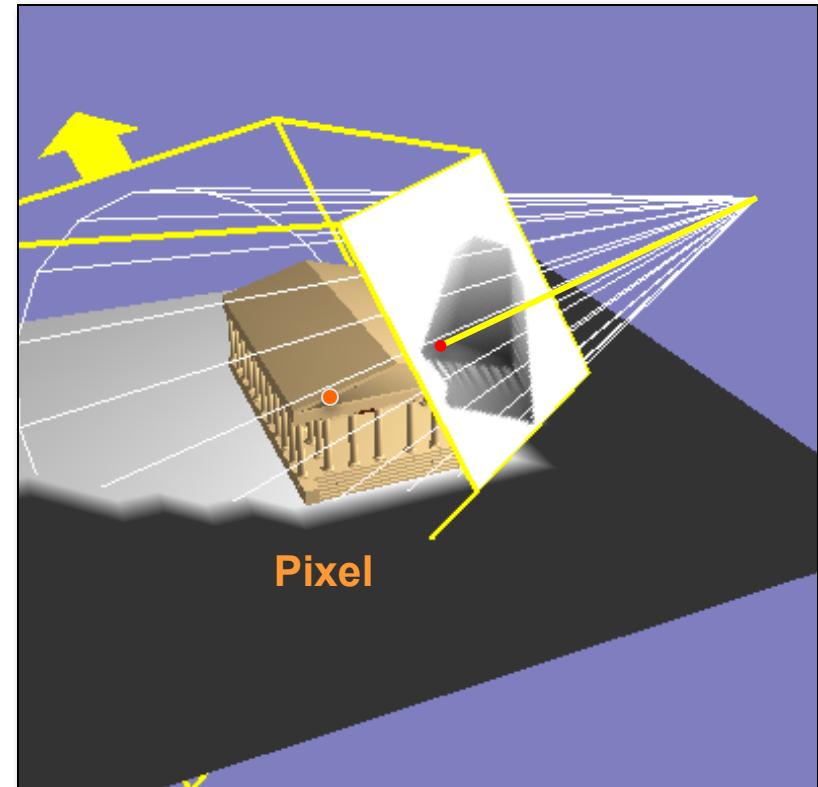
Shadow maps : step 2

- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $= [0, 1]^2$



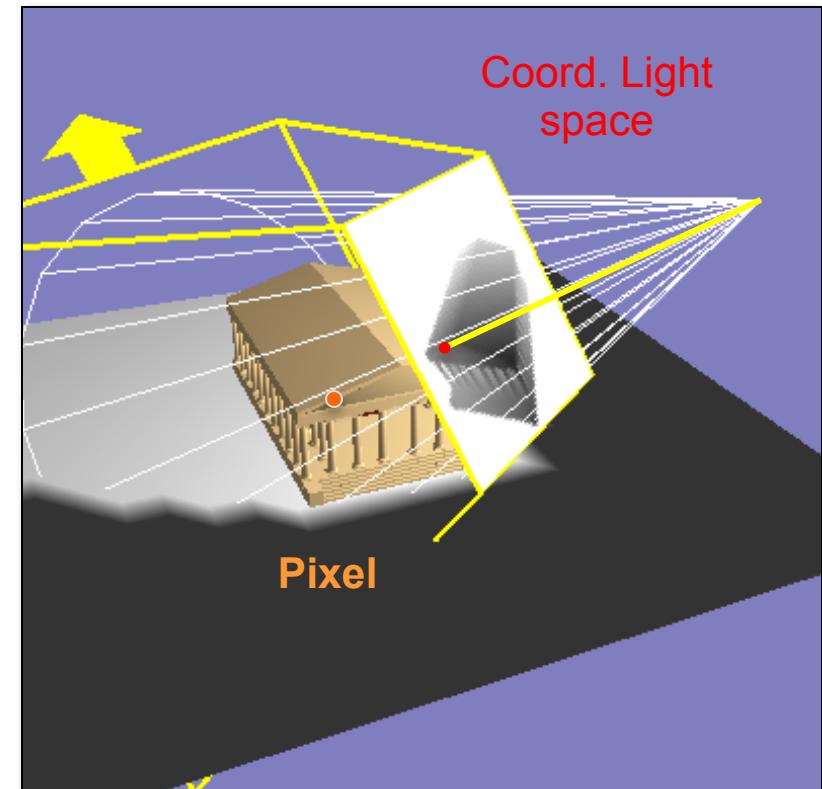
Shadow maps : step 2

- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $= [0, 1]^2$



Shadow maps : step 2

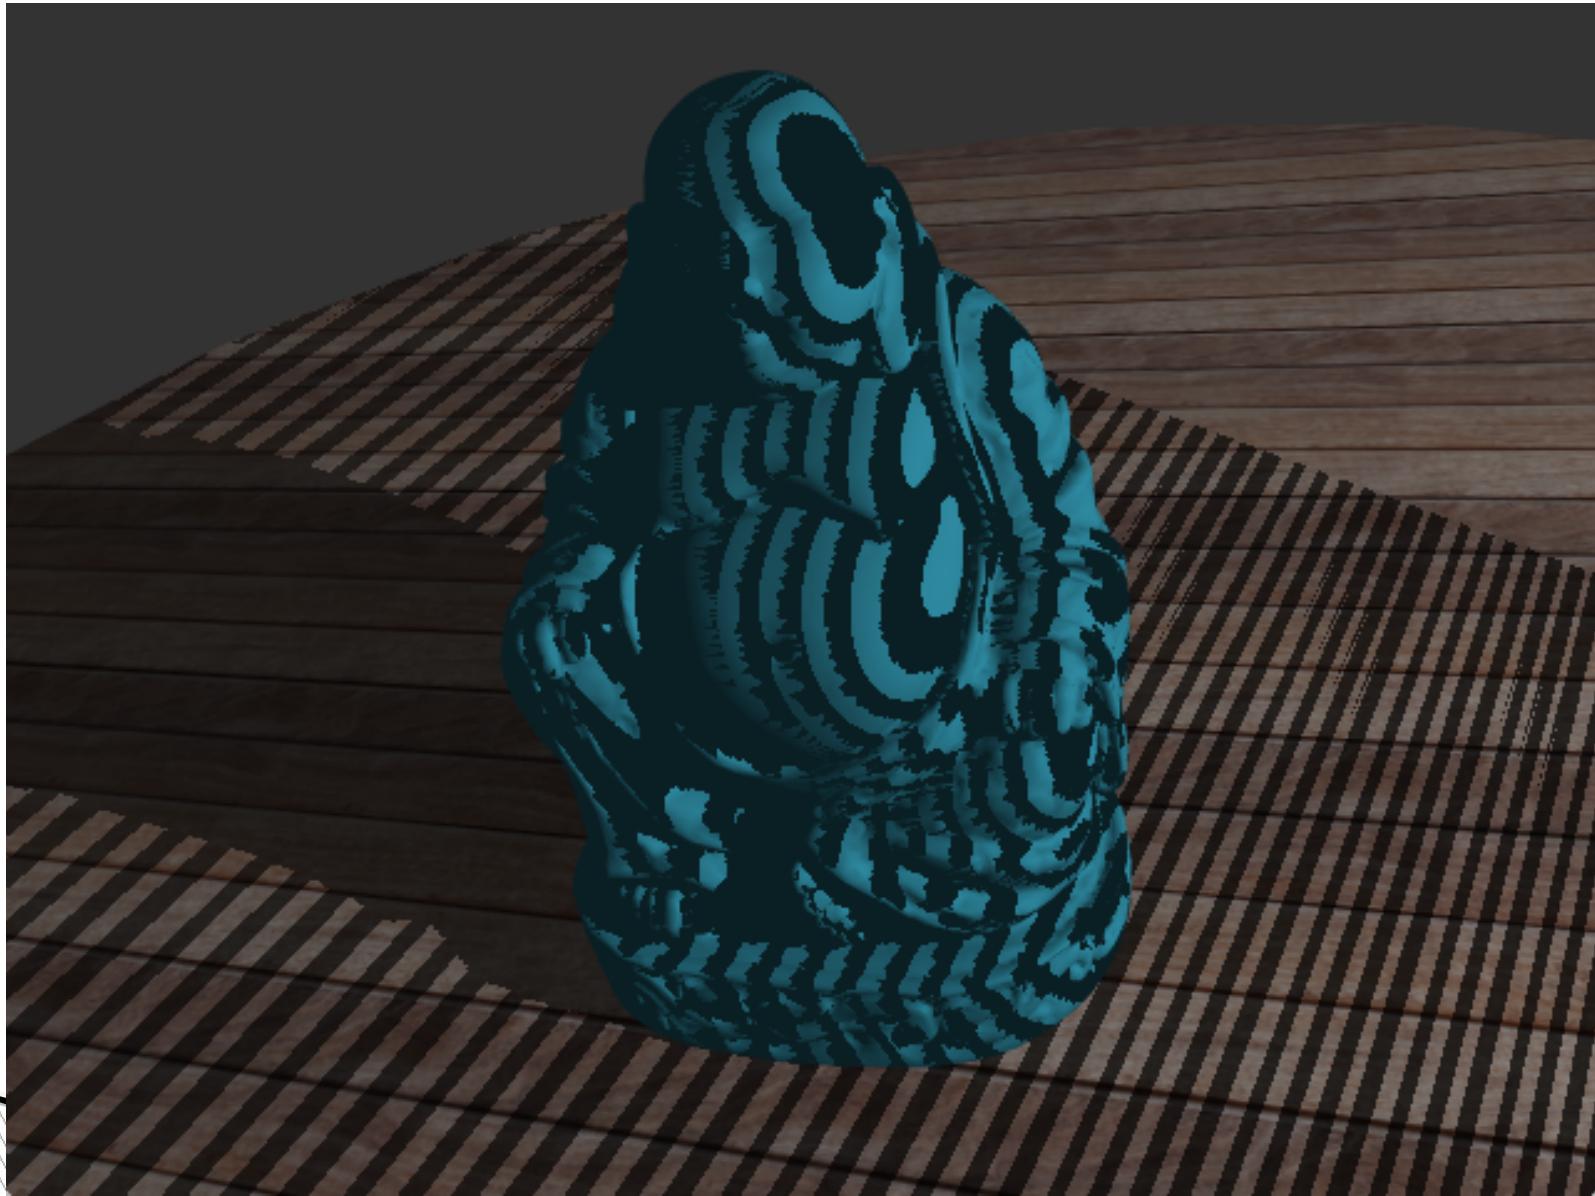
- ▶ Standard rendering
- ▶ Vertex shader:
 - Compute projection in screen space
 - And in light space
- ▶ Fragment shader :
 - Interpolate coordinates
 - Coord. texture shadow map
 - $z = \text{distance light source}$
 - z from shadow map
 - Comparison
 - ⚠ Coord. texture $= [0, 1]^2$



Shadow maps: comparison

- ▶ $z_{\text{shadowMap}} < z_{\text{computed}}$
 - In shadow
 - Ambient lighting only
- ▶ $z_{\text{shadowMap}} == z_{\text{computed}}$
 - Lit
 - Ambient + Diffuse + Specular
- ▶ $z_{\text{shadowMap}} > z_{\text{computed}}$
 - Should not happen, in theory

Shadow maps: 1st picture



Shadow maps: 1st picture

- ▶ “it’s not a bug, it’s a feature” *

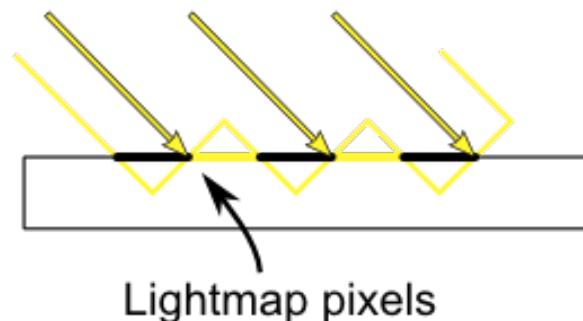


- ▶ What’s happening?

- Comparison $z_{\text{stored}}/\text{interpolated } z$
- z value constant for each pixel
- Self-shadowing

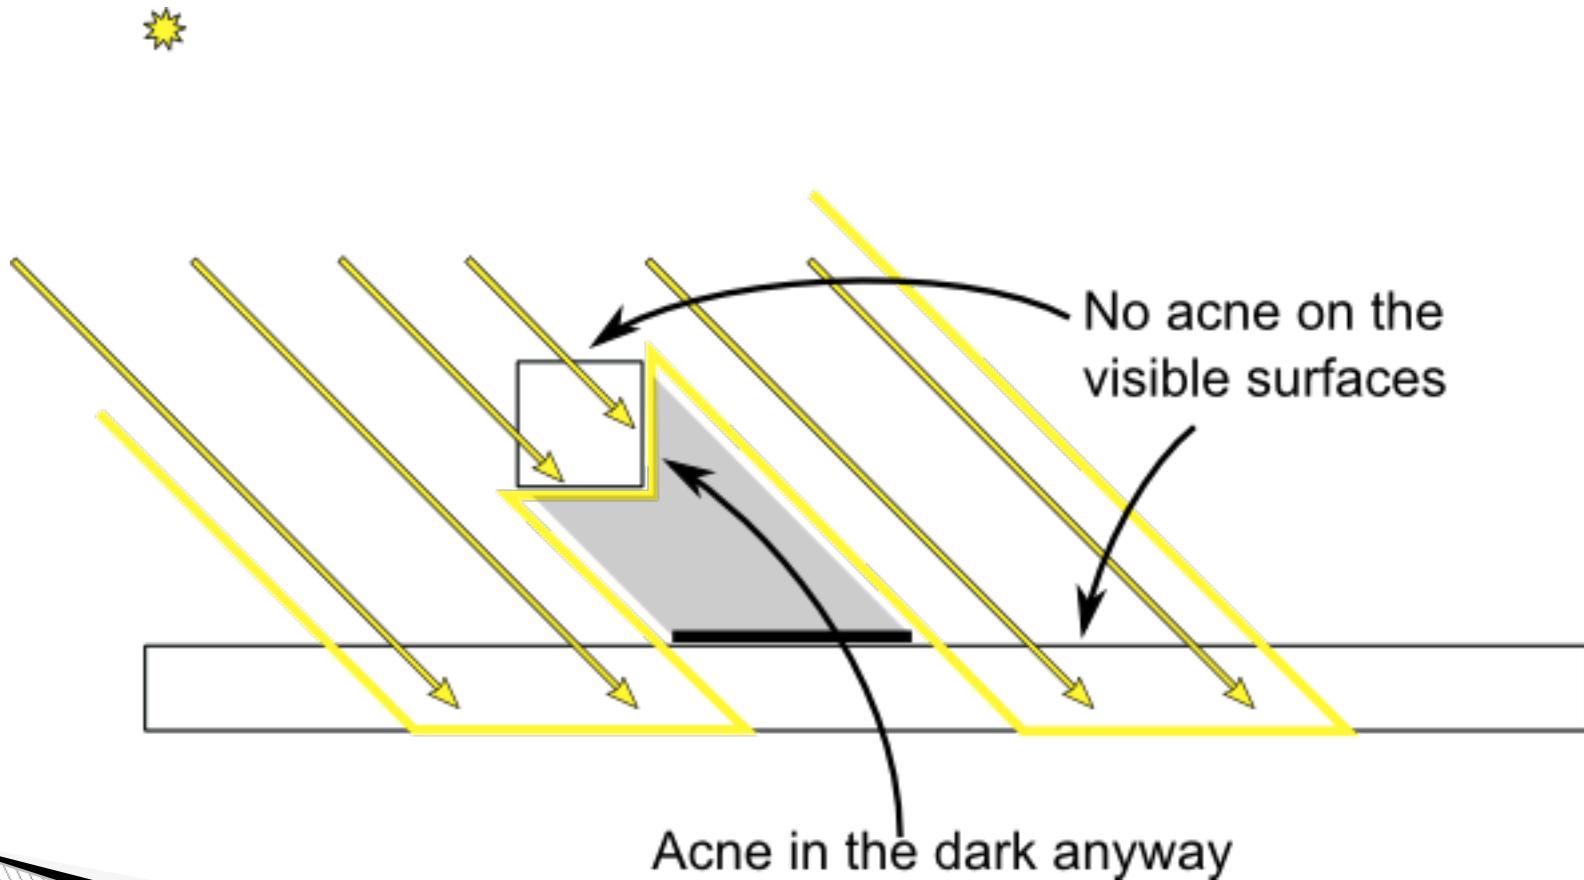
- ▶ Solutions:

- Comparison with $z + \text{epsilon}$ (bias)
- Draw only back-sided surfaces

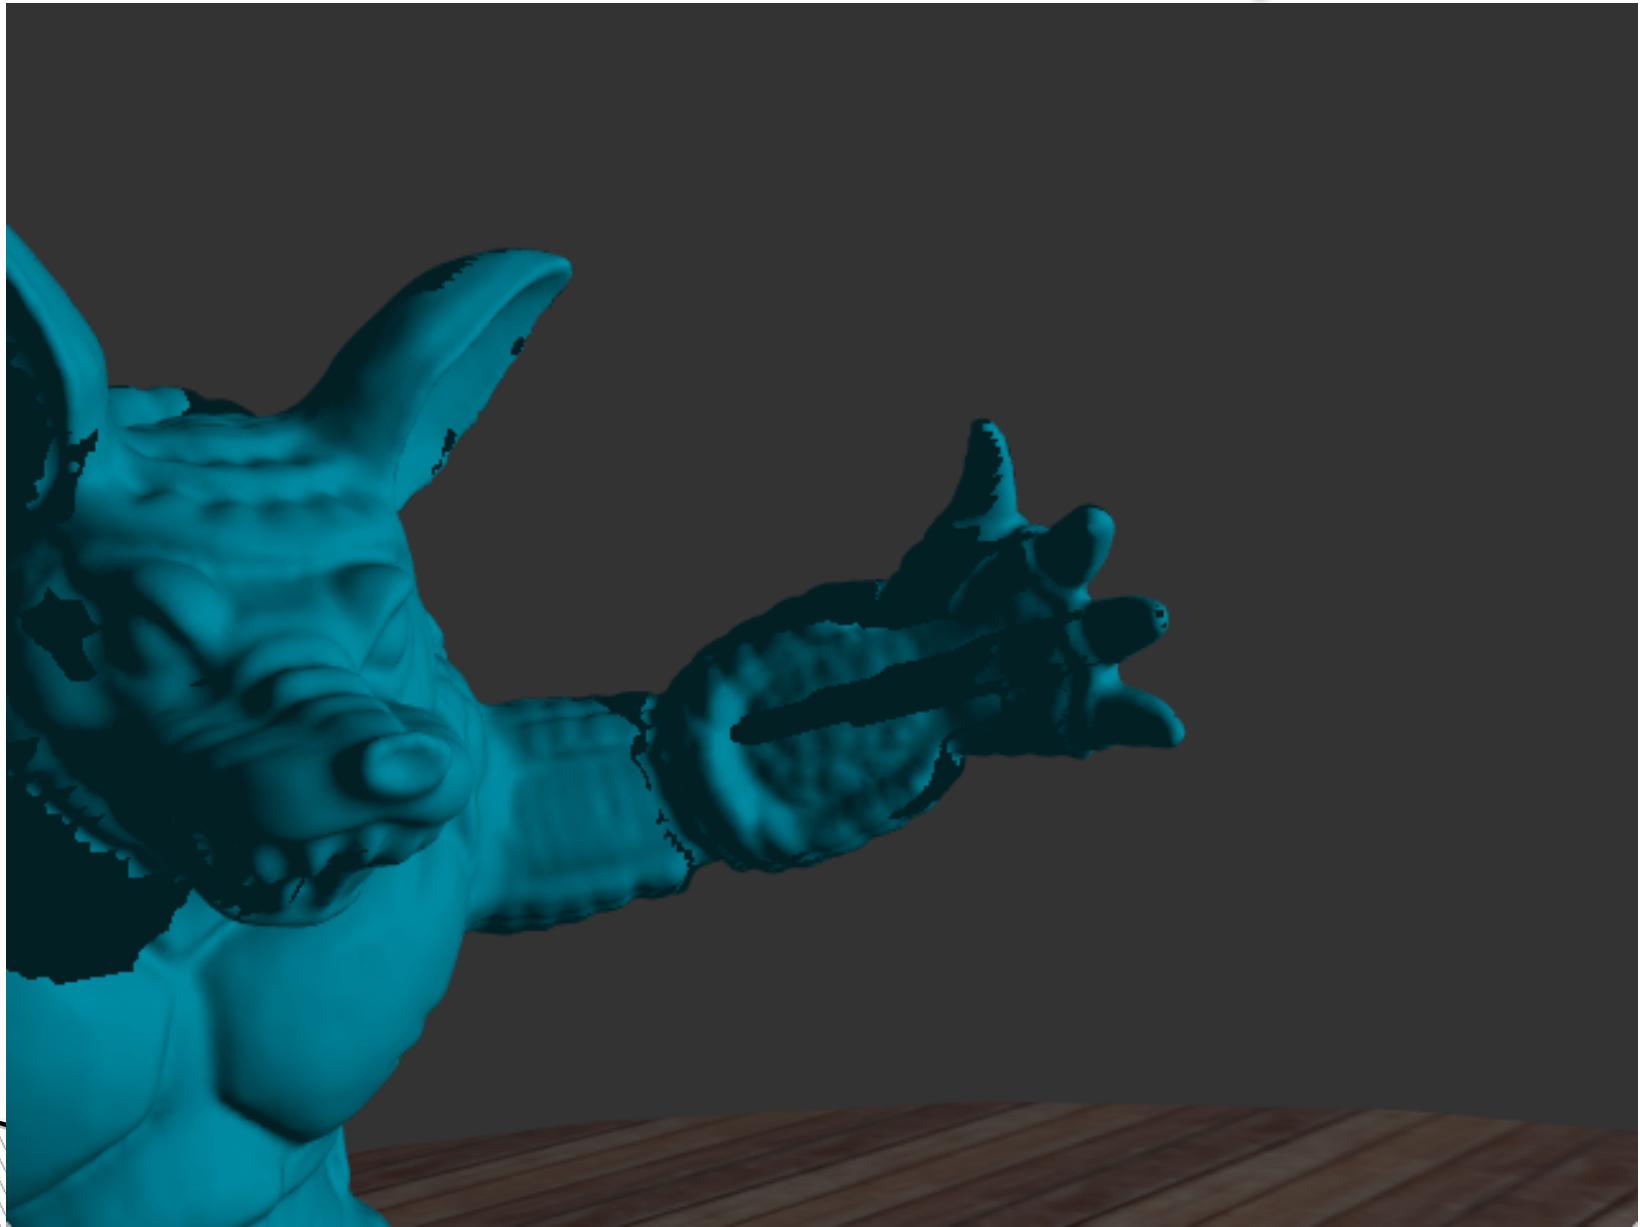


Back-sided surfaces only

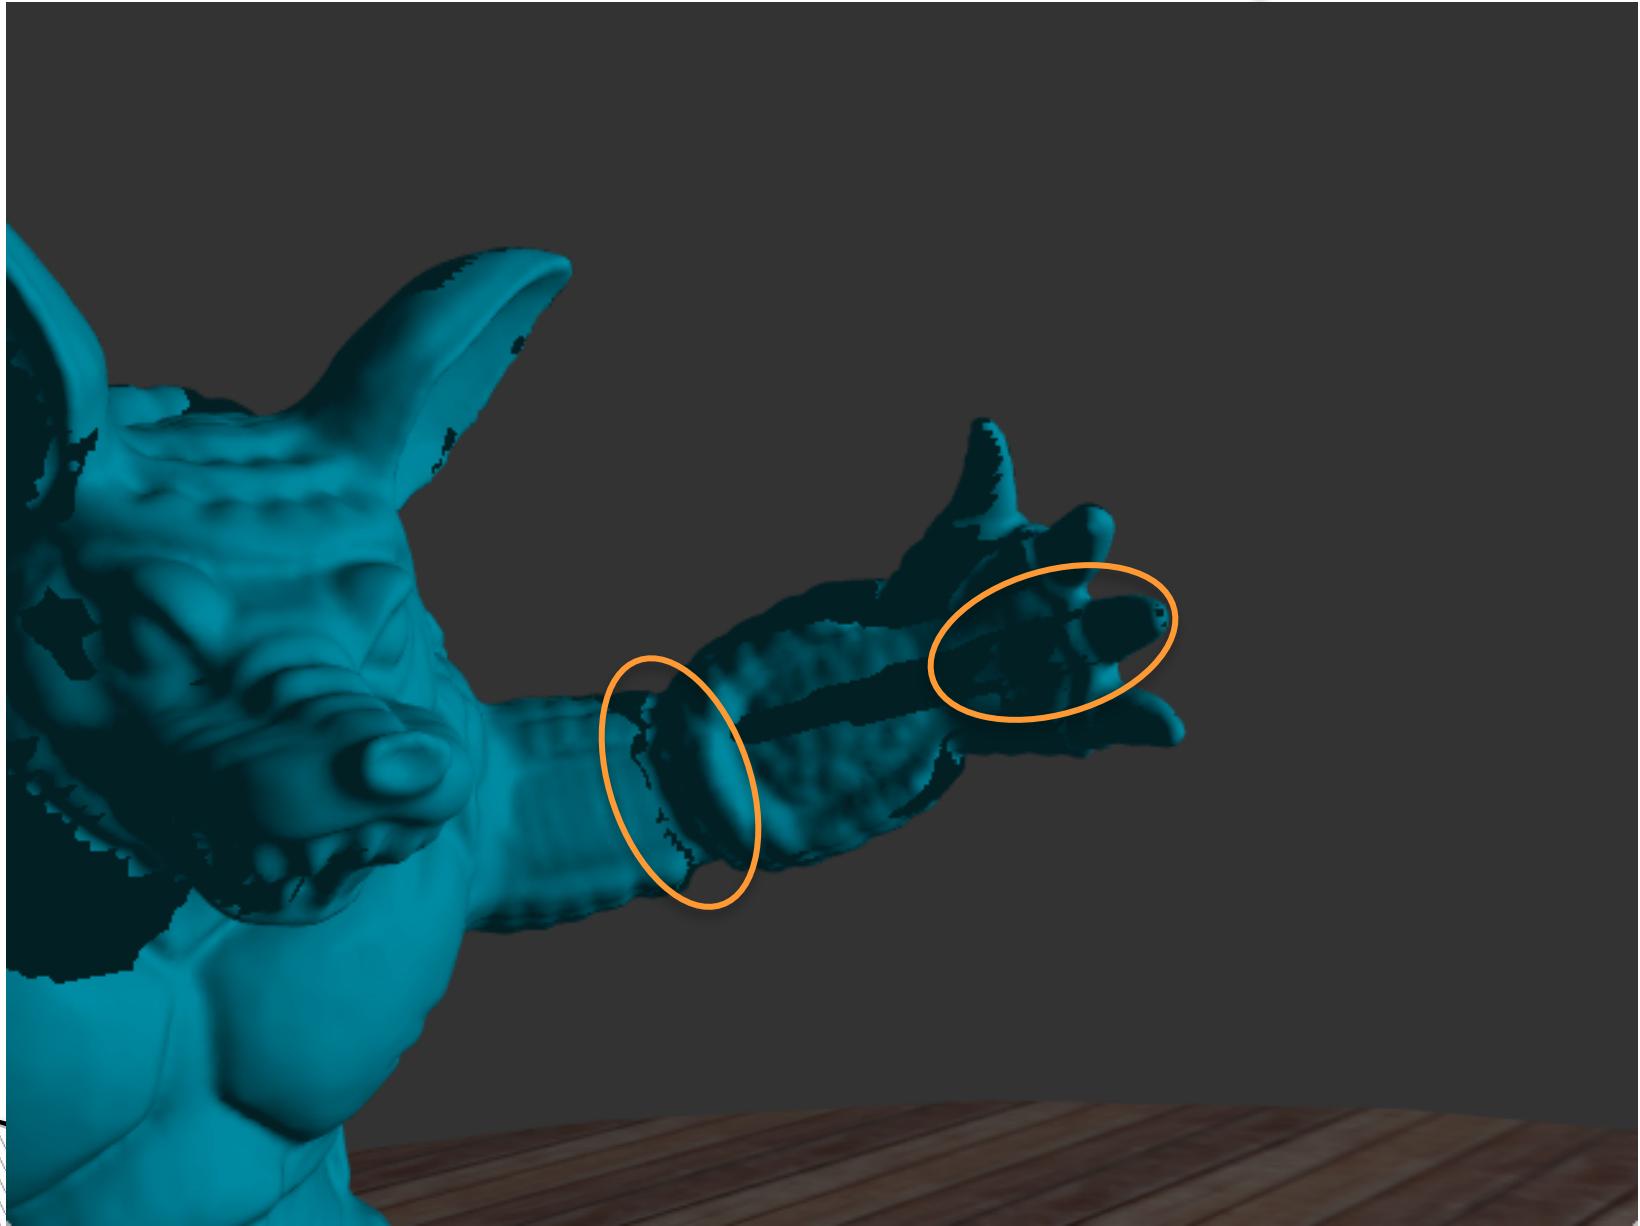
- ▶ Easy: `glCullFace(GL_FRONT);`



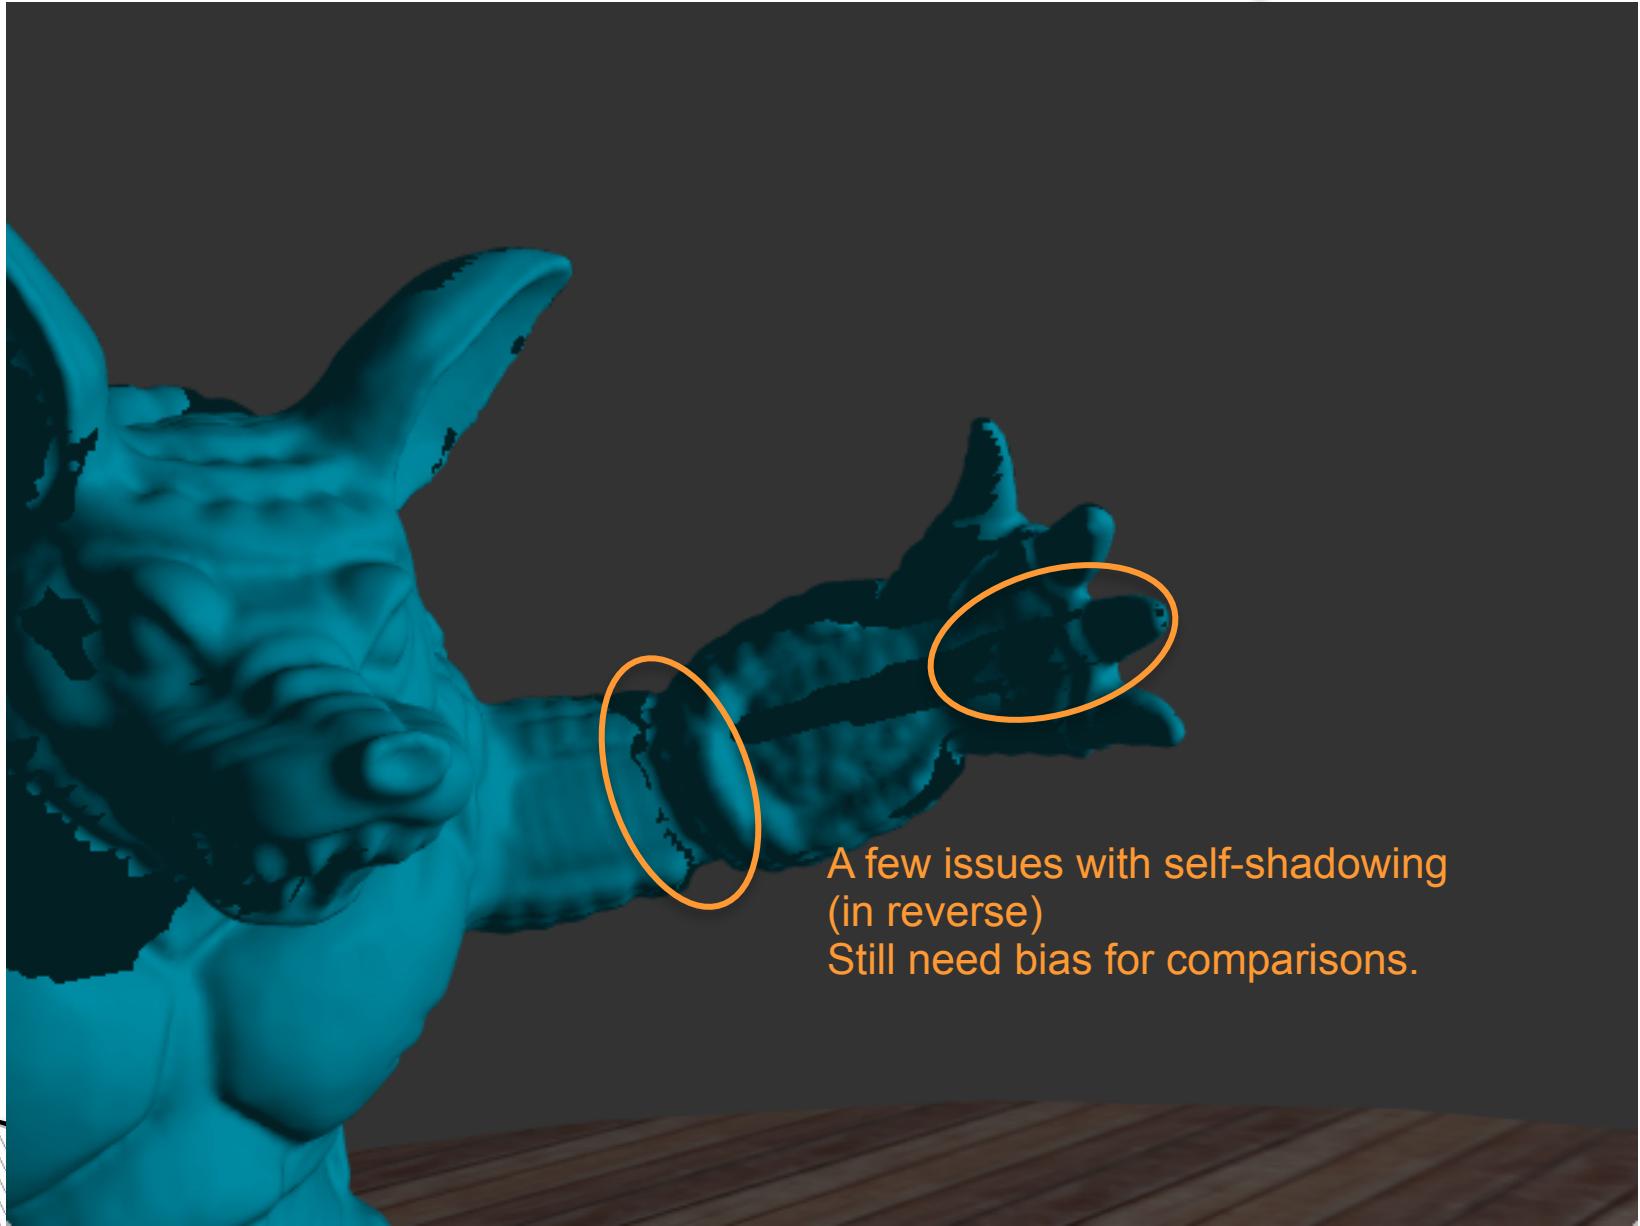
Back-sided surfaces only



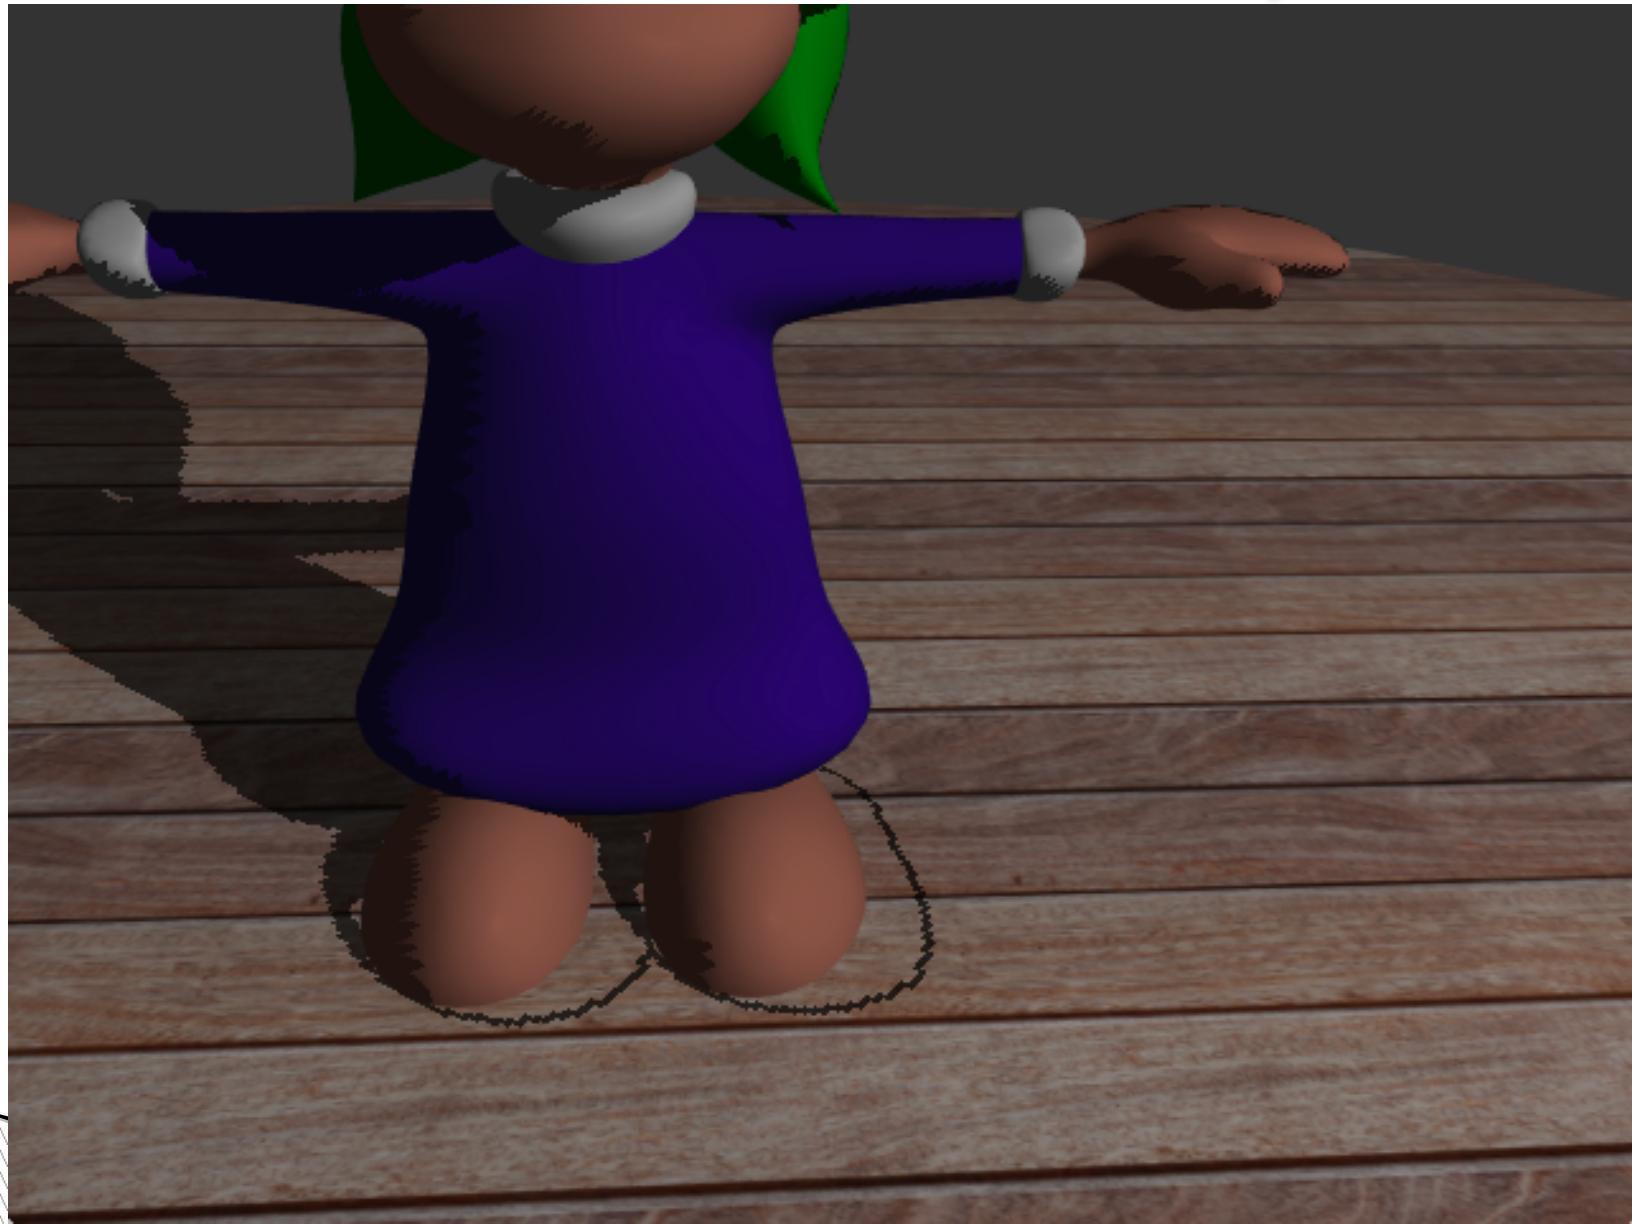
Back-sided surfaces only



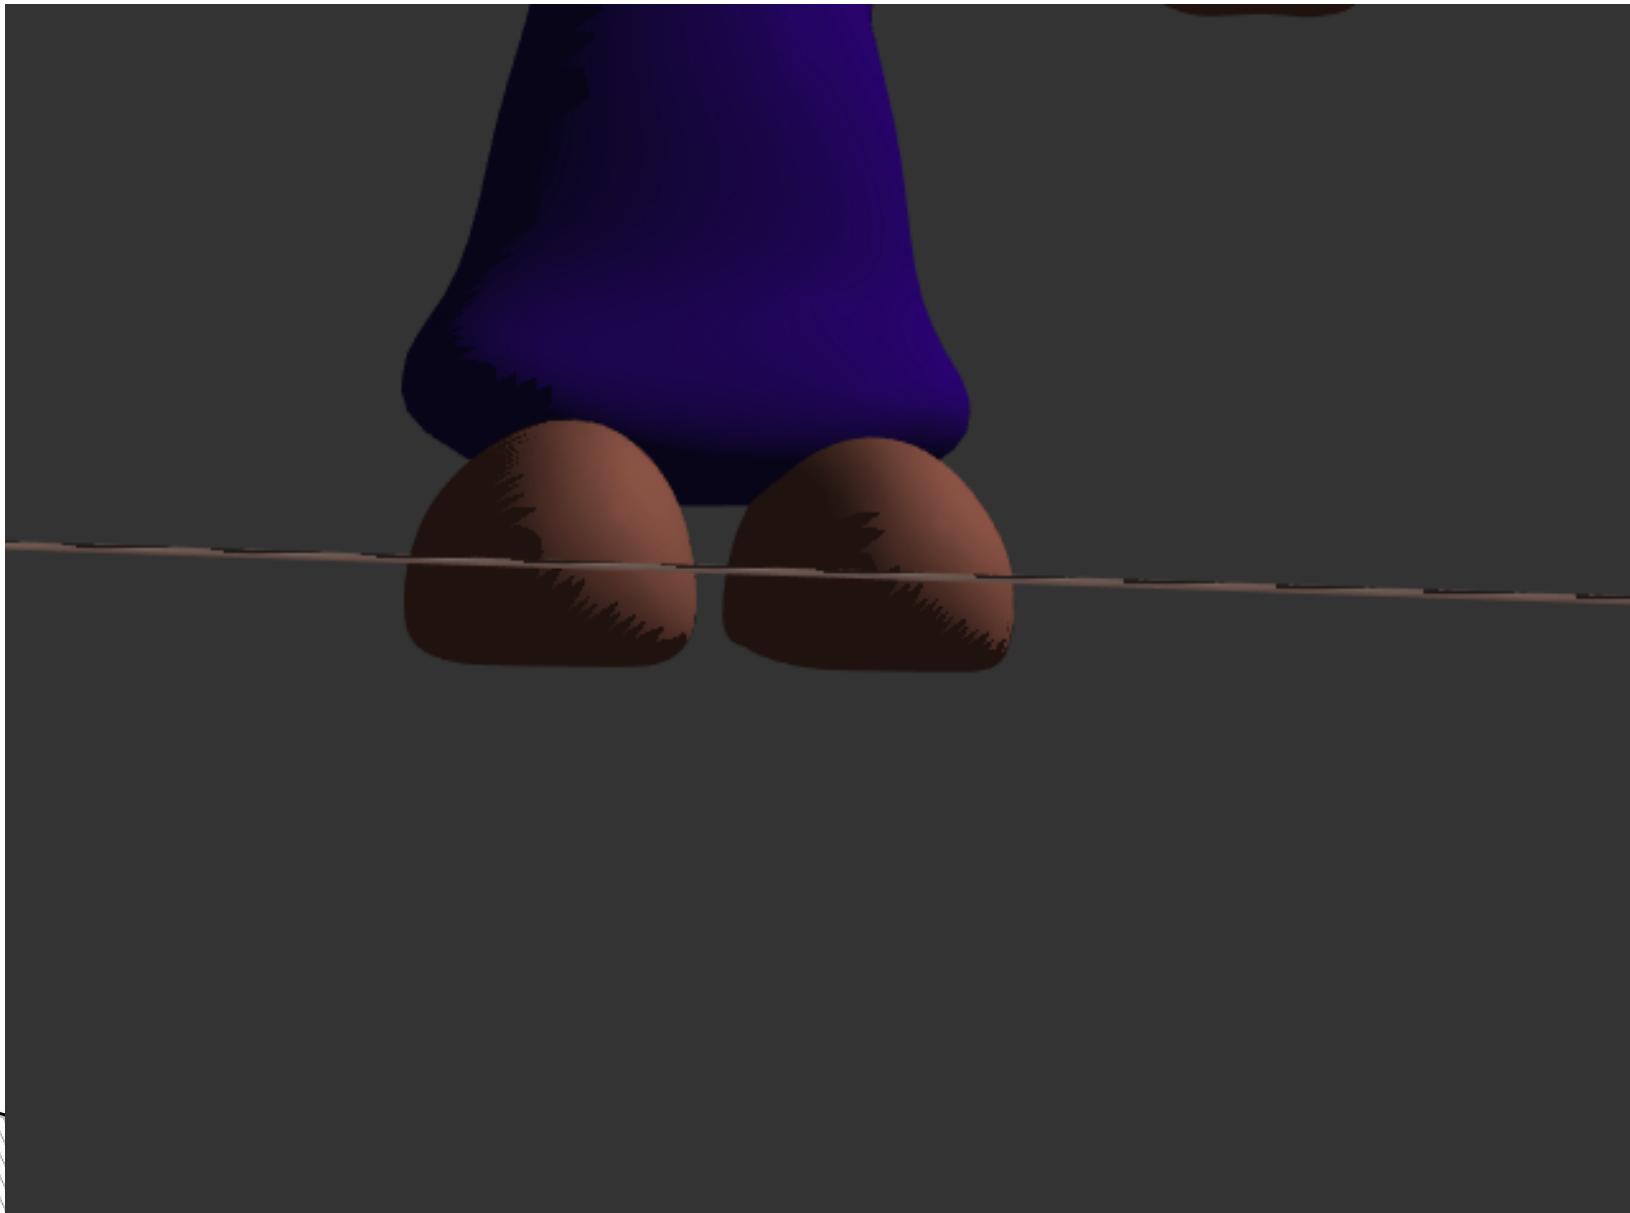
Back-sided surfaces only



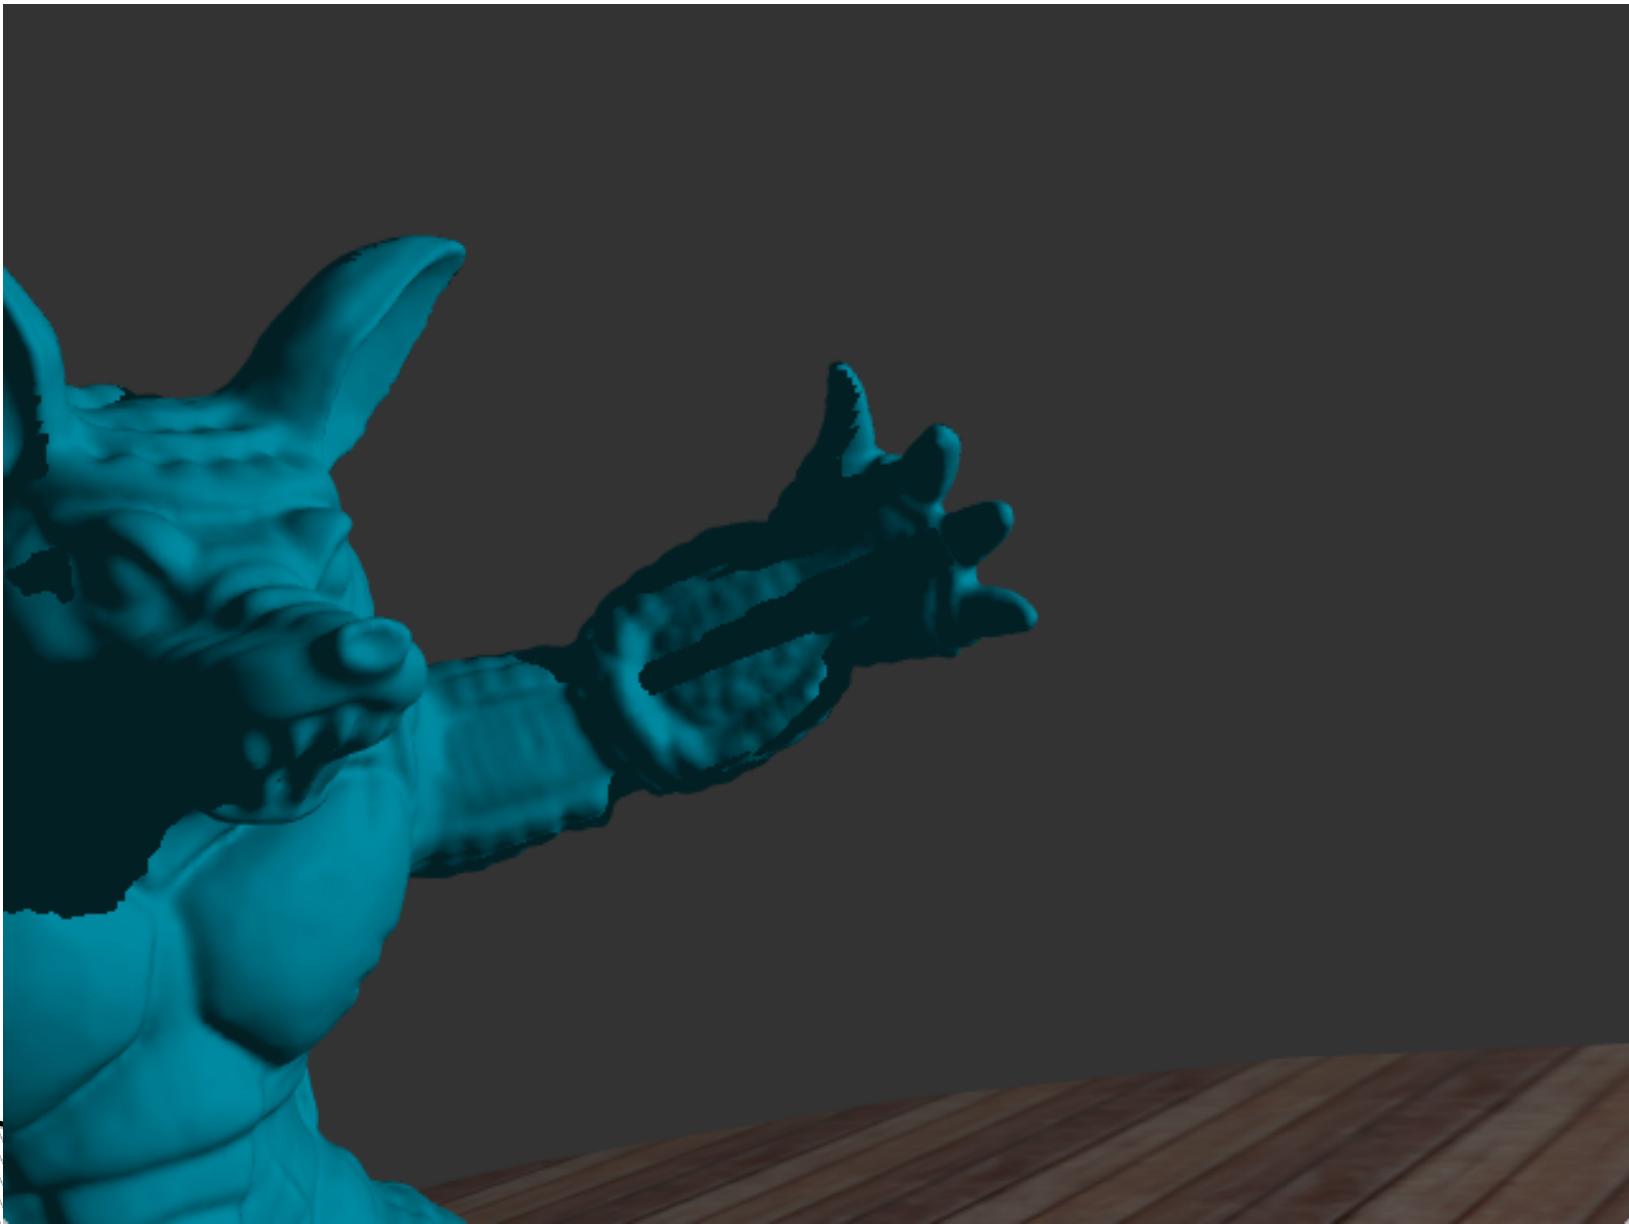
Back-sided surfaces only



Behind the Scenes



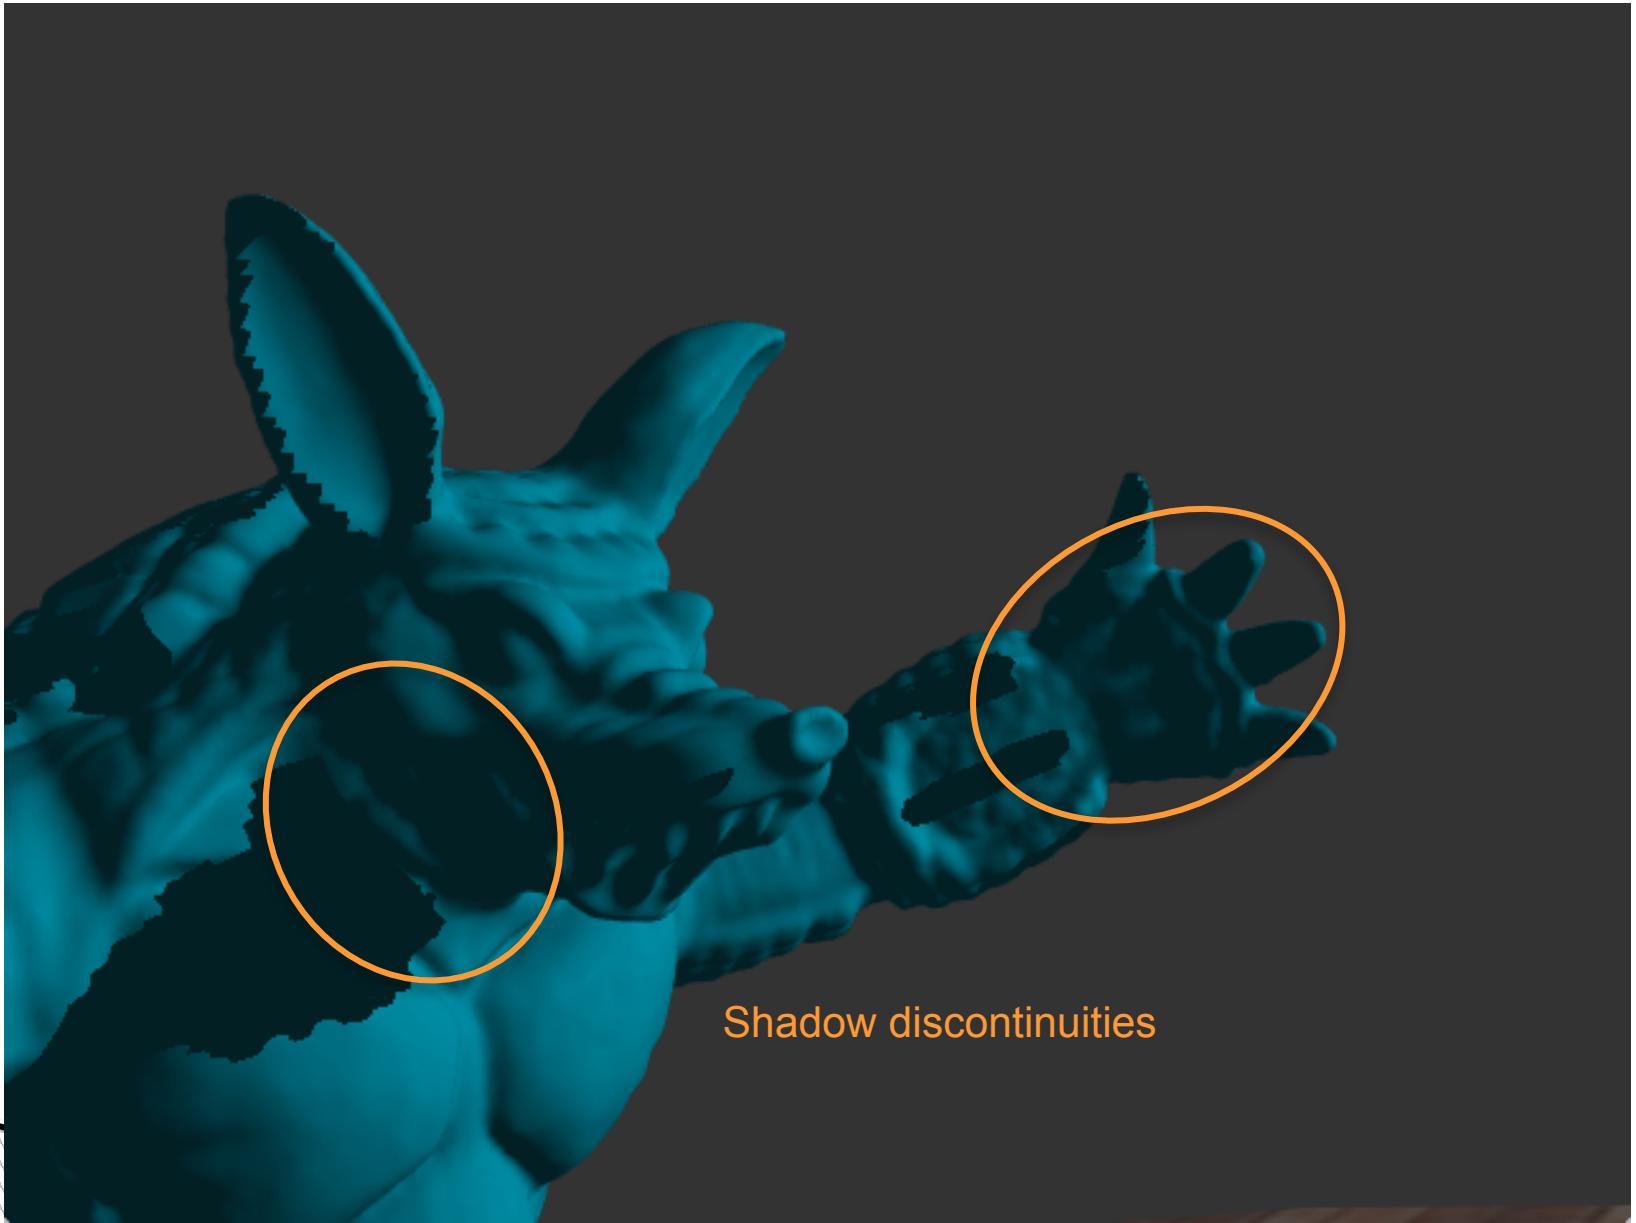
Small constant bias ($5e-3$)



Medium constant bias (1.5e-2)



Medium constant bias ($1.5\text{e-}2$)



Shadow discontinuities

Other bias methods

- ▶ Slope-dependent: $\tan(\text{angle N,L}) * a + b$
 - $b > 0, a = ?$
- ▶ Relative: $z1 * (1 - \text{epsilon}) < z2$

Projection / light source

- ▶ It's a projection:
 - Must divide by w
- ▶ What does it mean if $w < 0$?
 - What should we do?
- ▶ What should we do if we're outside shadow map?
 - How can we check?

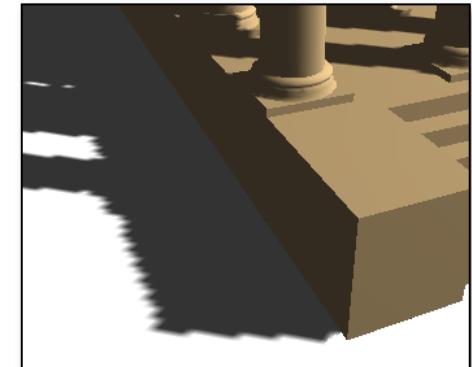
Shadow maps: pro & cons

▶ Pros

- Easy to implement
- Works, regardless of the geometry of the scene
- Cost does not depend on scene complexity

▶ Cons

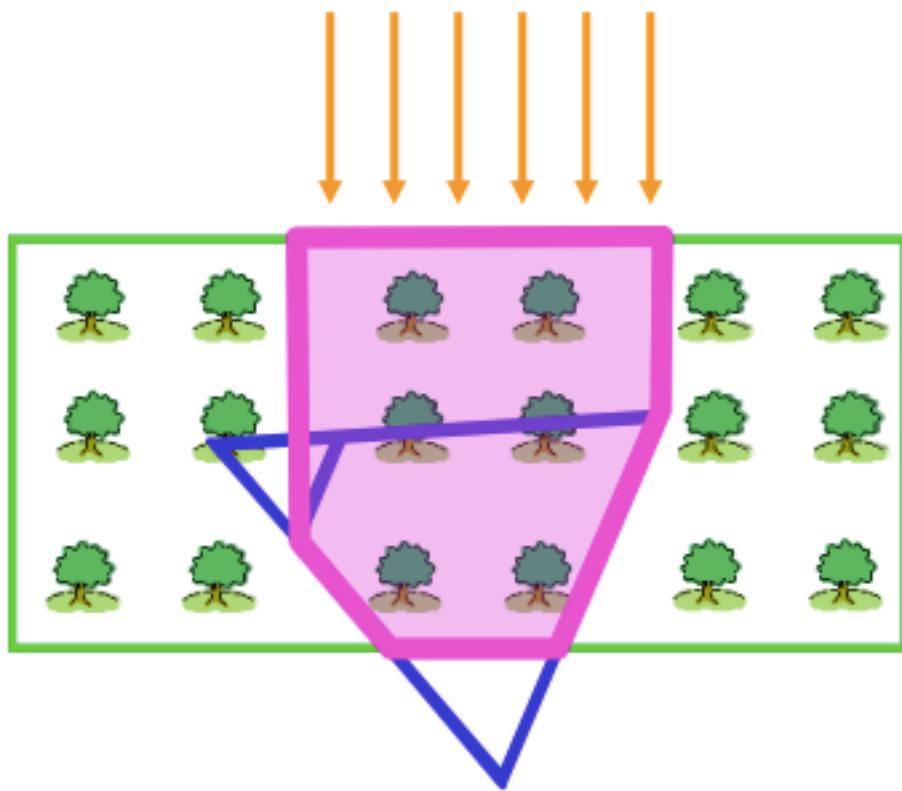
- Several ($>= 2$) scene rendering
- Omni-directional light sources?
- Sampling/aliasing
 - Increasing shadow map resolution is not enough (light source facing viewer)



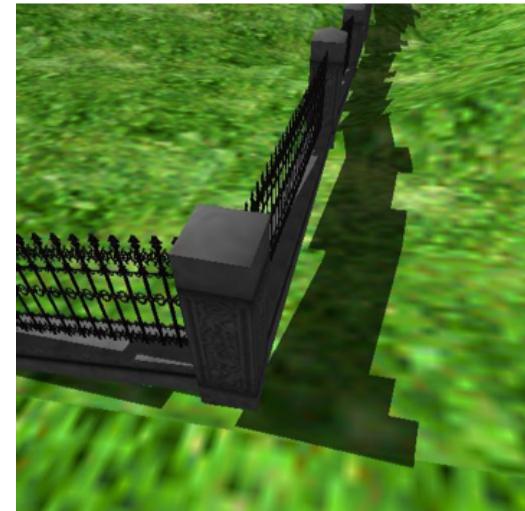
Aliasing issues: solutions

- ▶ Increase shadow map resolution
- ▶ Focus shadow map on visible parts of the scene
- ▶ Adapt sampling (warping)
 - Depending on light-source distance
- ▶ Multi-resolution Shadow maps
 - Cascading shadow map

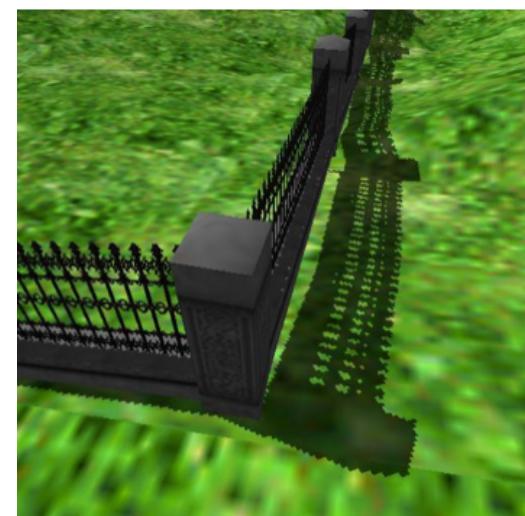
Focus the shadow map



Increases the practical resolution

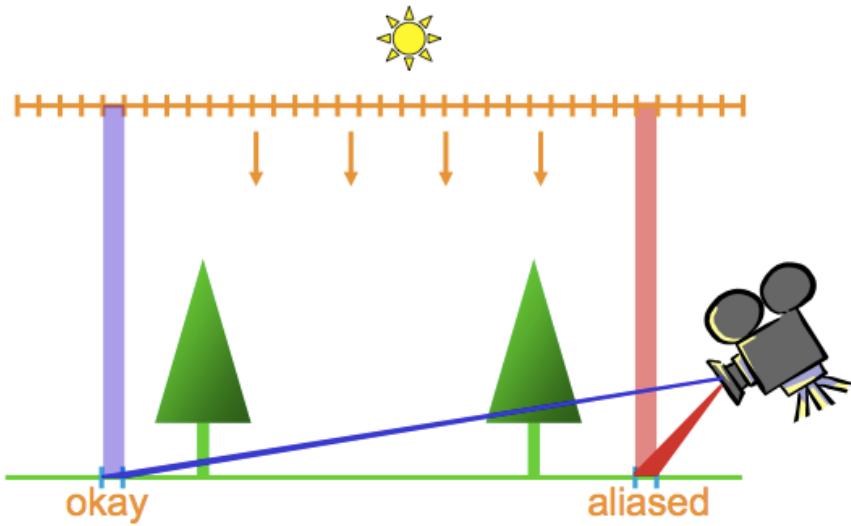


Unfocused

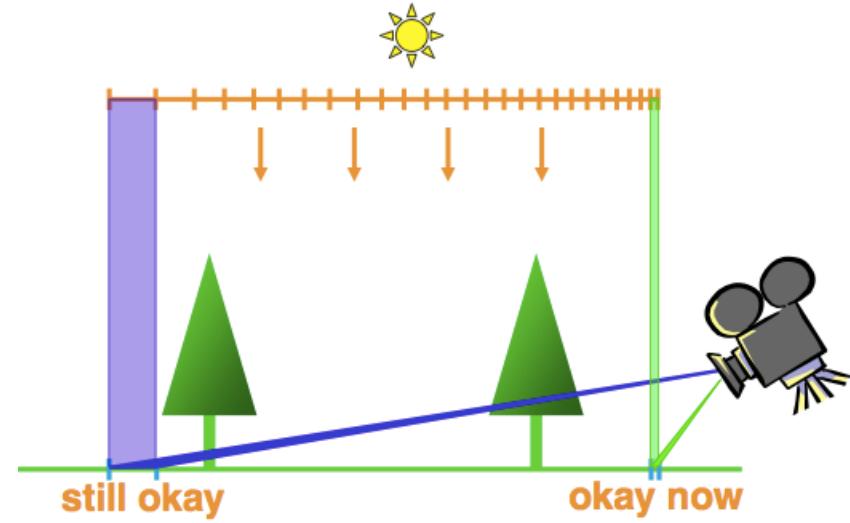


Focused

Warping for shadow maps



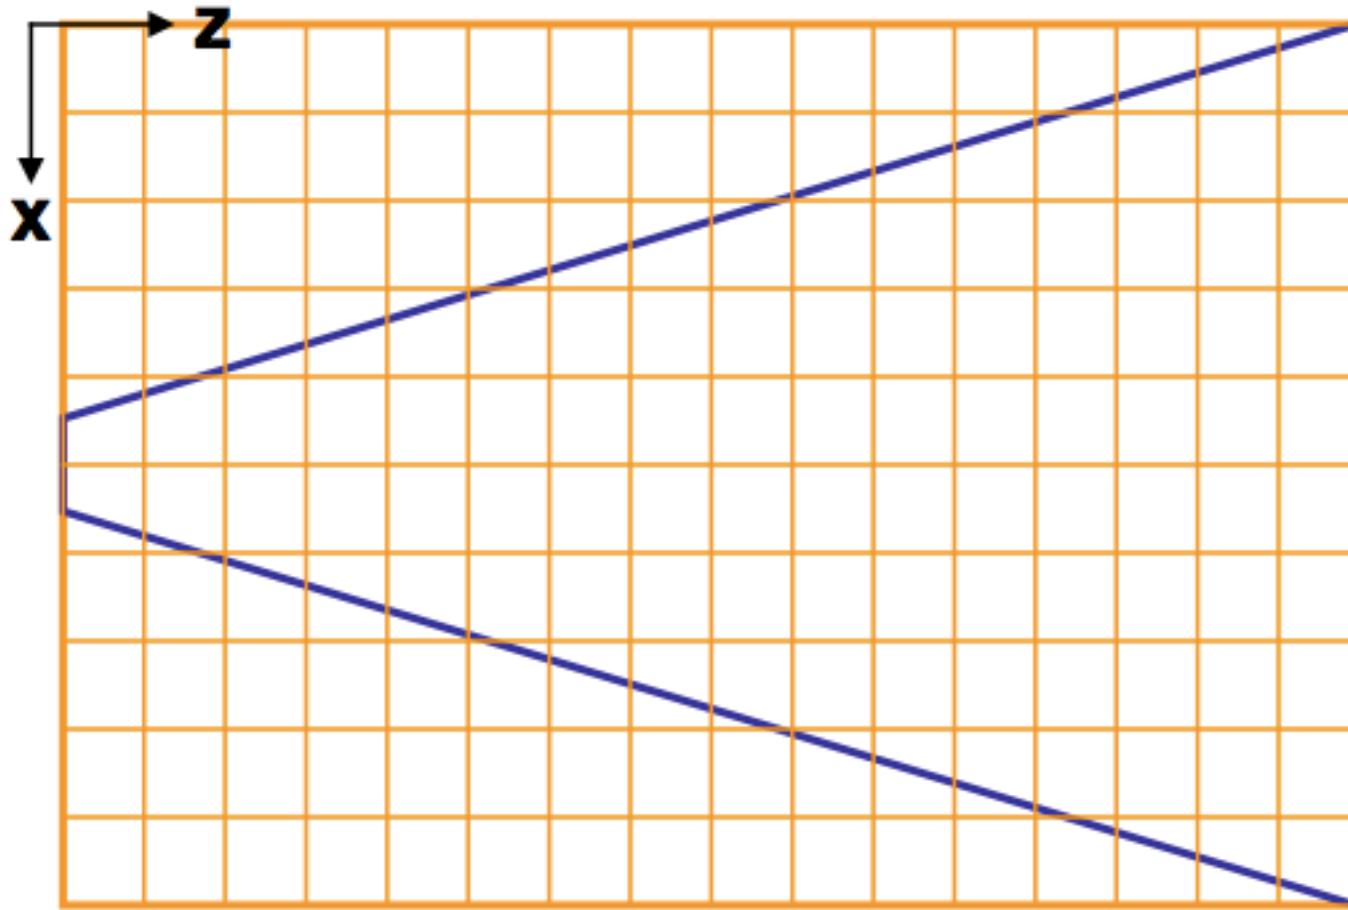
Uniform sampling in z



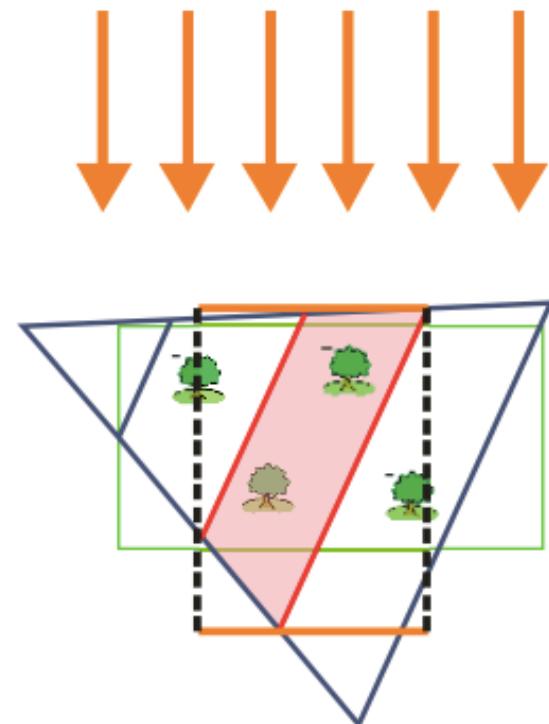
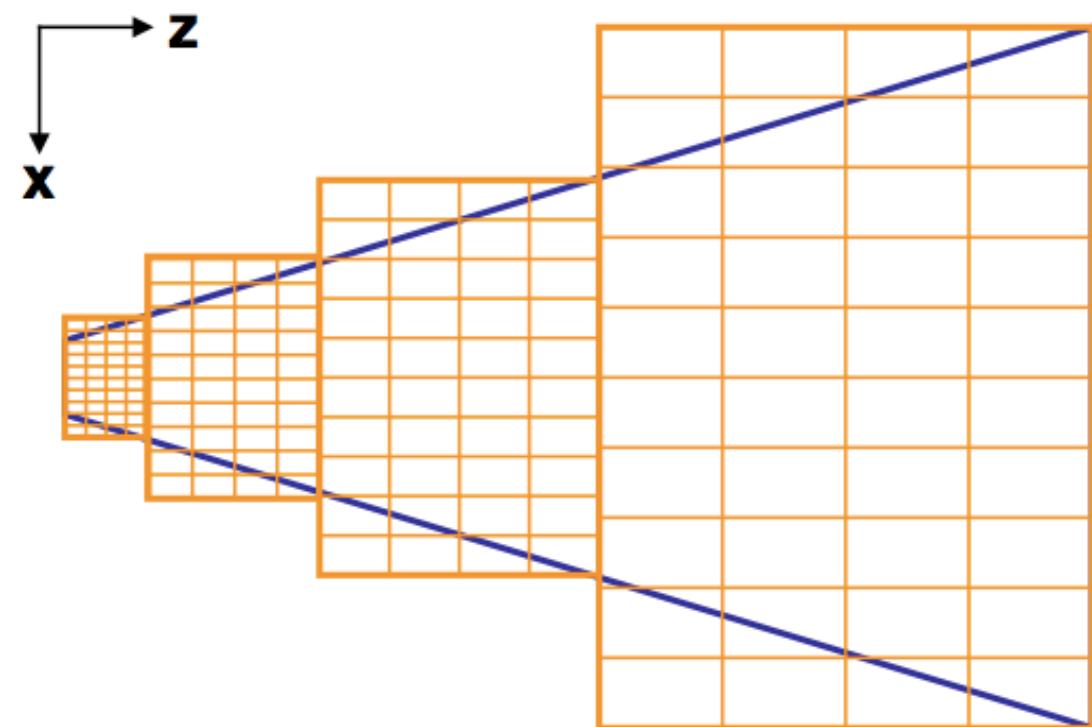
Variable sampling in z

- ▶ How?
- ▶ Linear projection
 - Not centered on the light source
 - Optimized based on view frustum + LS position
- ▶ TSM, LiSPSM...

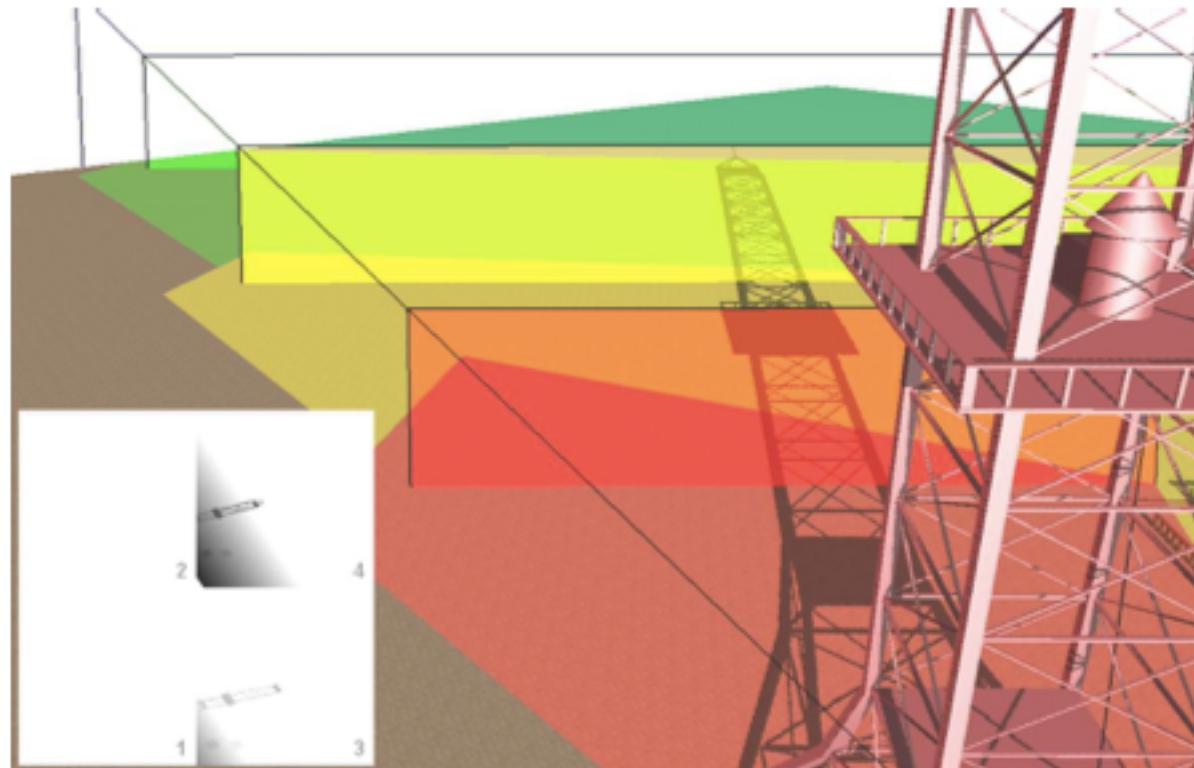
Cascading shadow maps



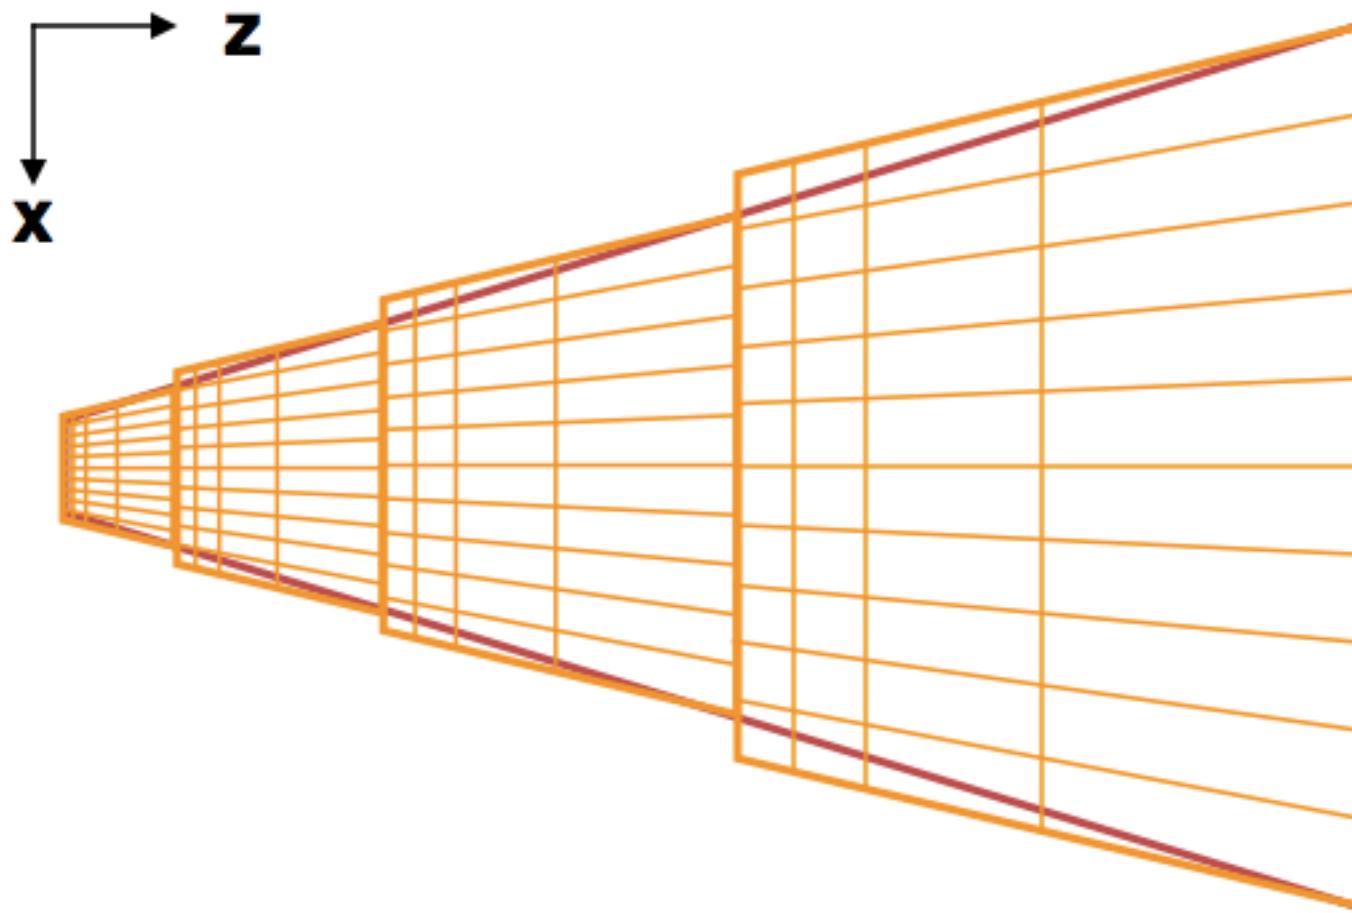
Cascading shadow maps

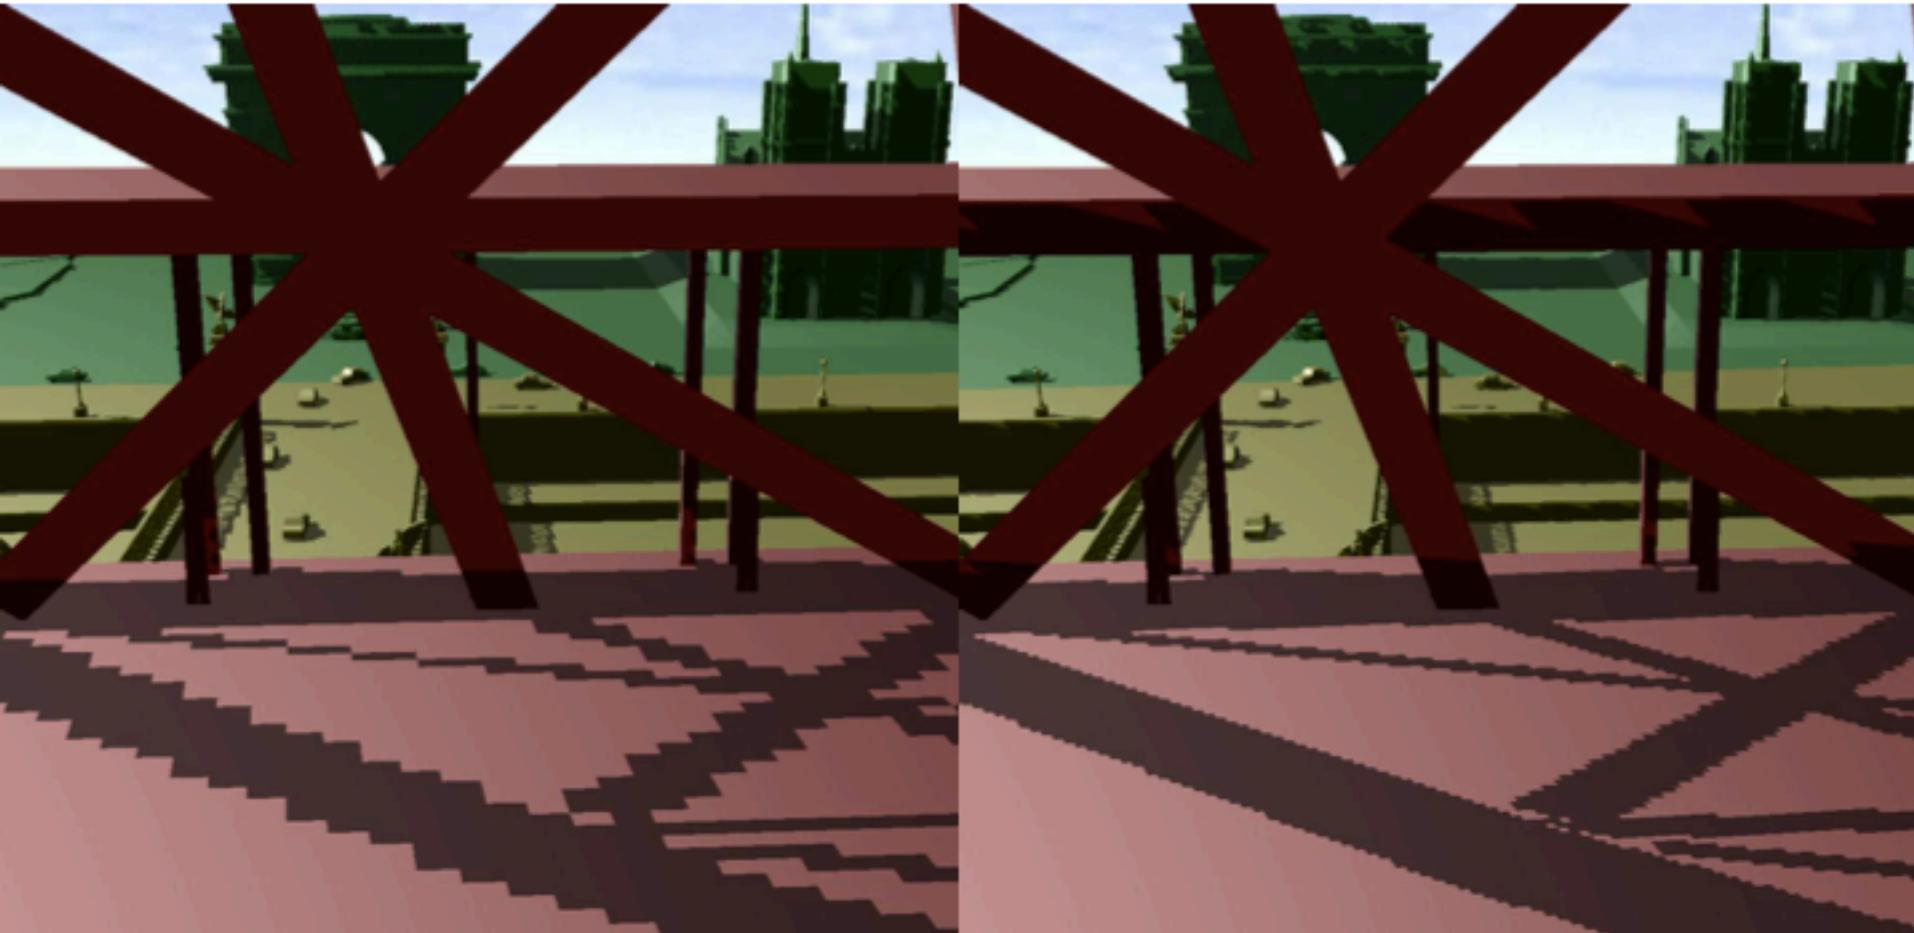


Cascading shadow maps



Cascading + warping





Partitioning

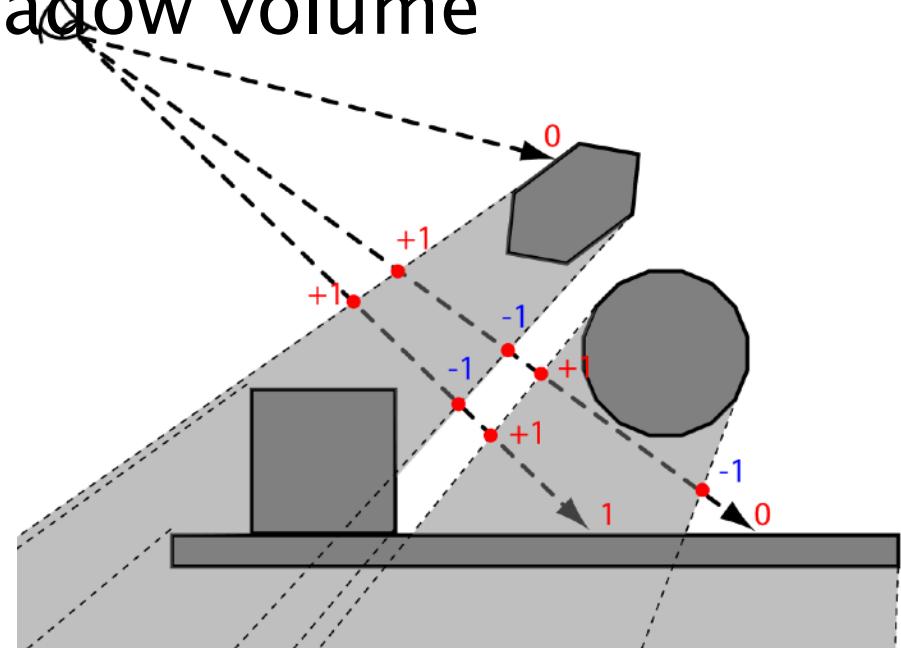
Partitioning + warping

Shadow Volumes algorithm

1. For each shadow casters, build a **shadow volume**
2. For each fragment, **count** how many times we enter/leave a shadow volume

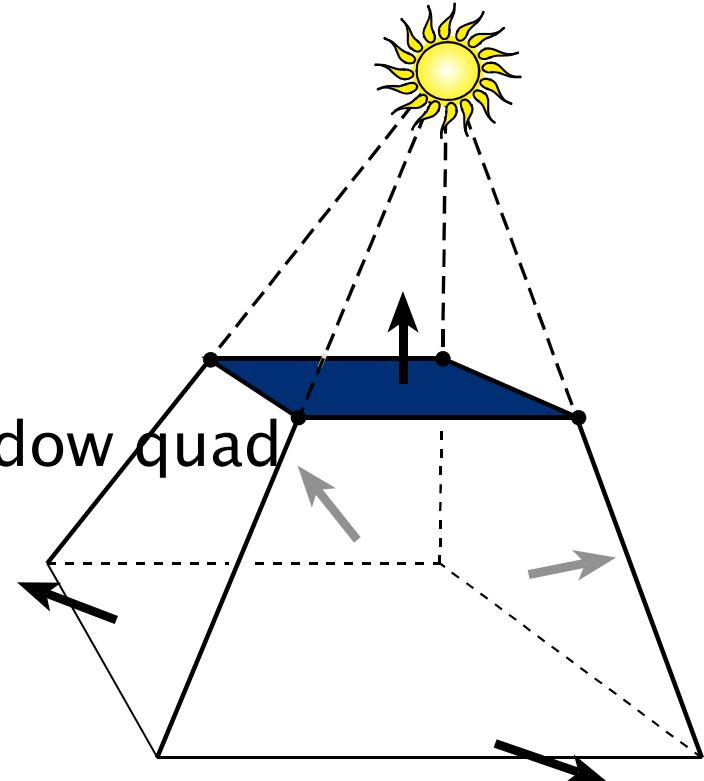
> 0 : in shadow

= 0 : lit



Shadow Volumes algorithm

- ▶ Building a shadow volume
 - Silhouette of each object from the light source
 - Infinite quads touching
 - the light source
 - Each silhouette edge
- ▶ Counting entering/leaving
 - Use the stencil buffer
 - Use the orientation of each shadow quad for the sign



Extract the silhouette?

- ▶ Silhouette of each object from the light source

How? 1mn



Building semi-infinite quads?

How? 1mn



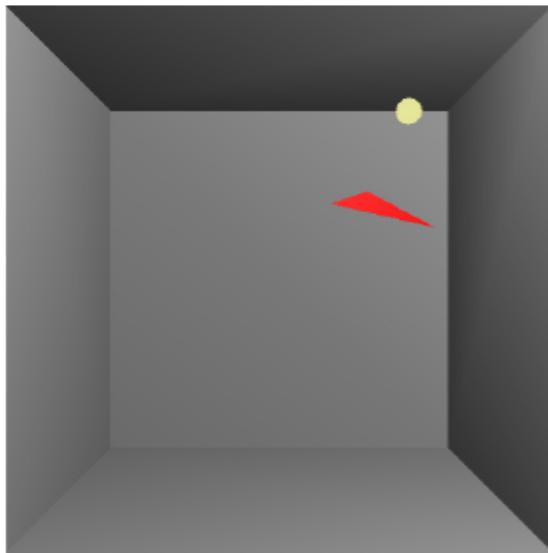
How do we count?

- ▶ Use the Stencil buffer
 - Shadow volume side visible, front-facing: +1
 - Shadow volume side visible, back-facing: -1
- ▶ 2 rendering pass:
 - First front-facing, then back-facing
 - glCullFace(...)
- ▶ 1 rendering pass:

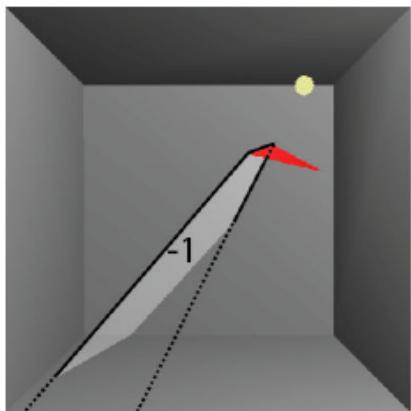
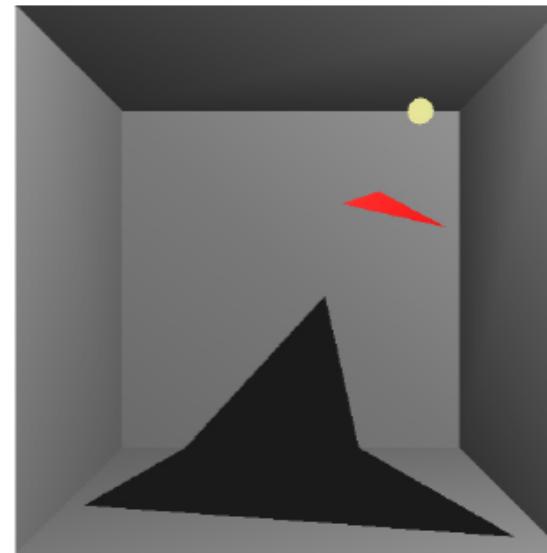
```
glStencilOpSeparate(GL_FRONT, GL_KEEP, GL_INCR_WRAP, GL_KEEP);  
glStencilOpSeparate(GL_BACK, GL_KEEP, GL_DECR_WRAP, GL_KEEP);
```

Z-pass by example: how the stencil buffer is used

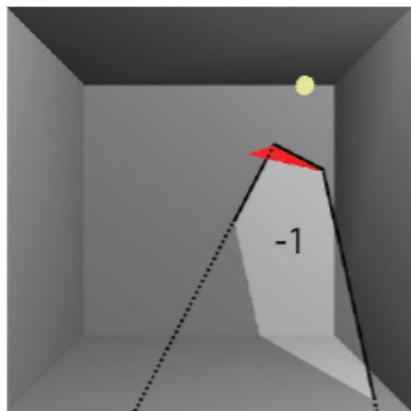
What we have...



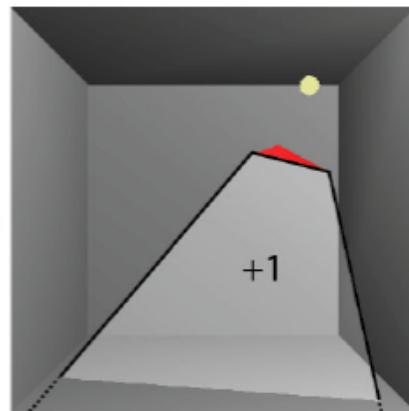
What we want...



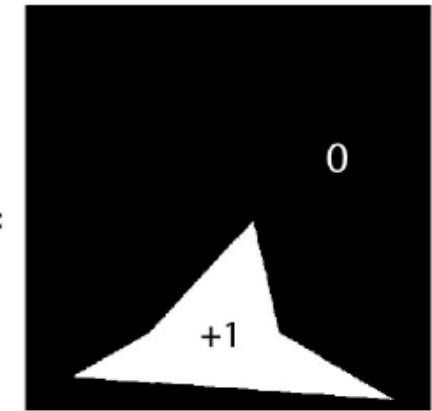
+



+

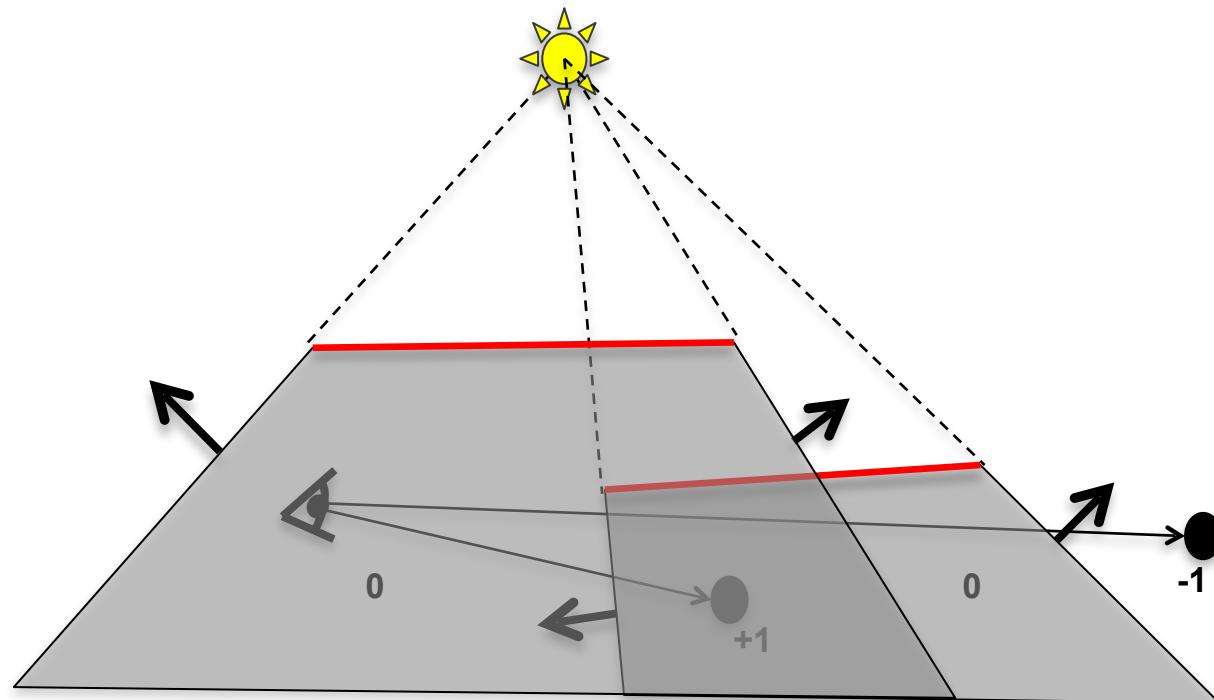


=



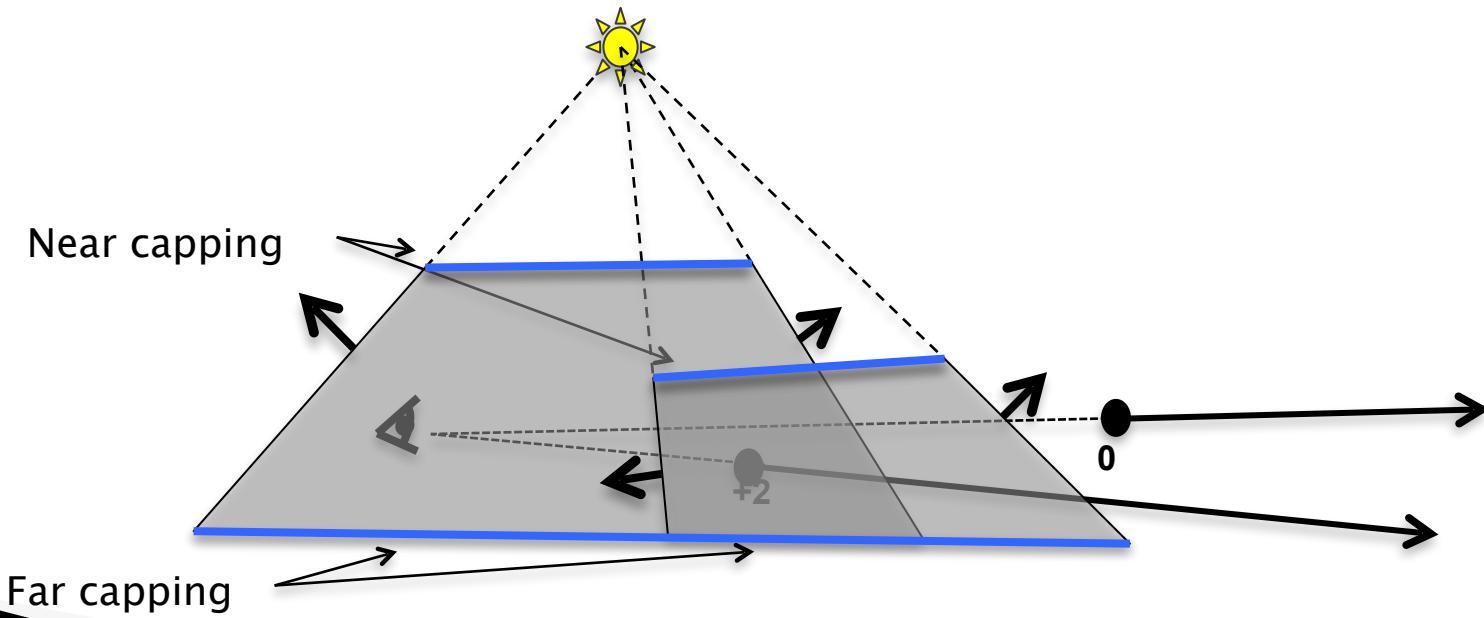
Z-pass: issue

- ▶ What if the eye is in shadow?



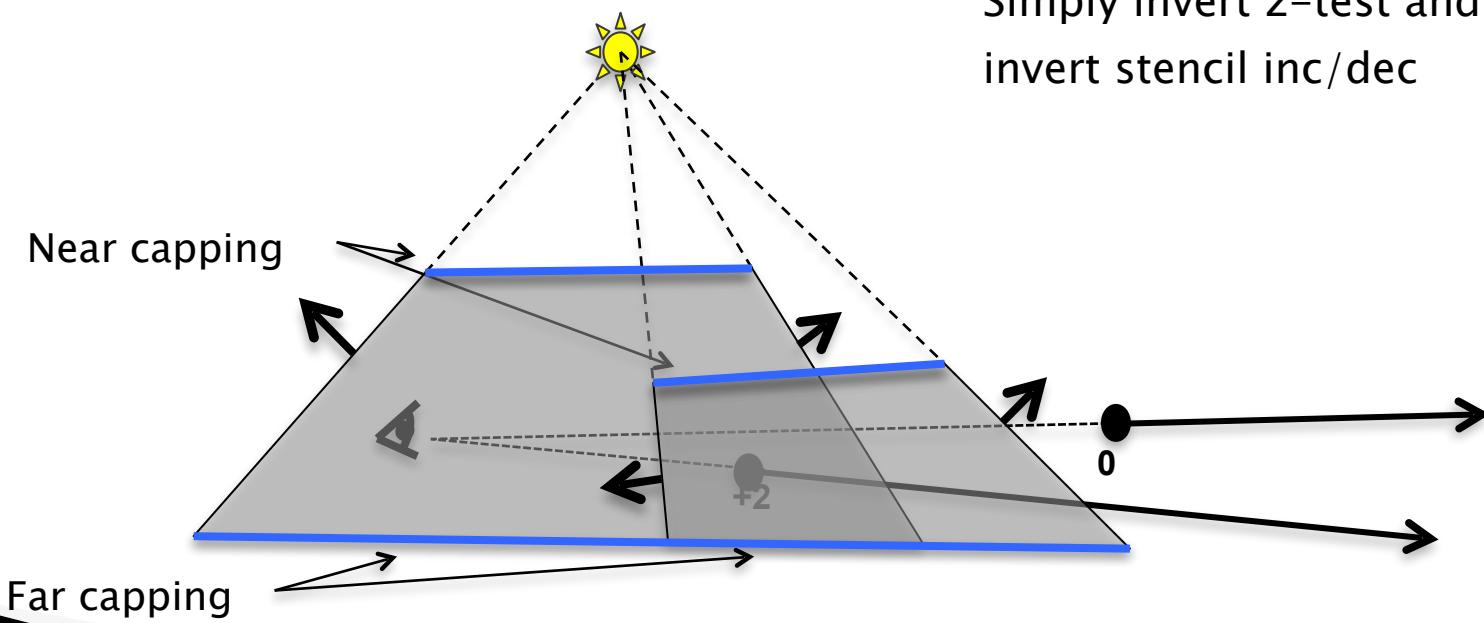
Z-fail

- ▶ Have a lit point as reference
- ▶ A point at infinity must be lit
- ▶ Need to cap the shadow volume

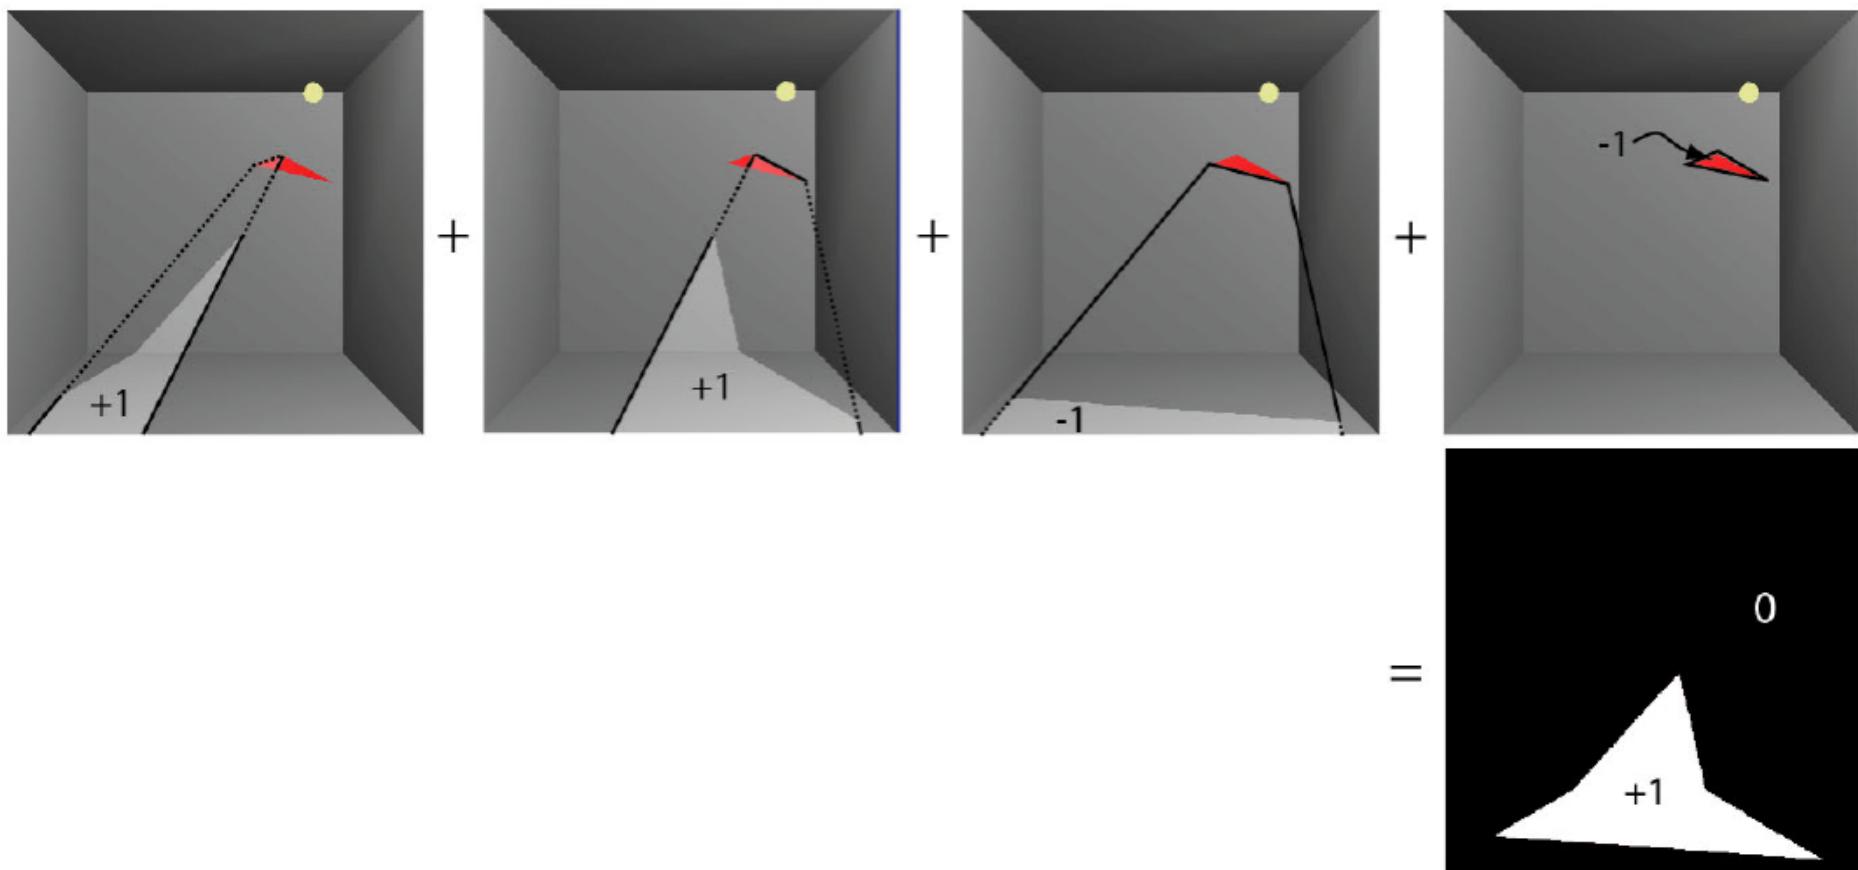


Z-fail

- ▶ Have a lit point as reference
- ▶ A point at infinity must be lit
- ▶ Need to cap the shadow volume



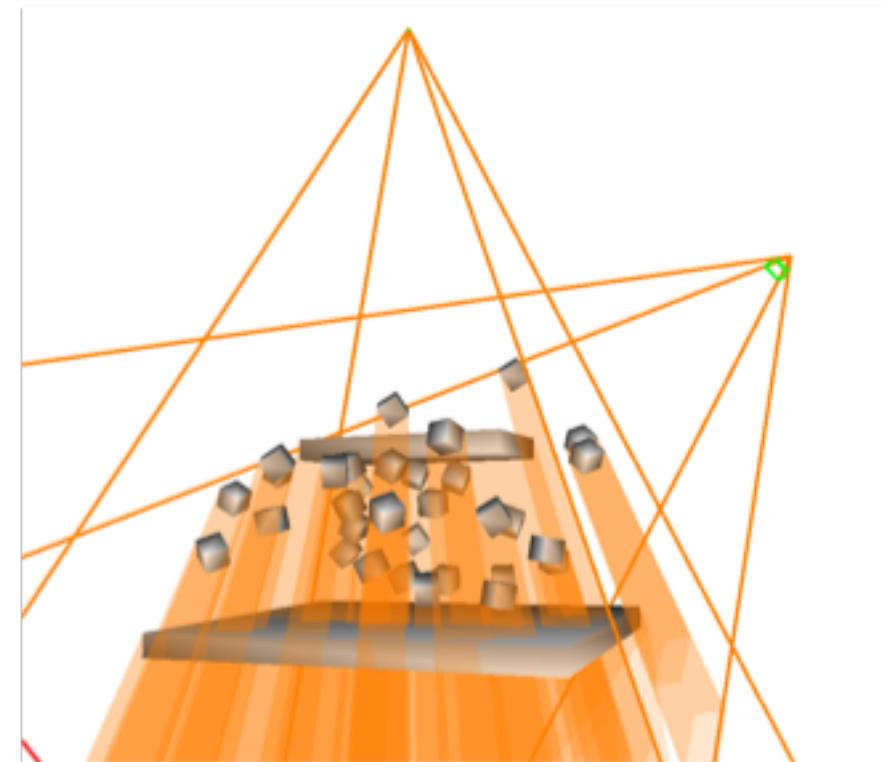
Z-fail by example



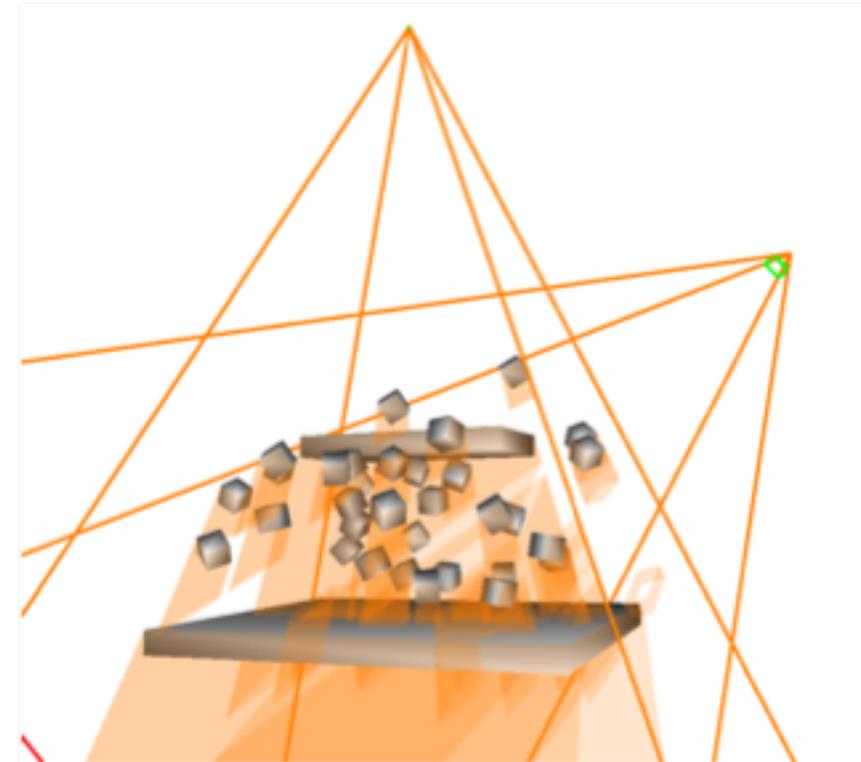
Shadow volumes: pro&cons

- ▶ Pros:
 - Sharp shadows
 - Arbitray positions for light source/caméra
 - Robust (if well programmed)
- ▶ Cons:
 - silhouette computation (CPU/GPU)
 - requirements on scene geometry (manifold, closed surfaces)
 - Rendering the scene twice, + the shadow volumes

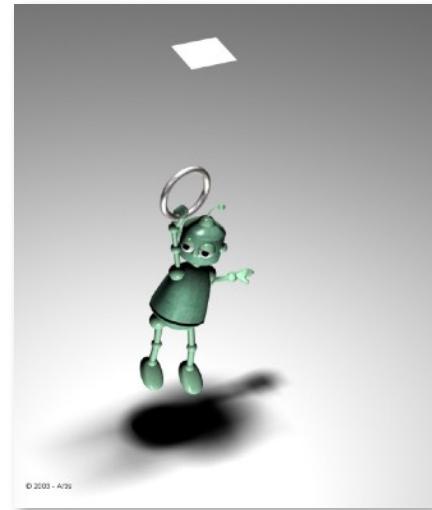
Overdraw



Shadow volumes



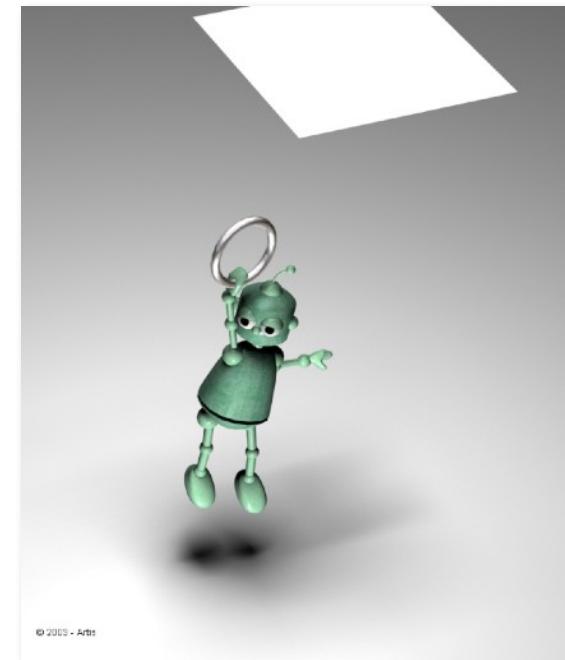
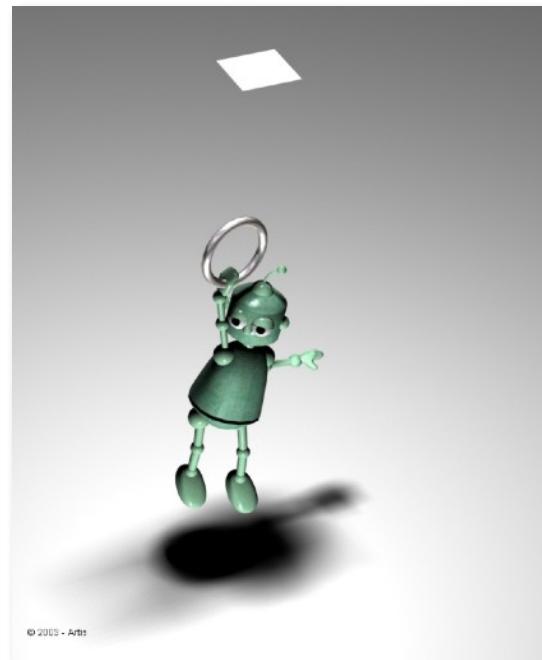
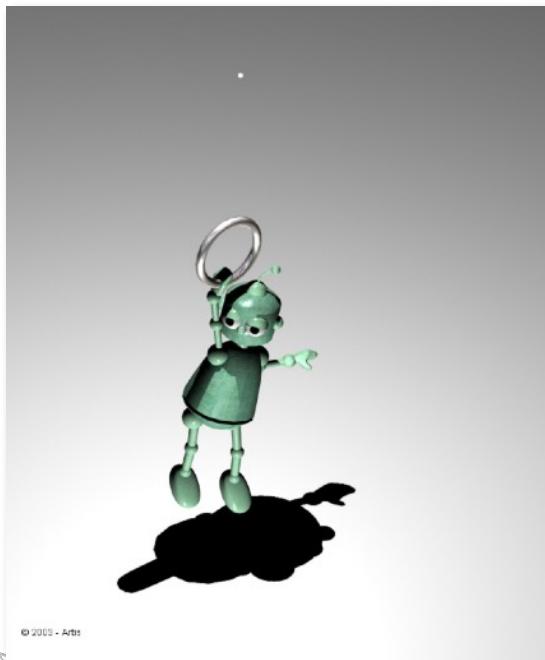
CC Shadow volumes



Soft shadow computations

Soft shadows

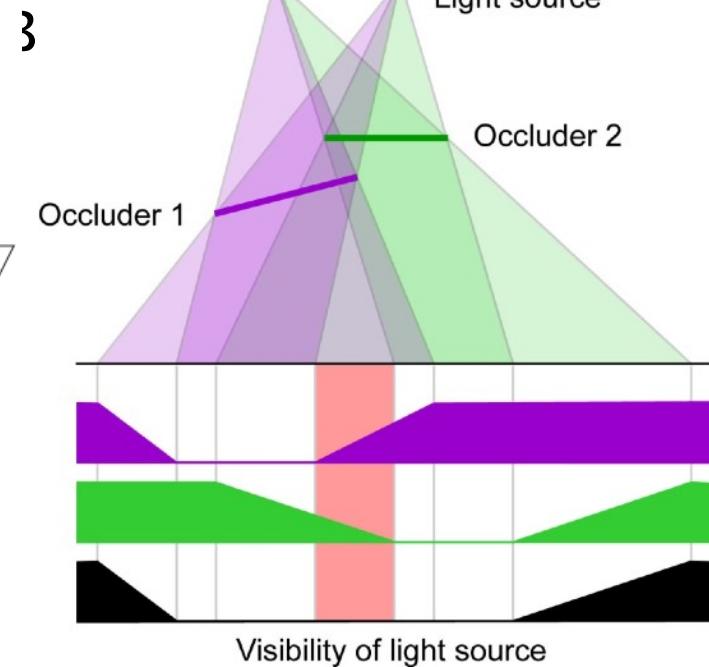
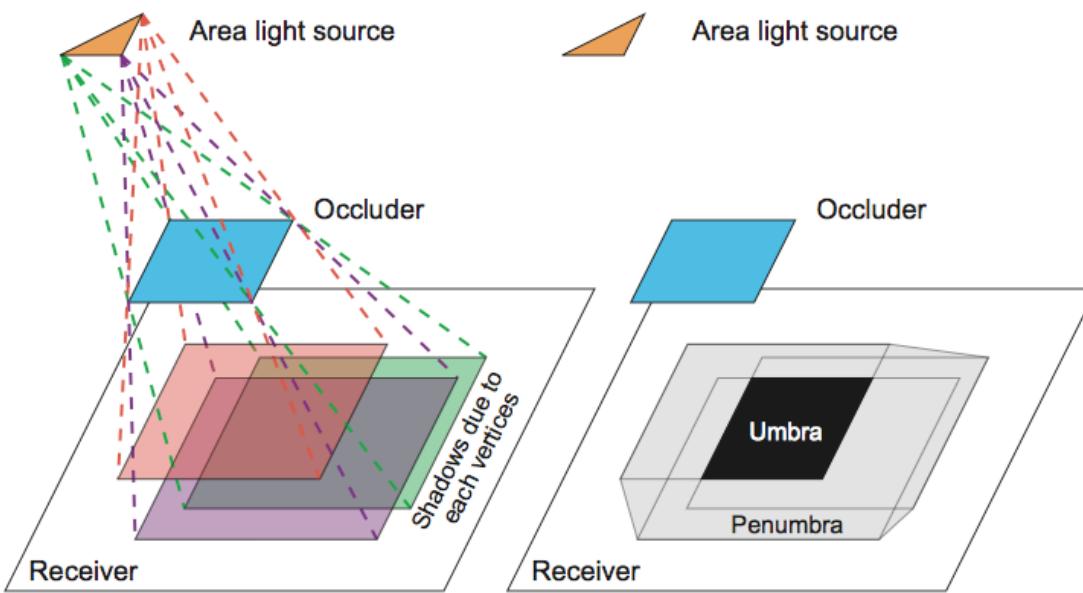
- ▶ More complex
 - Point-to-area visibility, with continuous value
 - Instead of binary point-point visibility
 - silhouette?



Soft shadows

- More complex

- Point-to-area visibility, with continuous value
 - Instead of binary point-point visibility
 - silhouette?
- Shadow of the sum \neq sum of shadows

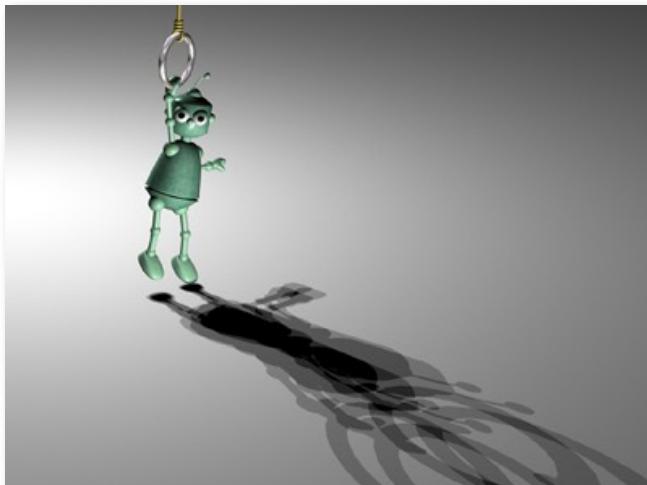


Soft shadows

- ▶ More complex
 - Point-to-area visibility, with continuous value
 - Instead of binary point-point visibility
 - silhouette?
 - Shadow of the sum \neq sum of shadows
 - A hides 50% and B hides 50%, A+B doesn't hide 100%
- ▶ Many algorithms
 - With varying accuracy
 - Approximating the shadow casters
 - Precomputations (Precomputed Radiance Transfert)
 - With varying speed

Soft shadows through sampling

- ▶ Accumulating shadows:
 - Compute several hard shadows
 - Add them, average the results
 - accumulation buffer
- Needs many samples
 - Computation time proportional to # échantillons



4 samples



1024 samples

Soft-shadow volume

- ▶ For each silhouette edge:
 - Compute volume around penumbra (wedge)
 - For each pixel in this wedge
 - Compute attenuation coefficient
- ▶ Beautiful, realistic, expensive

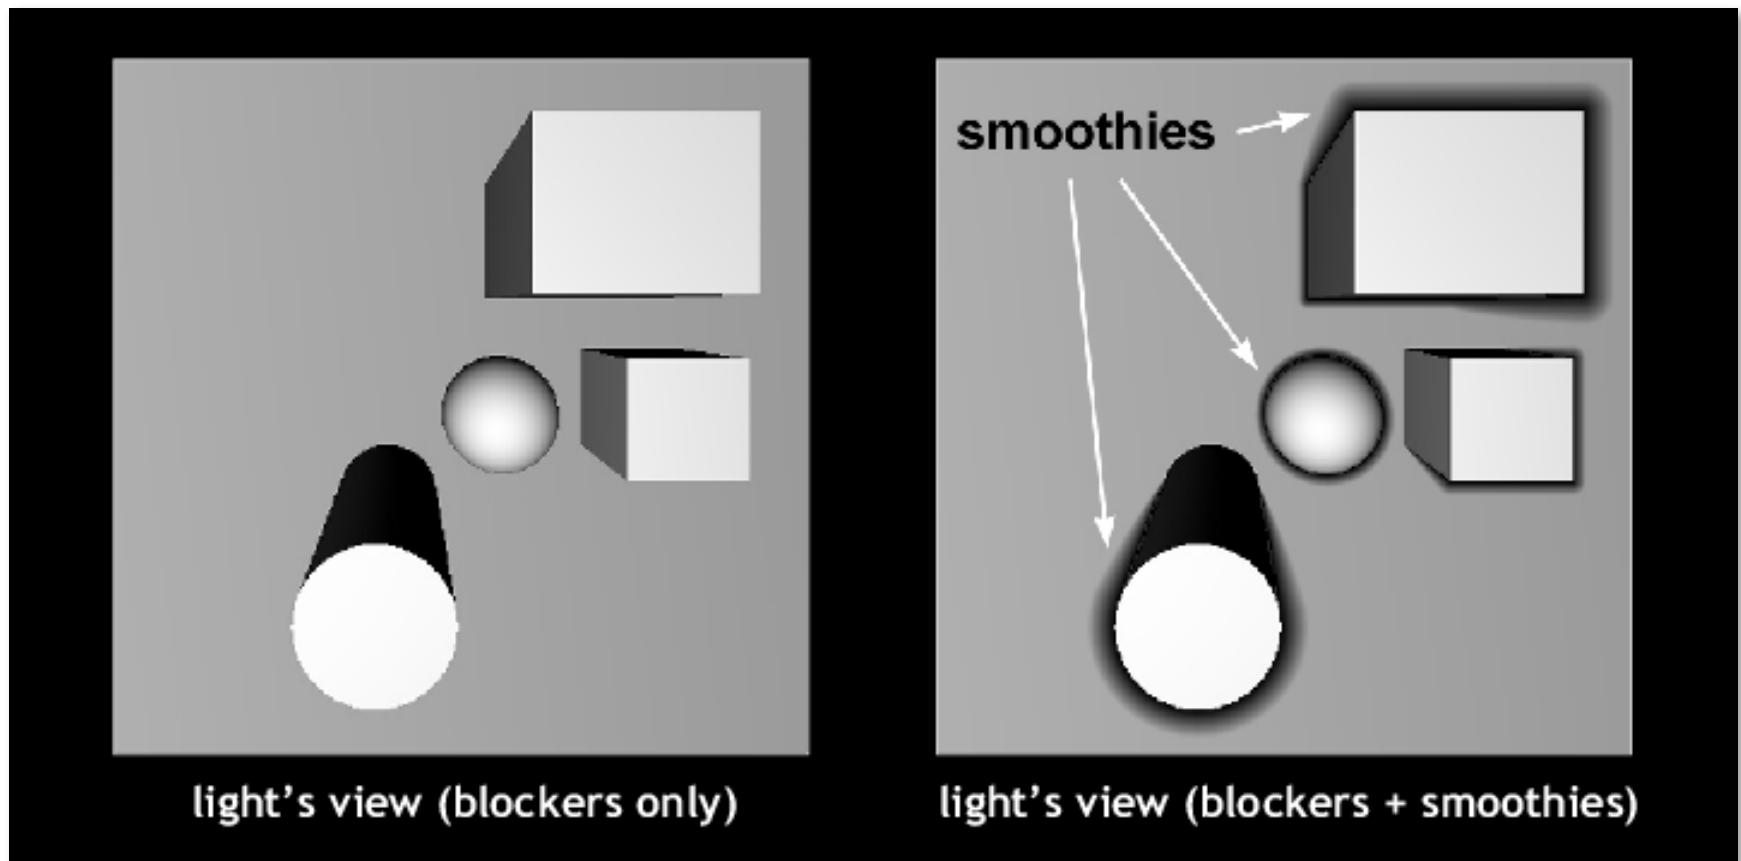


Penumbra wedges [Sig03] U. Assarson, T. Möller

Object/image methods

Rendering Fake Soft Shadows with Smoothies [SoR03]

E. Chan, F. Durand



Shadow mapping extension

- ▶ Percentage Close Filtering (PCF)
 - Filter shadow map around sampling point
 - Possible GPU speed-ups (2x2 kernel)
 - Pre-filtered, stored in mip-map



1 sample



9x9 kernel



17x17 kernel

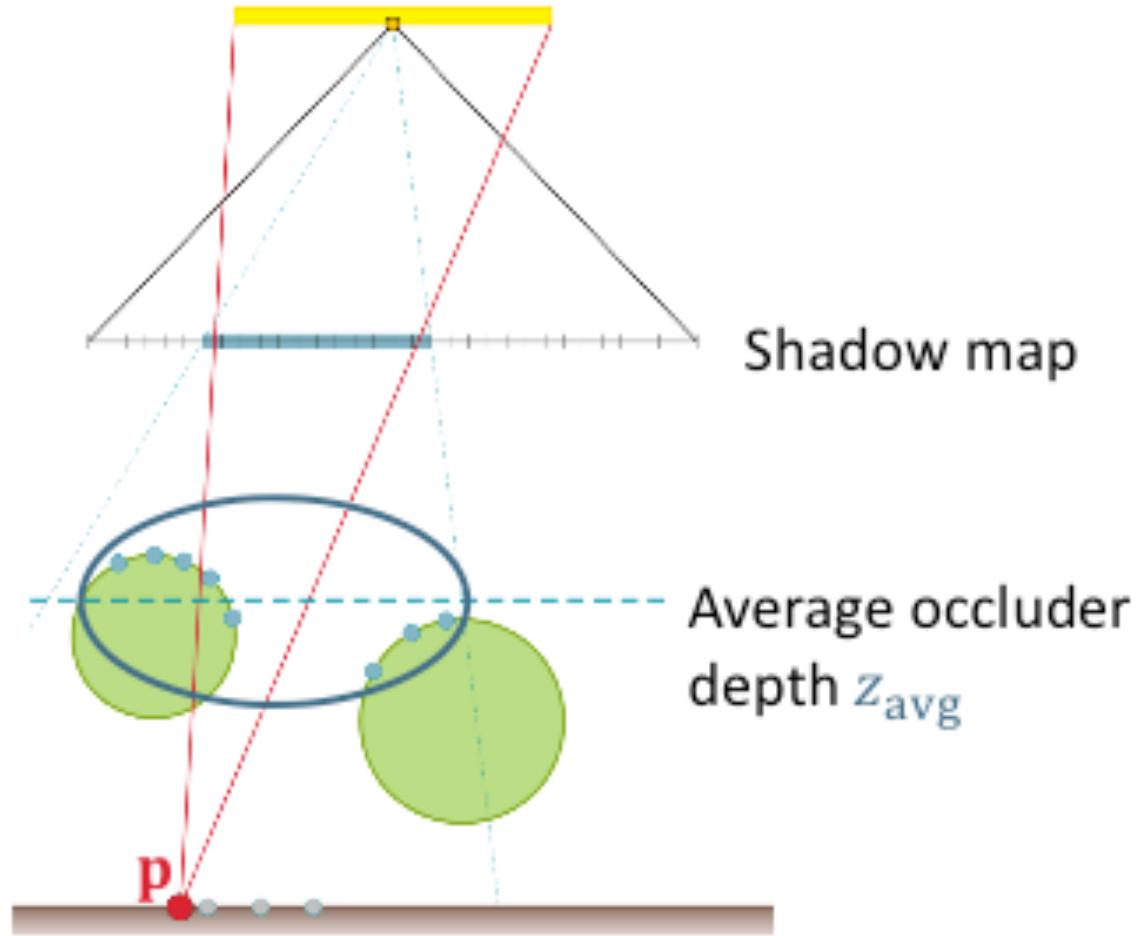
Shadow mapping extension

- ▶ Percentage Closer Soft Shadows (PCSS)
[Fernando 05]
 - Compute kernel size first, by sampling shadow map
 - Filter using PCF
(or extensions)



PCSS

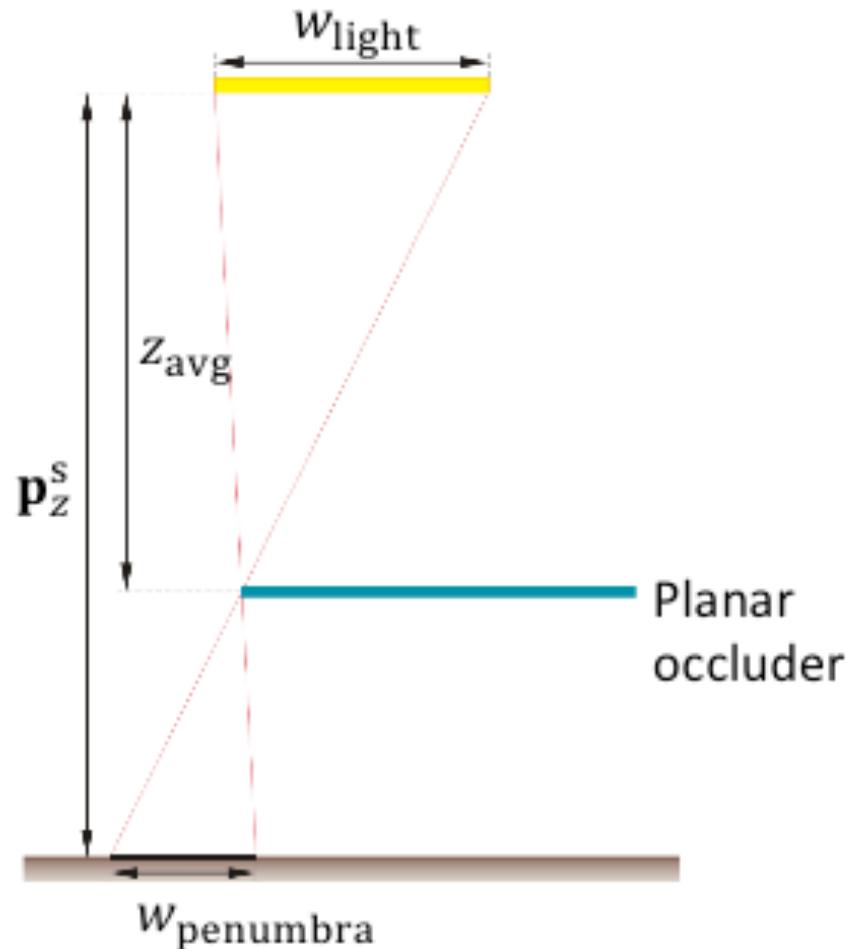
► 1. Blocker search



PCSS

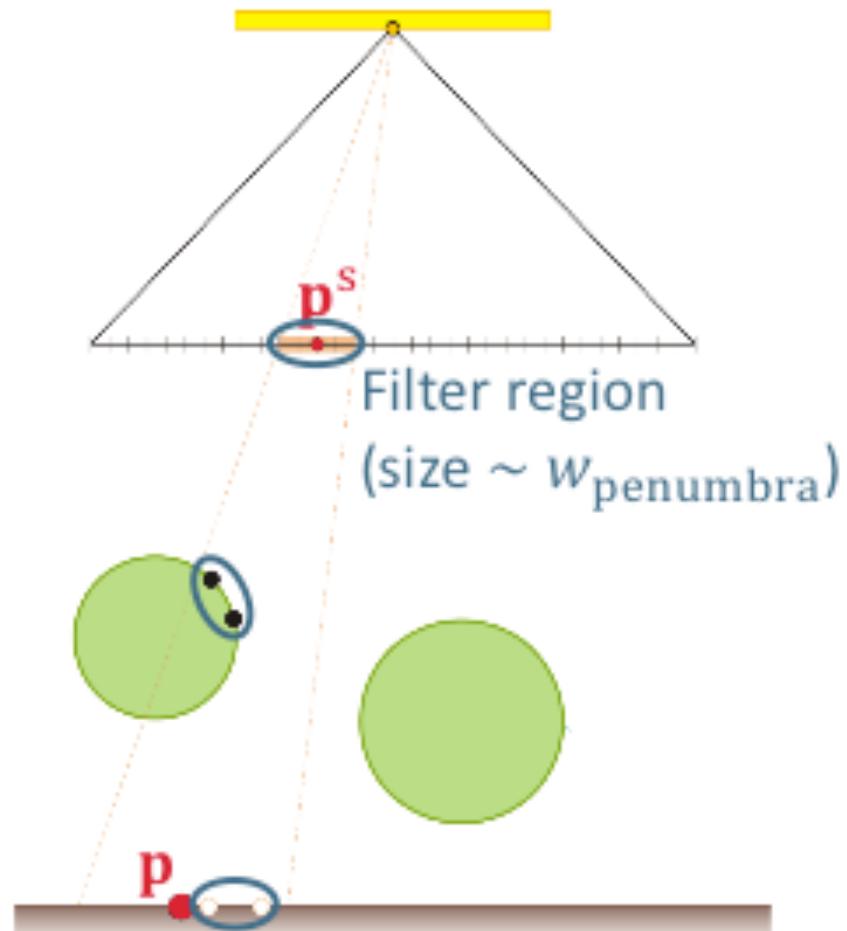
► 2. penumbra size

$$w_{\text{penumbra}} = \frac{\mathbf{p}_z^s - z_{\text{avg}}}{z_{\text{avg}}} w_{\text{light}}$$



PCSS

► 3. filtering



(here, occlusion = 50 %)

PCSS : issues (1)

- ▶ 1. blocker search
- ▶ 2. penumbra size
- ▶ 3. filtering

PCSS : issues (1)

- ▶ 1. blocker search
- ▶ 2. penumbra size
- ▶ 3. filtering

2 steps requiring several access to shadow map

PCSS : issues (2)

- ▶ Easy, quite fast
- ▶ Visually pleasing results
 - For a small light source
- ▶ No physical realism
- ▶ Visible artefacts
 - For large penumbra width
 - If occluders hidden from center of light source
 - For non-flat occluders

Shadow mapping extensions

- ▶ Percentage Closer Soft Shadows (PCSS)
[Fernando 05]



PCSS

Hellgate: London (2007)

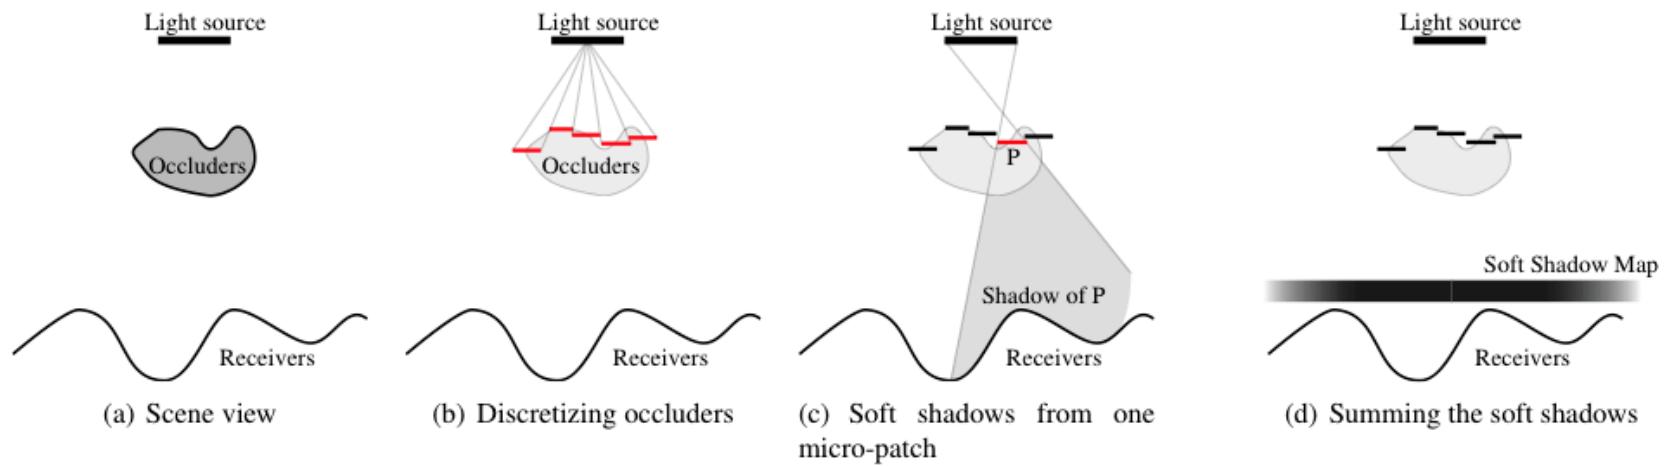


PCF

Soft shadow maps

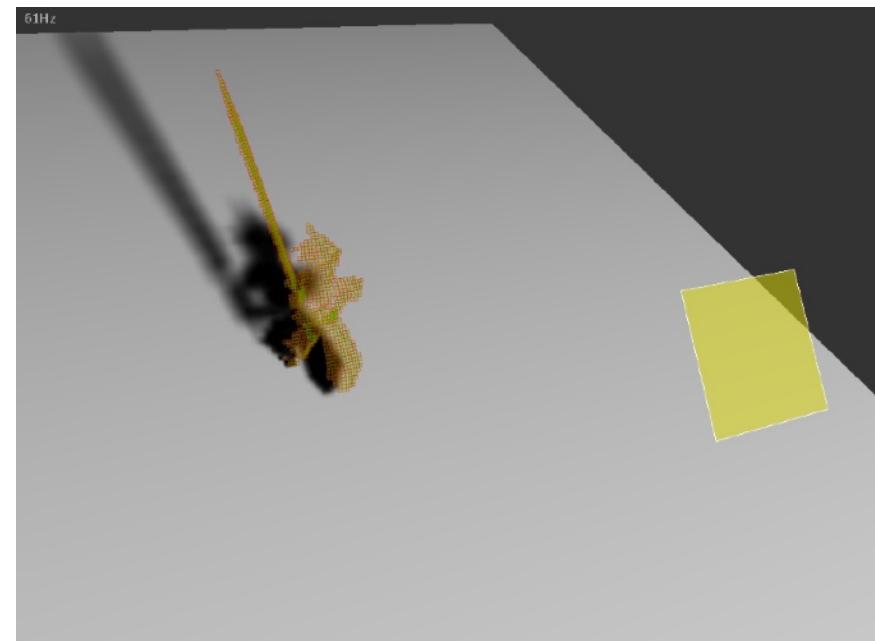
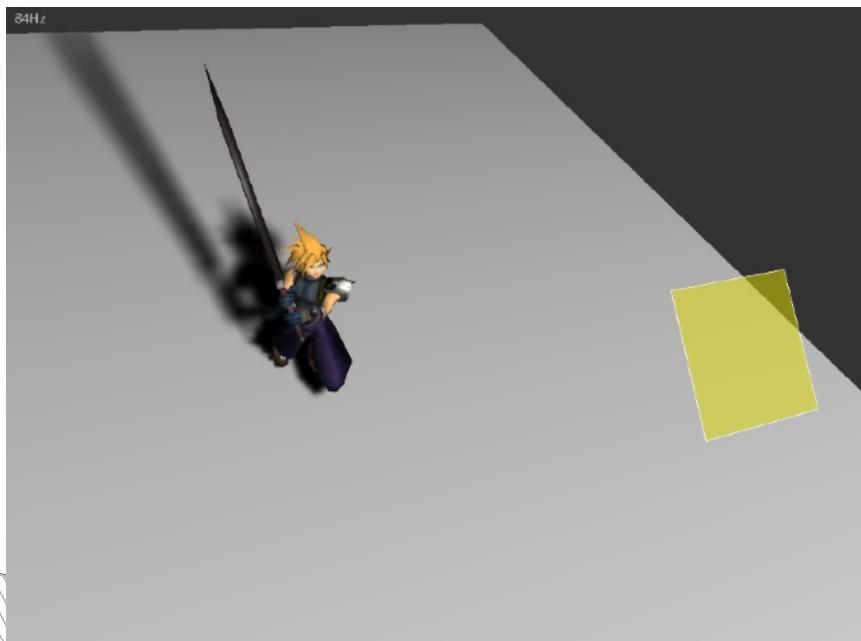
Soft Shadow Maps [AHLHHS06] Atty et al.

Real-time Soft Shadow Mapping by Backprojection [GBP06] Guennebaud et al..



Back-projection

- ▶ Shadow map = object discretization
- ▶ Compute shadow of discretized object
- ▶ Realistic, real-time, animated scene
- ▶ [Atty06] et [Guennebaud06]



Back-projection





Ambient occlusion

Motivation

- ▶ Approximating the occlusion under distant lighting
 - Ambient term taking visibility into account
- ▶ Perceptually related to depth, curvature and spatial proximity

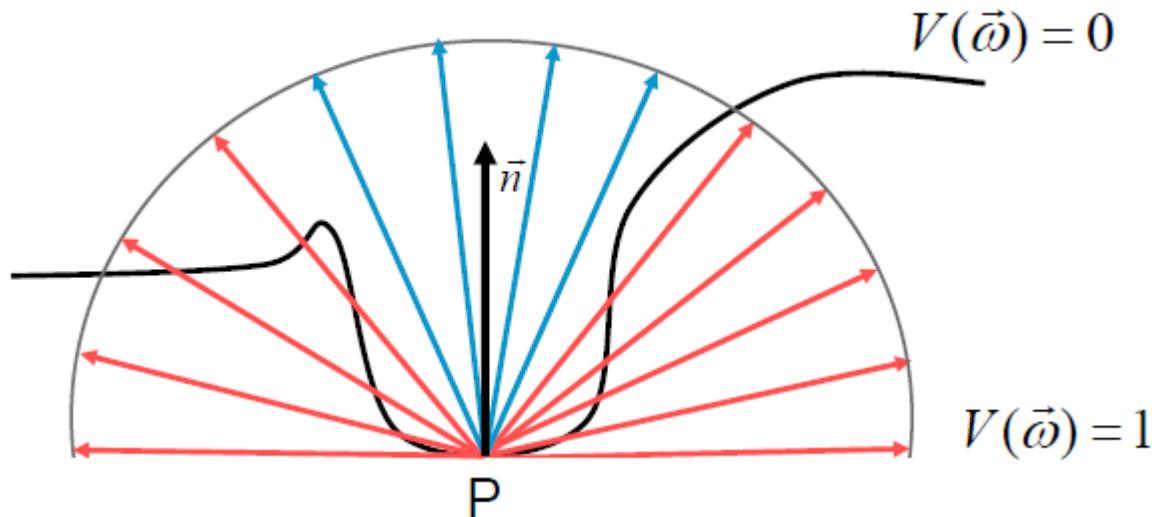


Definition

- ▶ Integral of visibility over hemisphere Ω :

$$A_P(\vec{n}) = \frac{1}{\pi} \int_{\Omega} V_P(\vec{\omega})(\vec{n} \cdot \vec{\omega}) d\omega$$

- Cosine term \Rightarrow diffuse lighting
- Usually, attenuation depending on the distance to P



Computing the integral

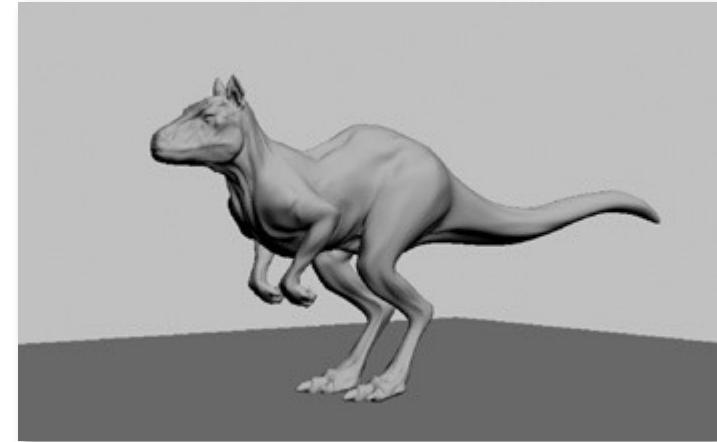
GPU Gems, chap 17

▶ Sampling

- Precomputation (ray-casting)
- Store in a texture

+ Rendering at no extra cost

- Slow precomputation
- Static scene



Diffuse



Diffuse + AO

Computing the integral

- ▶ Screen-Space Ambiant Occlusion (SSAO)
 - Use the depth buffer as an approximation of the scene
 - For each pixel, sample the hemisphere on the GPU
 - Filtering for noise reduction
- ✚ Independent form scene complexity
- ✚ No pre-computation
- ✚ Dynamic scene
- − Longer rendering

Cry Engine 2

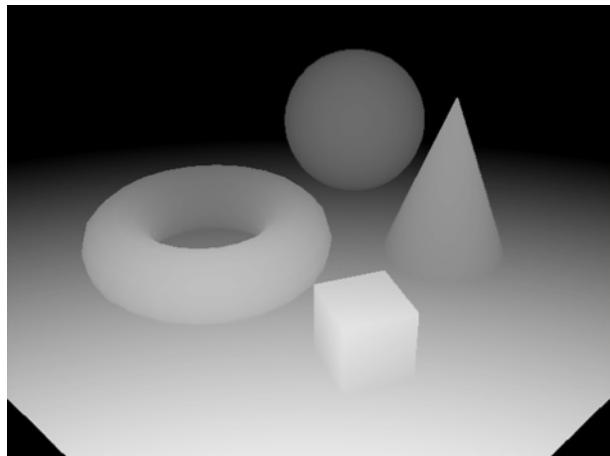


Deferred shading

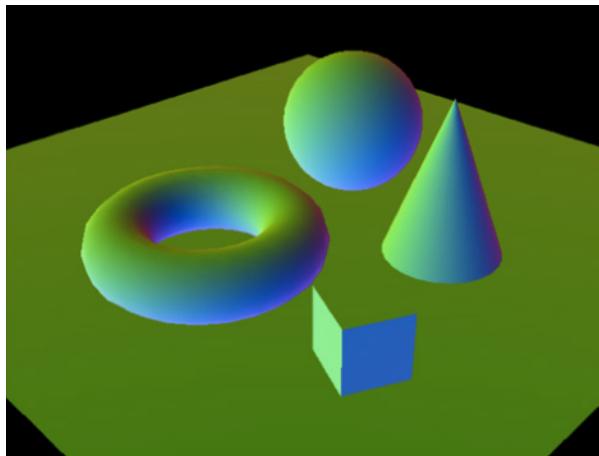
- ▶ Fragment shaders get expensive:
 - Complex materials, textures, indirect lighting...
- ▶ Pb. for complex scenes/multi-layers:
 - Shading for all surfaces
 - Even if they're invisible
 - Z-buffer test after the fragment shader
- ▶ Need: visibility before shading materials
 - Theoretically impossible
 - Solution: deferred shading

Deferred shading

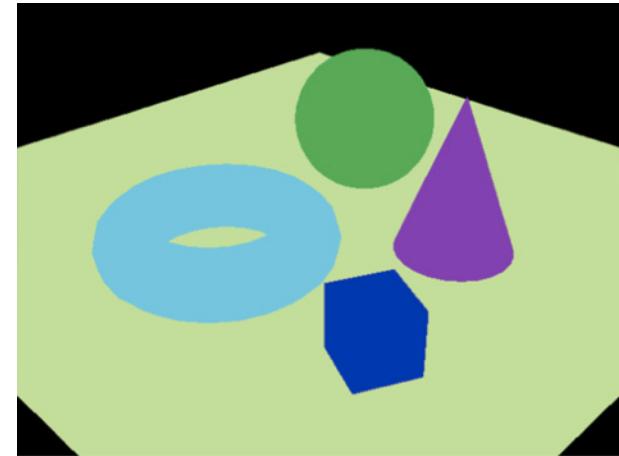
- ▶ 1st pass: rendering into 3–4 aux. buffers



Position (x,y,z)

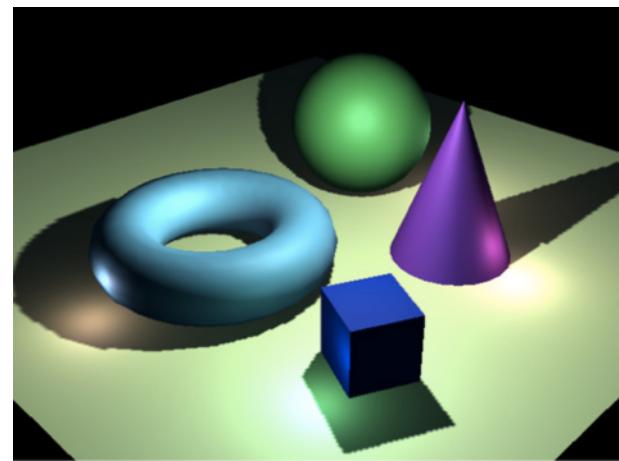


Normals



Colors, materials, textures

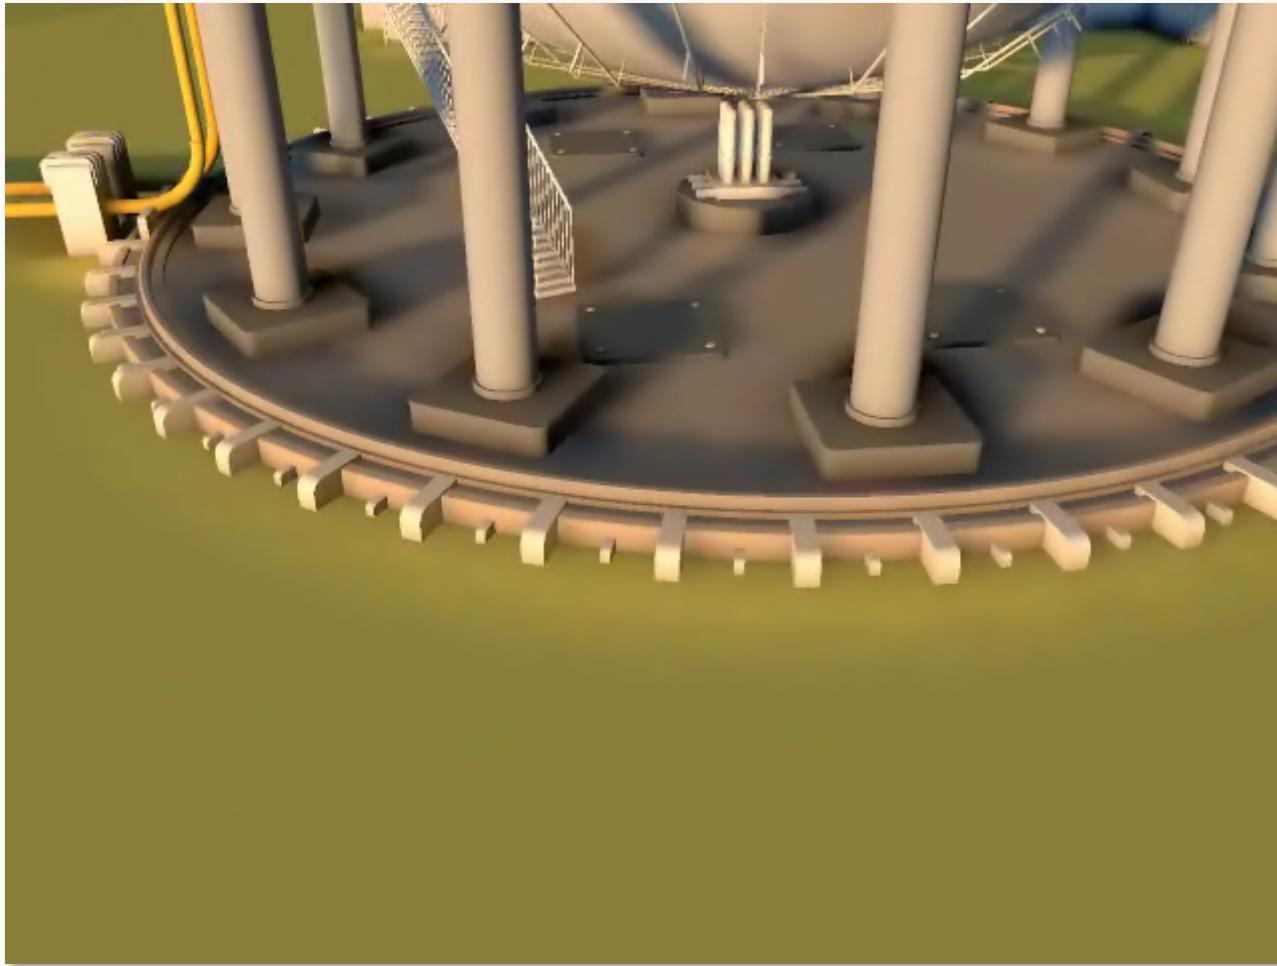
- ▶ 2nd pass: compute shading using these buffers



Deferred shading + SSAO

- ▶ SSAO :
 - Needs a geometry buffer
 - Is expensive: must reduce number of calls
- ▶ Deferred shading :
 - Has a geometry buffer
 - Did visibility as pre-computation
- ▶ Good match of algorithms

Beyond SSAO



Approximating Dynamic Global Illumination in Image Space
Ritschel et al. 2009